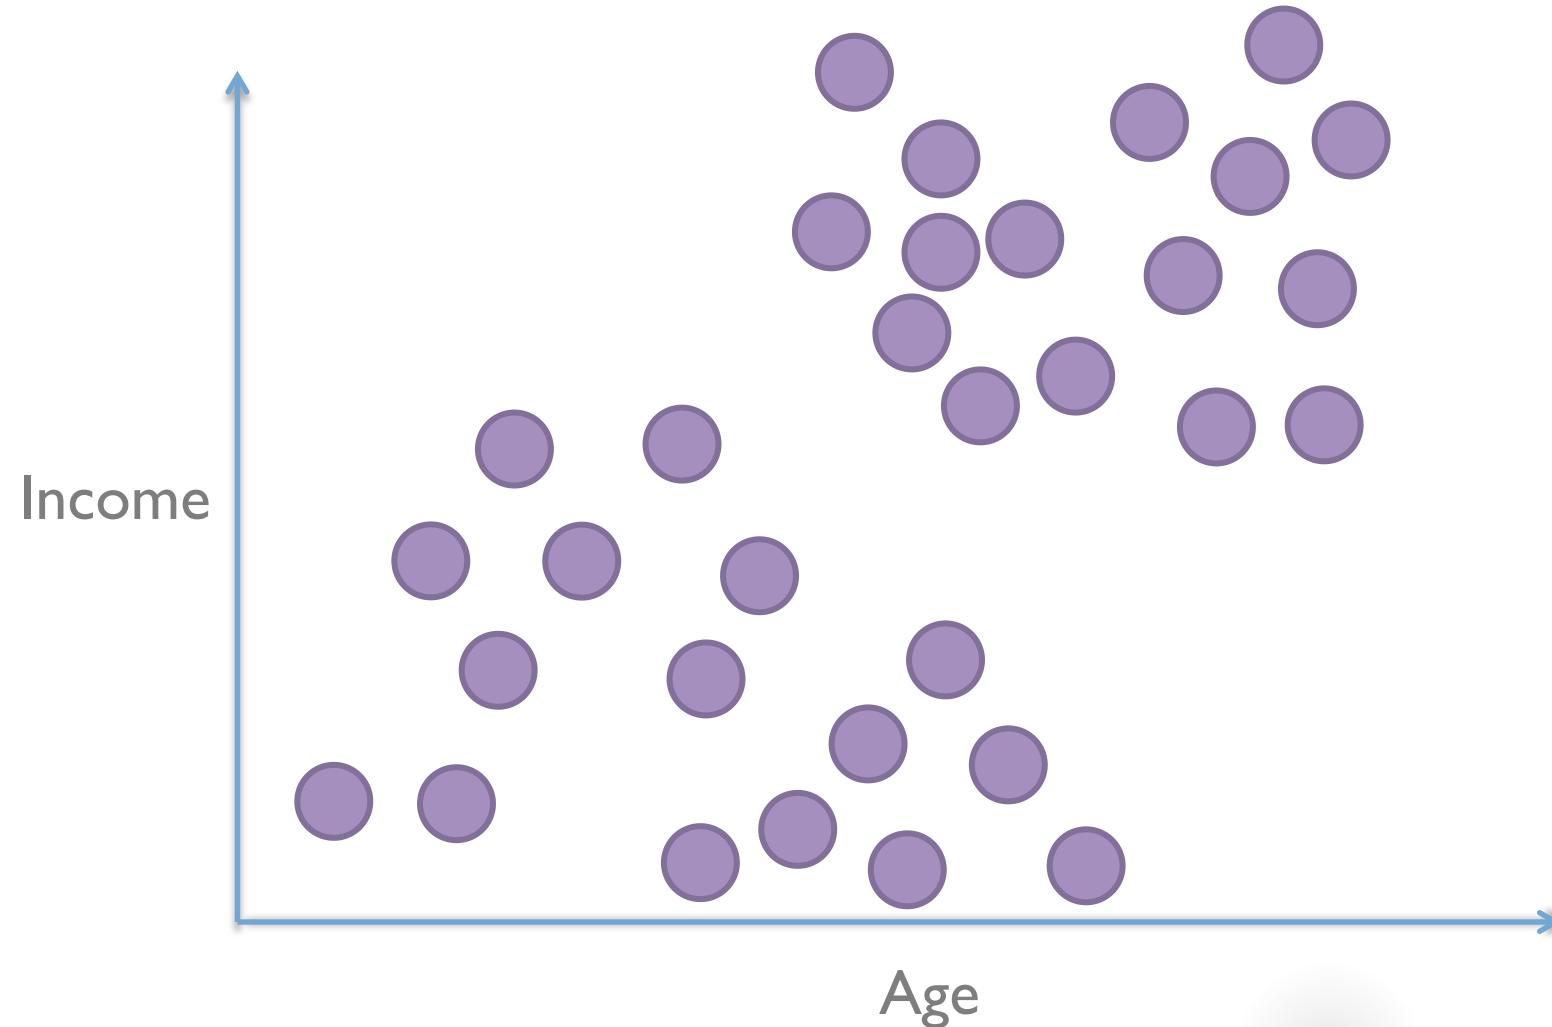


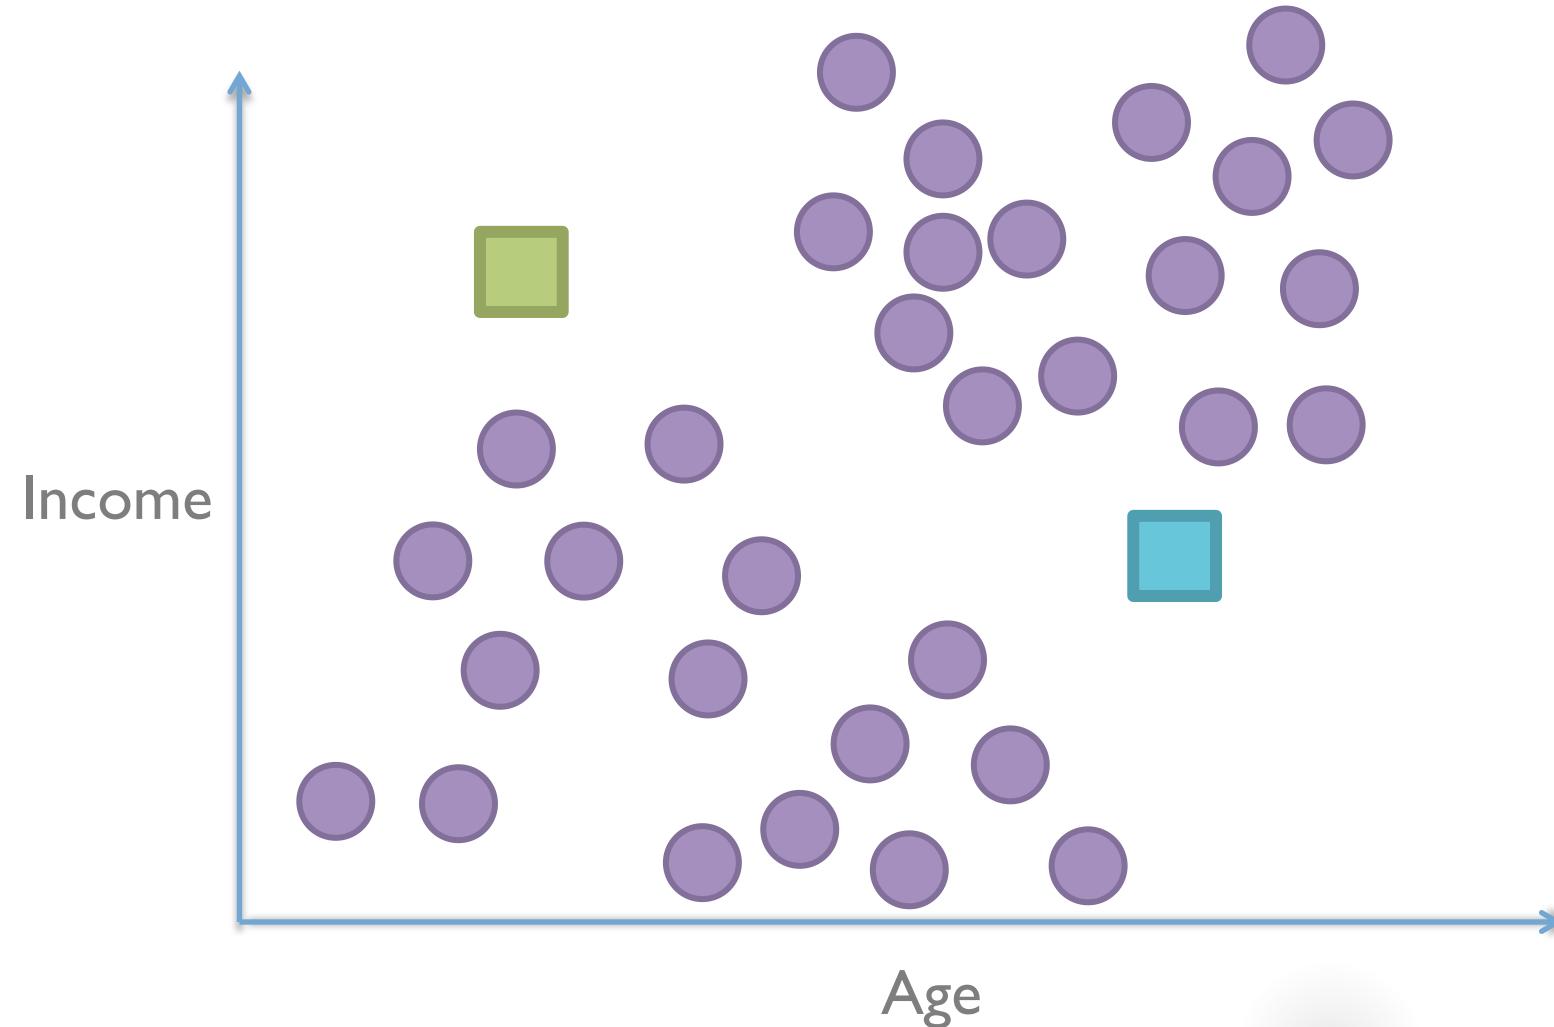
# Other Clustering Algorithms

# Remember K--Means?



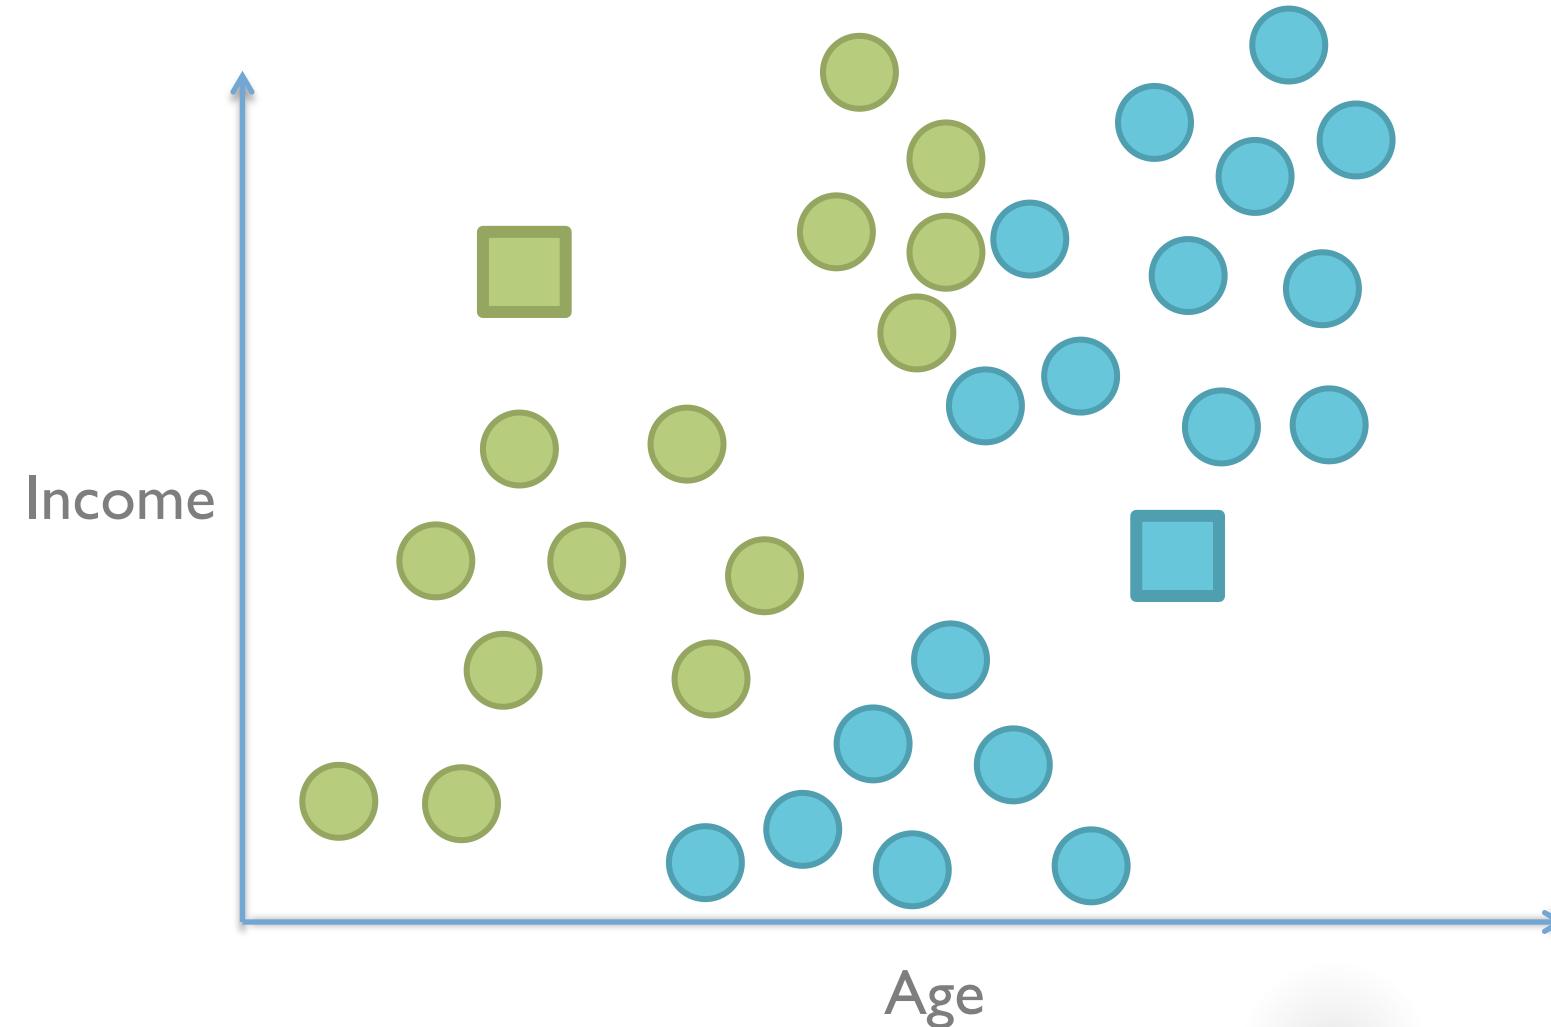
## K-Means K=2

Randomly assign two cluster centers



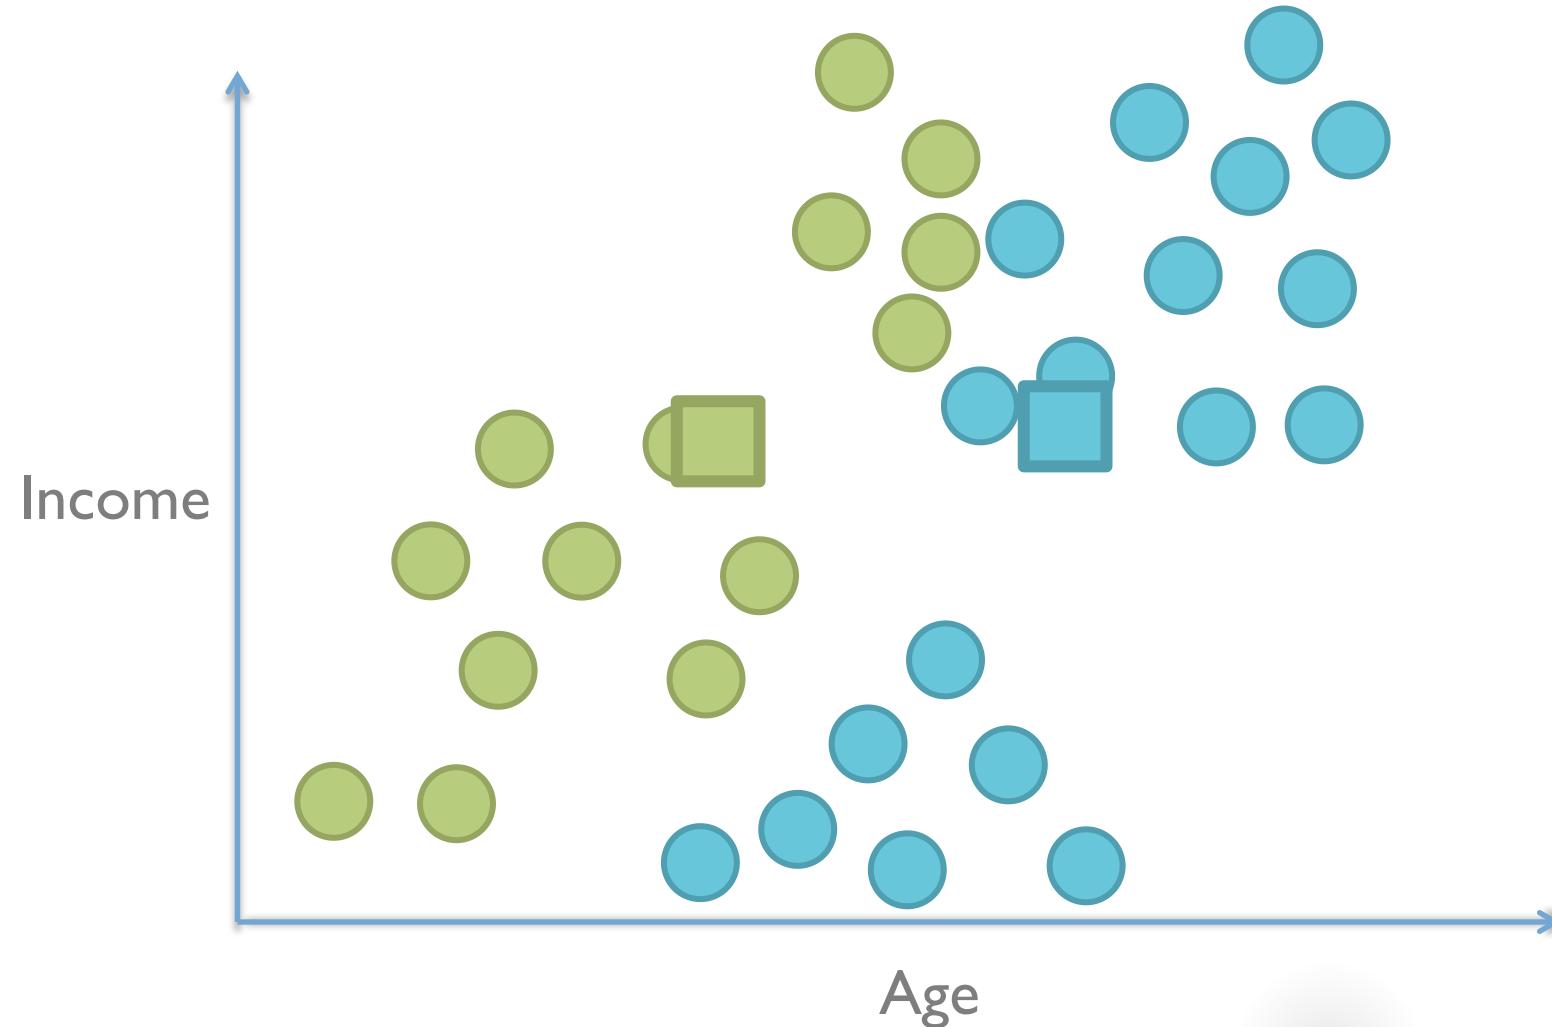
## K-Means K=2

Each point belongs to closest center



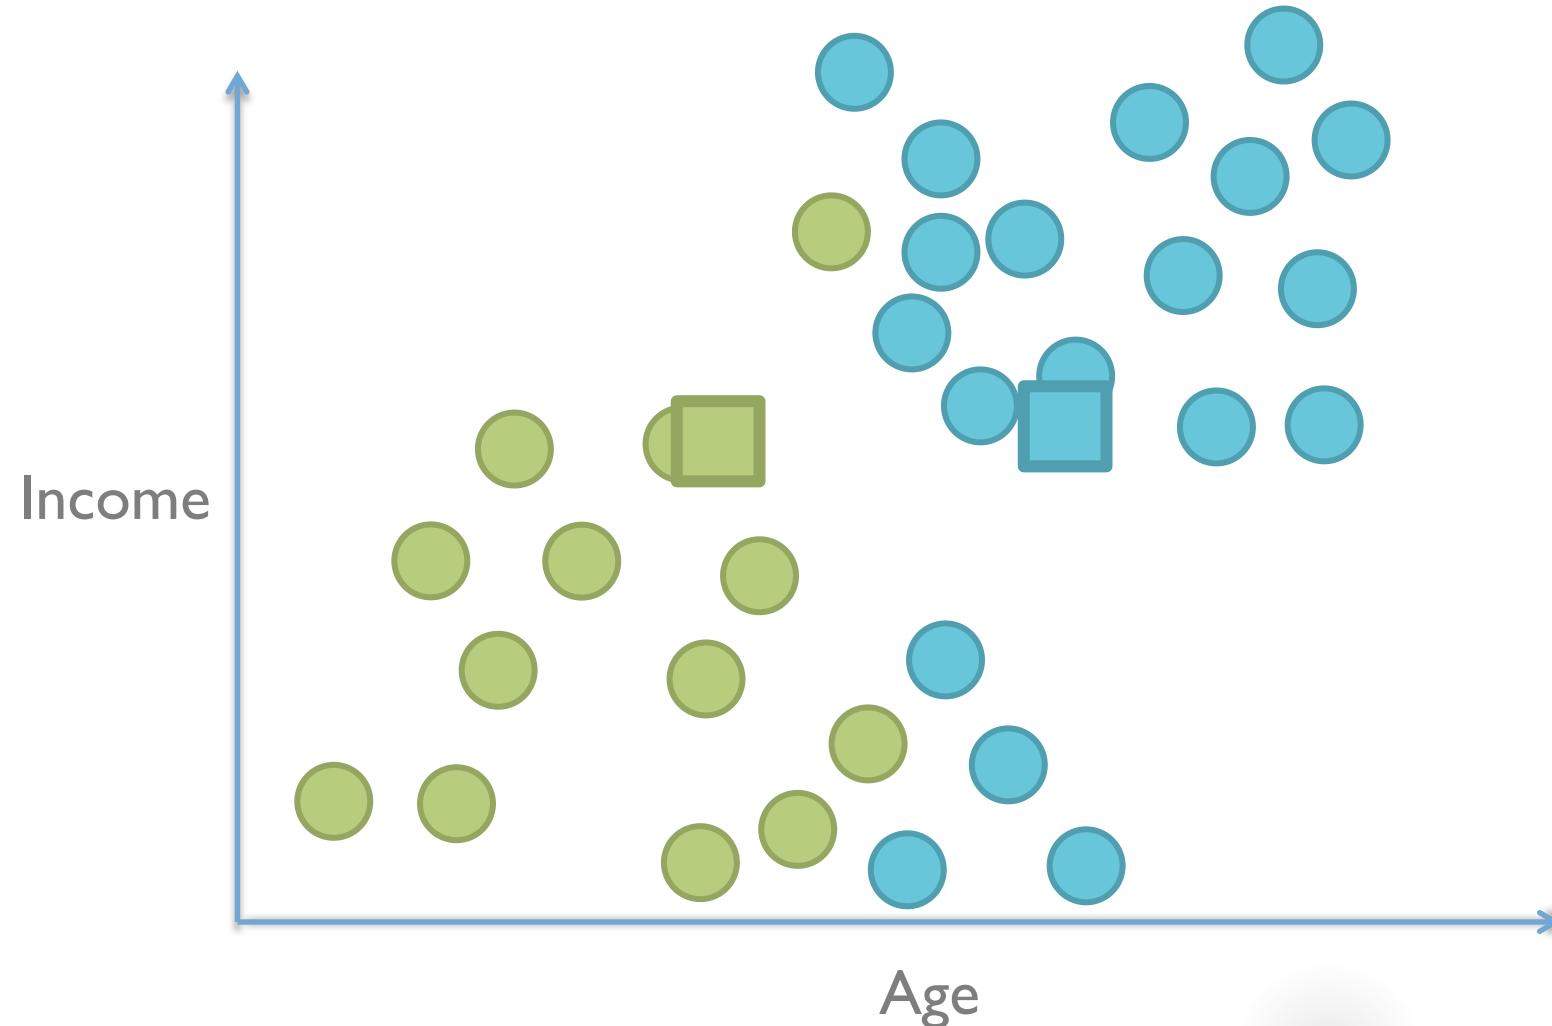
## K-Means K=2

Move each center to the cluster's mean



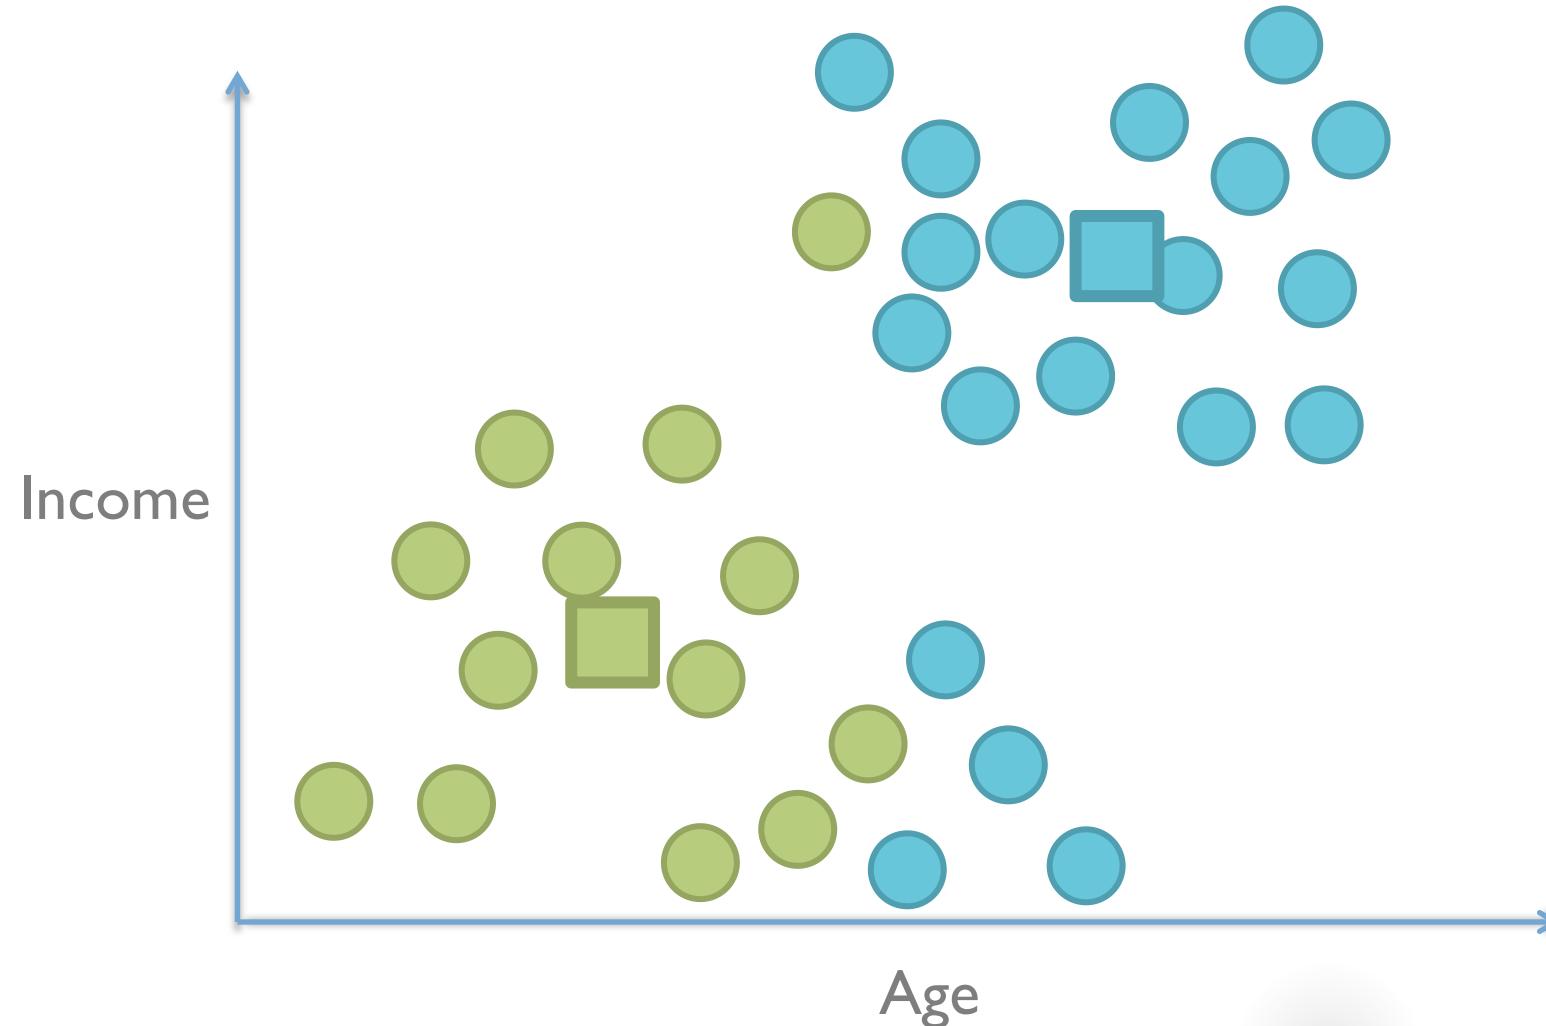
## K--Means K=2

Each point belongs to the closest center



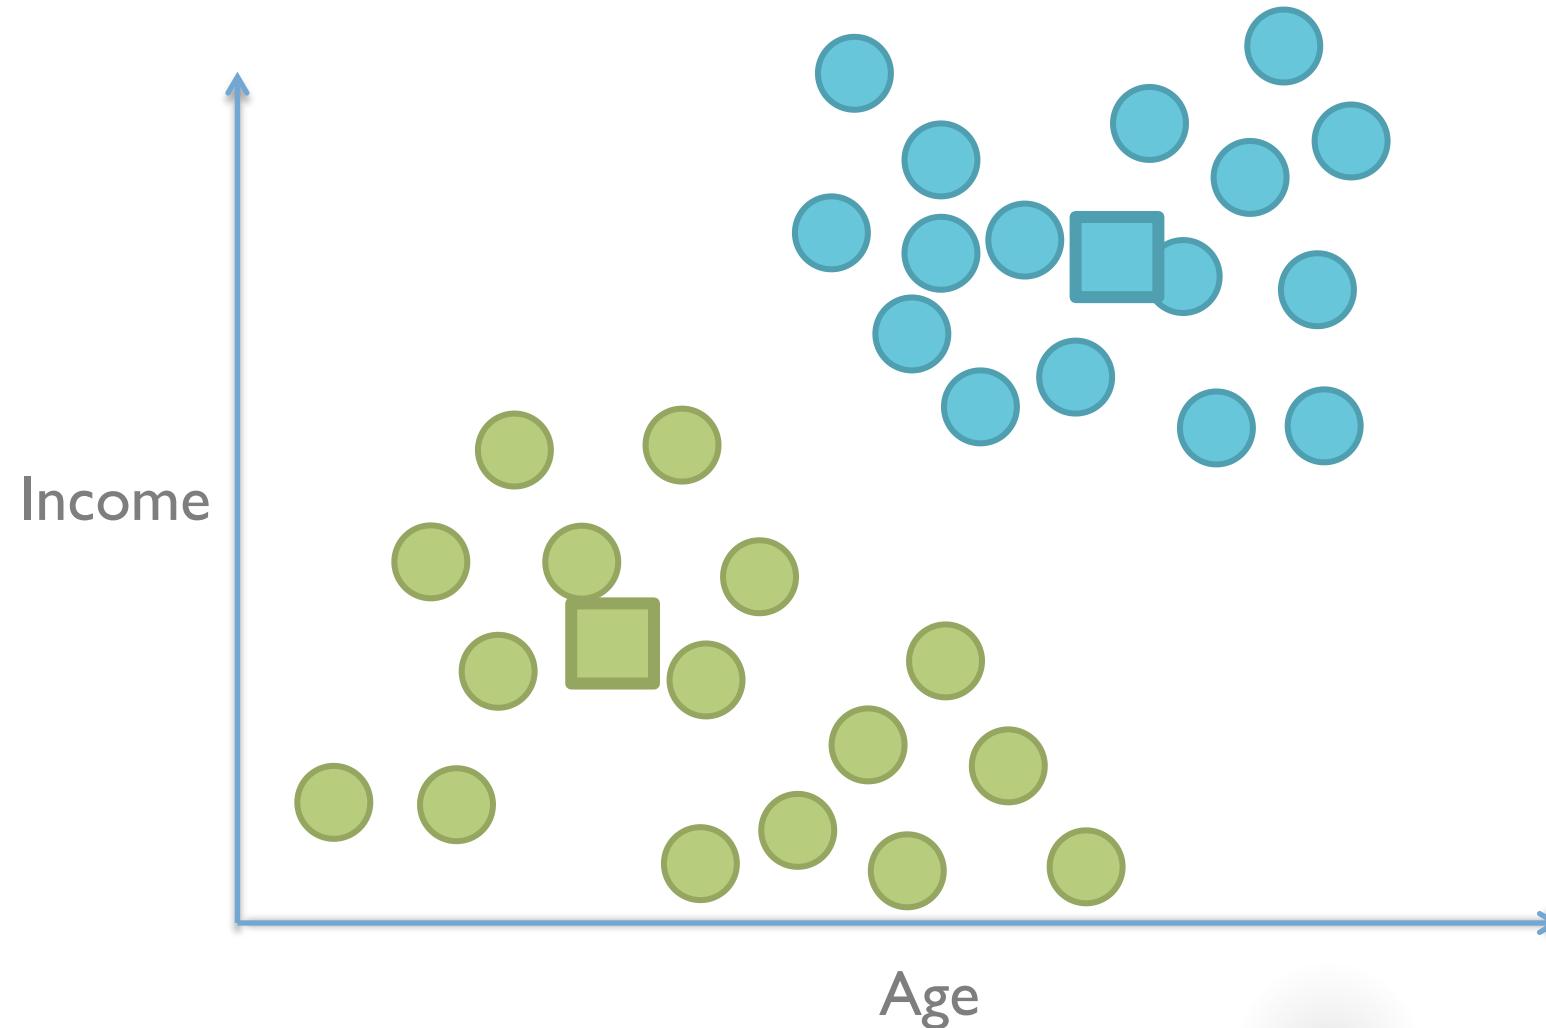
## K-Means K=2

Move each center to the cluster's mean



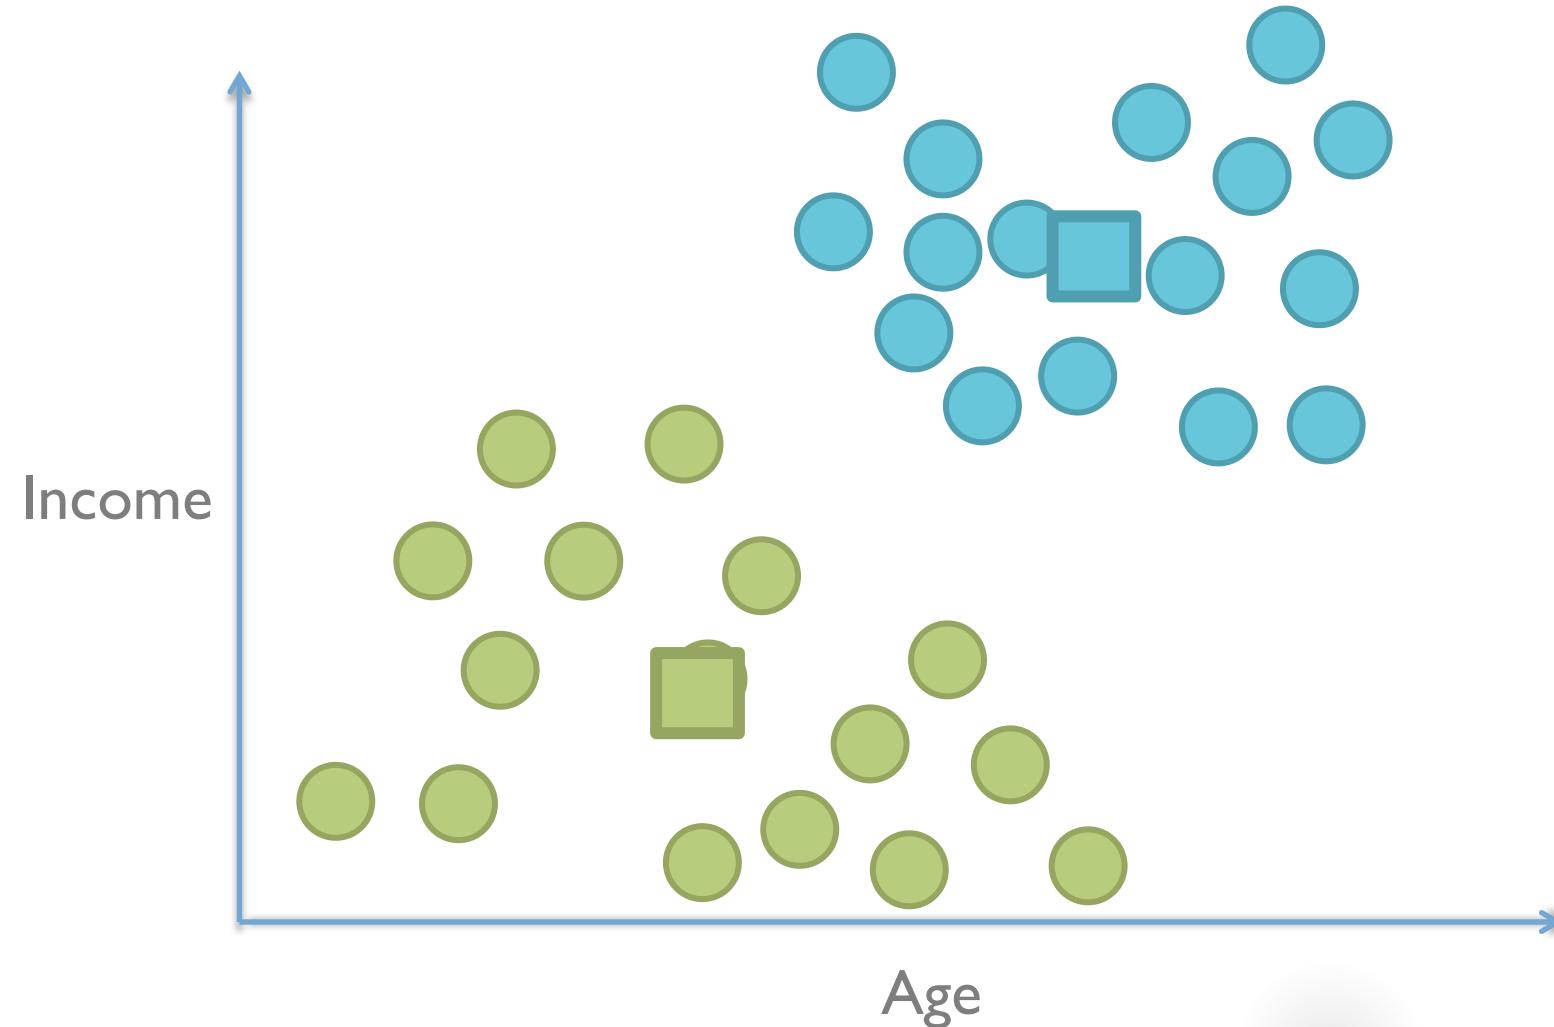
## K-Means K=2

Each point belongs to closest center



## K--Means K=2

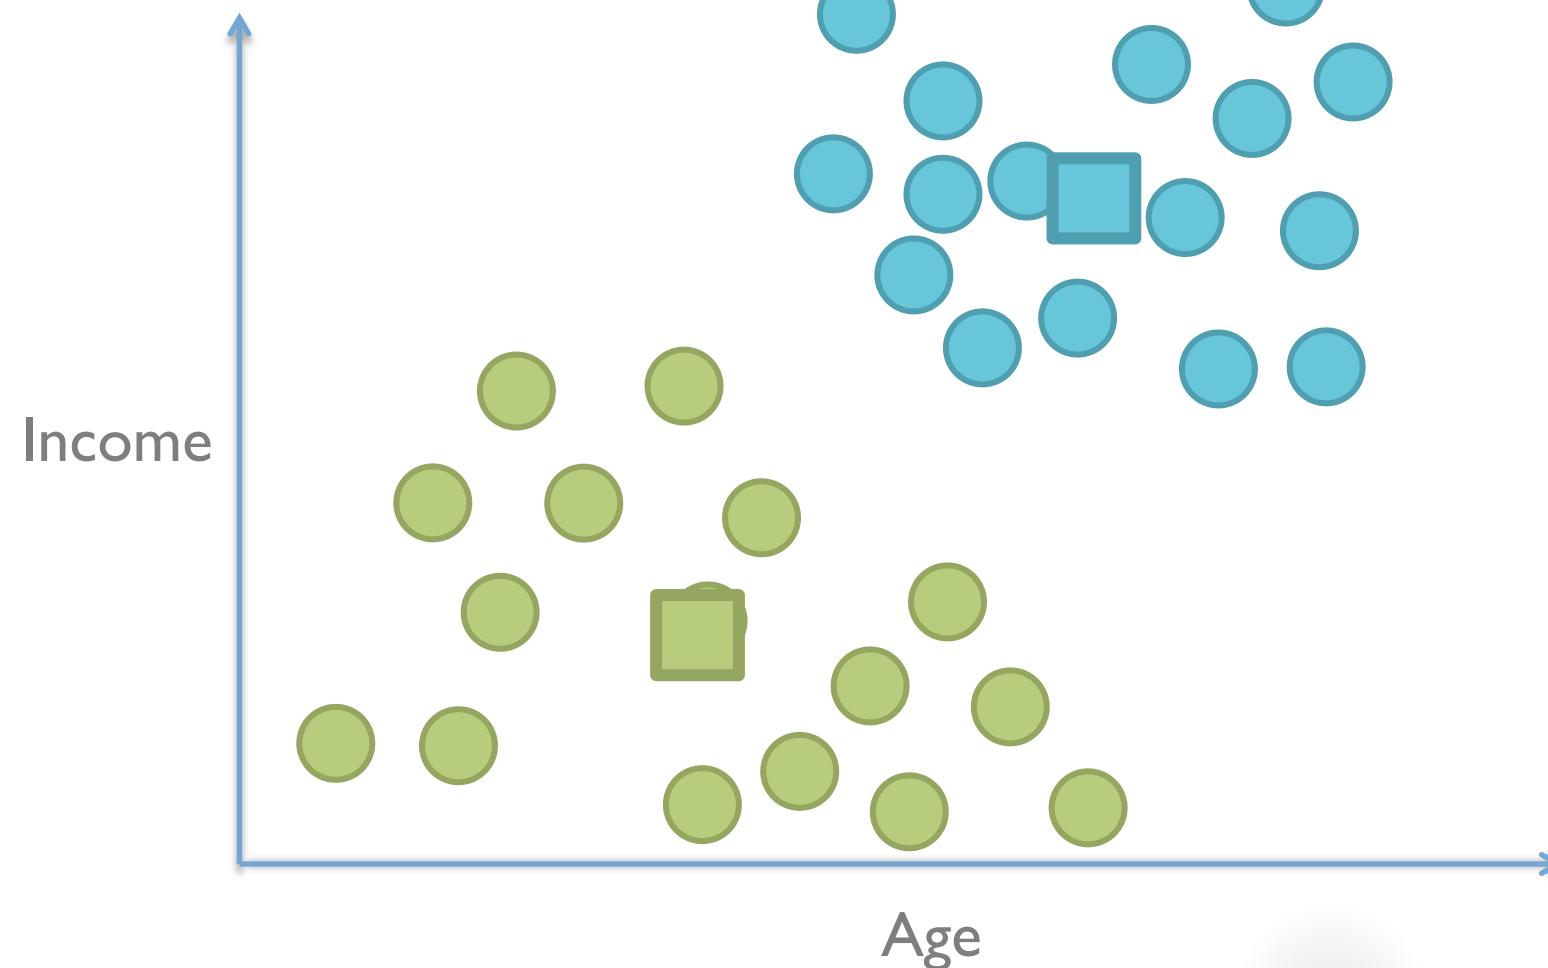
Move each center to the cluster's mean



## K-Means K=2

Points don't change anymore.

Converged!

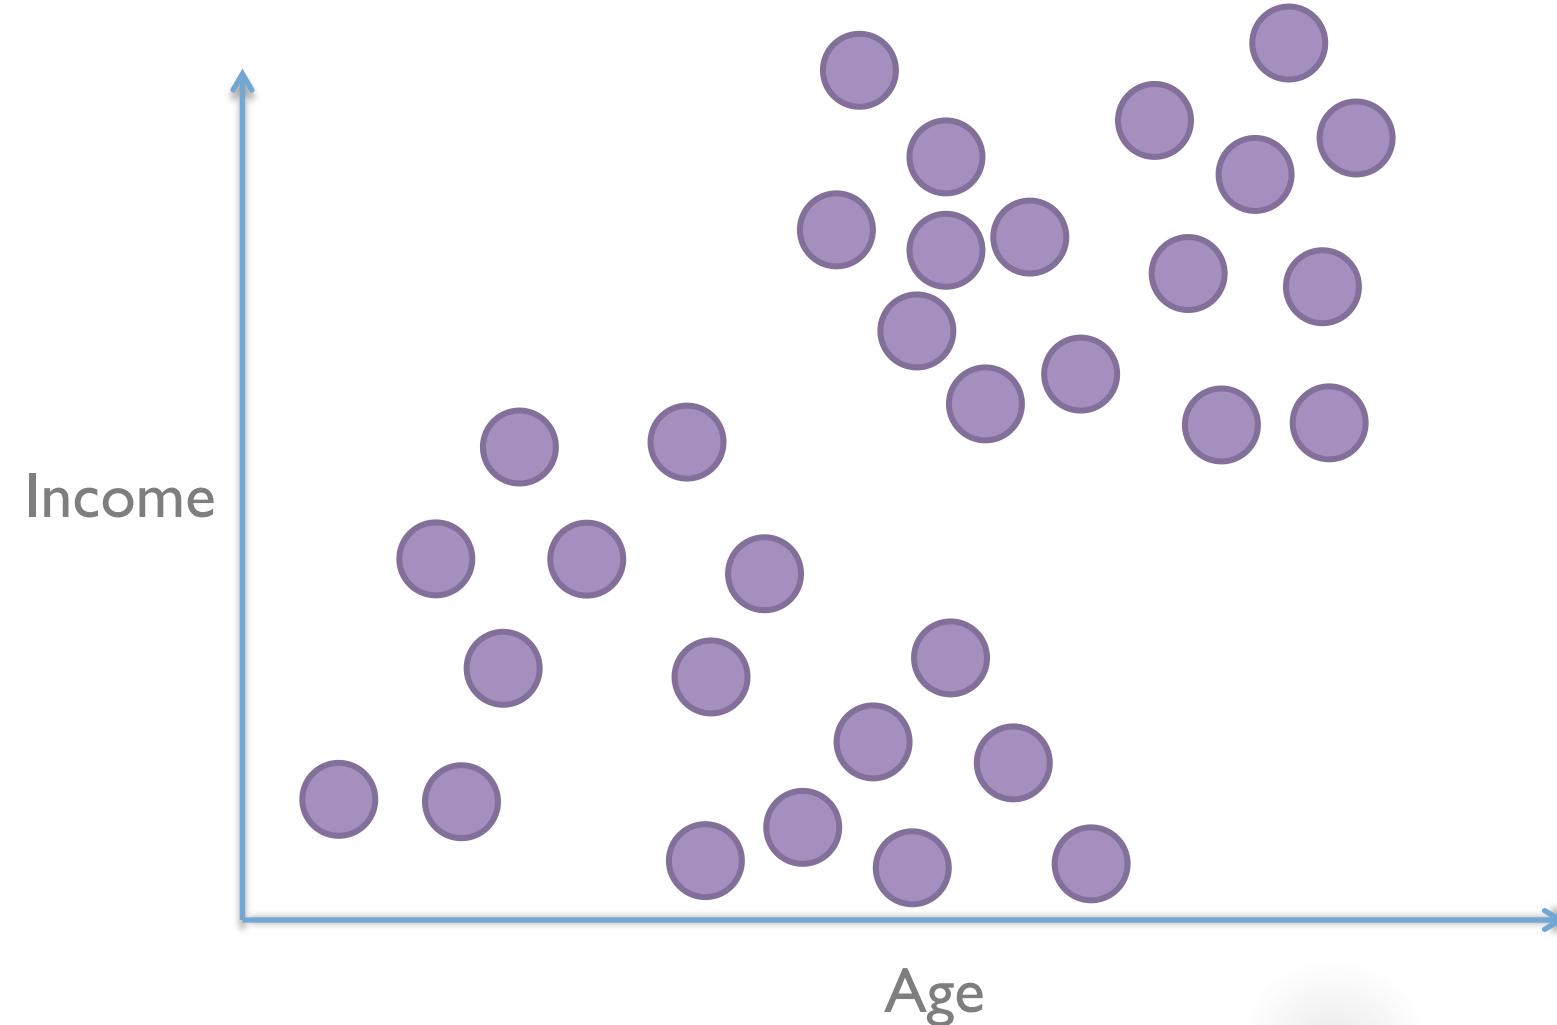




# Hierarchical Agglomerative Clustering

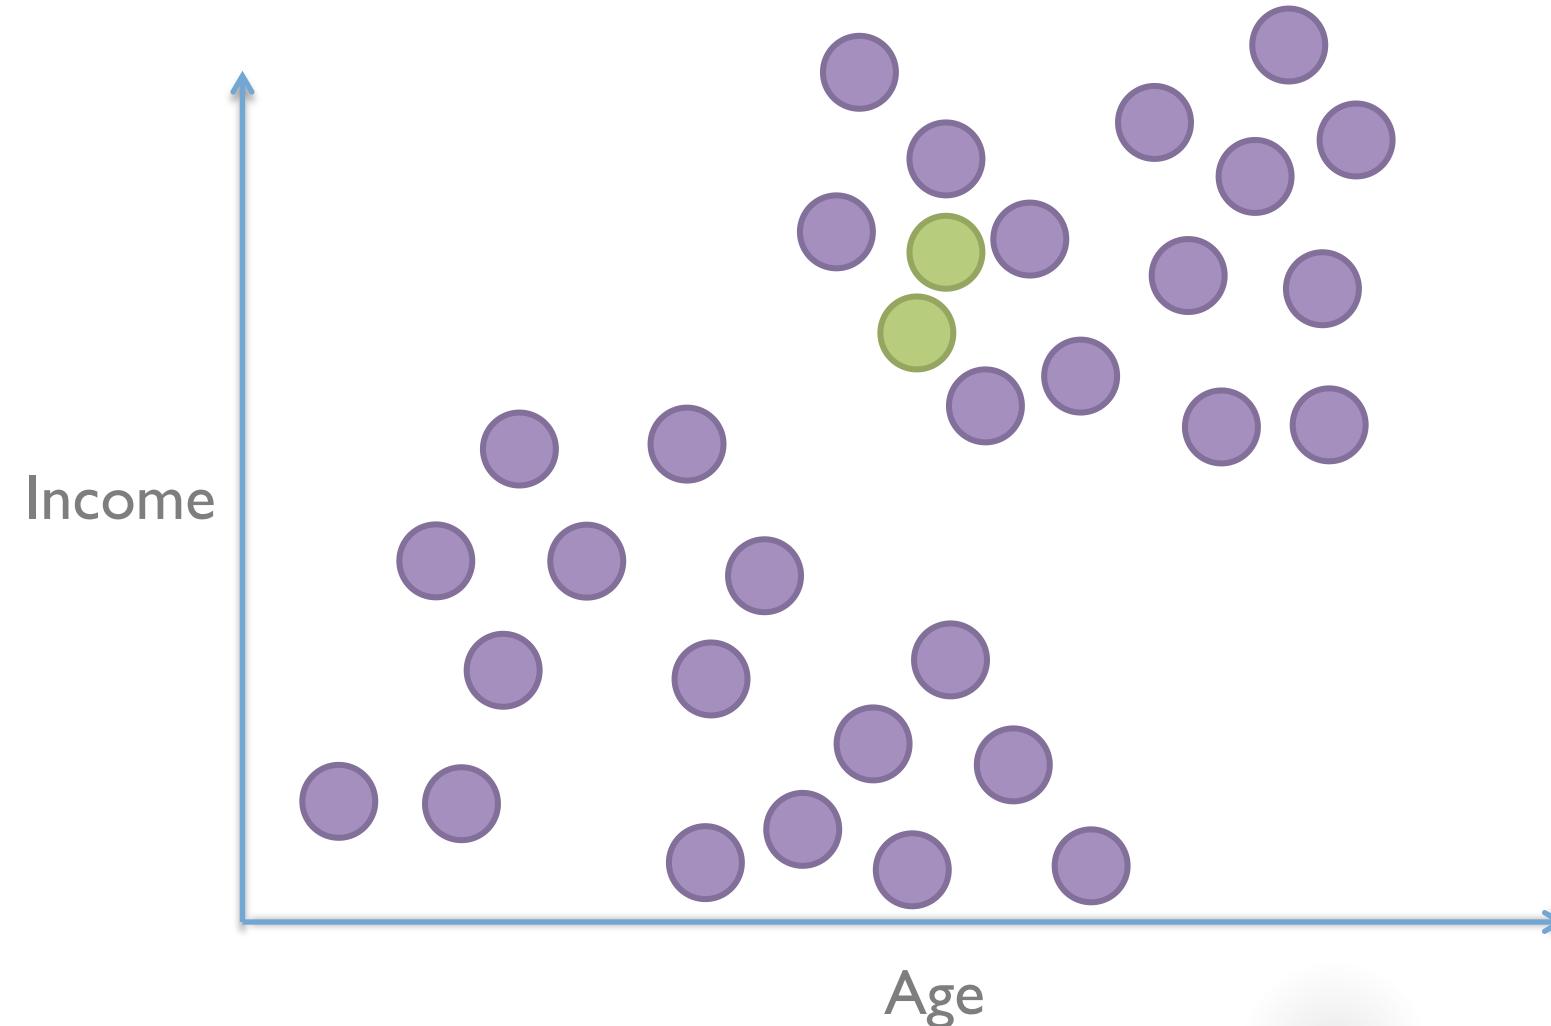
---

# Hierarchical Agglomerative Clustering



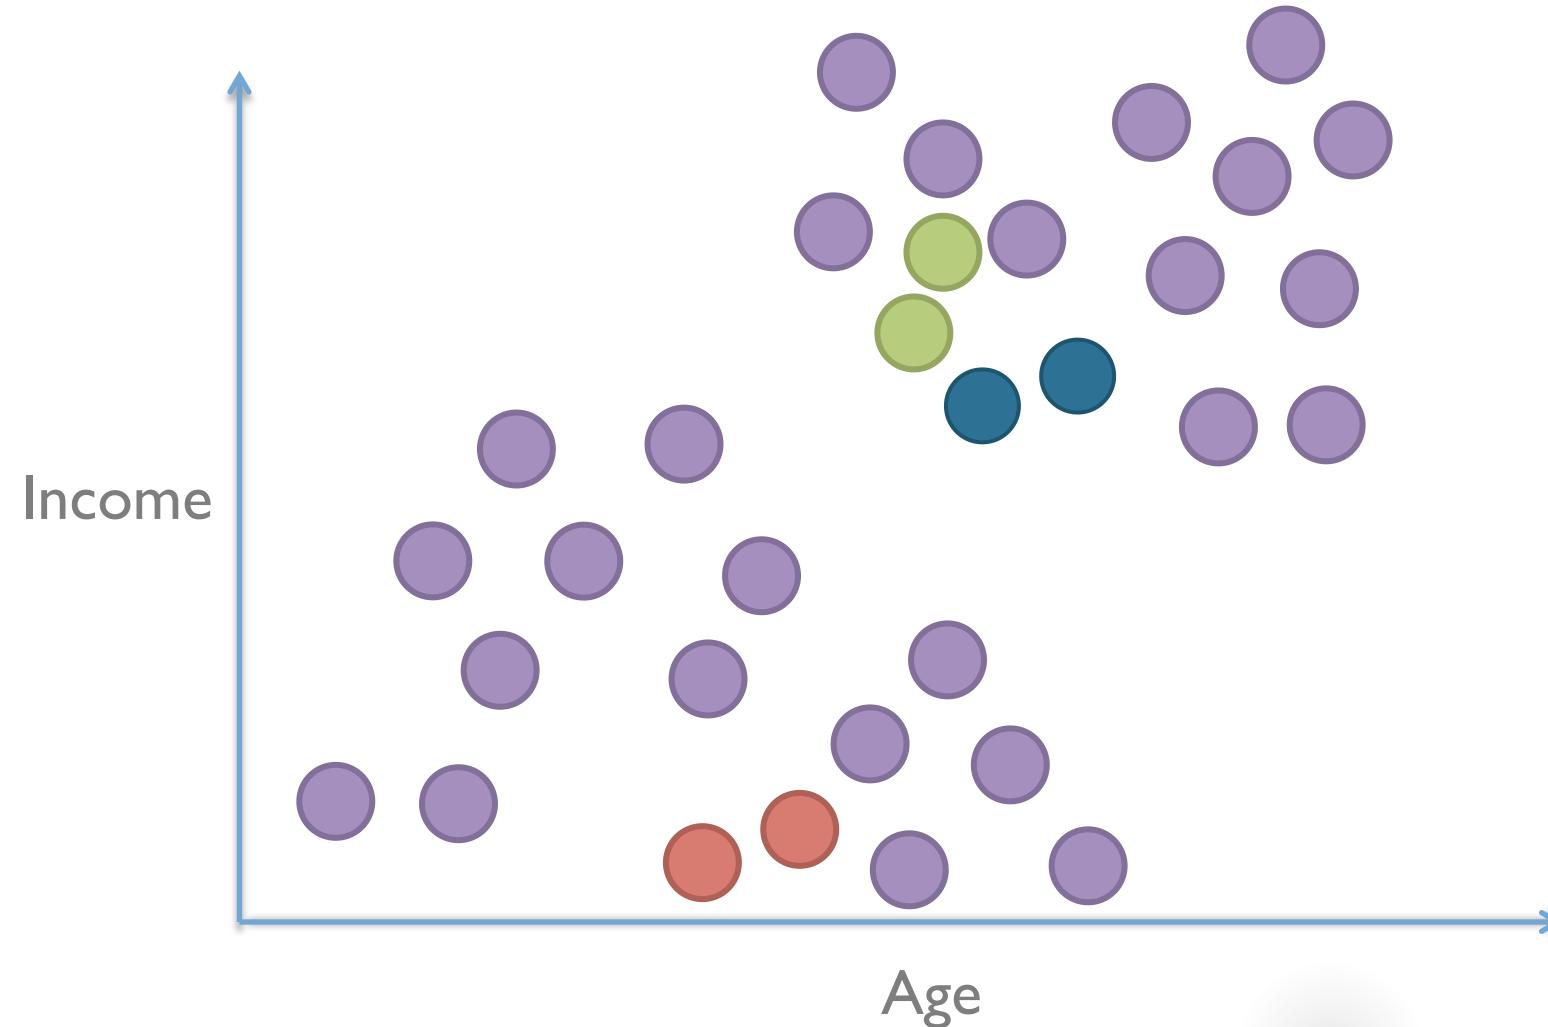
# Agglomerative Hierarchical Clustering

Find closet pair, merge into a cluster



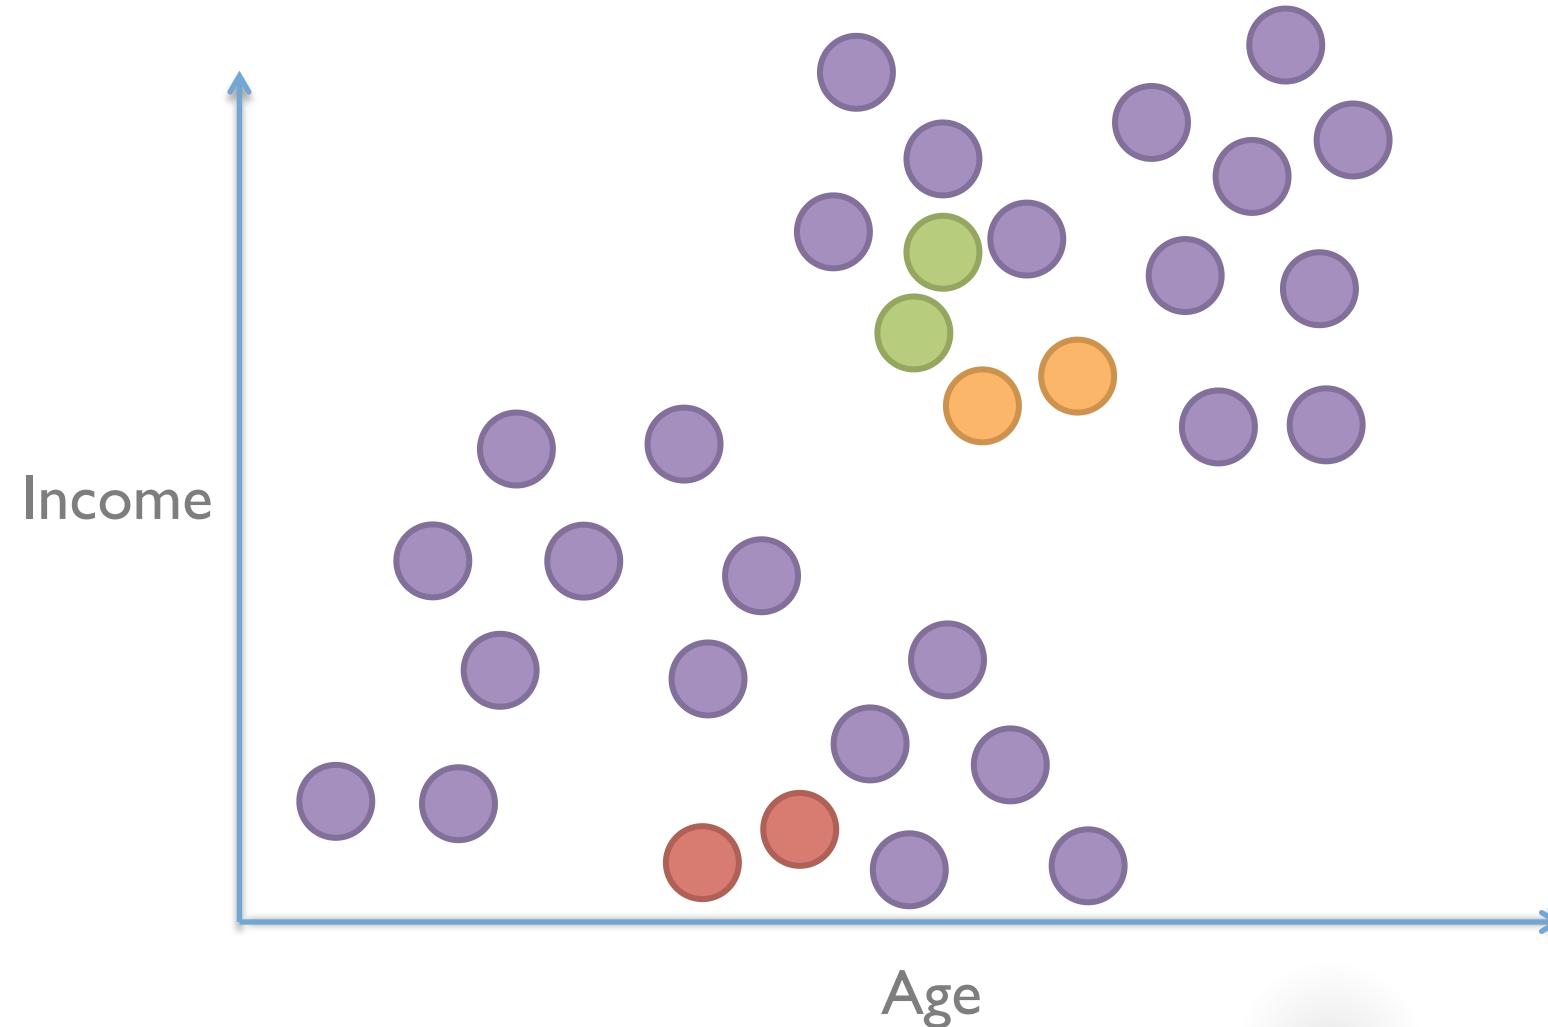
# Agglomerative Hierarchical Clustering

Find next closet pair, merge



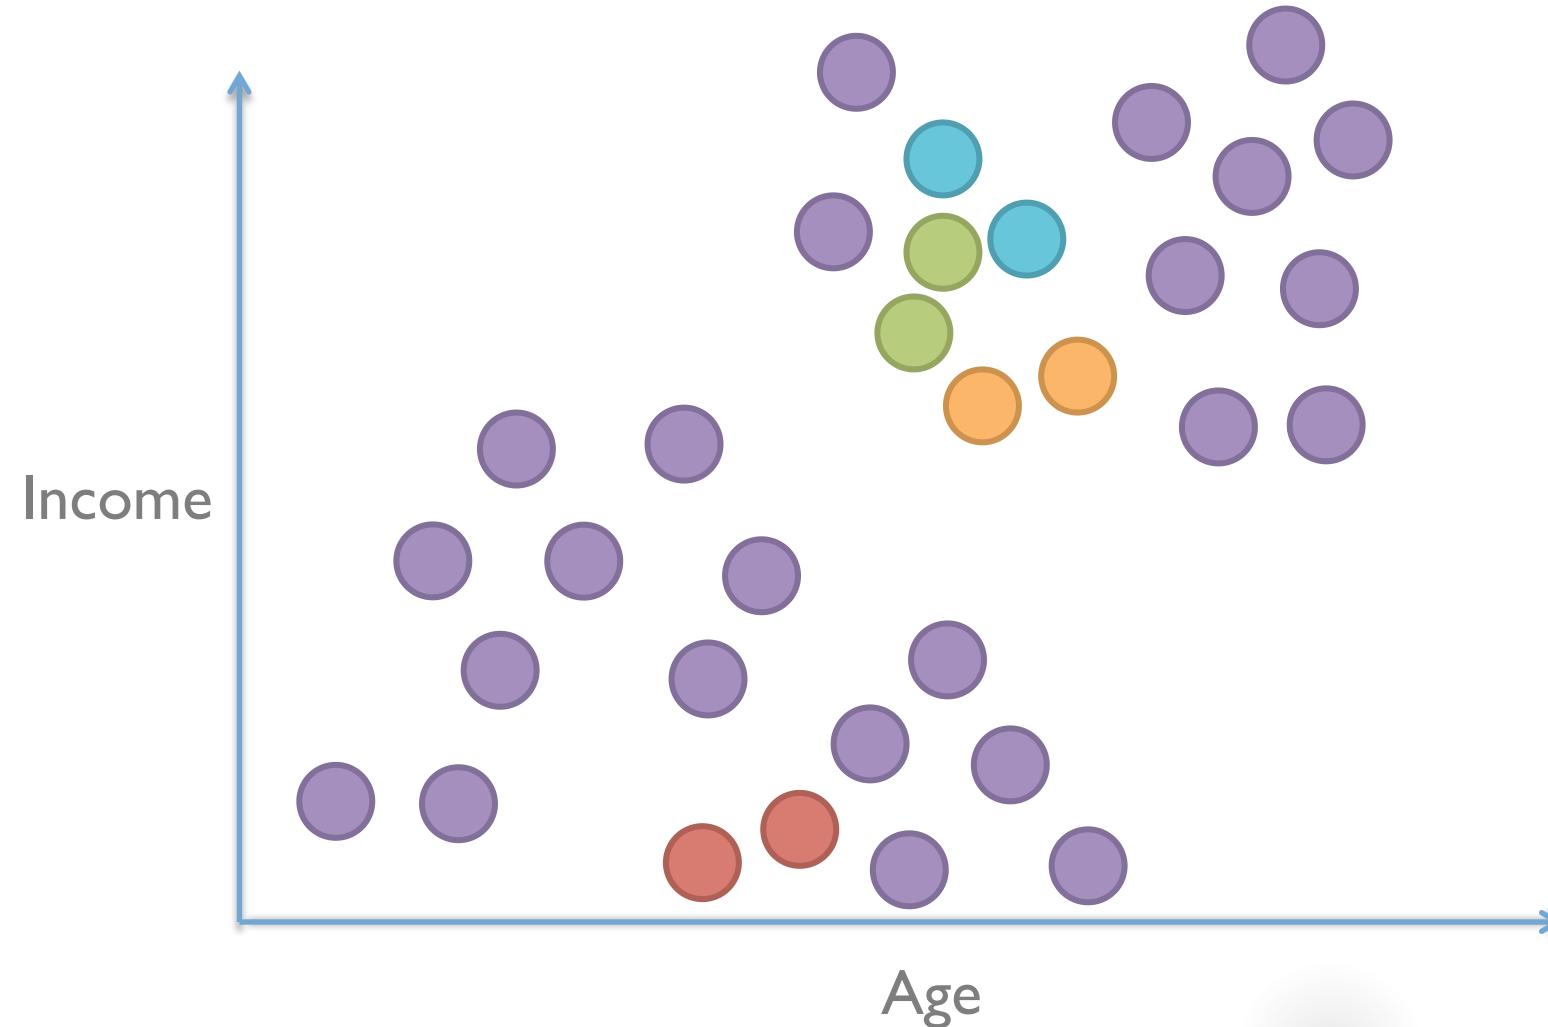
# Agglomerative Hierarchical Clustering

Find next closet pair, merge



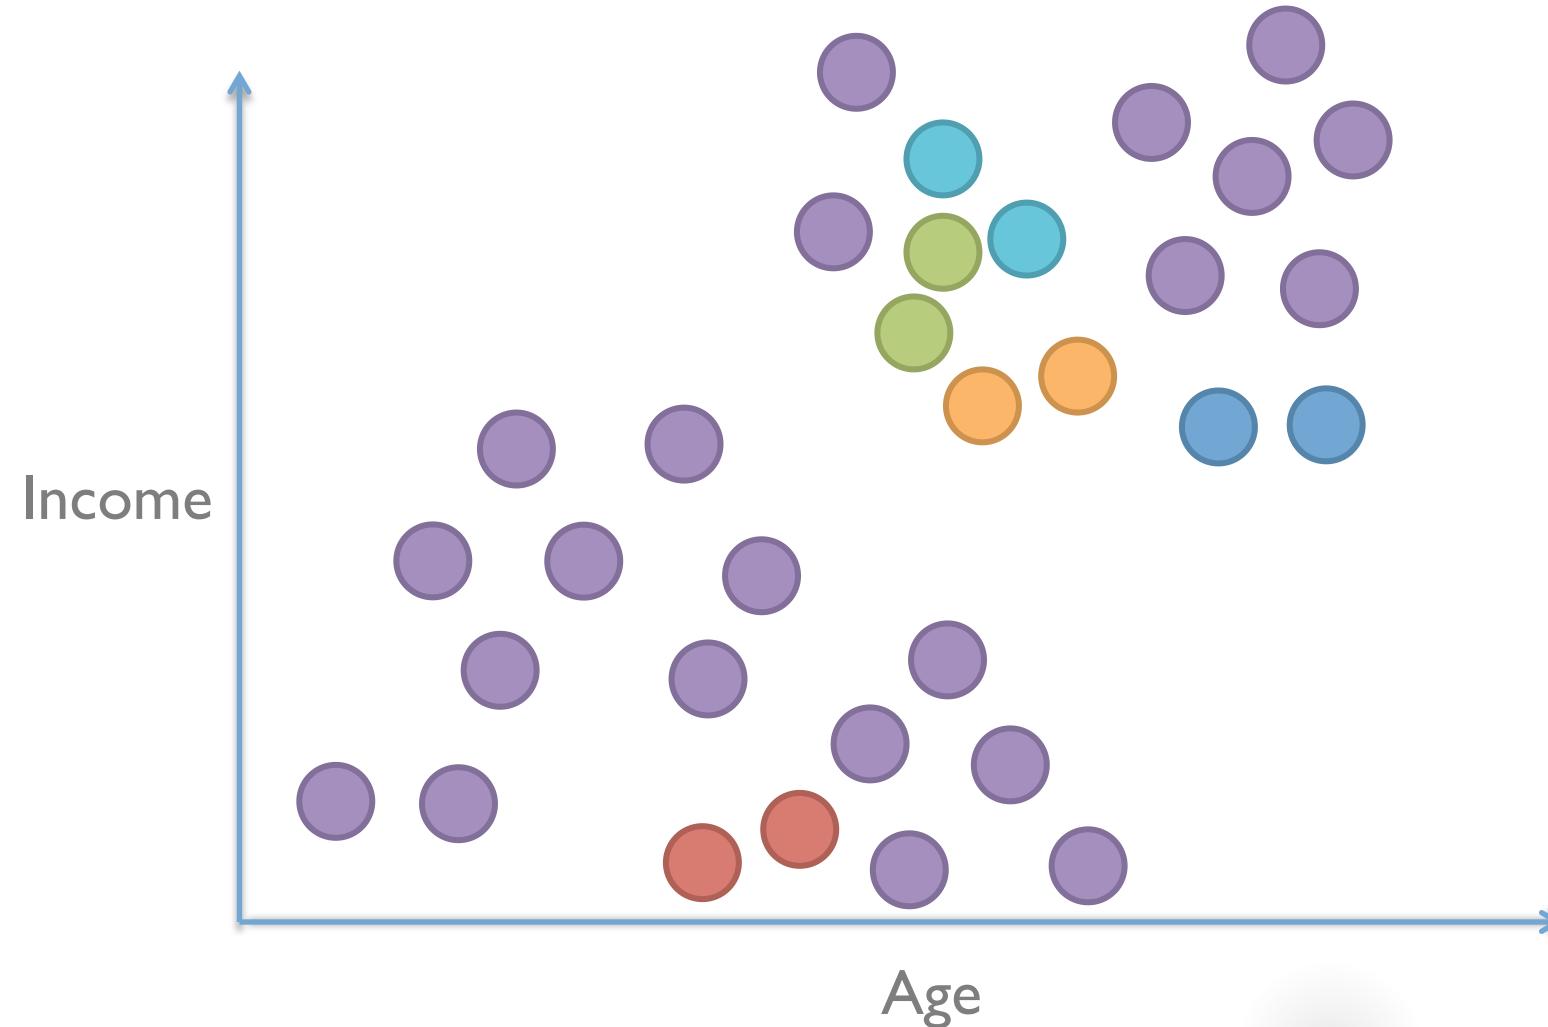
# Agglomerative Hierarchical Clustering

Find next closet pair, merge



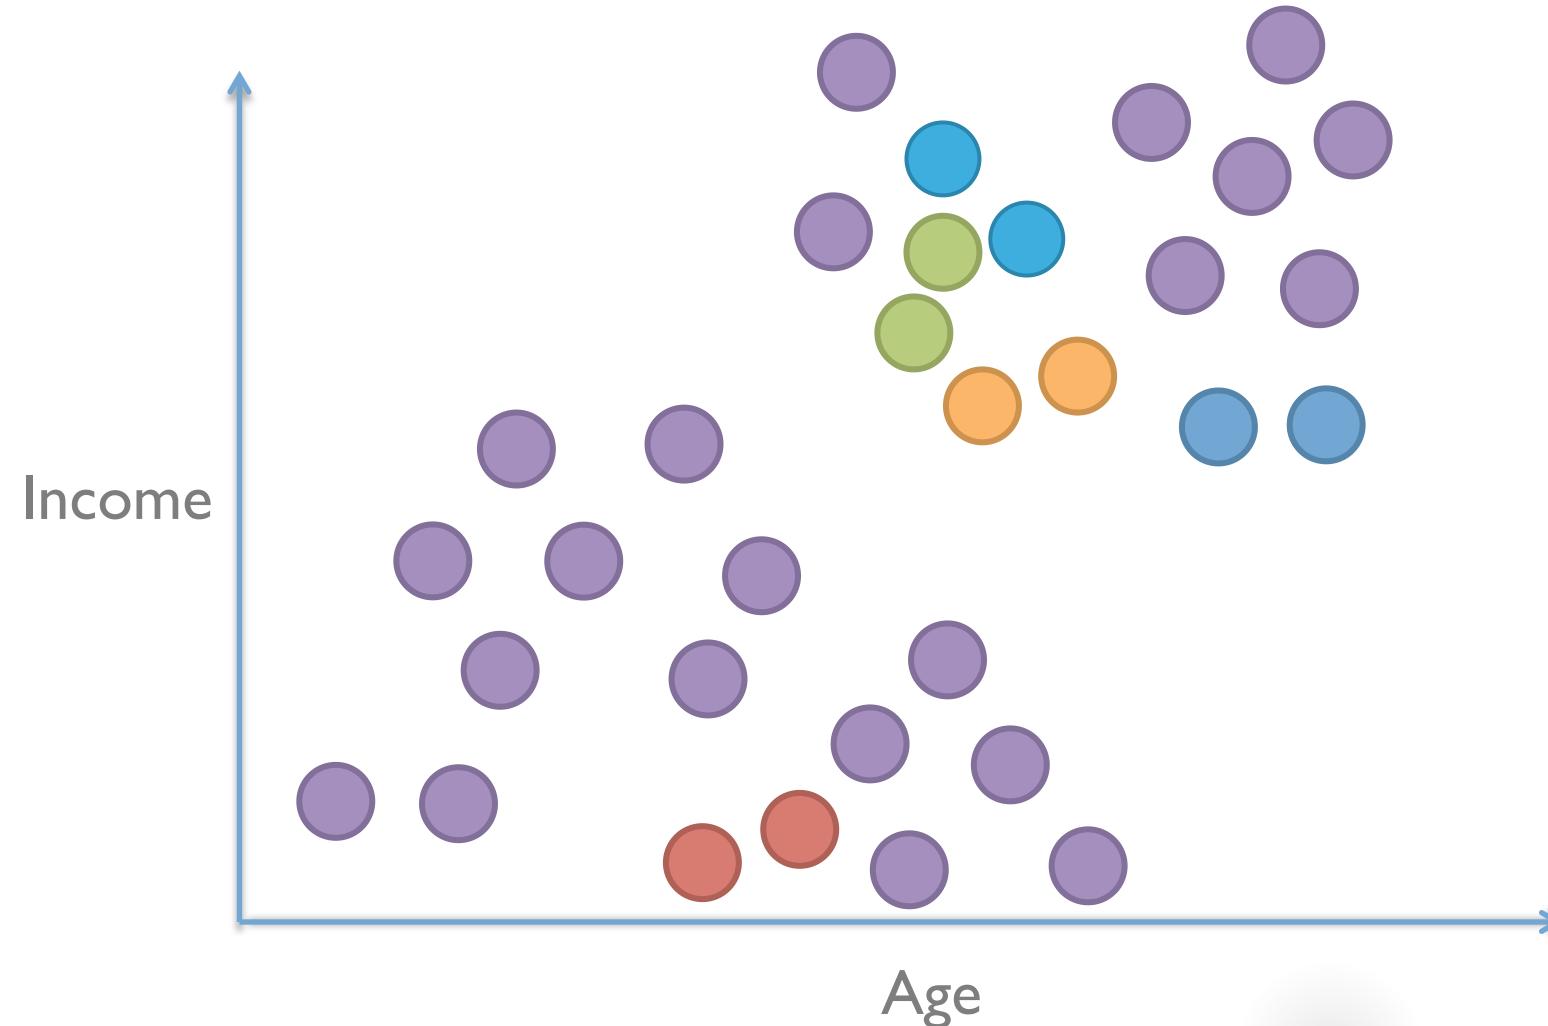
# Agglomerative Hierarchical Clustering

Find next closet pair, merge



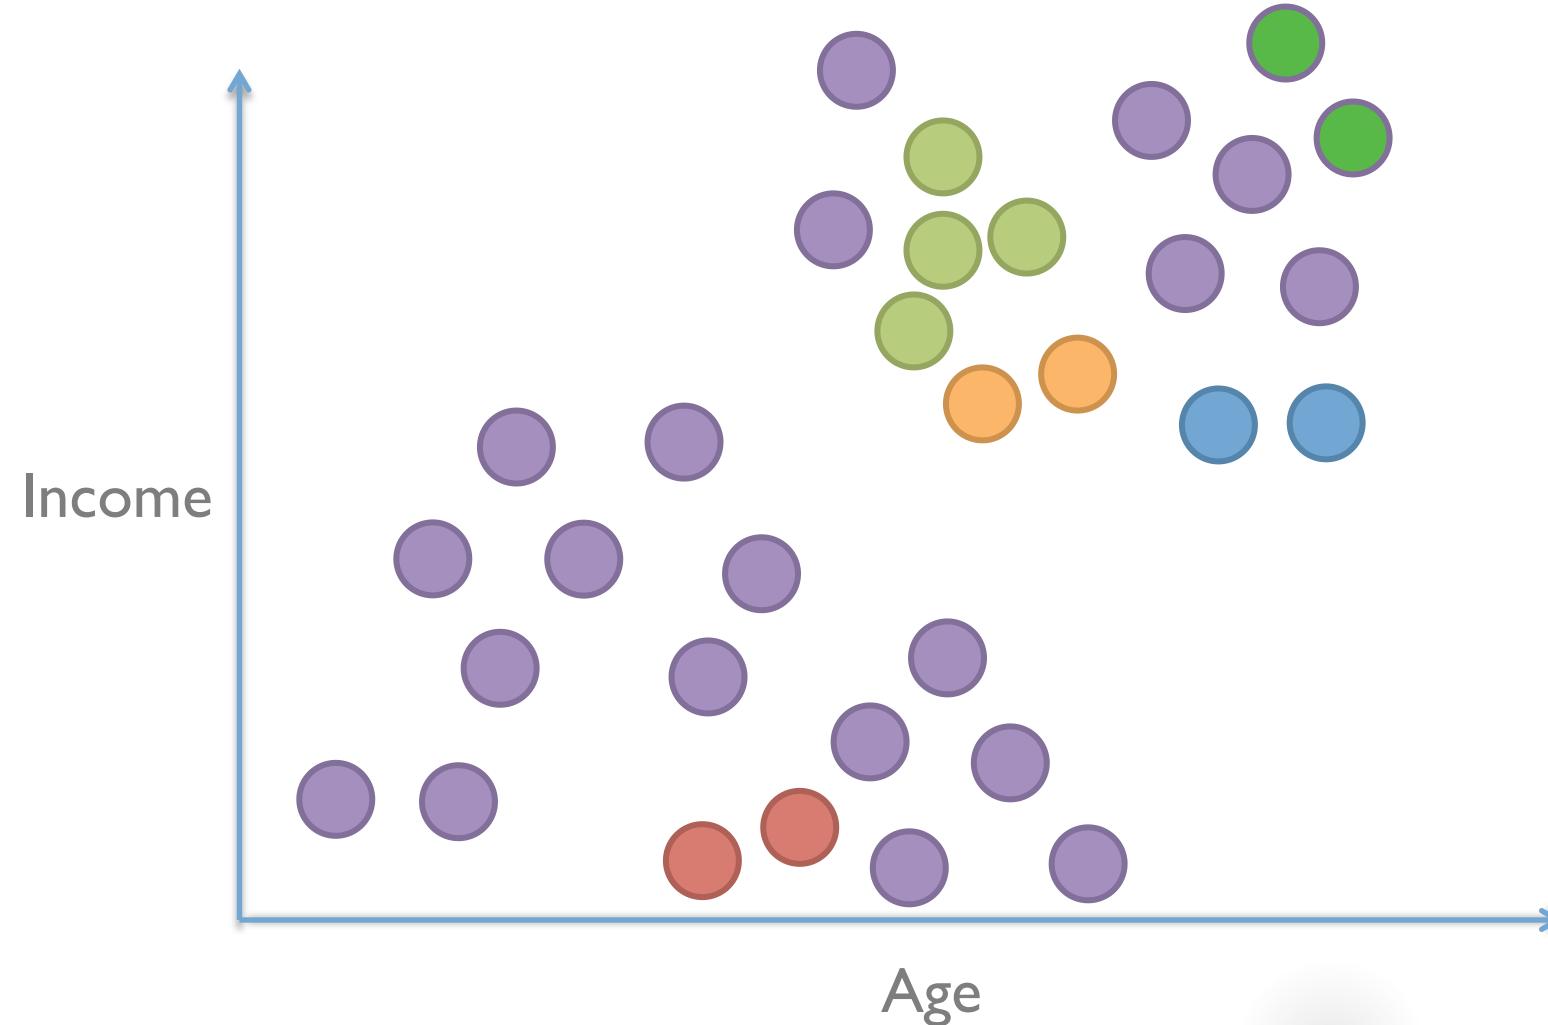
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



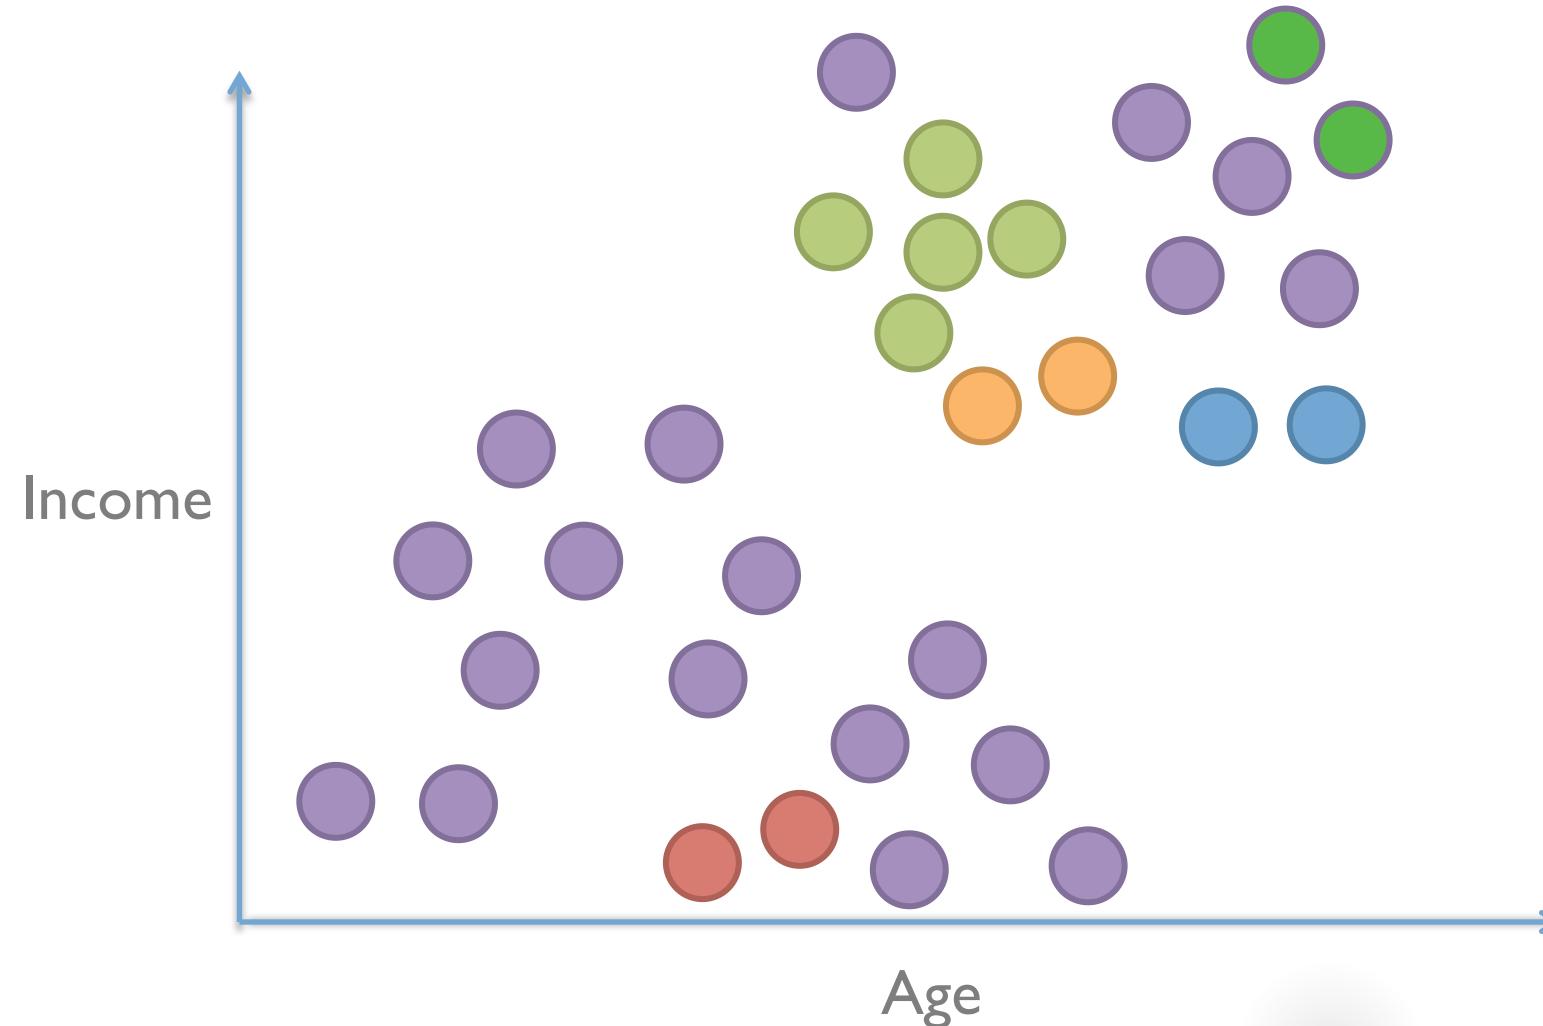
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



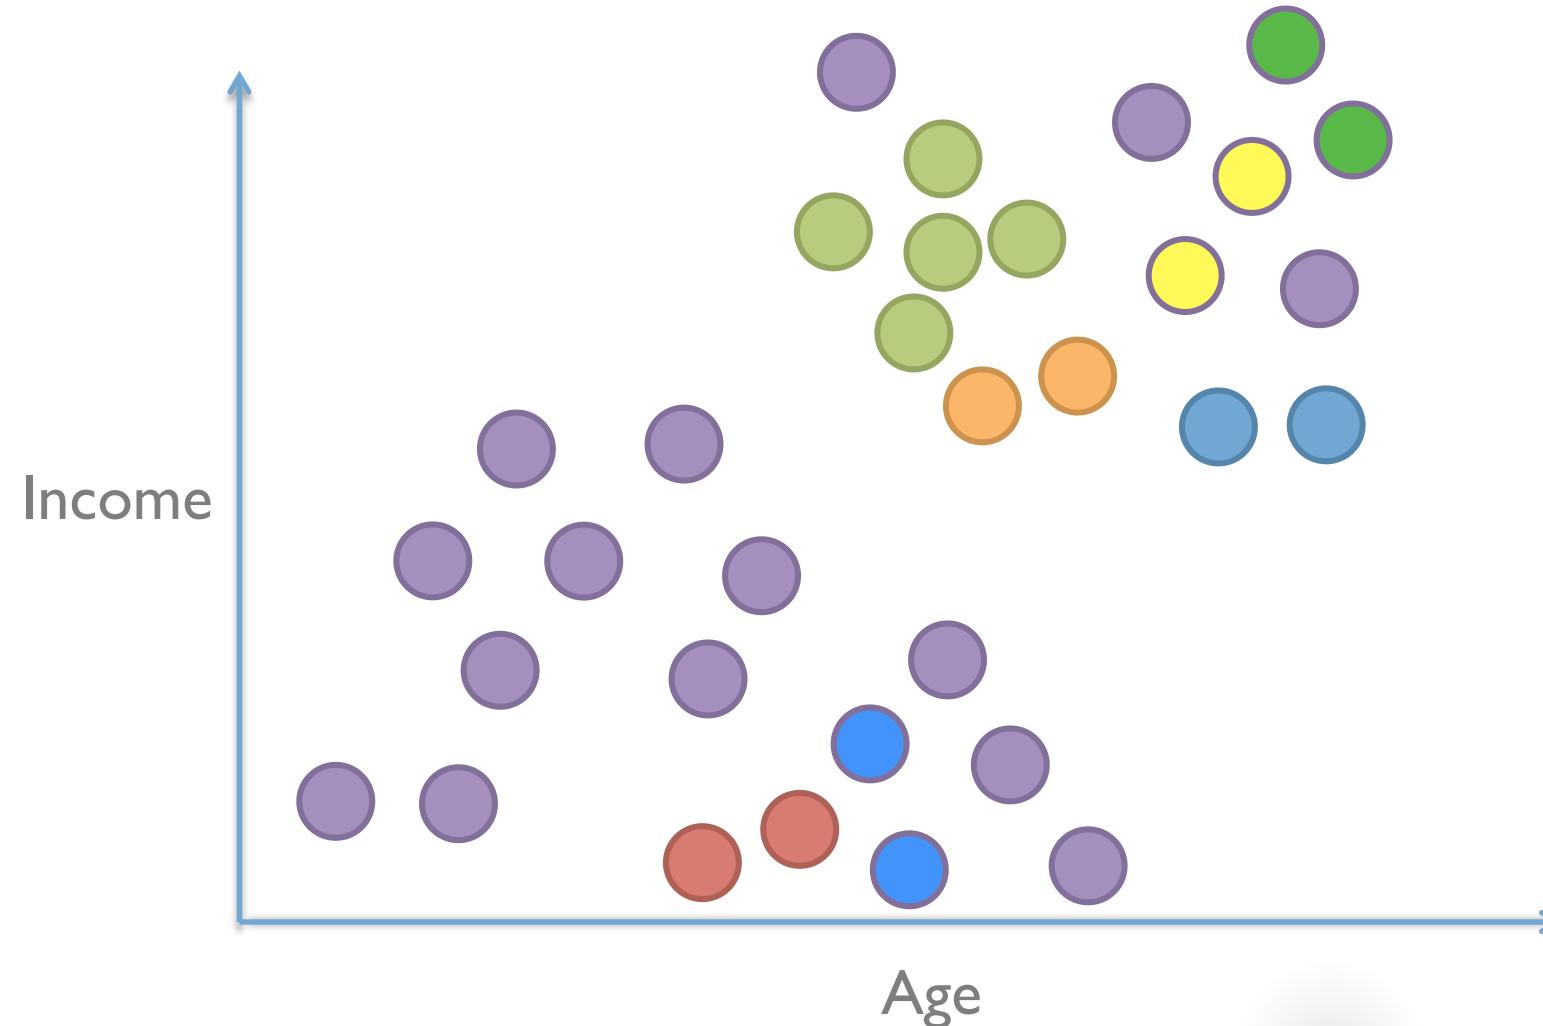
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



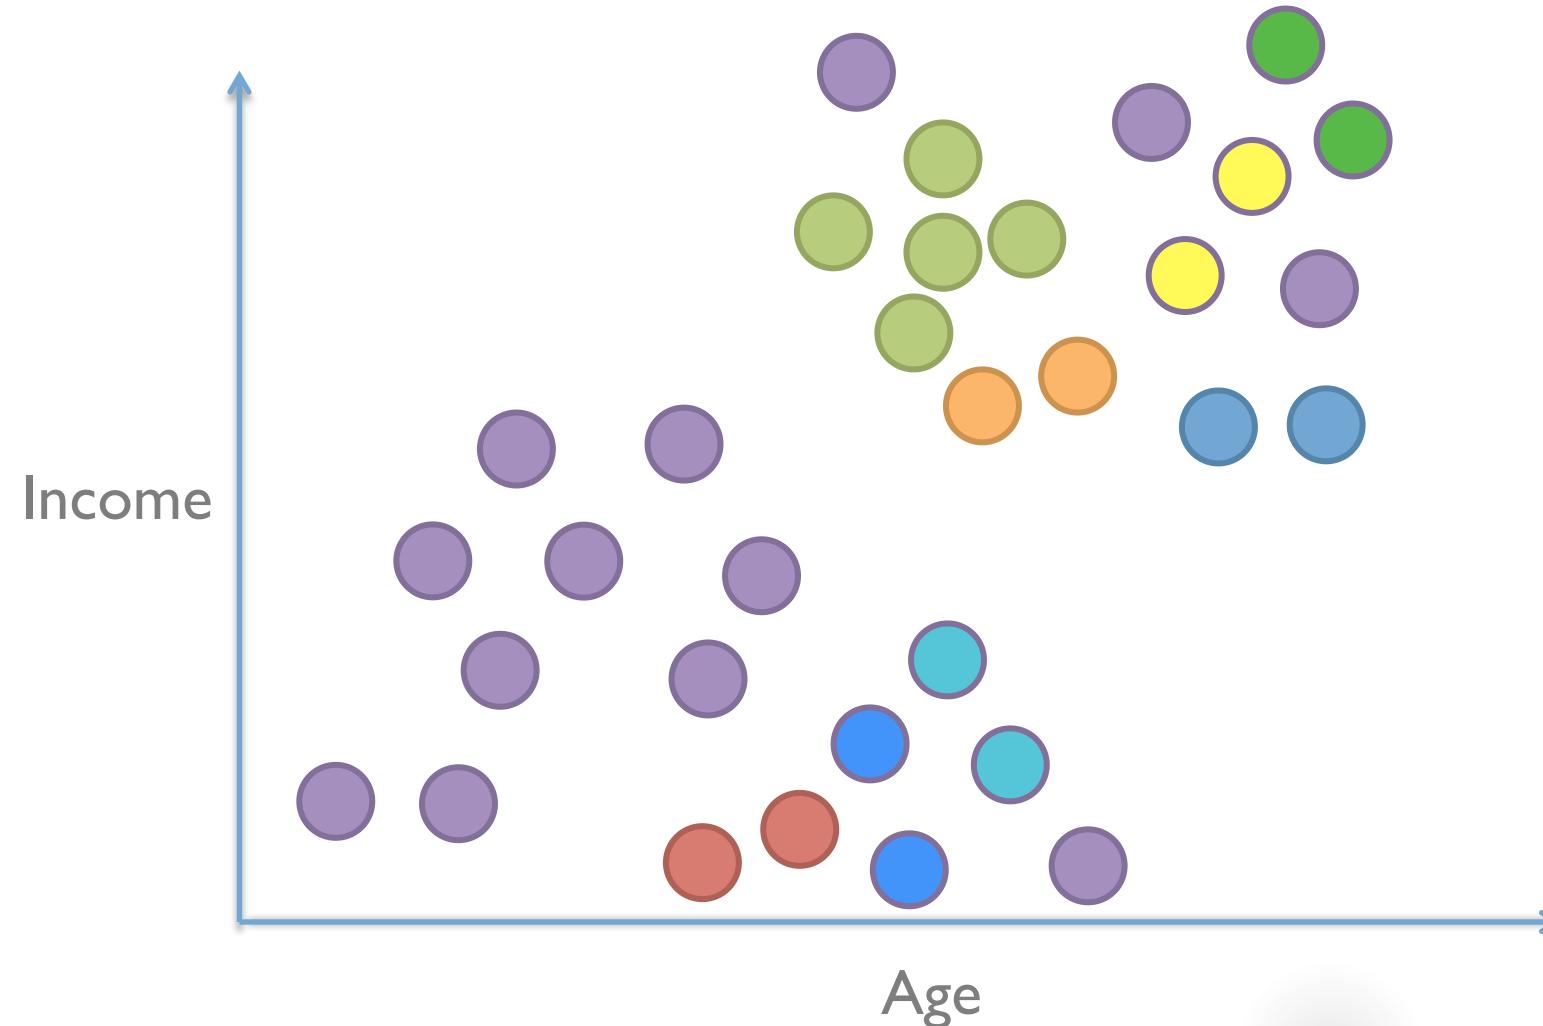
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



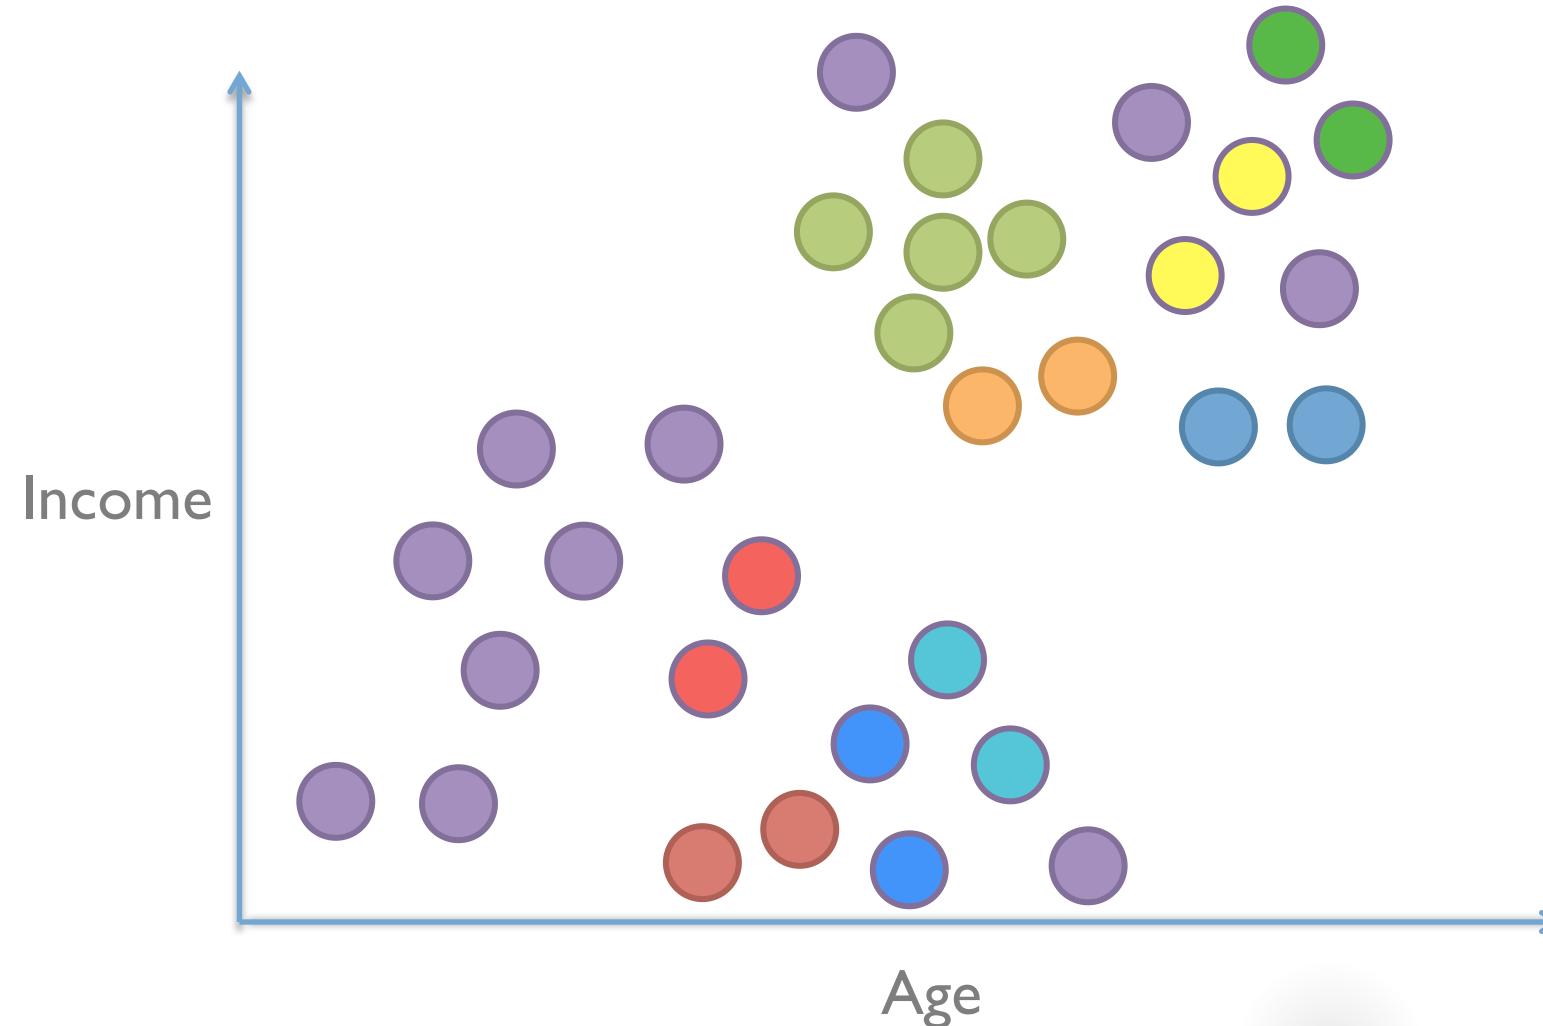
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



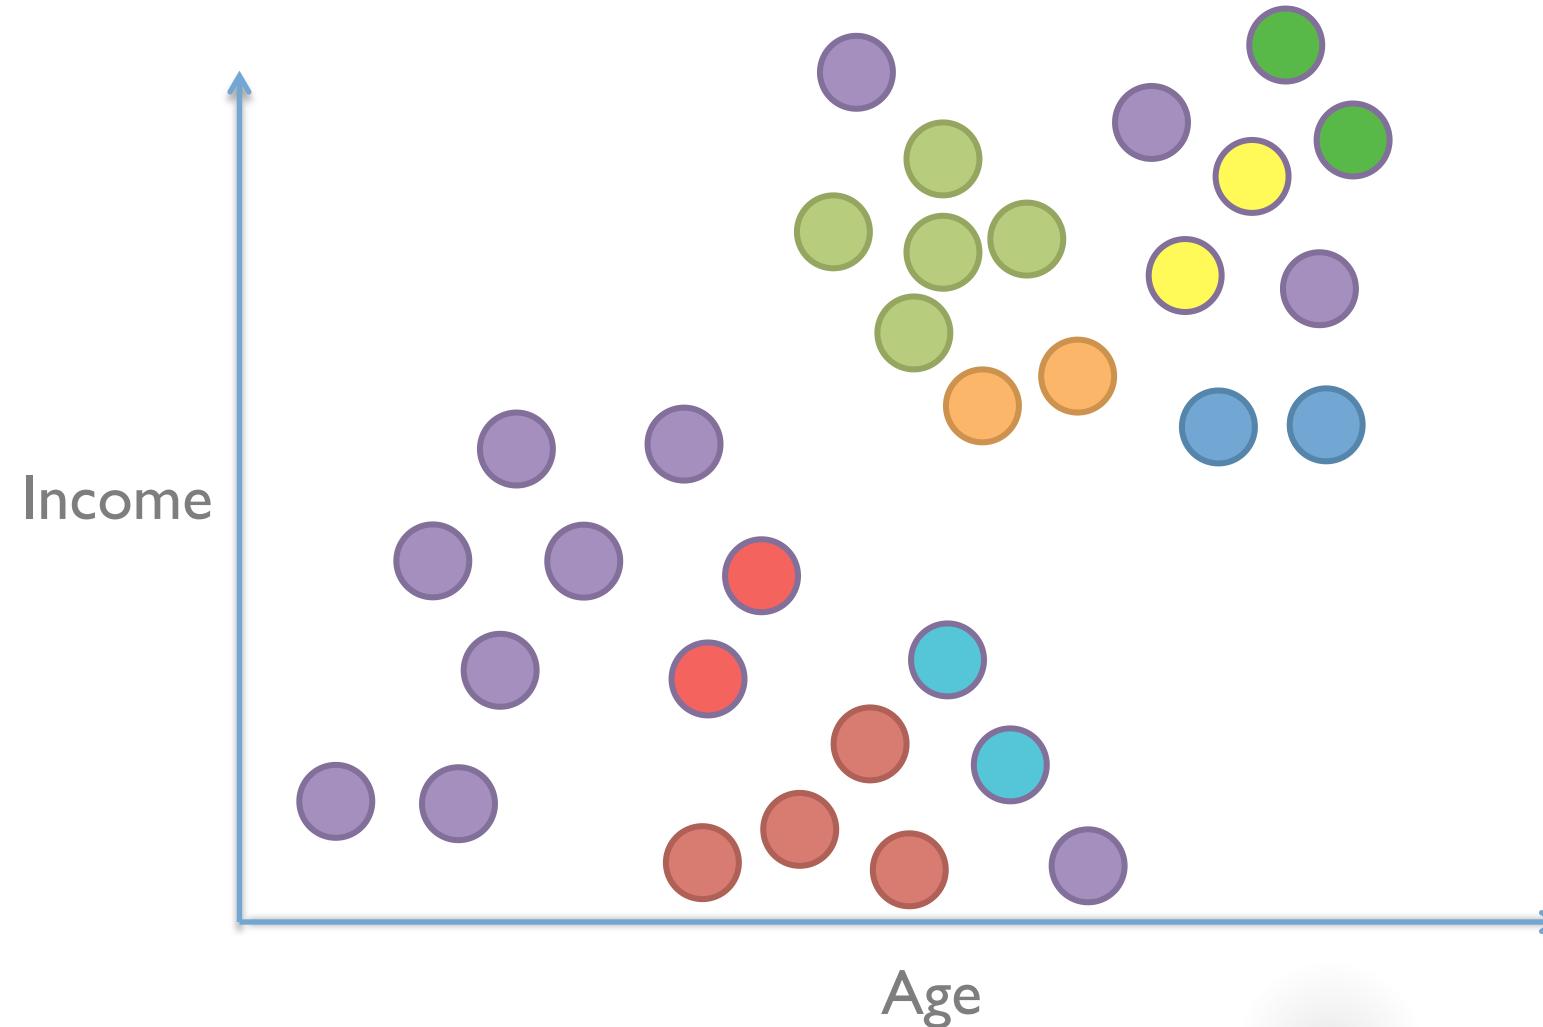
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



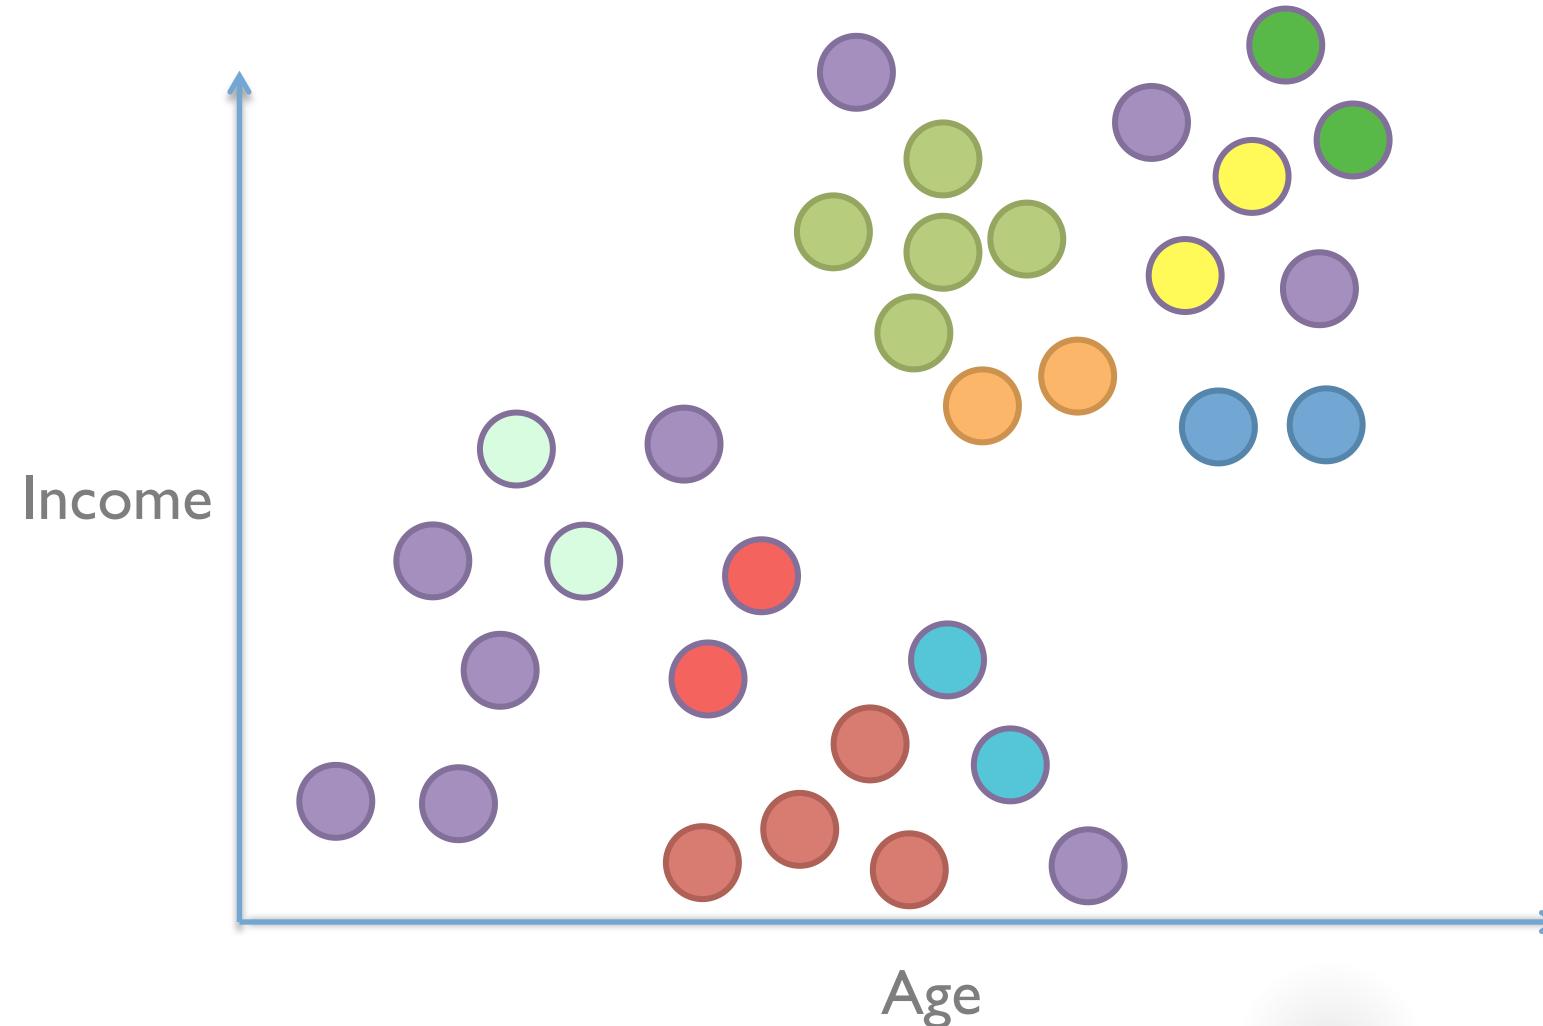
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



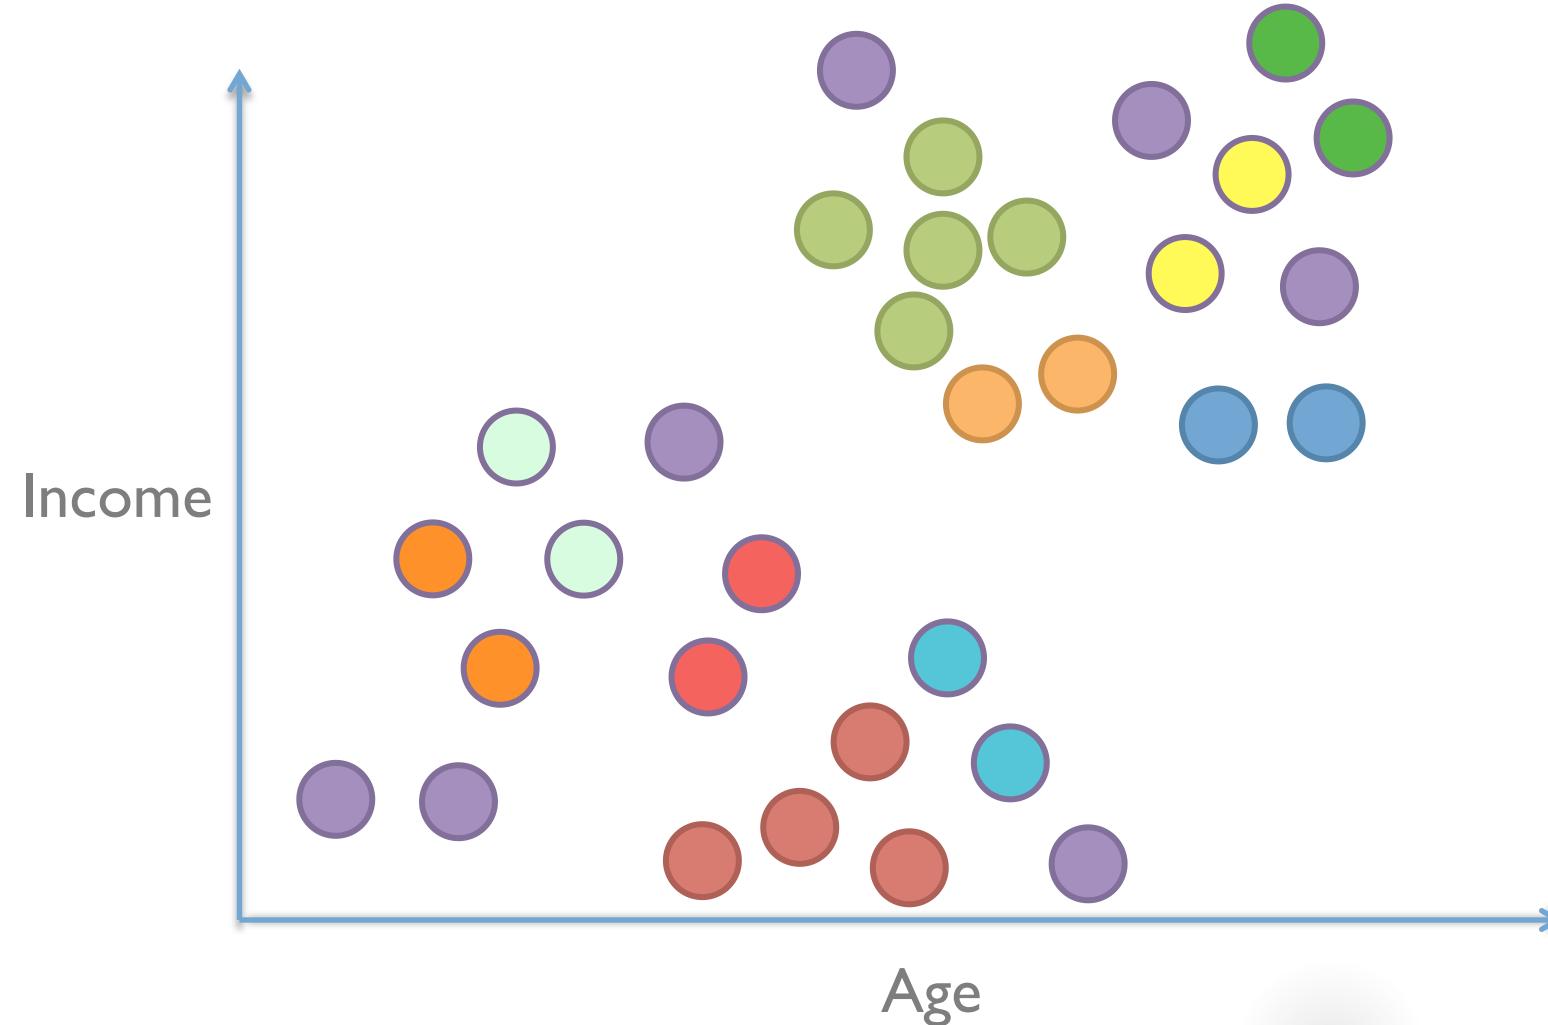
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



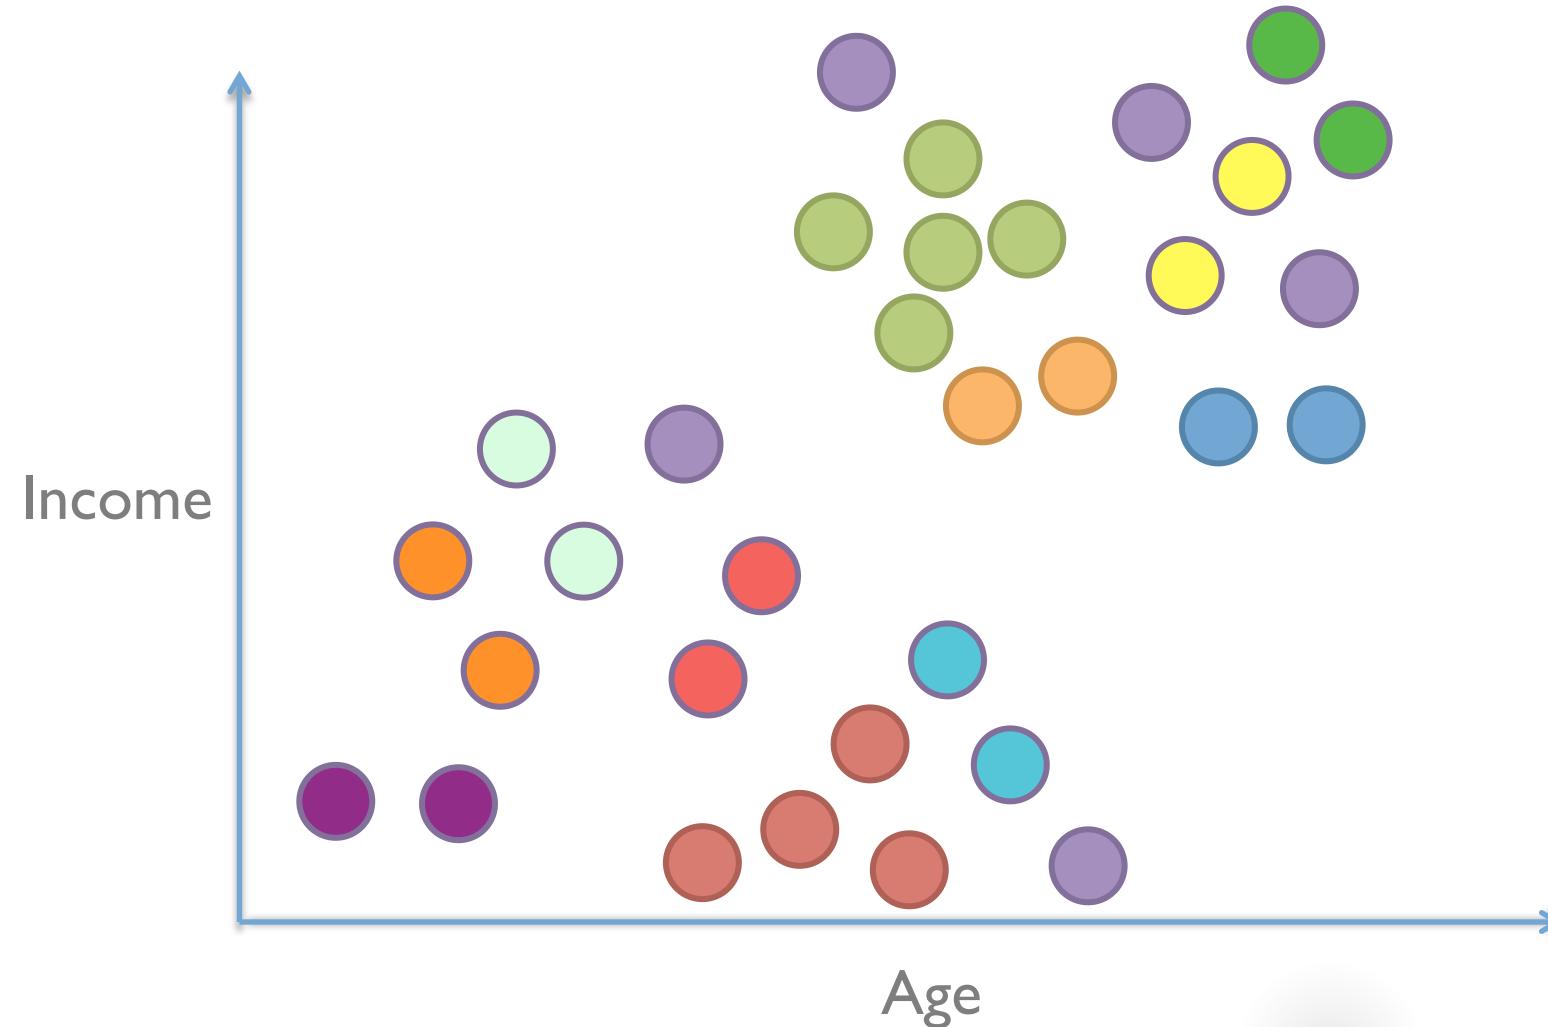
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



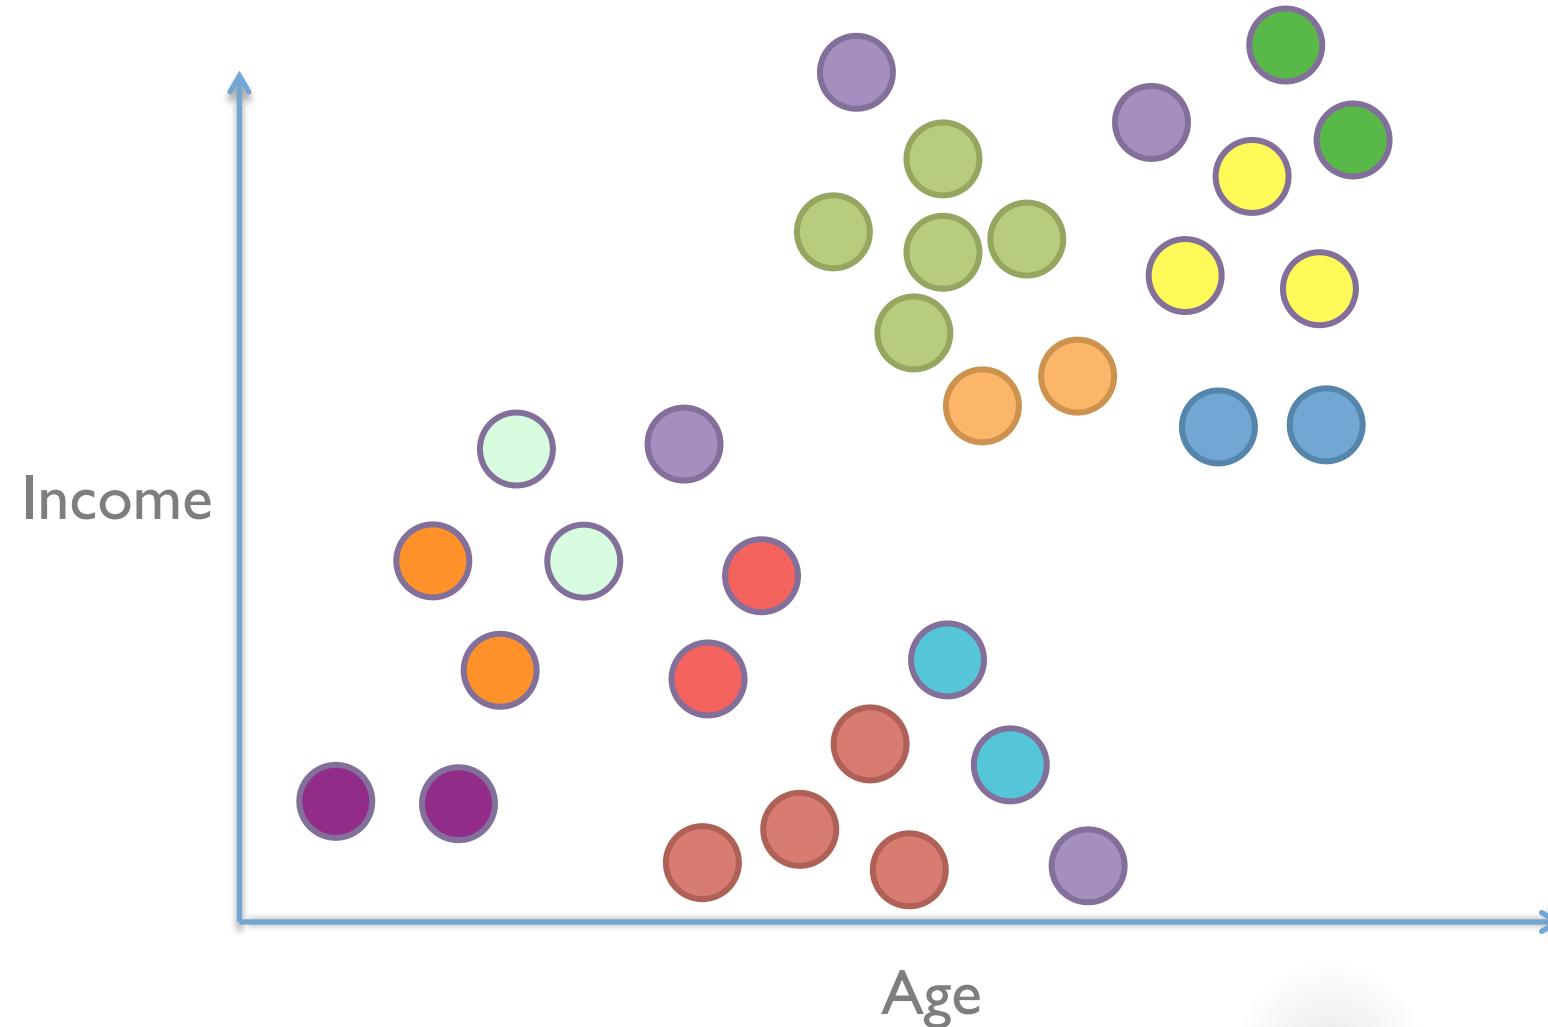
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



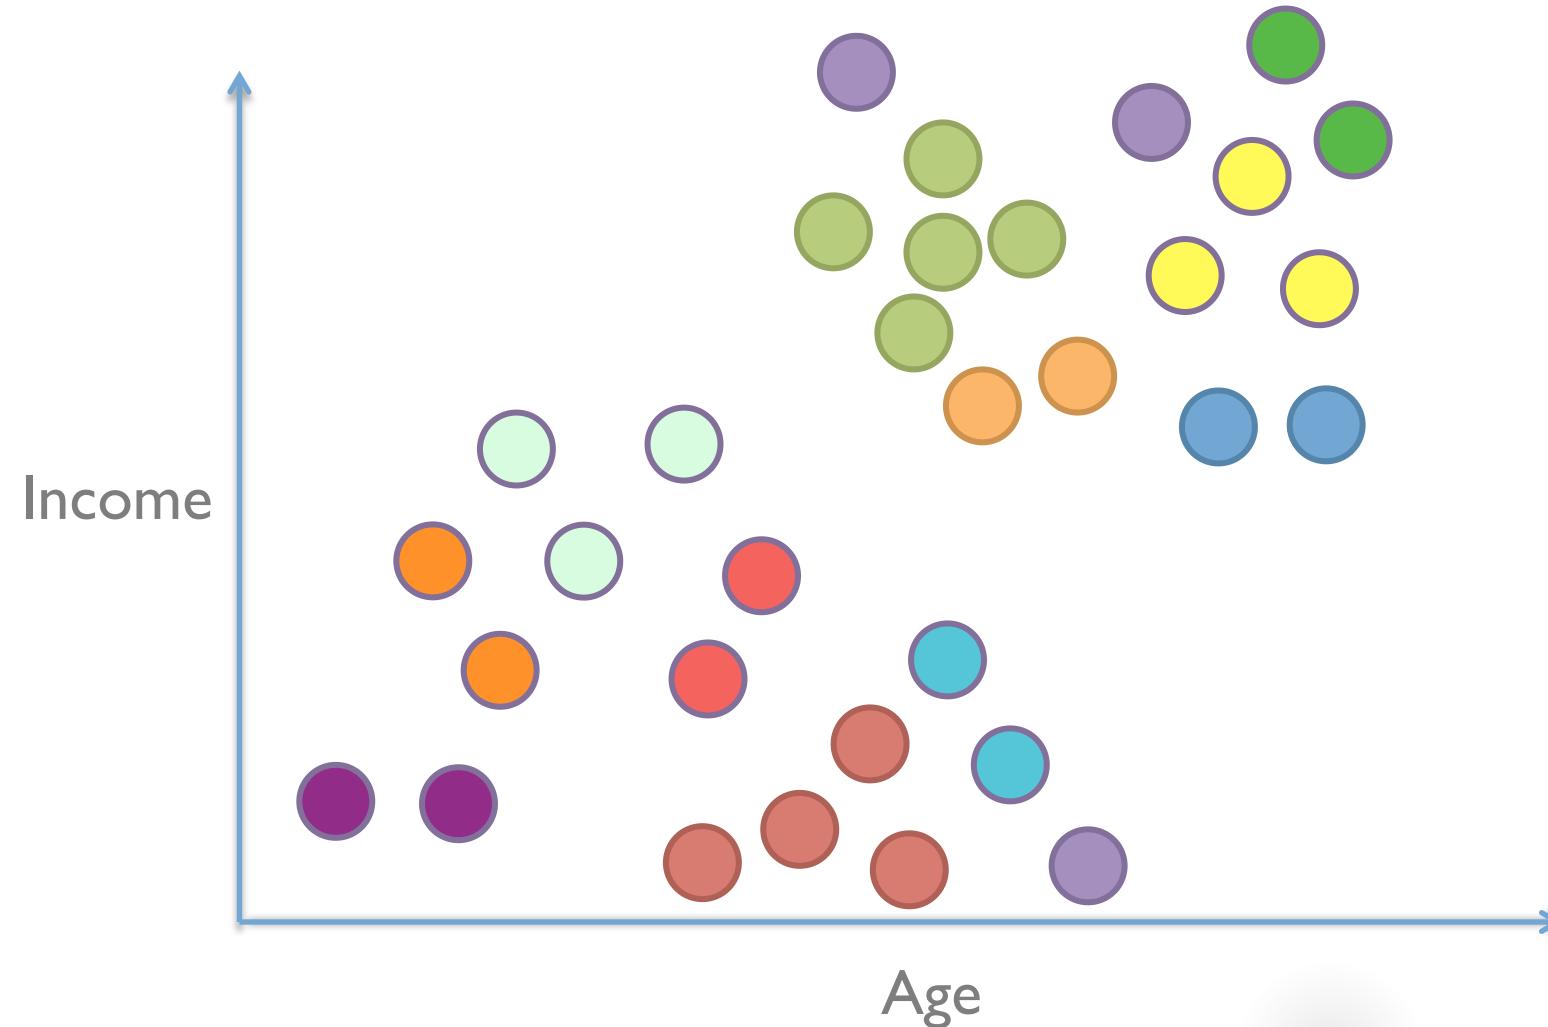
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



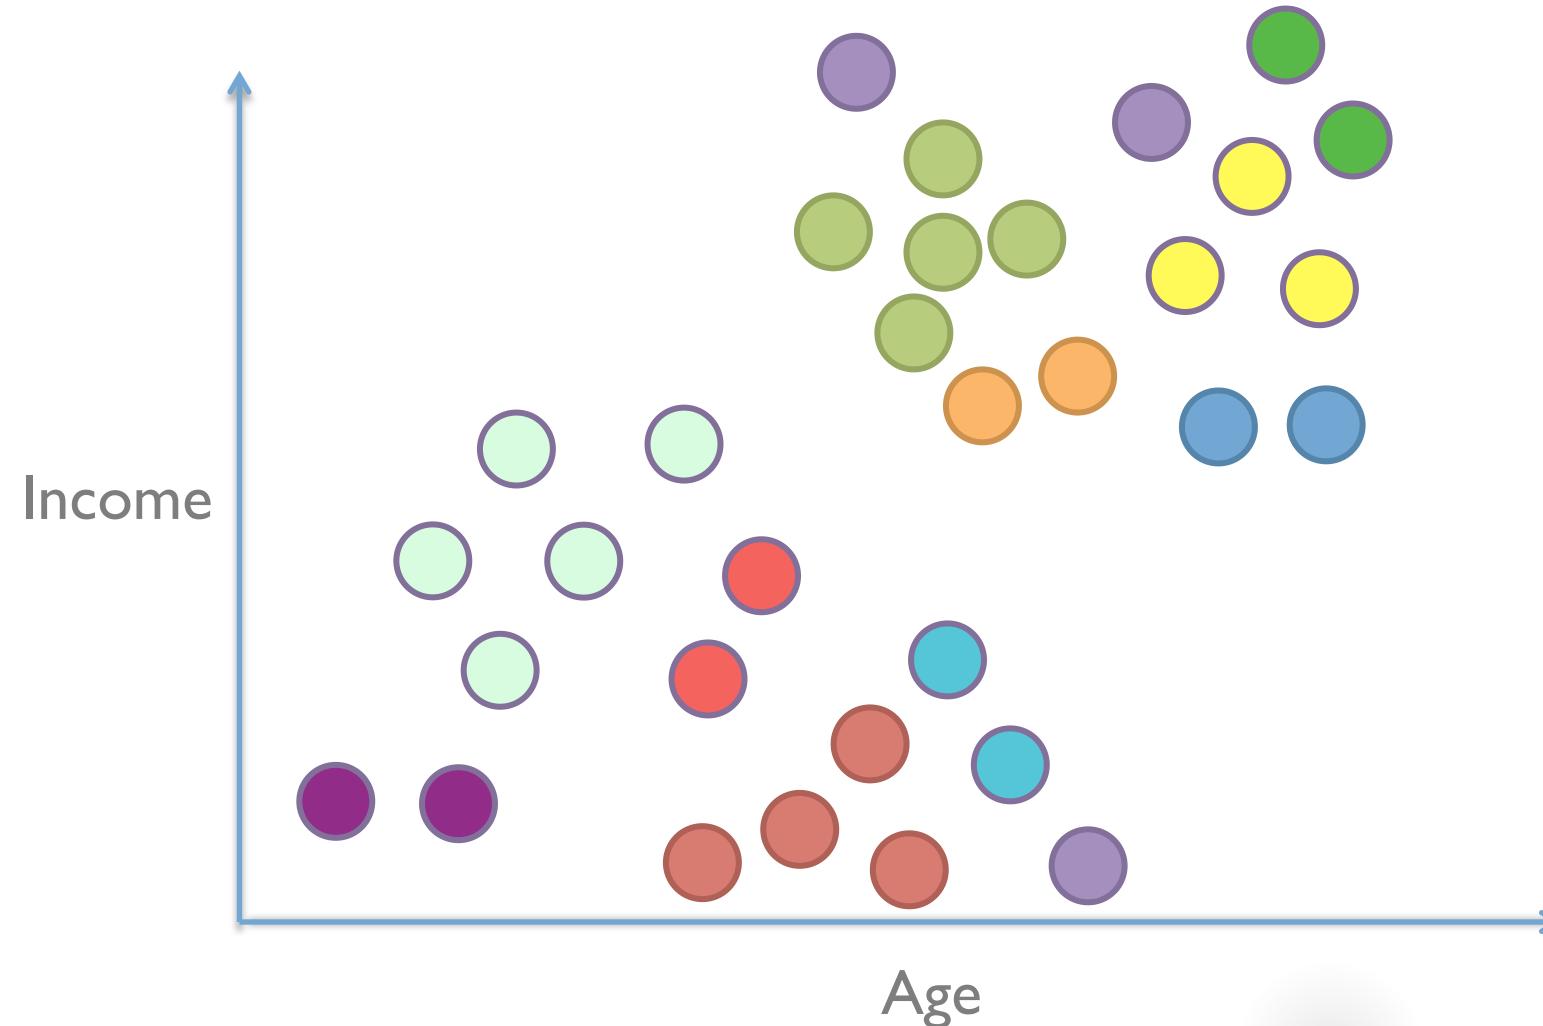
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



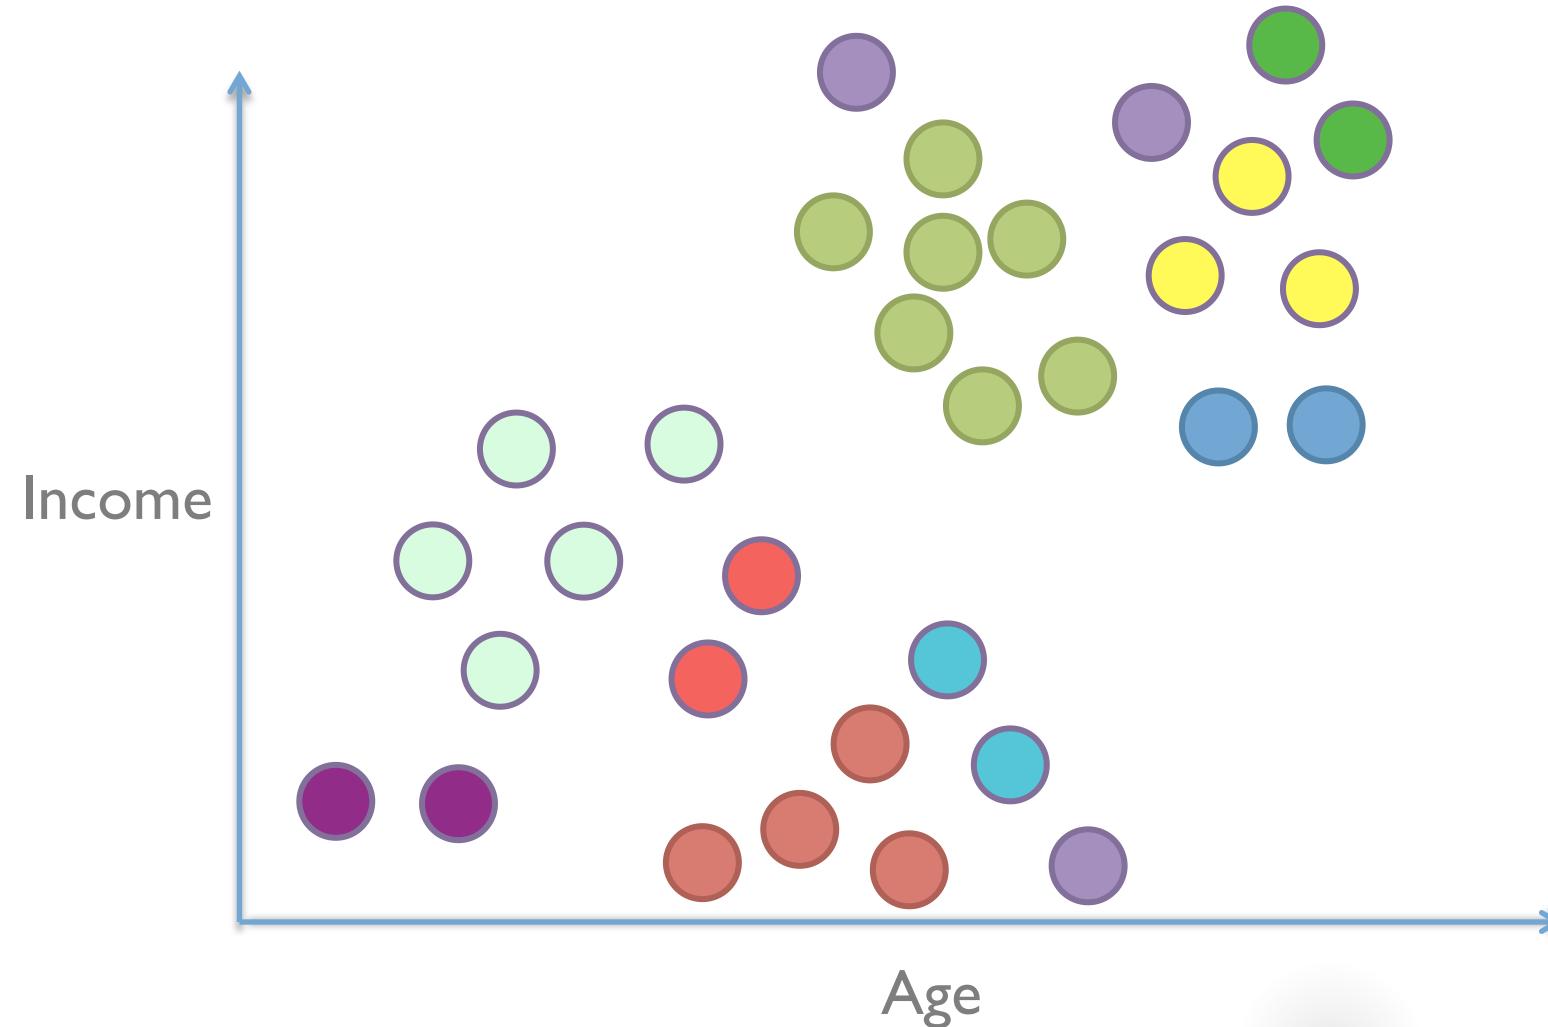
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



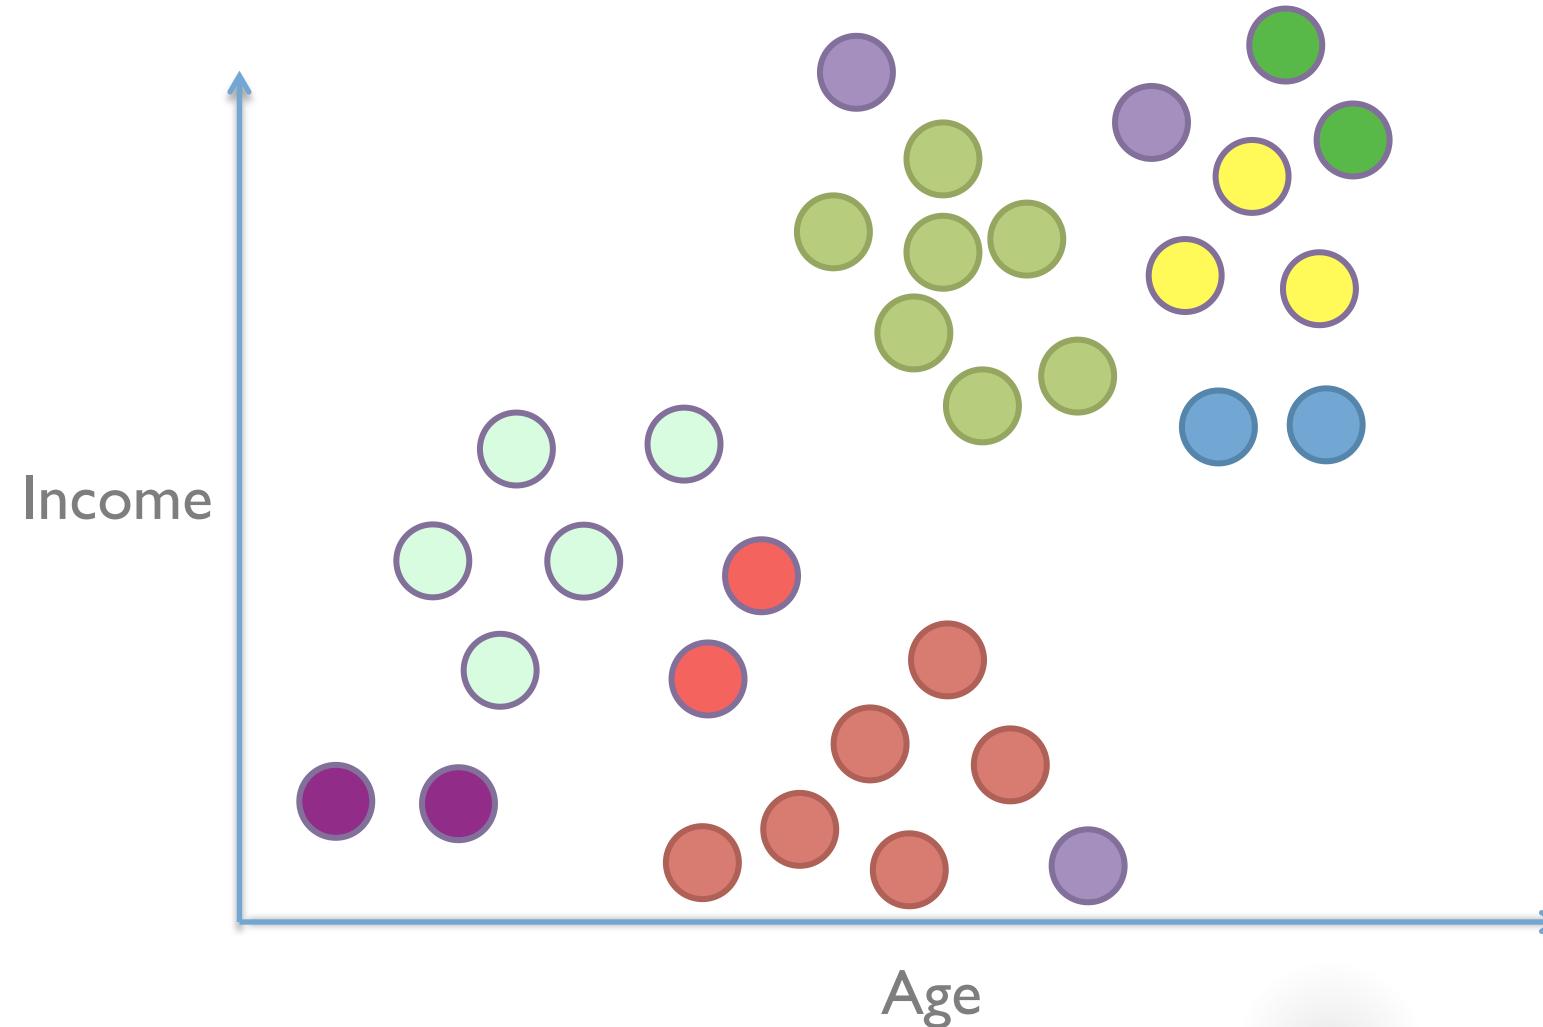
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



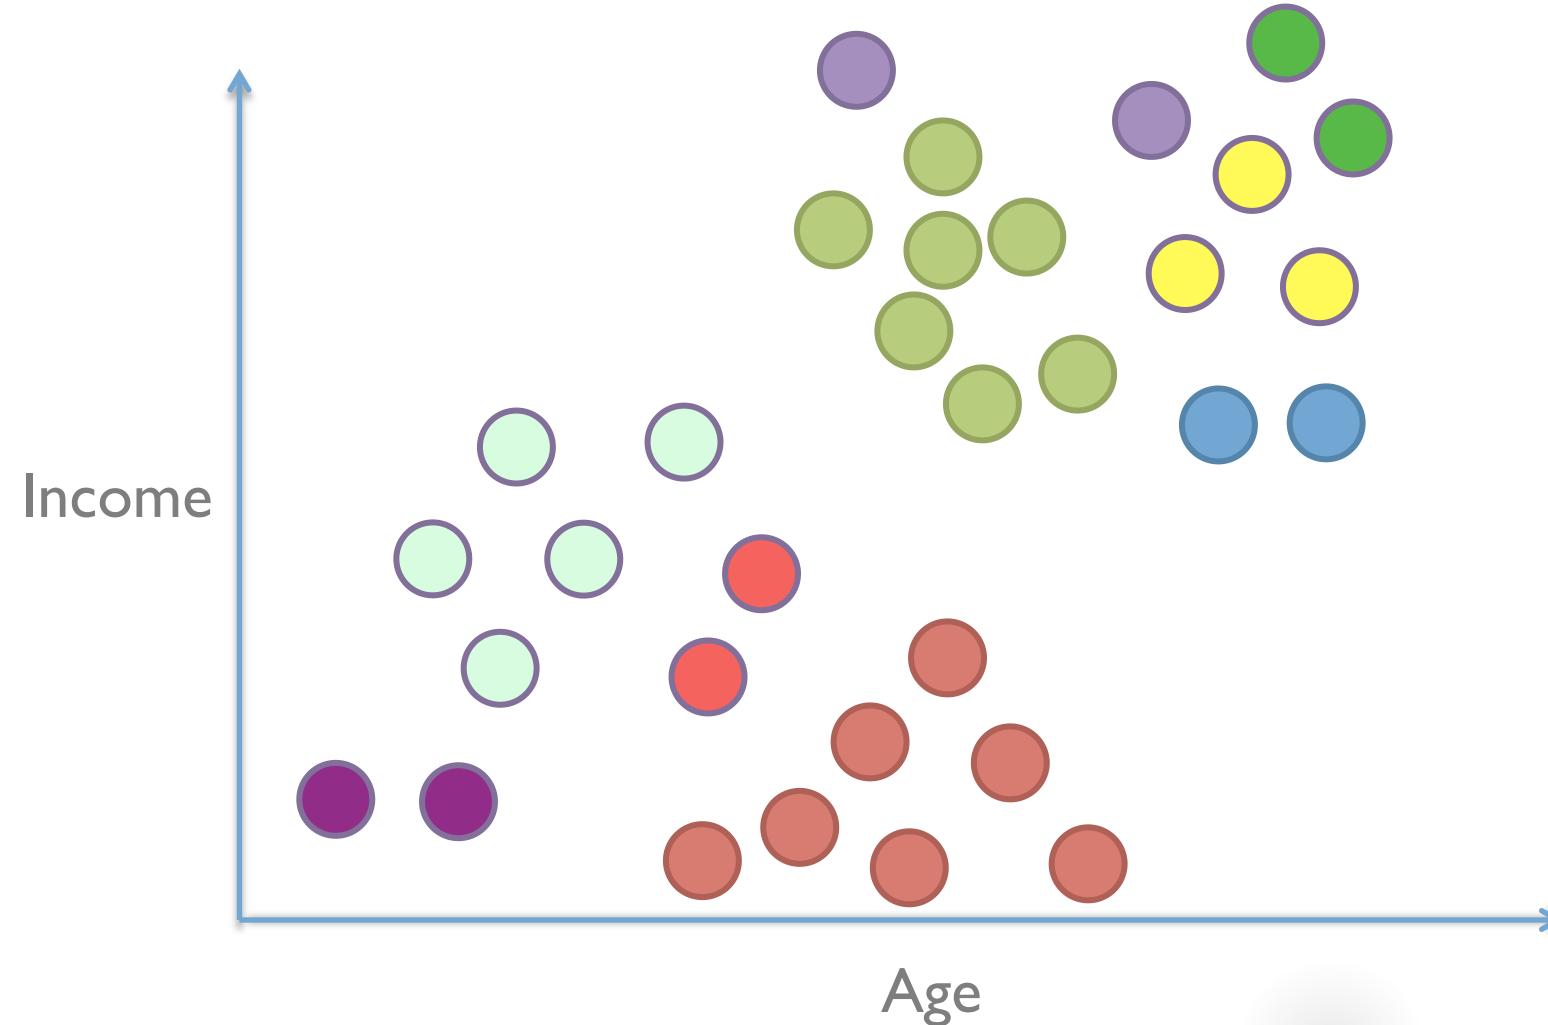
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



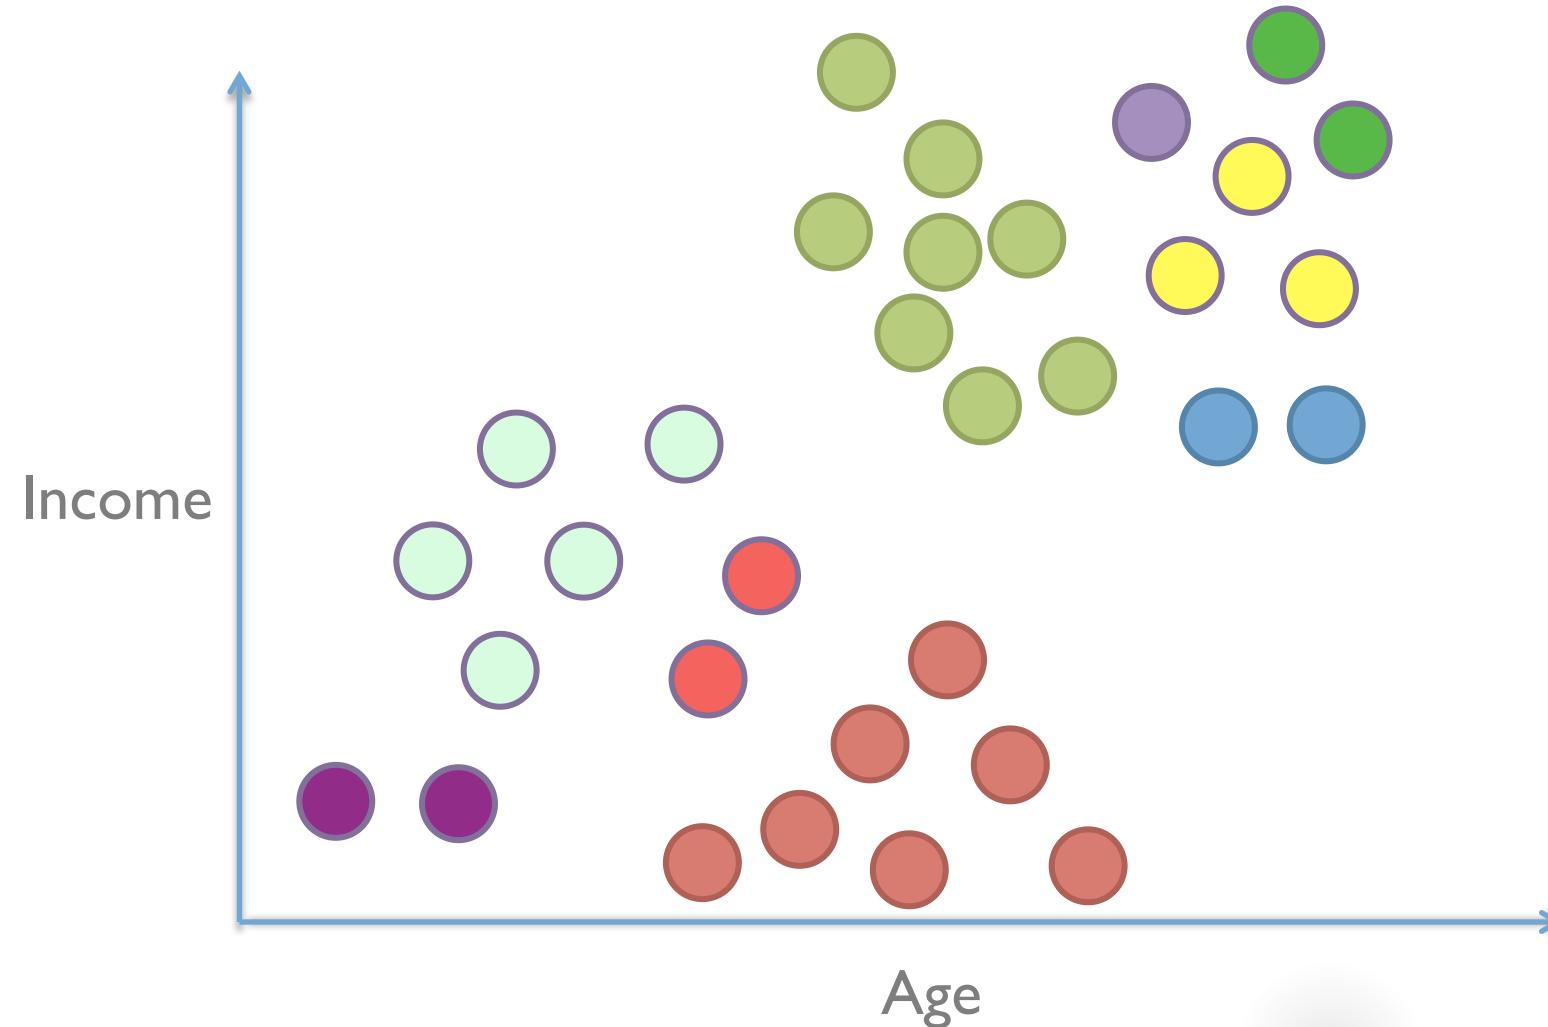
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



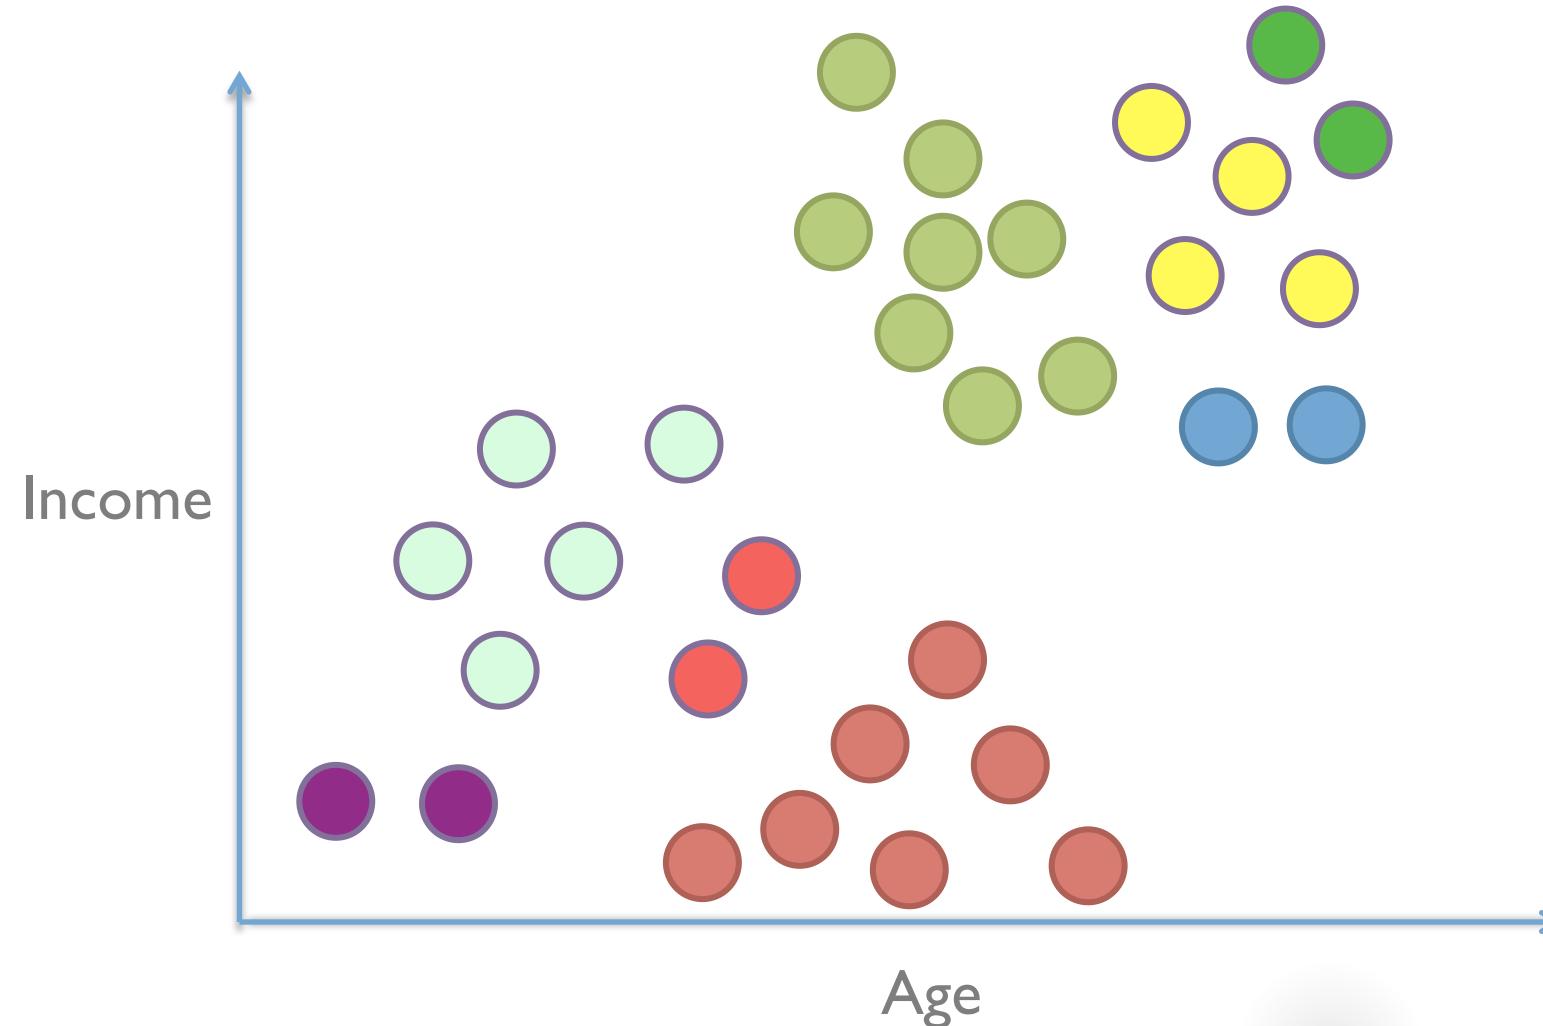
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



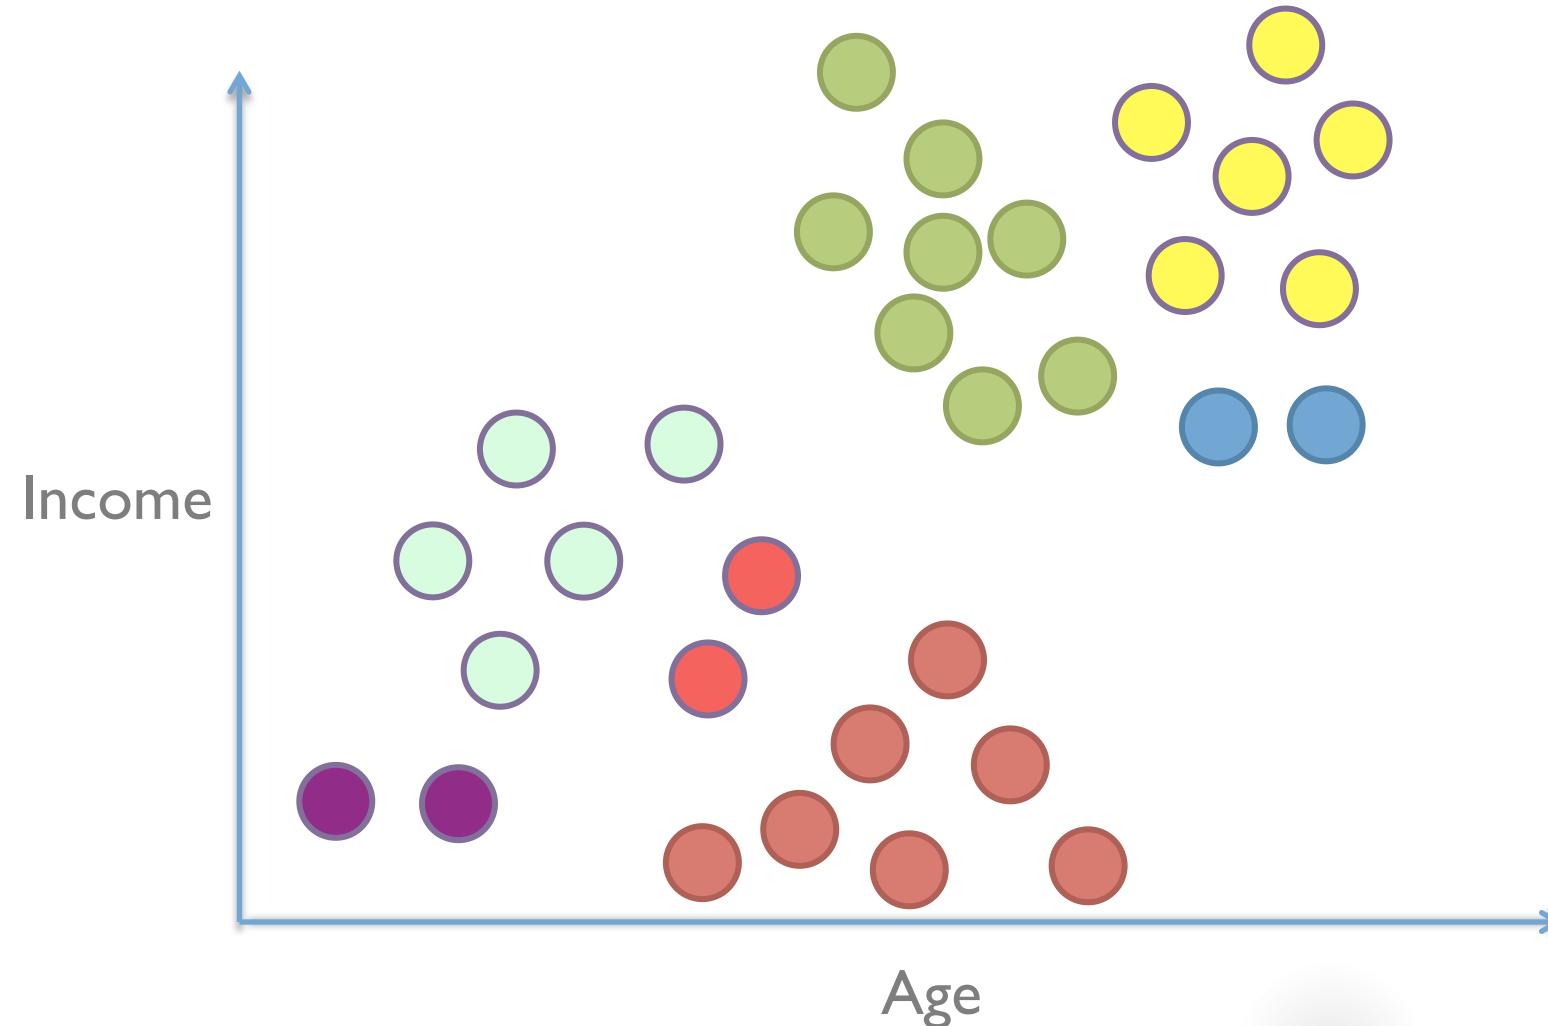
# Agglomerative Hierarchical Clustering

If closest pair are clusters, same: merge



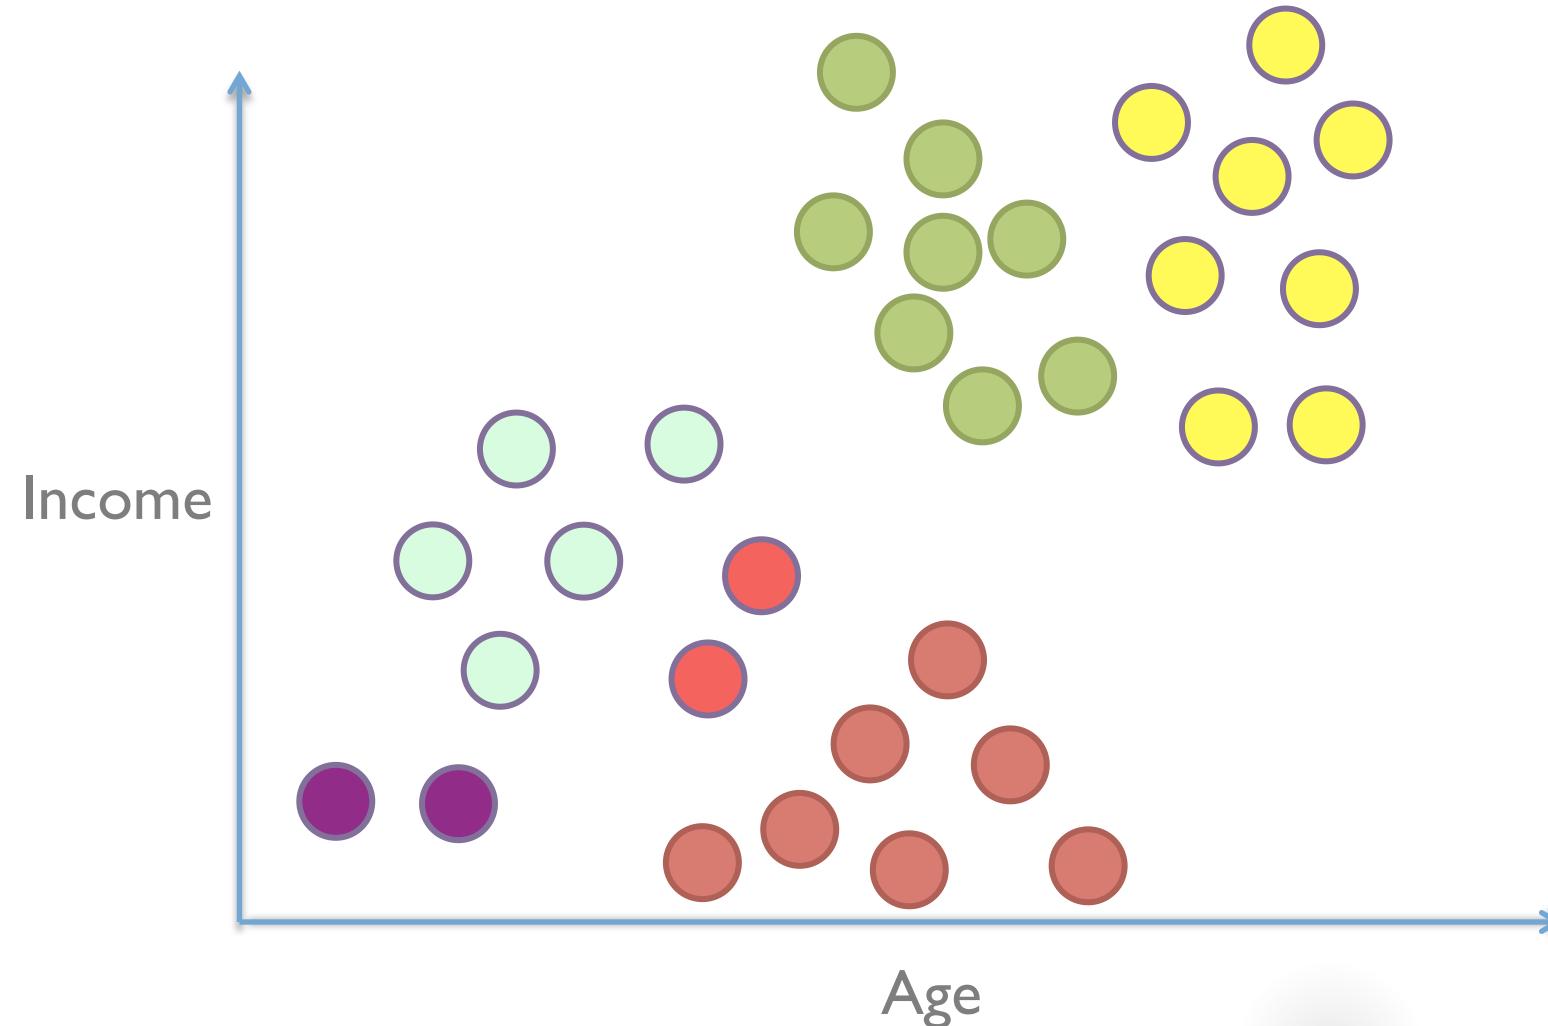
# Agglomerative Hierarchical Clustering

Current num\_clusters = 7



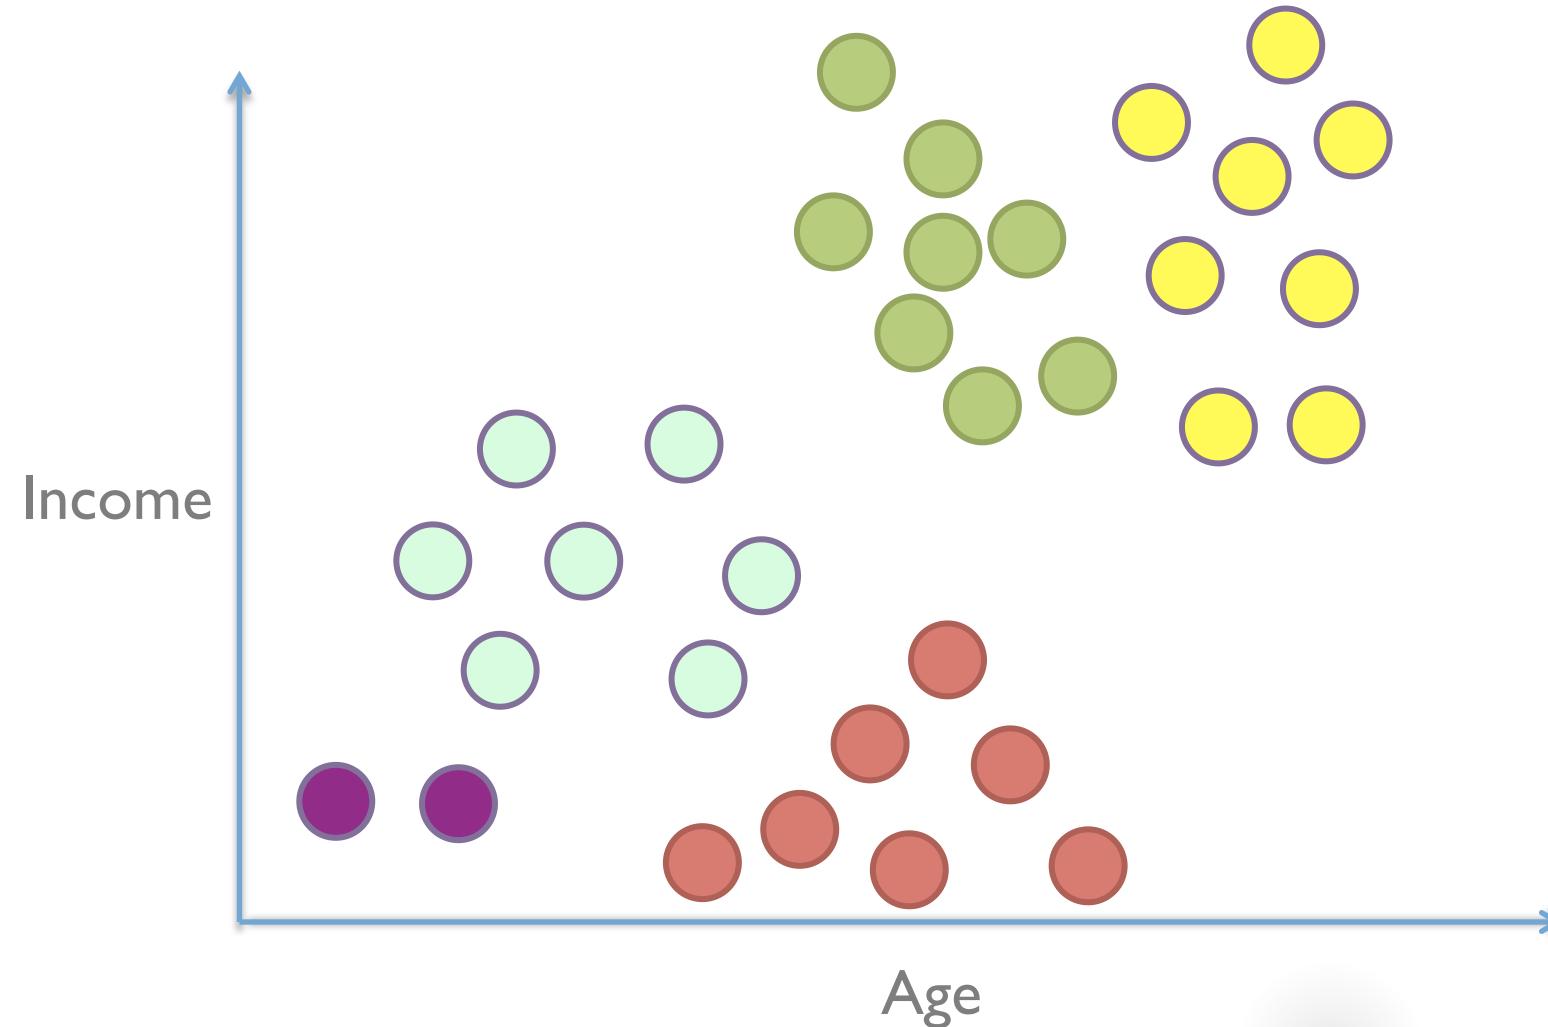
# Agglomerative Hierarchical Clustering

Current num\_clusters = 6



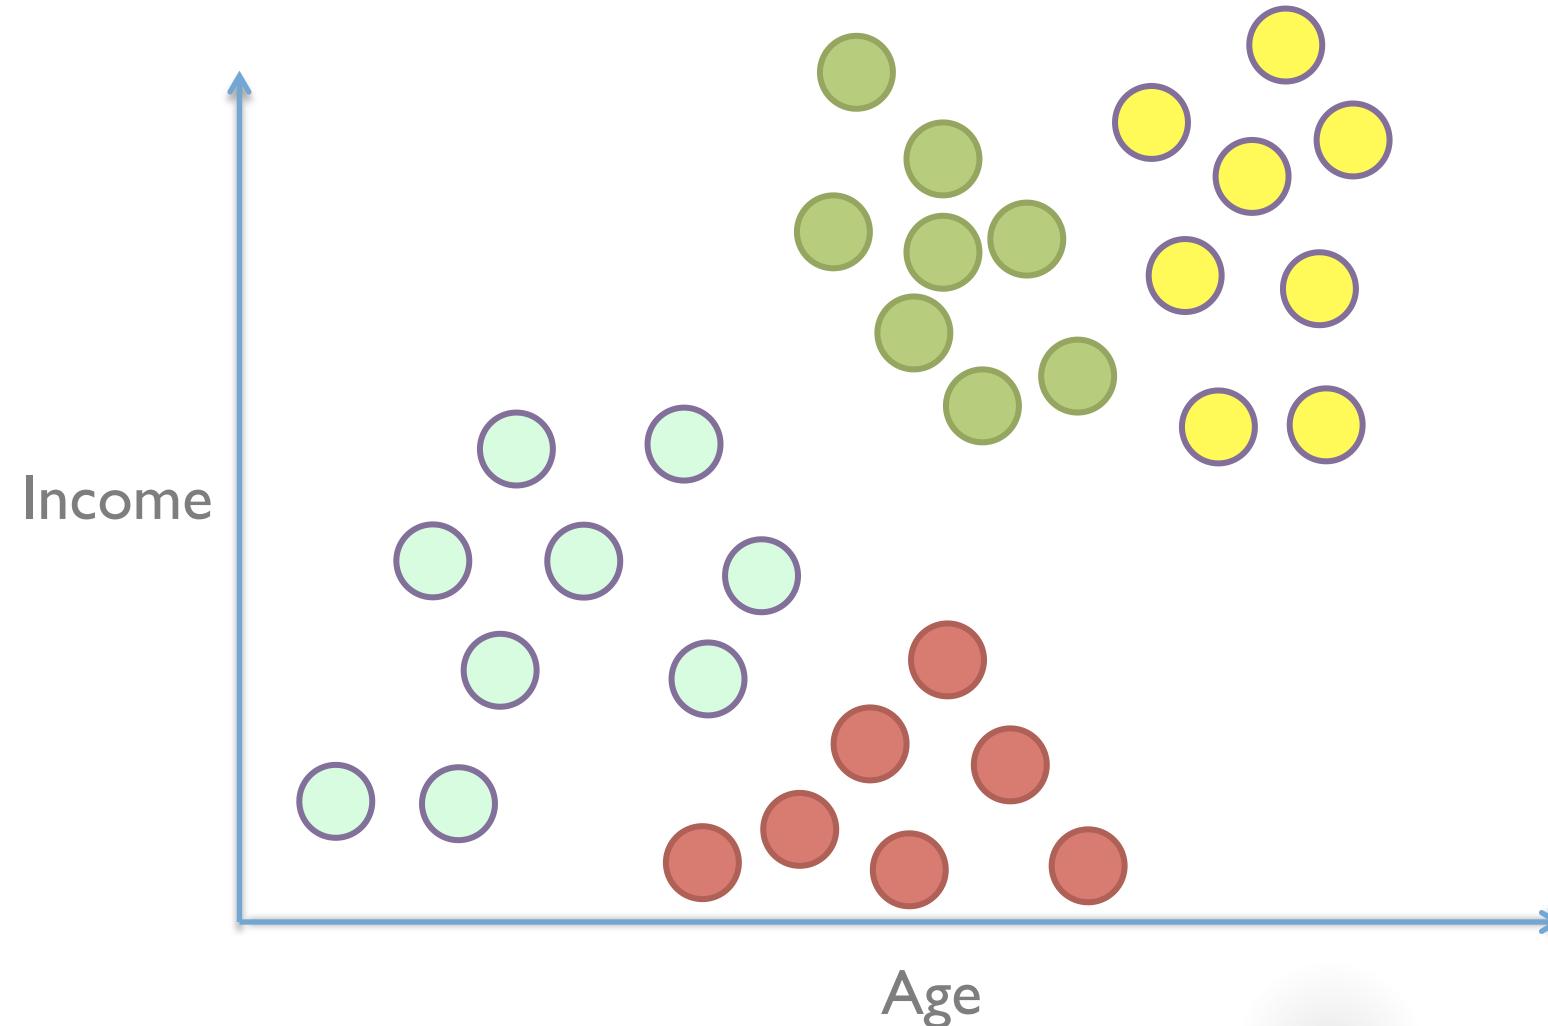
# Agglomerative Hierarchical Clustering

Current num\_clusters = 5



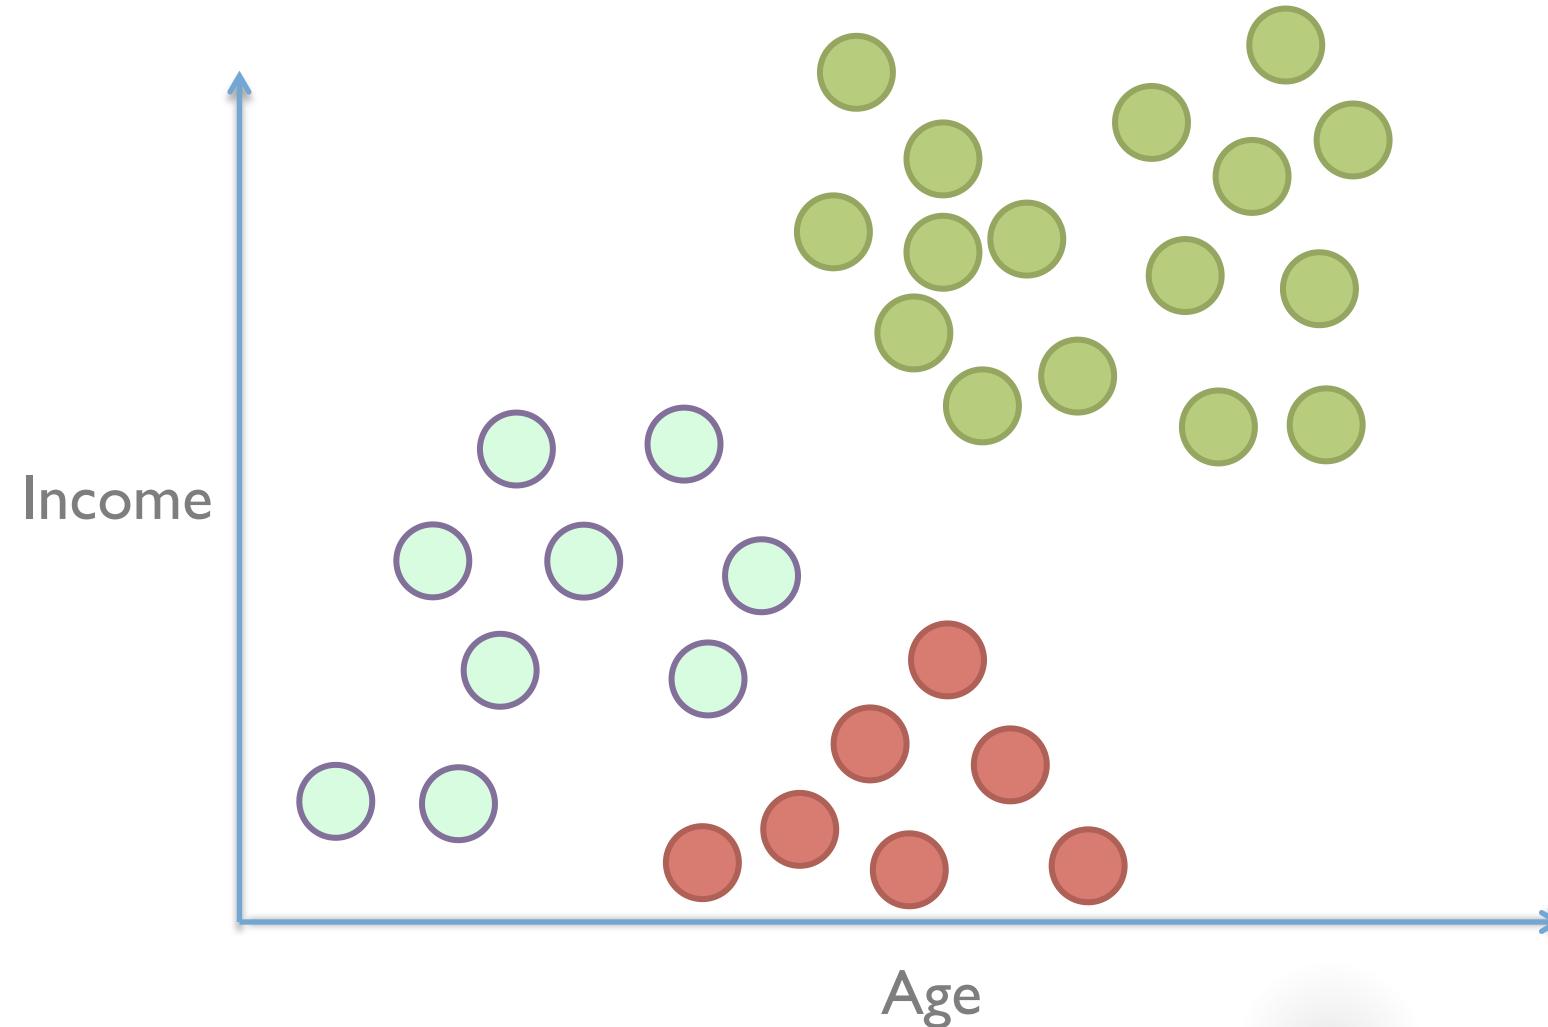
# Agglomerative Hierarchical Clustering

Current num\_clusters = 4



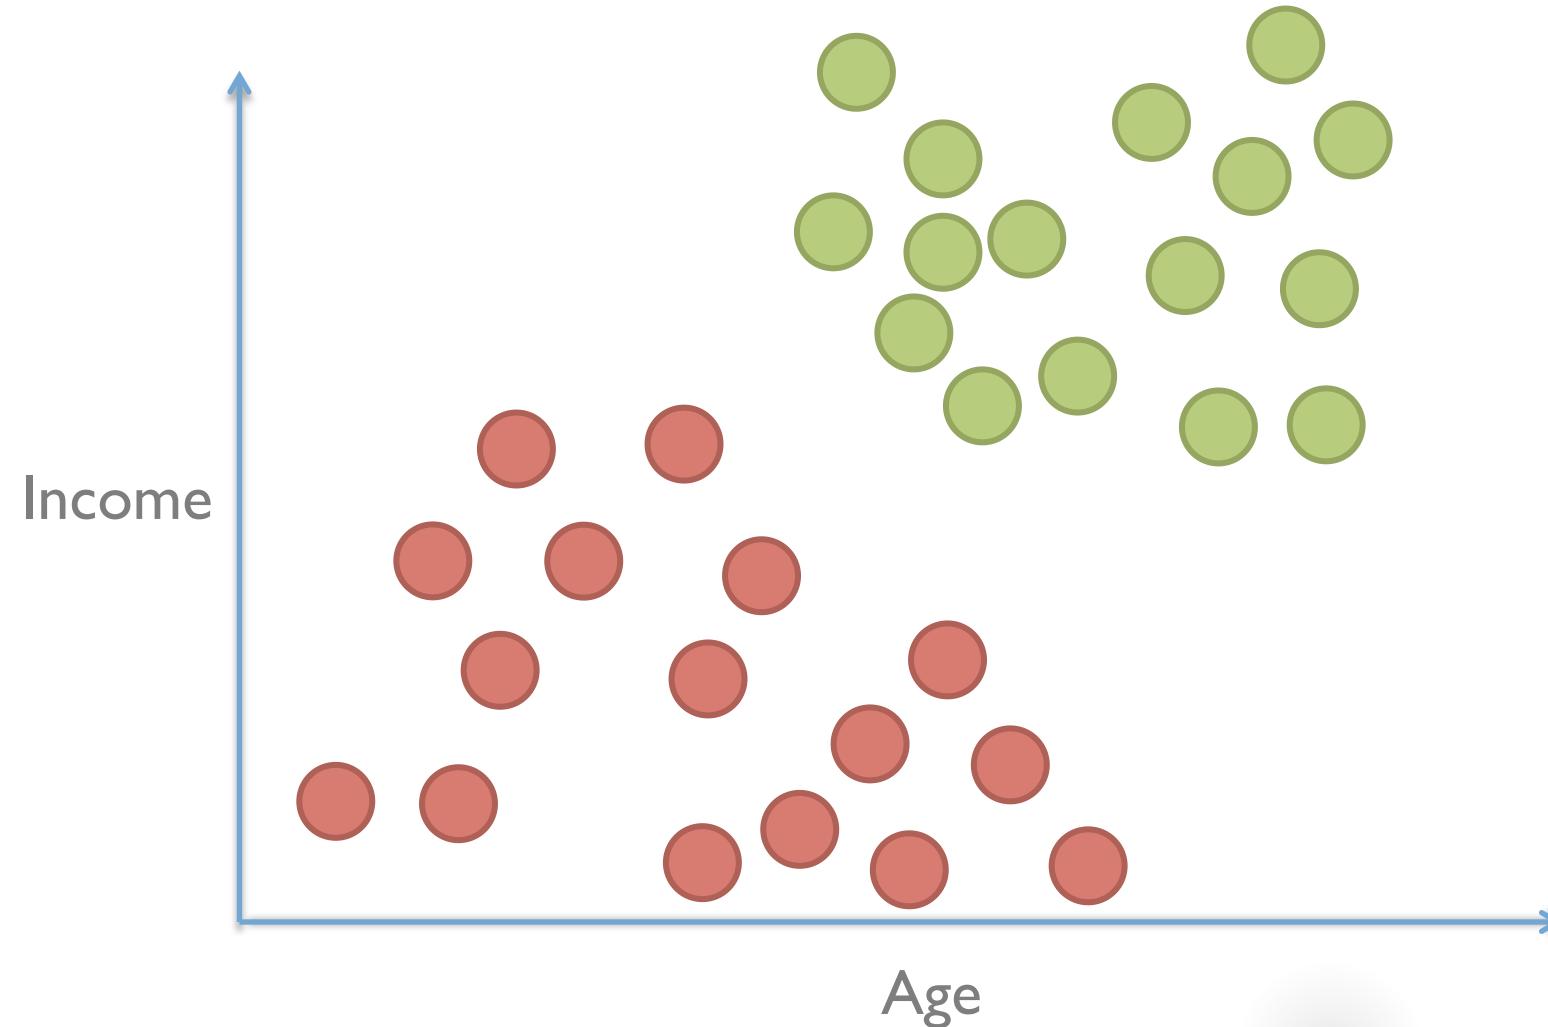
# Agglomerative Hierarchical Clustering

Current num\_clusters = 3



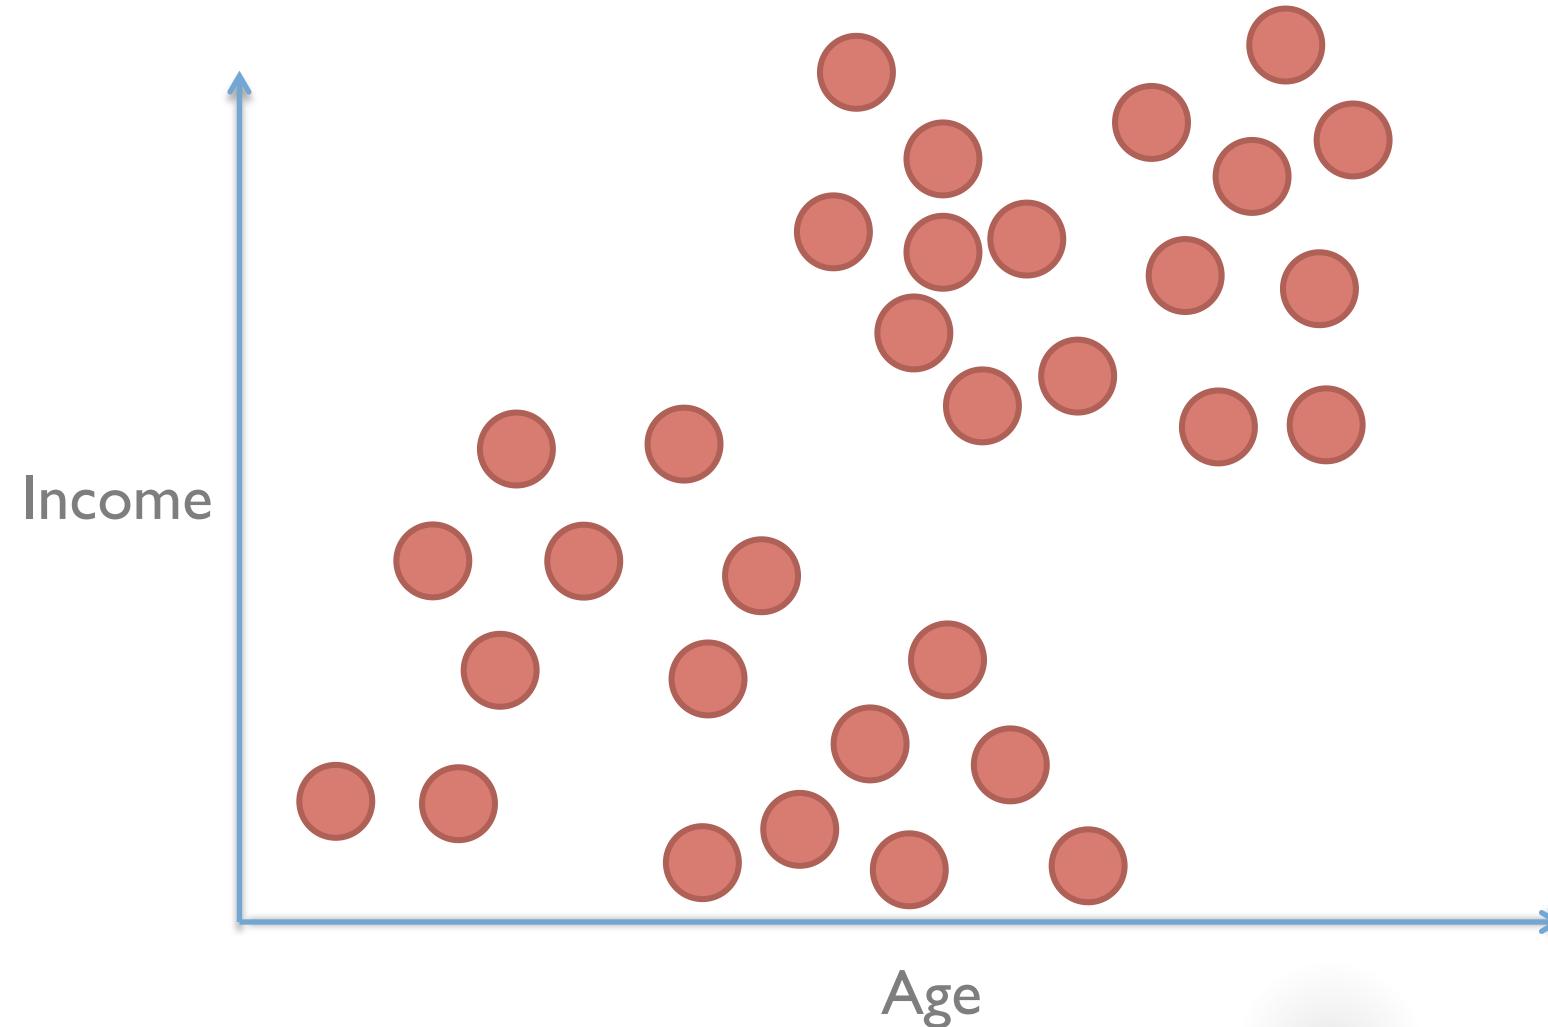
# Agglomerative Hierarchical Clustering

Current num\_clusters = 2



# Agglomerative Hierarchical Clustering

Current num\_clusters = 1



# Agglomerative Hierarchical Clustering

Stop either

- when you reach the correct **num\_clusters** (next step would mean fewer clusters than desired)

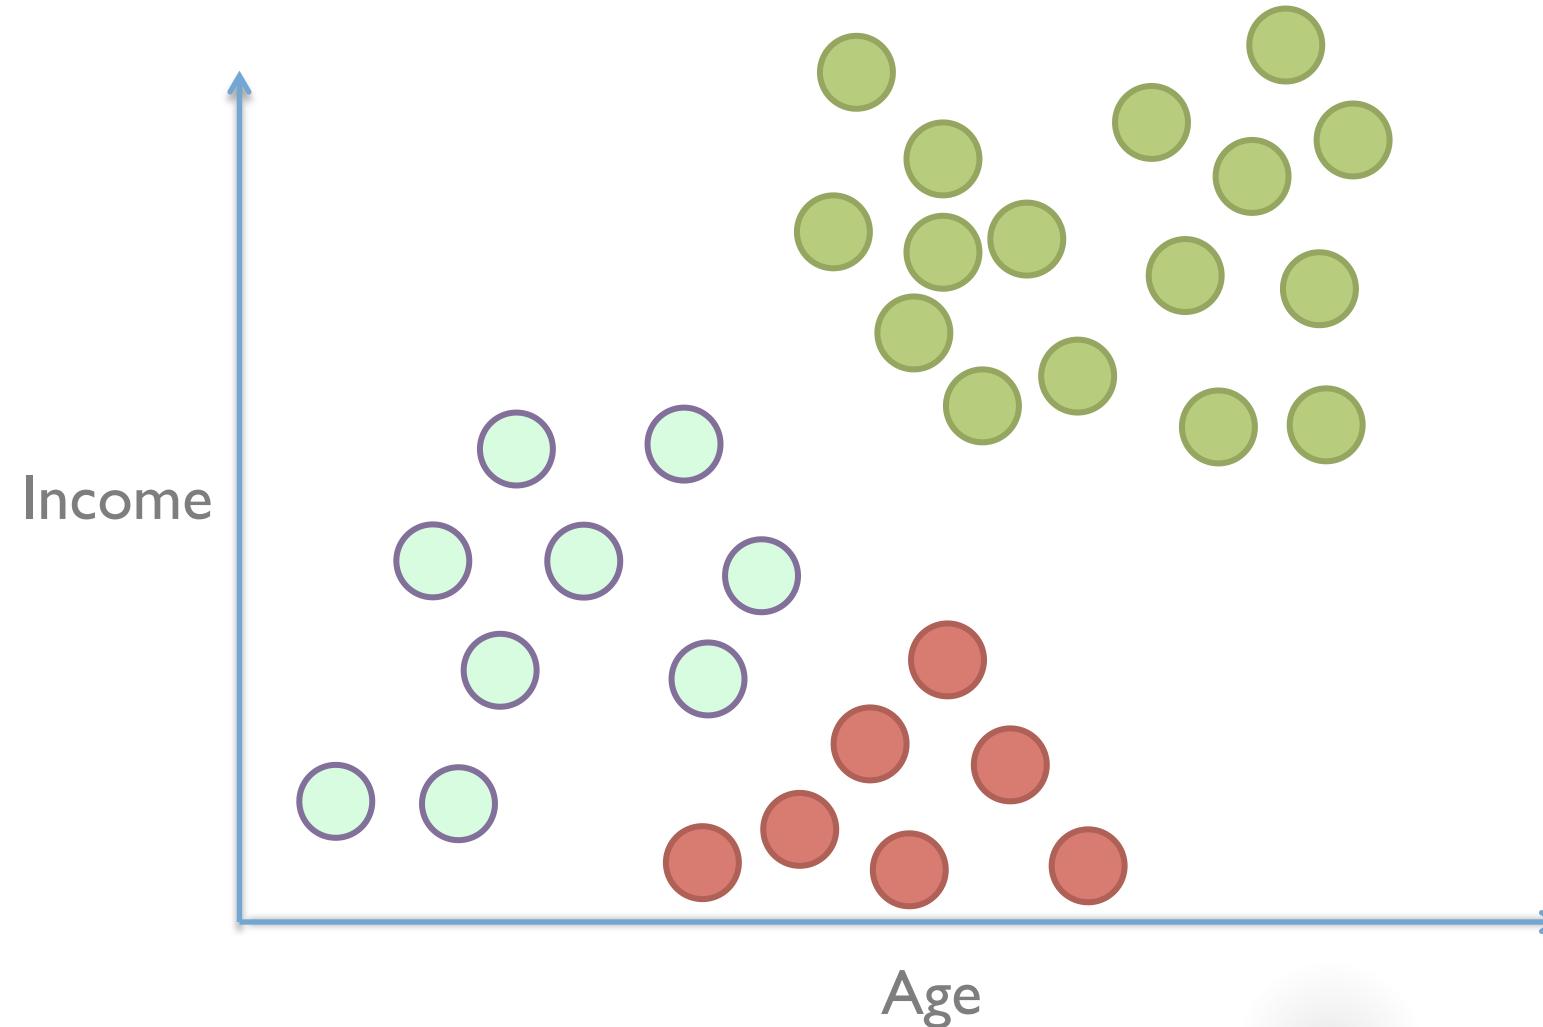
Or

- when **min cluster distance** reaches a set value (next step would mean merging along too large a distance)



# Agglomerative Hierarchical Clustering

Current num\_clusters = 3



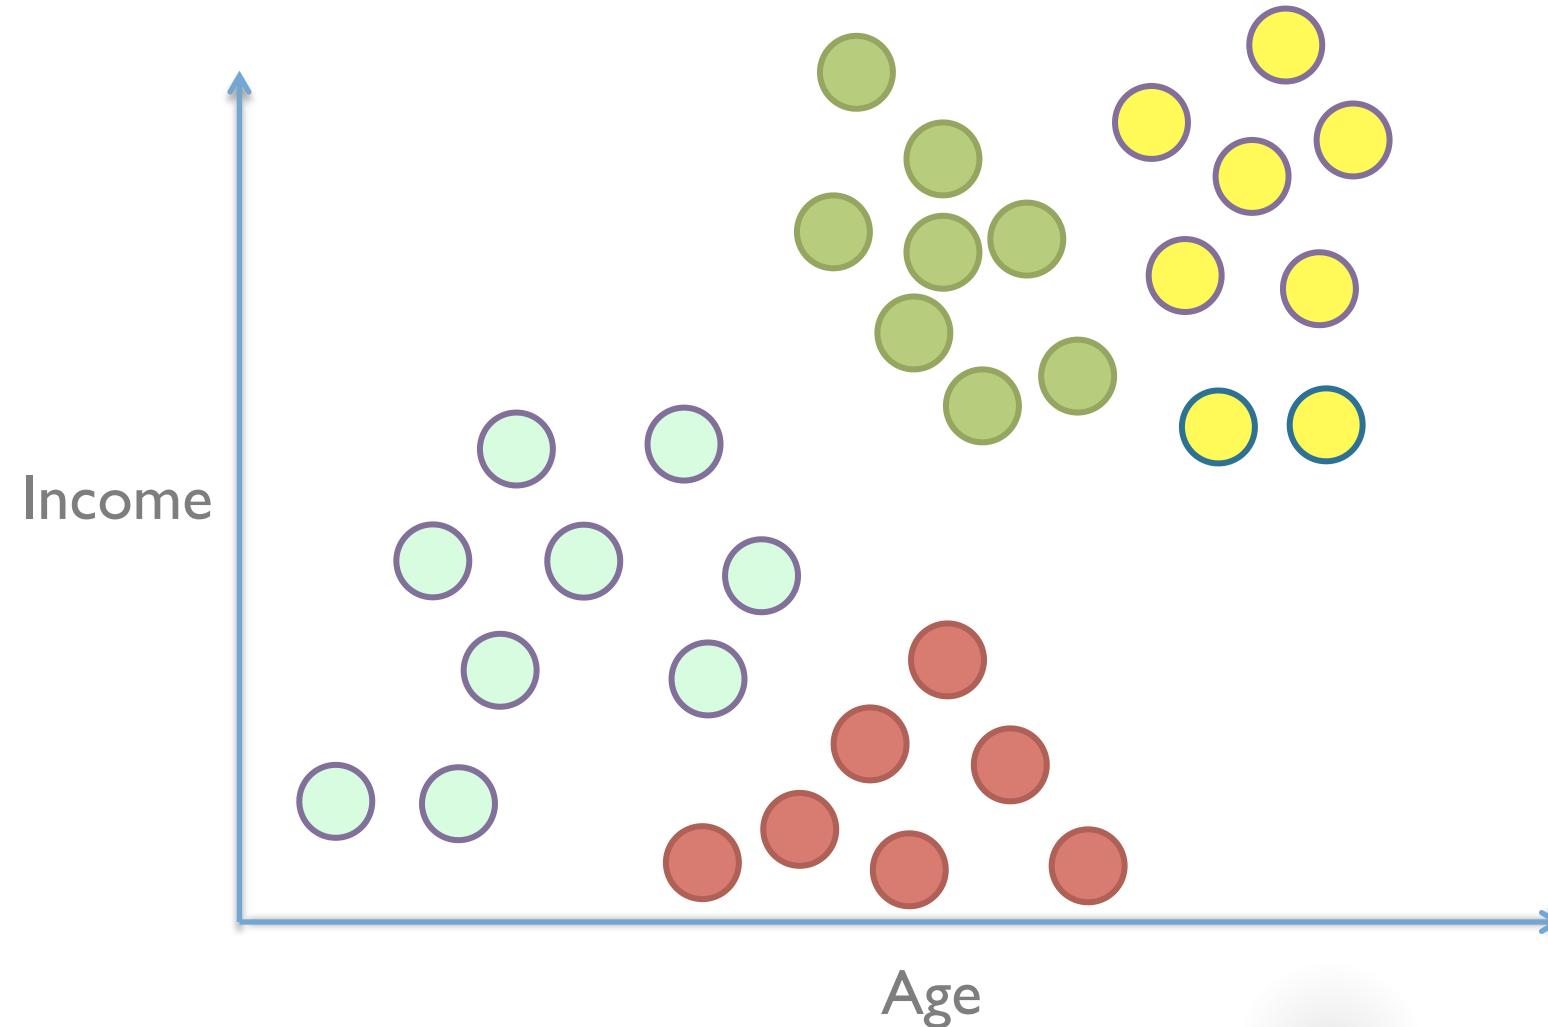
# Agglomerative Hierarchical Clustering

Current `num_clusters` = 3



# Agglomerative Hierarchical Clustering

Current num\_clusters = 4



# Agglomerative Hierarchical Clustering

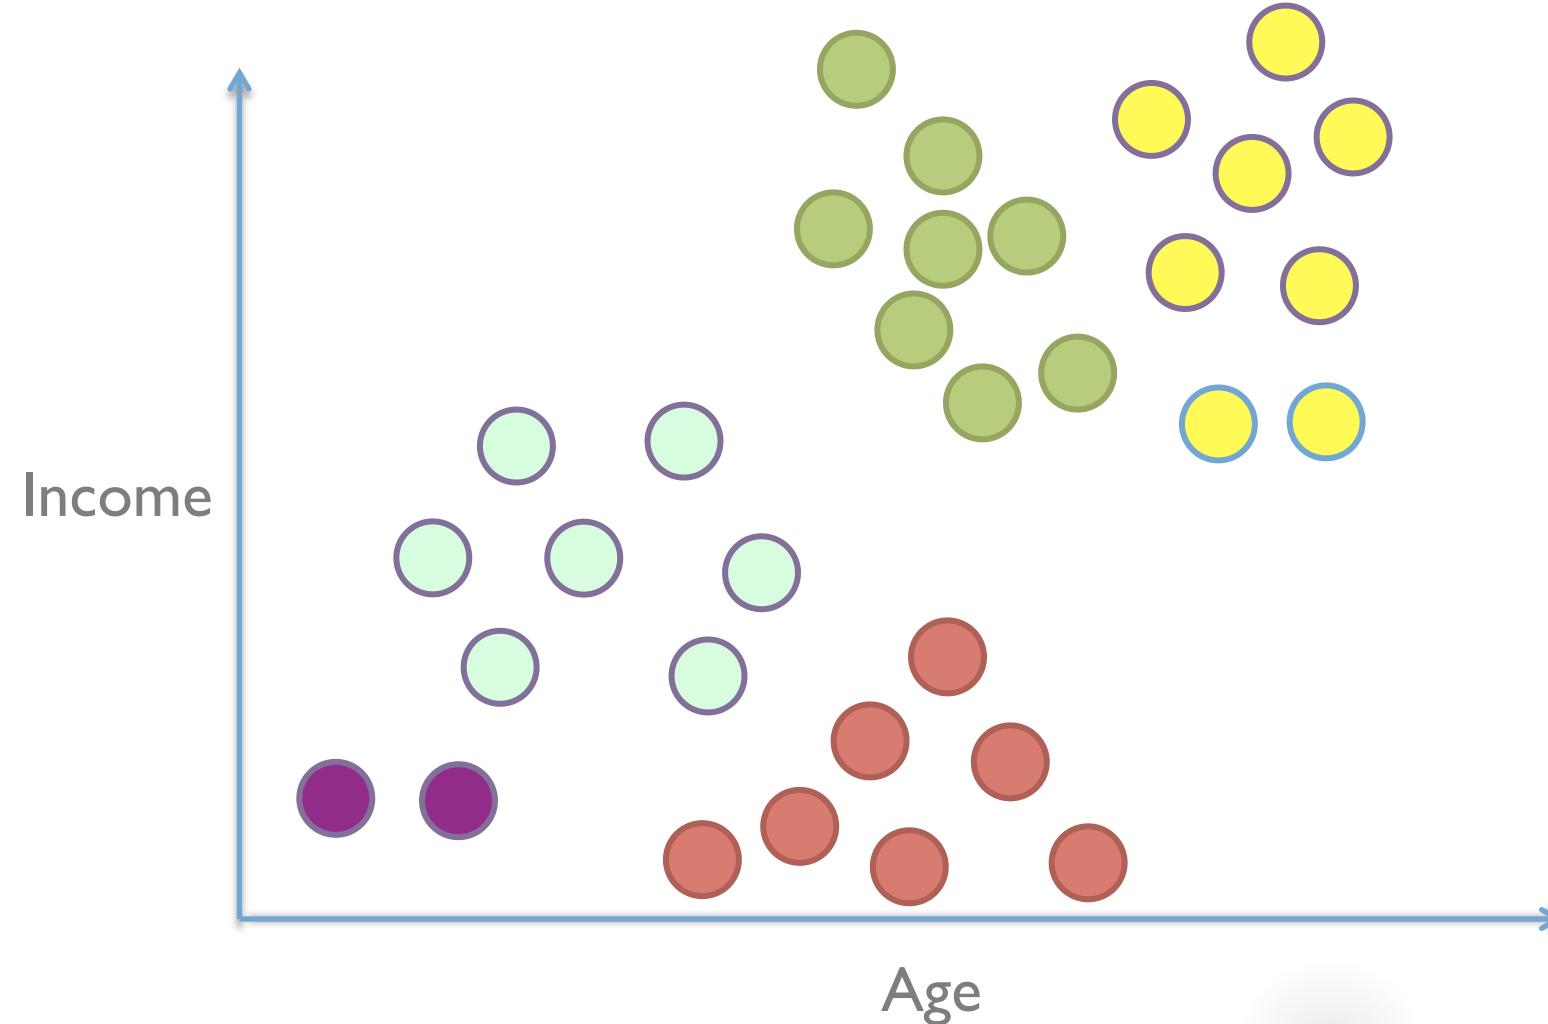
Current `num_clusters = 4`



# Agglomerative Hierarchical Clustering



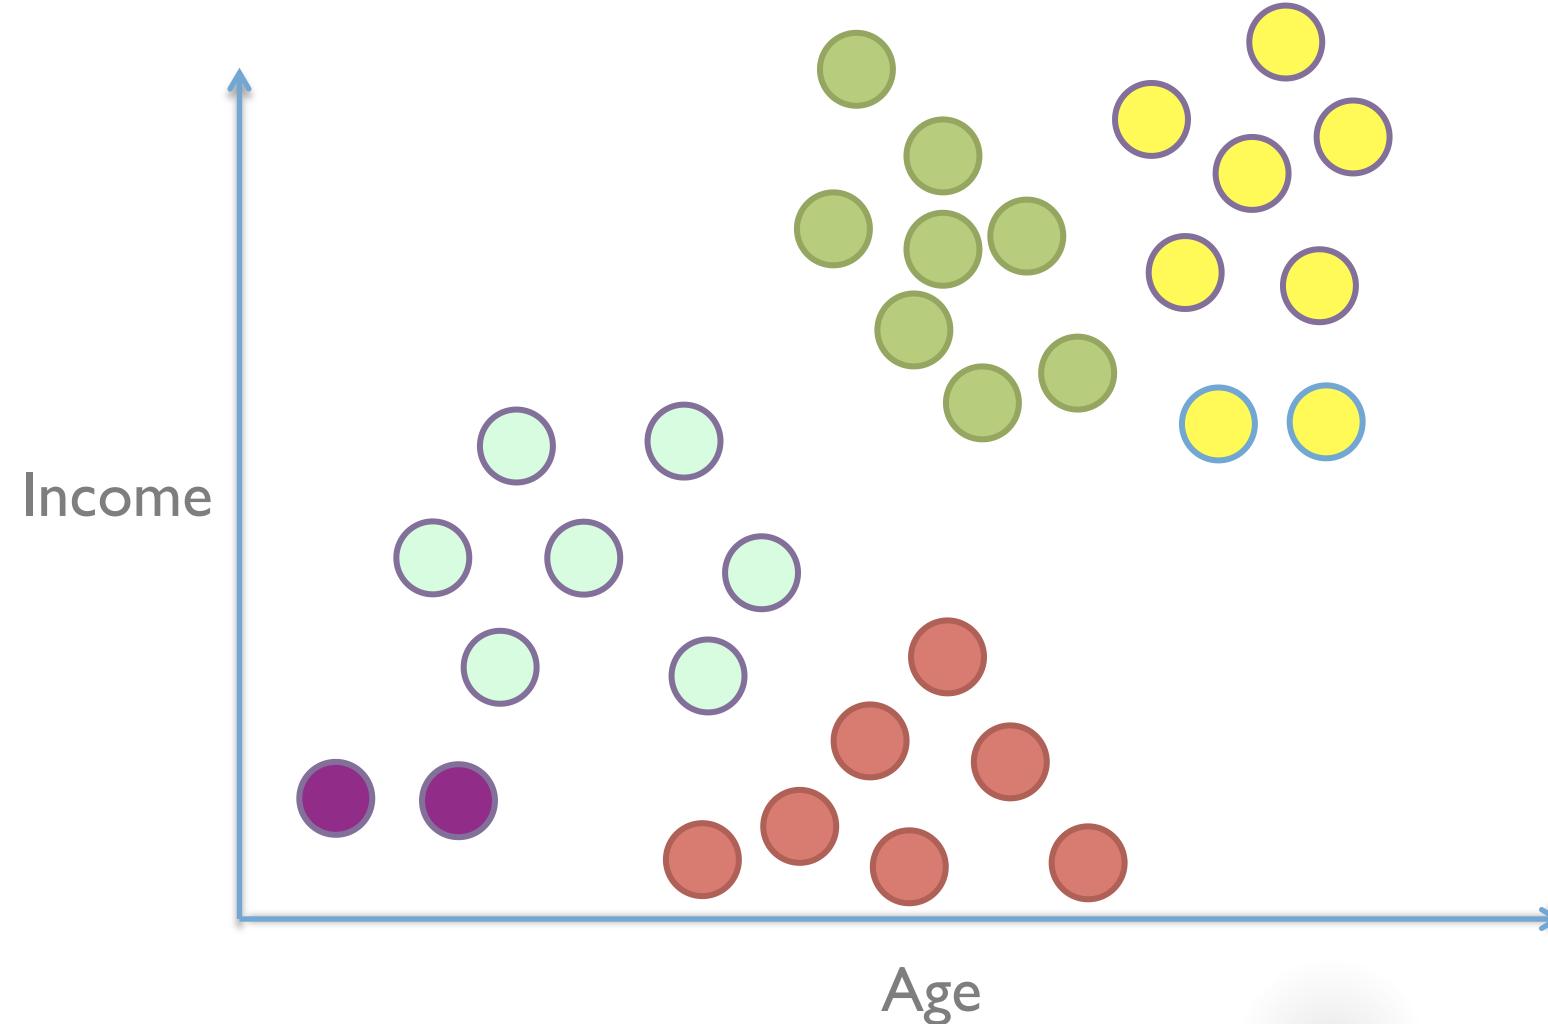
# Agglomerative Hierarchical Clustering



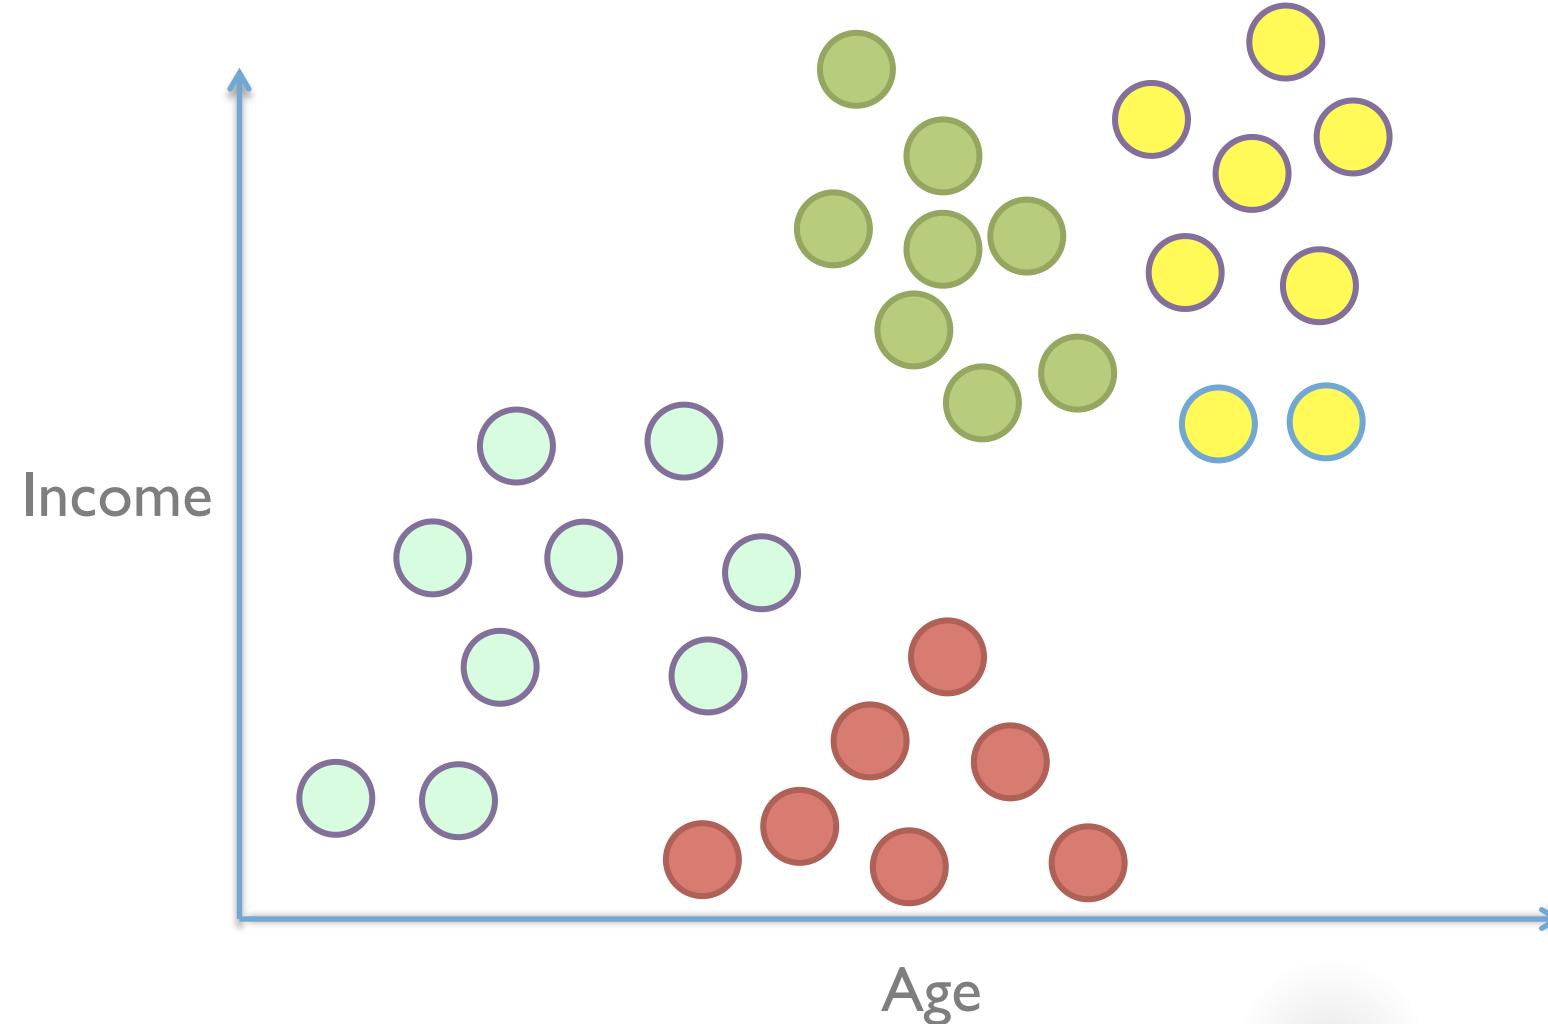
# Agglomerative Hierarchical Clustering



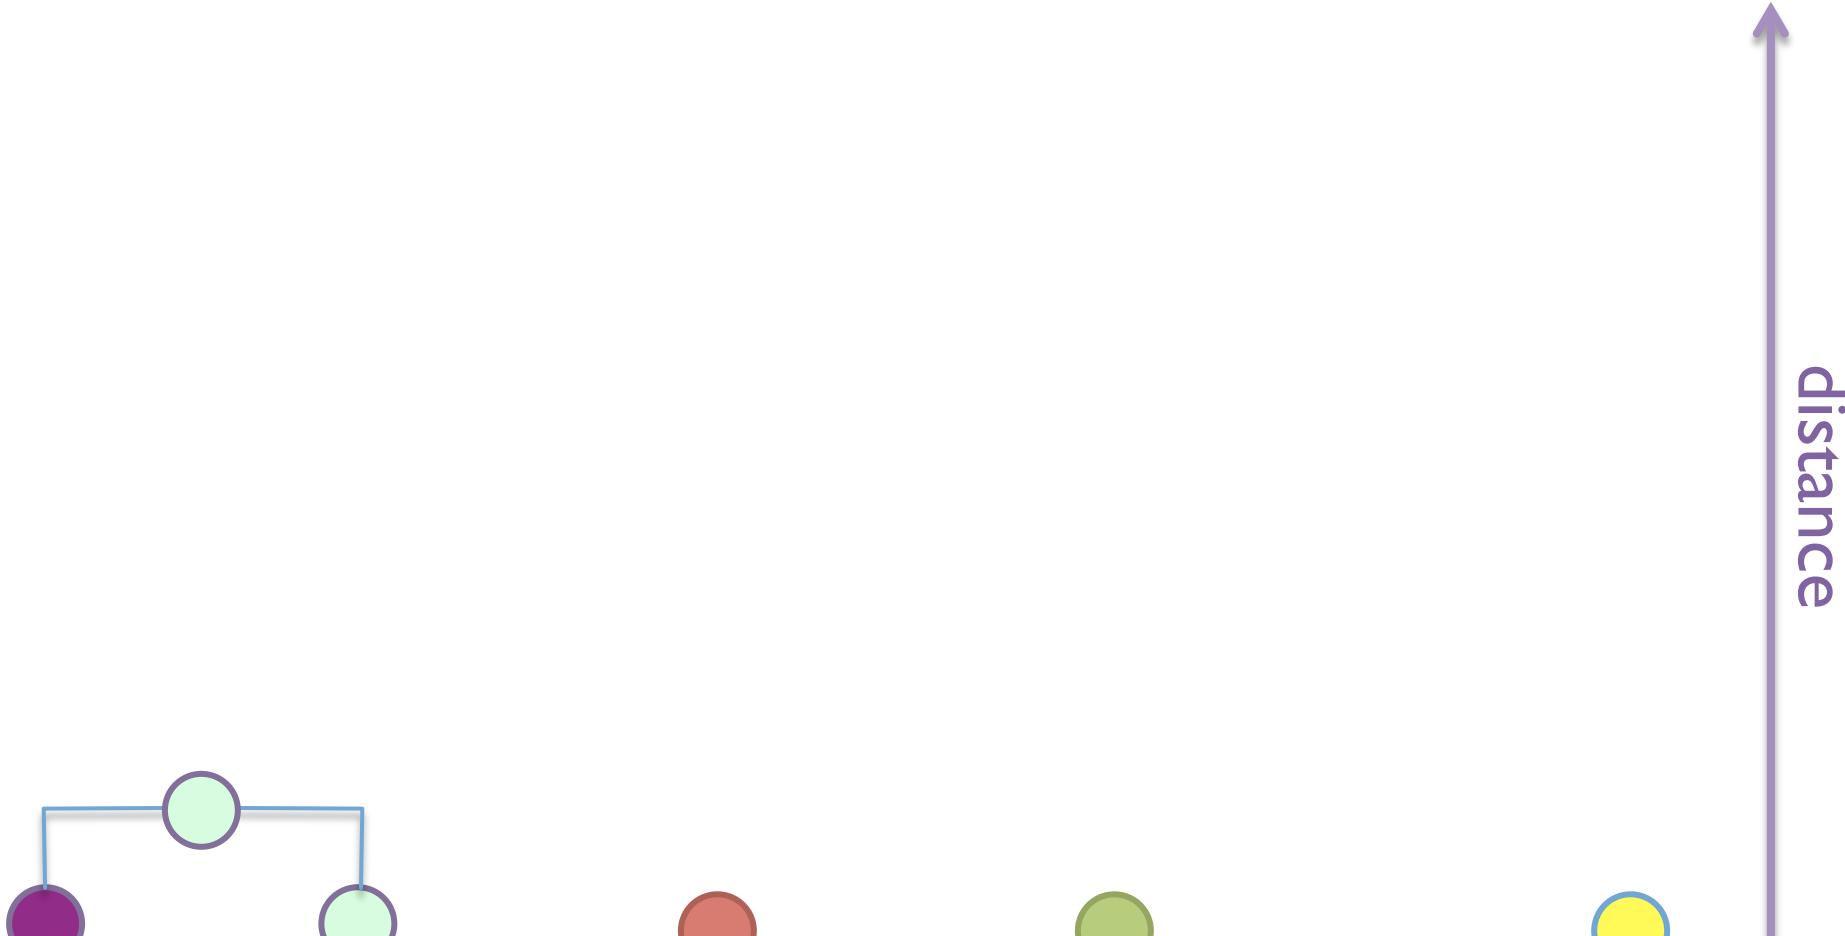
# Agglomerative Hierarchical Clustering



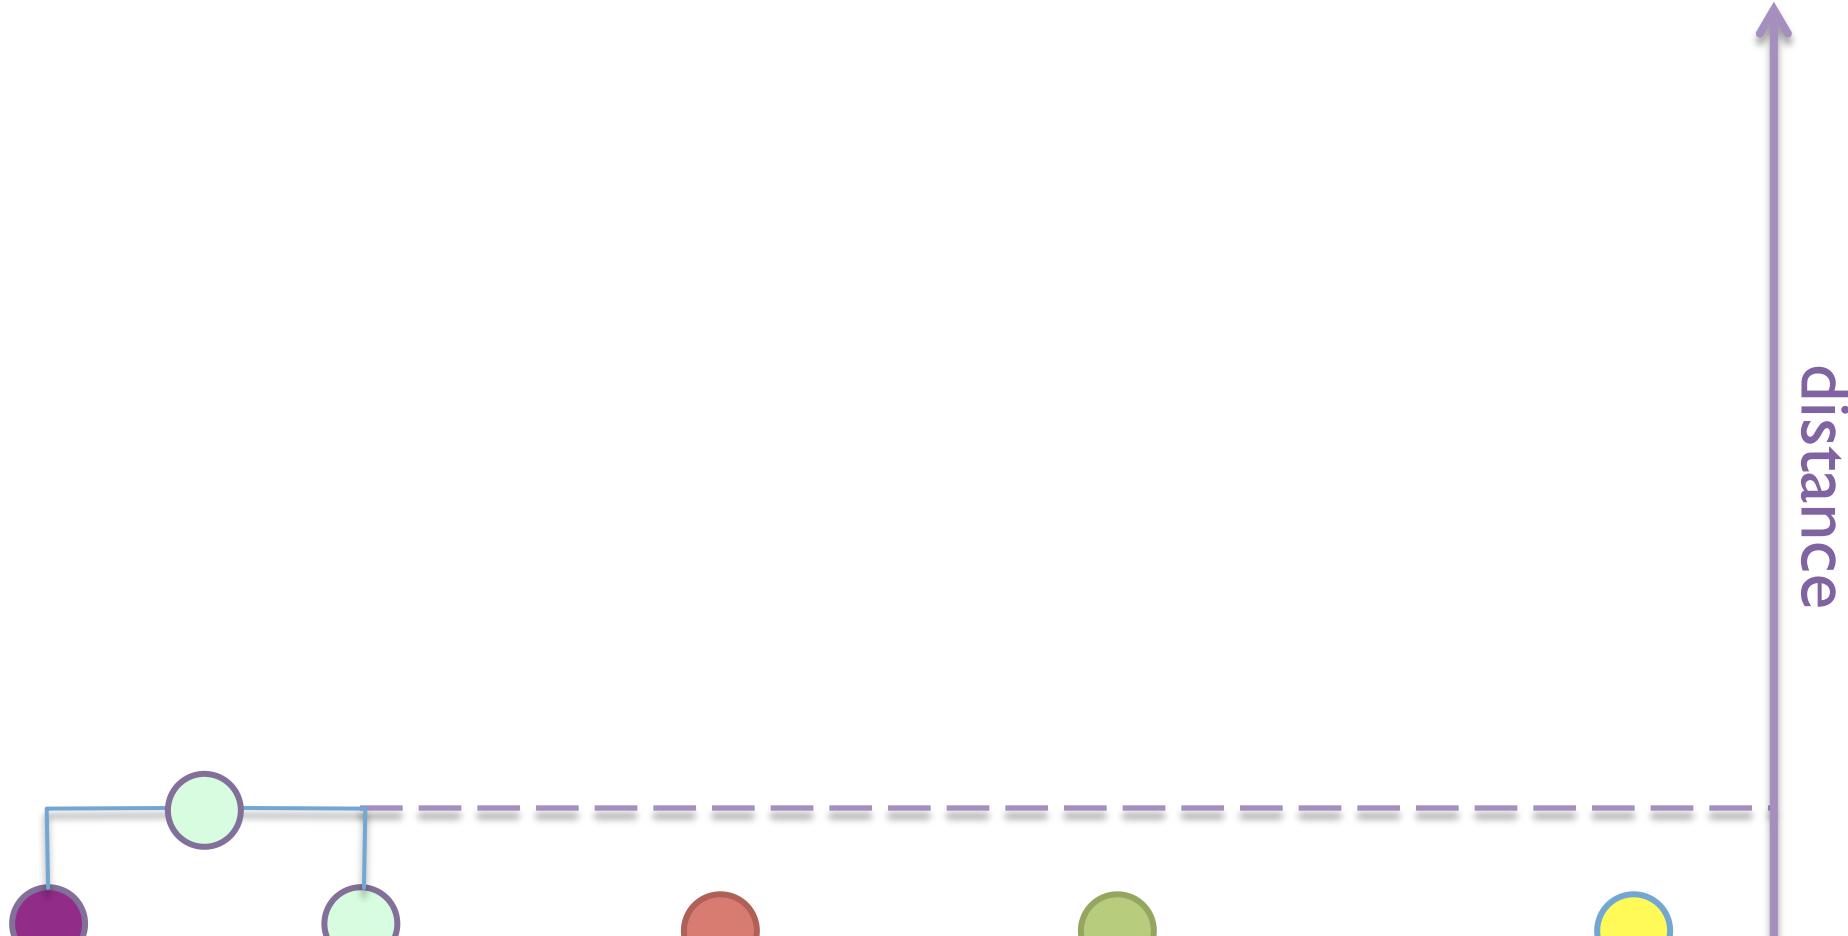
# Agglomerative Hierarchical Clustering



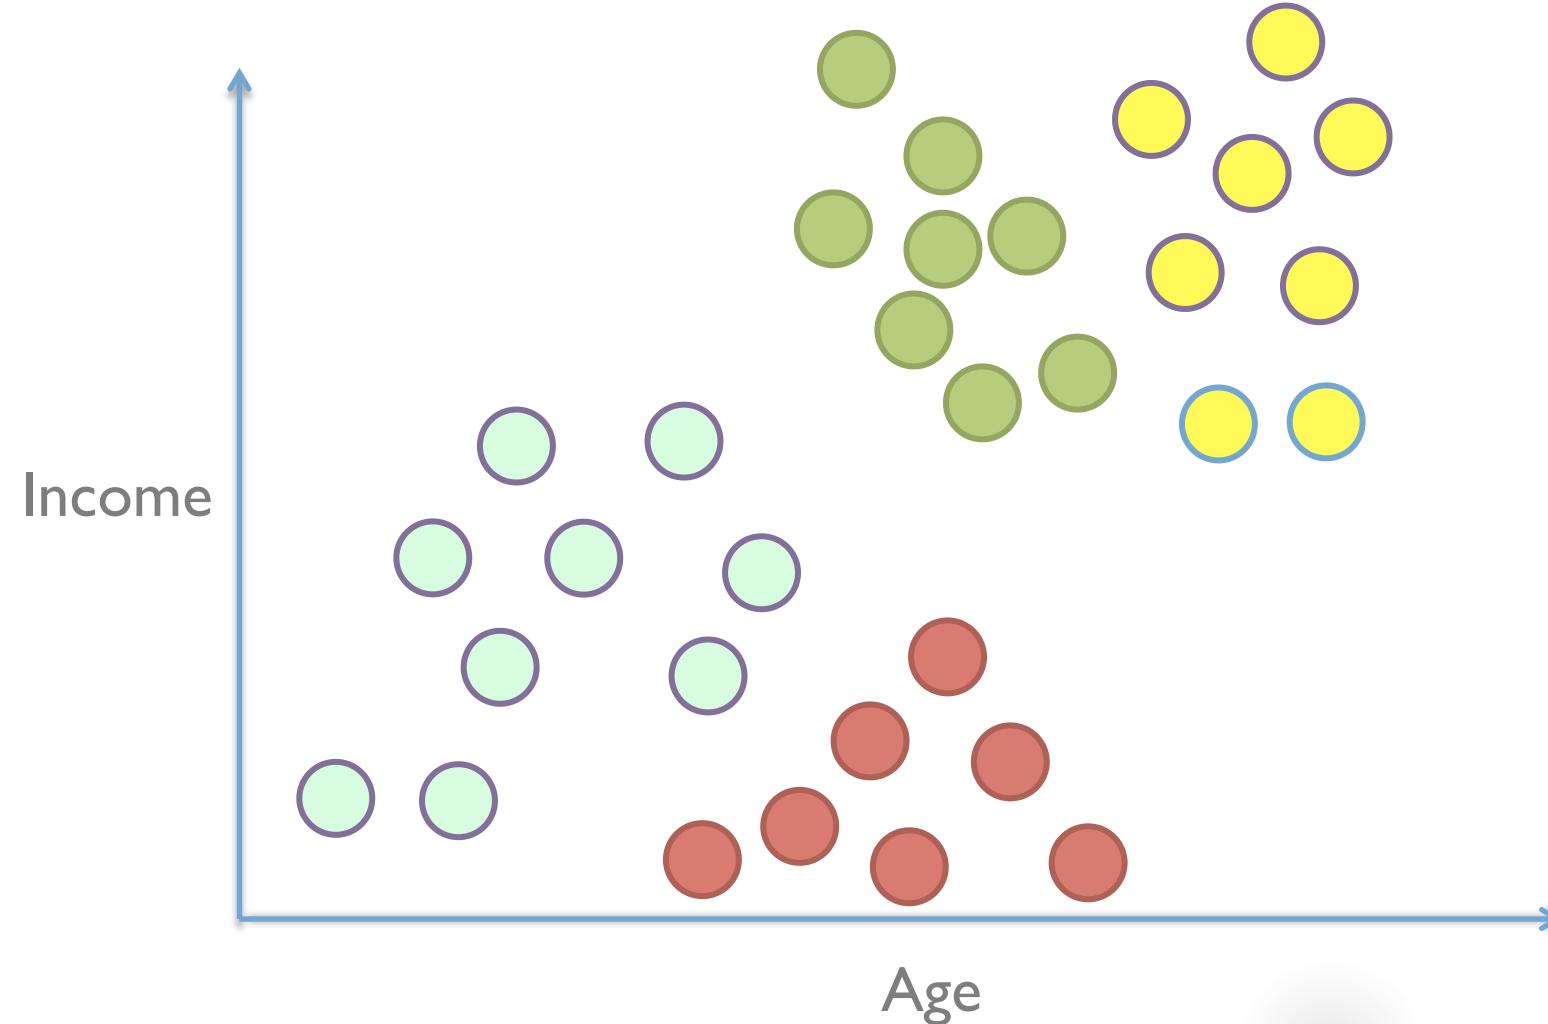
# Agglomerative Hierarchical Clustering



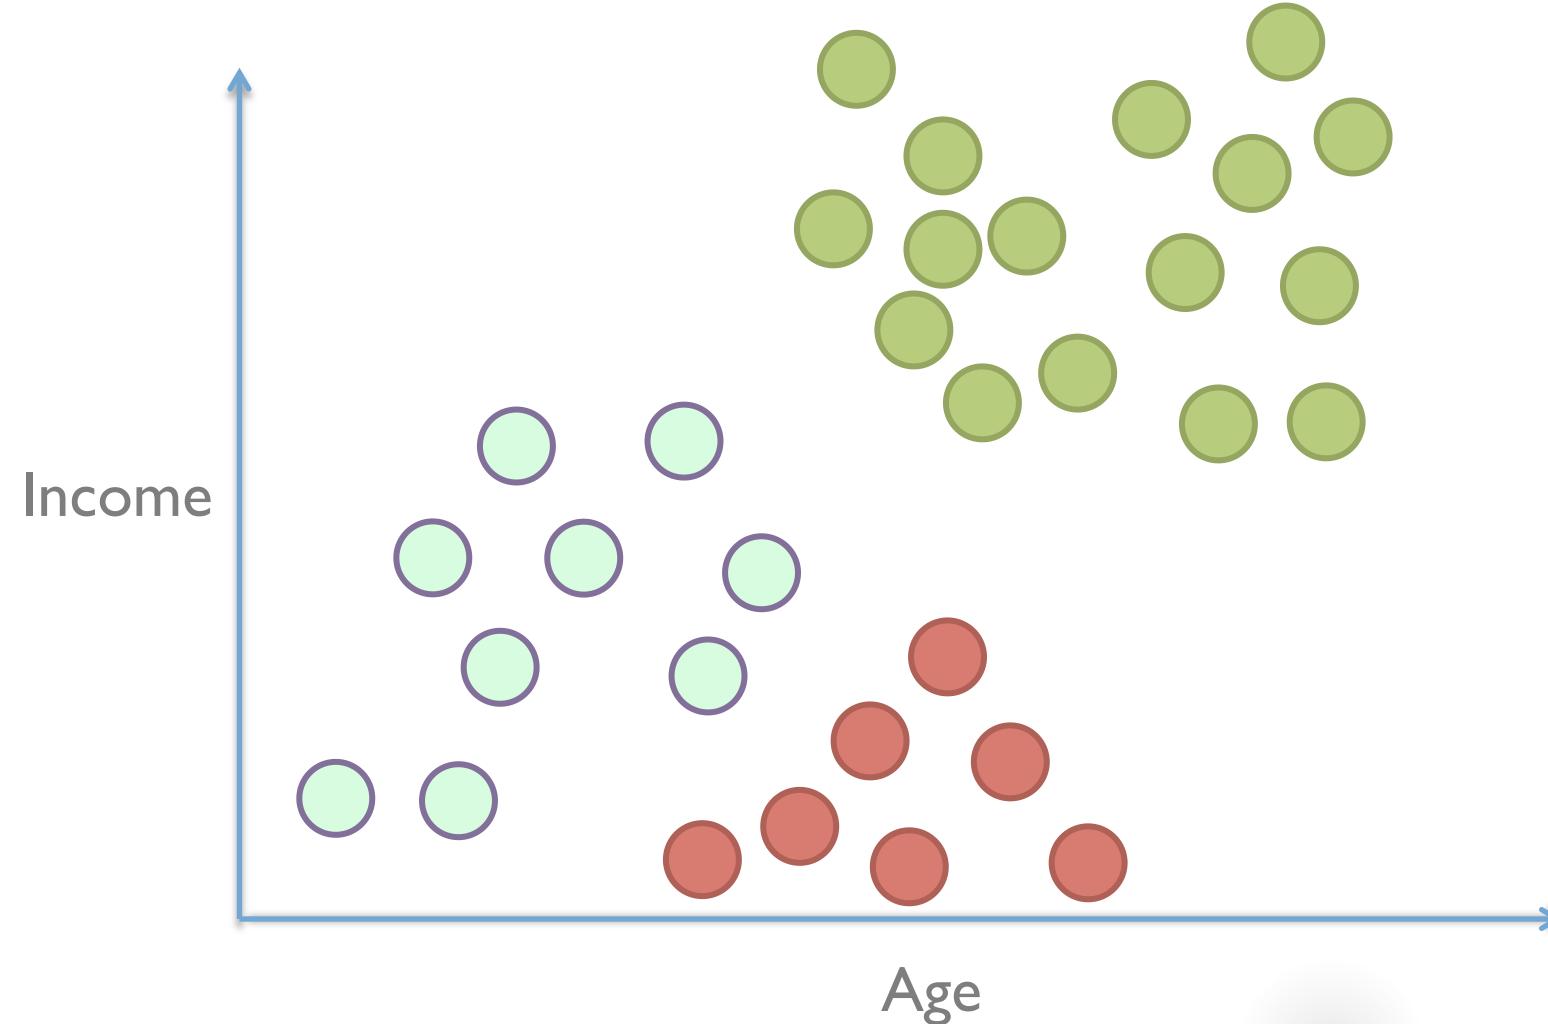
# Agglomerative Hierarchical Clustering



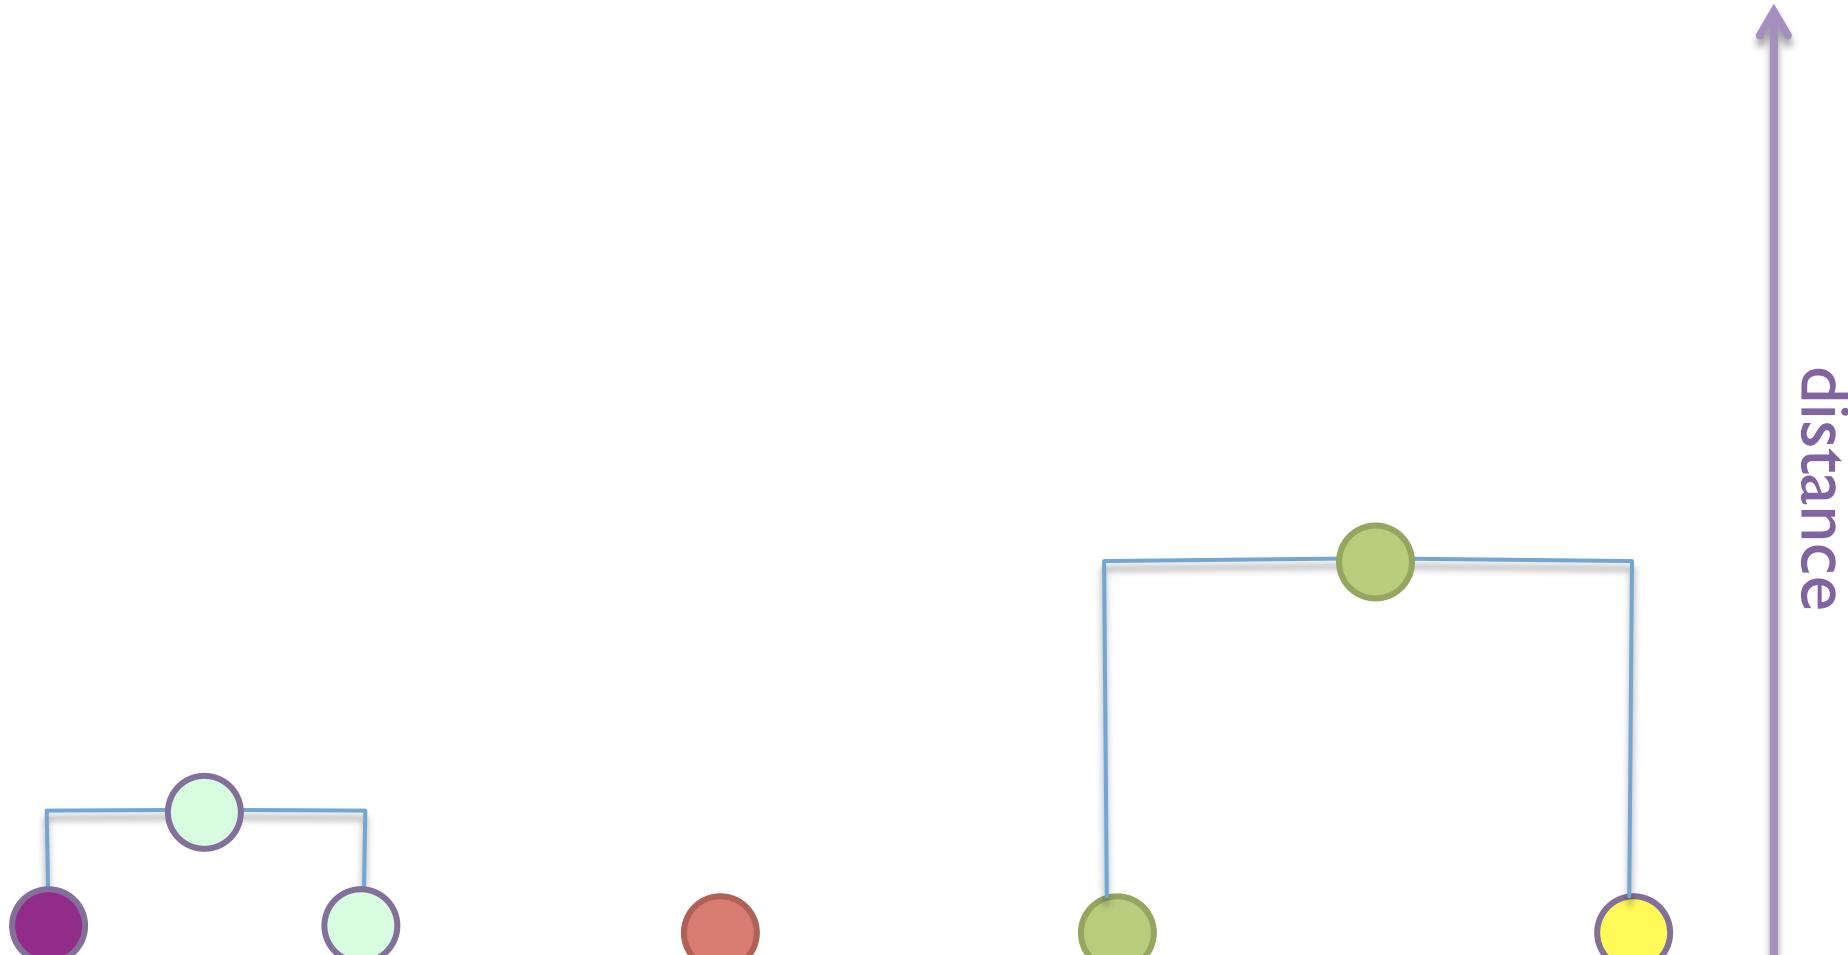
# Agglomerative Hierarchical Clustering



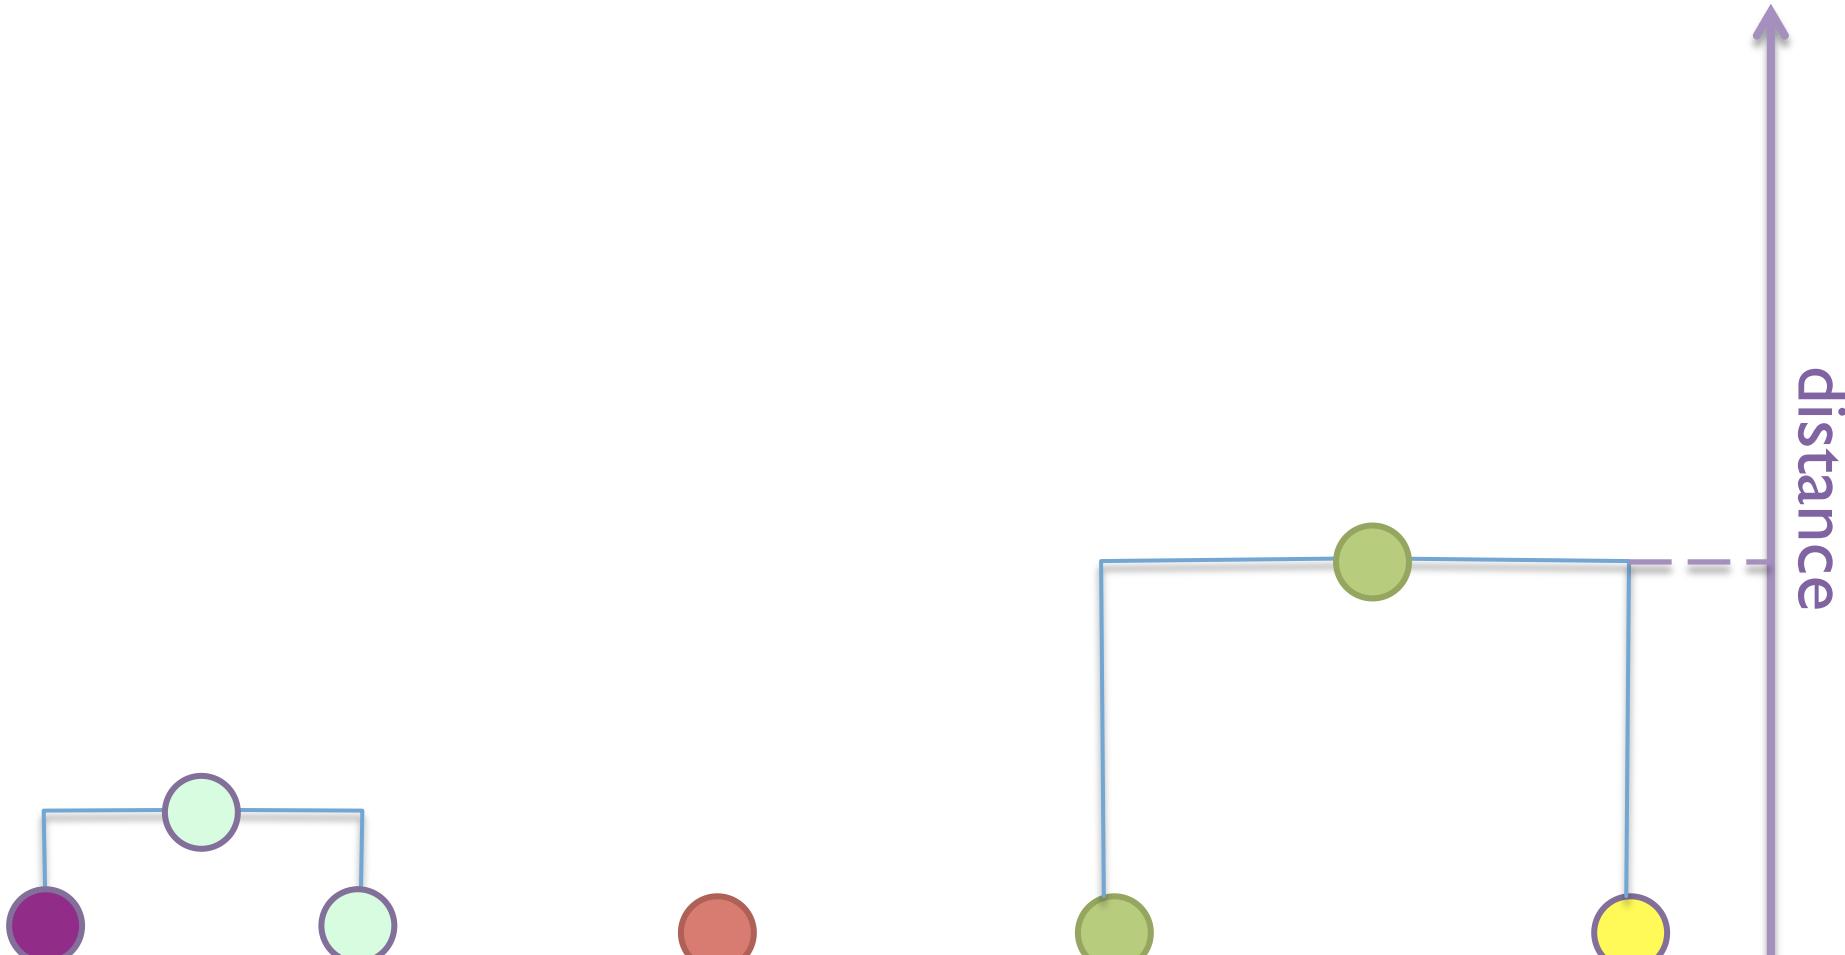
# Agglomerative Hierarchical Clustering



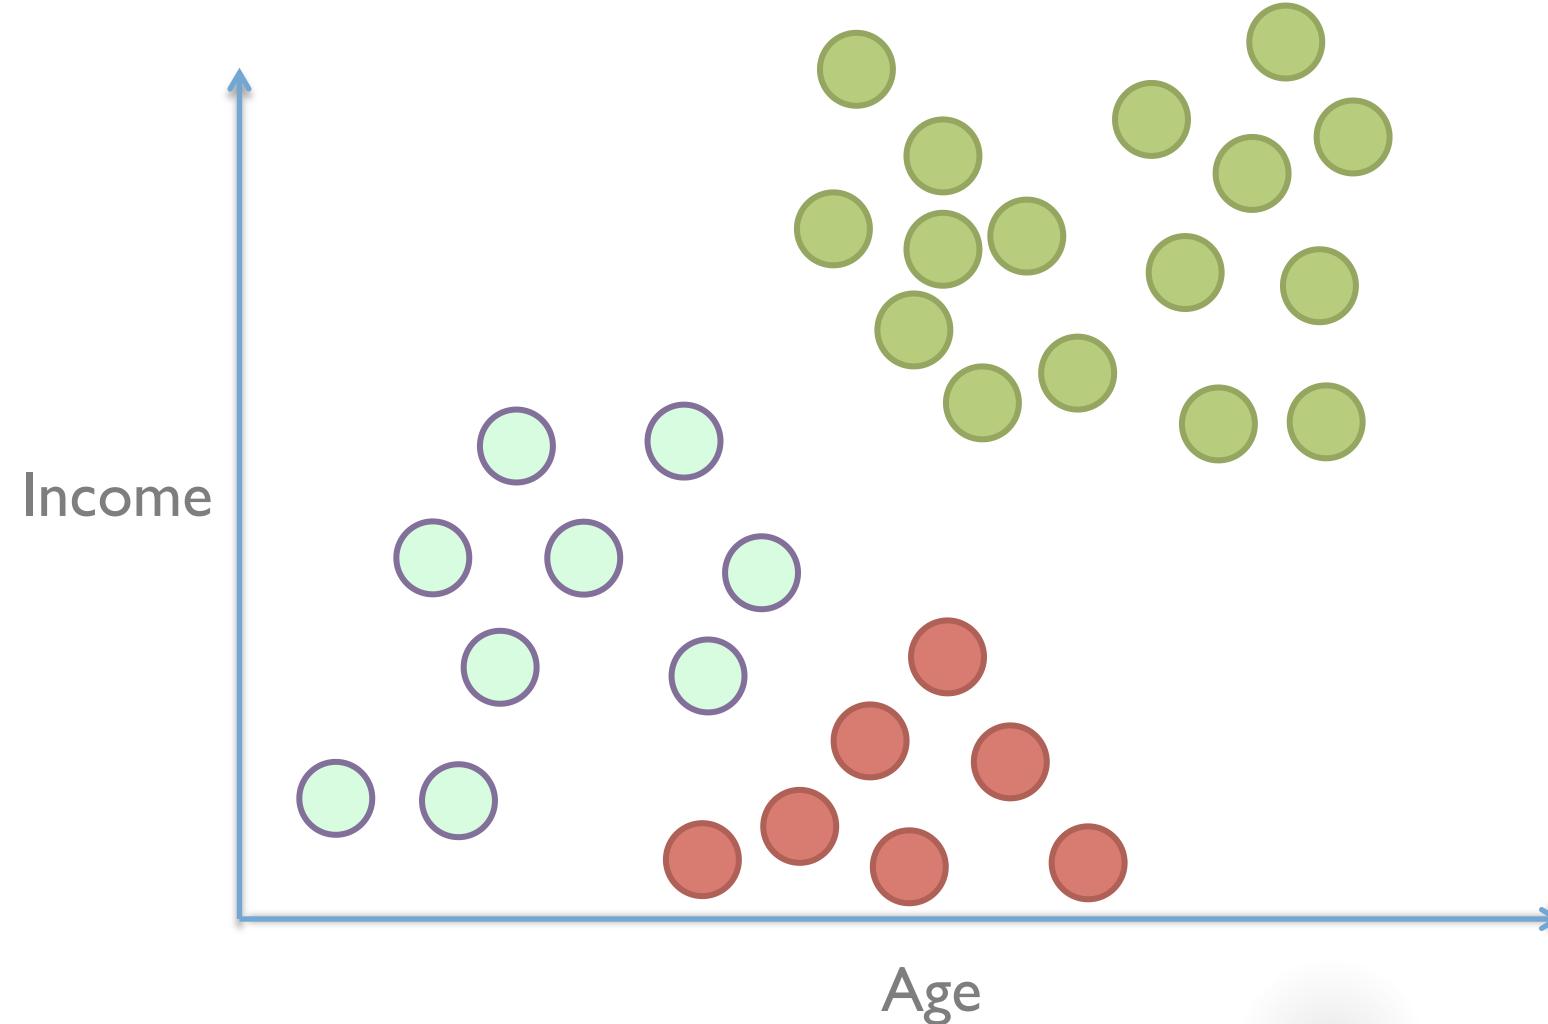
# Agglomerative Hierarchical Clustering



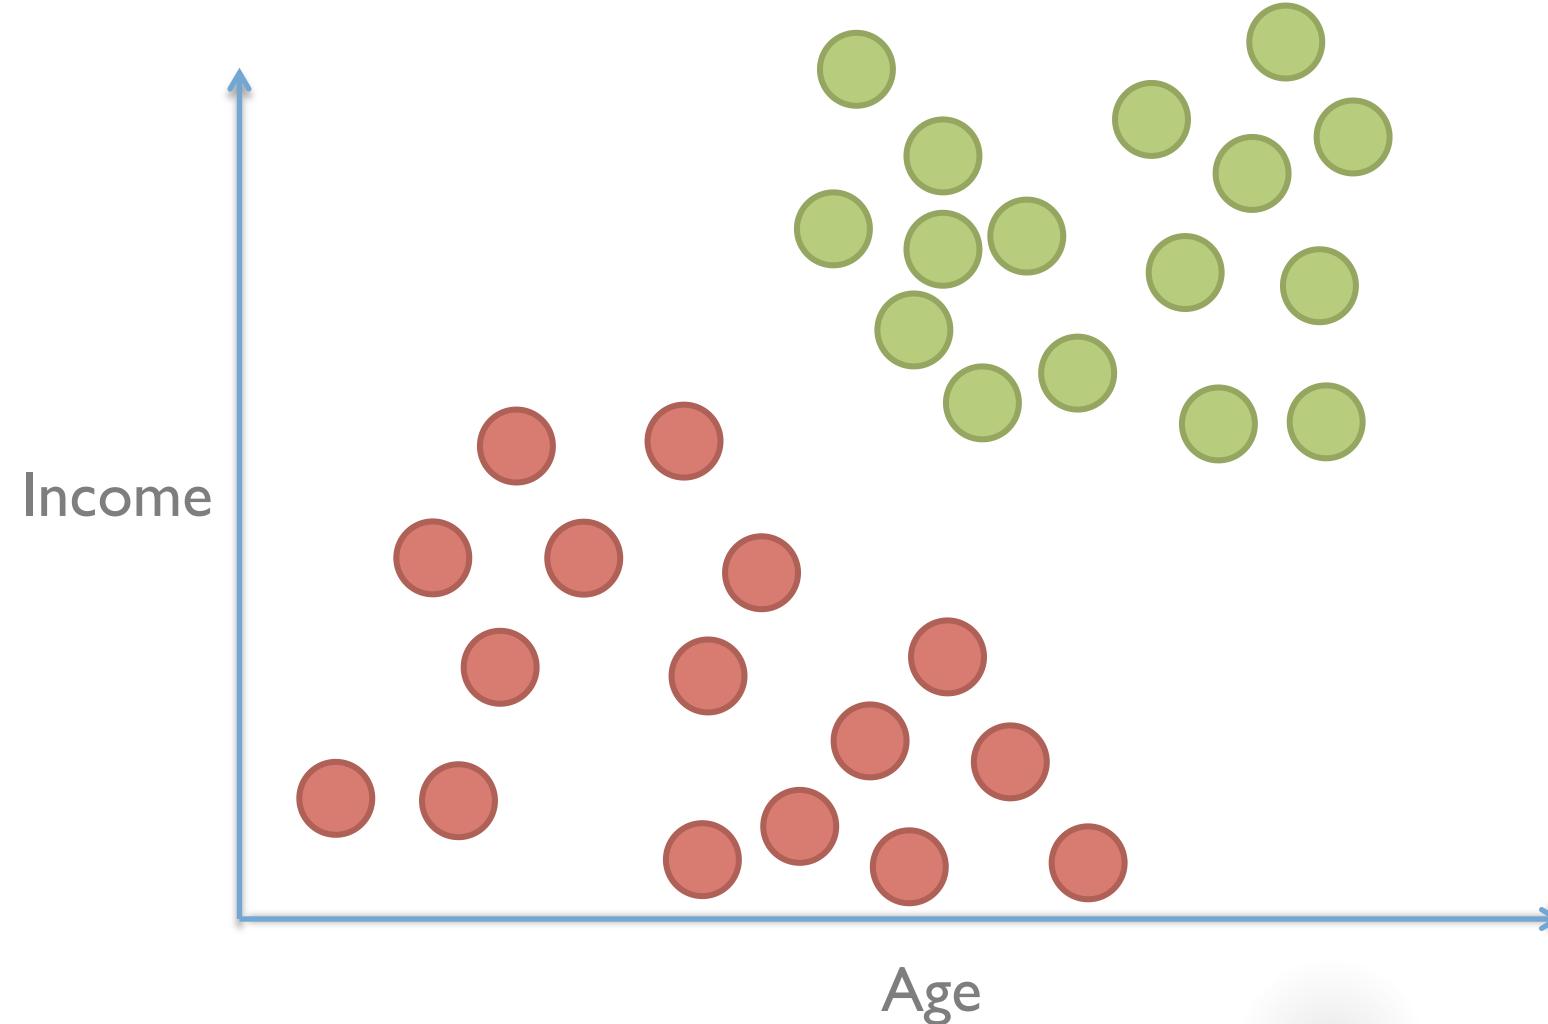
# Agglomerative Hierarchical Clustering



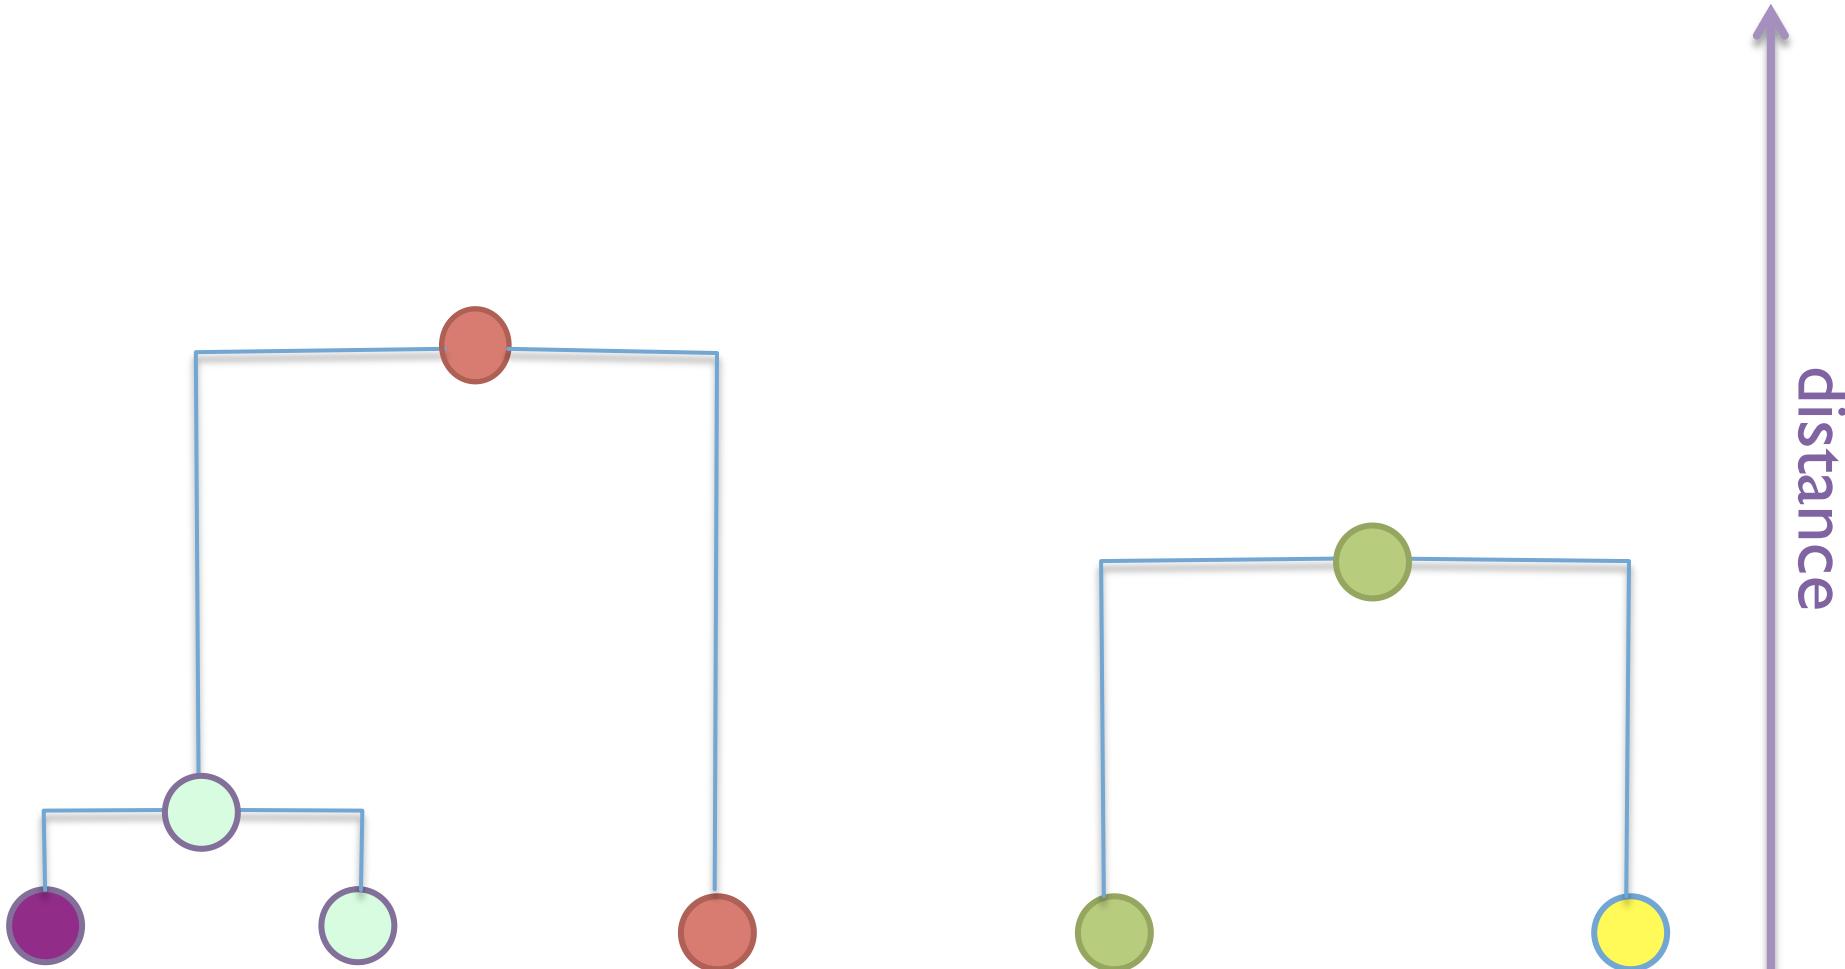
# Agglomerative Hierarchical Clustering



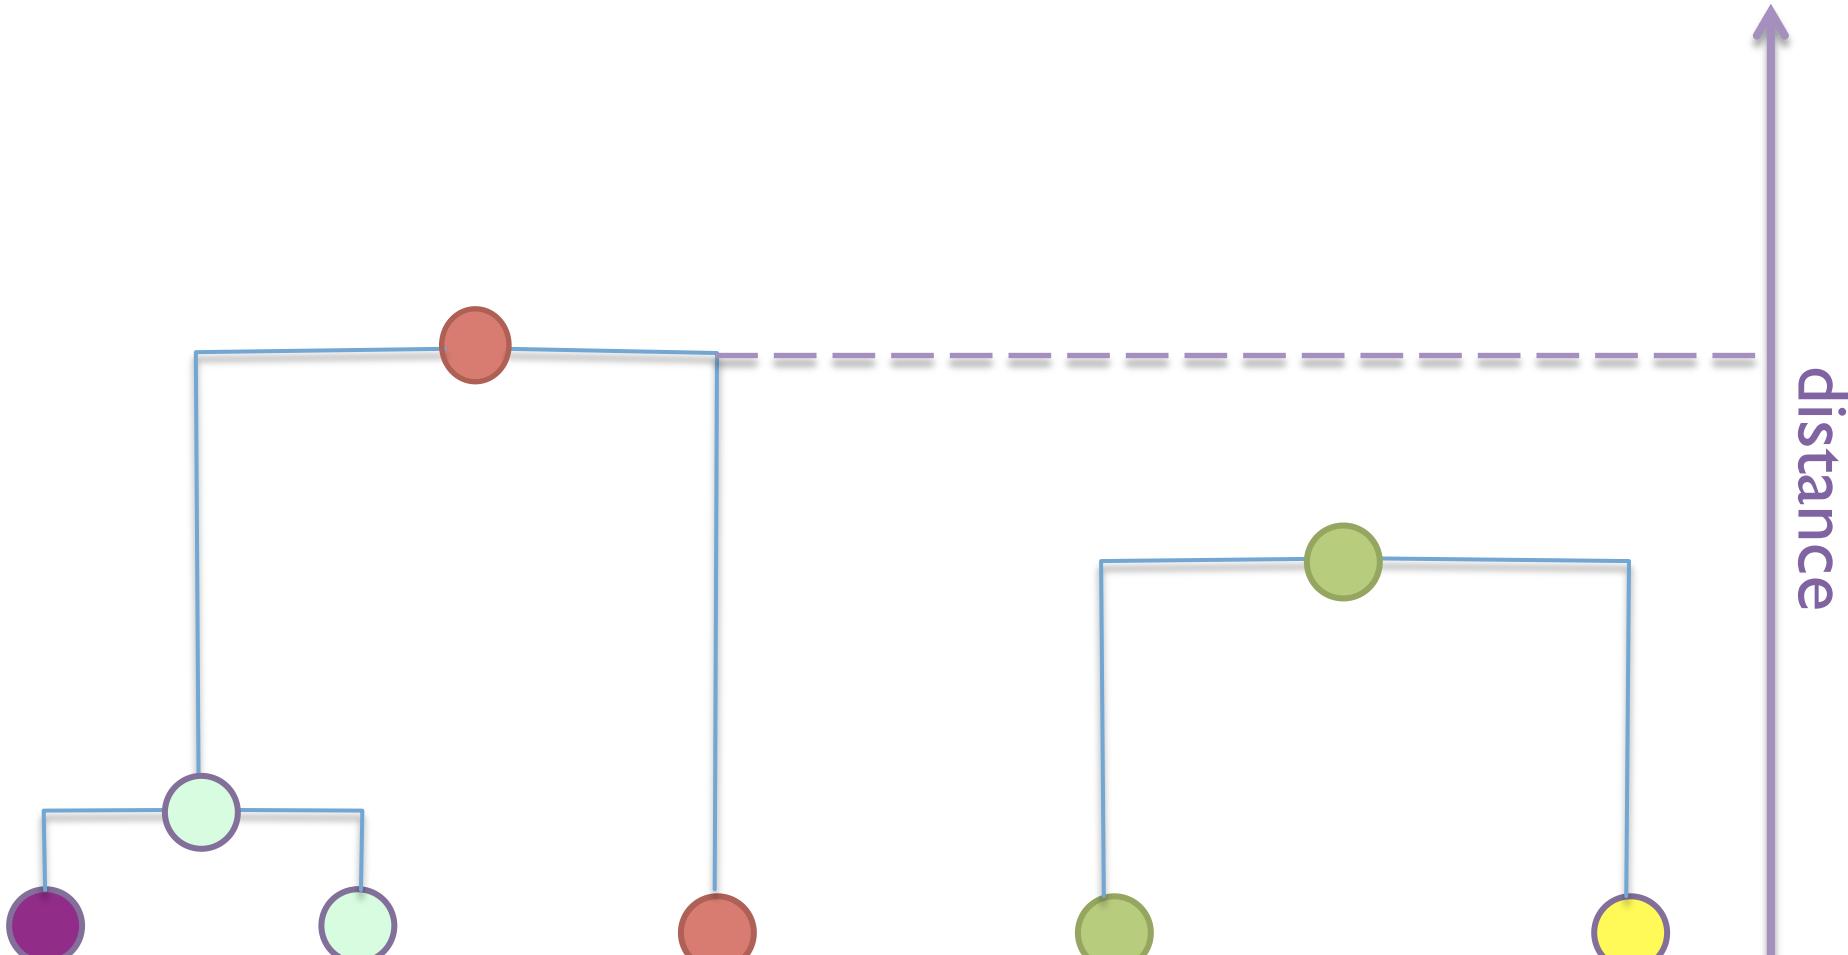
# Agglomerative Hierarchical Clustering



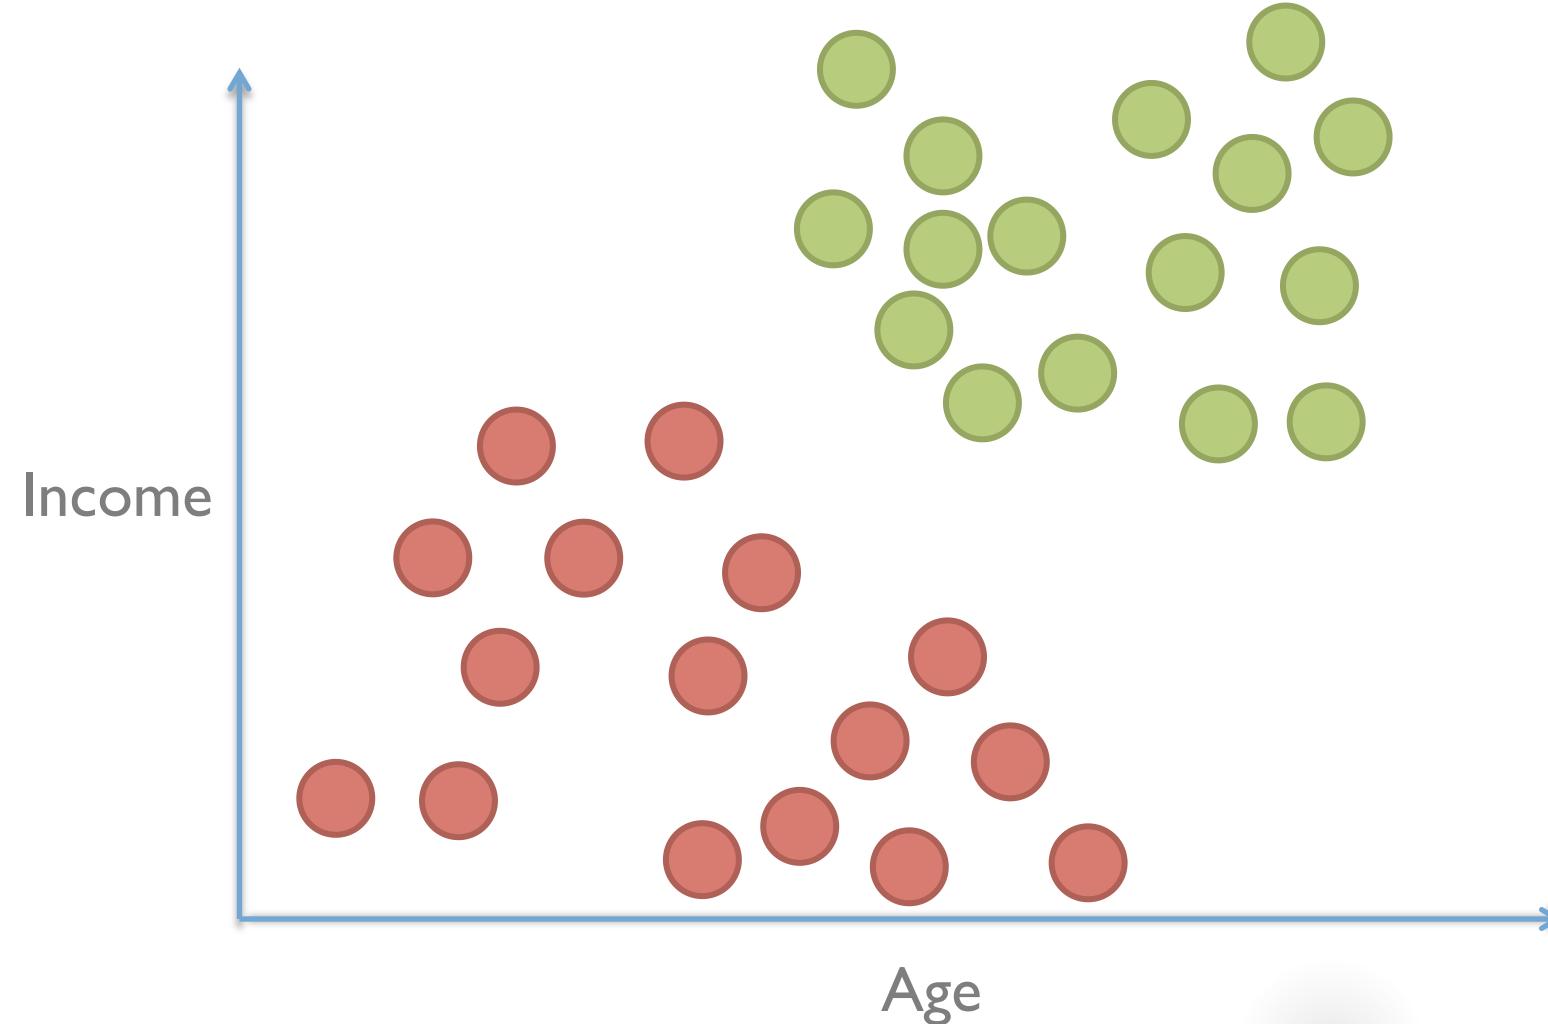
# Agglomerative Hierarchical Clustering



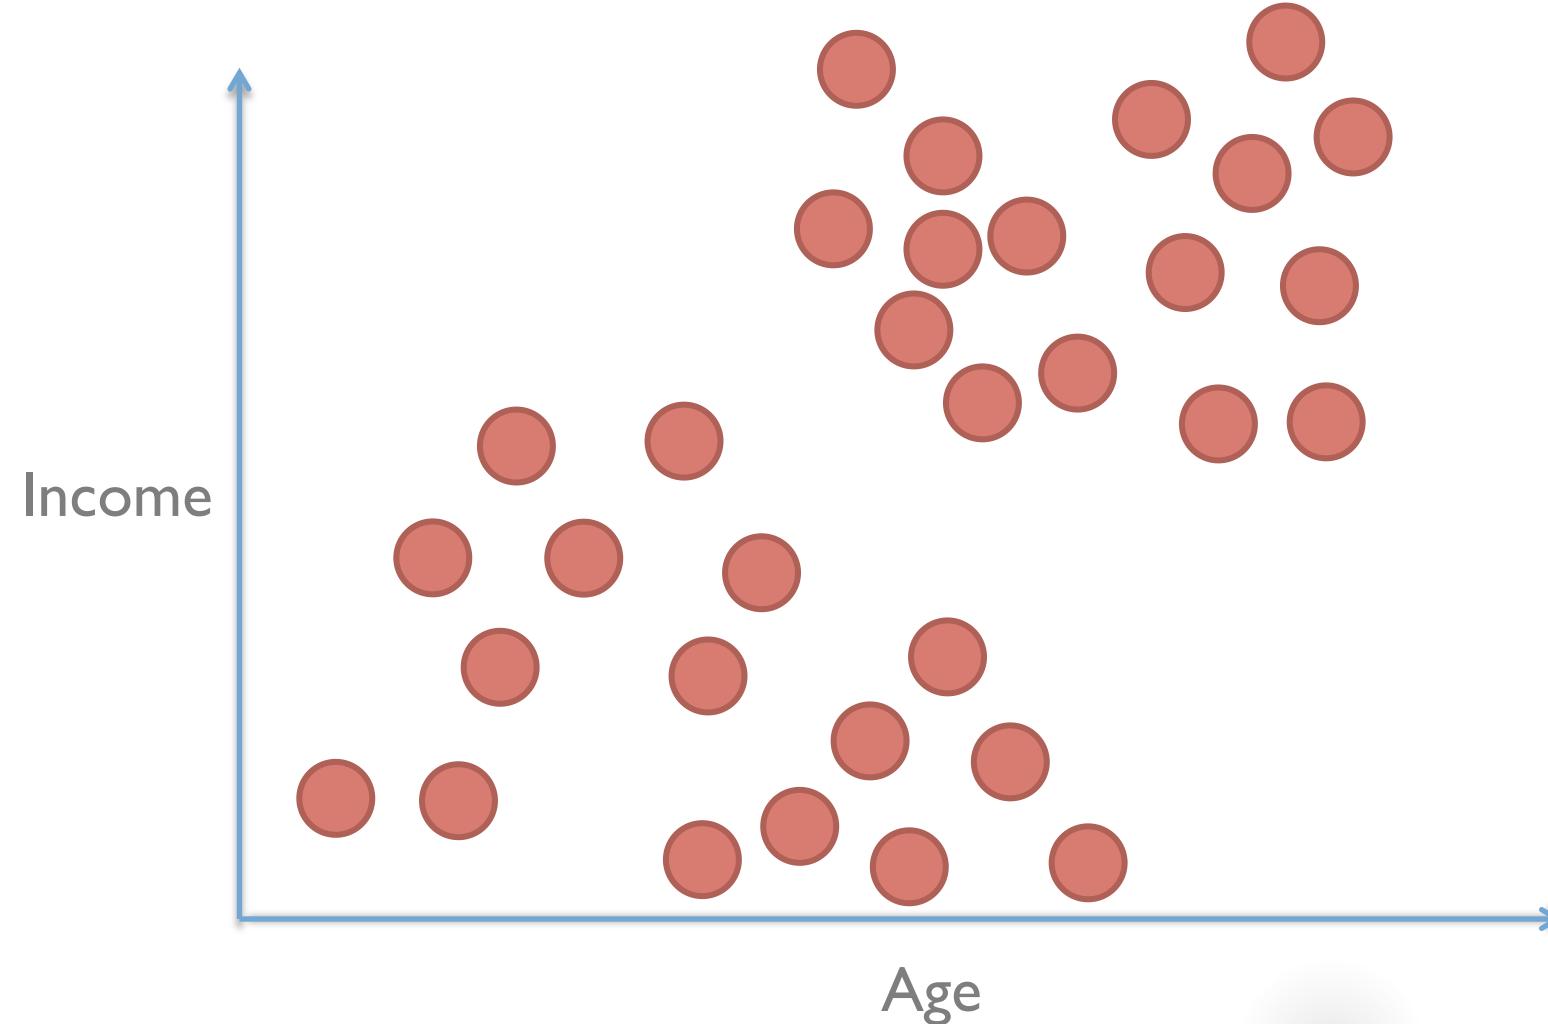
# Agglomerative Hierarchical Clustering



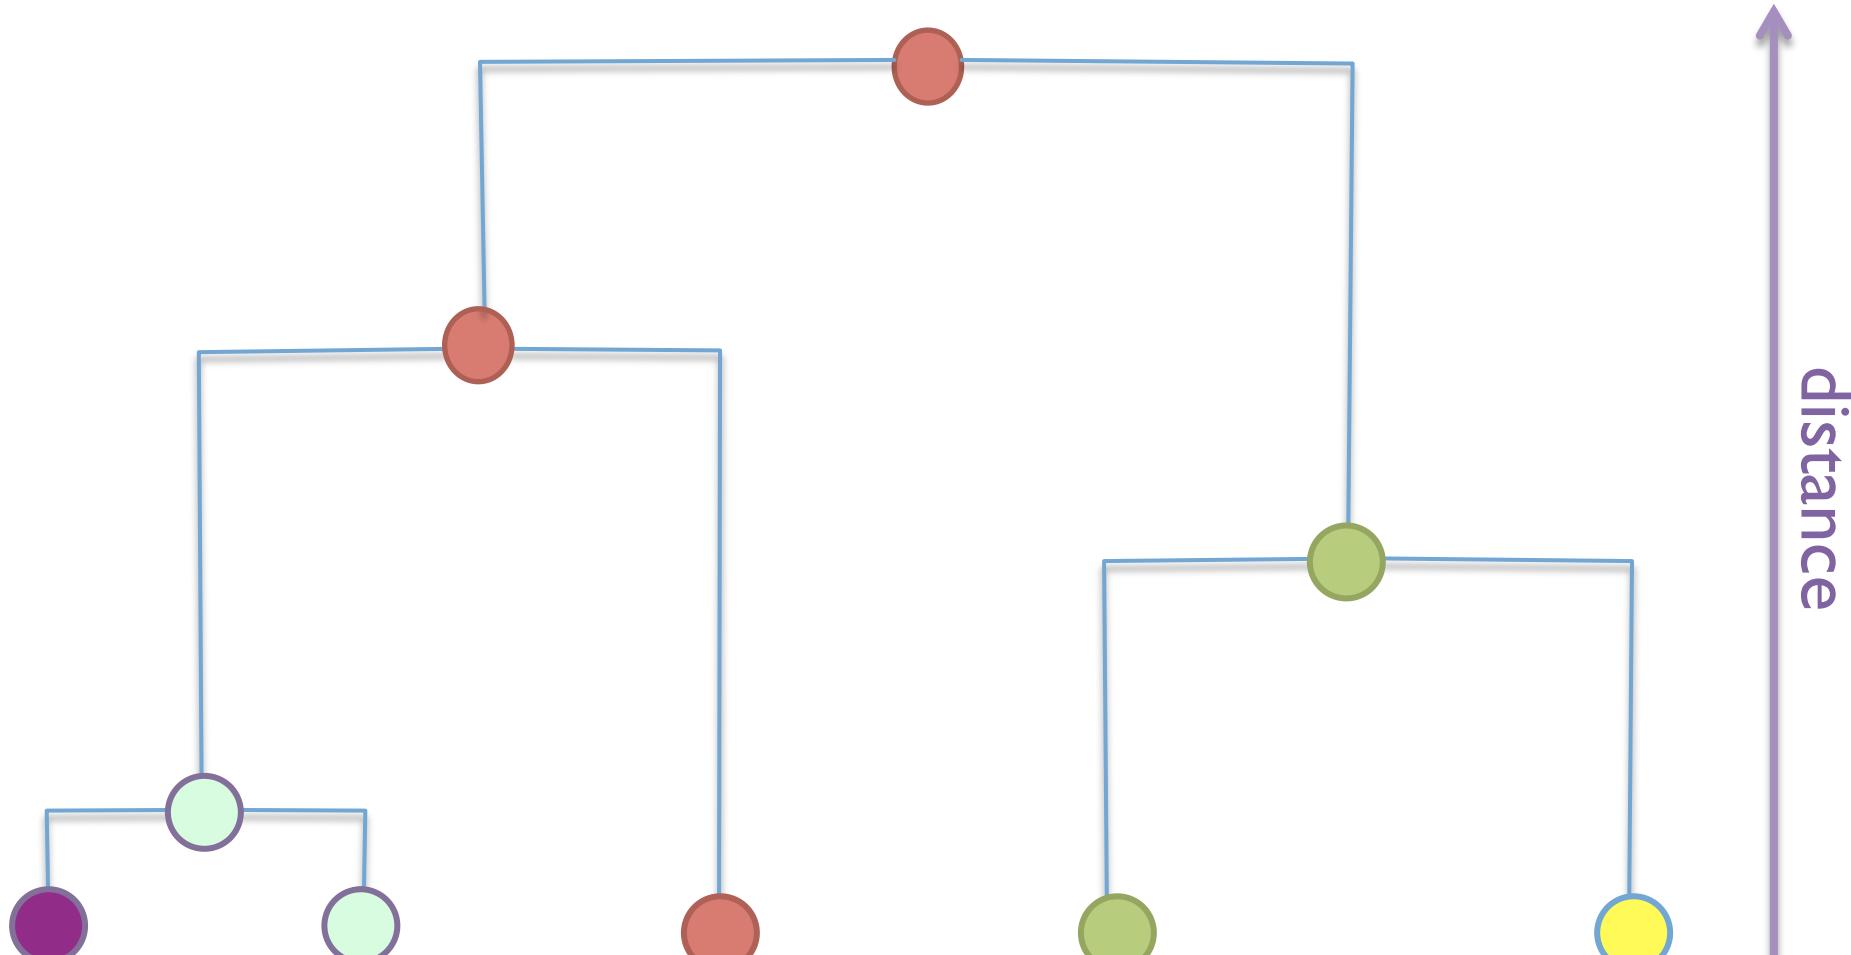
# Agglomerative Hierarchical Clustering



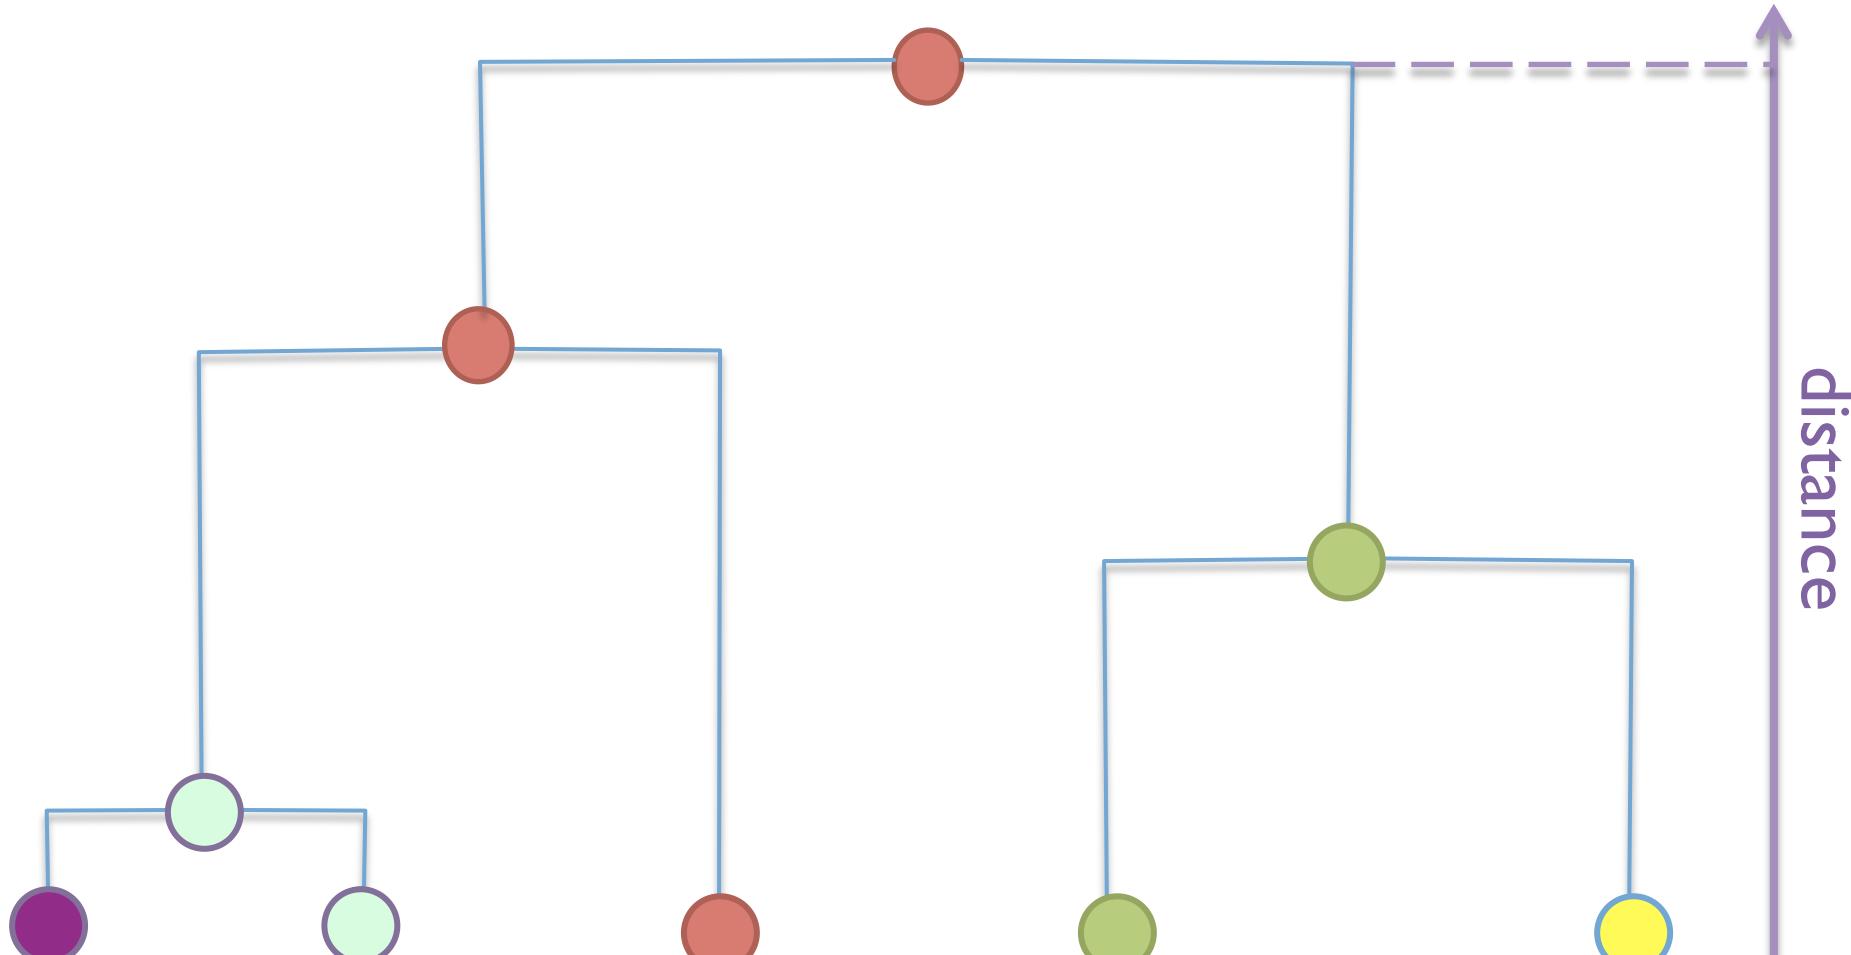
# Agglomerative Hierarchical Clustering



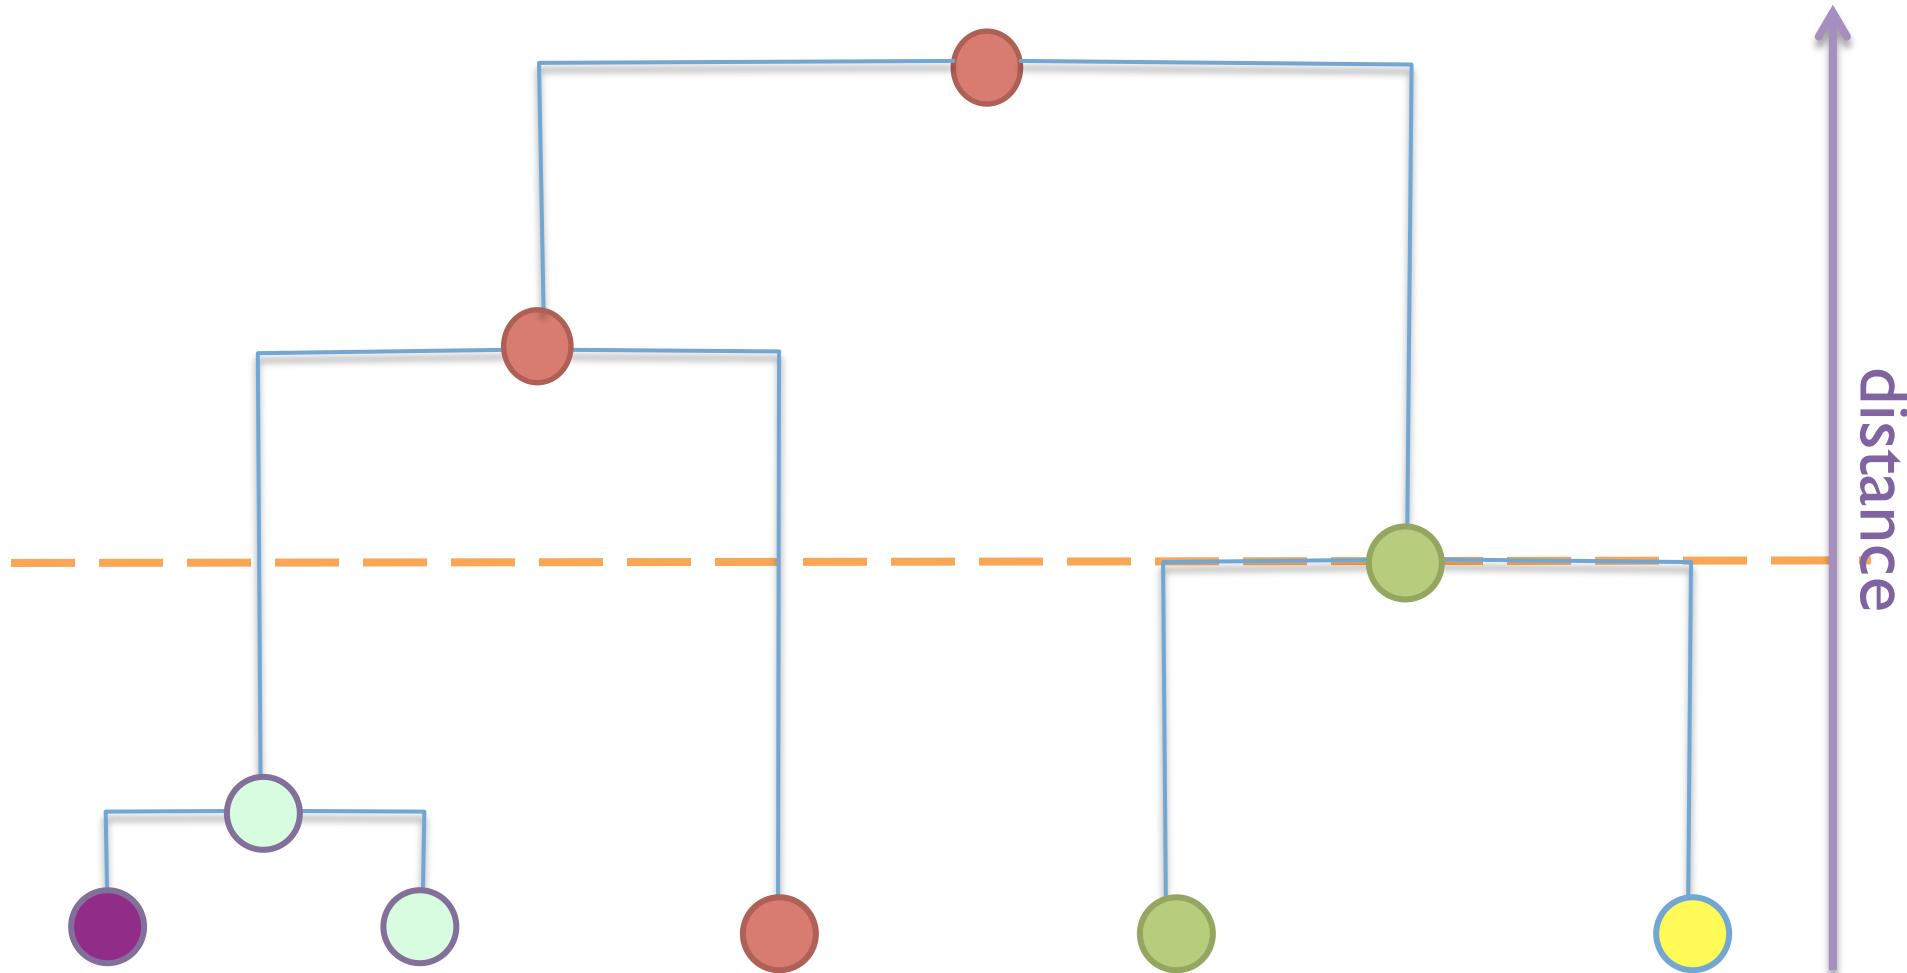
# Agglomerative Hierarchical Clustering



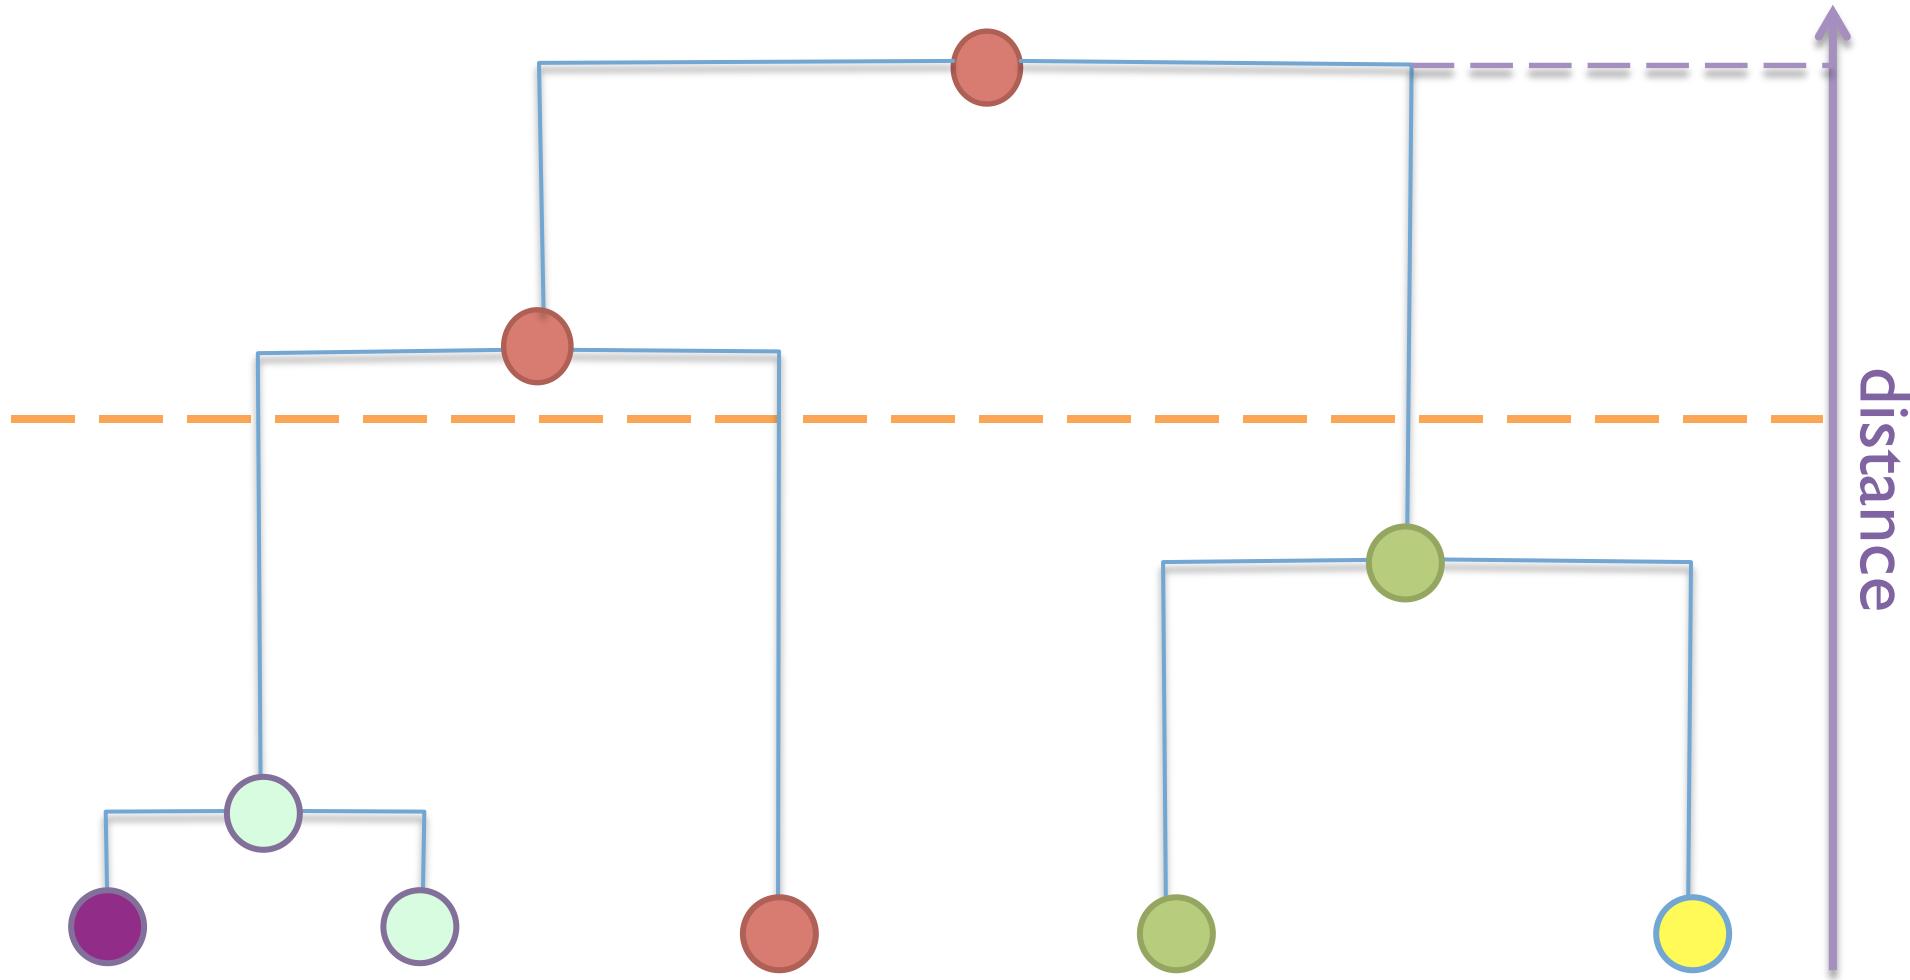
# Agglomerative Hierarchical Clustering



# Criterion: Stop at $k=3$ clusters!



# Criterion: Stop at $\text{dist}=3 \dots \dots !$



```
from sklearn.cluster import  
AgglomerativeClustering
```

```
AgglomerativeClustering(....,
```

```
....
```

```
,
```

```
linkage = "ward",
```

```
....
```

```
,
```

```
affinity = "euclidian",
```

```
....
```

```
)
```



```
from sklearn.cluster import  
AgglomerativeClustering
```

```
AgglomerativeClustering(....,  
                      ...,  
                      linkage = "ward",  
                      ...,  
                      affinity = "euclidian",  
                      ....)
```



## Agglomerative Clustering **linkage**

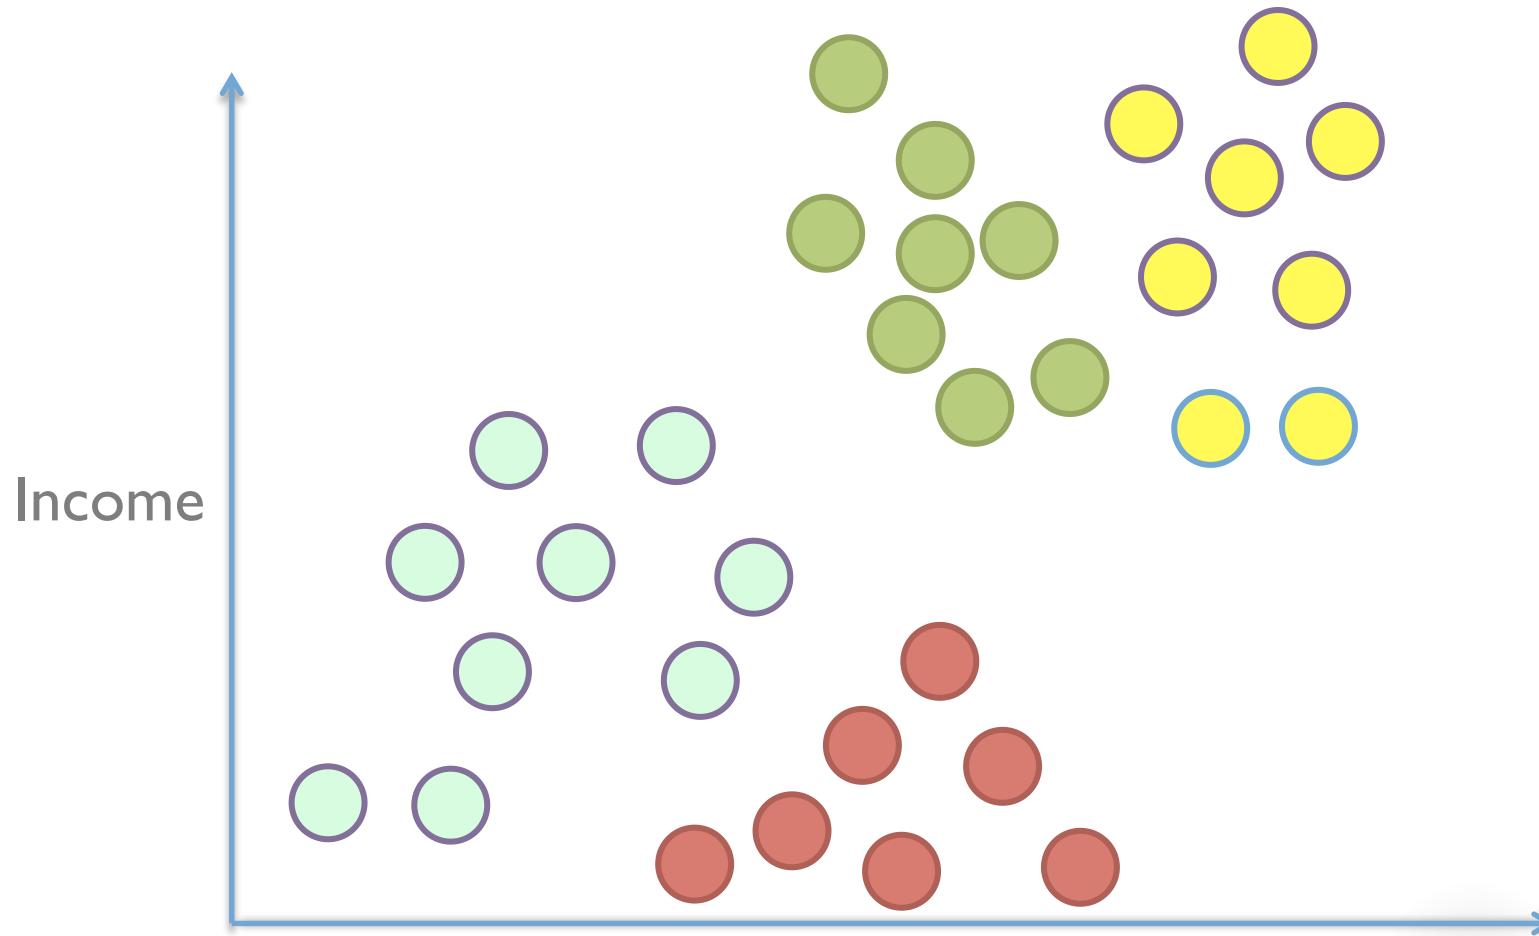
Which clusters to merge next.

Definition of “shortest distance”



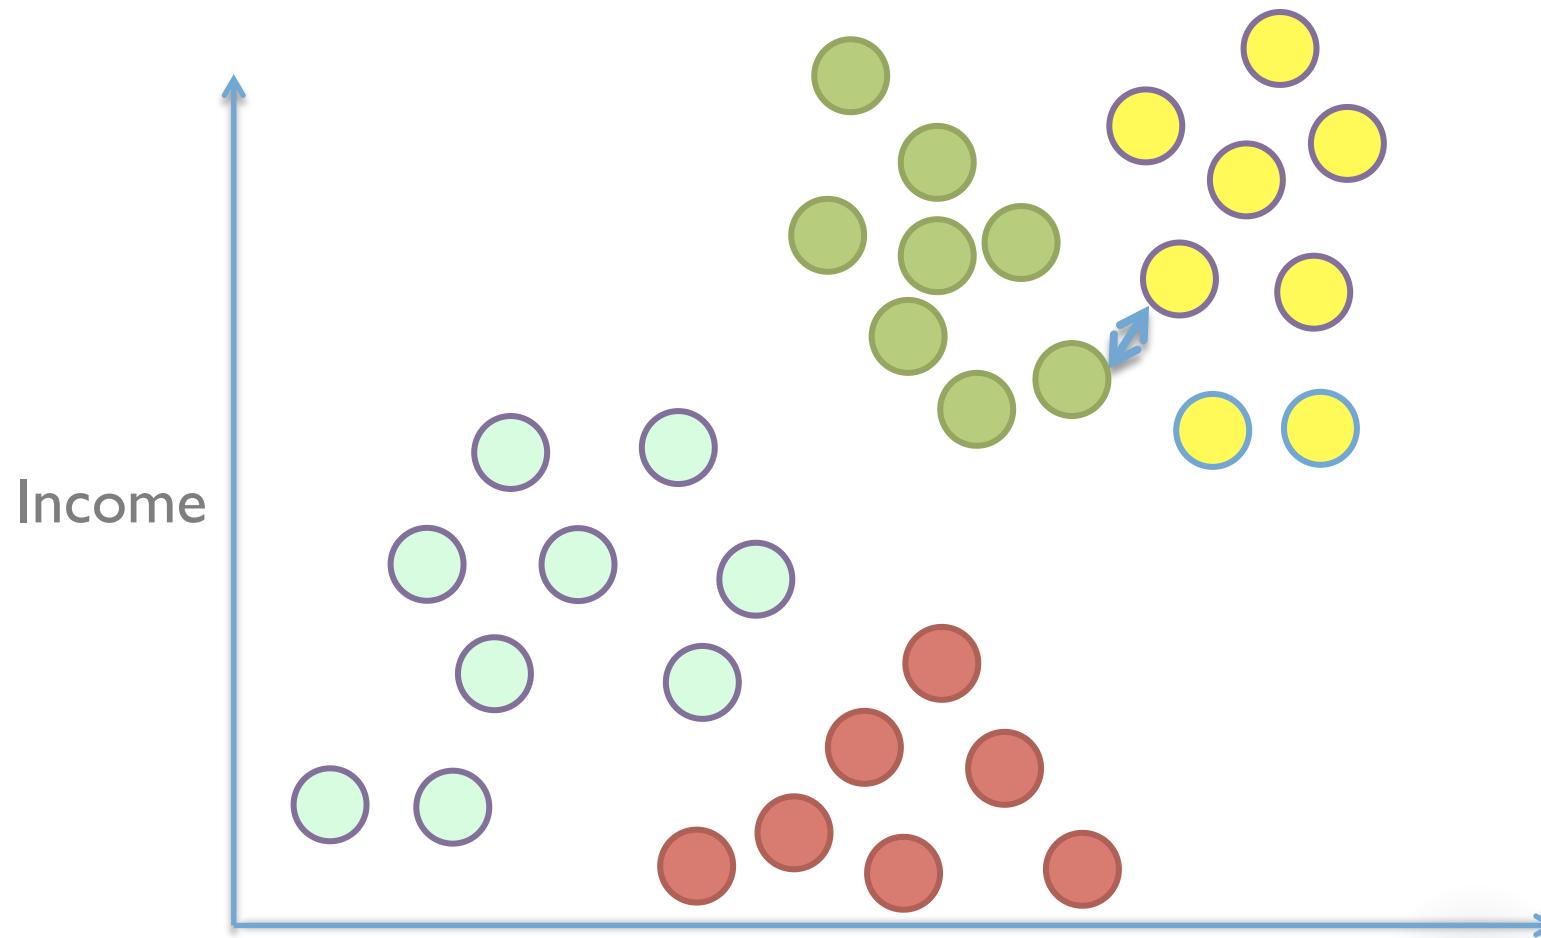
linkage = “single”

Distance between two cluster is the minimum pairwise distance



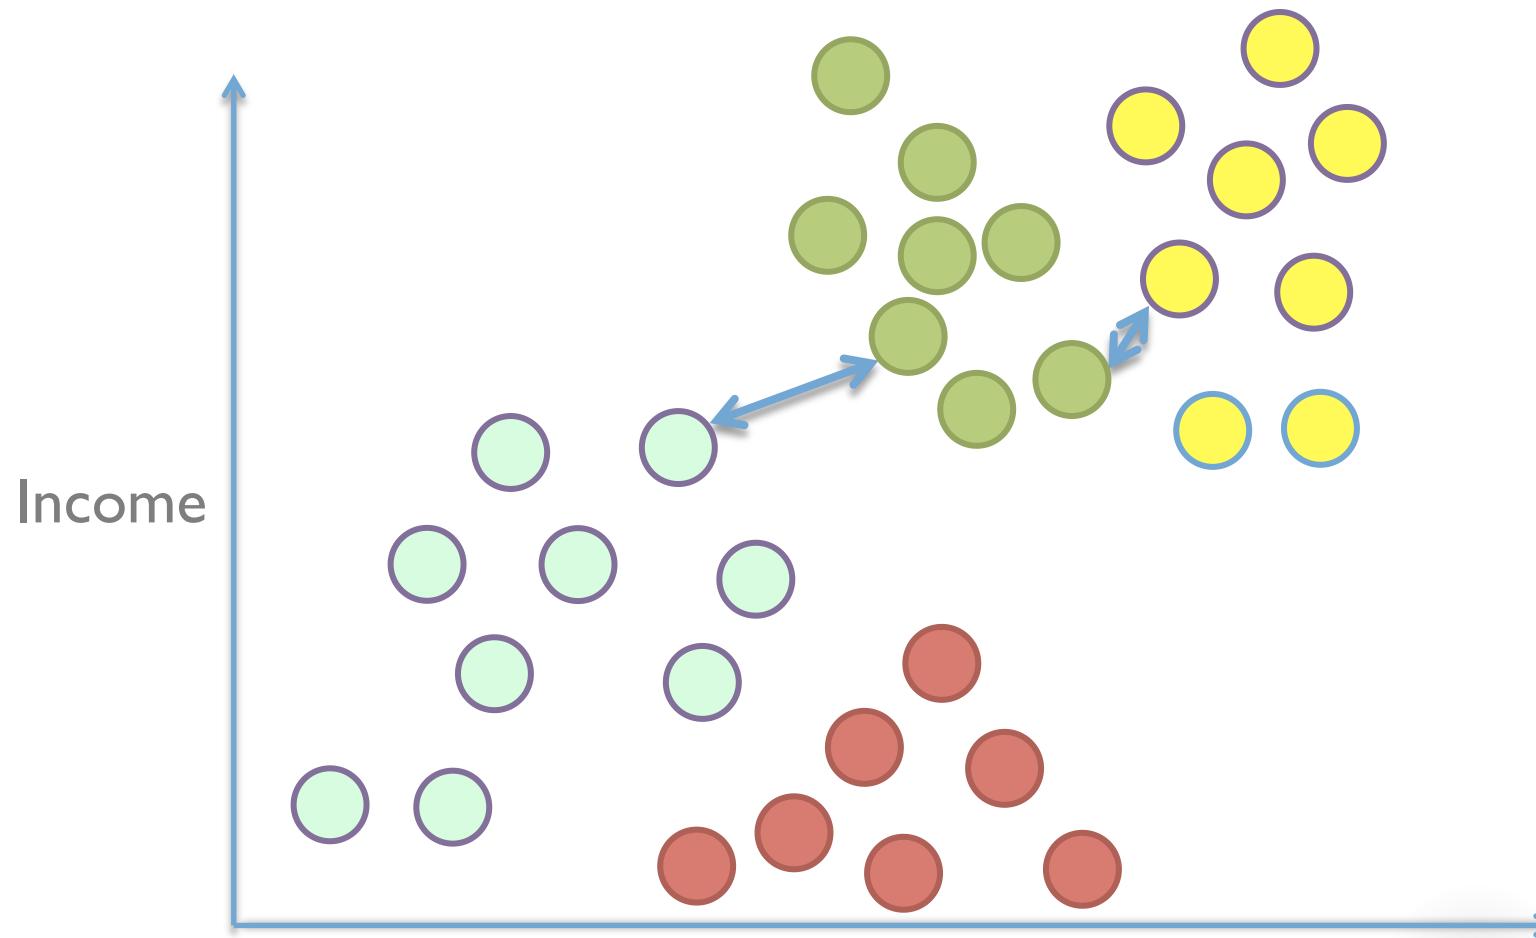
linkage = “single”

Distance between two cluster is the minimum pairwise distance



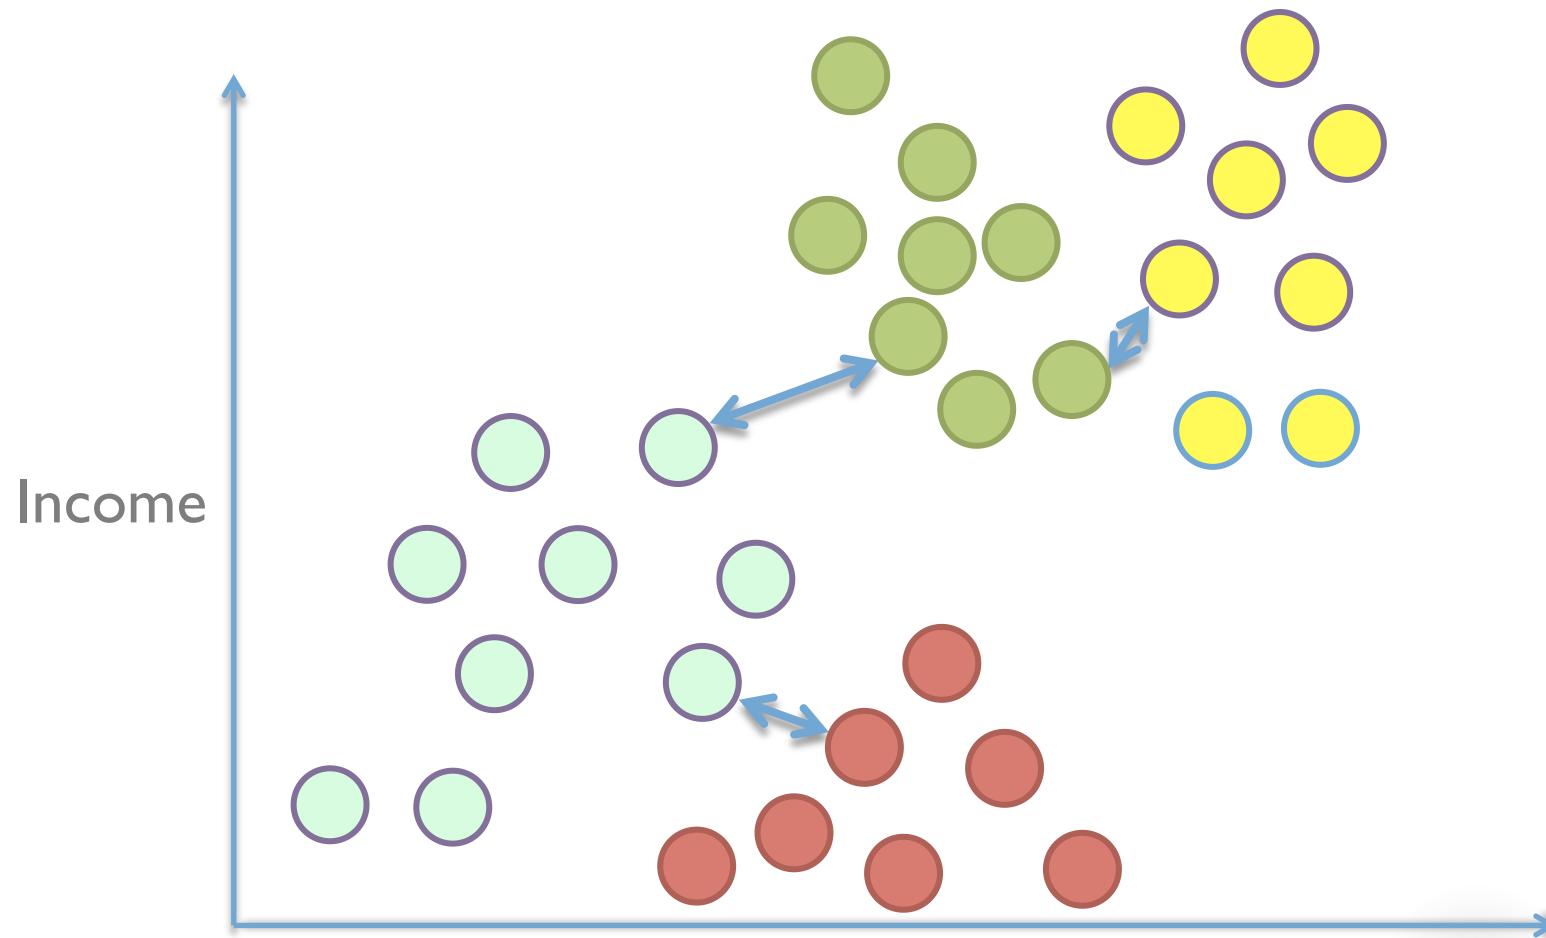
linkage = “single”

Distance between two cluster is the minimum pairwise distance



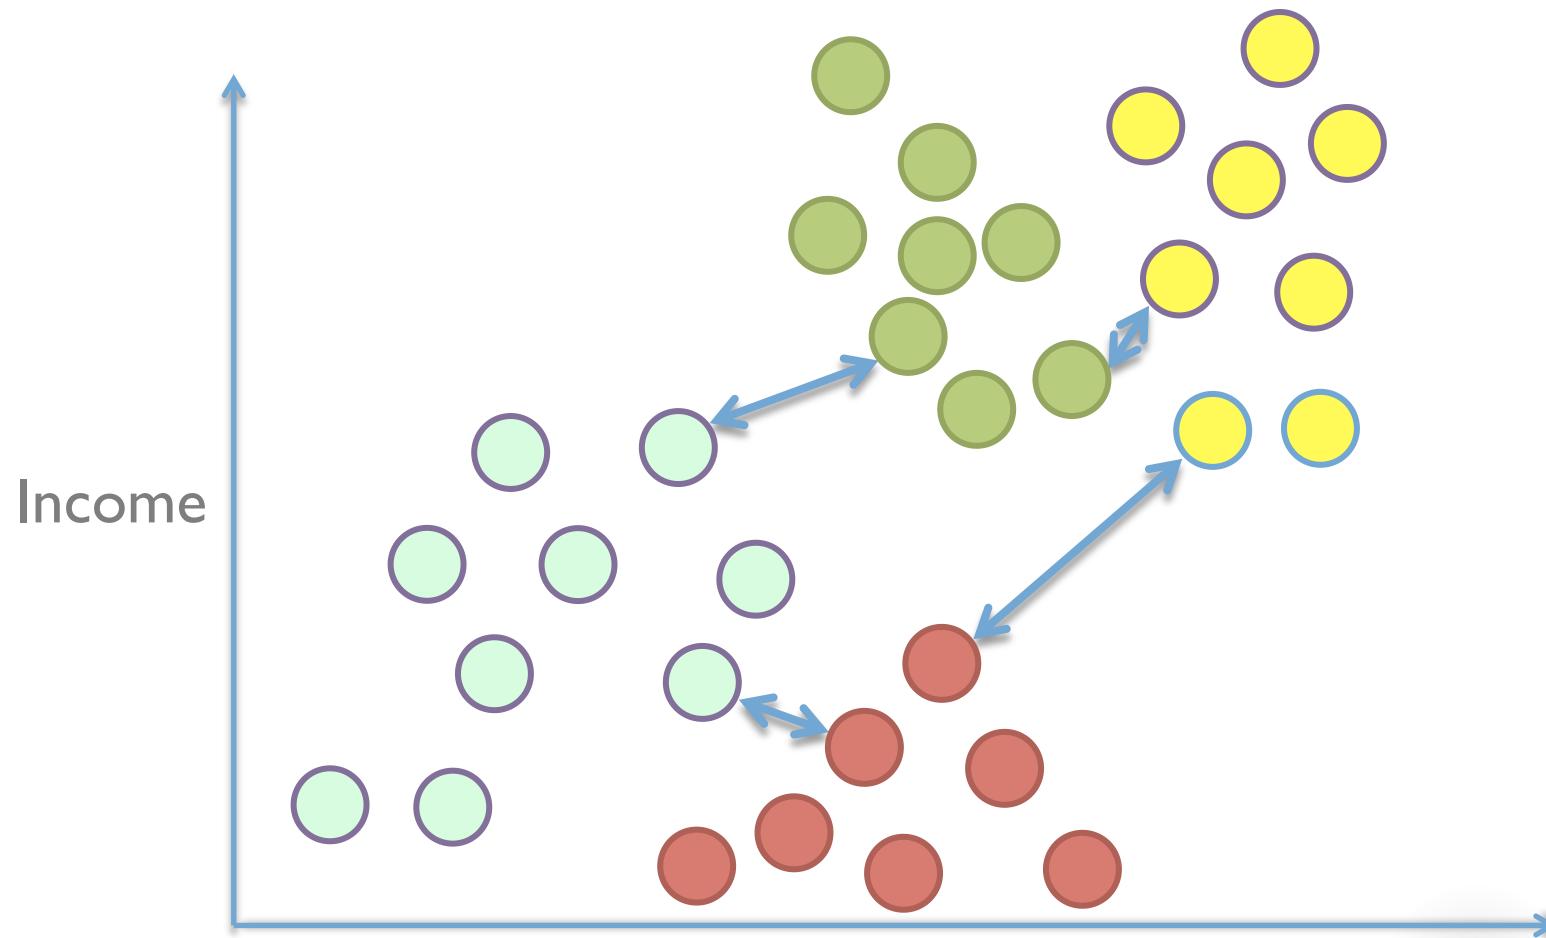
linkage = “single”

Distance between two cluster is the minimum pairwise distance



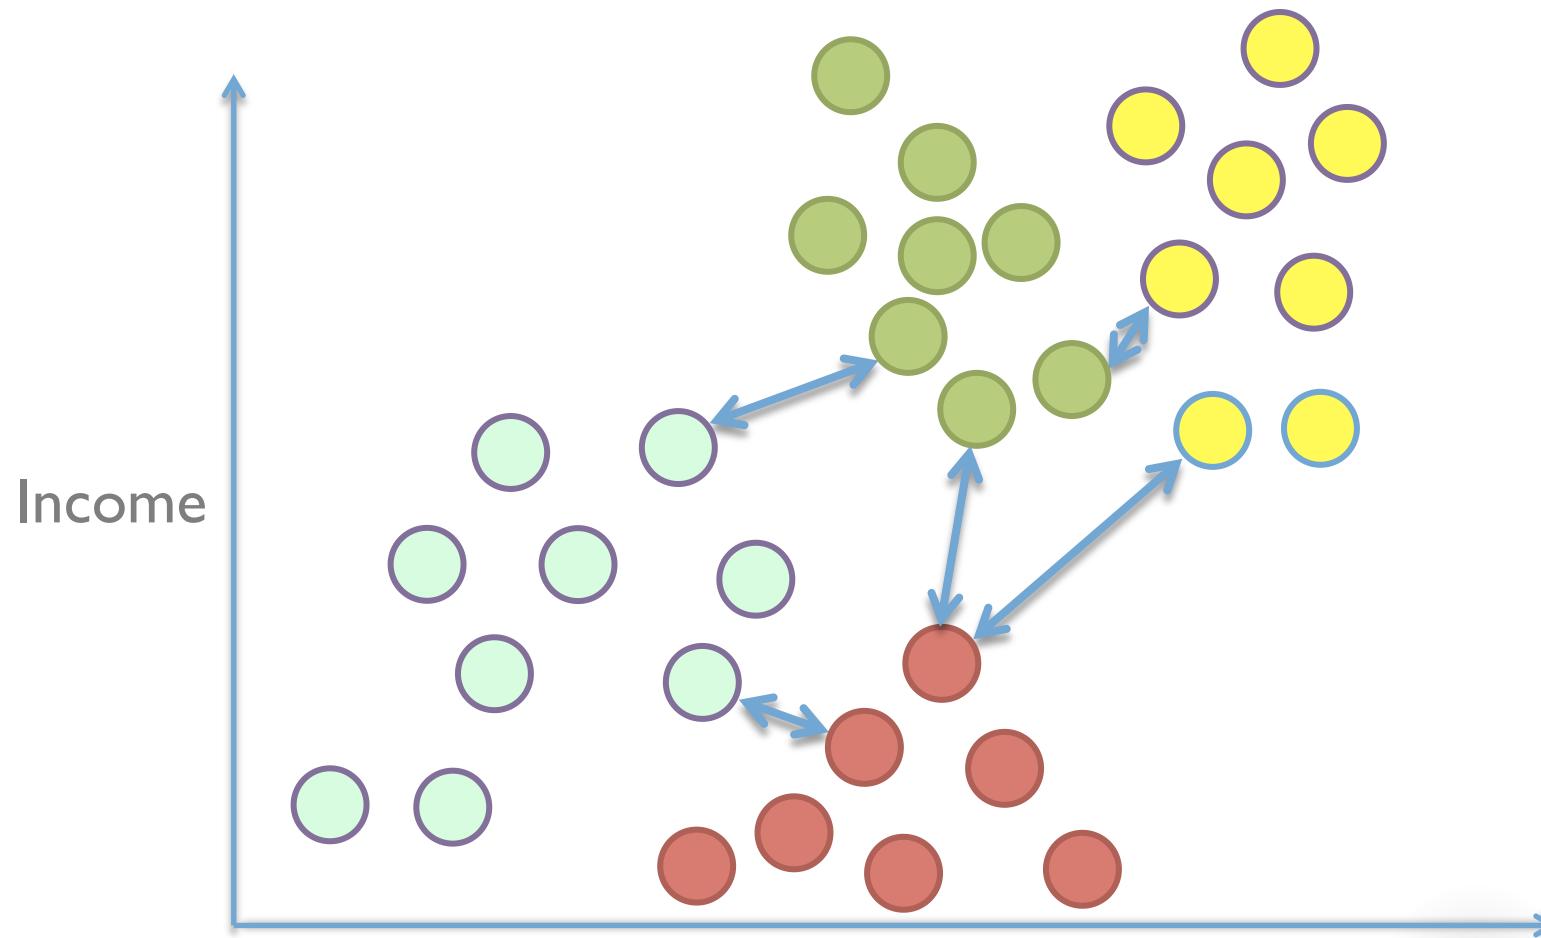
linkage = “single”

Distance between two cluster is the minimum pairwise distance



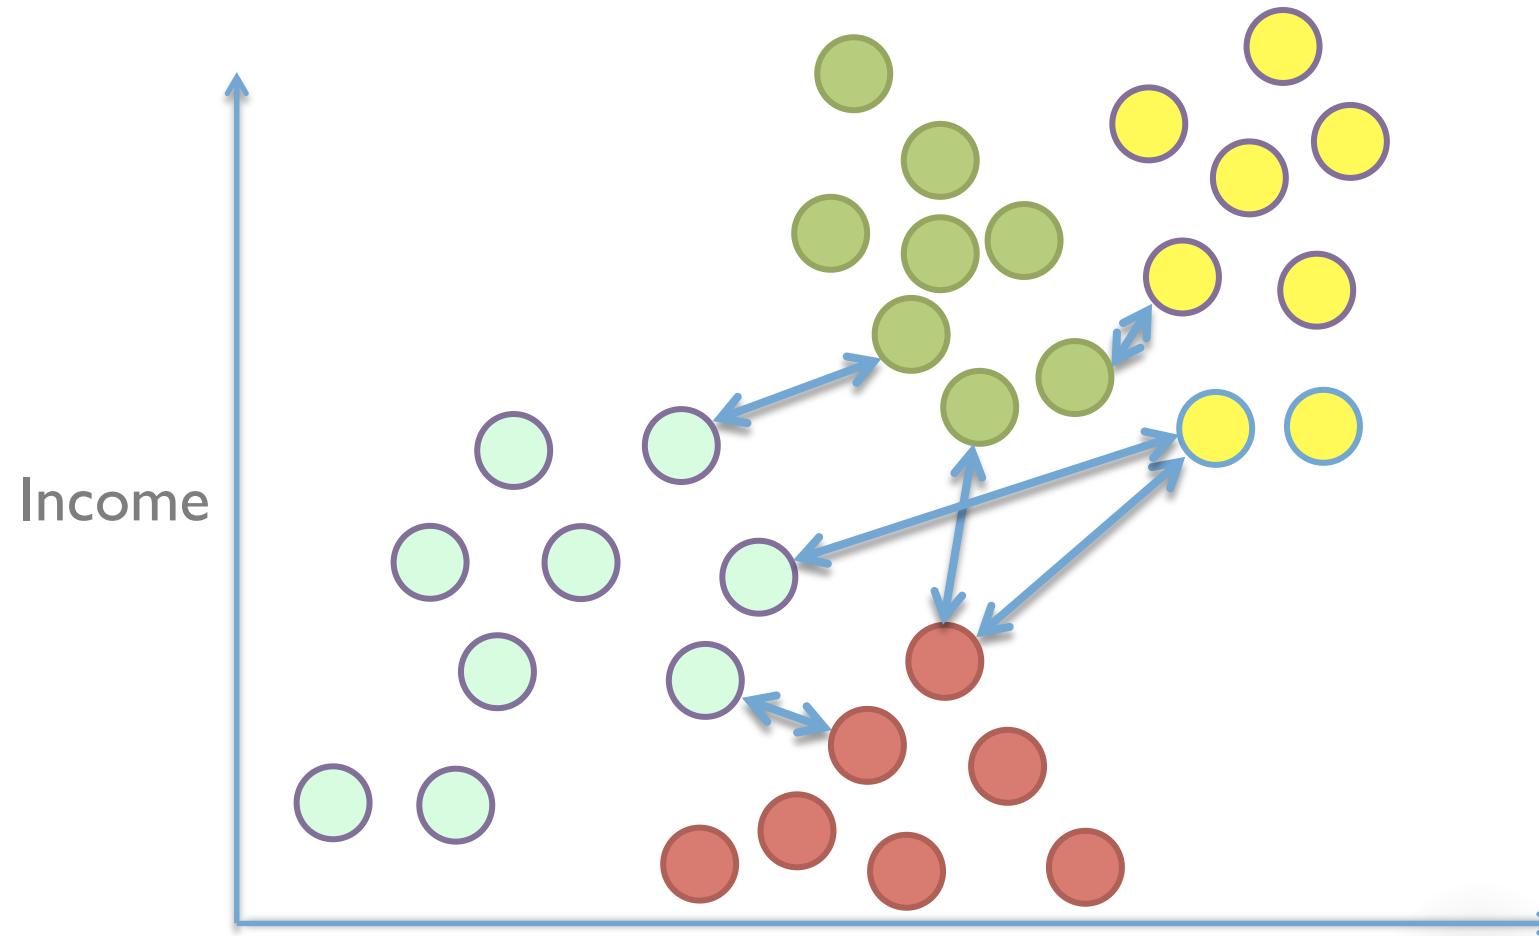
linkage = “single”

Distance between two cluster is the minimum pairwise distance



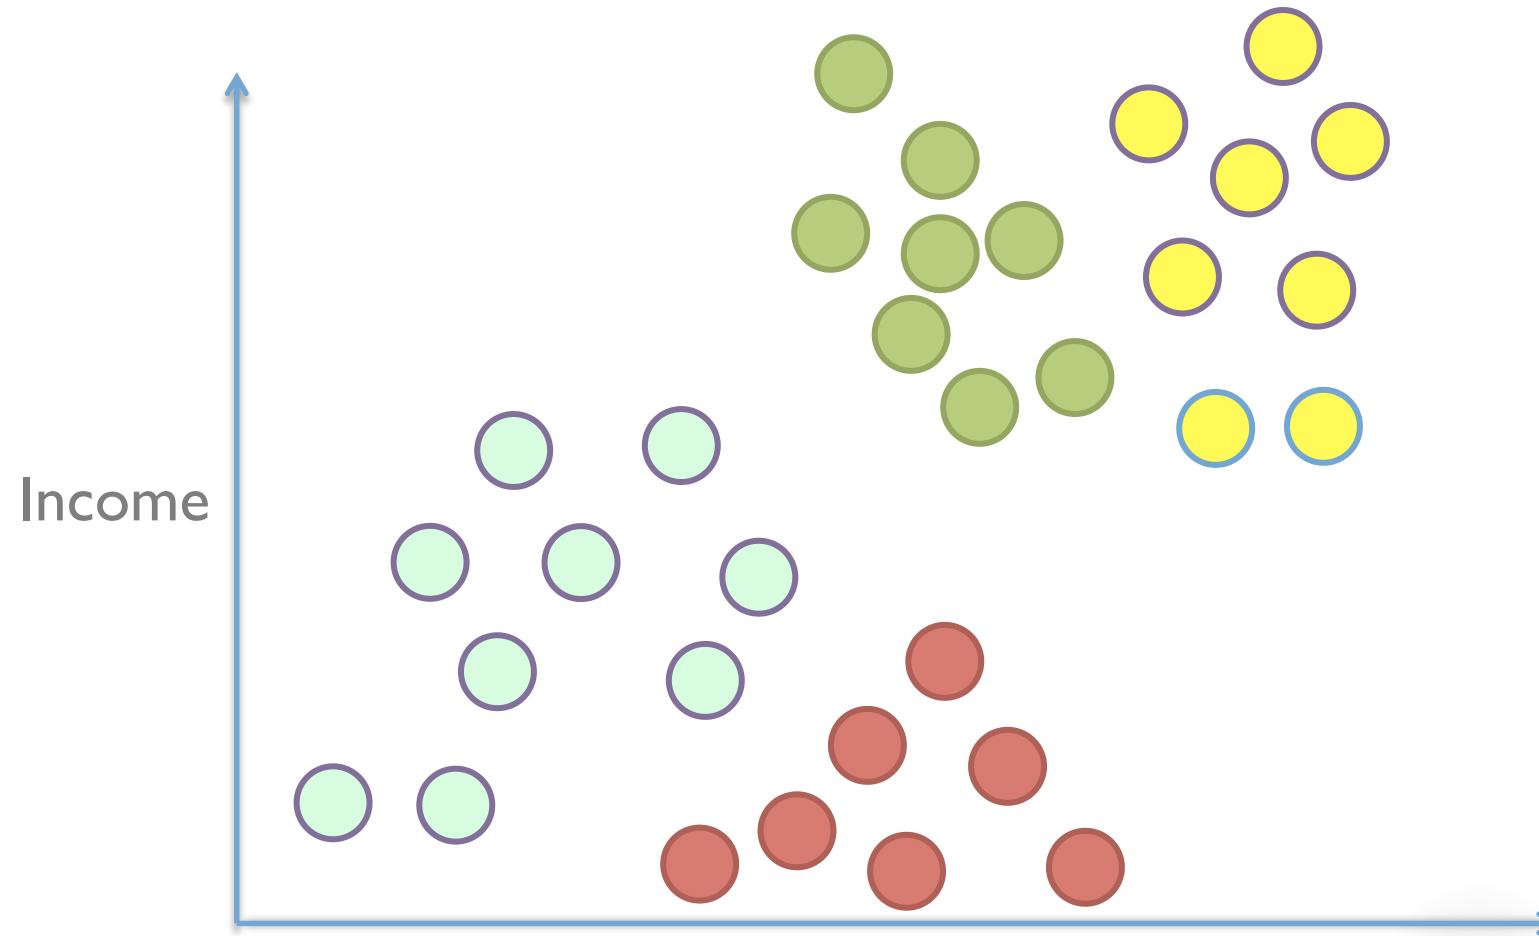
linkage = “single”

Fast. Can form long chains instead of clusters.



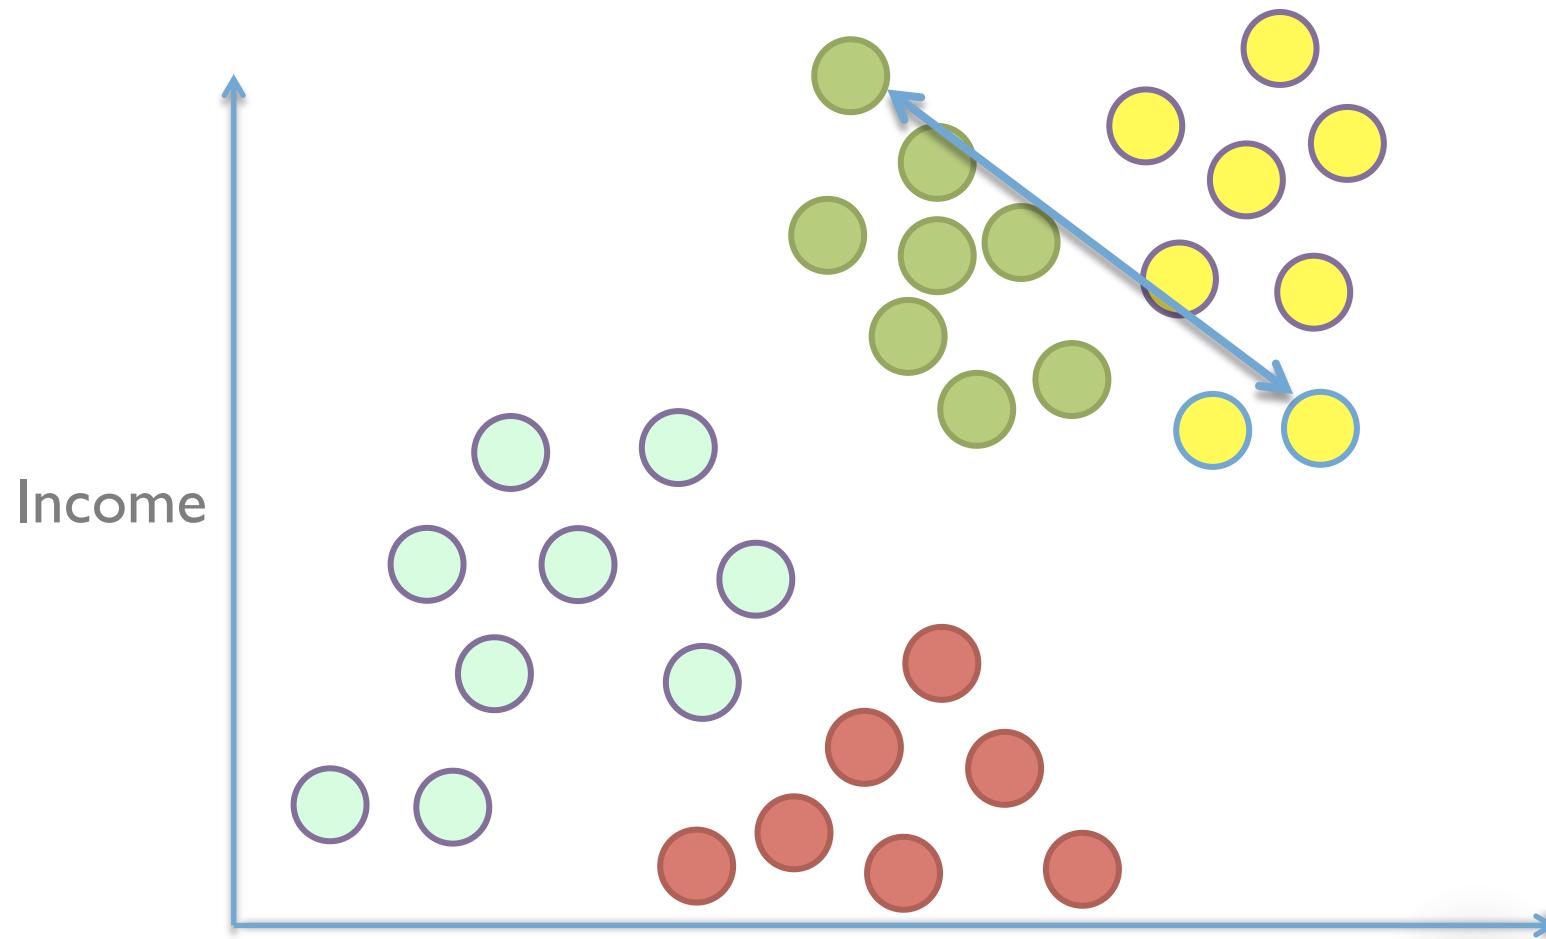
linkage = “complete”

Distance between two clusters is the maximum pairwise distance



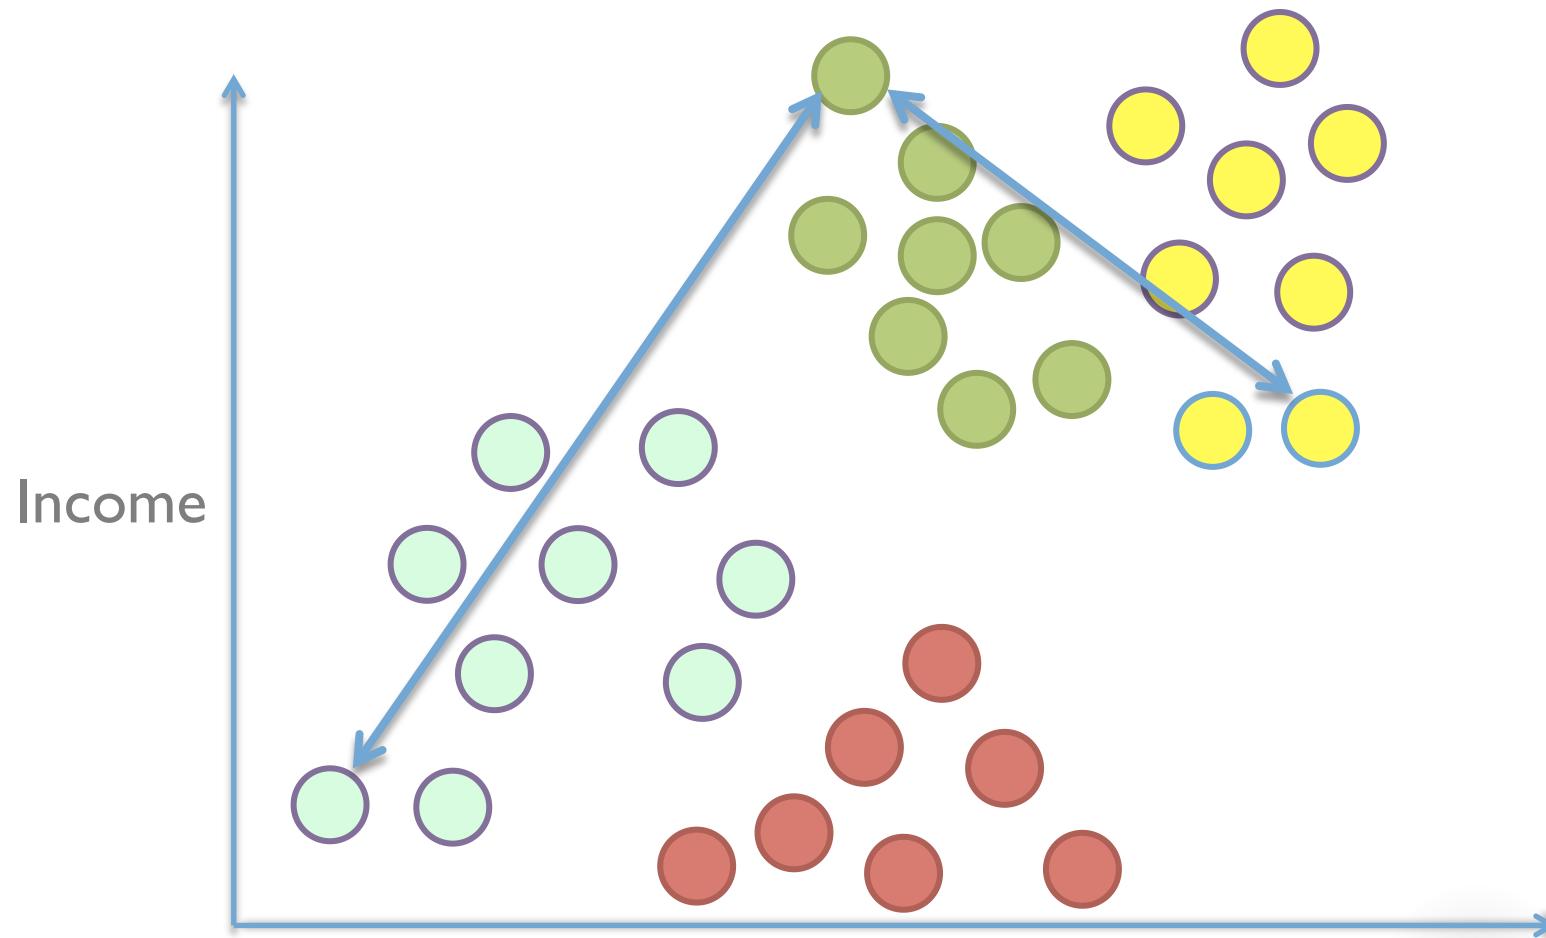
linkage = “complete”

Distance between two clusters is the maximum pairwise distance



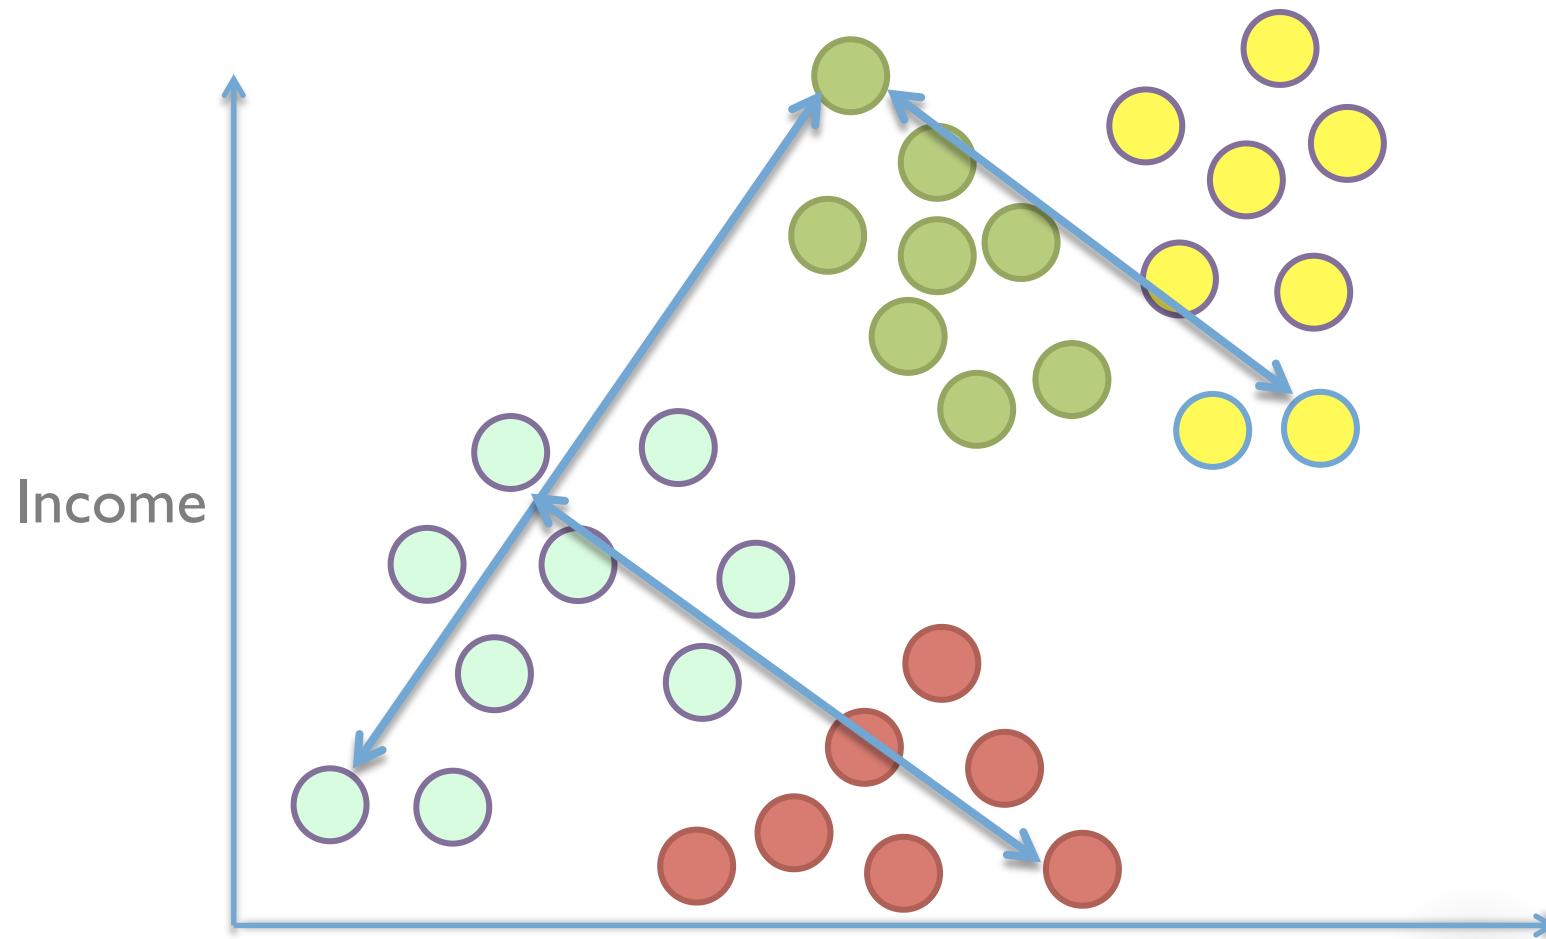
linkage = “complete”

Distance between two clusters is the maximum pairwise distance



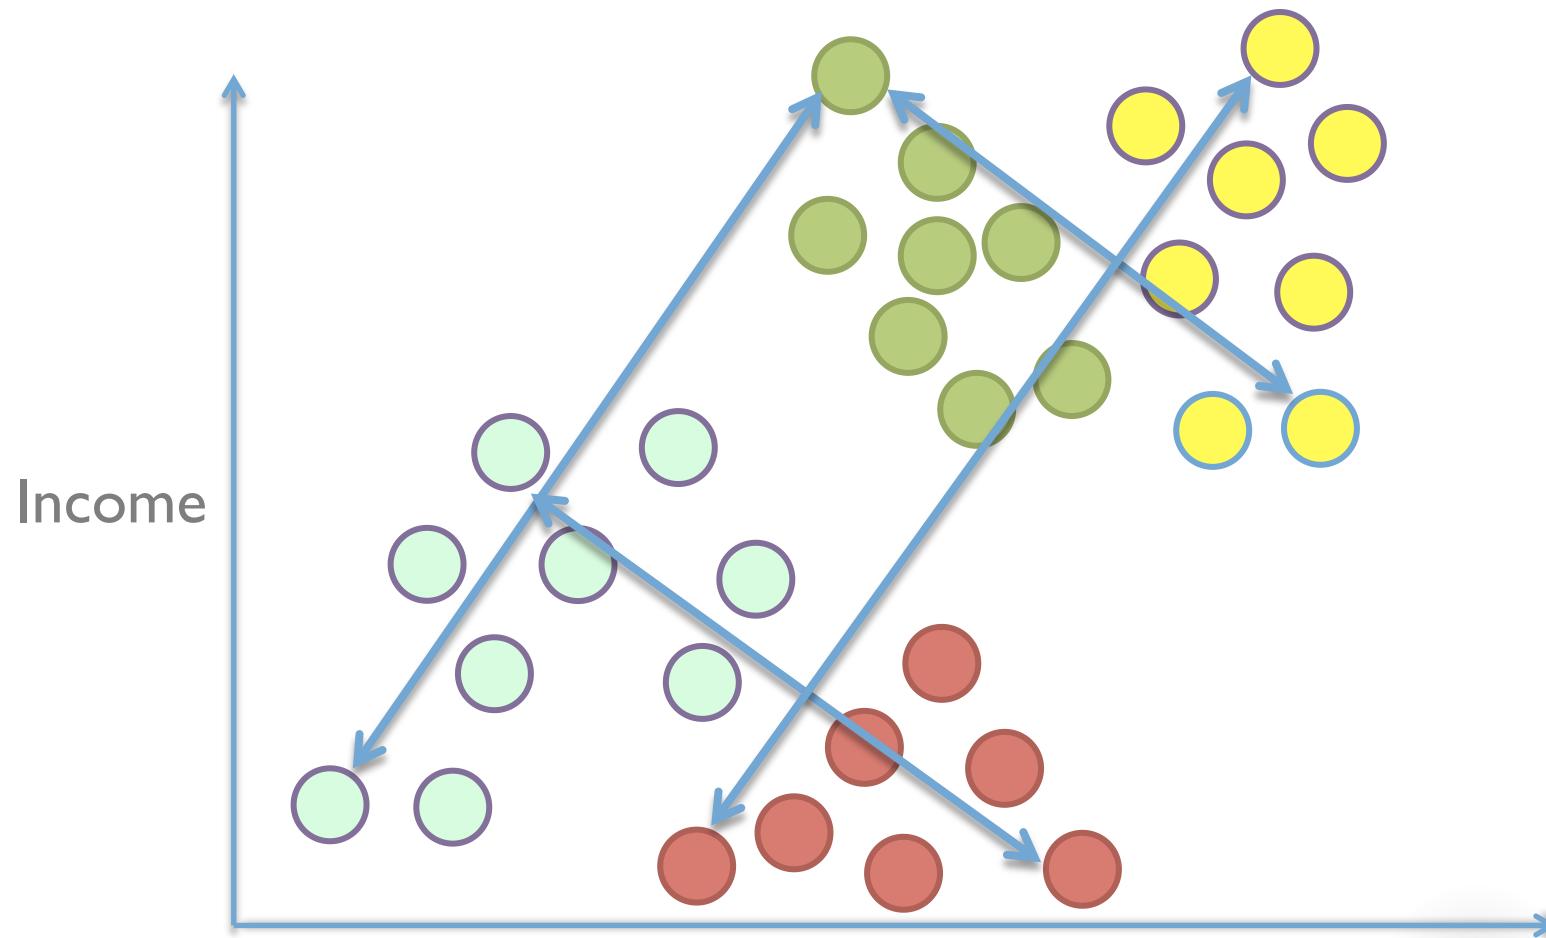
linkage = “complete”

Distance between two clusters is the maximum pairwise distance



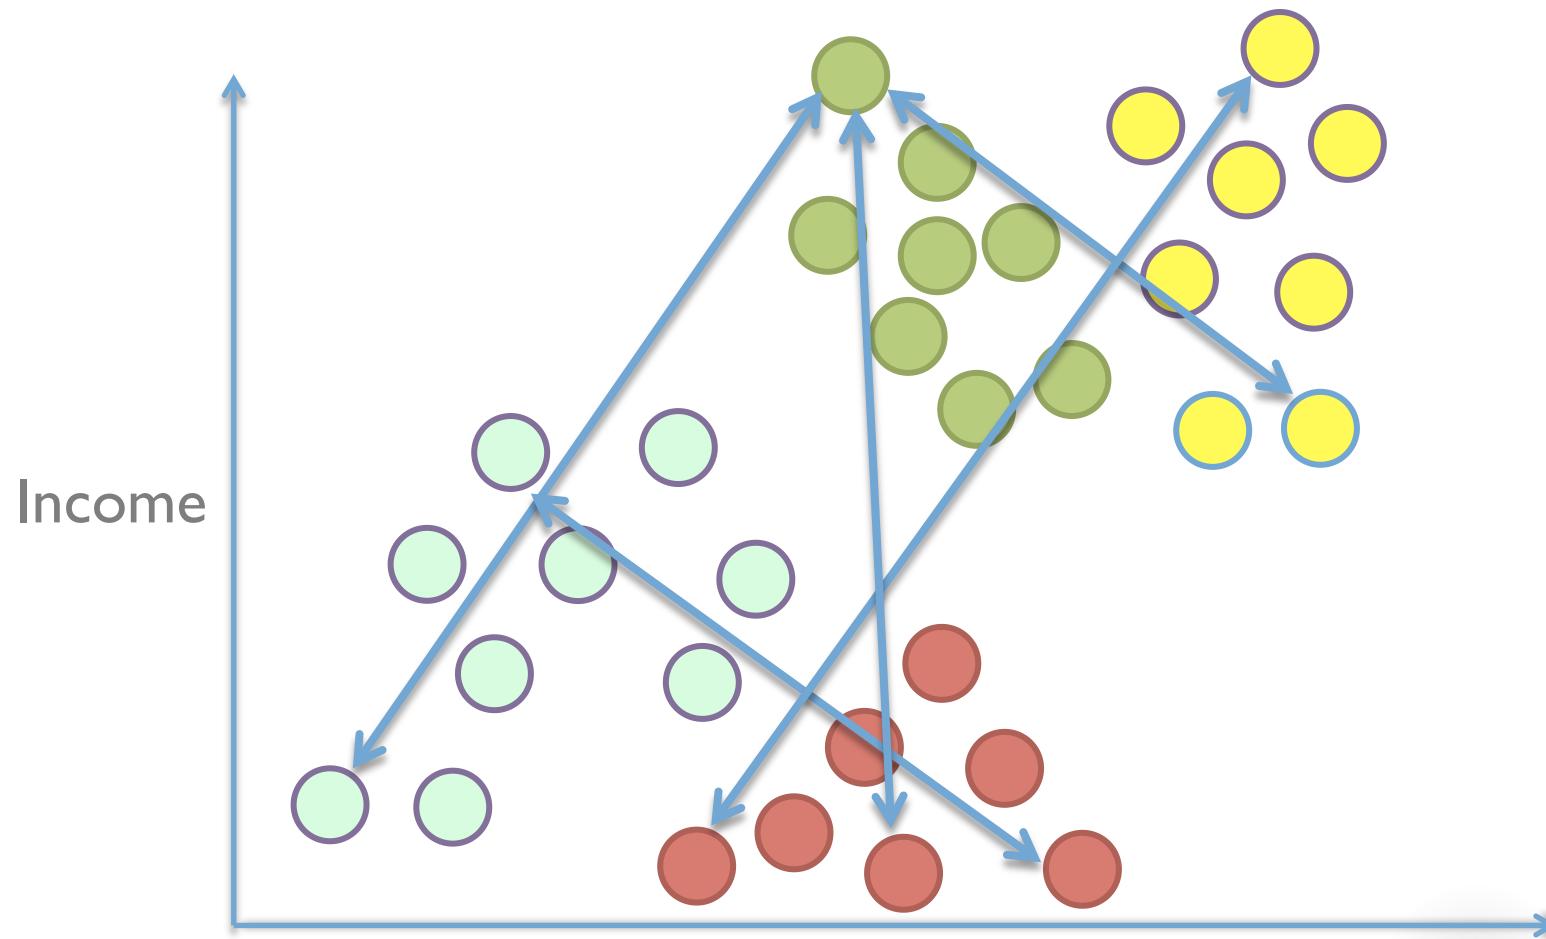
linkage = “complete”

Distance between two clusters is the maximum pairwise distance



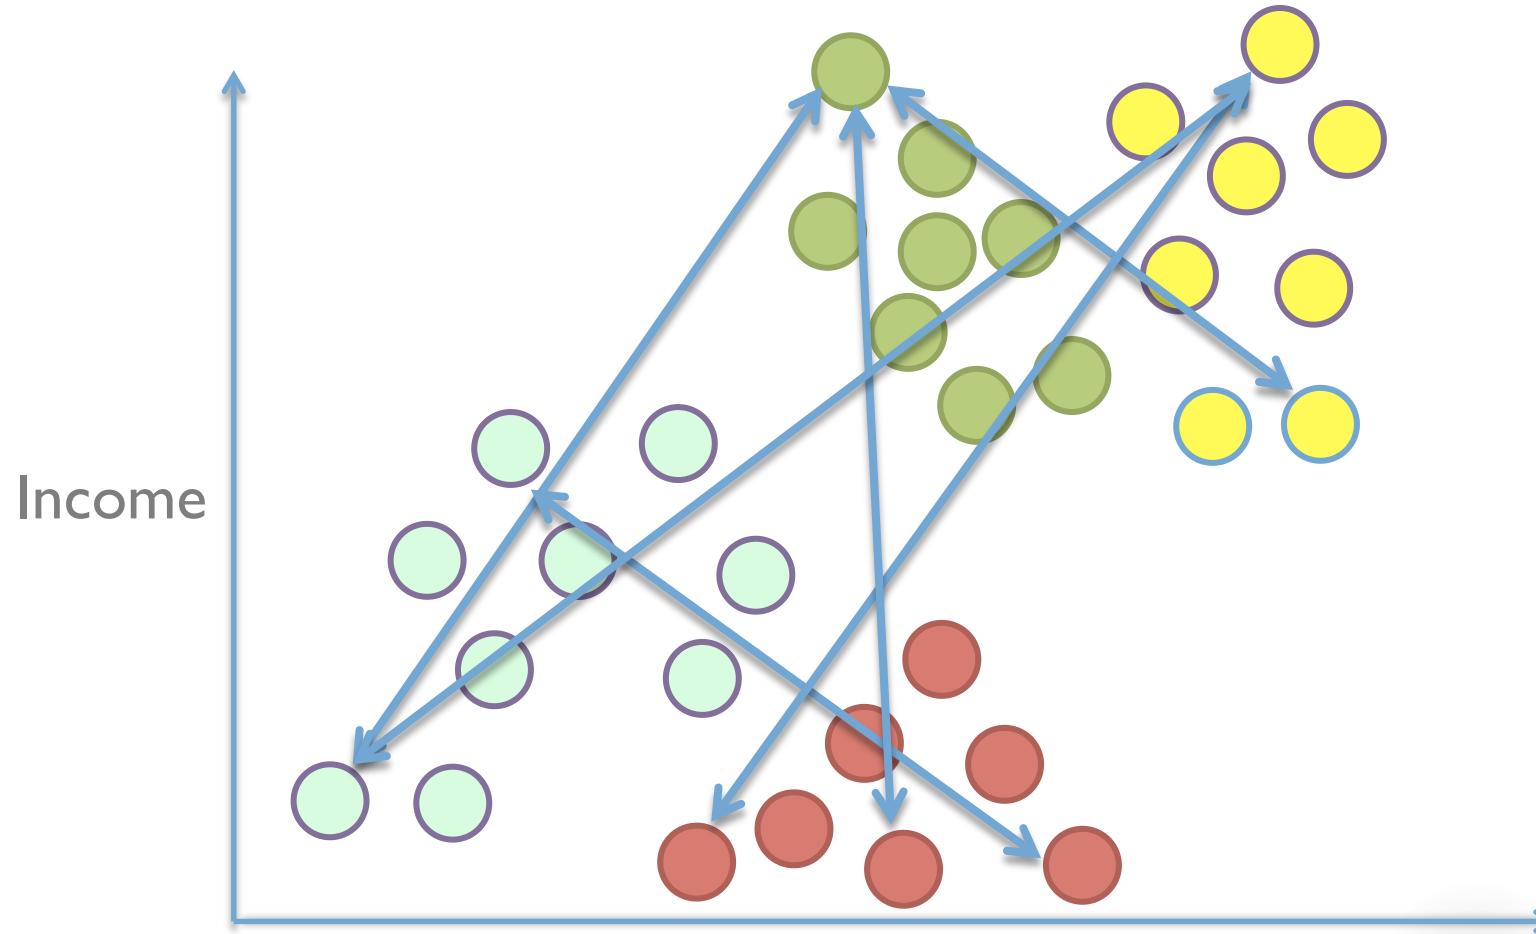
linkage = “complete”

Distance between two clusters is the maximum pairwise distance



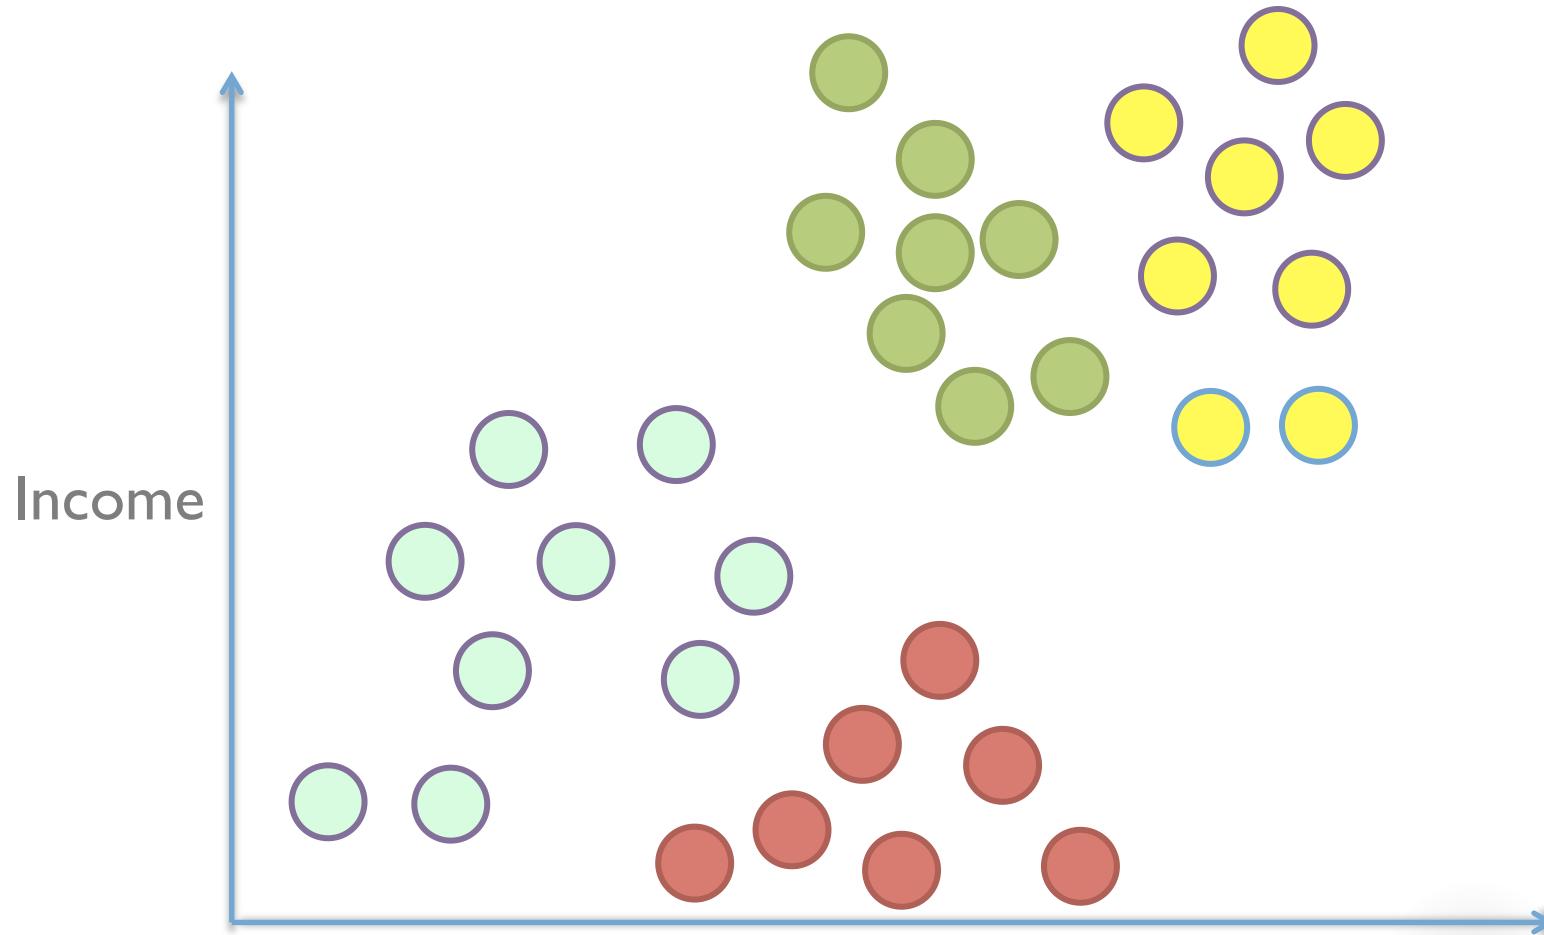
linkage = “complete”

Fast. No chains. Sensitive to outliers.



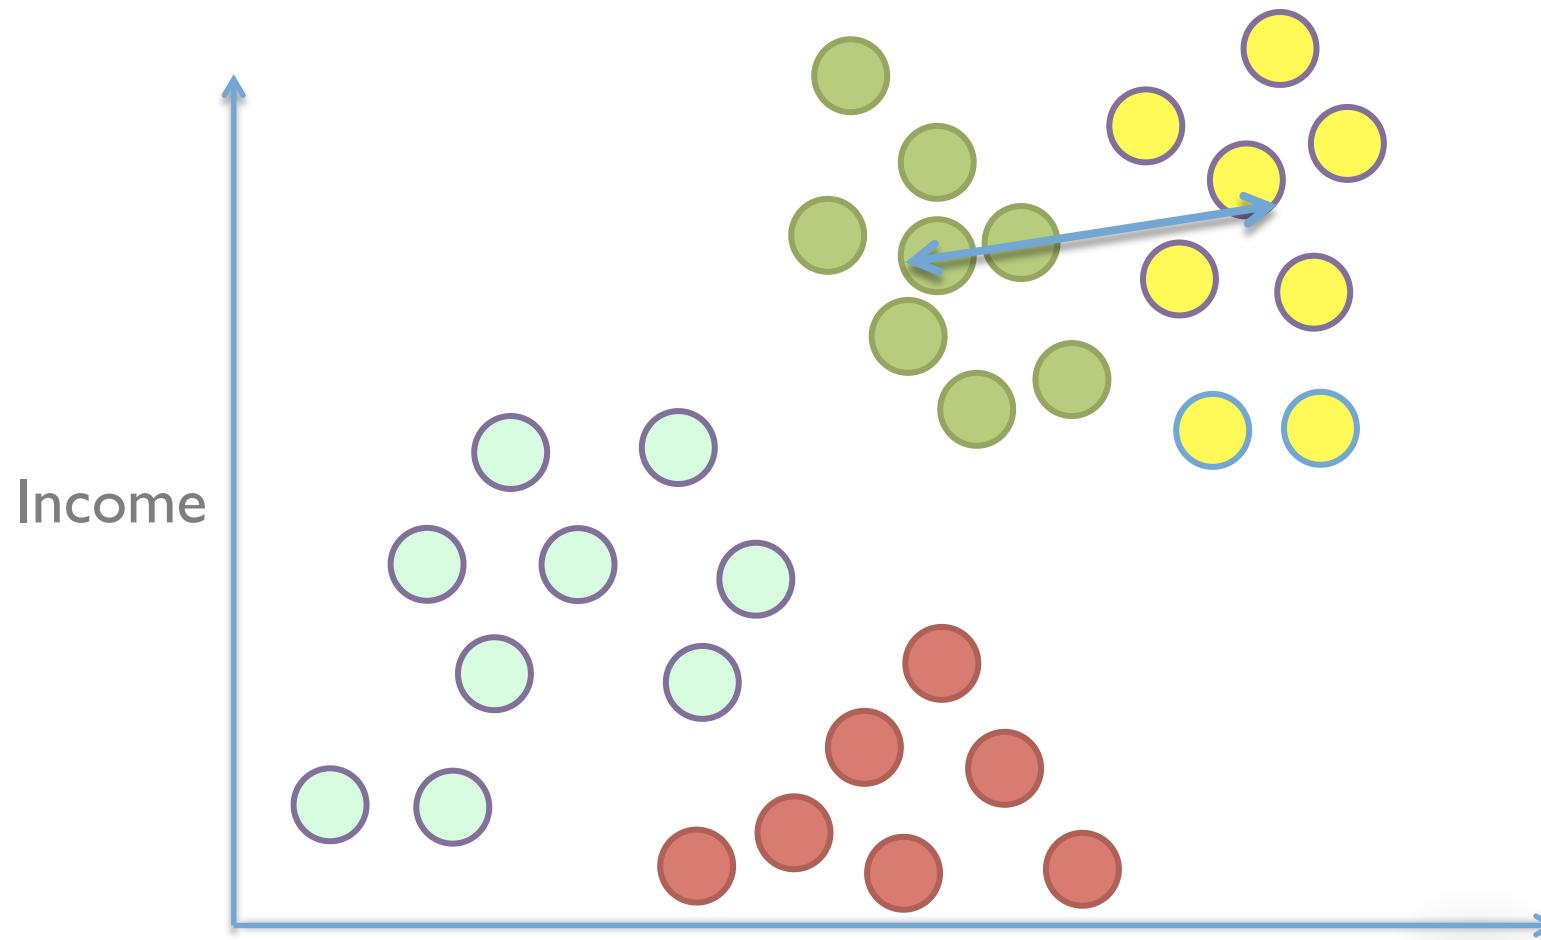
linkage = “average”

Distance between two cluster is average of all pairwise distance



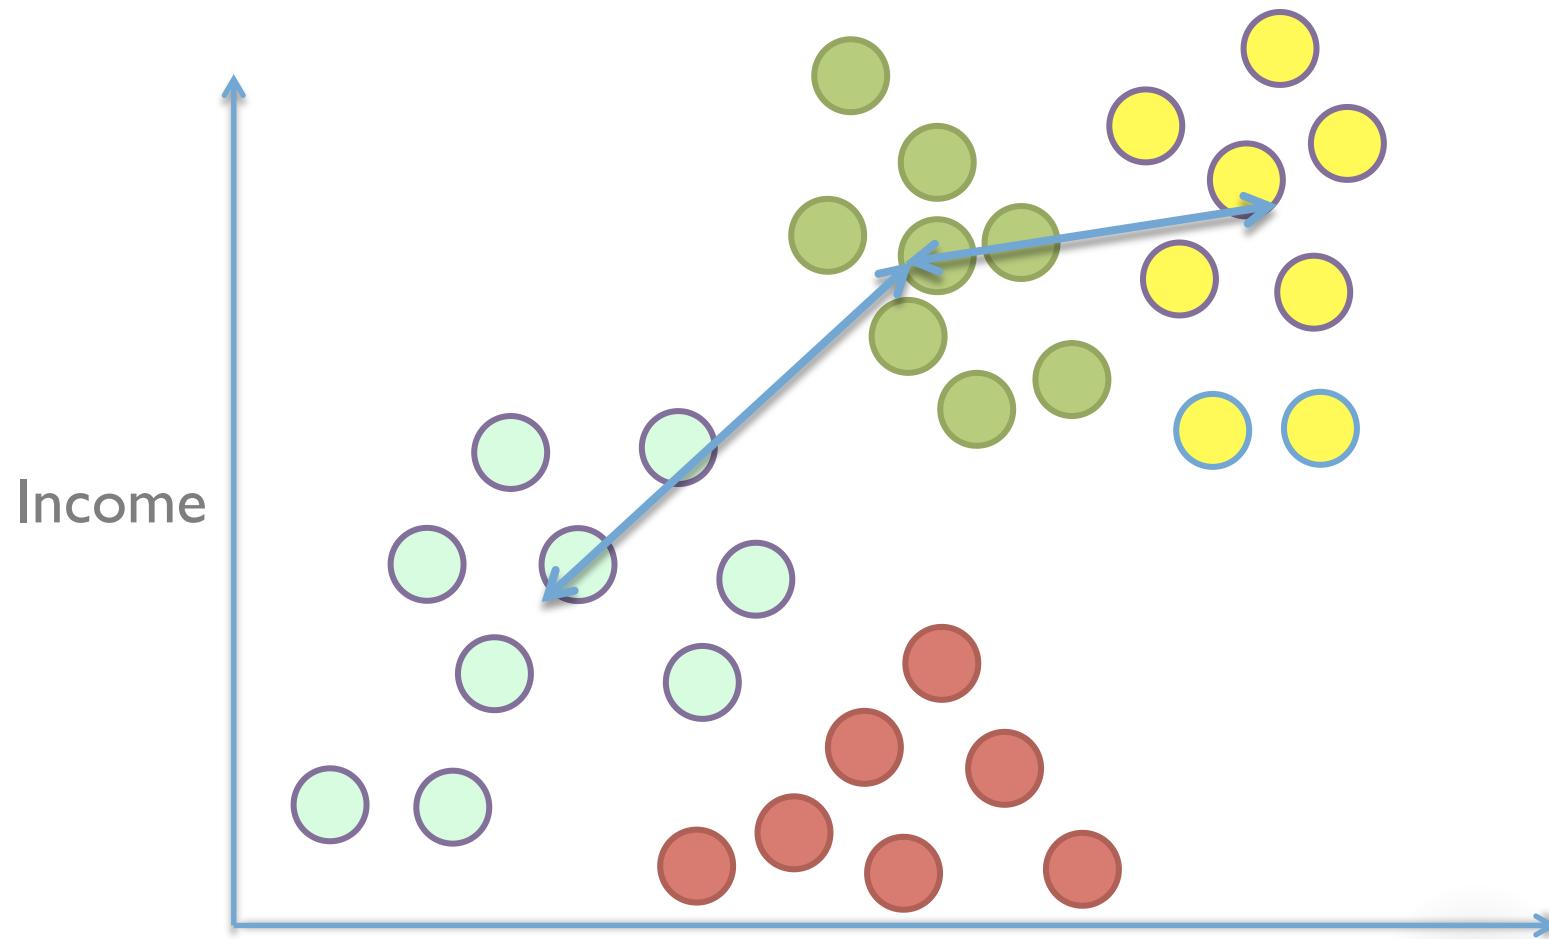
linkage = “average”

Distance between two cluster is average of all pairwise distance



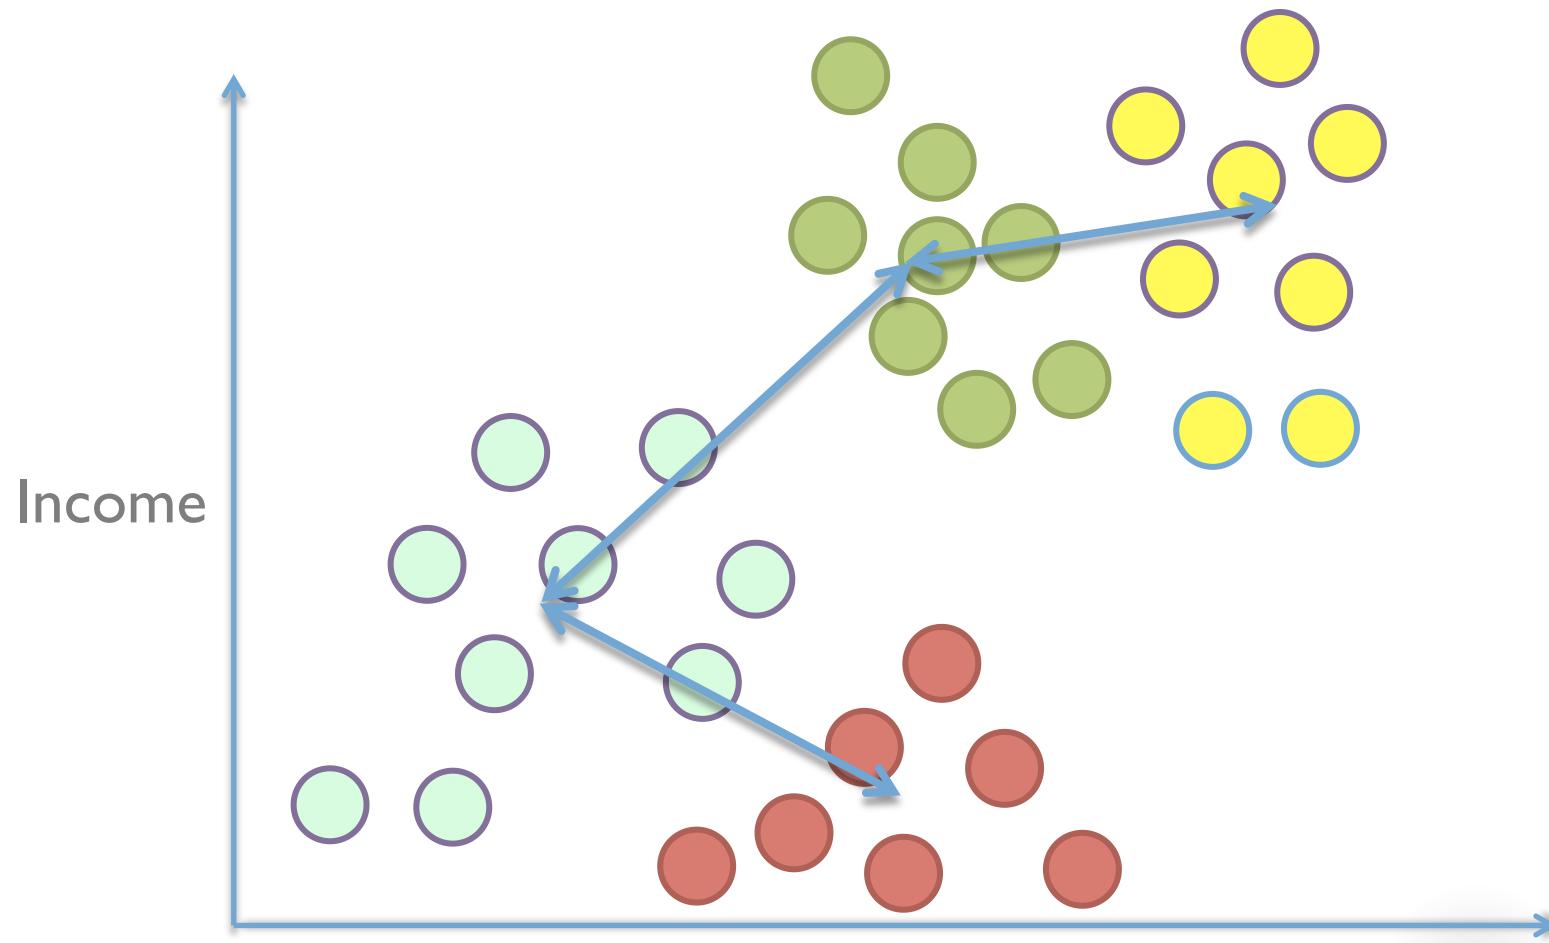
linkage = “average”

Distance between two cluster is average of all pairwise distance



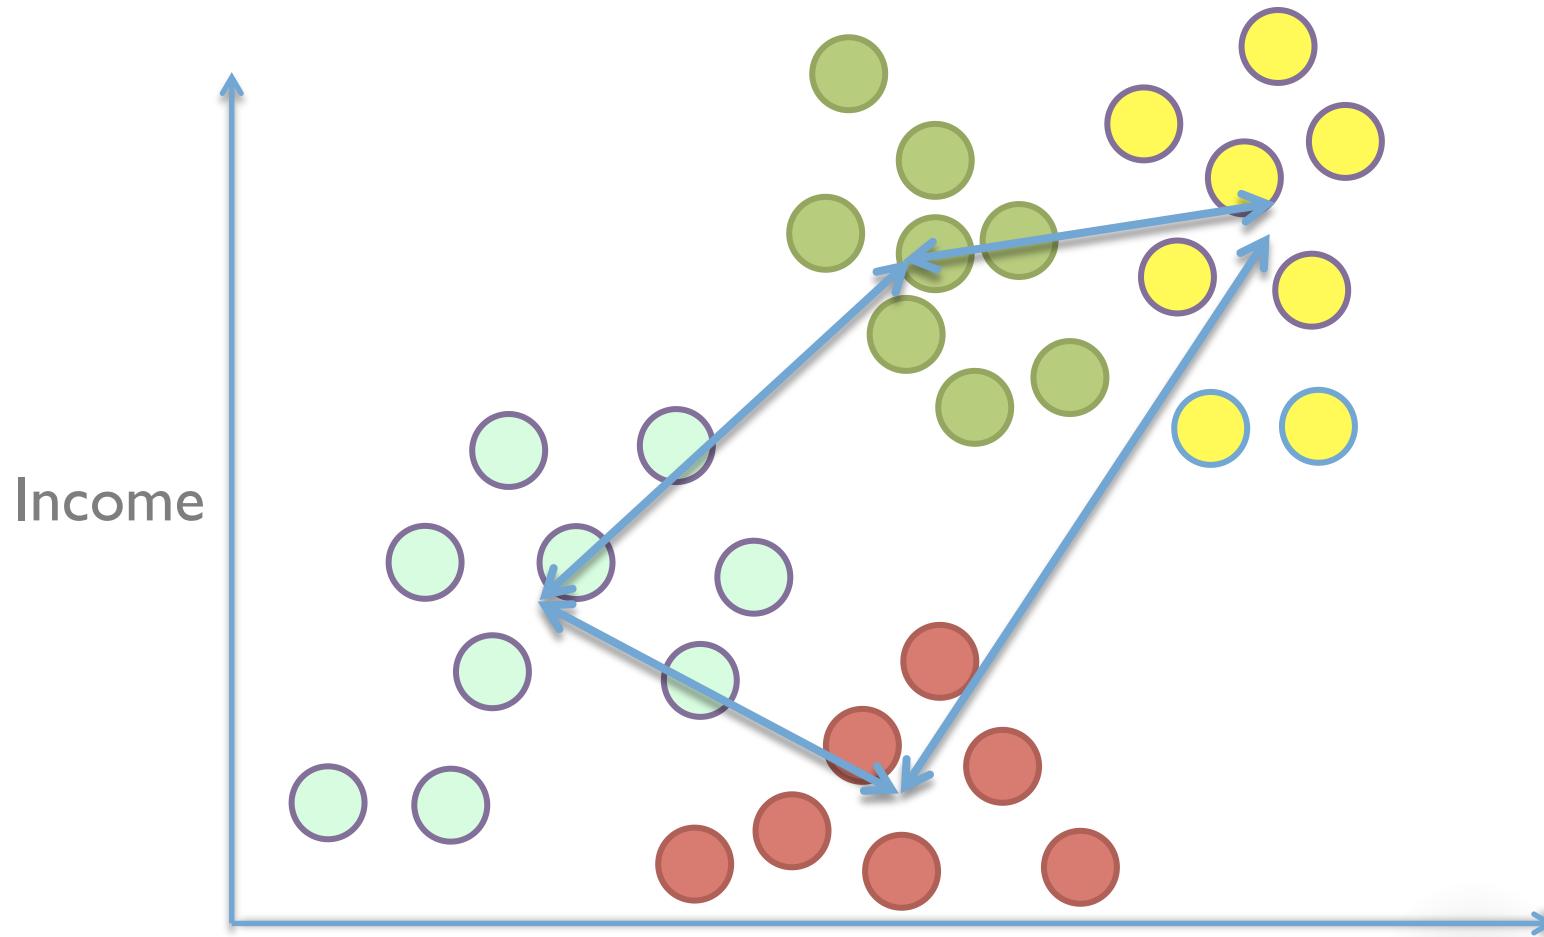
linkage = “average”

Distance between two cluster is average of all pairwise distance



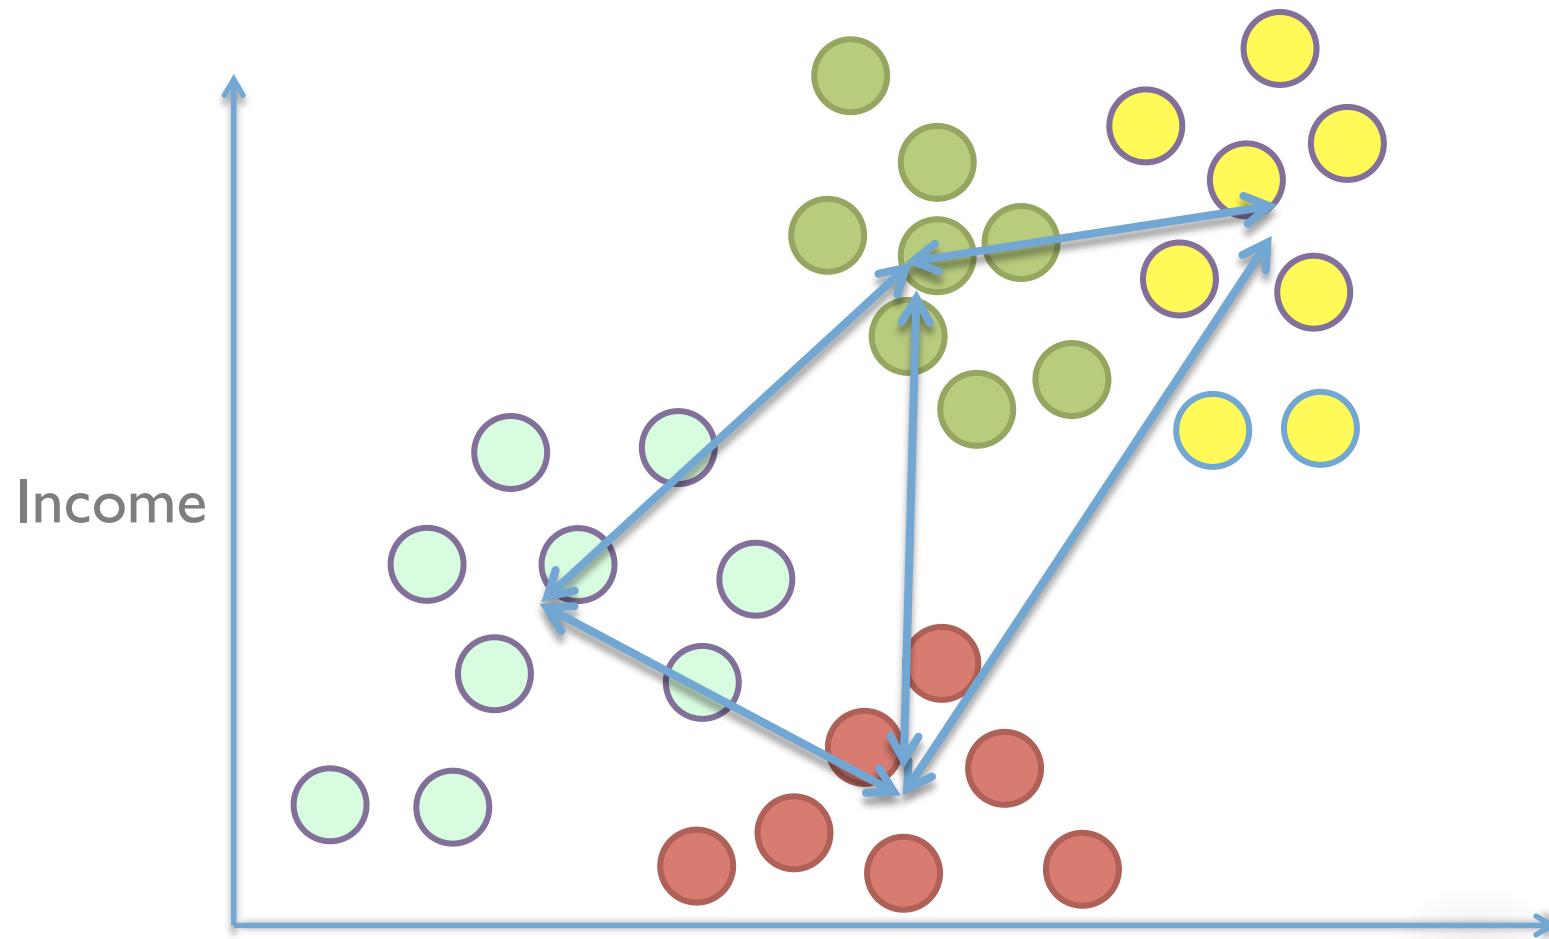
linkage = “average”

Distance between two cluster is average of all pairwise distance



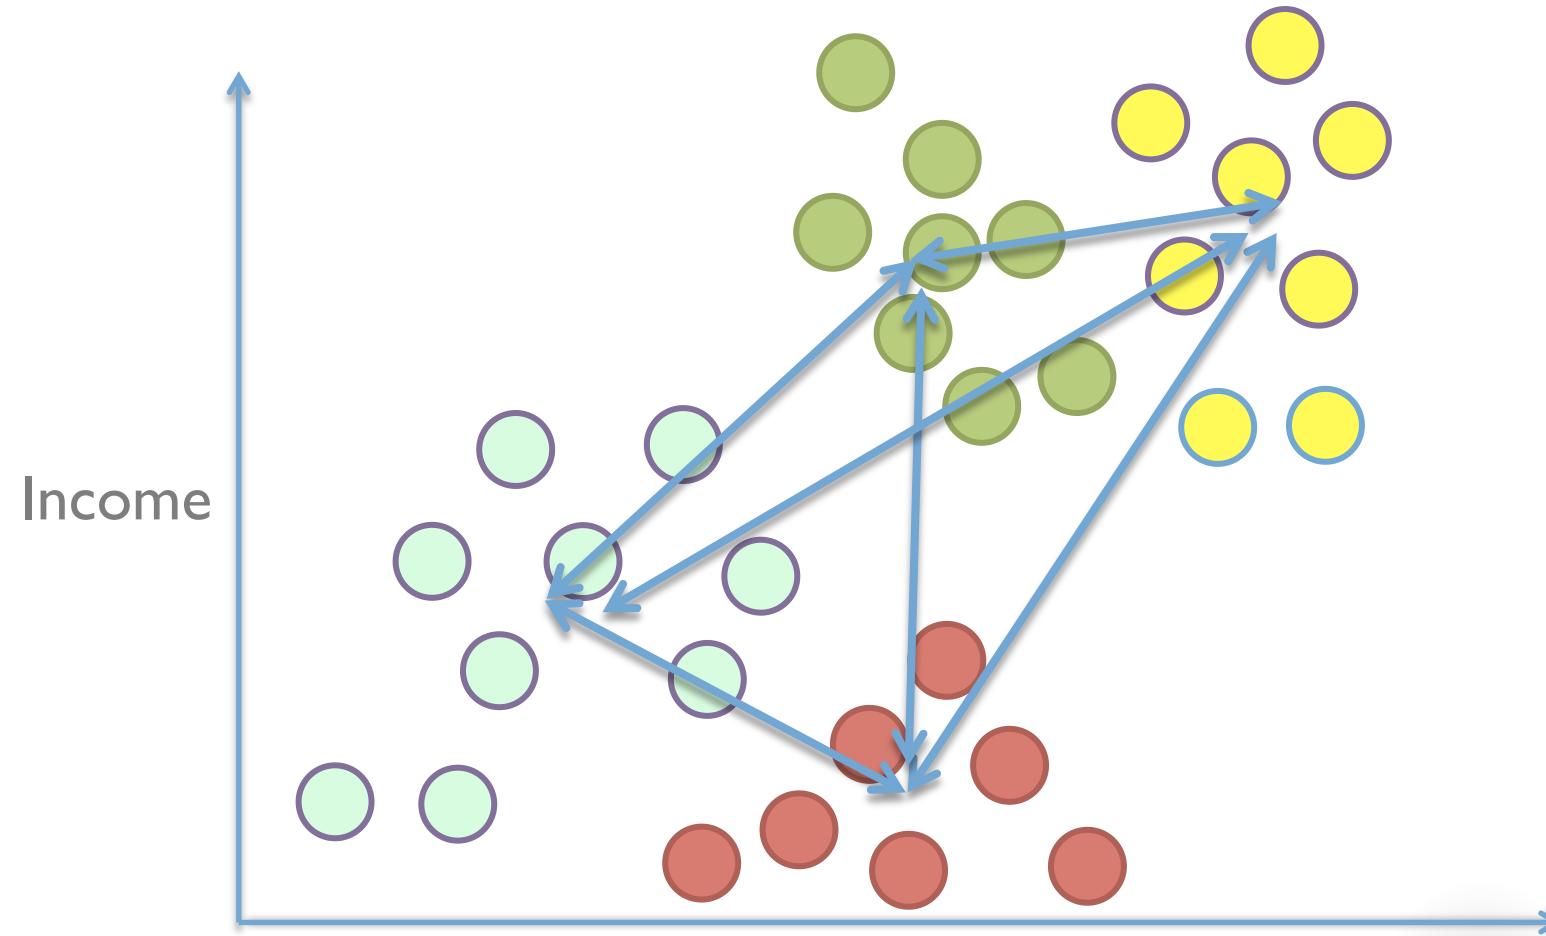
linkage = “average”

Distance between two cluster is average of all pairwise distance



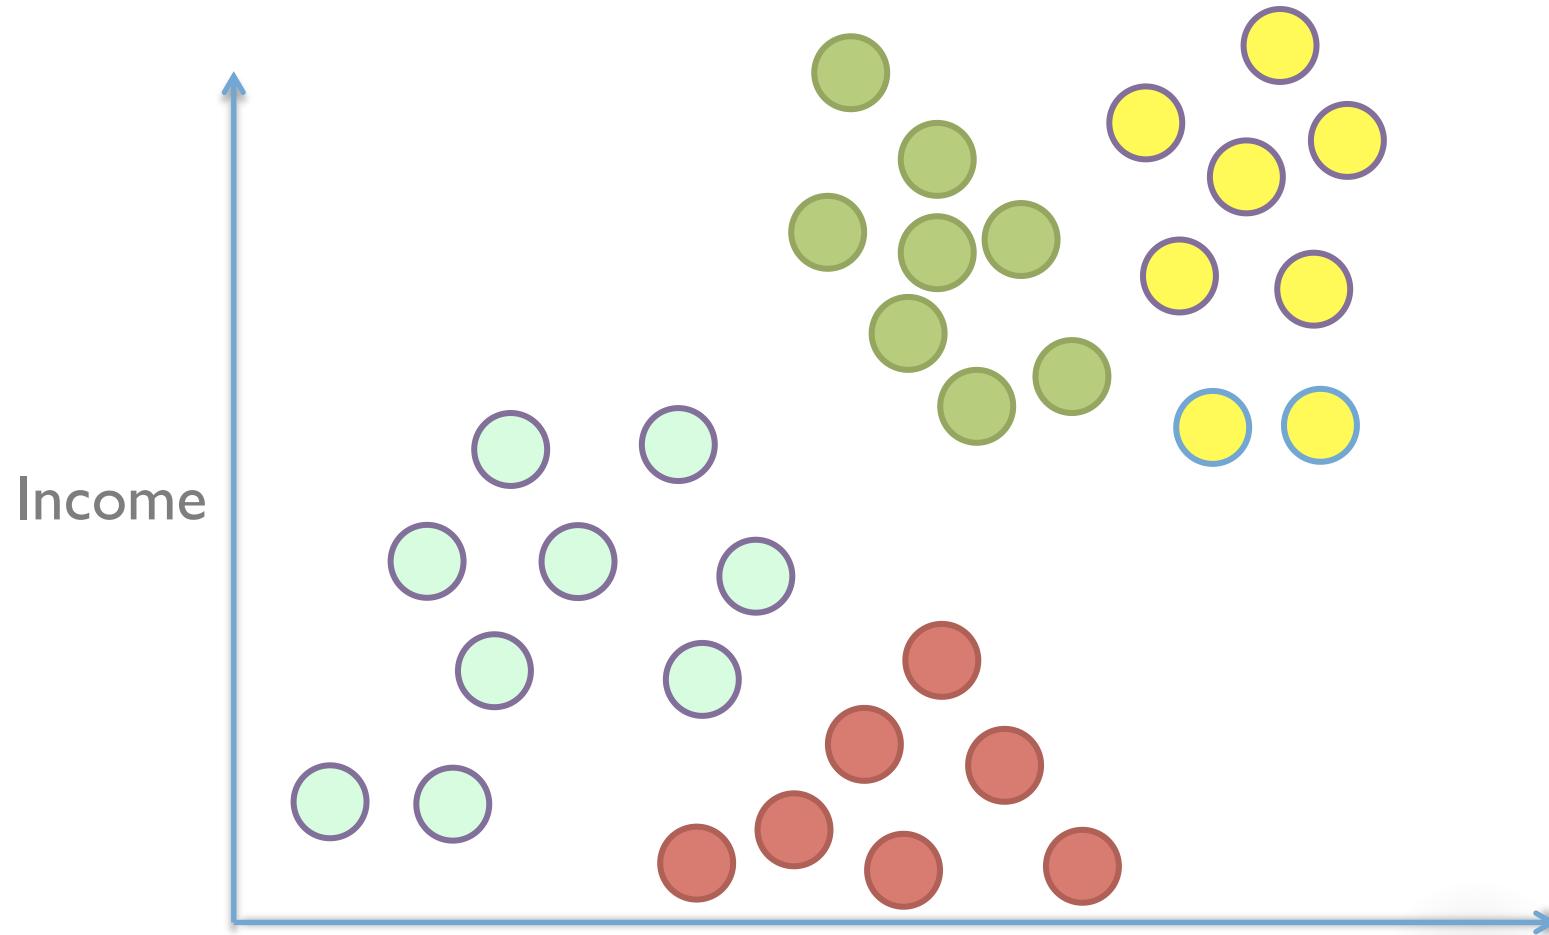
linkage = “average”

Slower. No chains. Not sensitive to outliers.



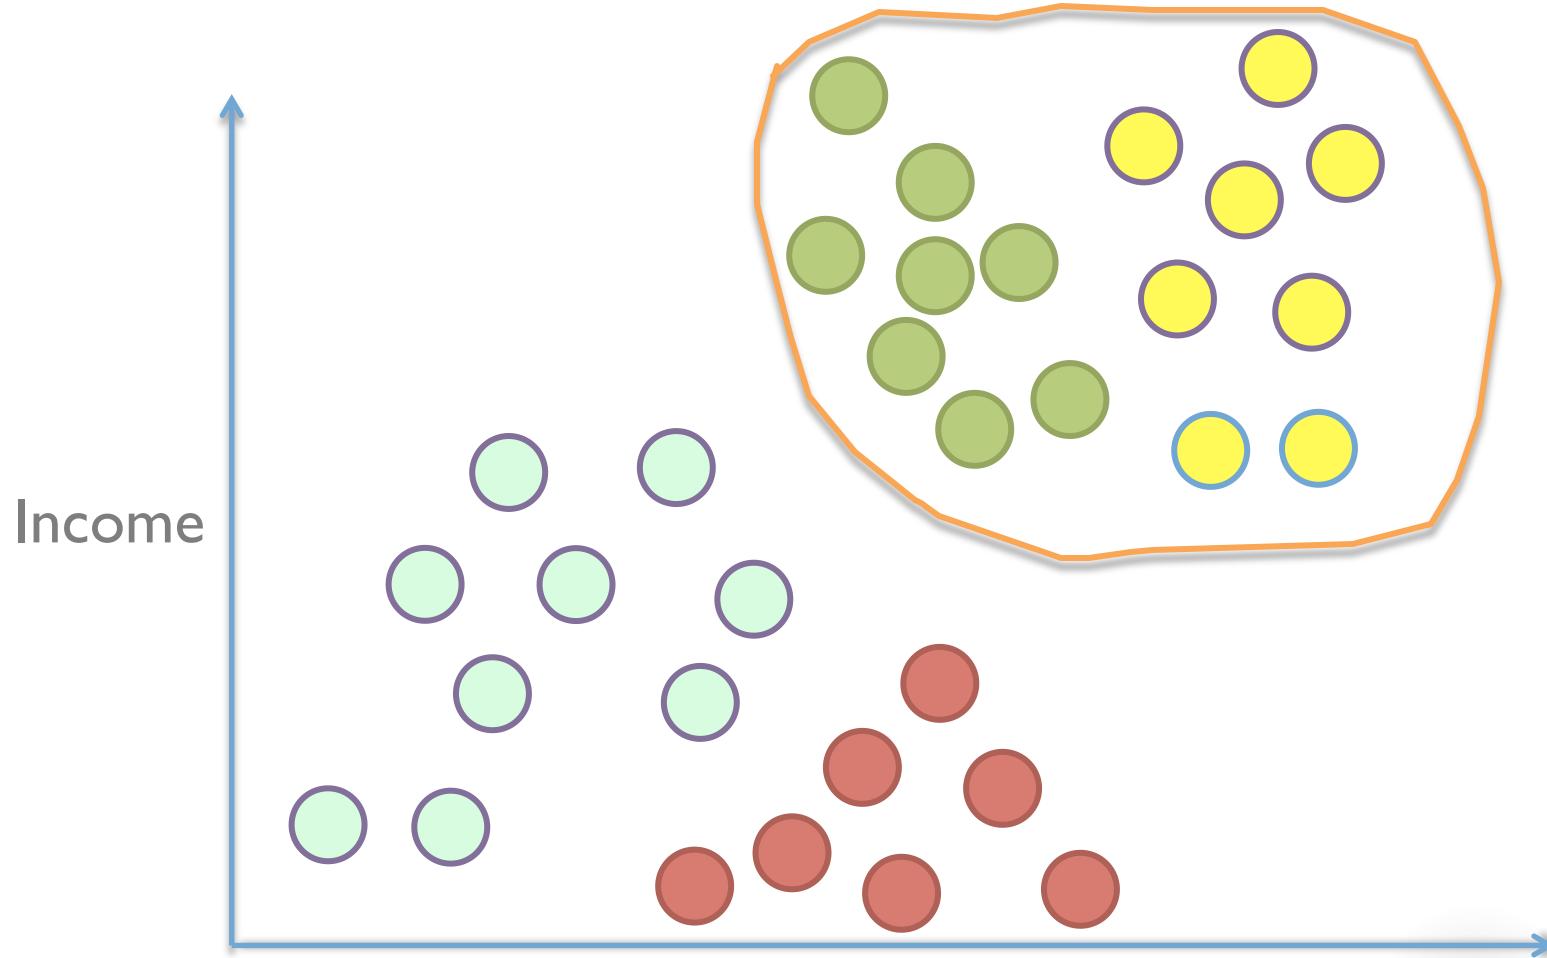
linkage = “ward”

Don’t merge based on distance. Merge the two whose combination will give the best score (lowest inertia).



linkage = “ward”

This will choose the merge with the minimum within-cluster variance (high density clusters).



```
from sklearn.cluster import AgglomerativeClustering  
Agglomerative Clustering(linkage="Ward").fit_predict(X)
```

or

```
from sklearn.cluster import Ward  
Ward().fit_predict(X)
```

(Same thing)



```
from sklearn.cluster import AgglomerativeClustering
```

```
AgglomerativeClustering(....,
```

```
....
```

```
,
```

```
linkage = "ward",
```

```
....
```

```
,
```

```
affinity = "euclidean",
```

```
....
```

```
)
```



```
from sklearn.cluster import AgglomerativeClustering
```

```
AgglomerativeClustering(....,
```

```
....
```

```
,
```

```
linkage = "ward",
```

```
....
```

```
,
```

```
affinity = "euclidian",
```

```
....
```

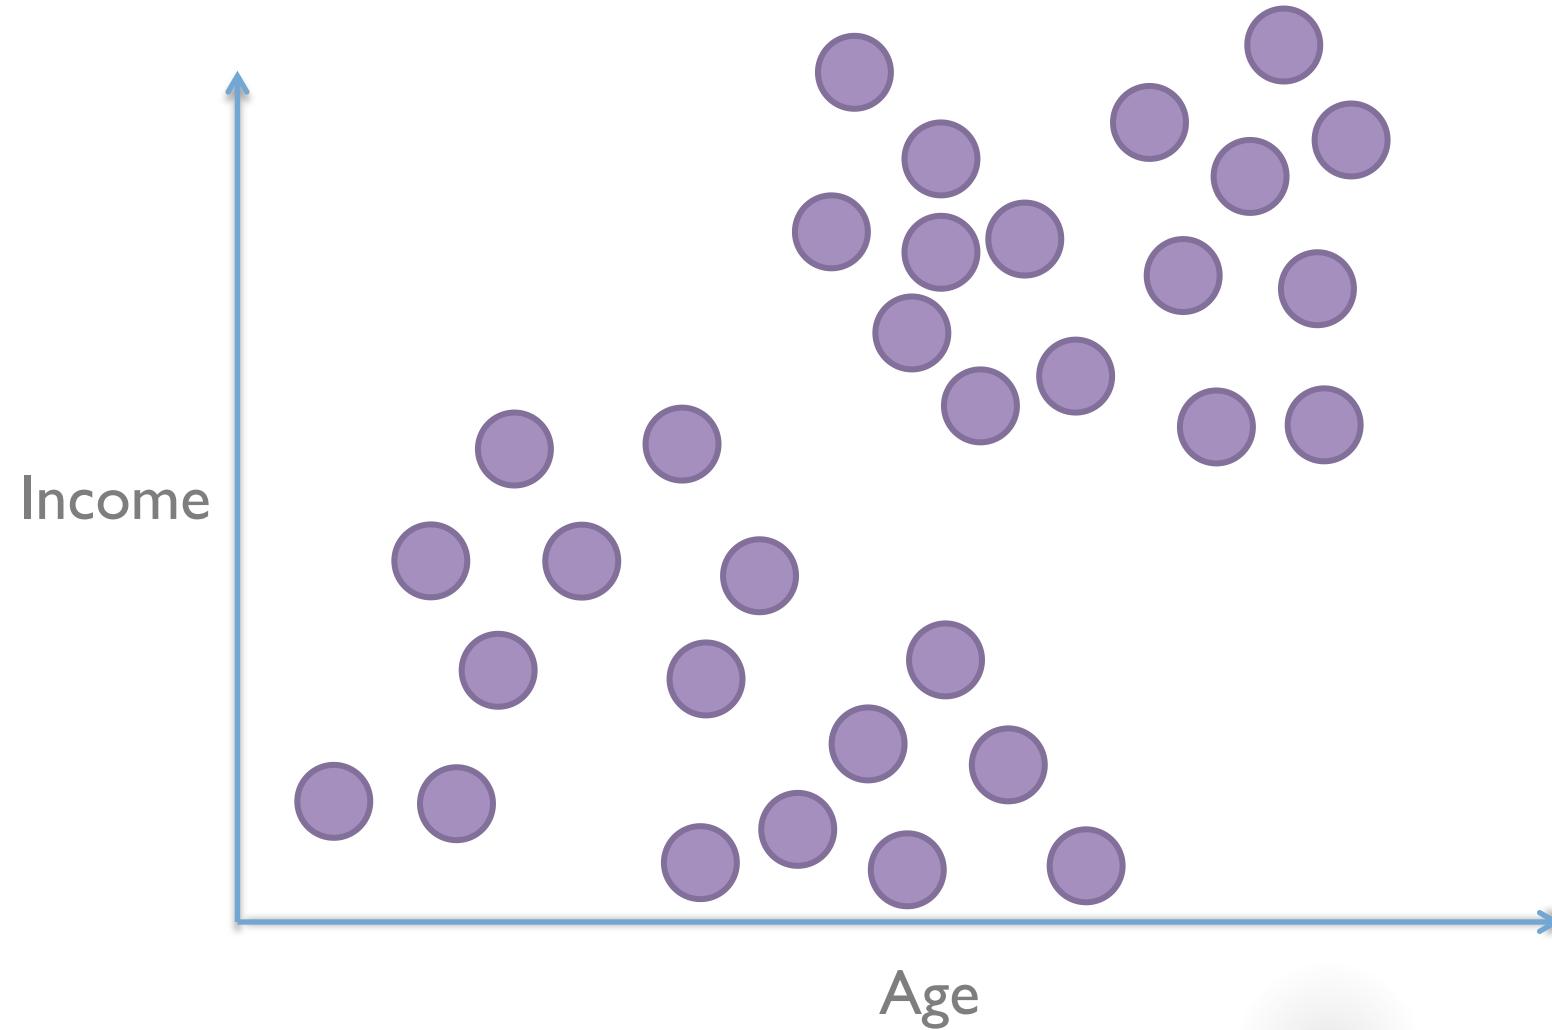
```
)
```



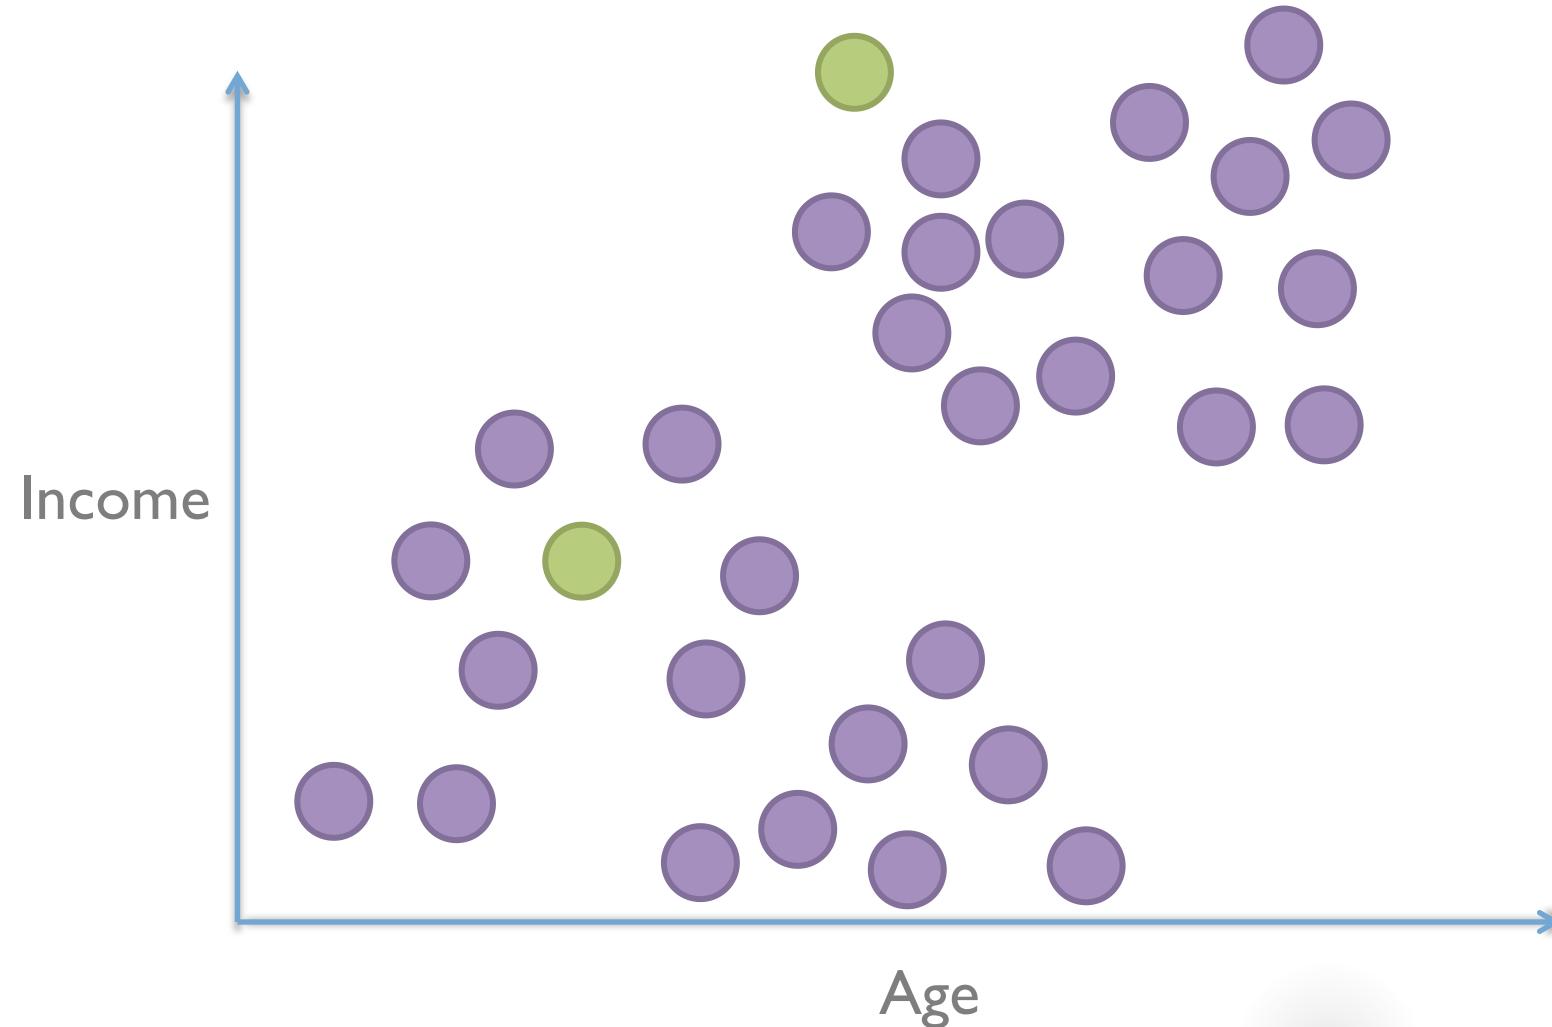
# Distance Metrics



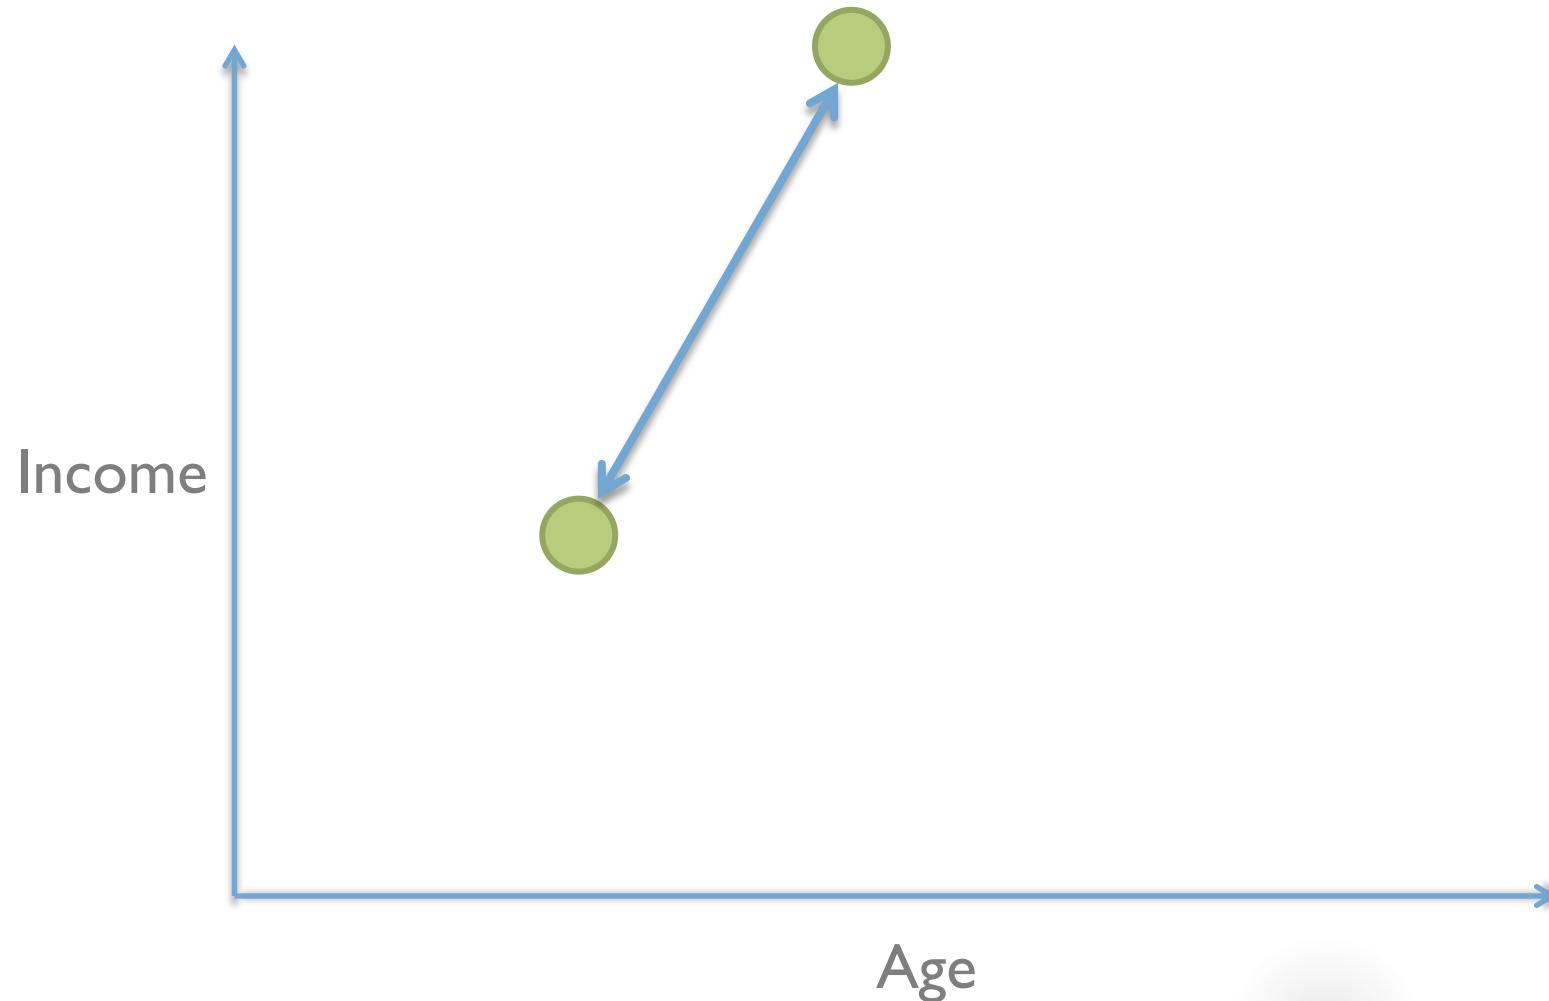
# Euclidean Distance



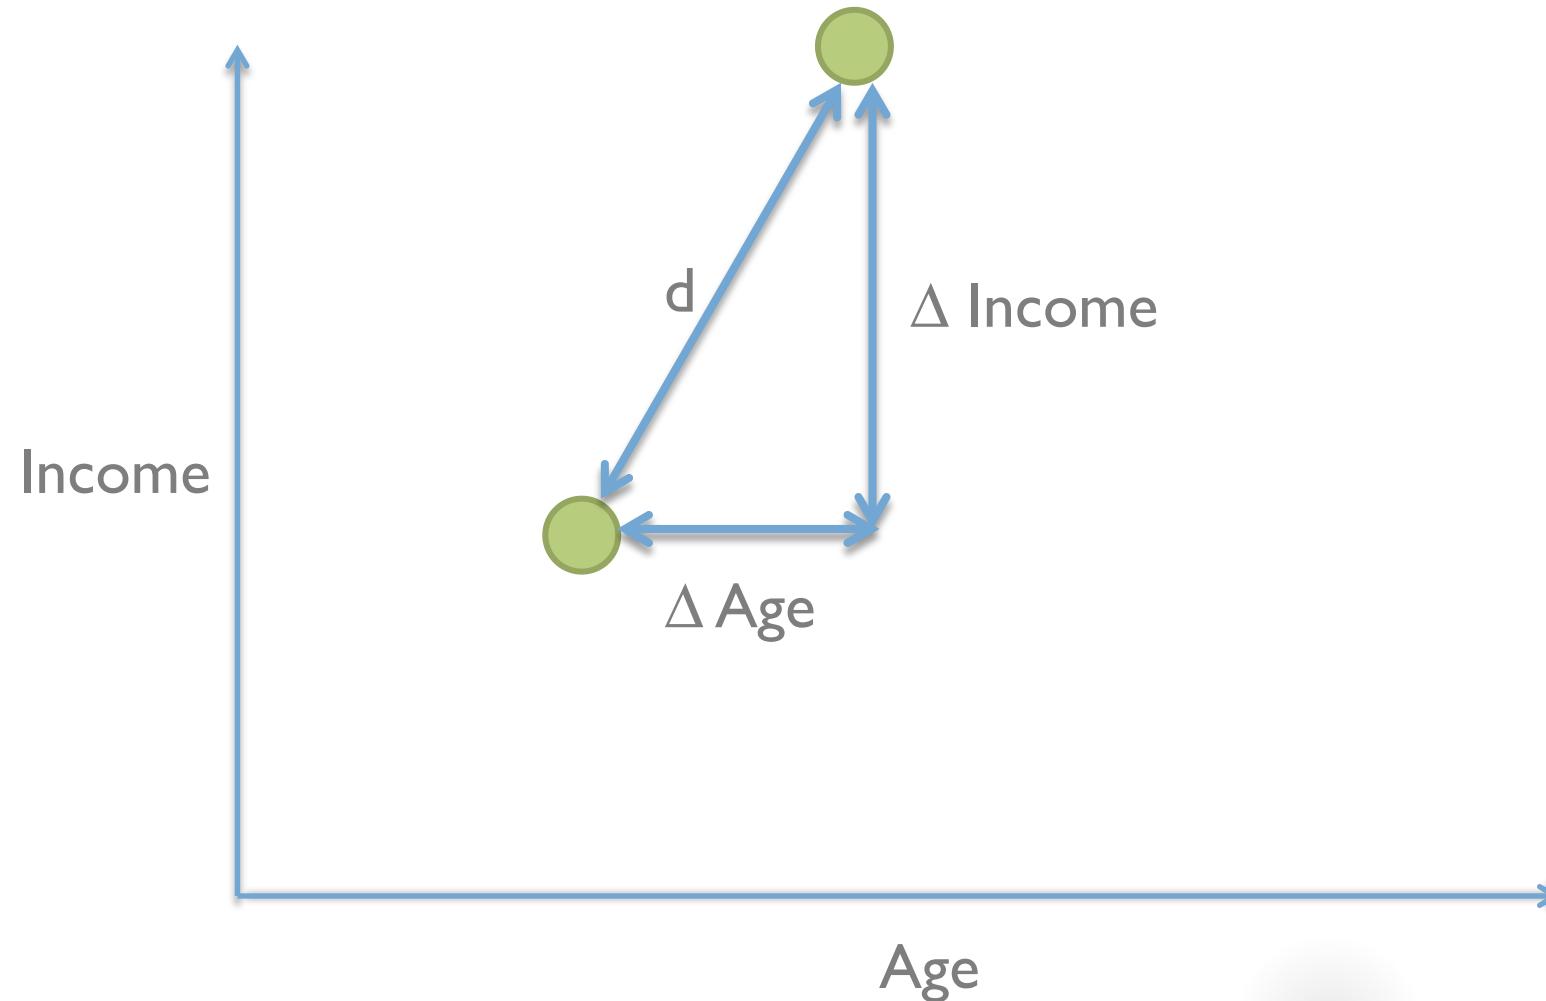
# Euclidean Distance



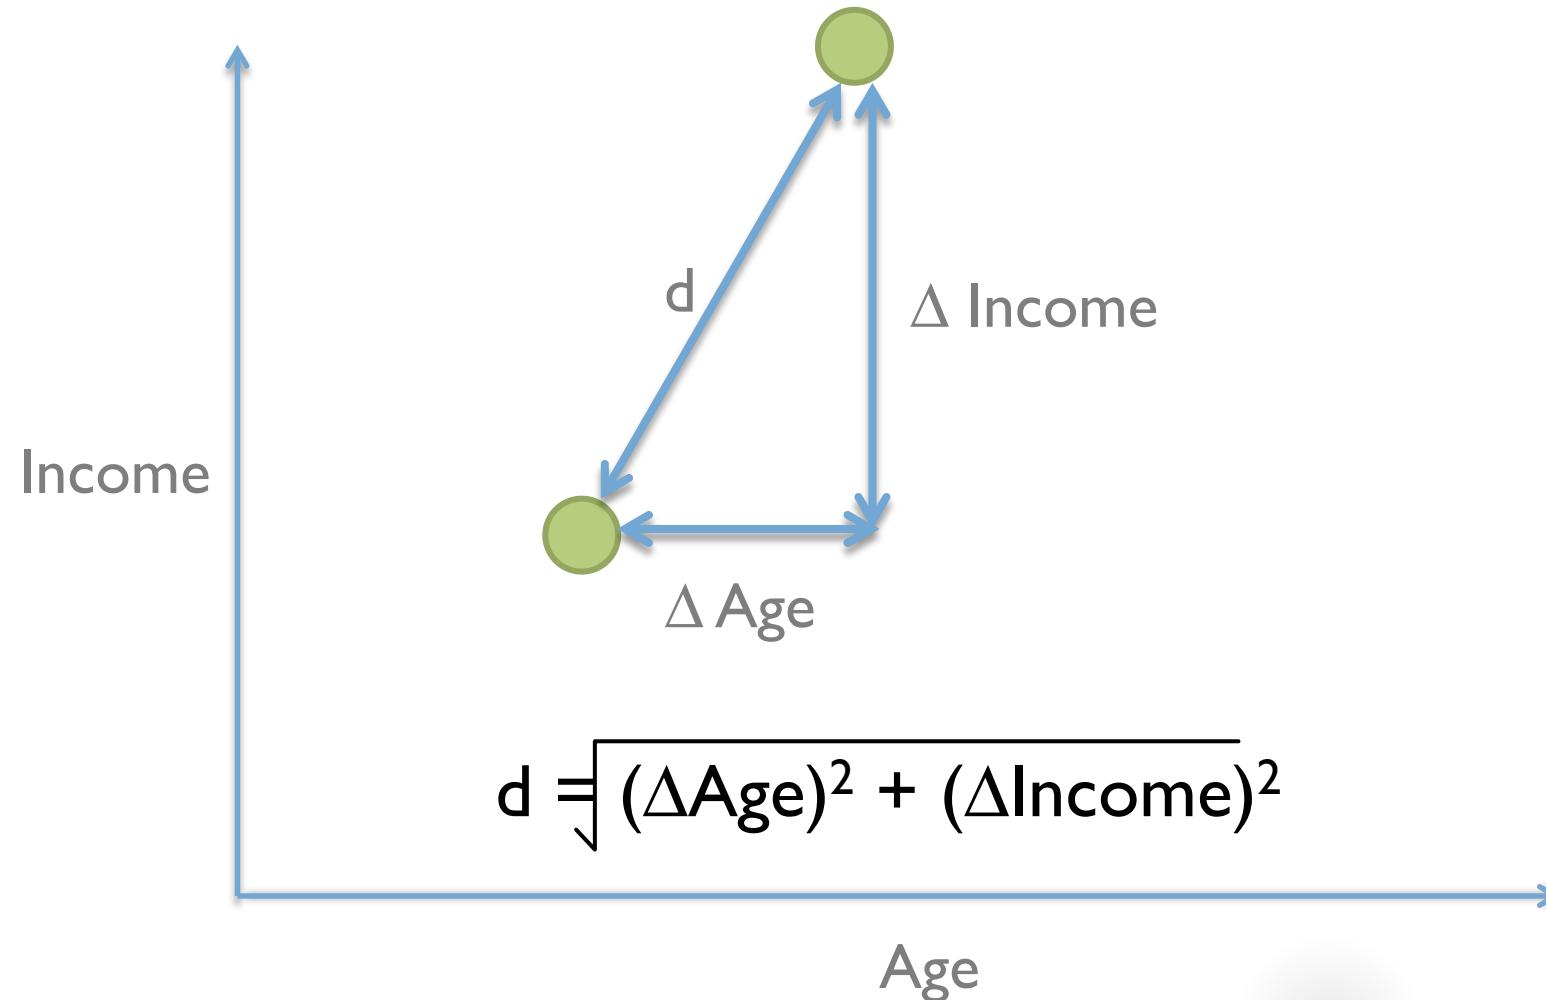
# Euclidean Distance



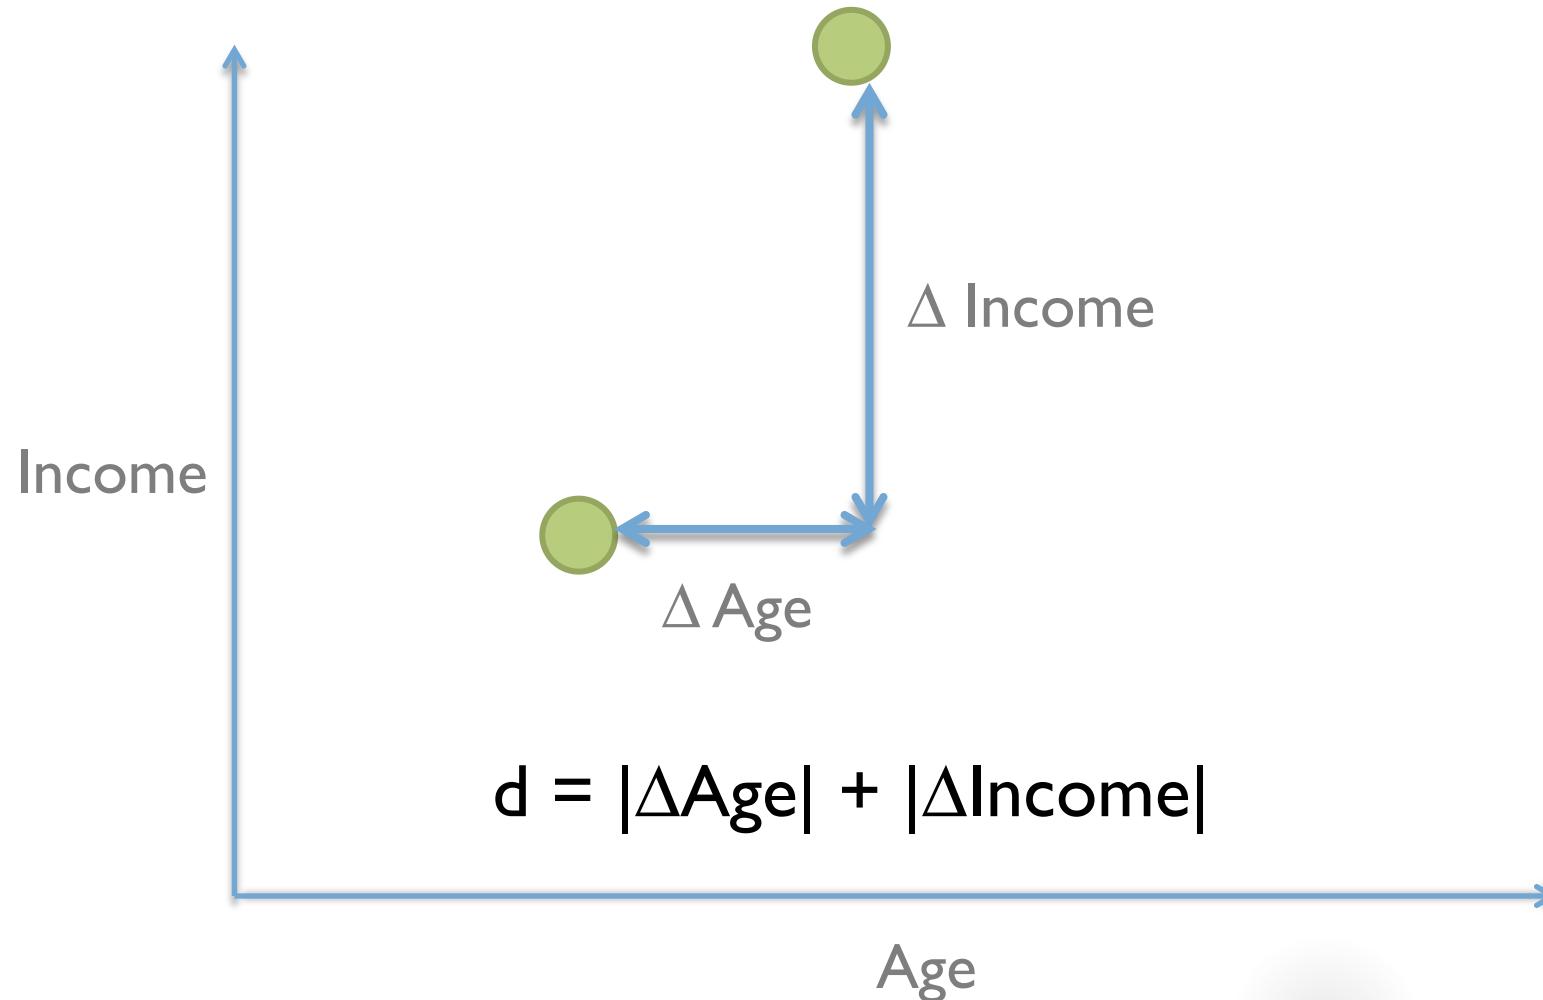
# Euclidean Distance



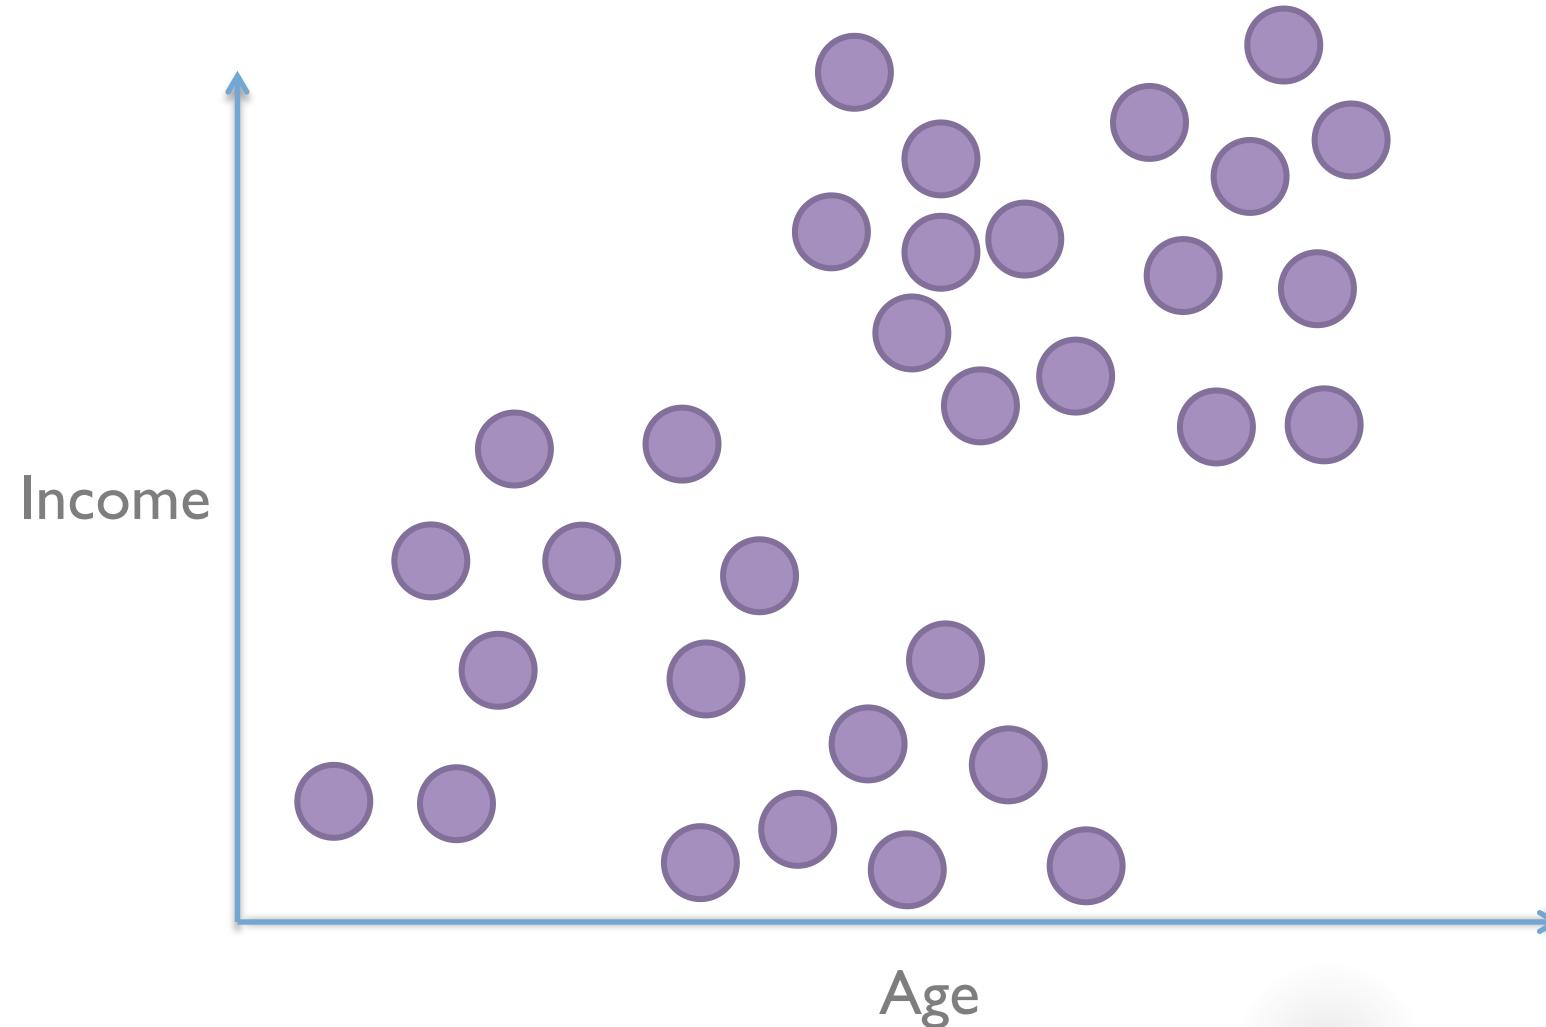
## Euclidean Distance (L2 dist)



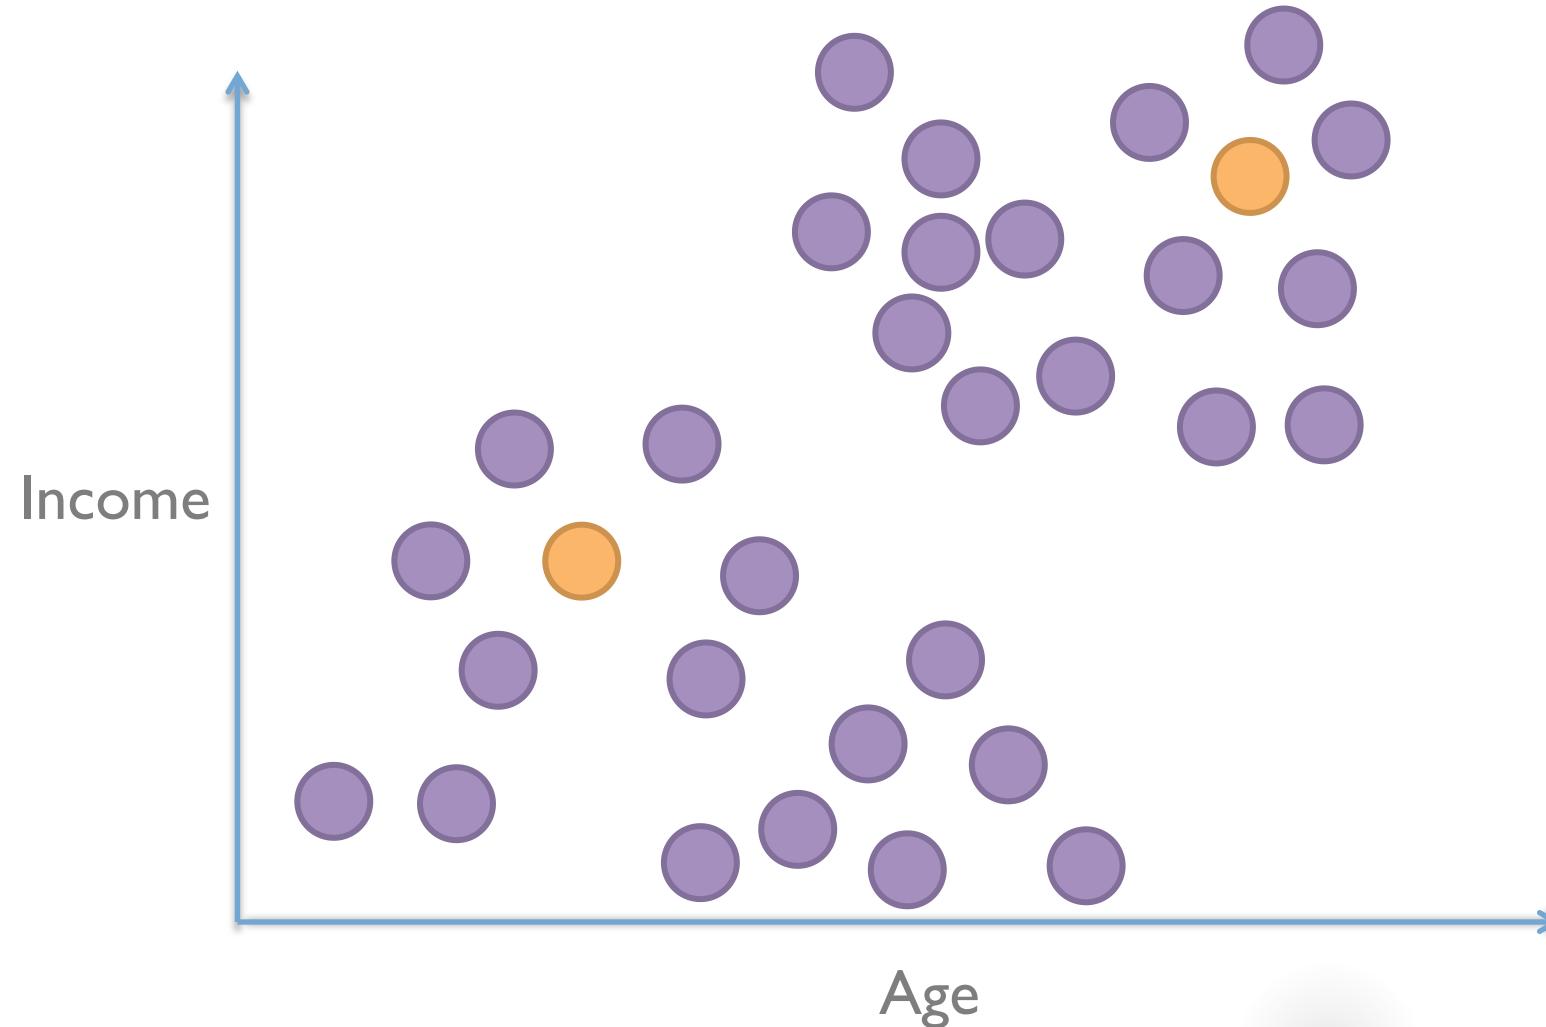
## Manhattan Distance (Cityblock dist, L1 dist)



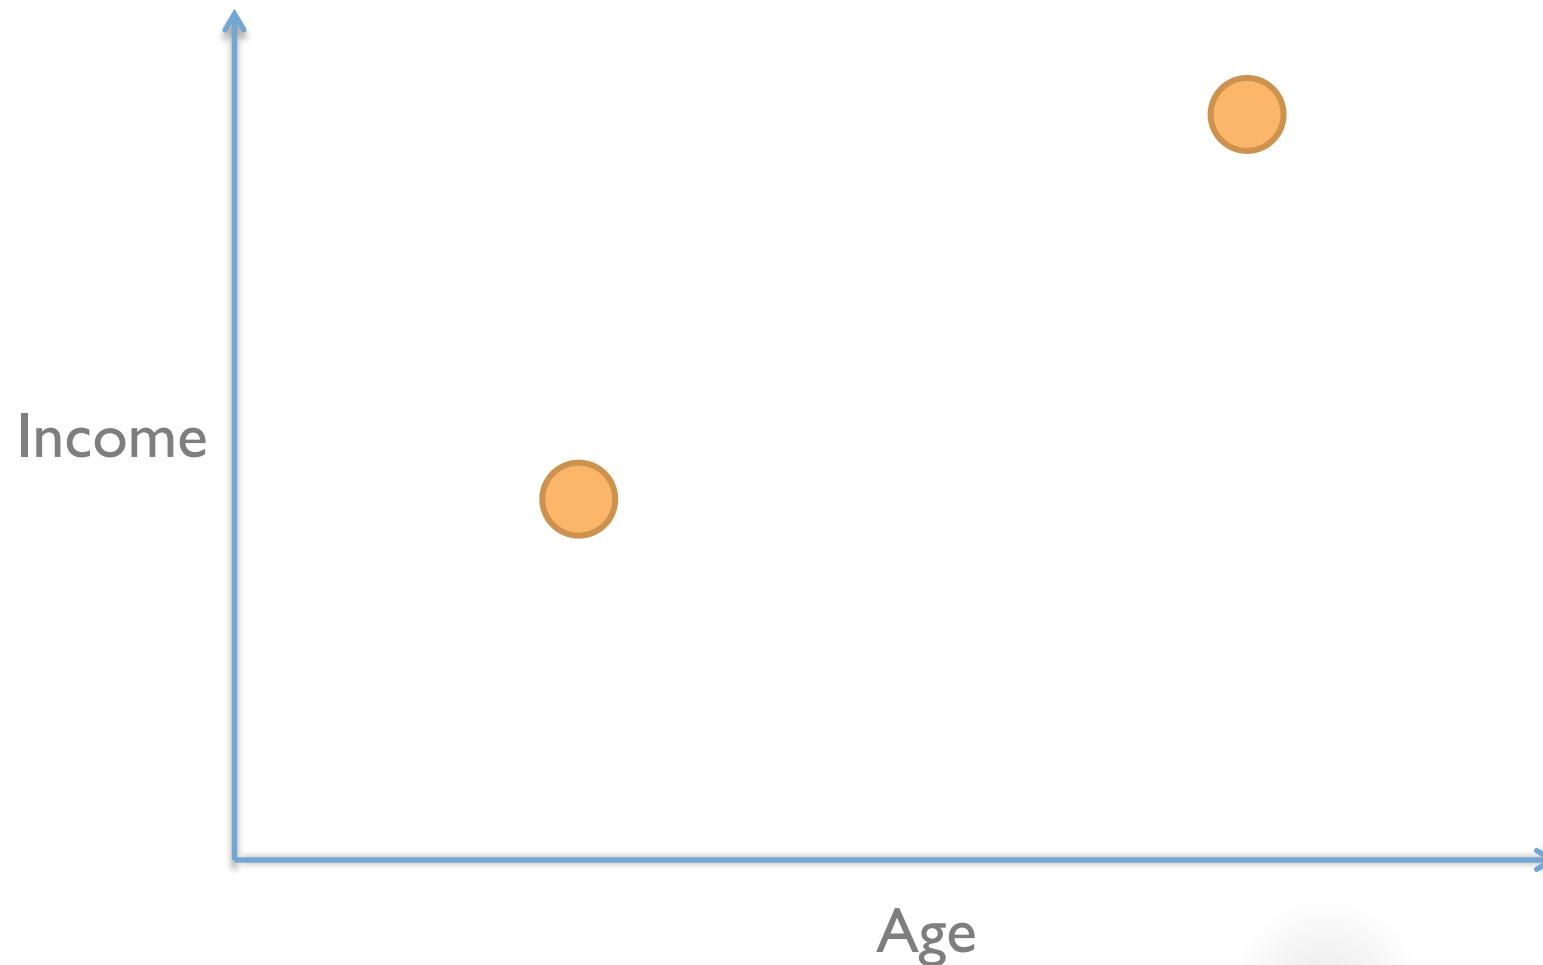
# Cosine Distance



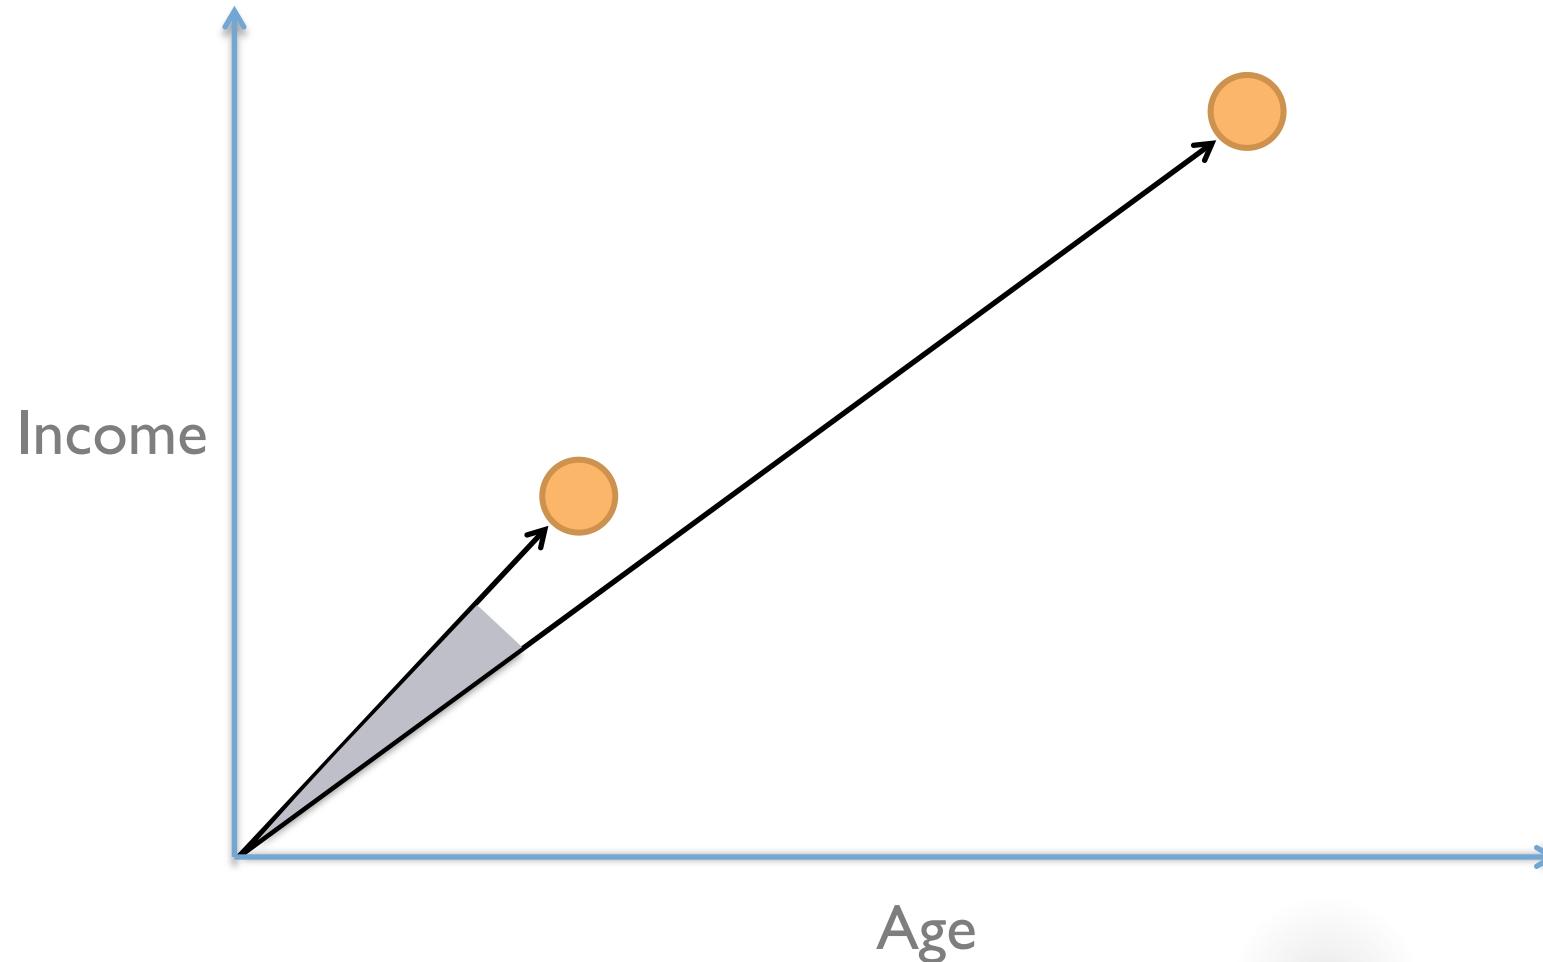
# Cosine Distance



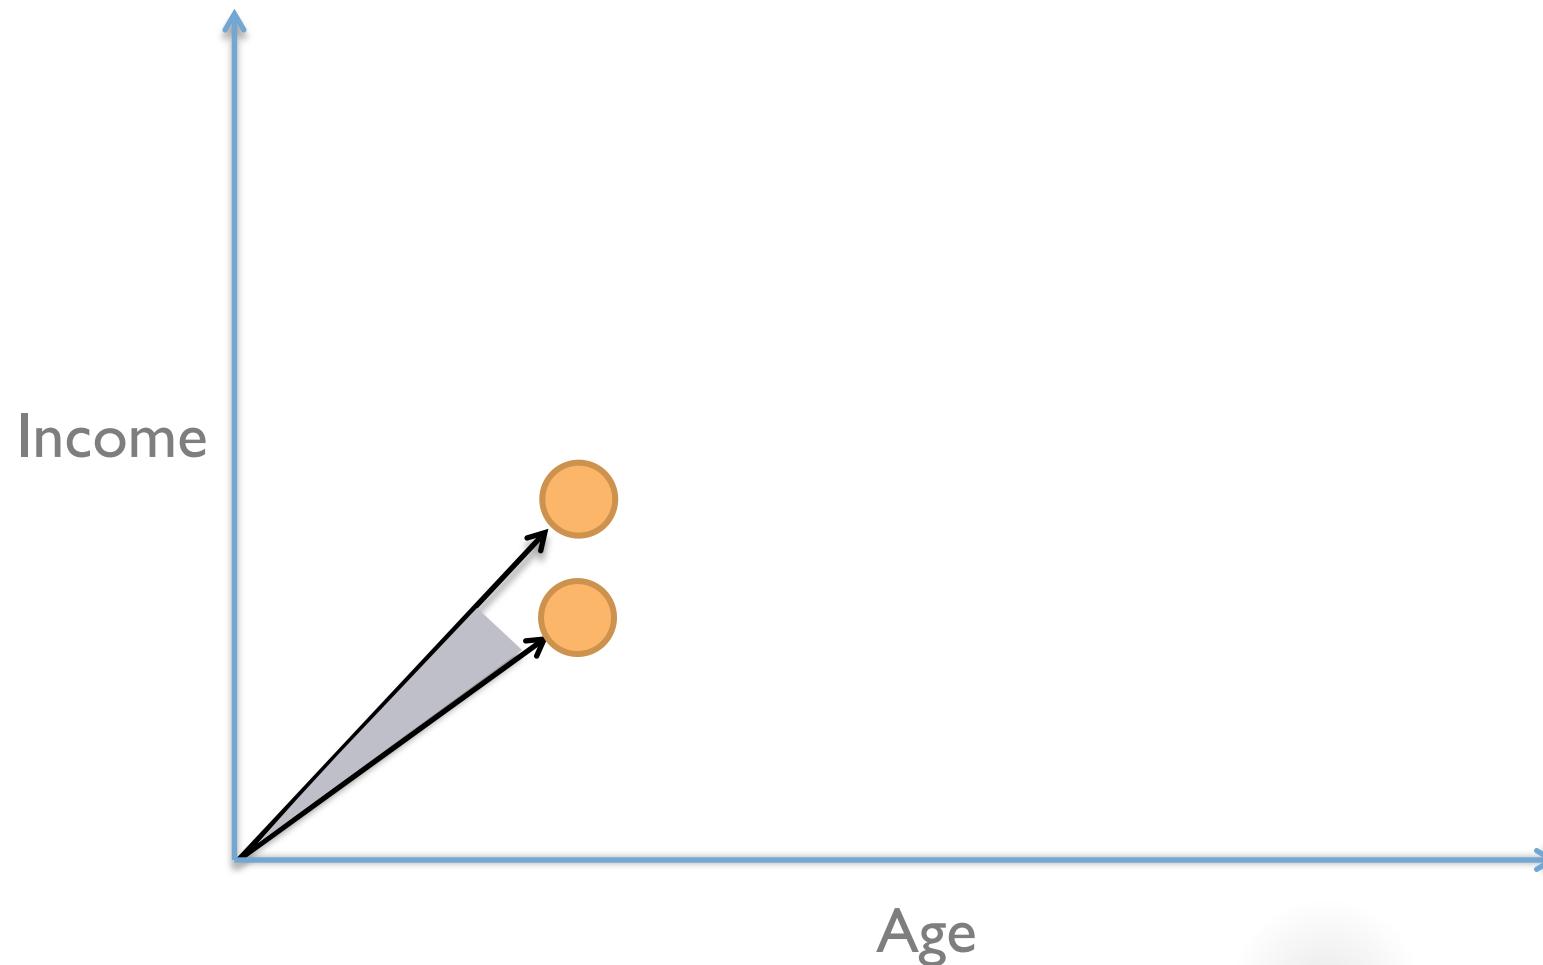
# Cosine Distance



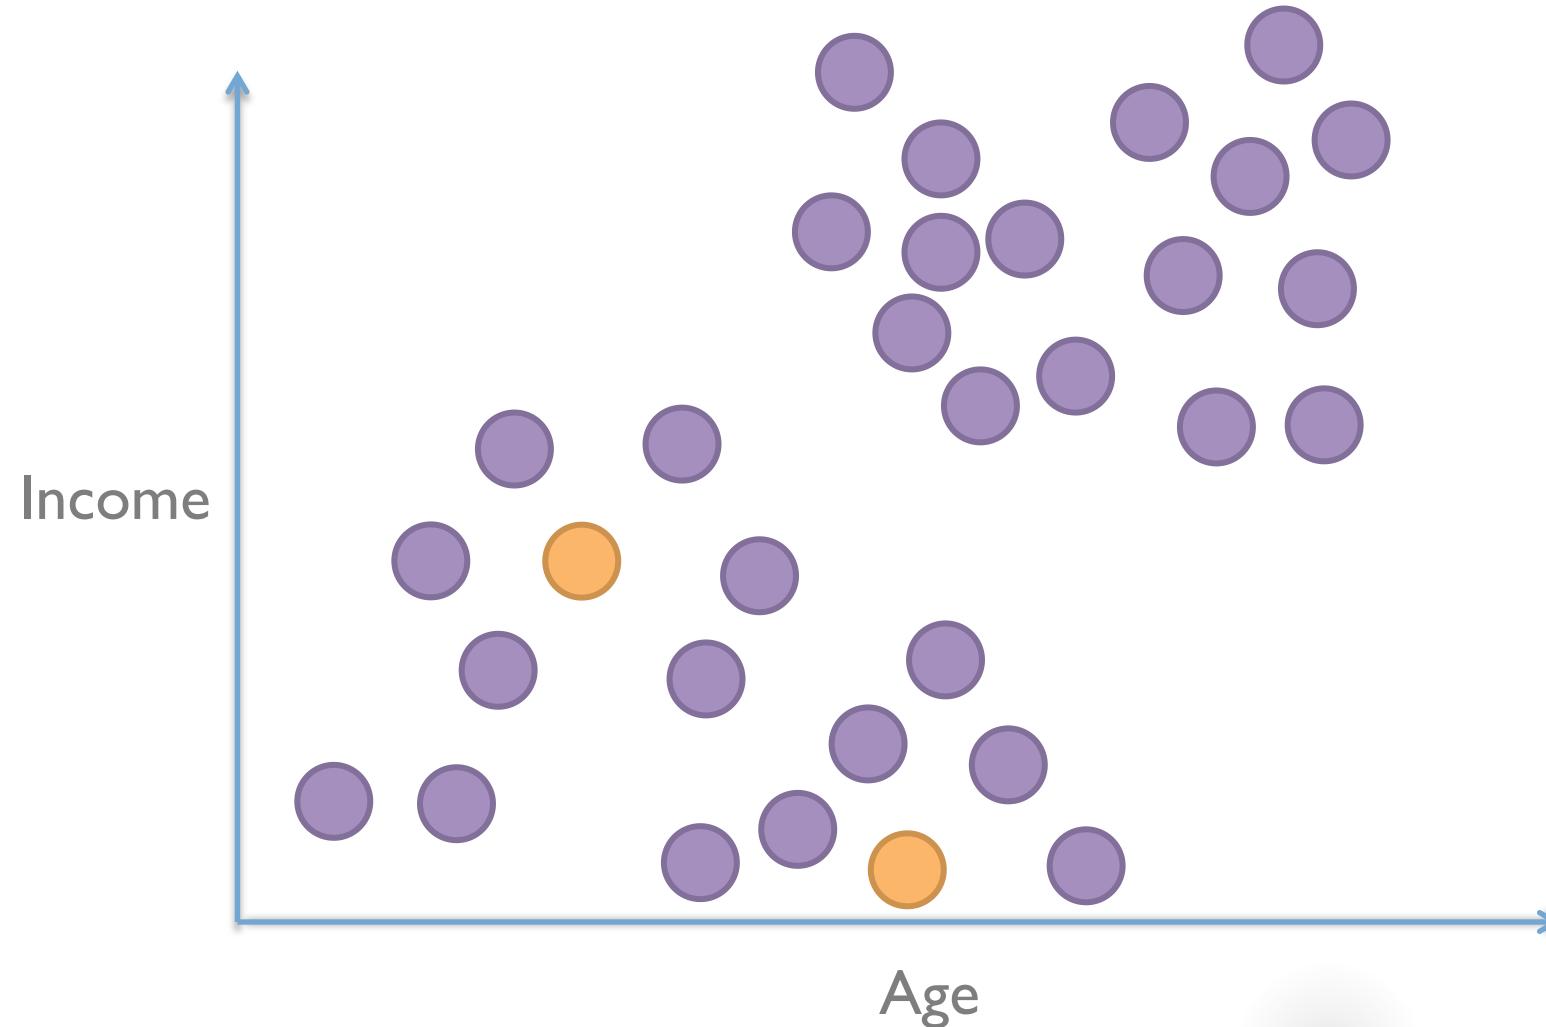
# Cosine Distance



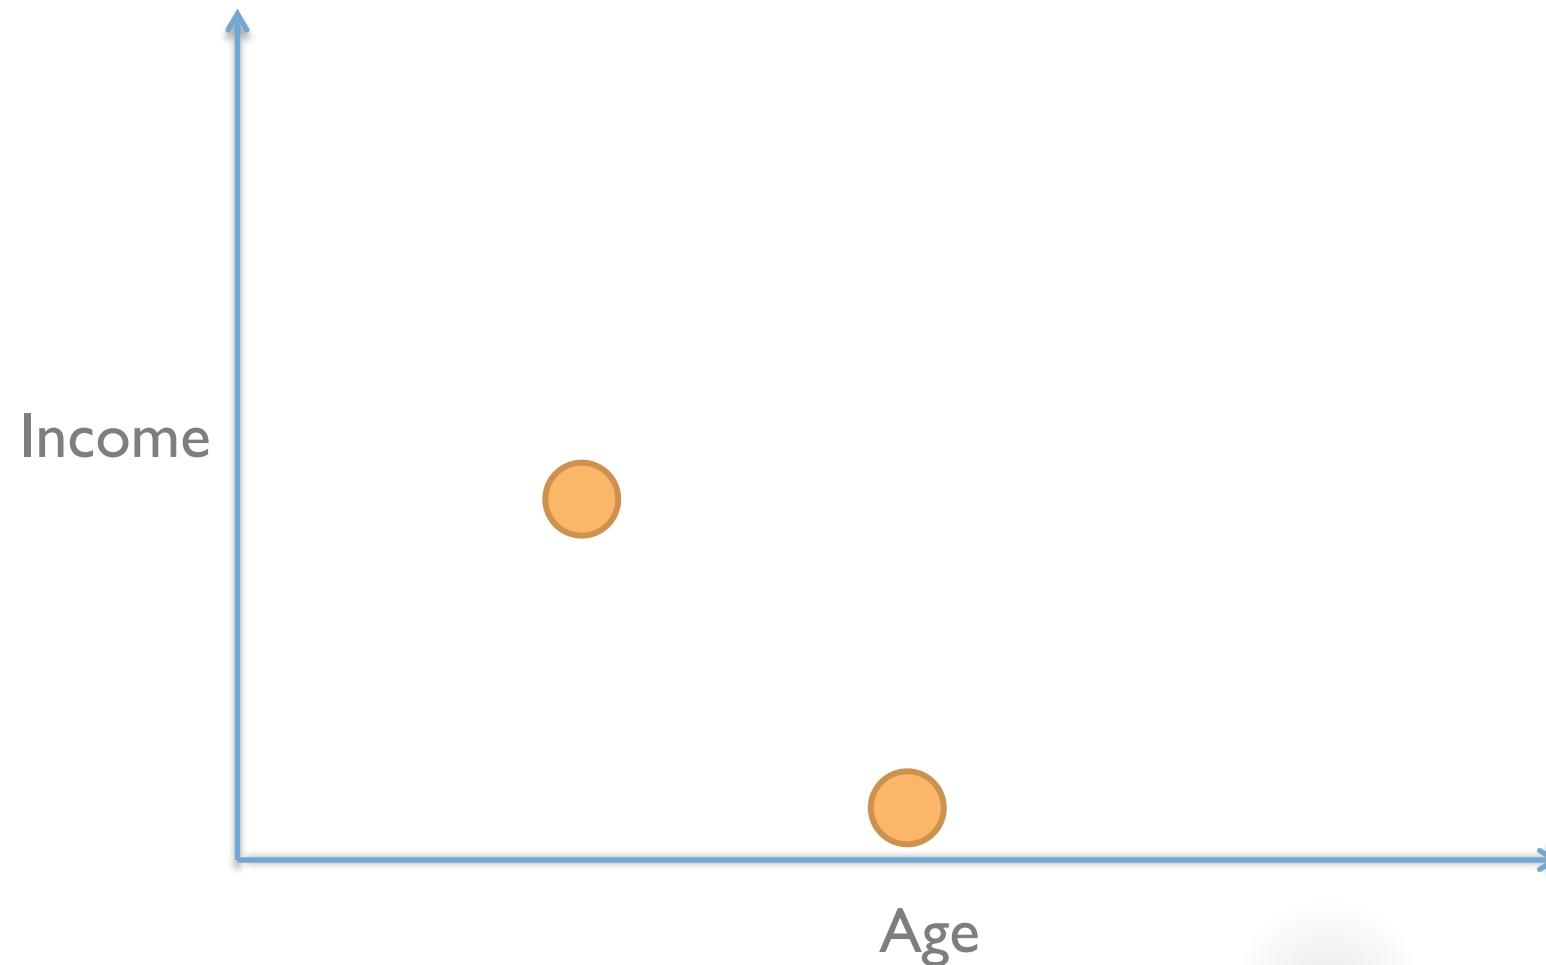
# Cosine Distance



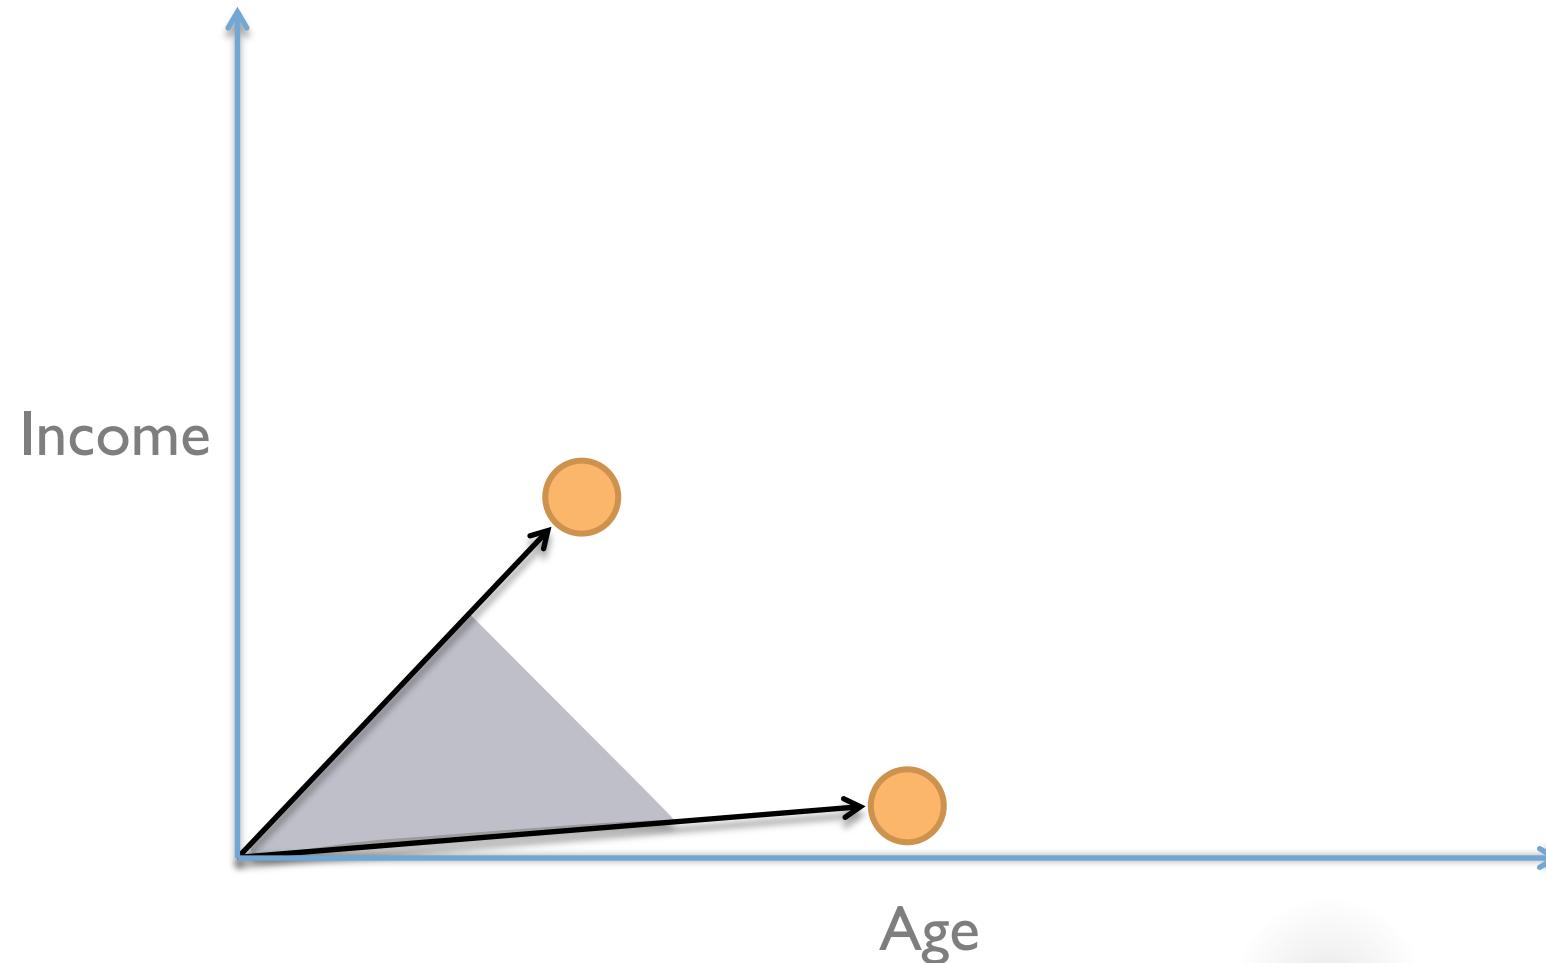
# Cosine Distance



# Cosine Distance

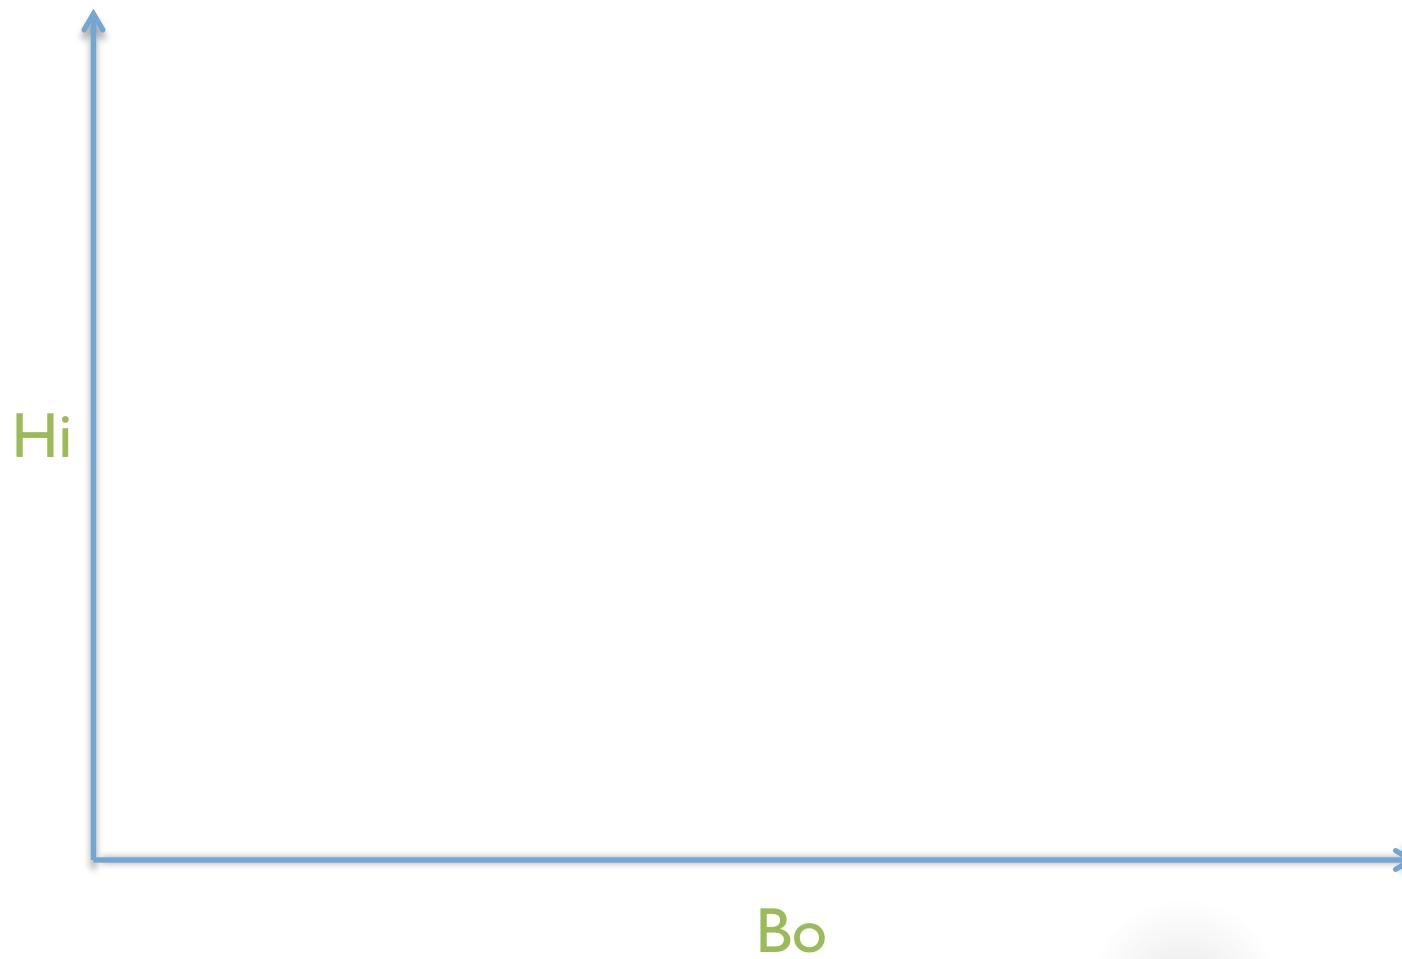


# Cosine Distance



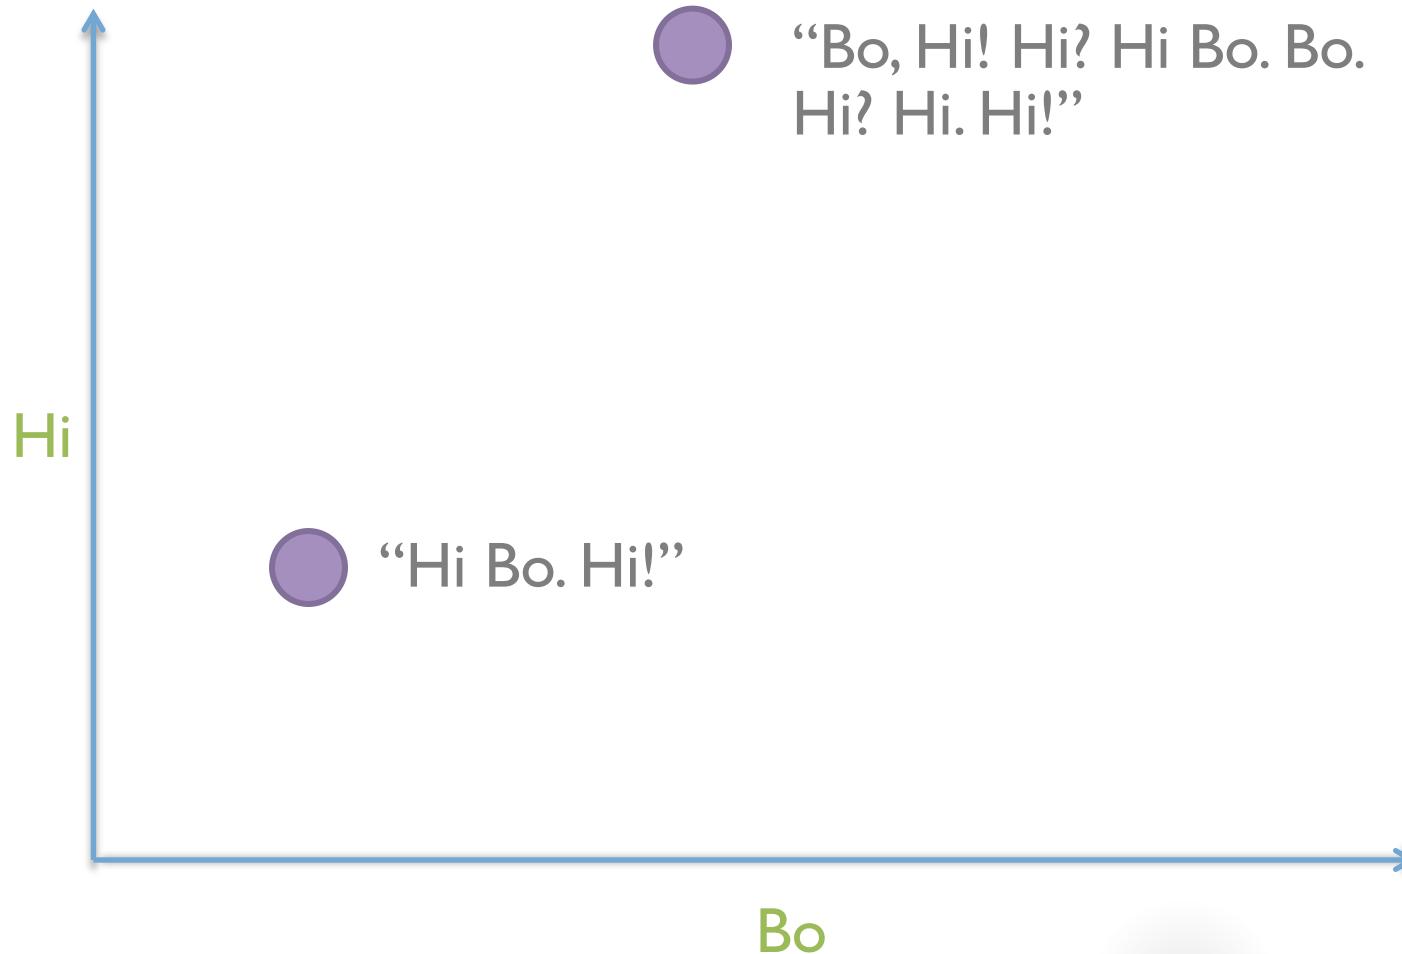
Cosine Distance is better for text data

Euclidean Distance is sensitive to doc. length



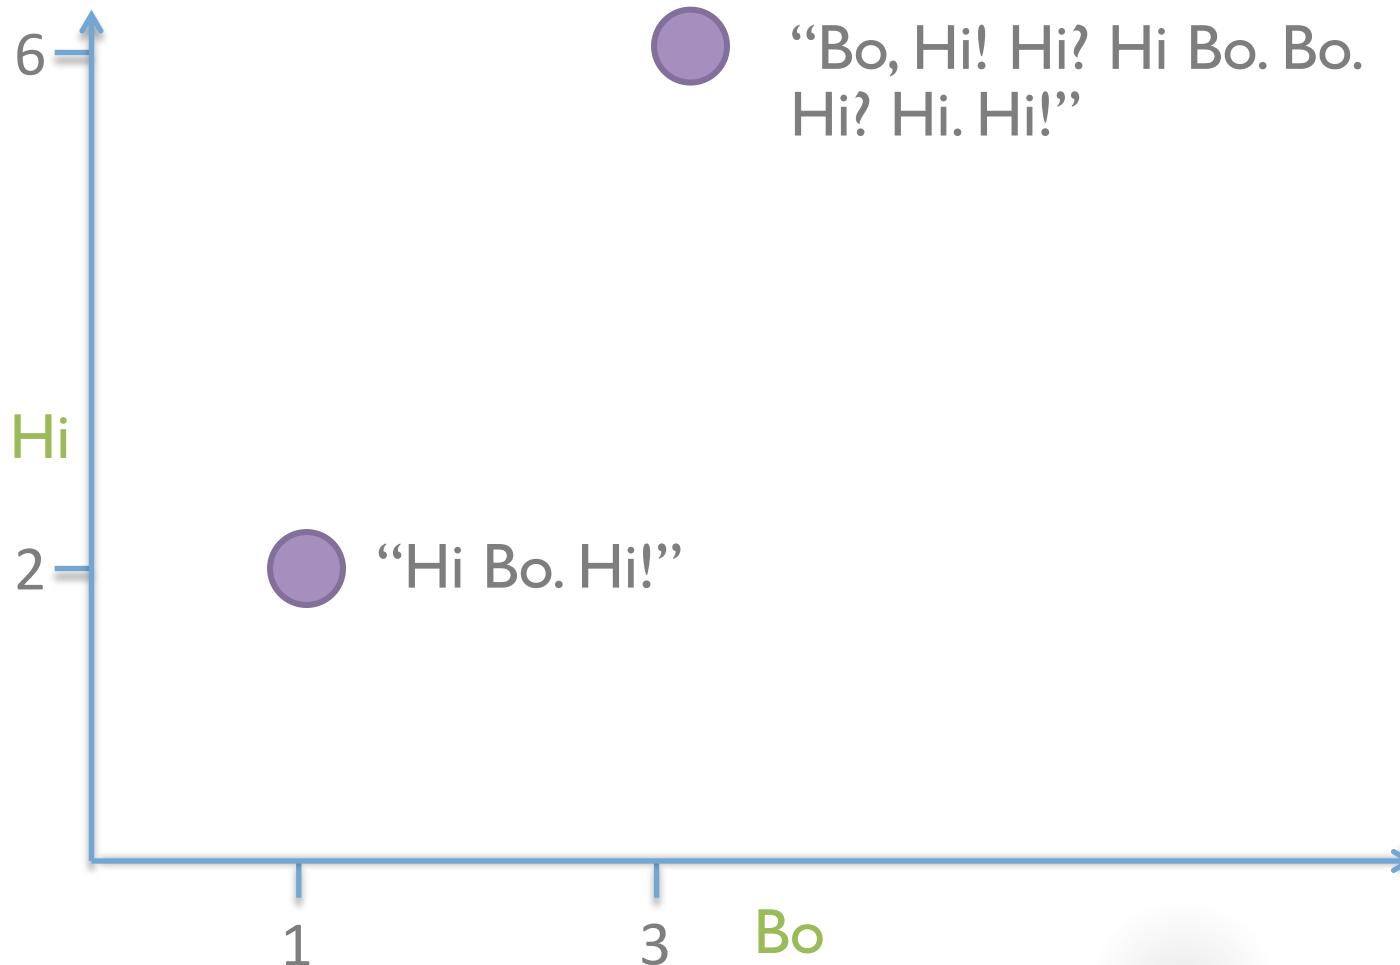
Cosine Distance is better for text data

Euclidean Distance is sensitive to doc. length



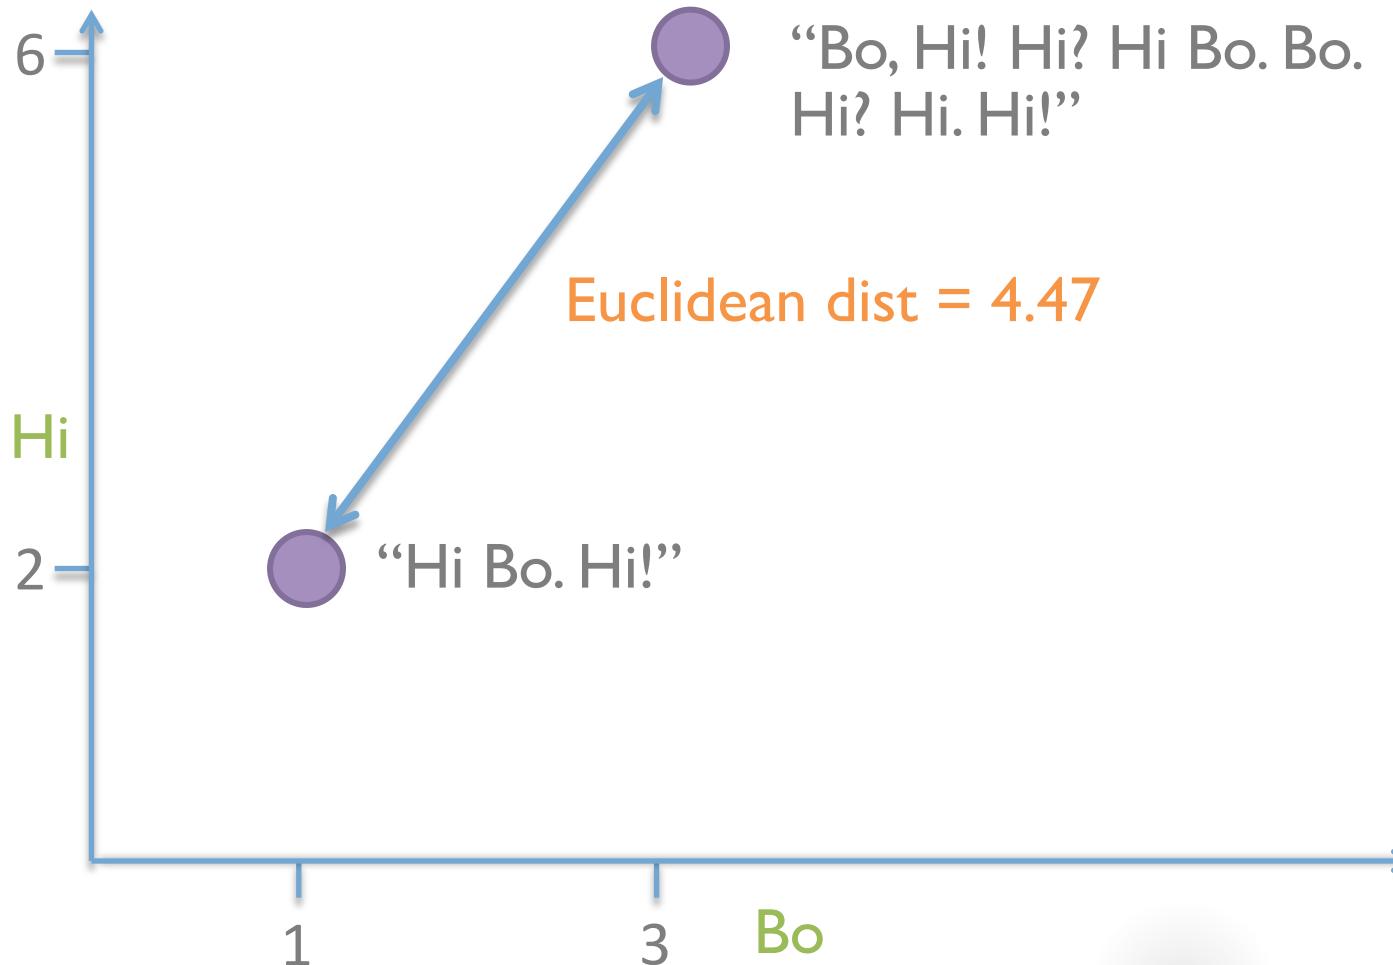
Cosine Distance is better for text data

Euclidean Distance is sensitive to doc. length



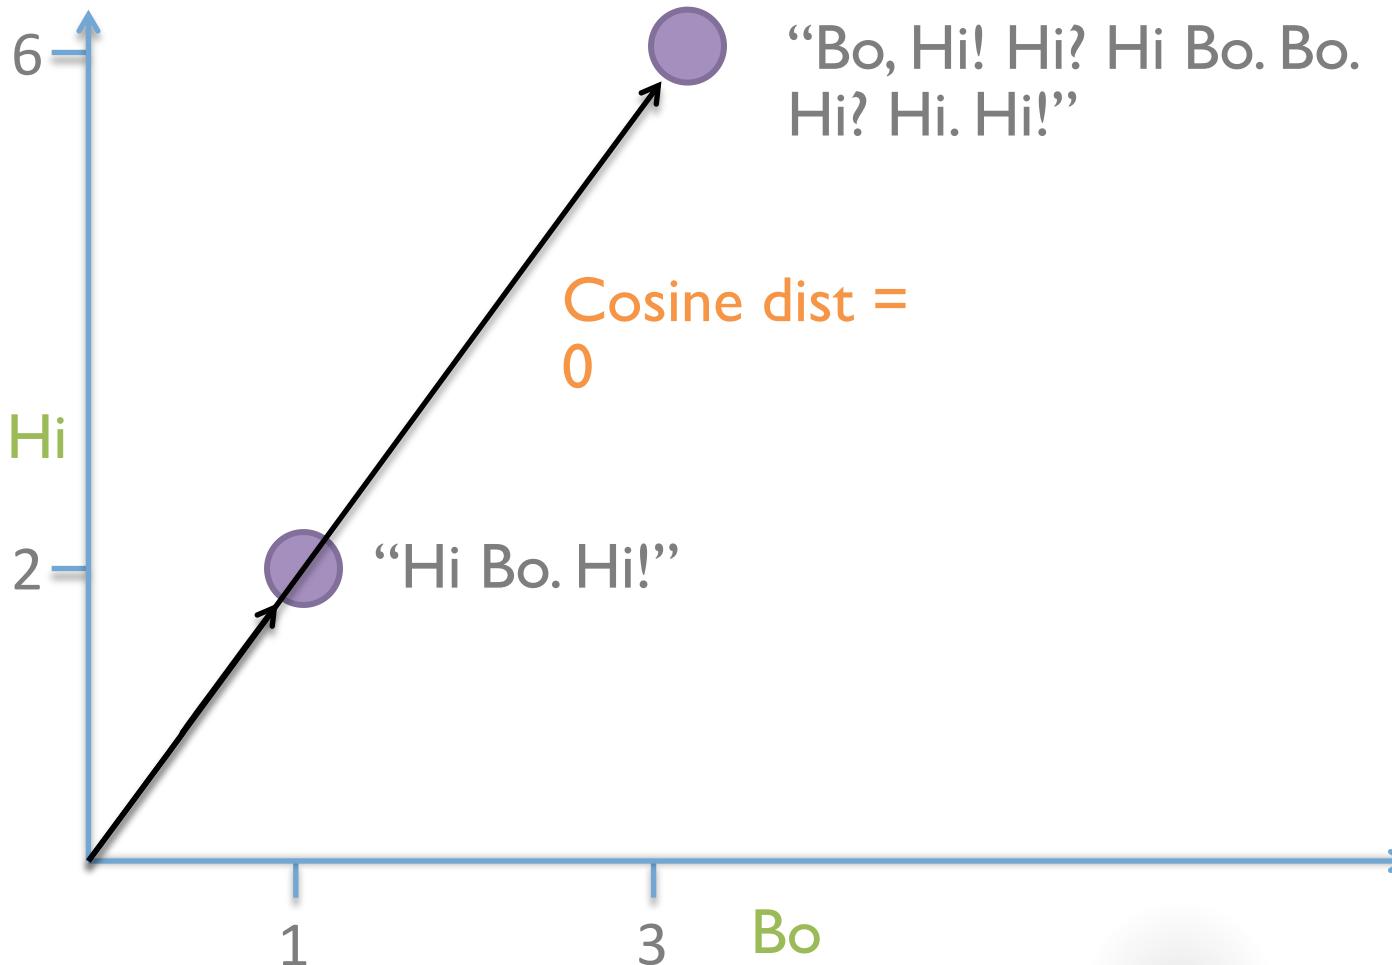
Cosine Distance is better for text data

Euclidean Distance is sensitive to doc. length

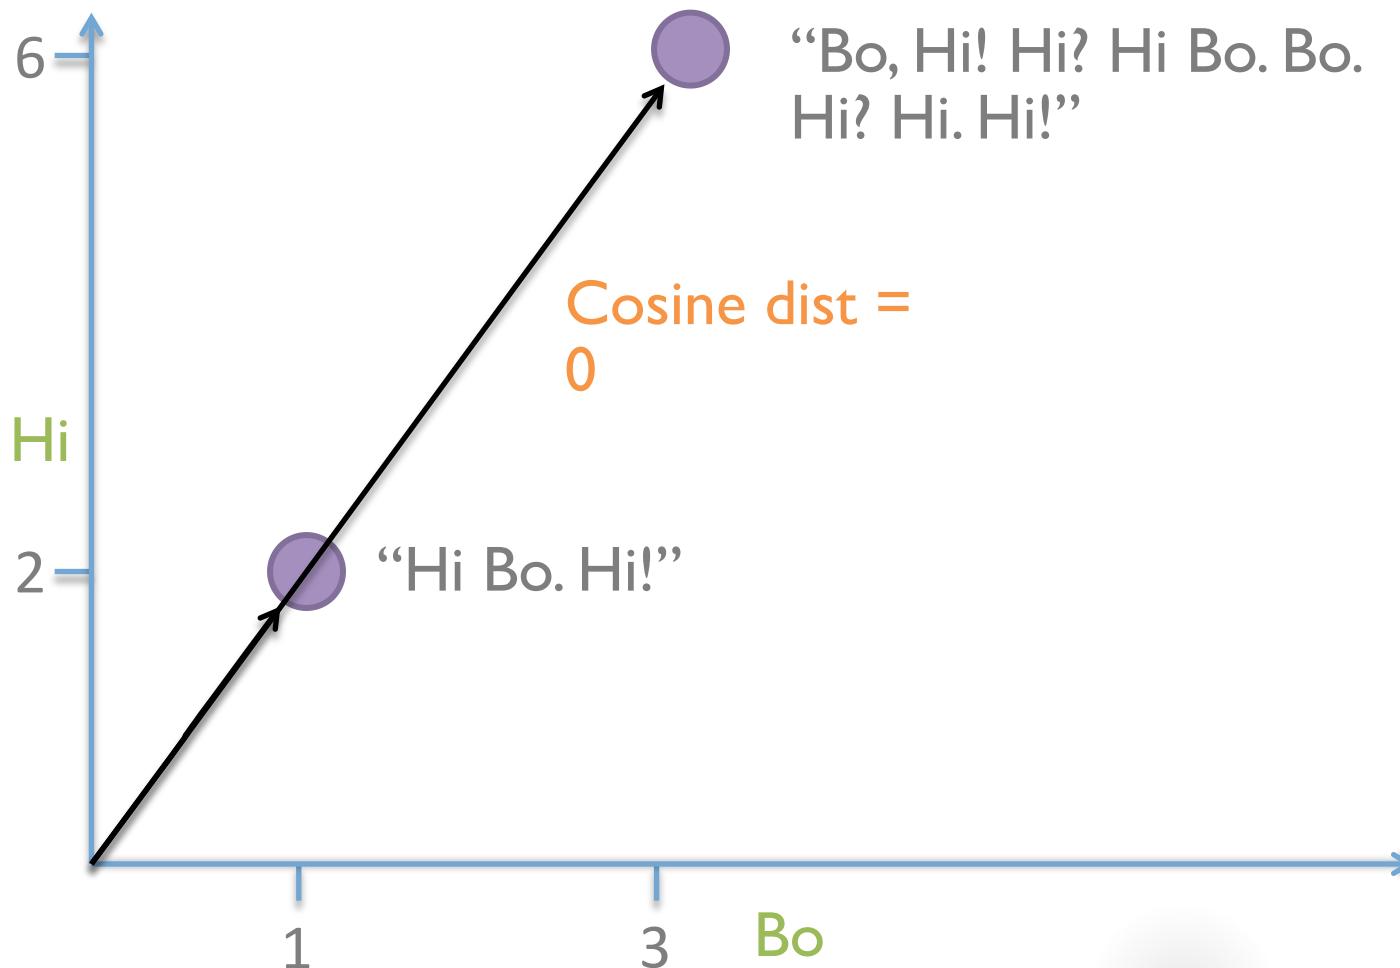


Cosine Distance is better for text data

Euclidean Distance is sensitive to doc. length



Cosine Distance captures the “meaning” distance better in a bag of words system



# Jaccard Distance

Applies to sets (like bag of words)



# Jaccard Distance

Applies to sets (like bag of words)

**Sentence A:** “I like chocolate ice cream.”

set A = {I, like, chocolate, ice, cream}

**Sentence B:** “Do I want chocolate cream or vanilla cream?”

set B = {Do, I, want, chocolate, cream, or, vanilla}

Jaccard dist



# Jaccard Distance

Applies to sets (like bag of words)

**Sentence A:** “I like chocolate ice cream.”

set A = {I, like, chocolate, ice, cream}

**Sentence B:** “Do I want chocolate cream or vanilla cream?”

set B = {Do, I, want, chocolate, cream, or, vanilla}

$$\text{Jaccard dist} \quad 1 - \frac{|A \cap B|}{|A \cup B|} = 1 - \frac{\text{len(shared)}}{\text{len(all)}}$$



# Jaccard Distance

Applies to sets (like bag of words)

Sentence A: "I like chocolate ice cream."

set A = {I, like, chocolate, ice, cream}

Sentence B: "Do I want chocolate cream or vanilla cream?"

set B = {Do, I, want, chocolate, cream, or, vanilla}

$$\text{Jaccard dist} \quad | - \frac{A \cap B}{A \cup B} = | \frac{\text{len(shared)}}{\text{len(all)}}$$



# Jaccard Distance

Applies to sets (like bag of words)

Sentence A: "I like chocolate ice cream."

set A = {I, like, chocolate, ice, cream}

Sentence B: "Do I want chocolate cream or vanilla cream?"

set B = {Do, I, want, chocolate, cream, or, vanilla}

Jaccard dist

$$| \frac{\text{len}(\{I, chocolate, cream\})}{\text{len}(\{I, like, chocolate, ice, cream, do, want, or, vanilla\})} |$$



# Jaccard Distance

Applies to sets (like bag of words)

Sentence A: "I like chocolate ice cream."

set A = {I, like, chocolate, ice, cream}

Sentence B: "Do I want chocolate cream or vanilla cream?"

set B = {Do, I, want, chocolate, cream, or, vanilla}

Jaccard dist       $= 1 - 3/9 / = 1 - 0.333 = 0.667$





# Back to Other Clustering Algorithms

---

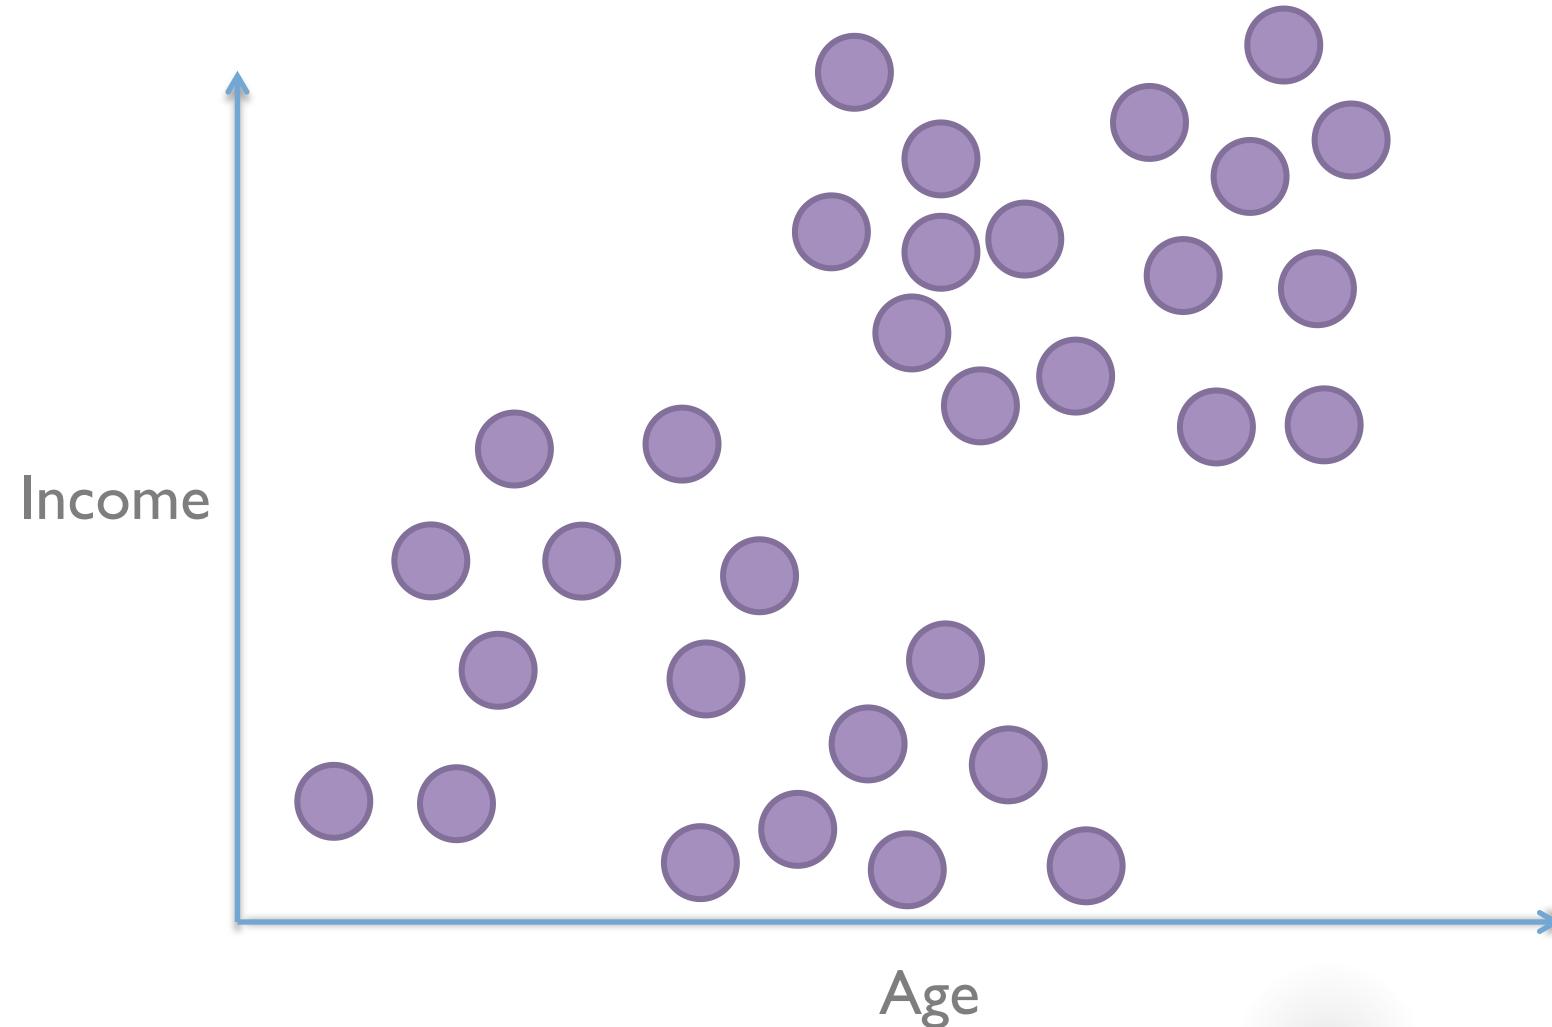


# DBSCAN



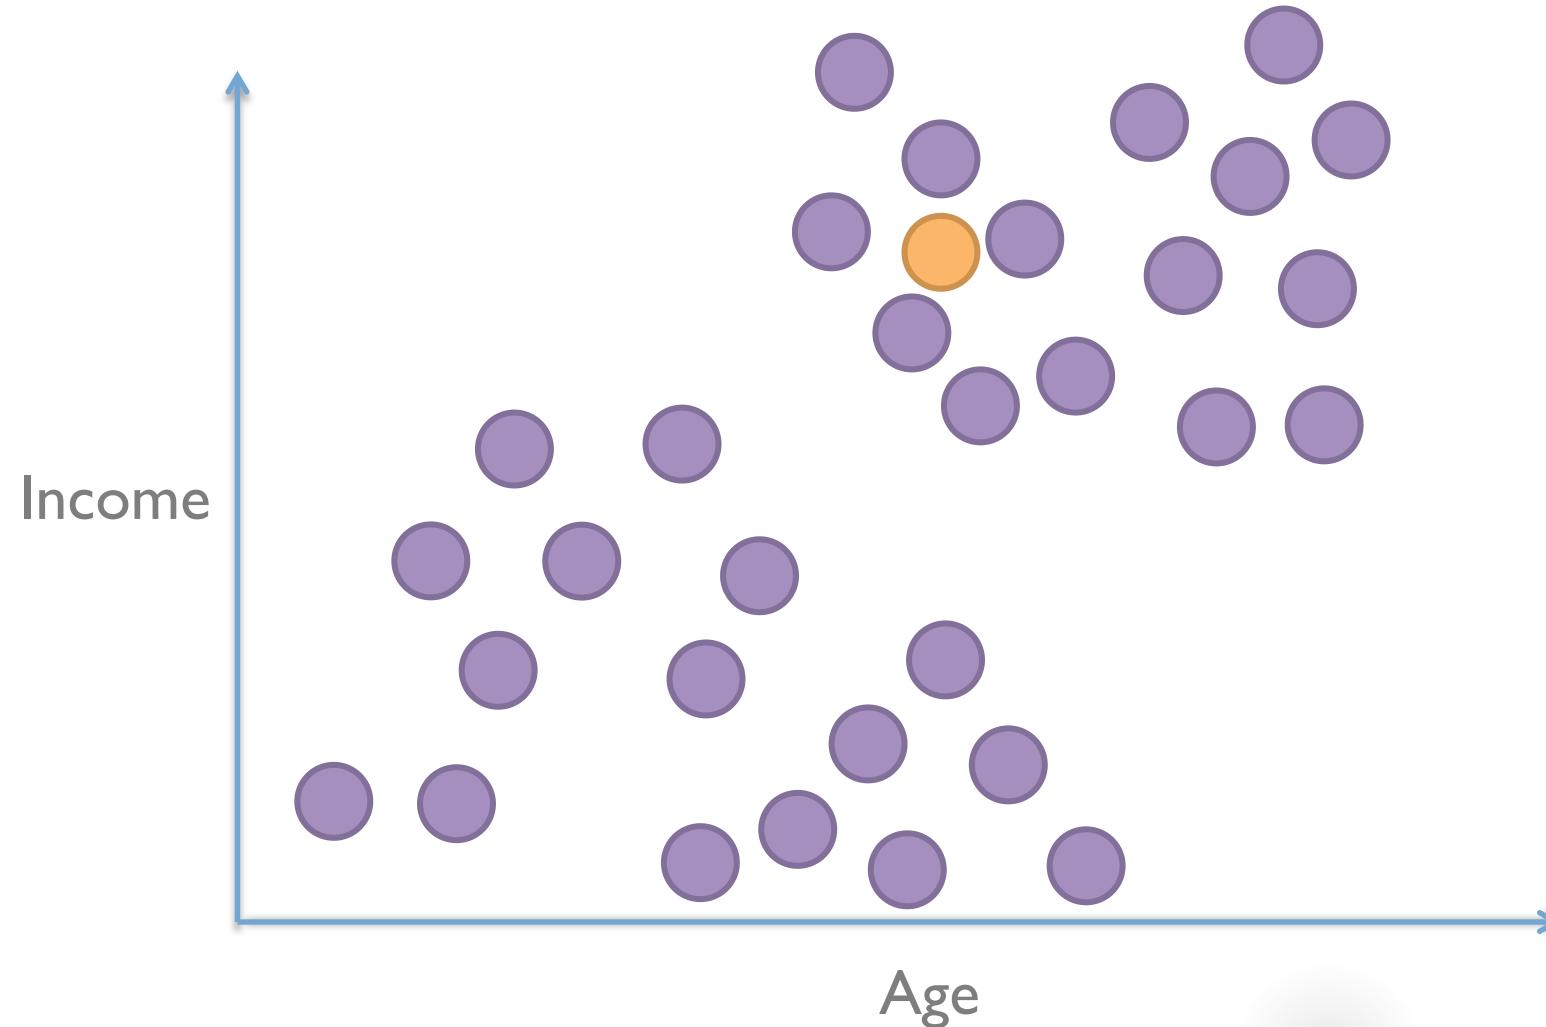
# DBSCAN

Density---based spatial clustering of applications with noise



# DBSCAN

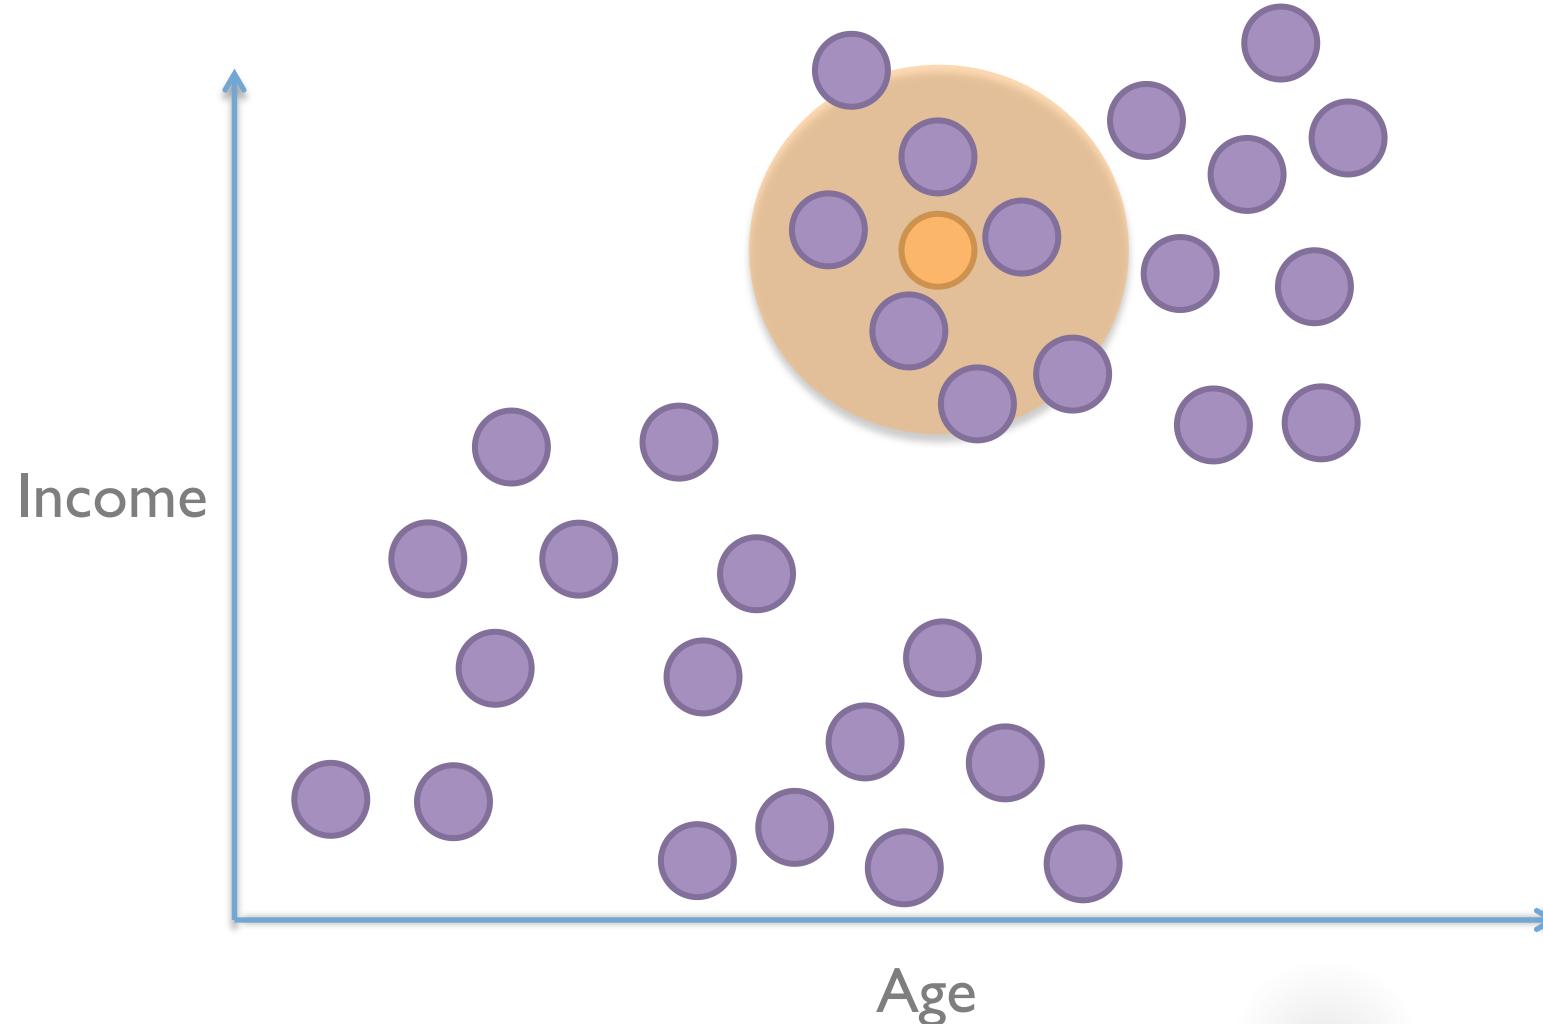
Start at a random point



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

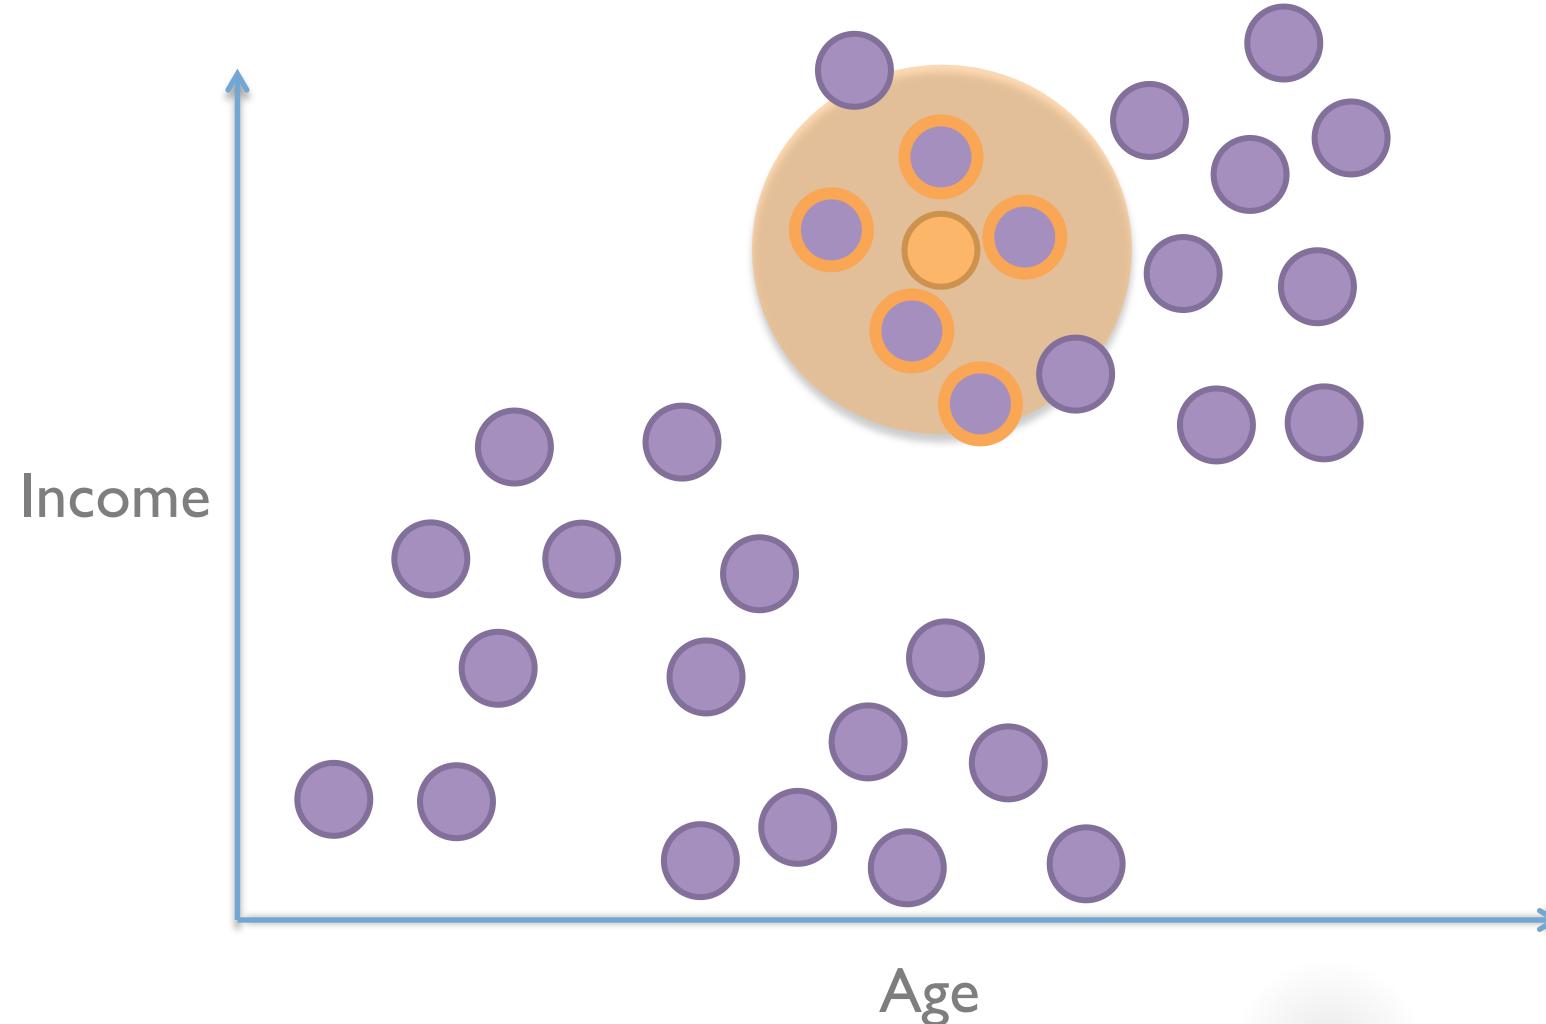
Look at the radius  $\text{epsilon}$  around point  
If enough points ( $n_{\text{clu}}$ ) within circle, start cluster



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

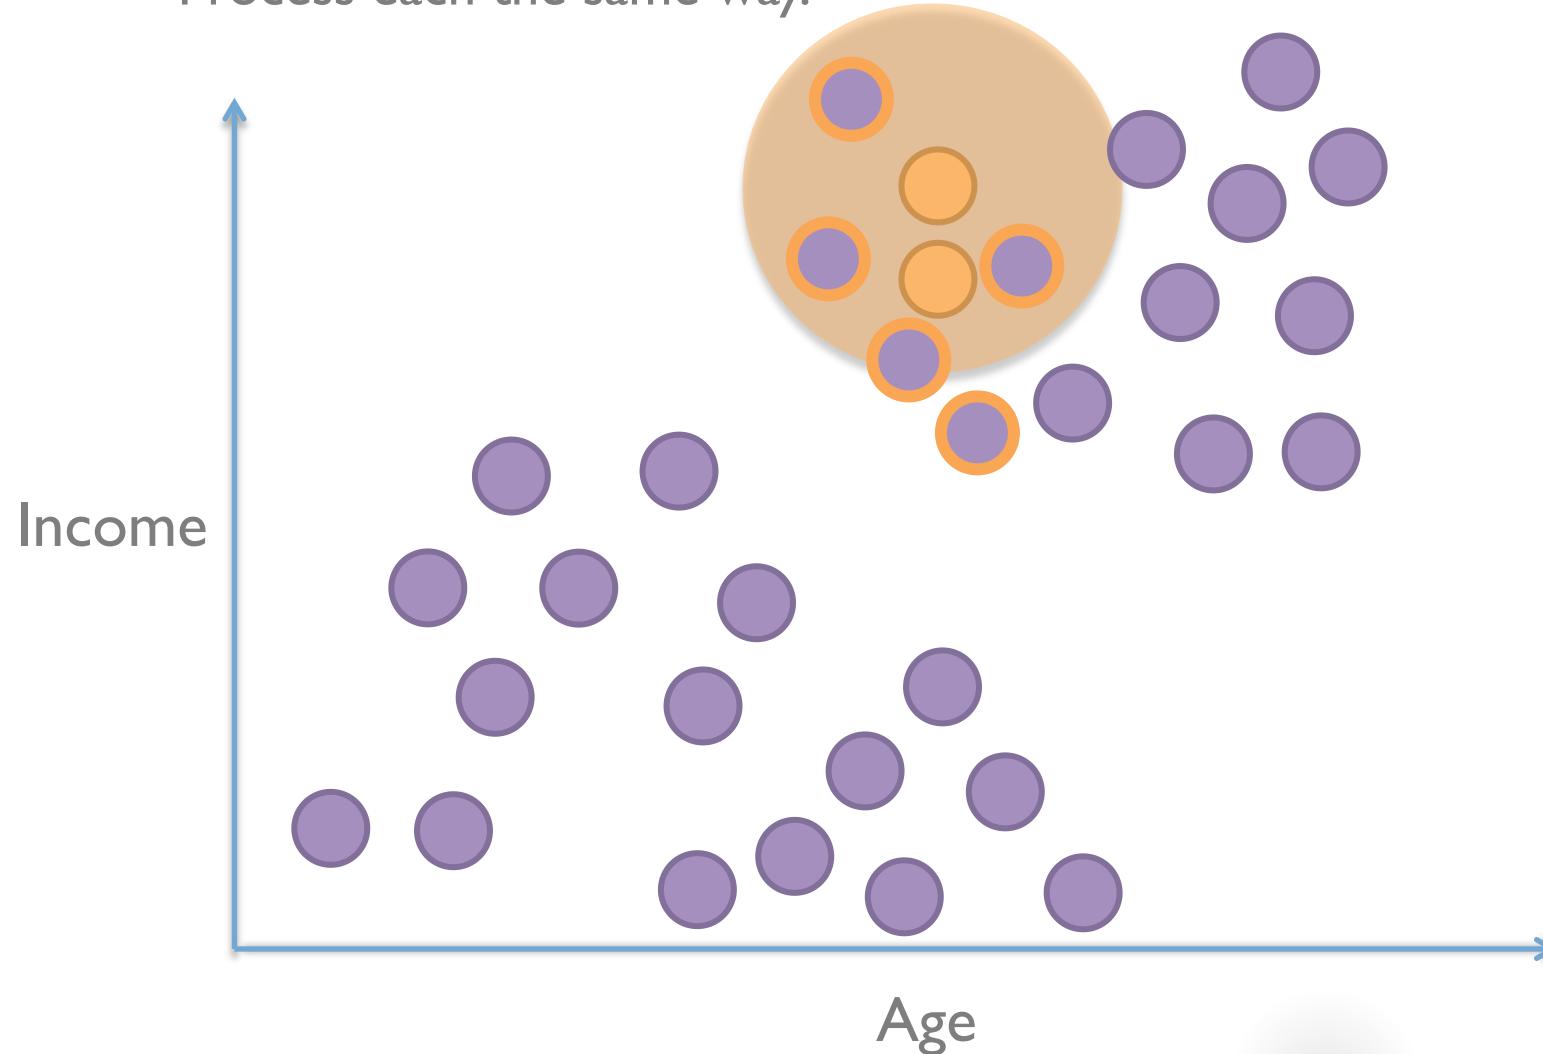
Everybody within  $\text{epsilon}$  are part of the cluster.  
Process each the same way.



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

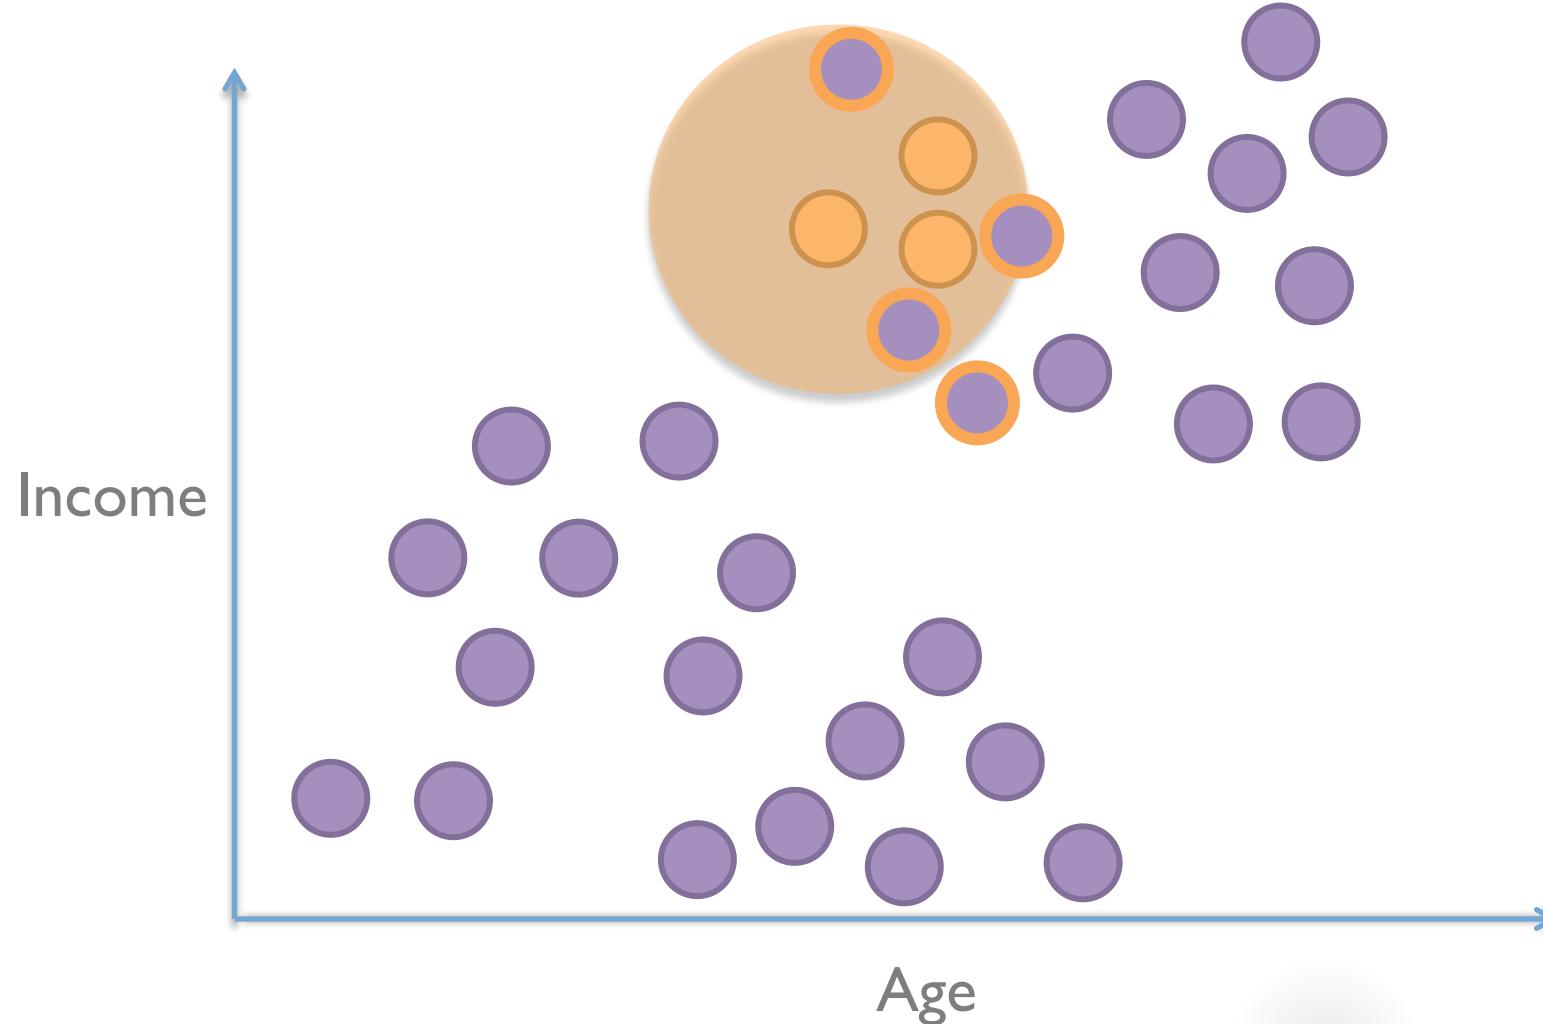
Everybody within  $\text{epsilon}$  are part of the cluster.  
Process each the same way.



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

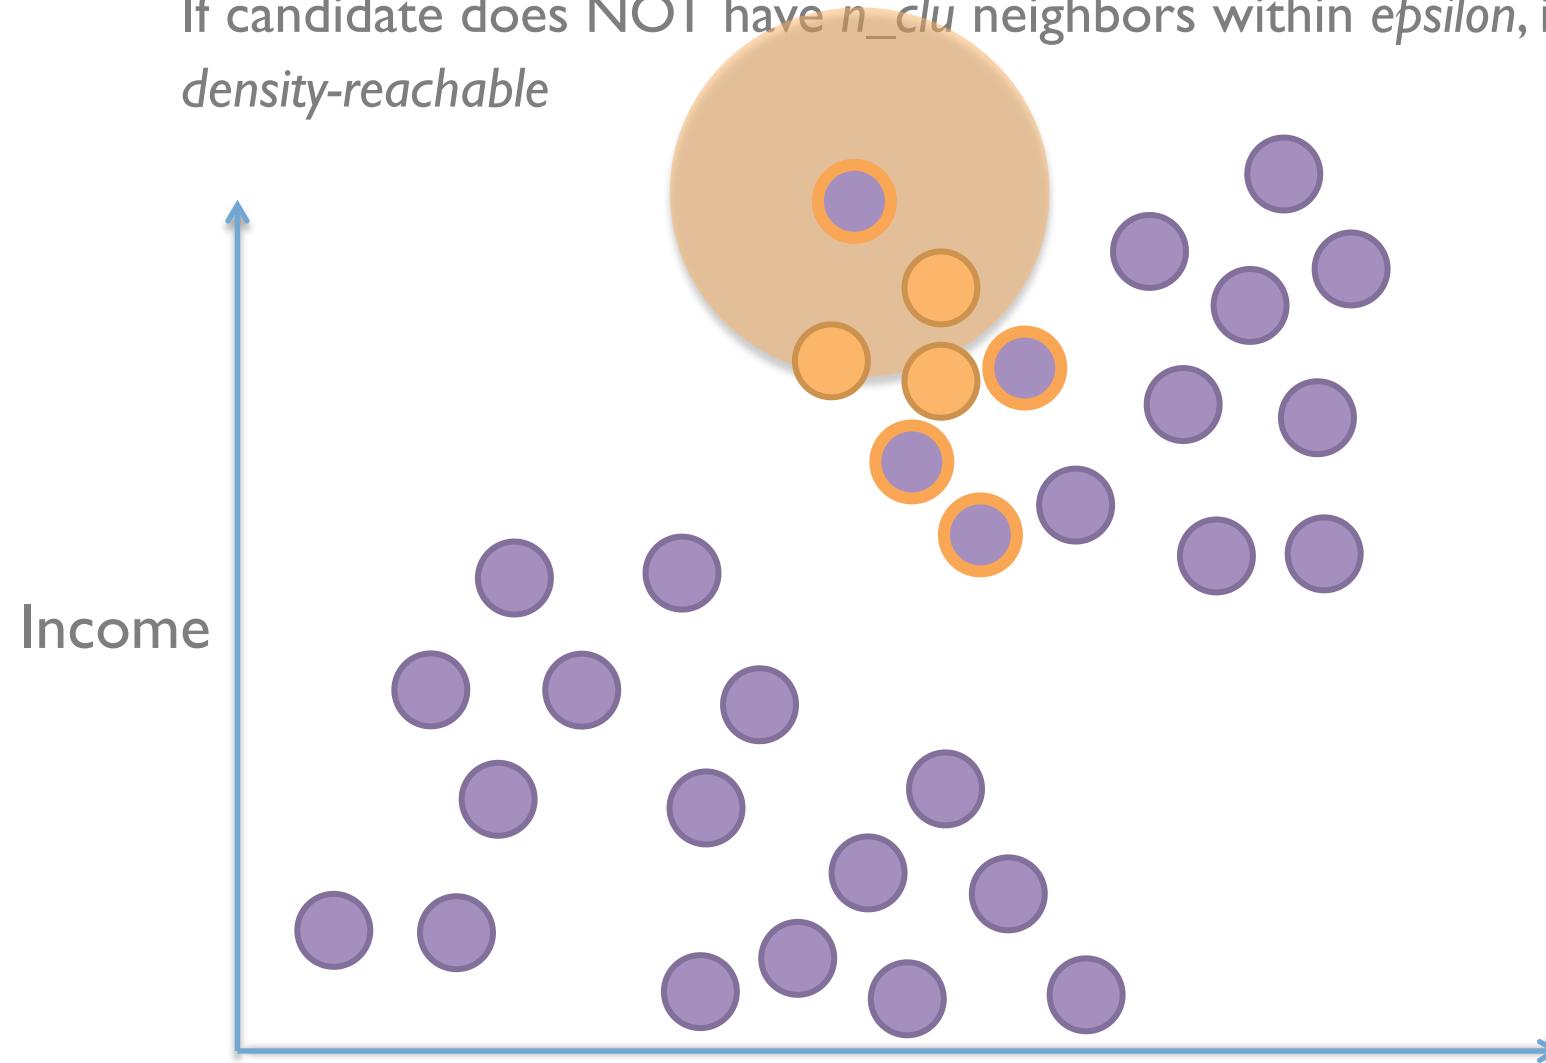
Everybody within  $\text{epsilon}$  are part of the cluster.  
Process each the same way.



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

If candidate does NOT have  $n_{\text{clu}}$  neighbors within  $\text{epsilon}$ , it is *density-reachable*



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

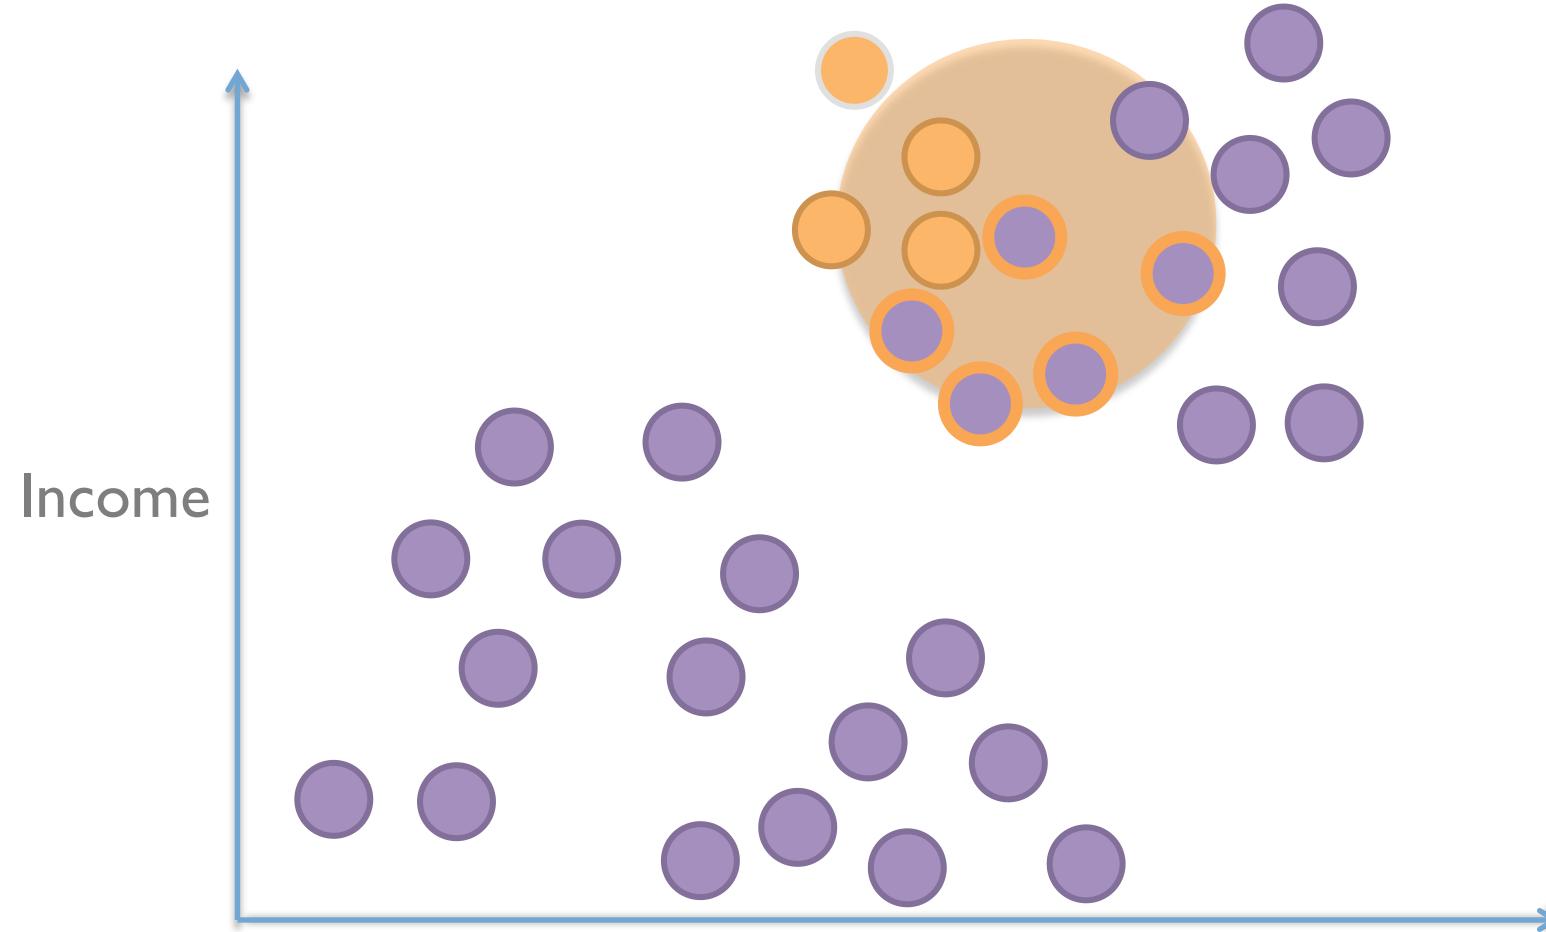
If candidate does NOT have  $n_{\text{clu}}$  neighbors within  $\text{epsilon}$ , it is *density-reachable*



epsilon = 1.75  
n\_clu = 3

## DBSCAN

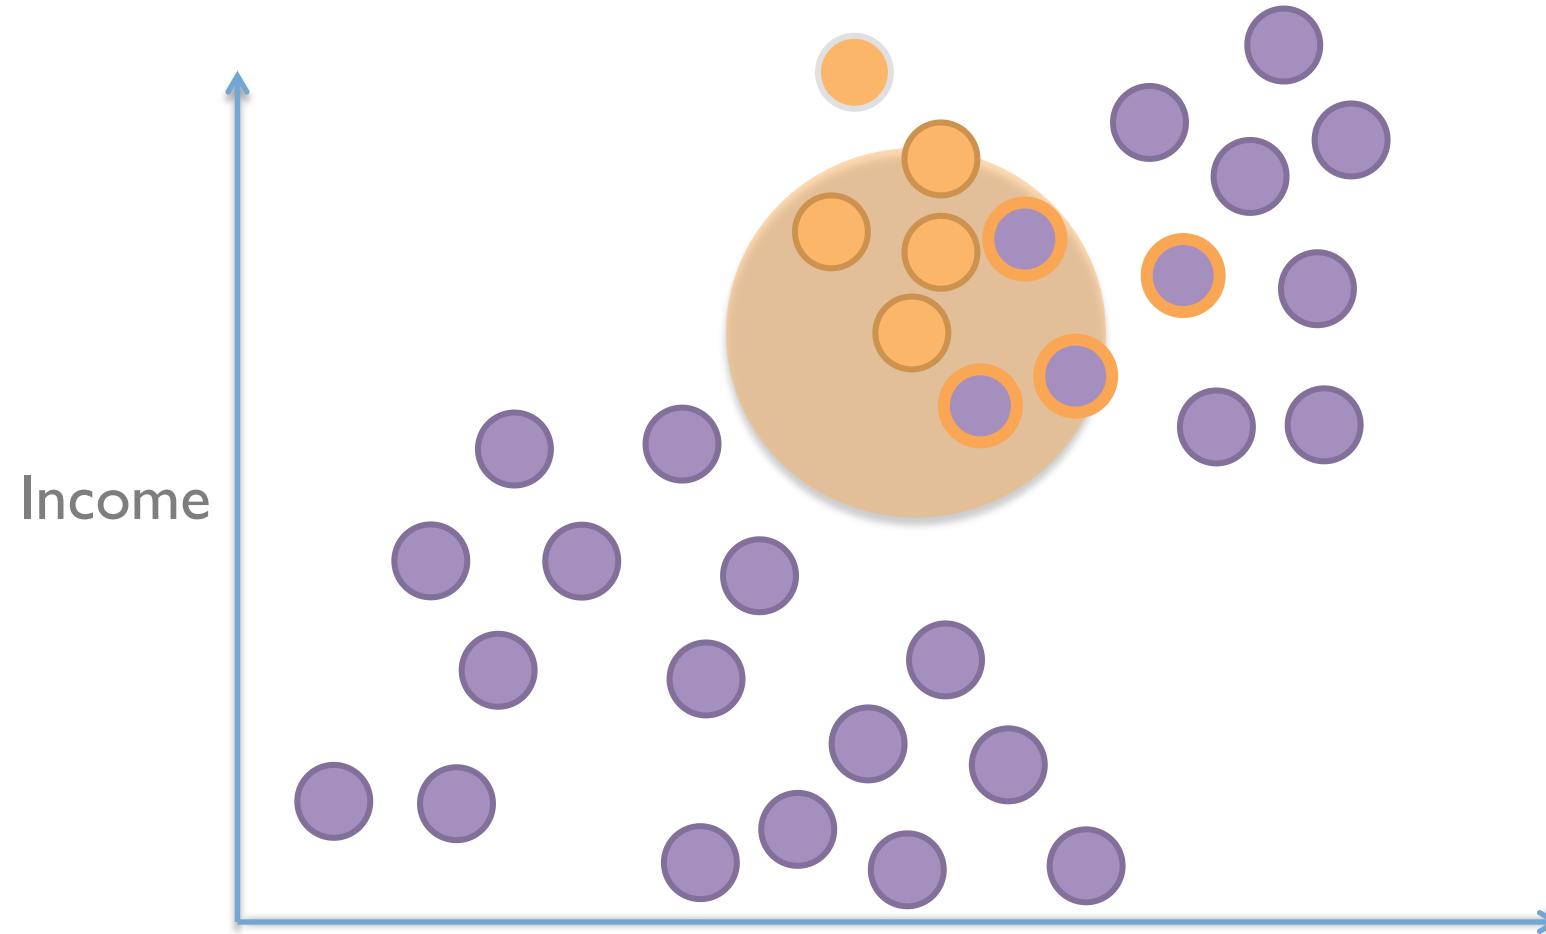
Keep going until you find the entire cluster and no point is left unvisited by this chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

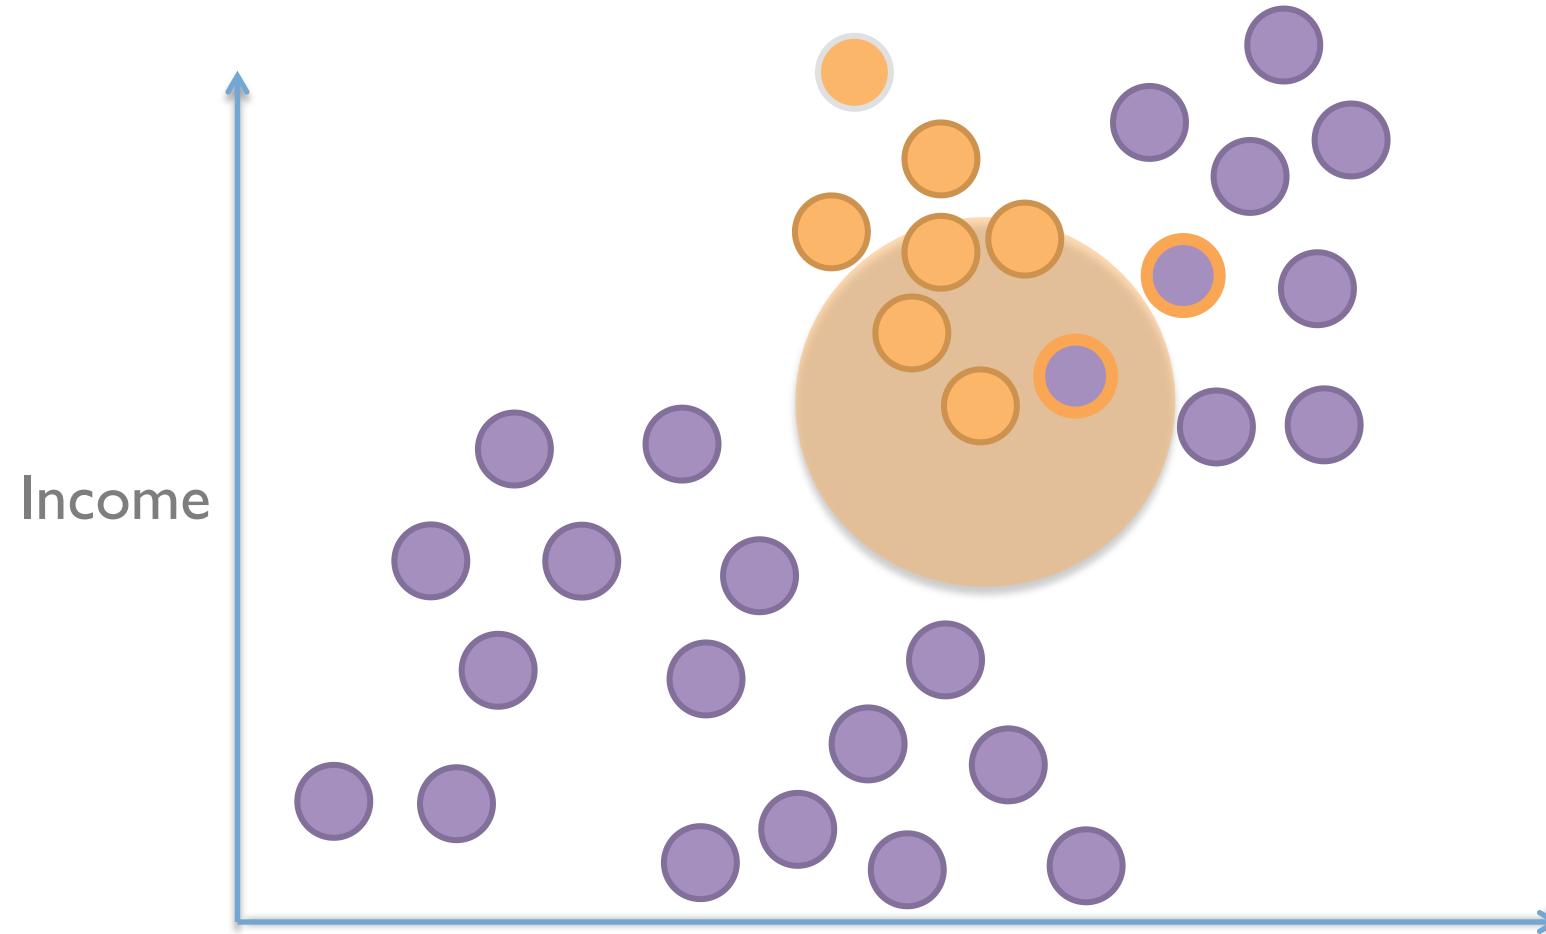
Keep going until you find the entire cluster and no point is left unvisited by this chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

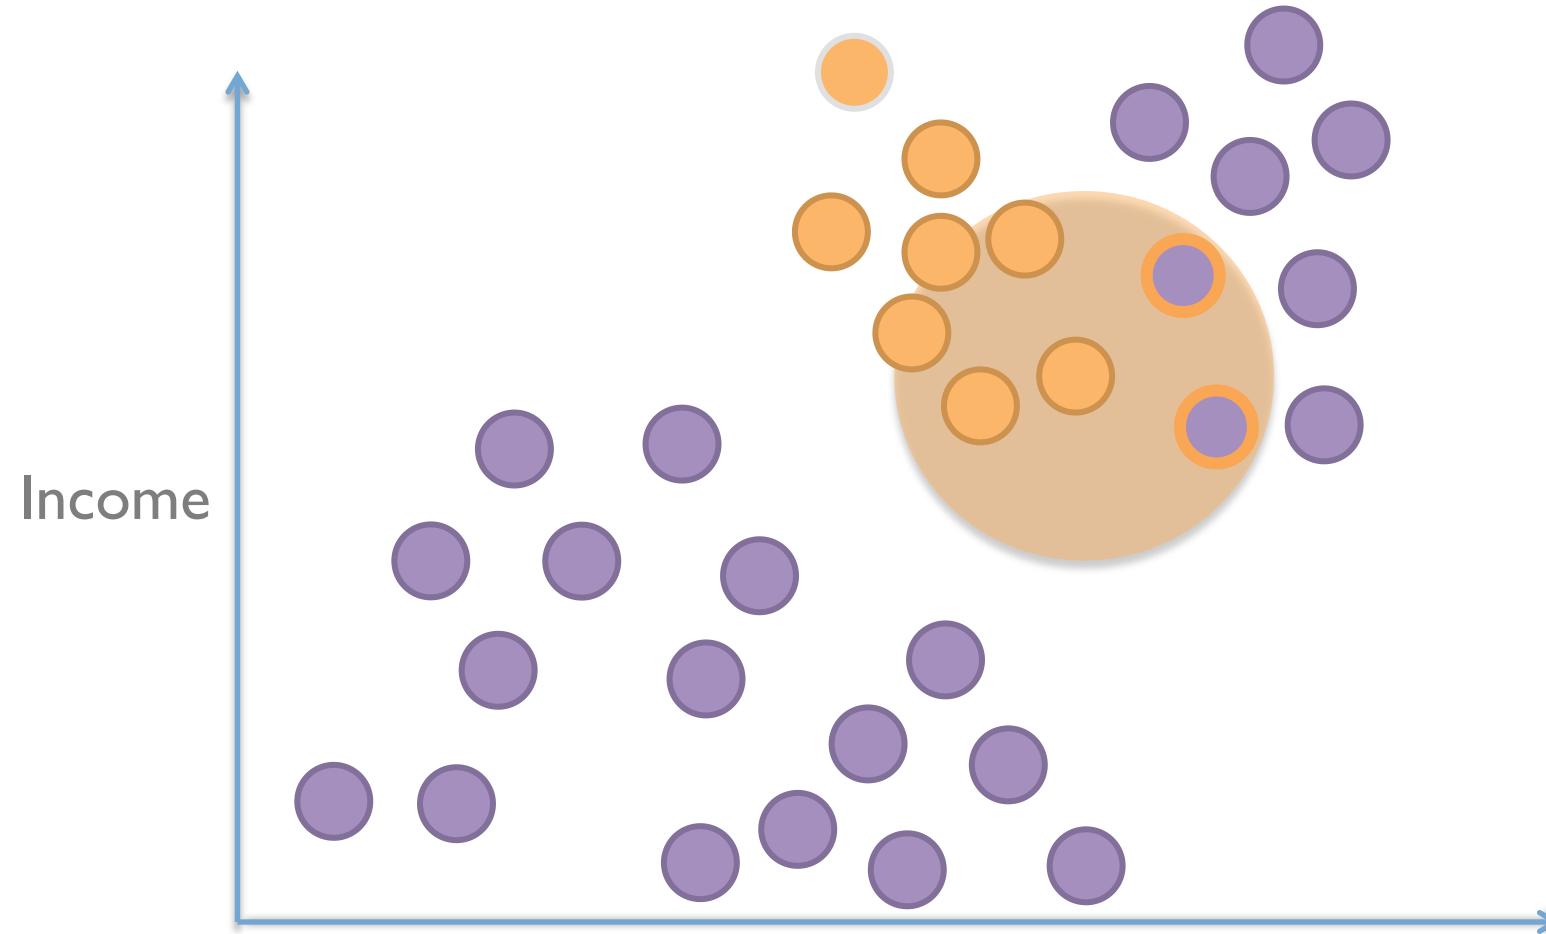
Keep going until you find the entire cluster and no point is left unvisited by this chain reaction



epsilon = 1.75  
n\_clu = 3

## DBSCAN

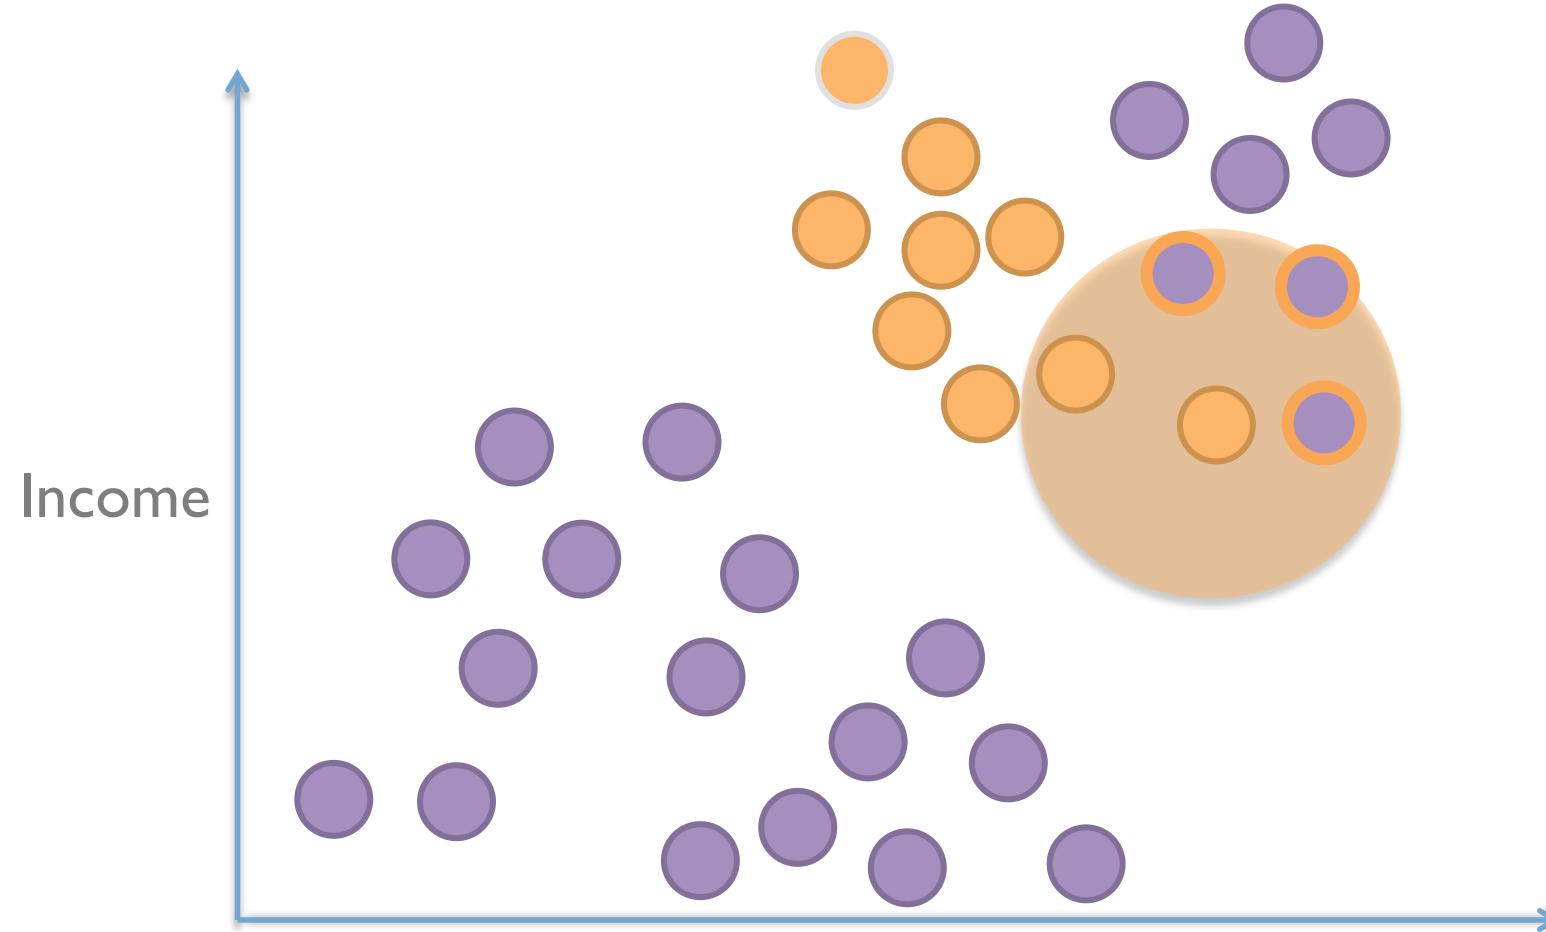
Keep going until you find the entire cluster and no point is left unvisited by this chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

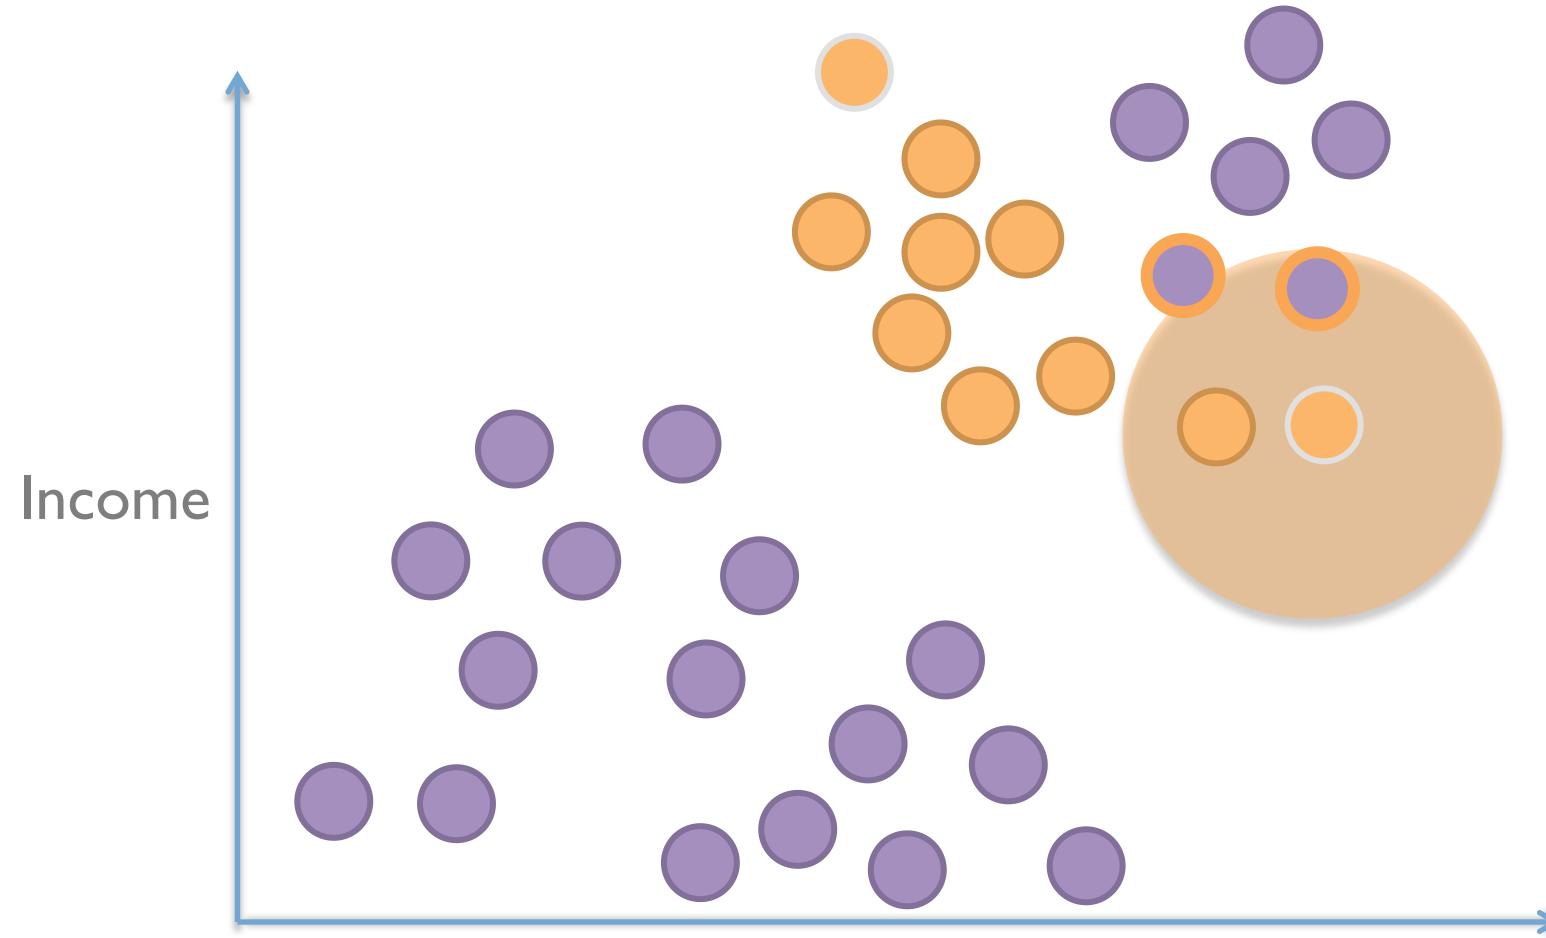
Keep going until you find the entire cluster and no point is left unvisited by this chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

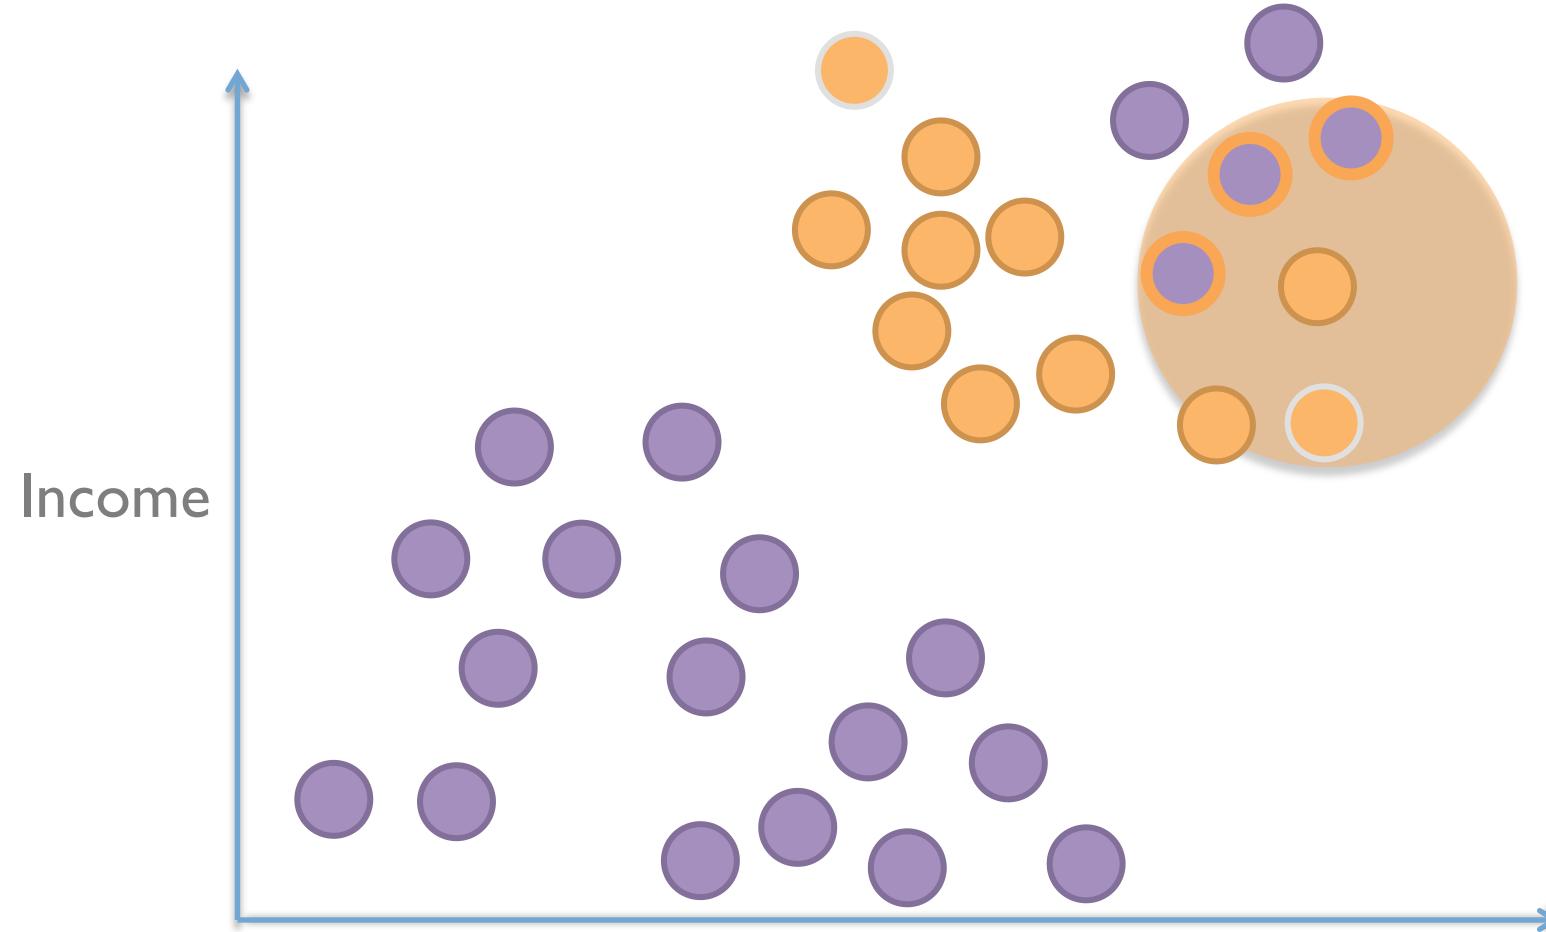
Keep going until you find the entire cluster and no point is left unvisited by this chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

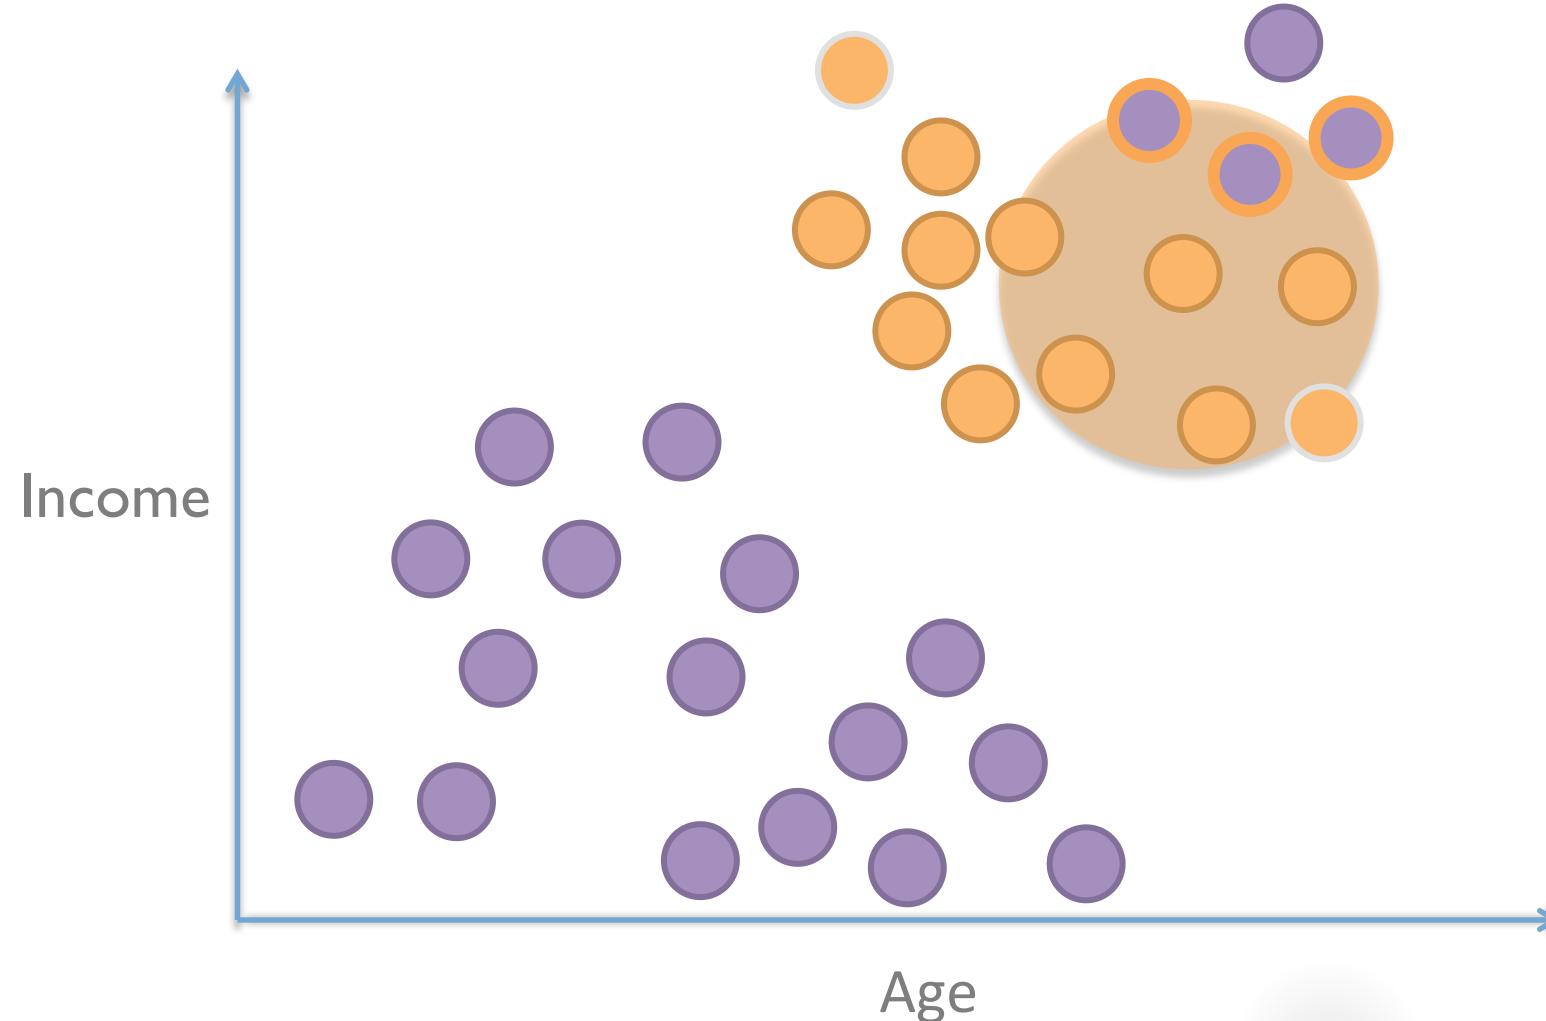
Keep going until you find the entire cluster and no point is left unvisited by this chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

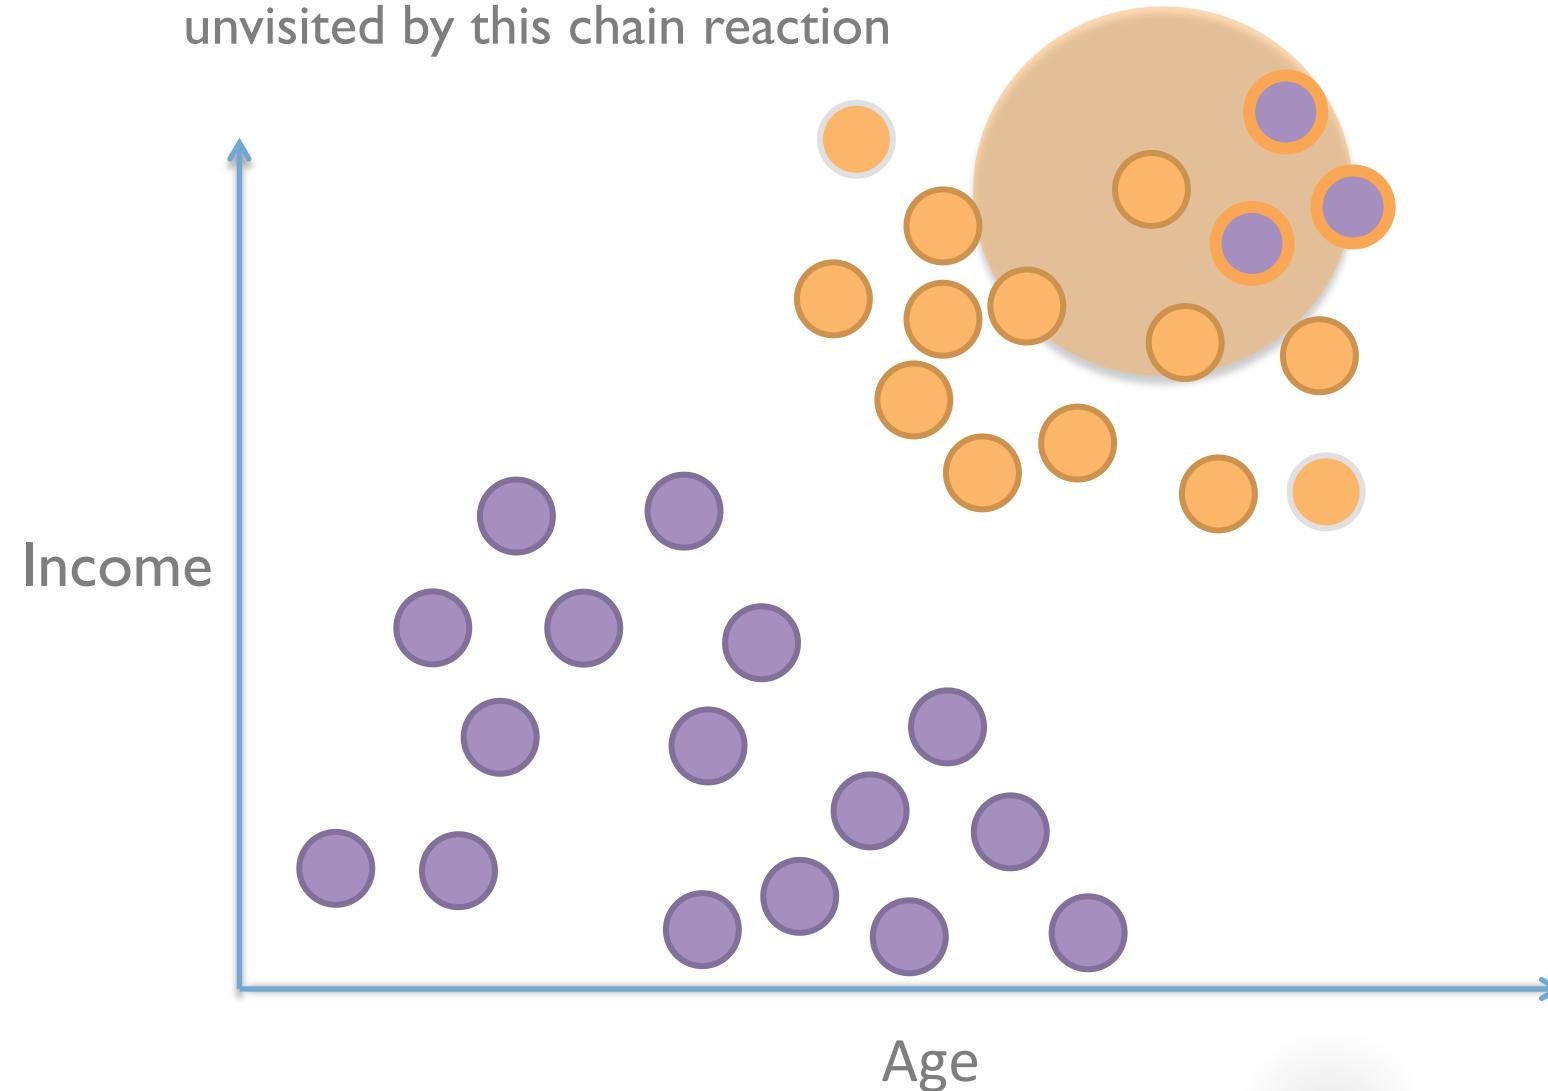
Keep going until you find the entire cluster and no point is left unvisited by this chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

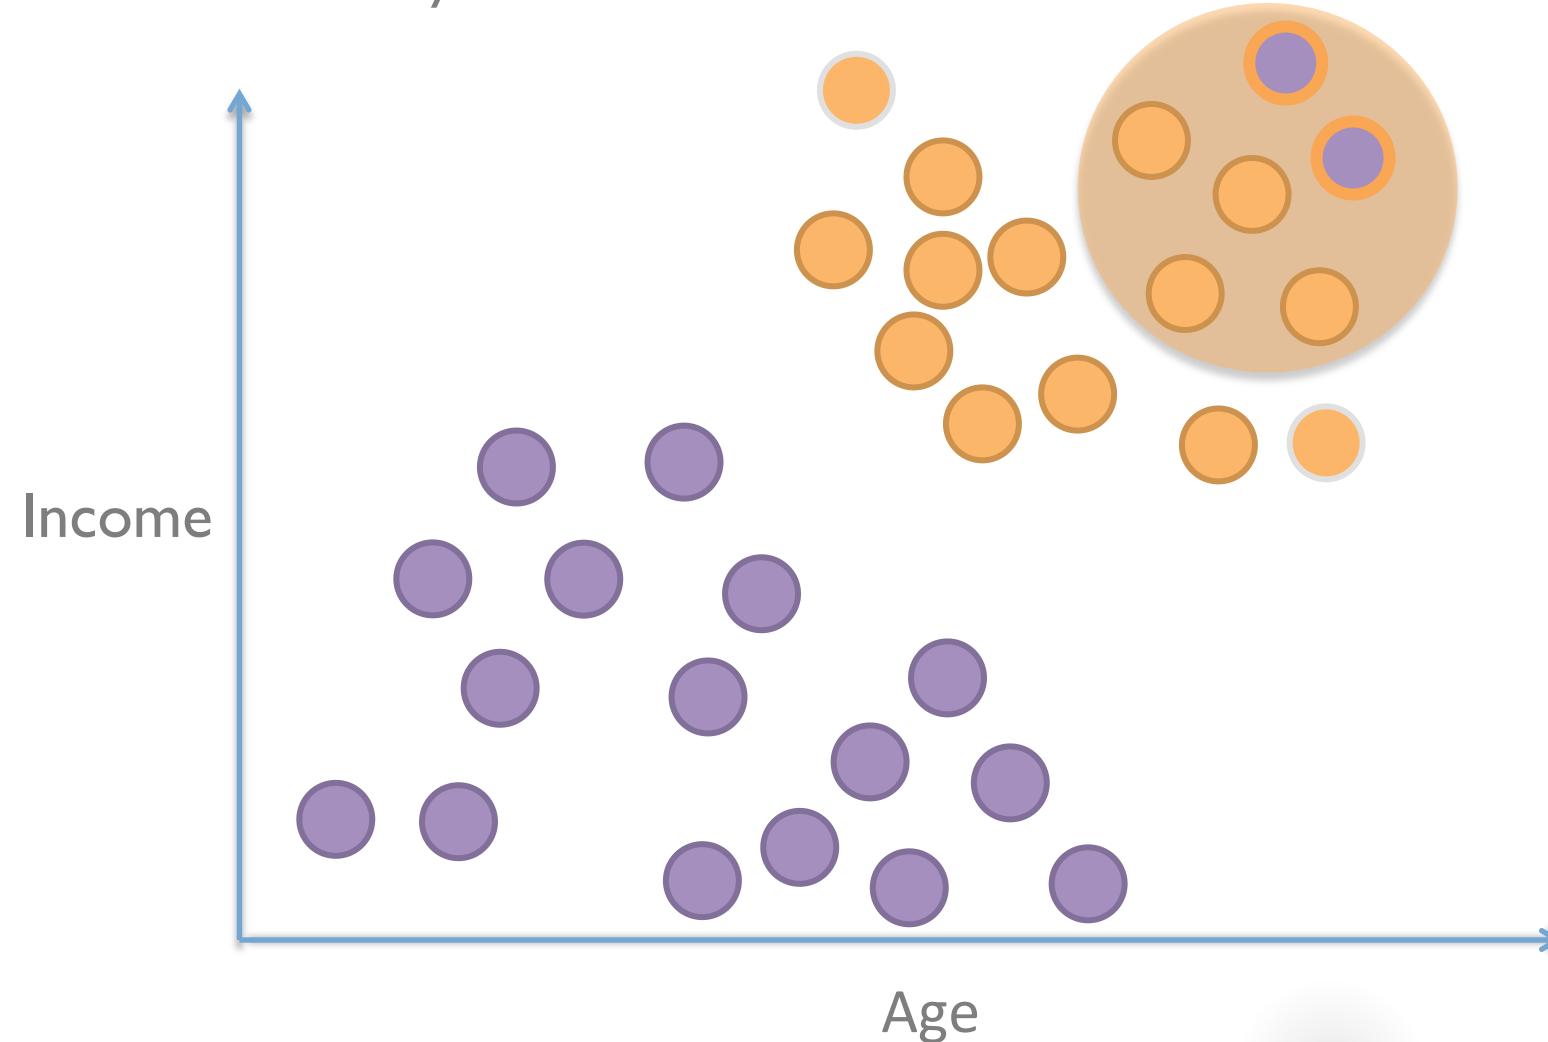
Keep going until you find the entire cluster and no point is left unvisited by this chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

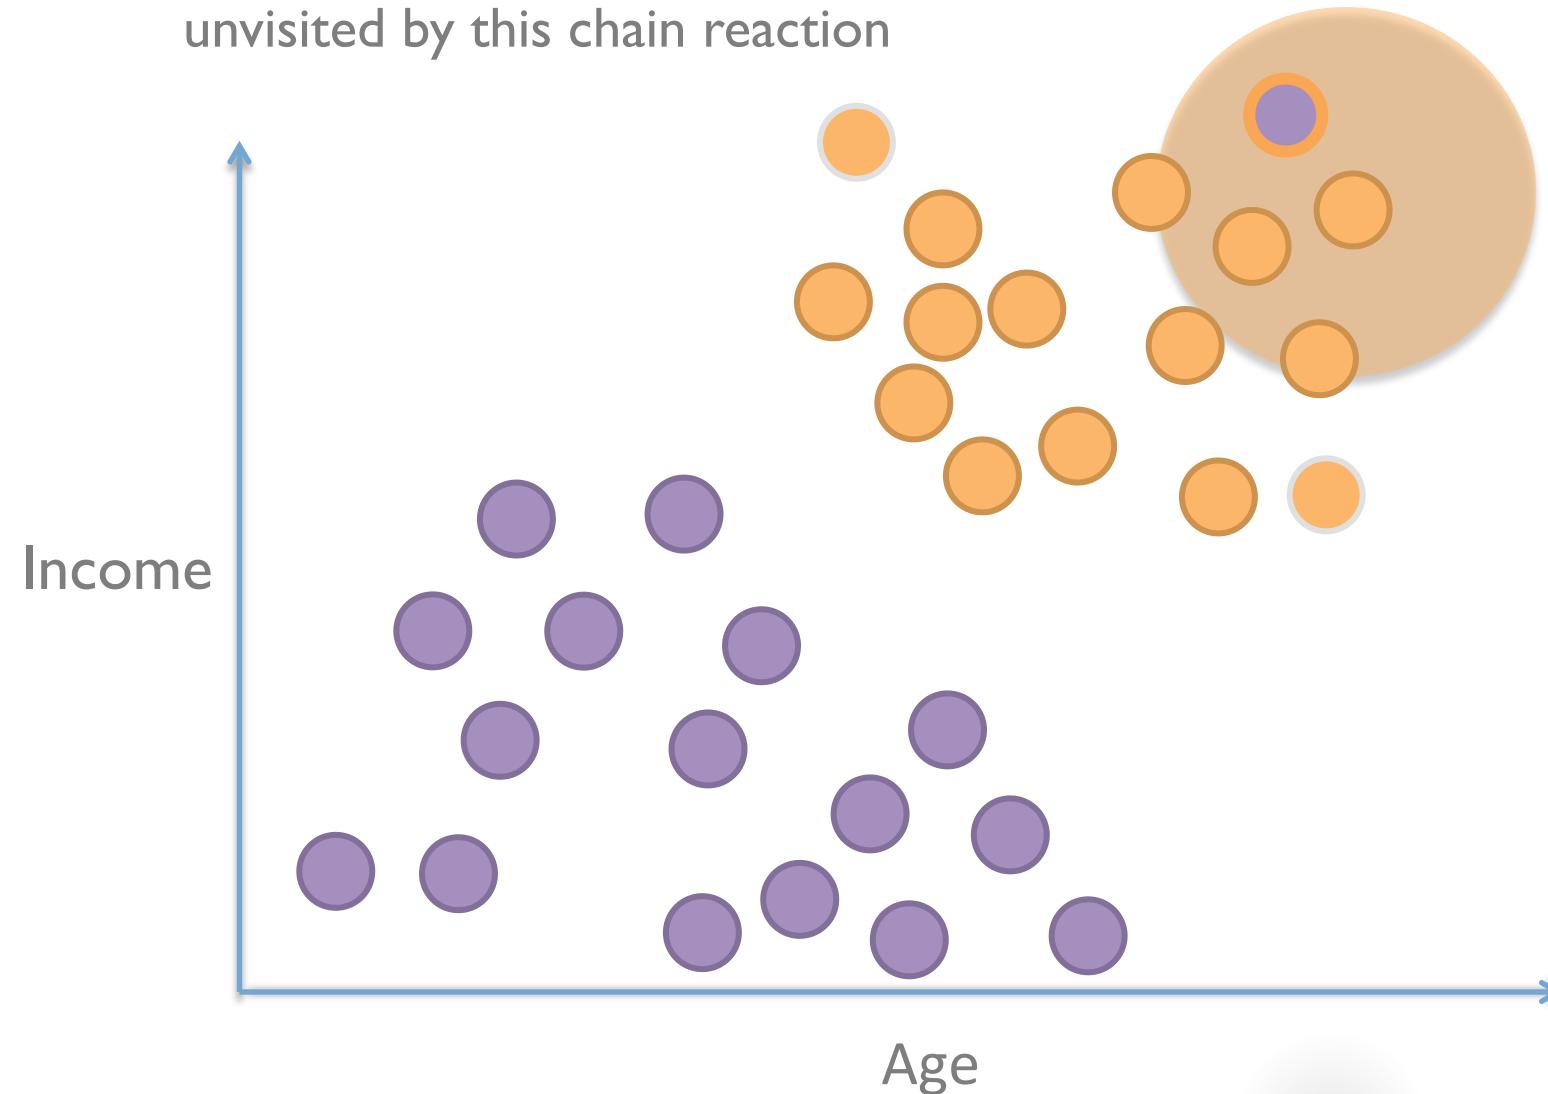
Keep going until you find the entire cluster and no point is left unvisited by this chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

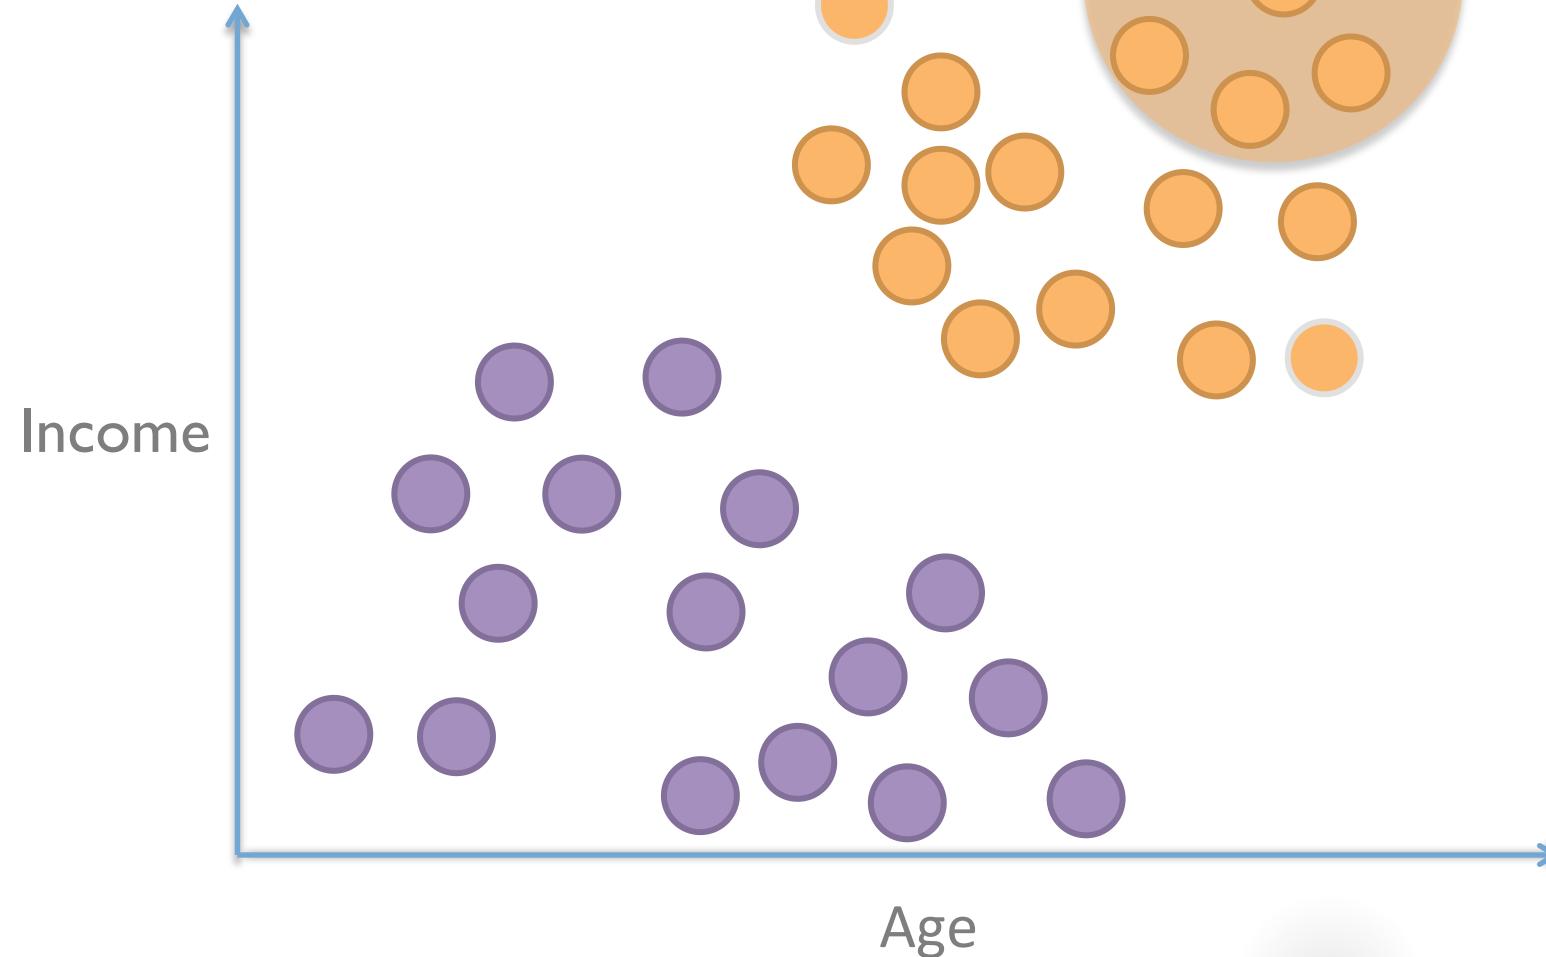
Keep going until you find the entire cluster and no point is left unvisited by this chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

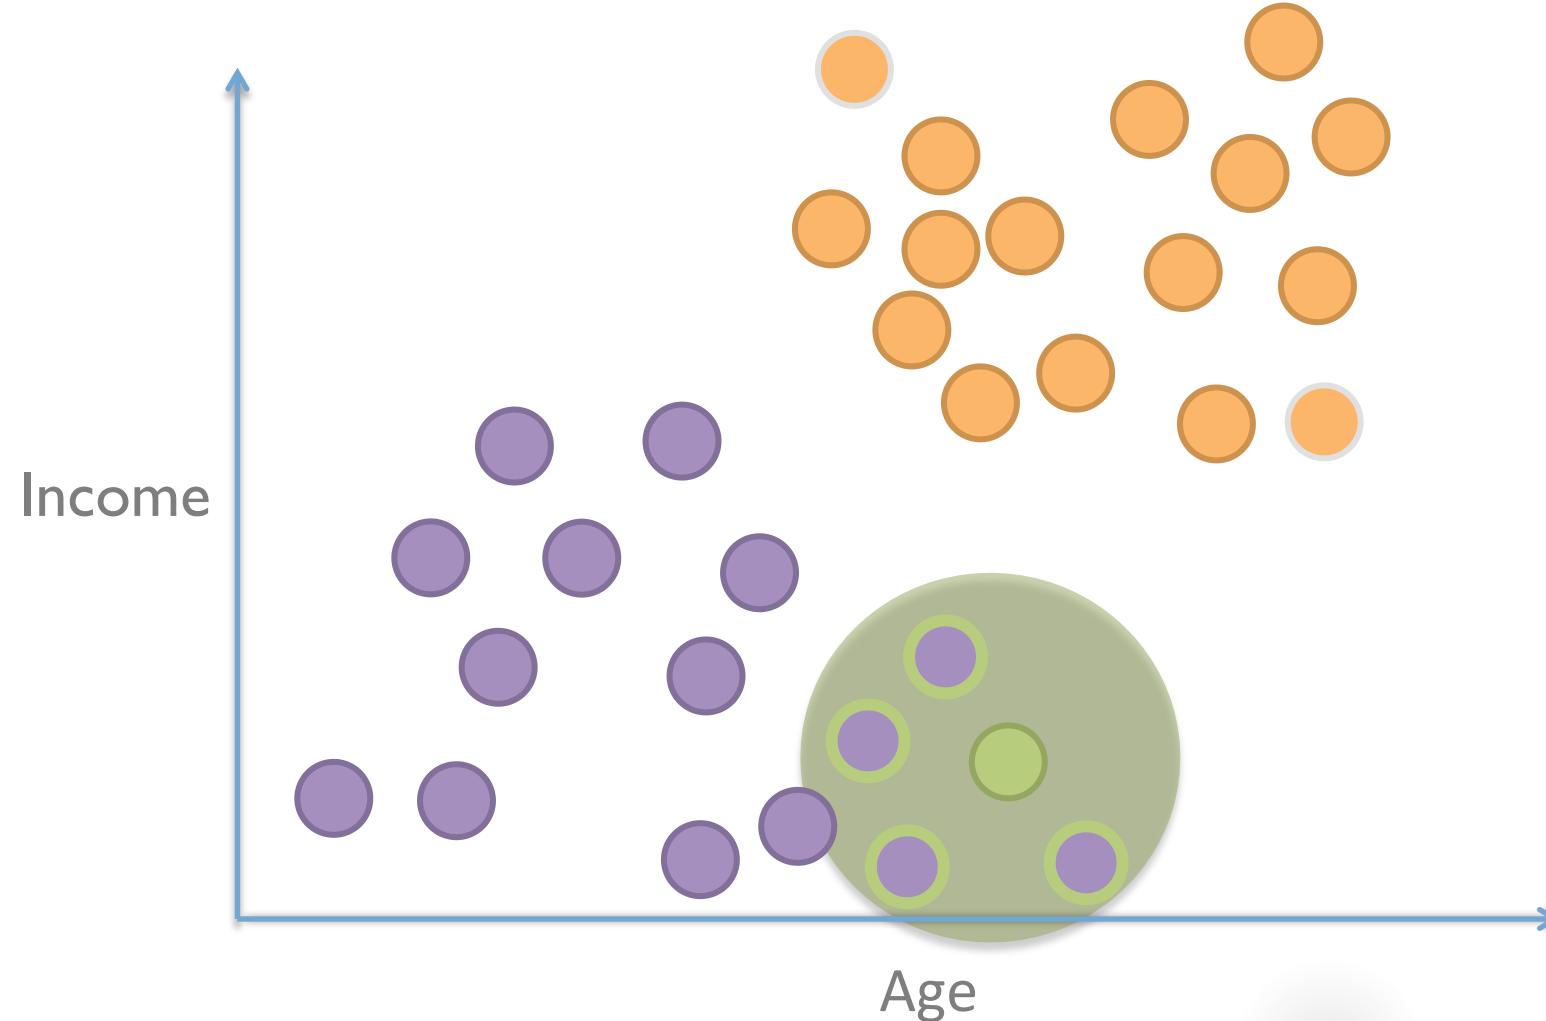
If no neighbors left, randomly try a new (unvisited) point to potentially start a new cluster



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

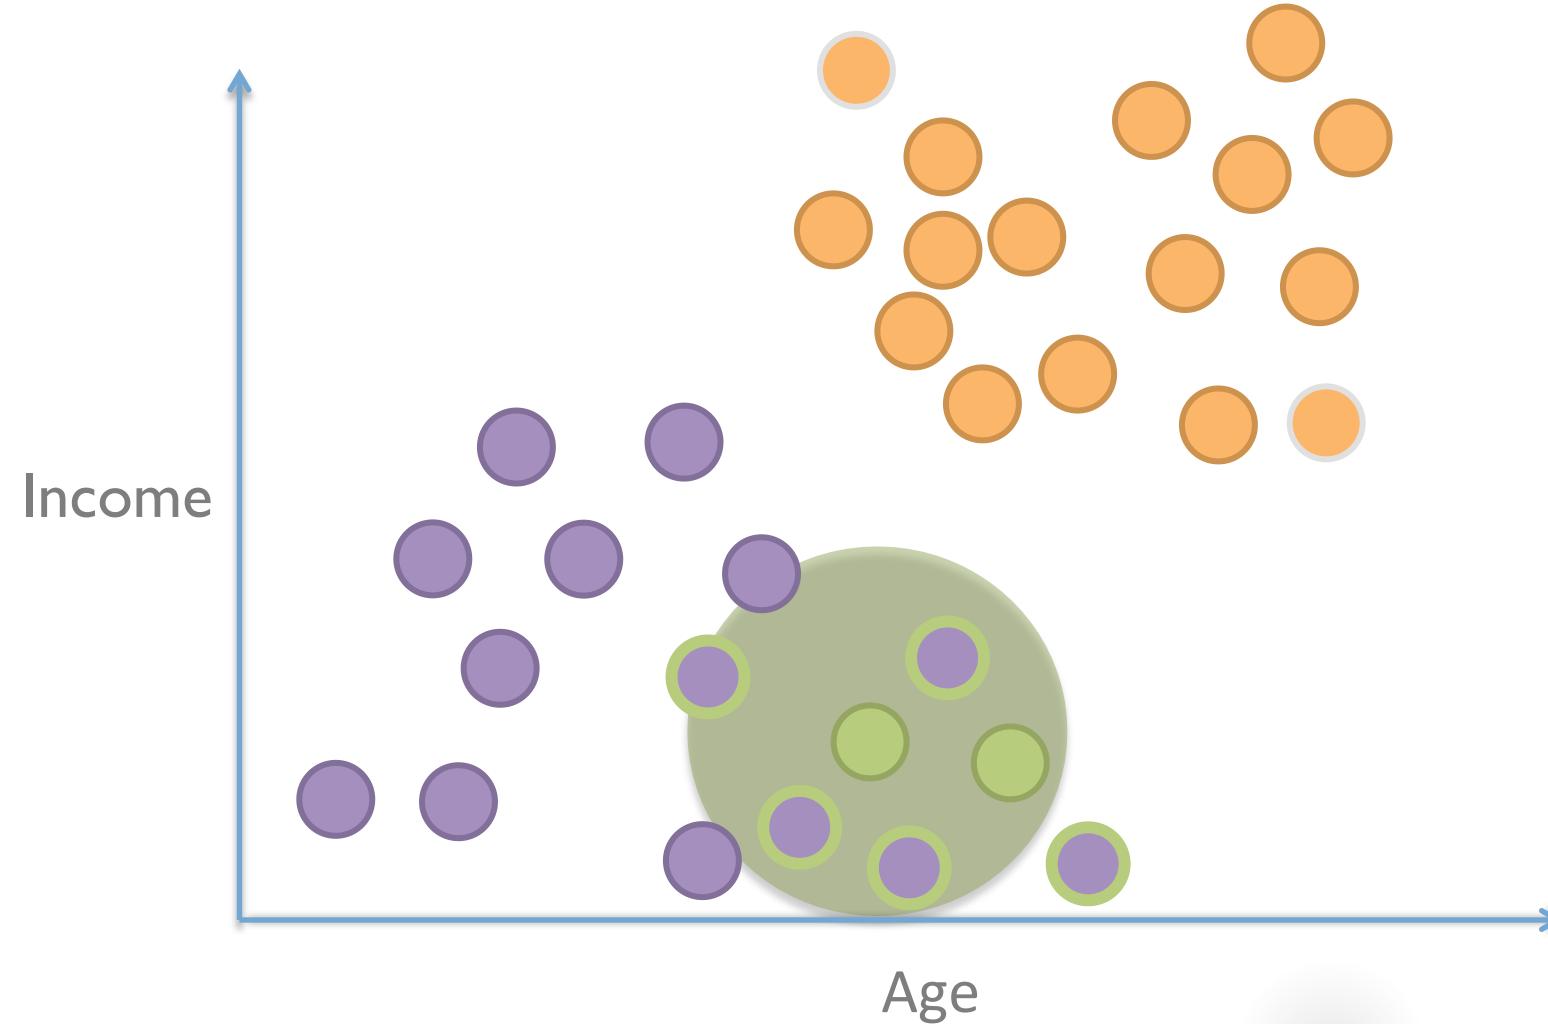
If no neighbors left, randomly try a new (unvisited) point to potentially start a new cluster



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

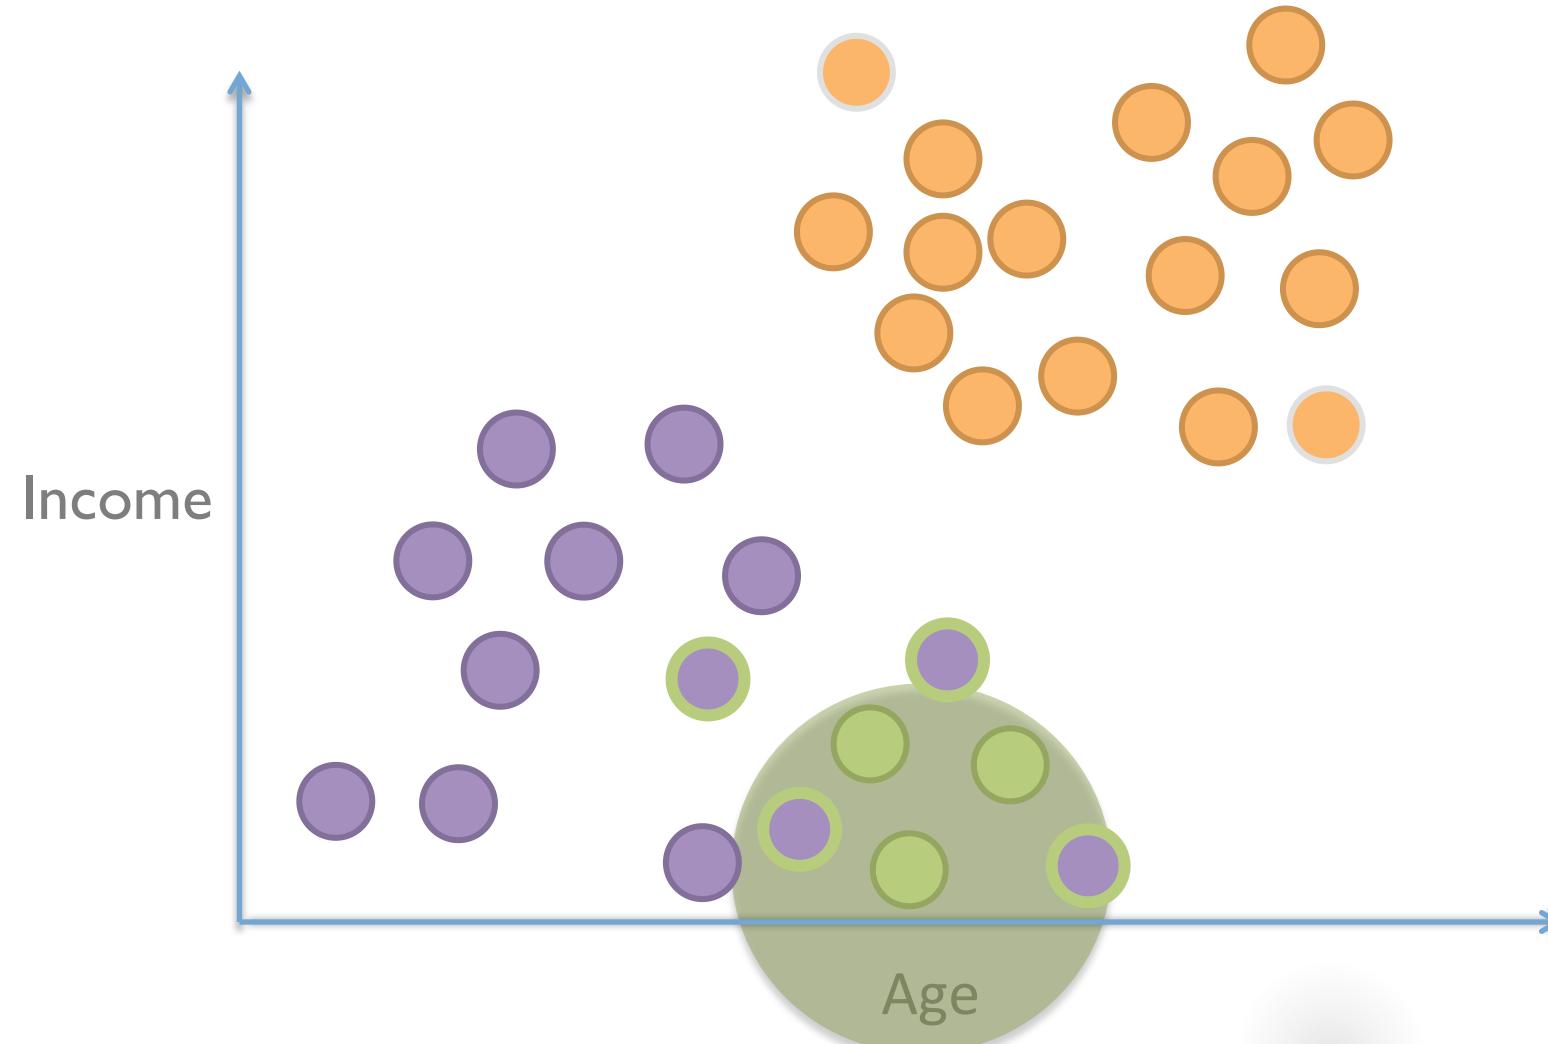
Keep adding neighbors within  $\text{epsilon}$  as chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

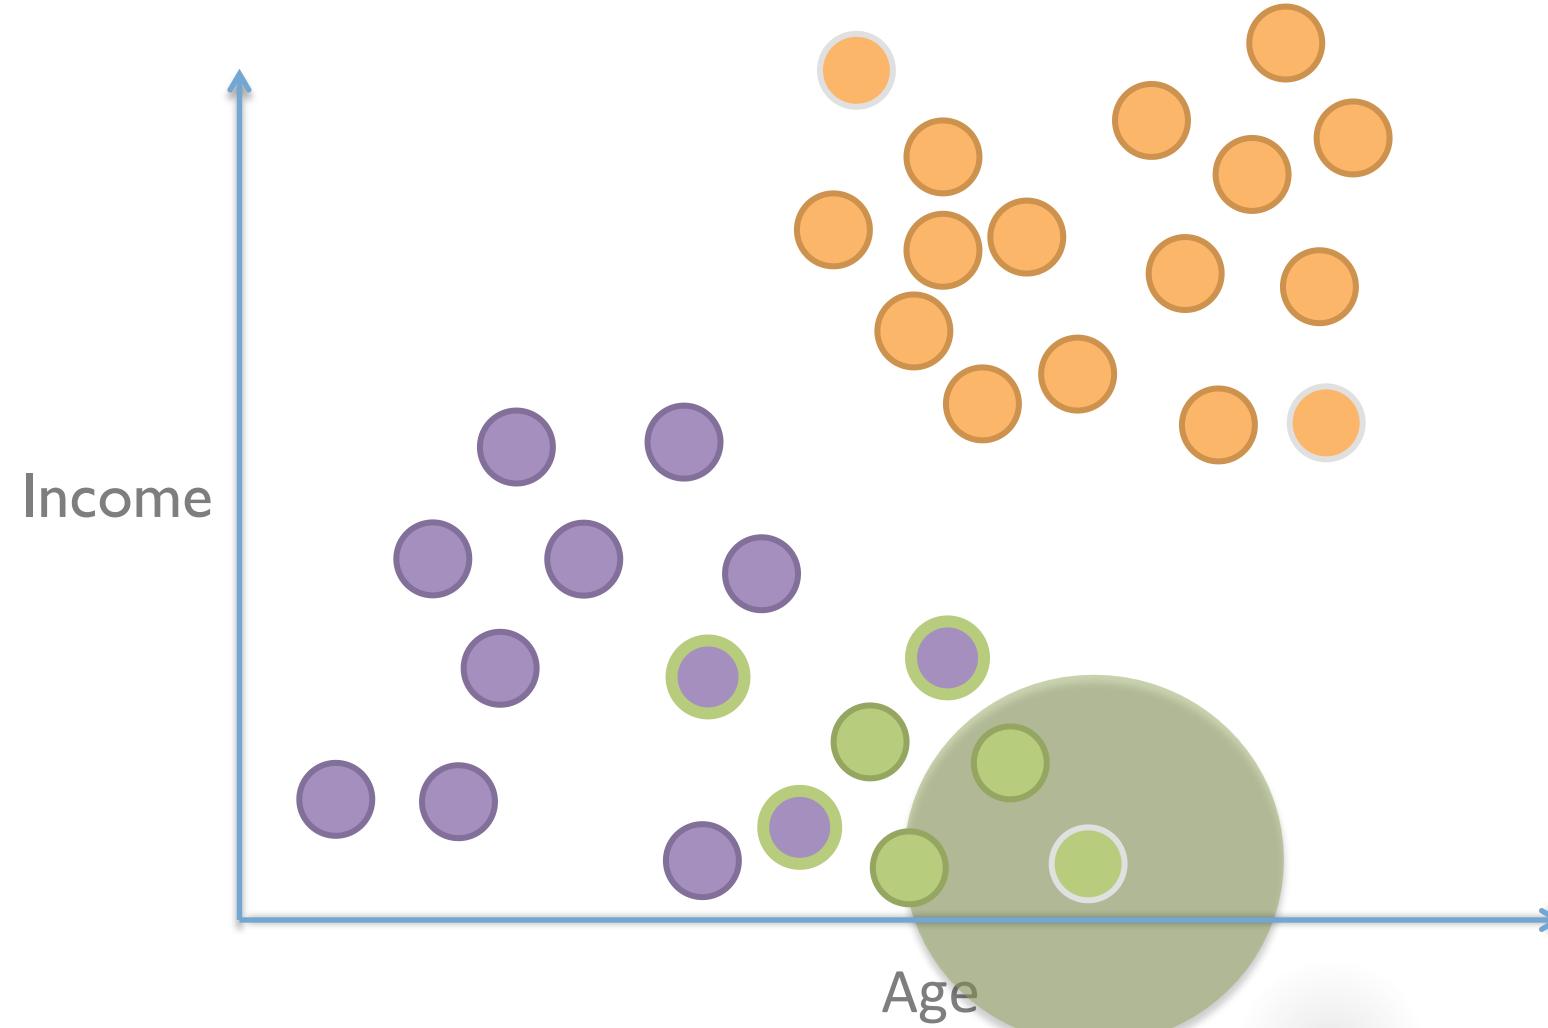
Keep adding neighbors within  $\text{epsilon}$  as chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

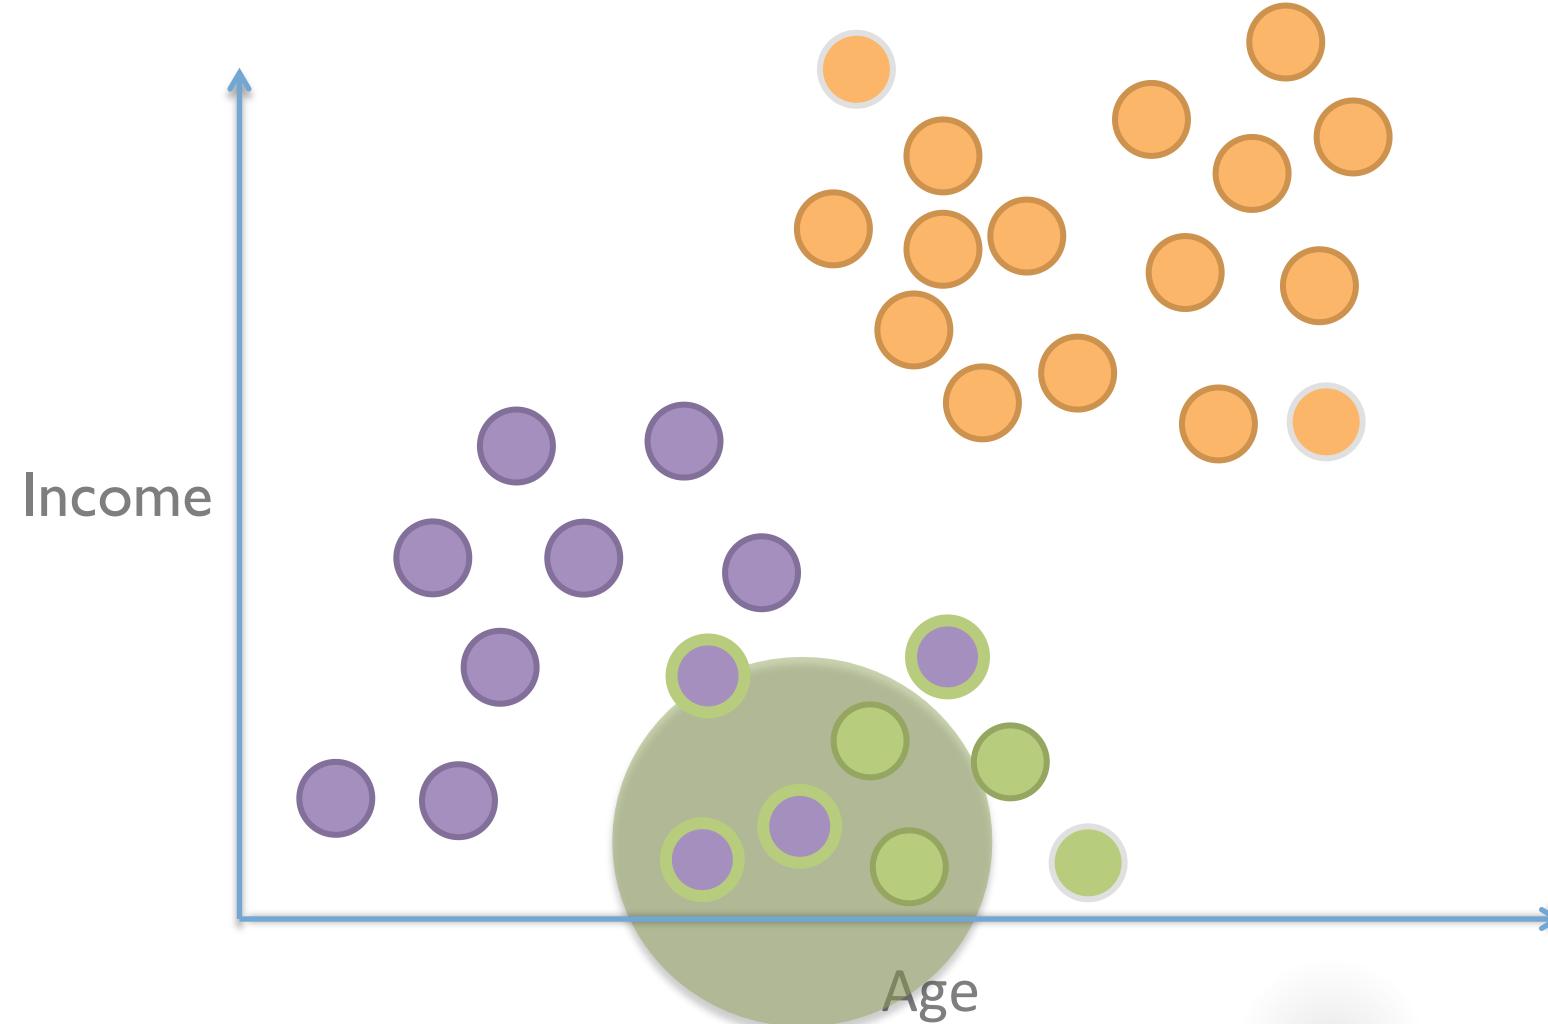
Keep adding neighbors within  $\text{epsilon}$  as chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

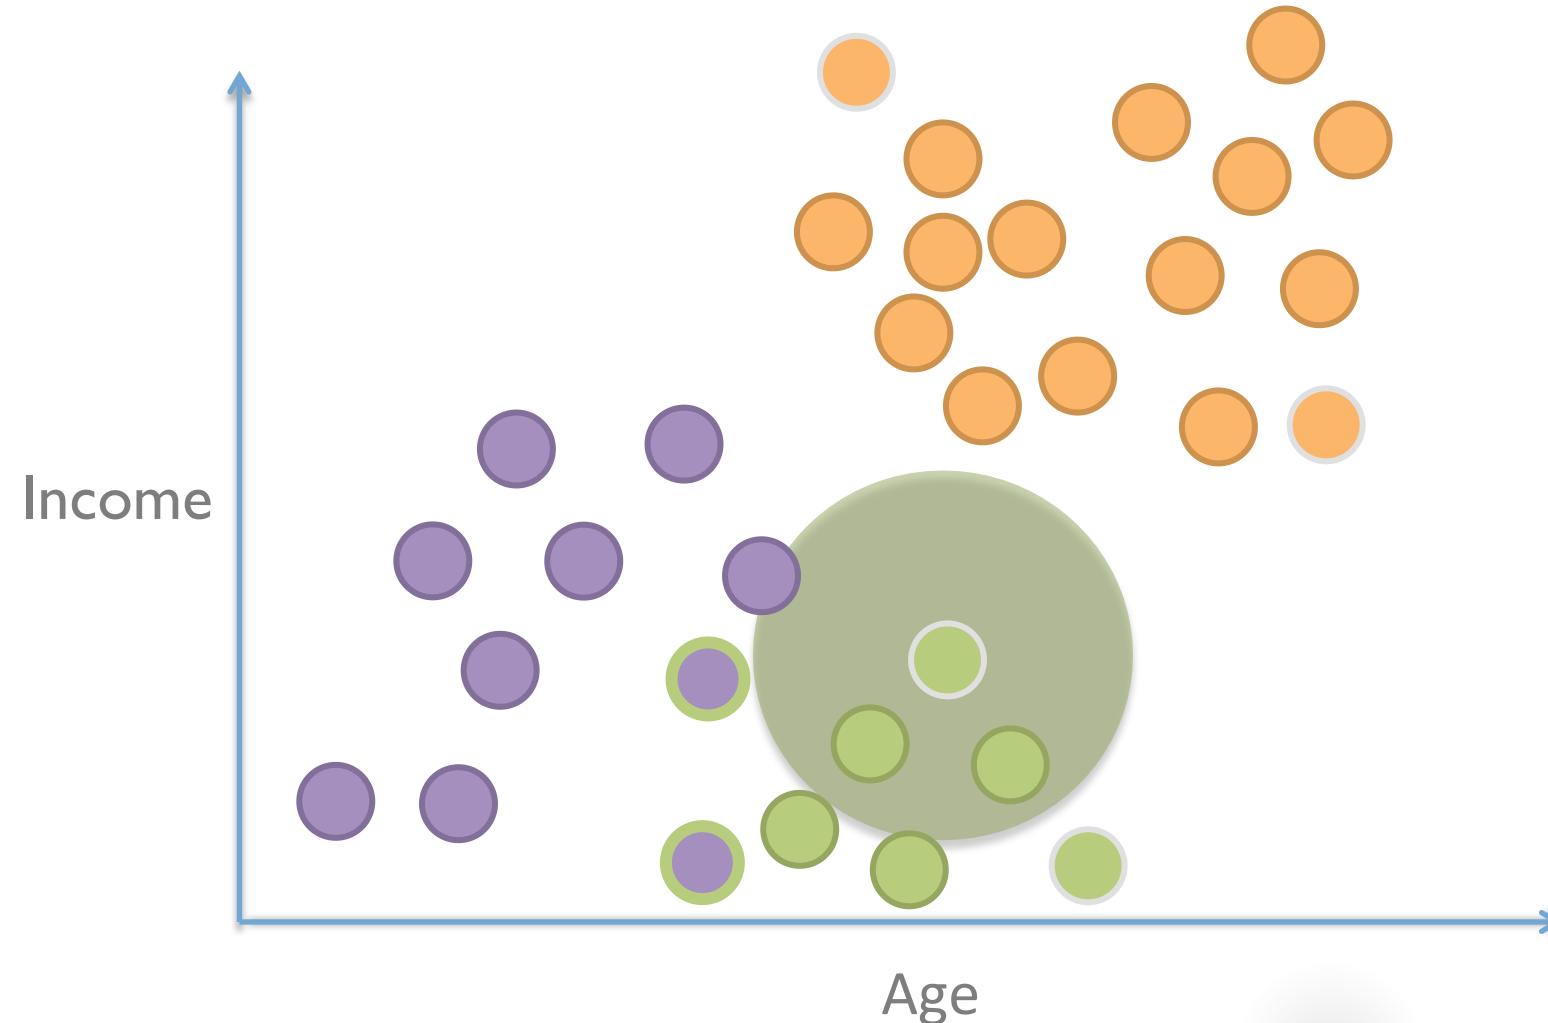
Keep adding neighbors within  $\text{epsilon}$  as chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

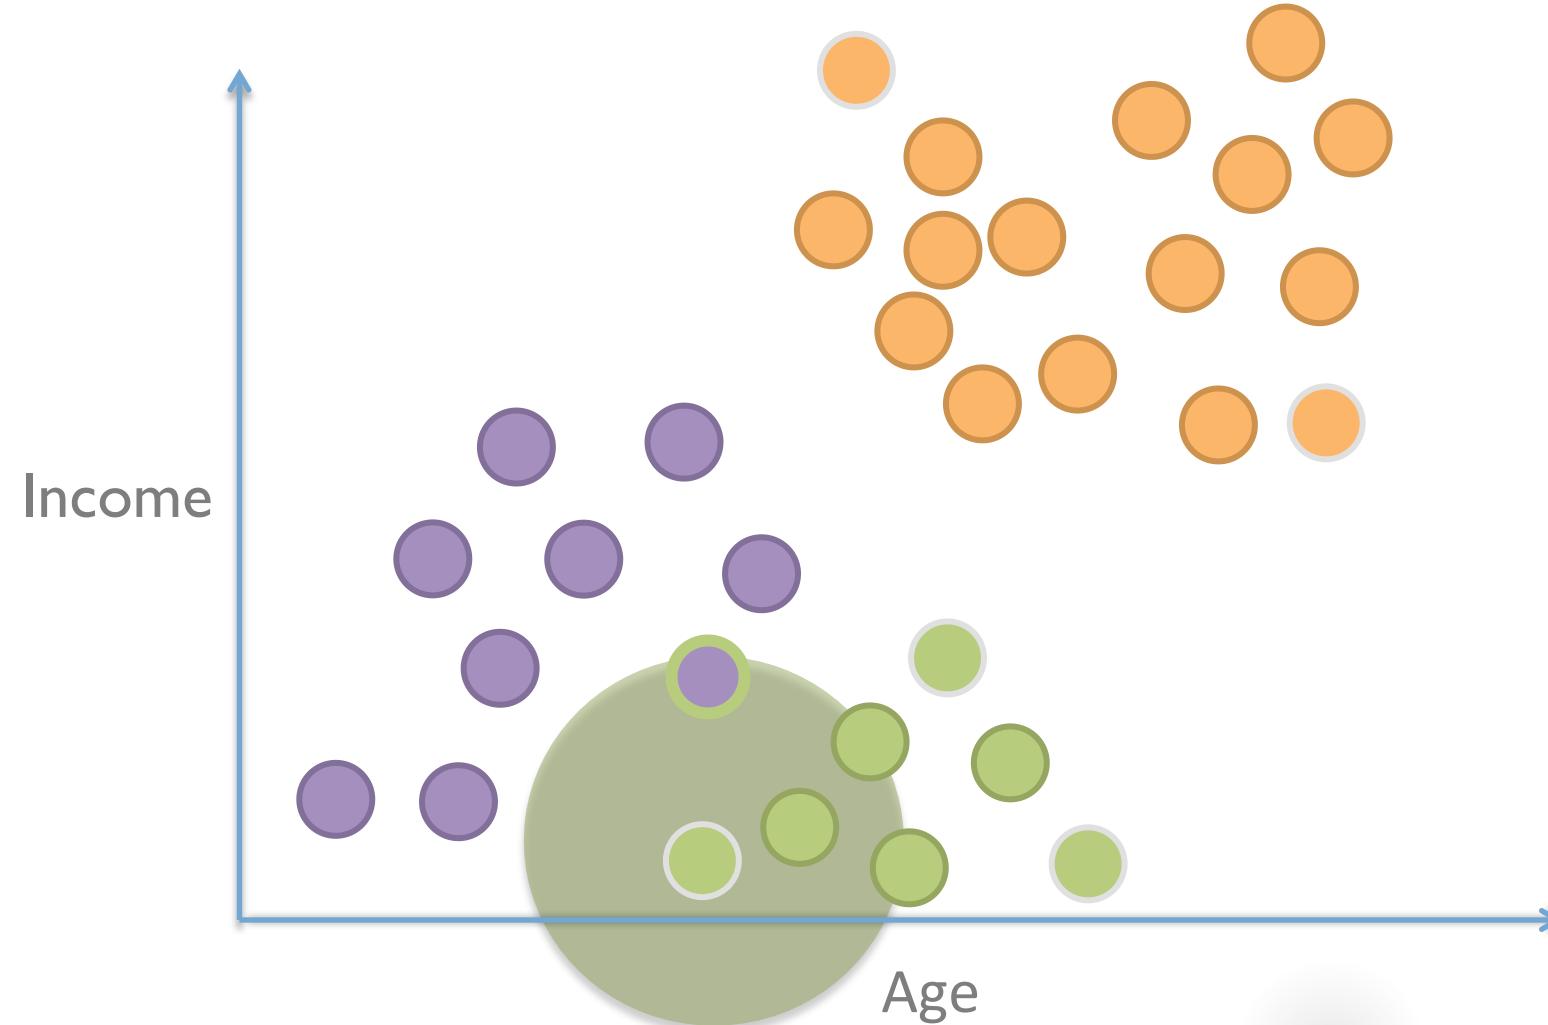
Keep adding neighbors within  $\text{epsilon}$  as chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

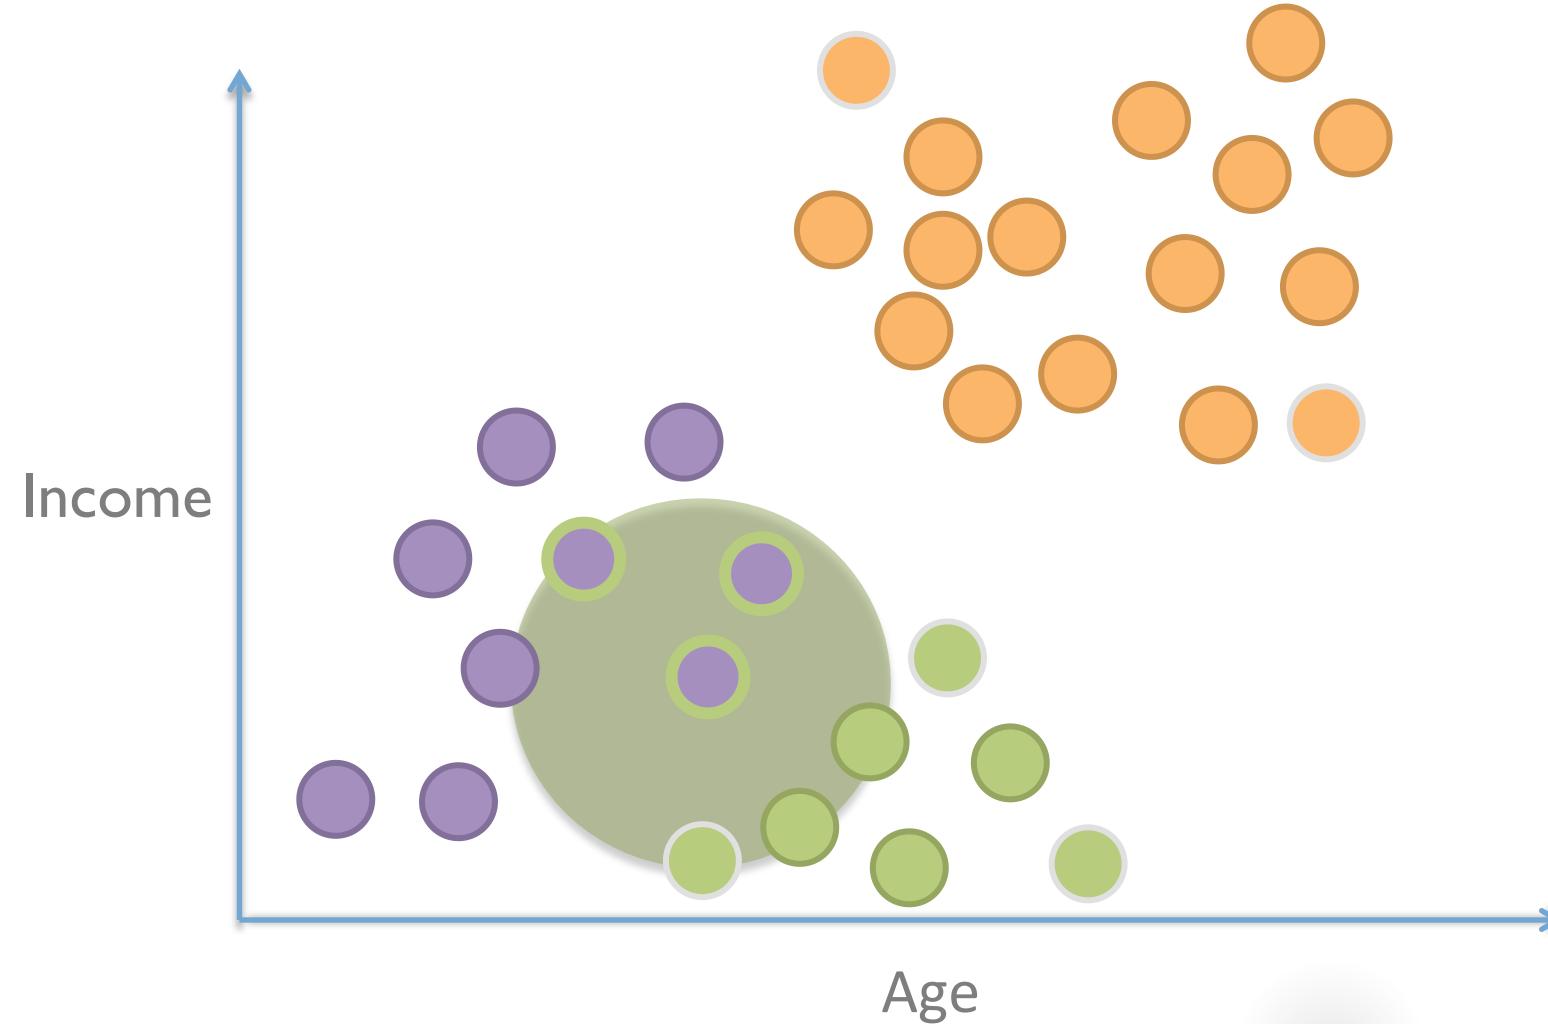
Keep adding neighbors within  $\text{epsilon}$  as chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

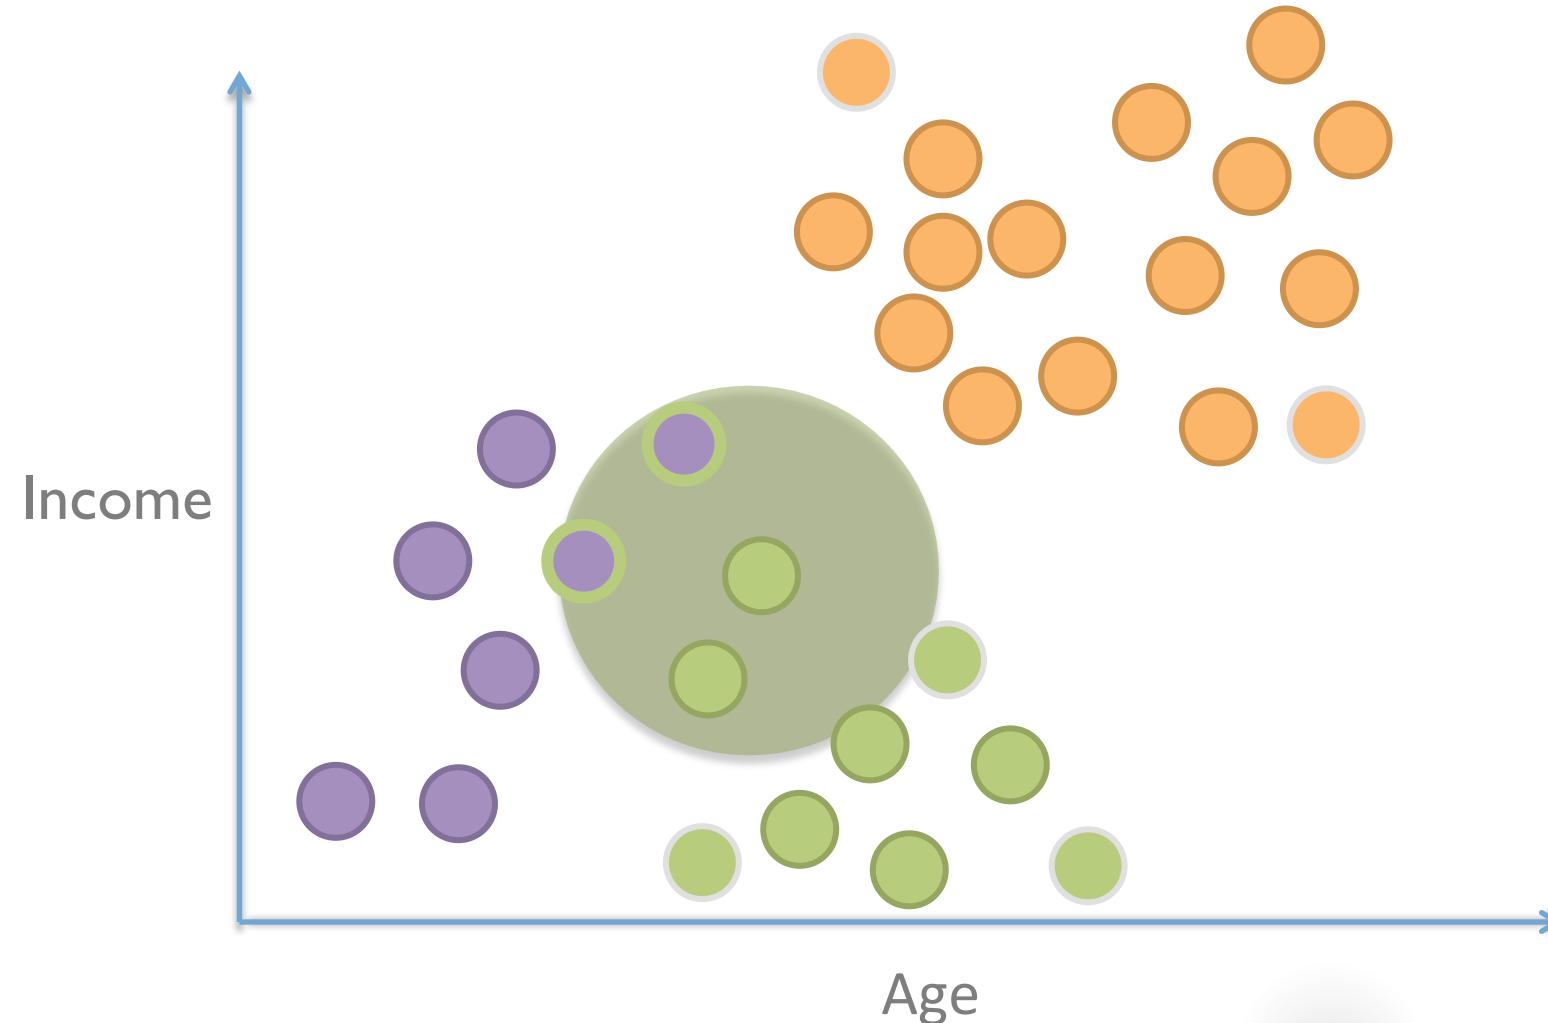
Keep adding neighbors within  $\text{epsilon}$  as chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

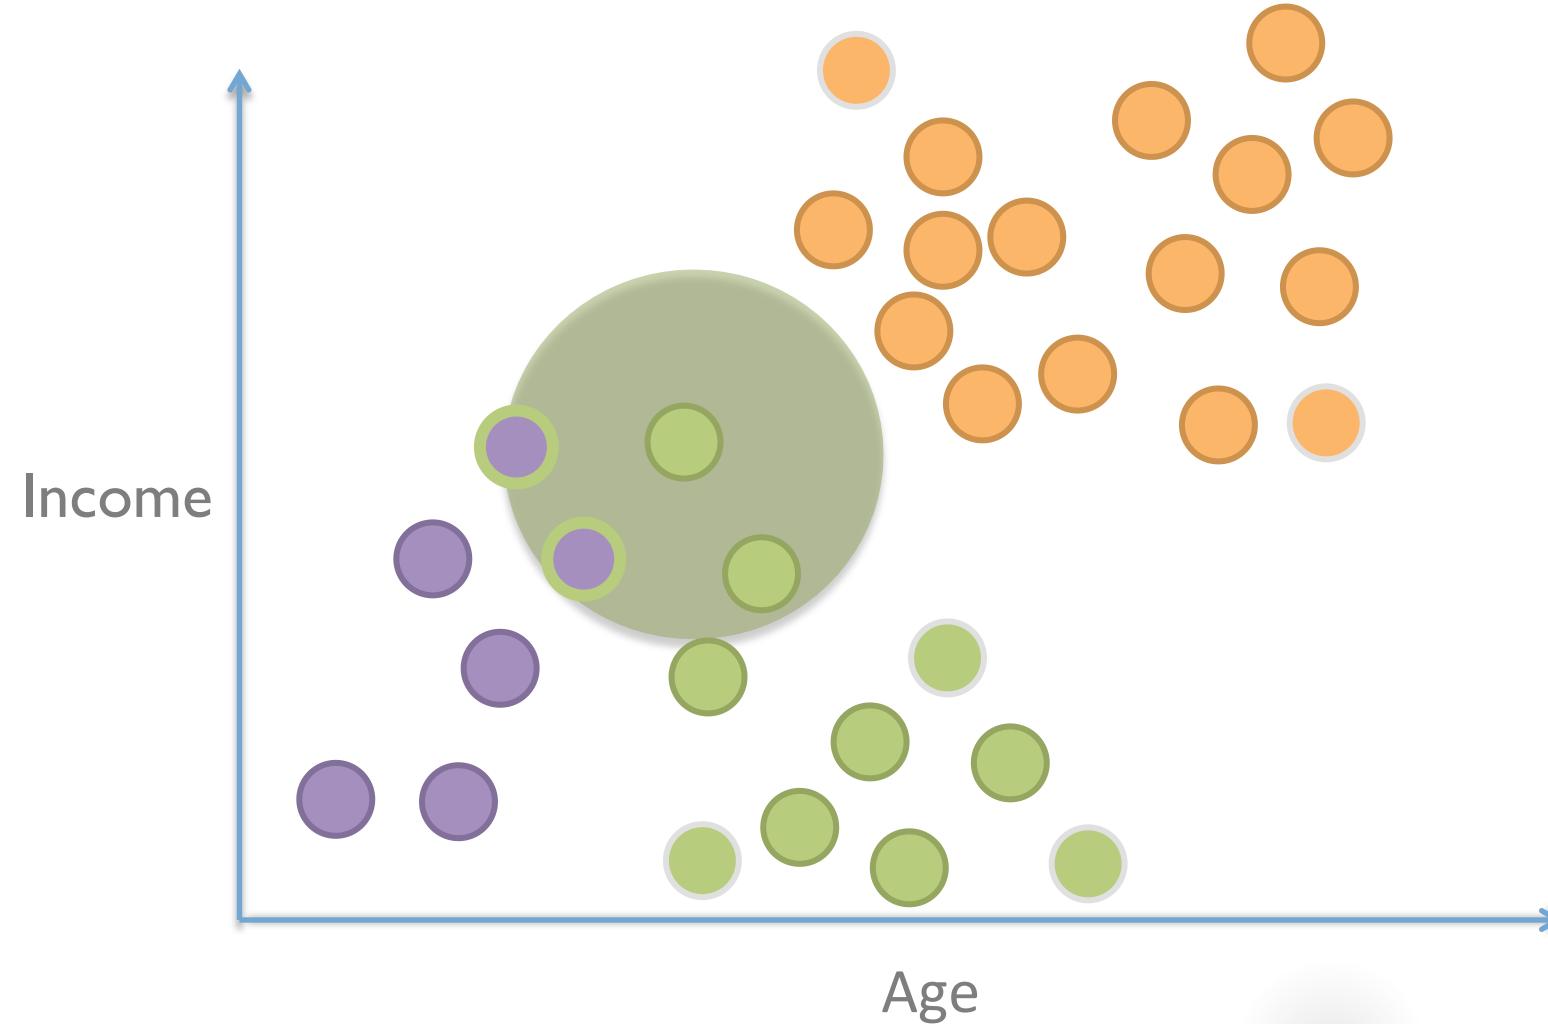
Keep adding neighbors within  $\text{epsilon}$  as chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

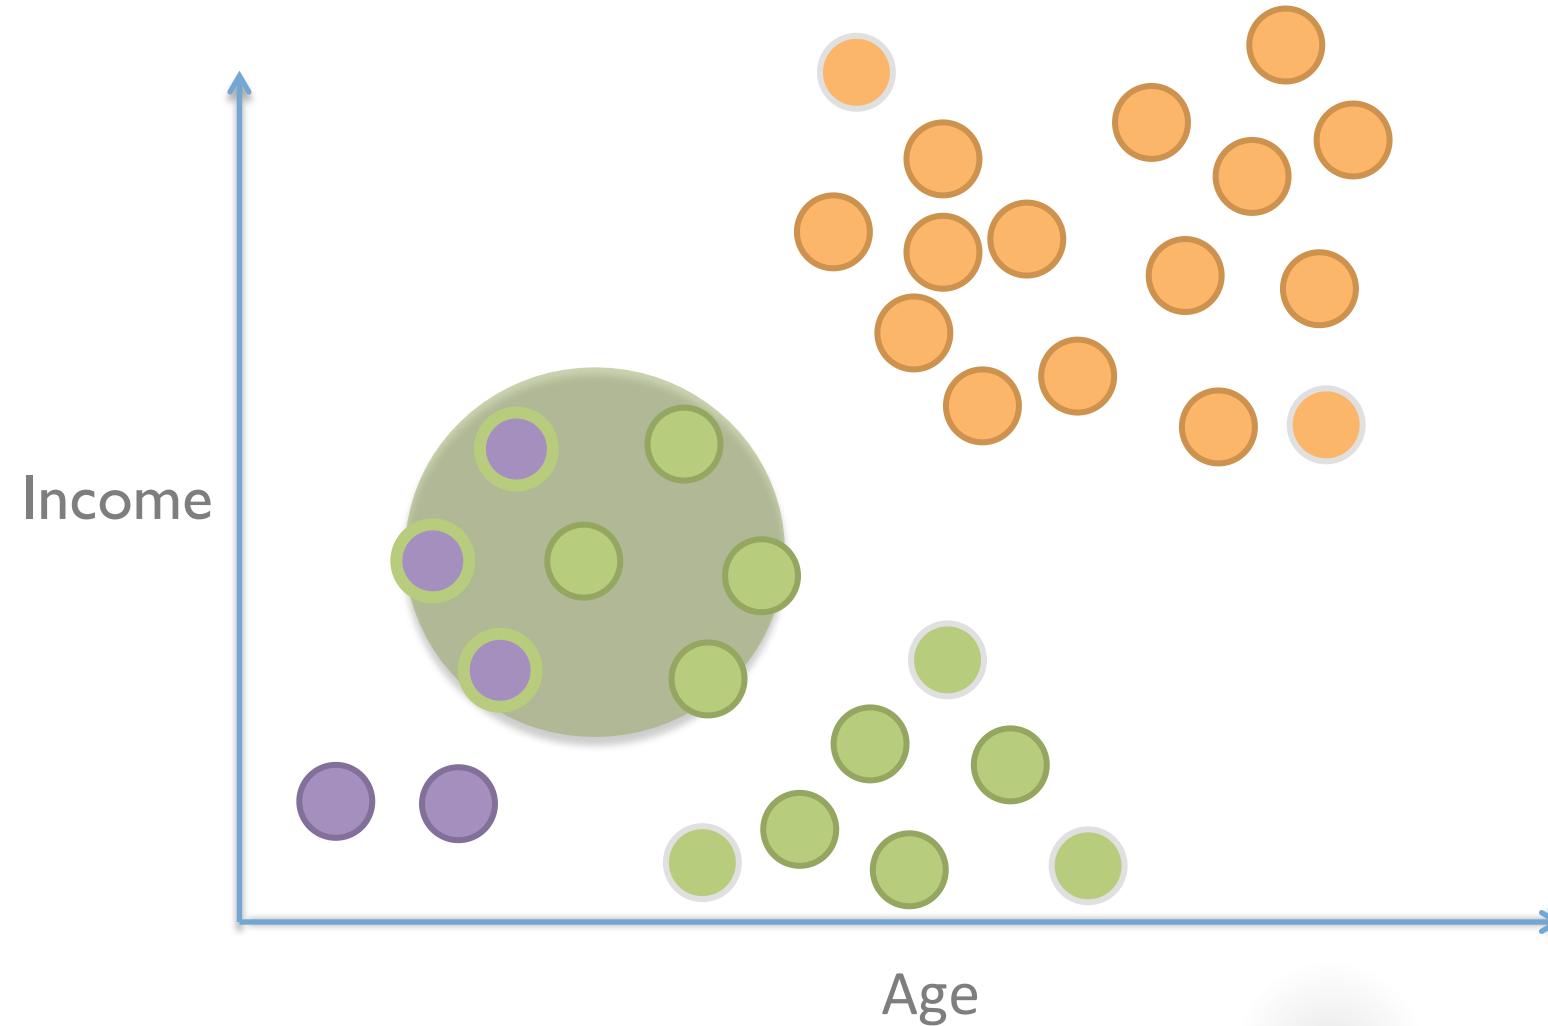
Keep adding neighbors within  $\text{epsilon}$  as chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

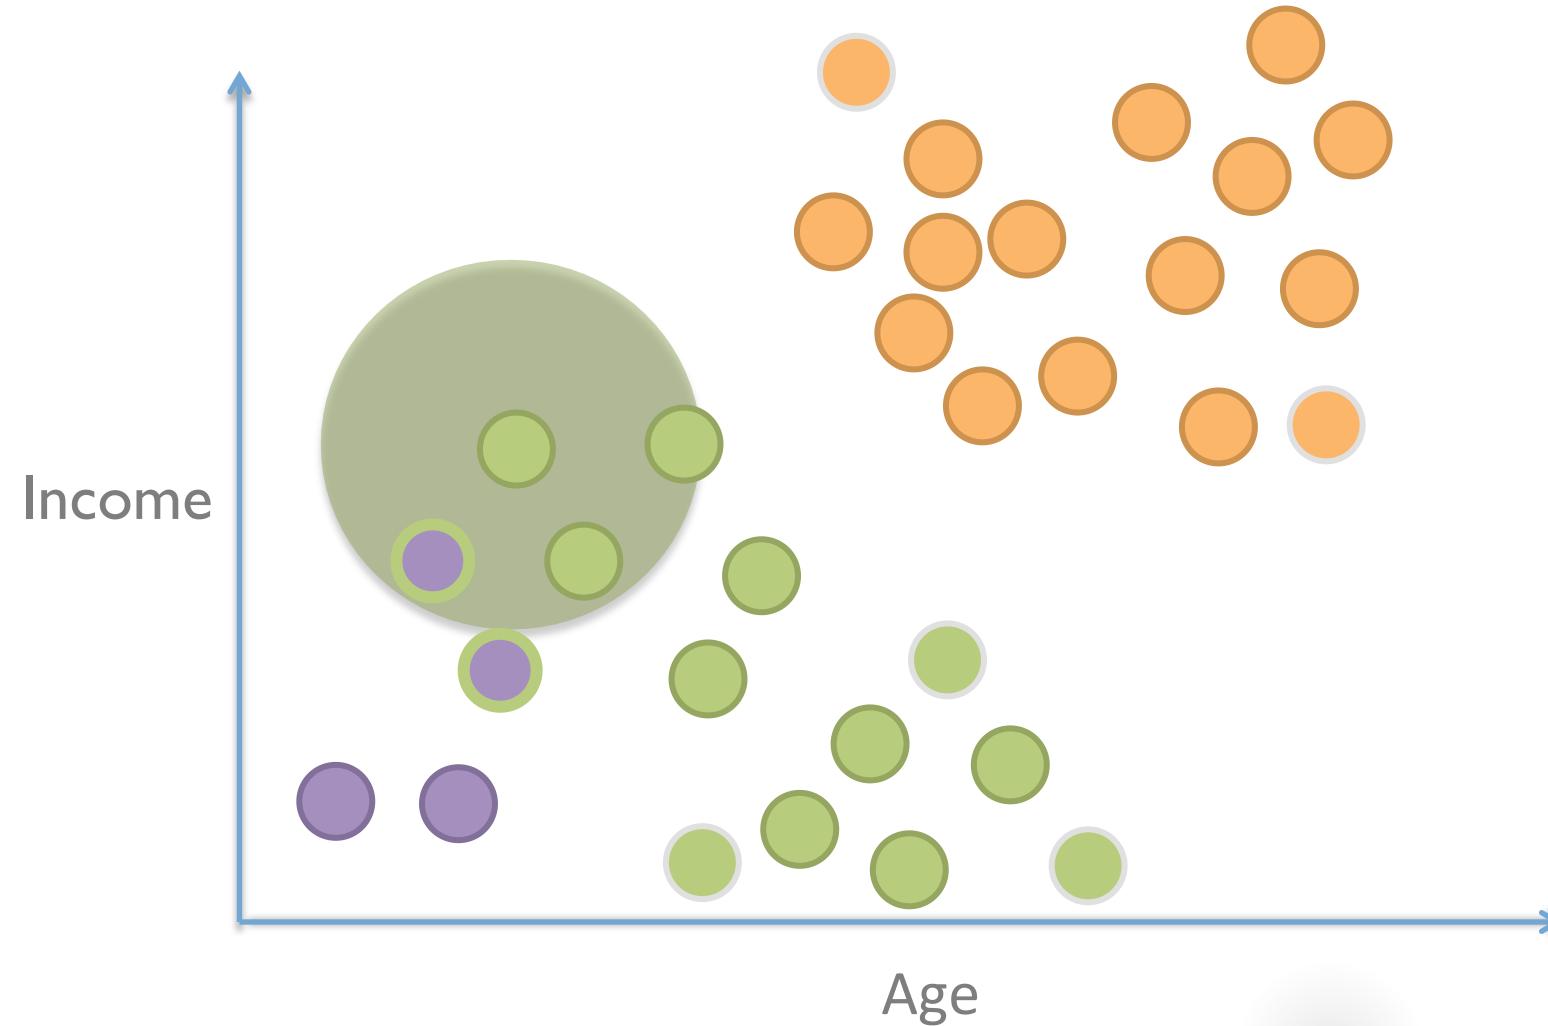
Keep adding neighbors within  $\text{epsilon}$  as chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

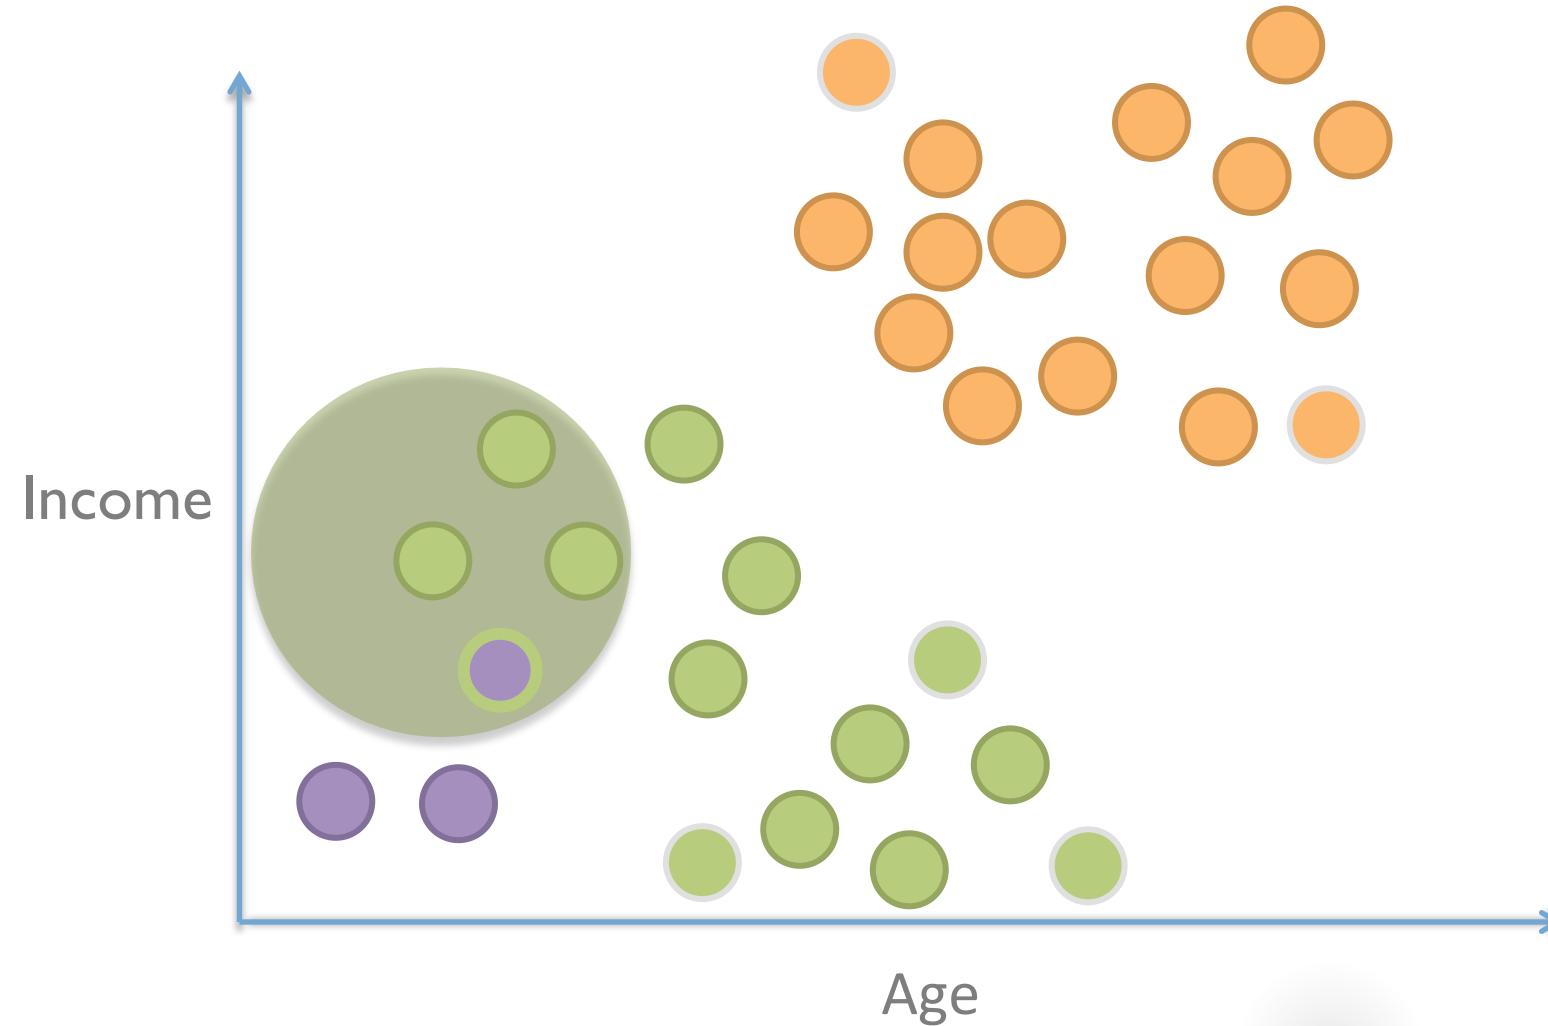
Keep adding neighbors within  $\text{epsilon}$  as chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

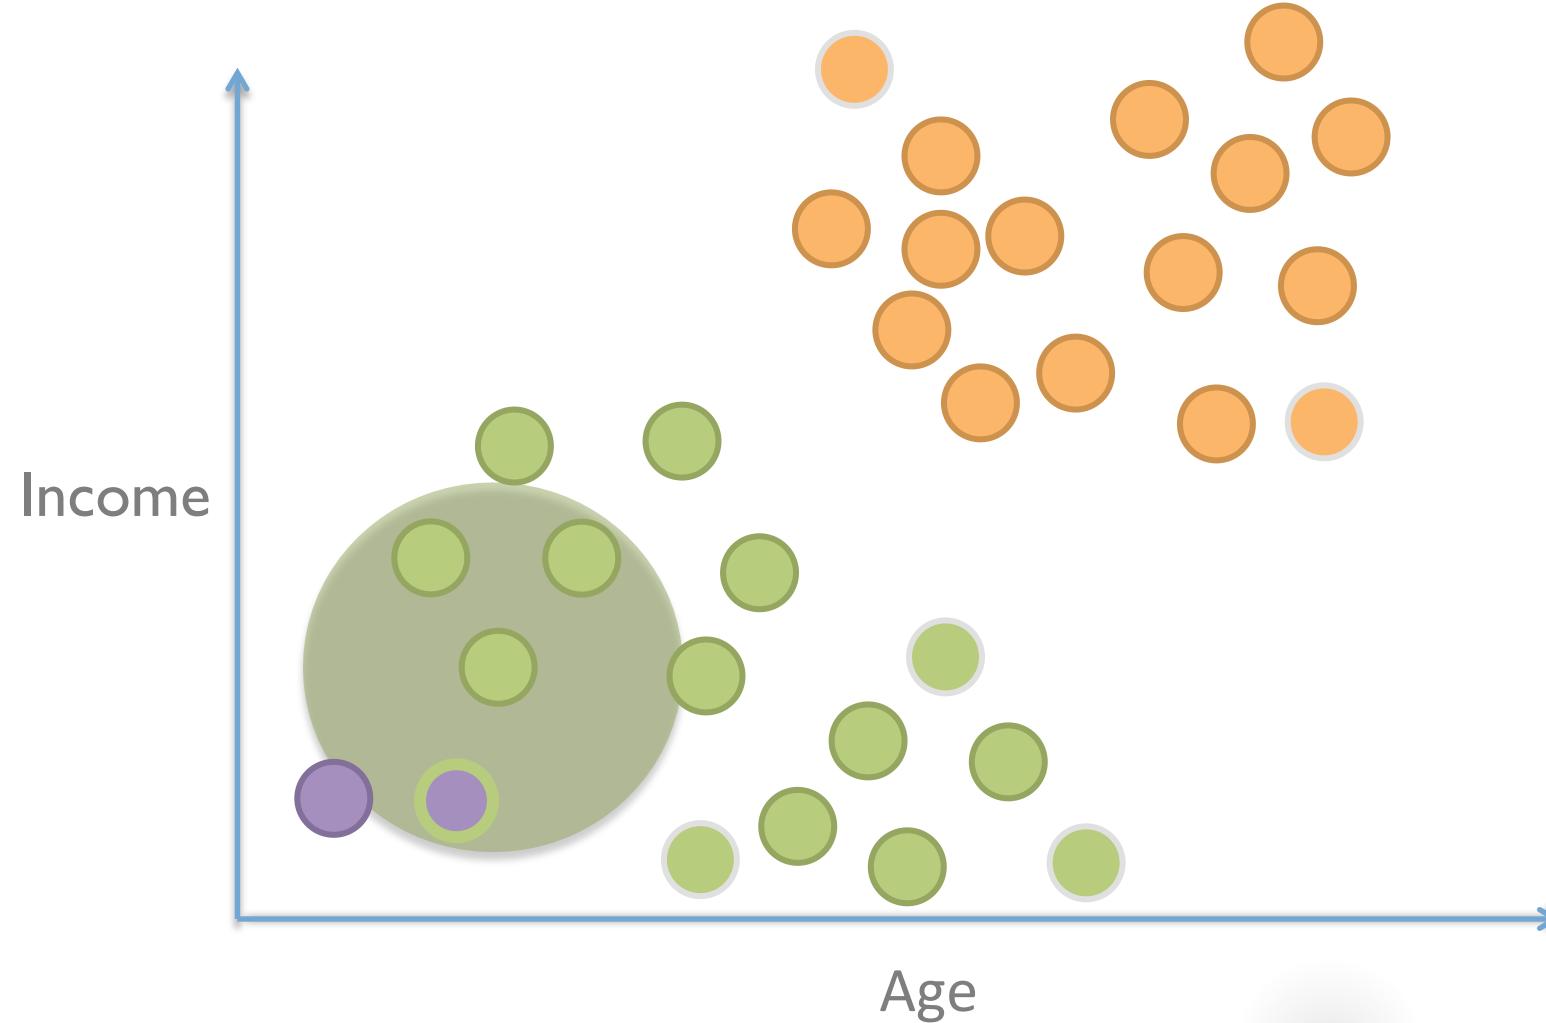
Keep adding neighbors within  $\text{epsilon}$  as chain reaction



epsilon = 1.75  
n\_clu = 3

## DBSCAN

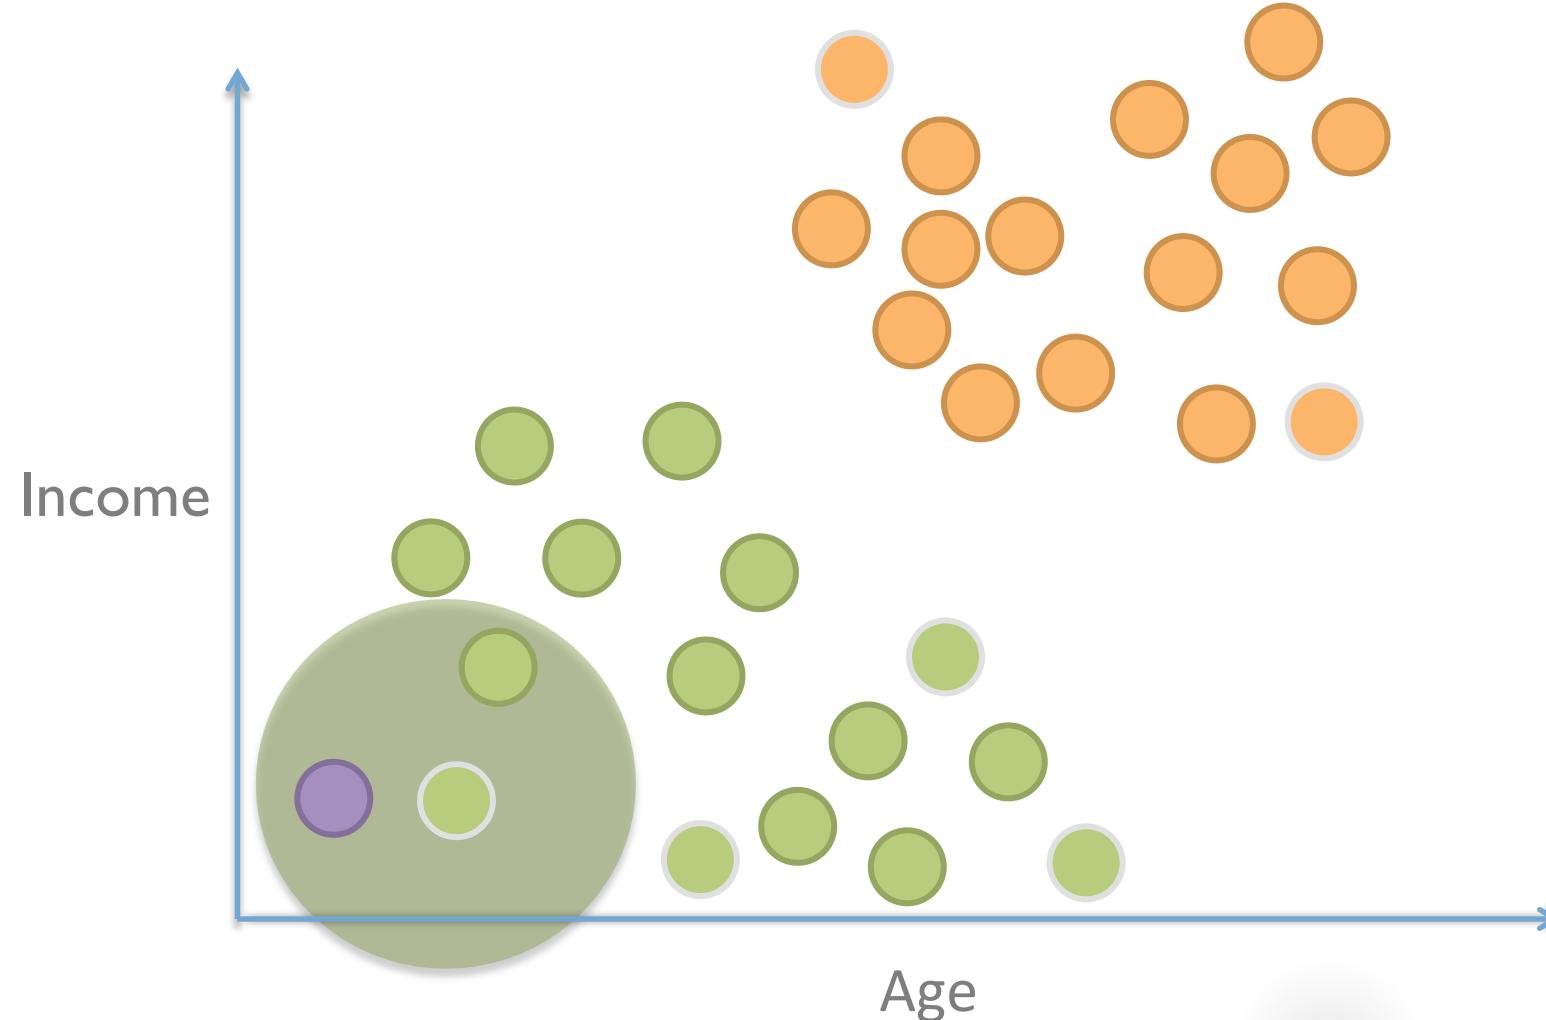
Keep adding neighbors within *epsilon* as chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

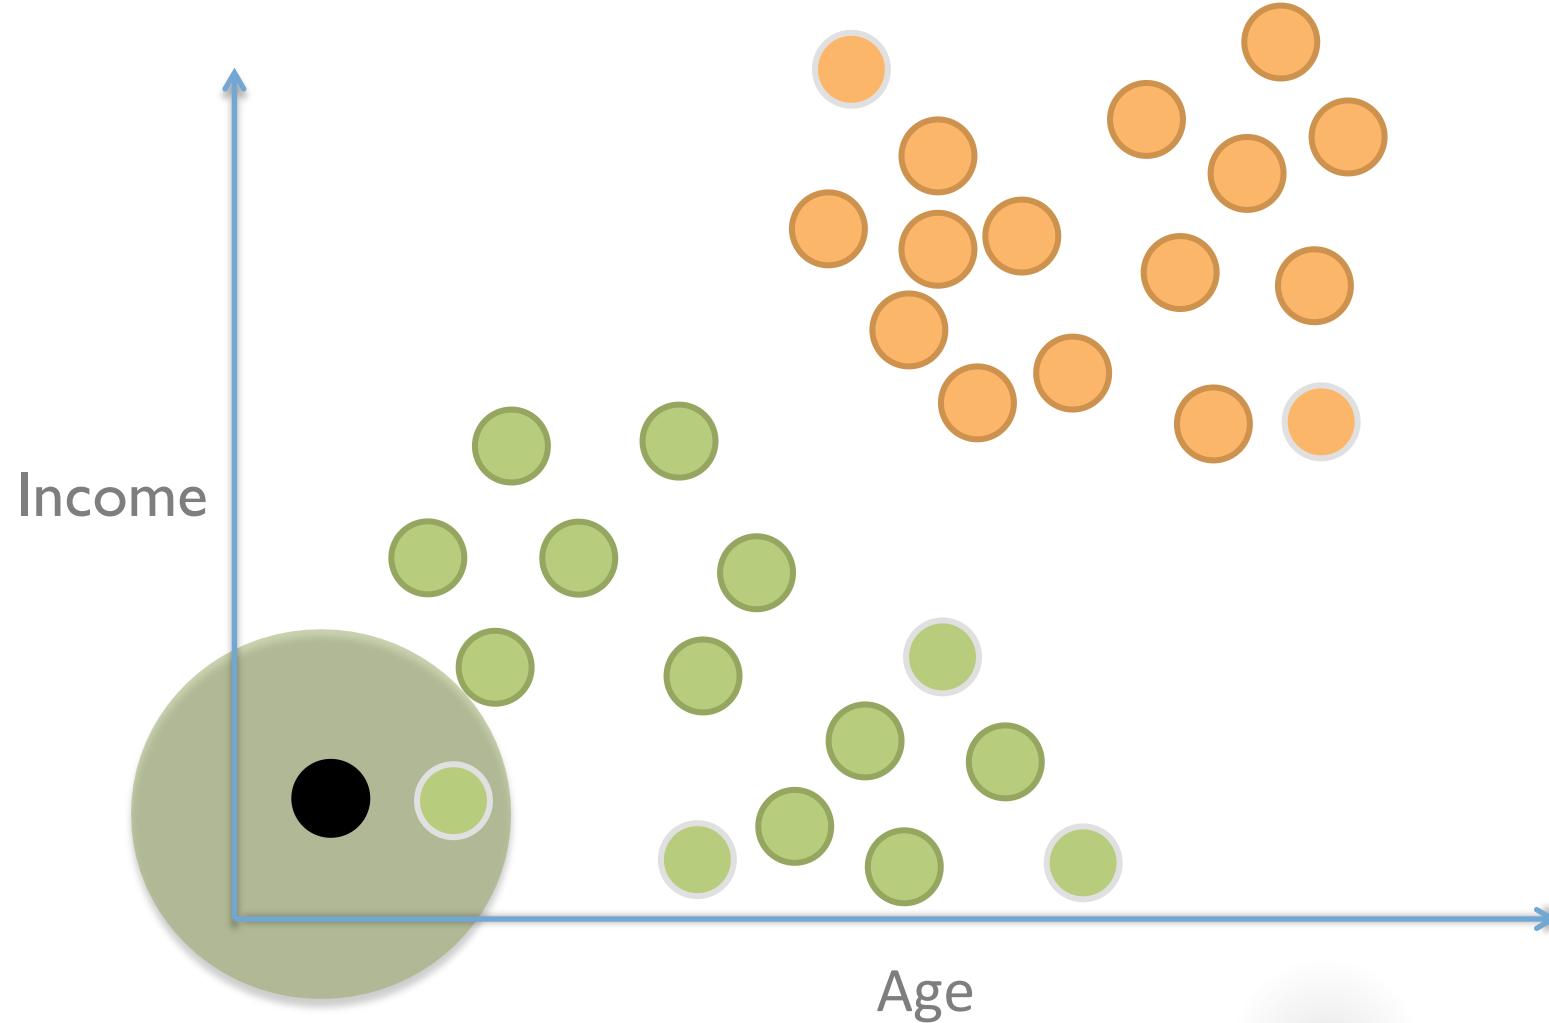
Keep adding neighbors within  $\text{epsilon}$  as chain reaction



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

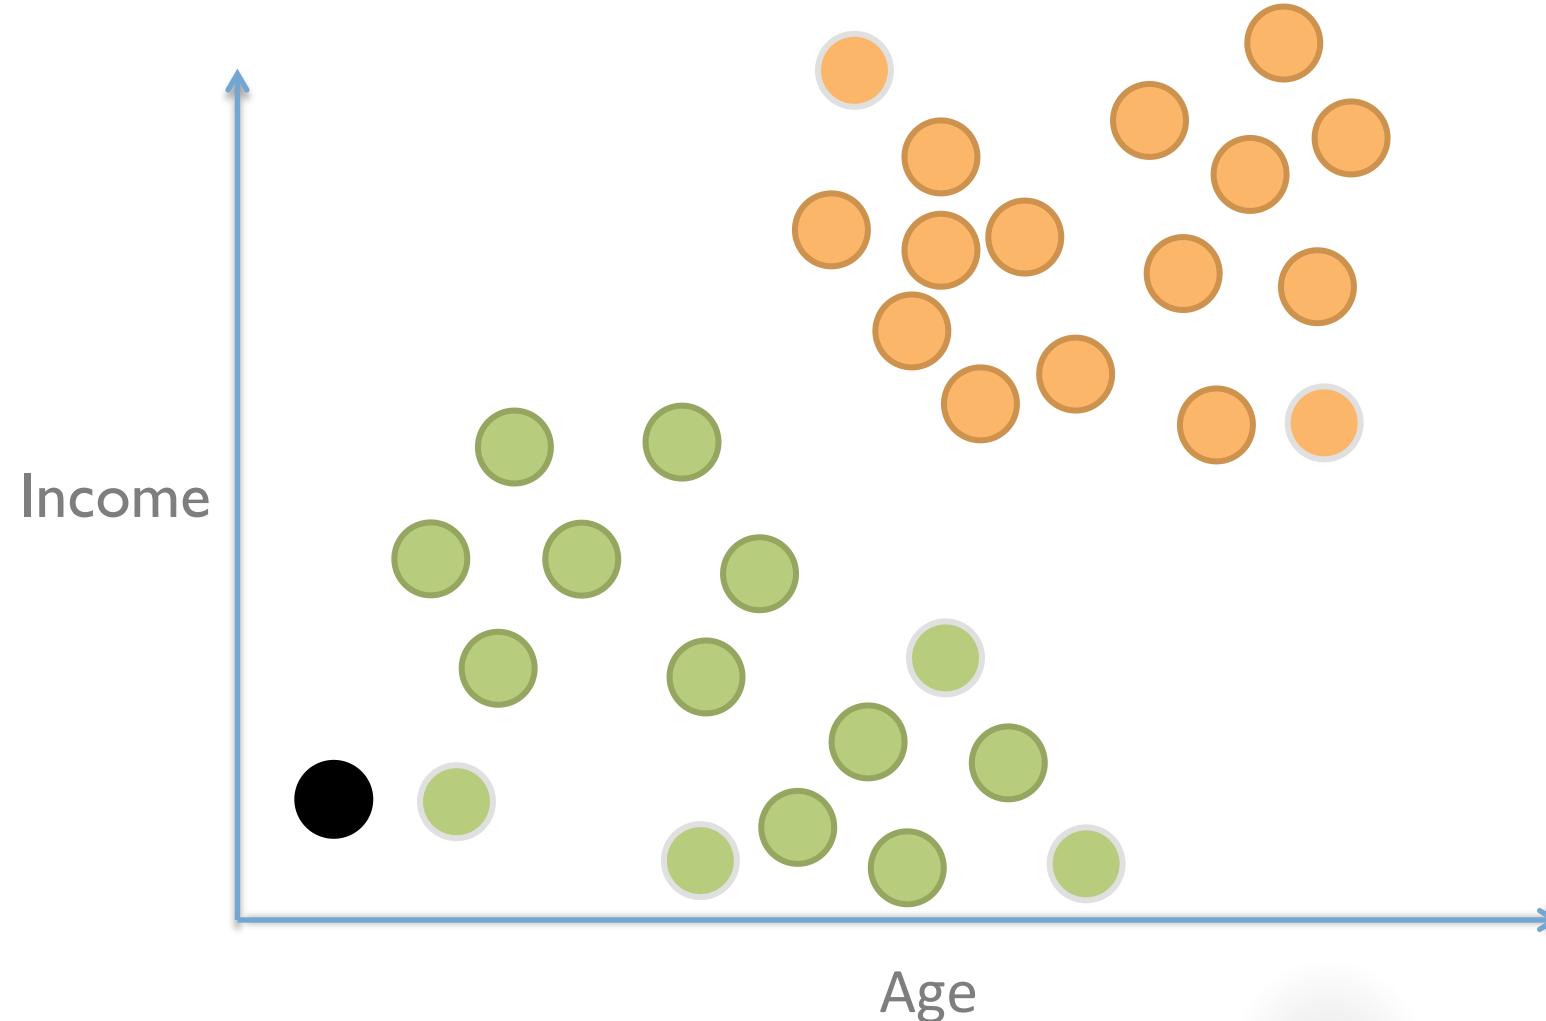
If there is a point that does not have  $n_{\text{clu}}$  neighbors and is not reached from a core point: It is a **noise** point.



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

## DBSCAN

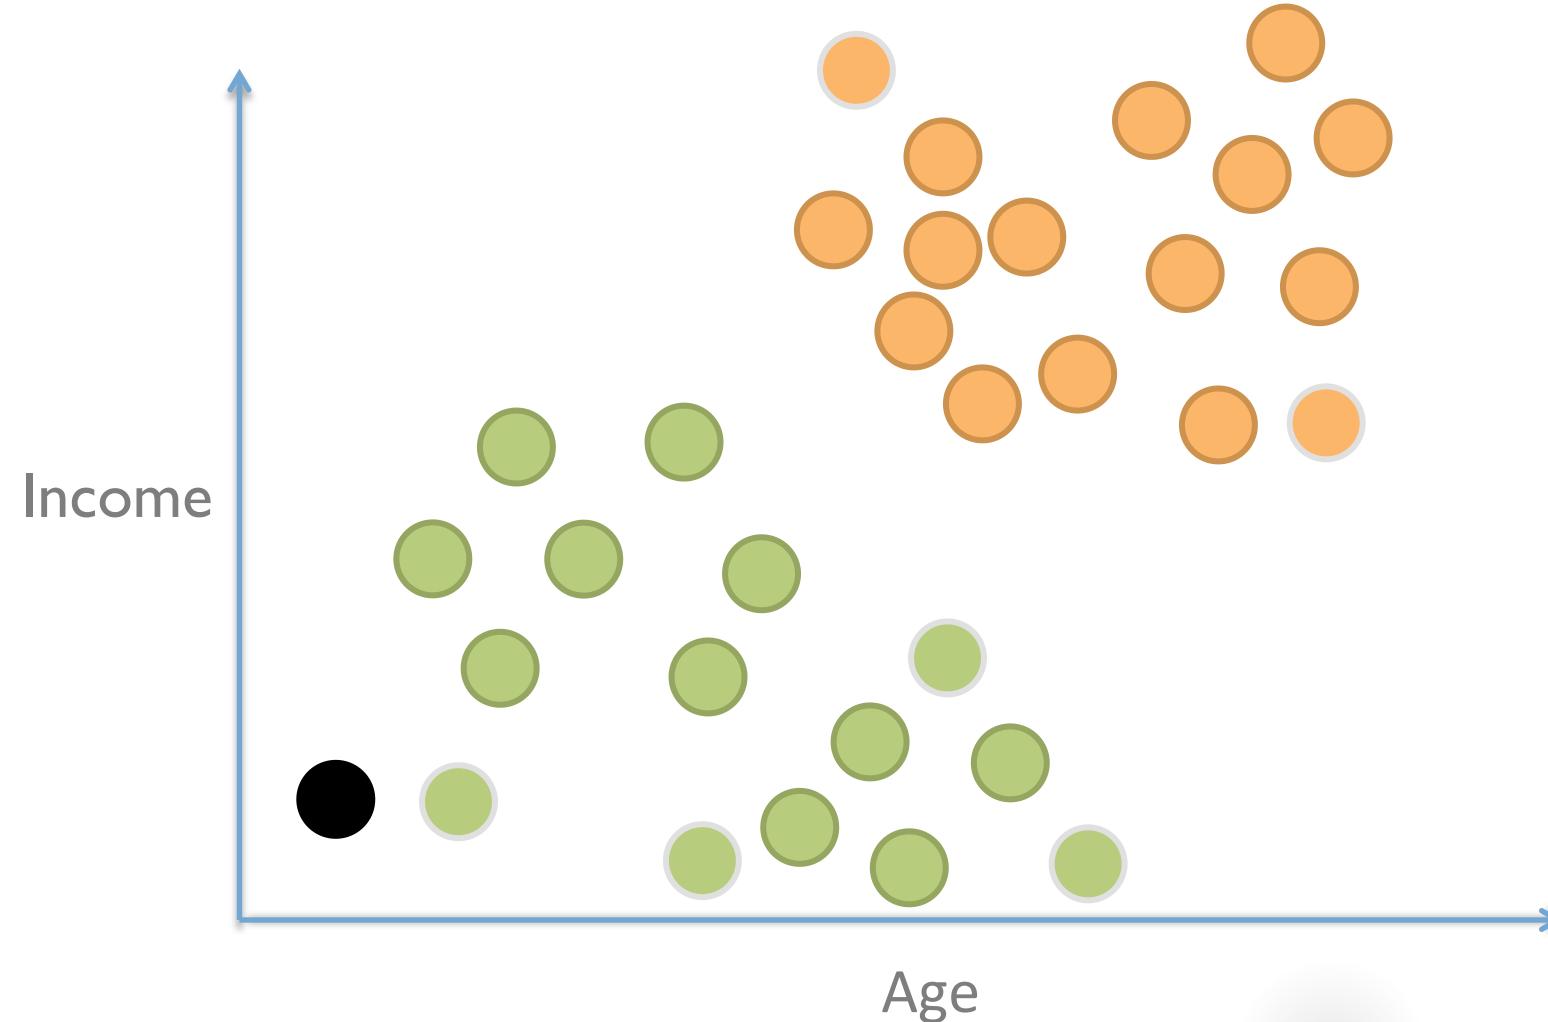
Ta--daa! Notice the core points, density-reachable points (at the borders) and the **noise** point.



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 3$

# DBSCAN

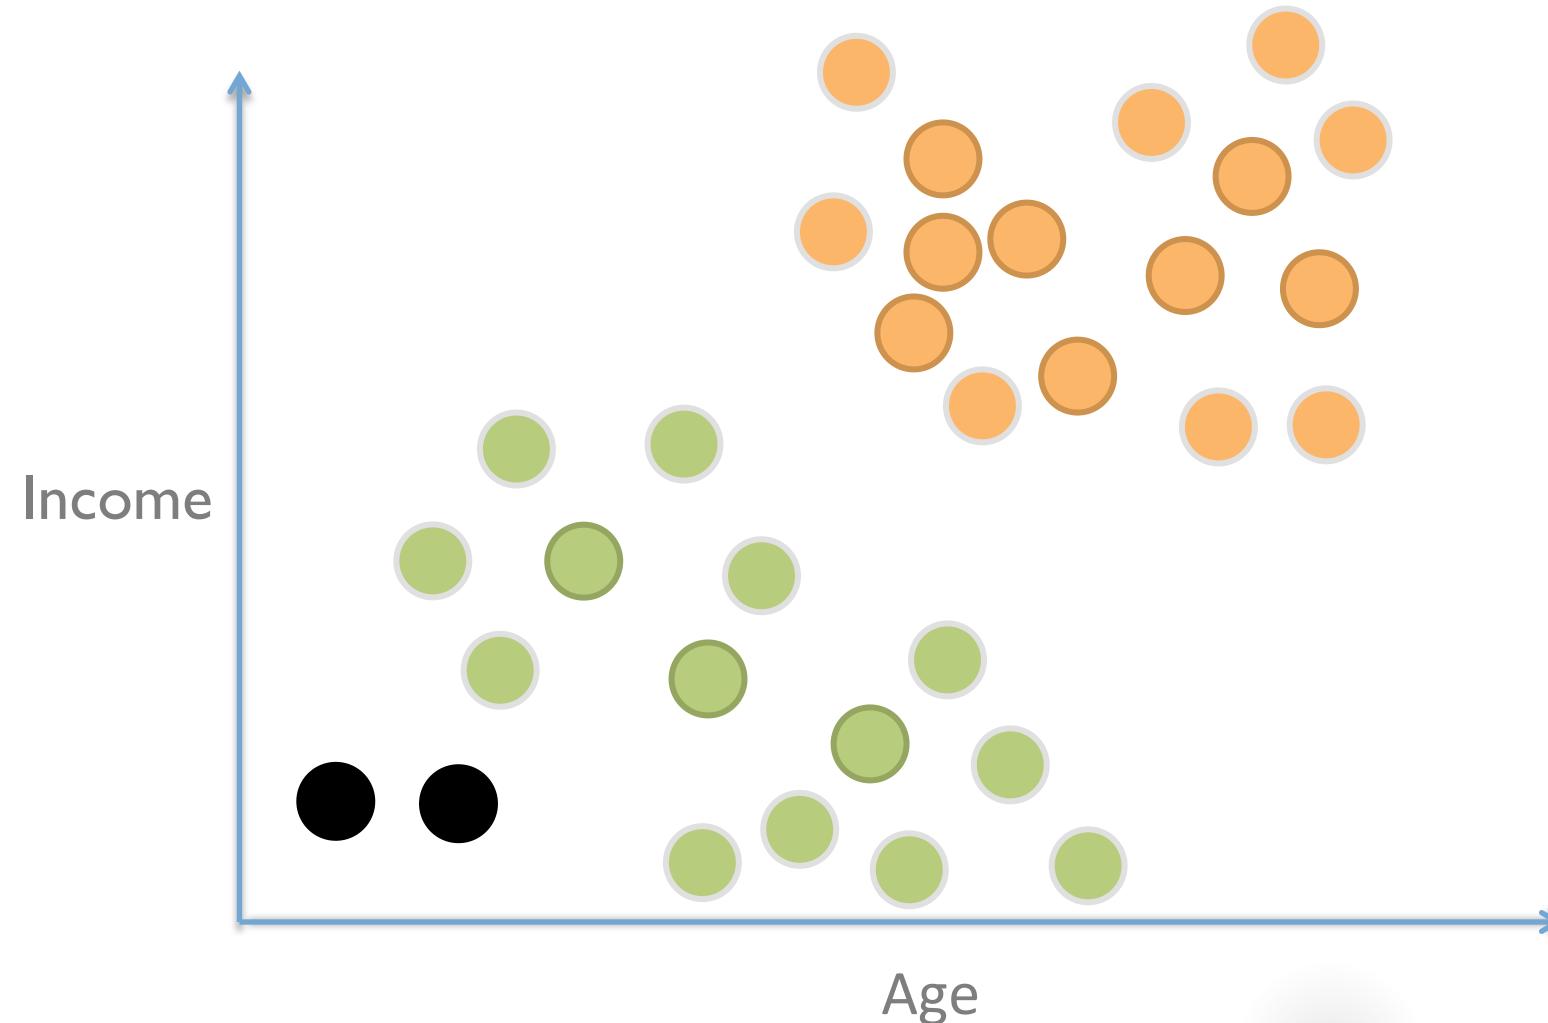
The effect of the parameters



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 5$

# DBSCAN

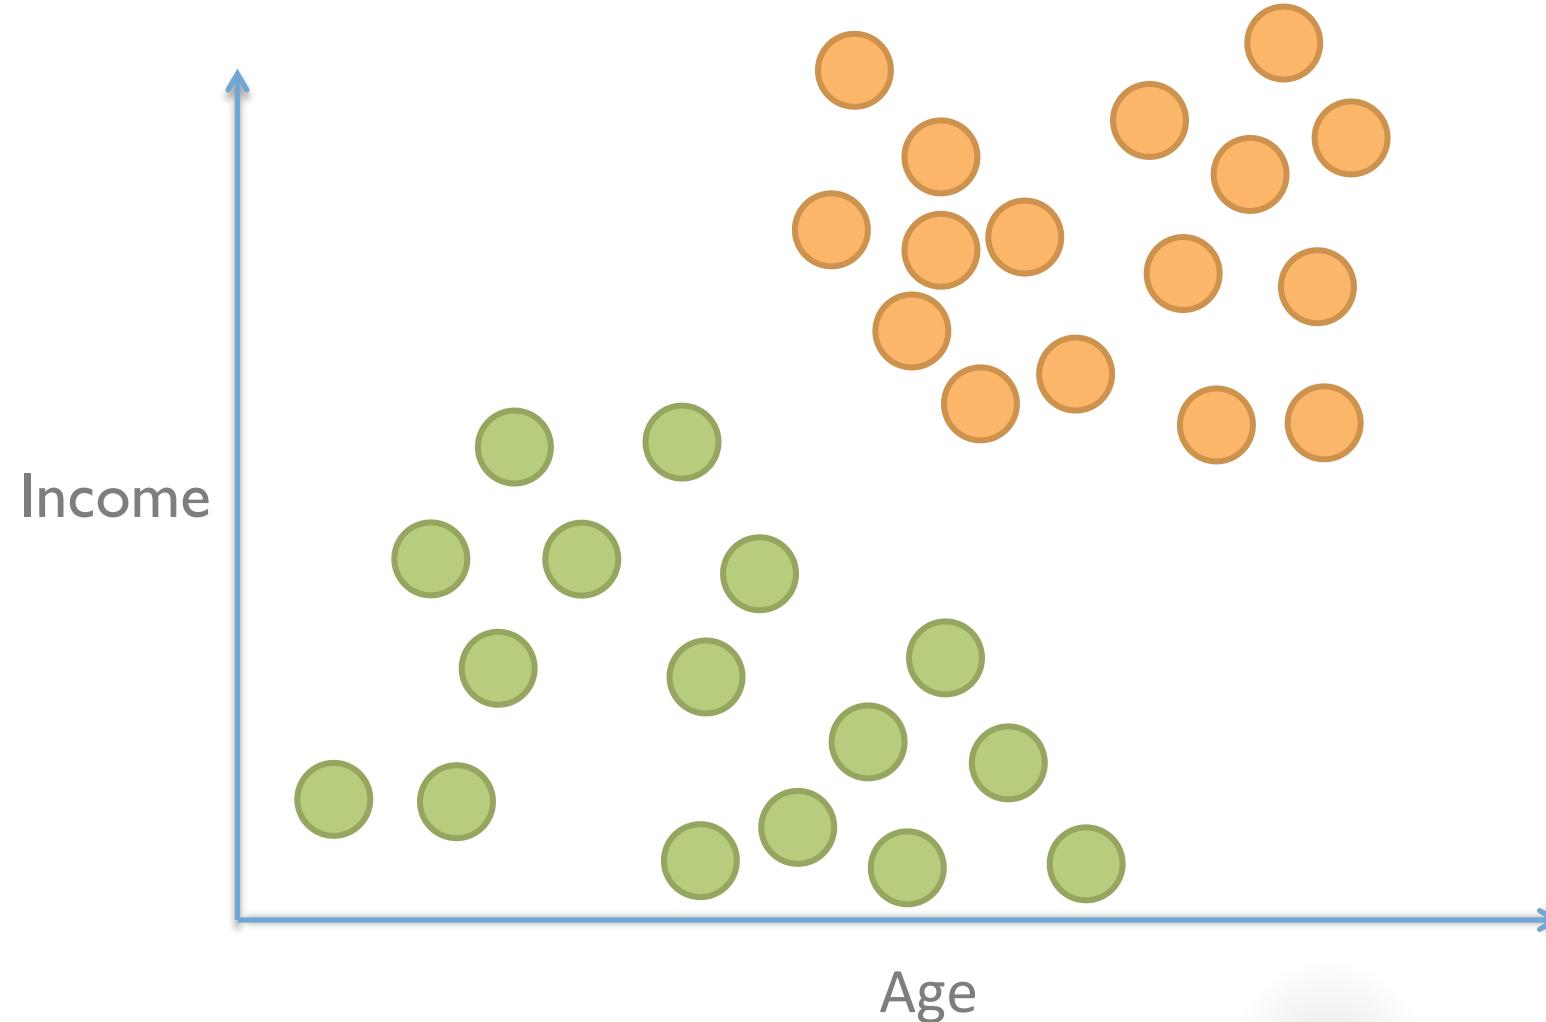
The effect of the parameters



$\text{epsilon} = 1.75$   
 $n_{\text{clu}} = 1$

# DBSCAN

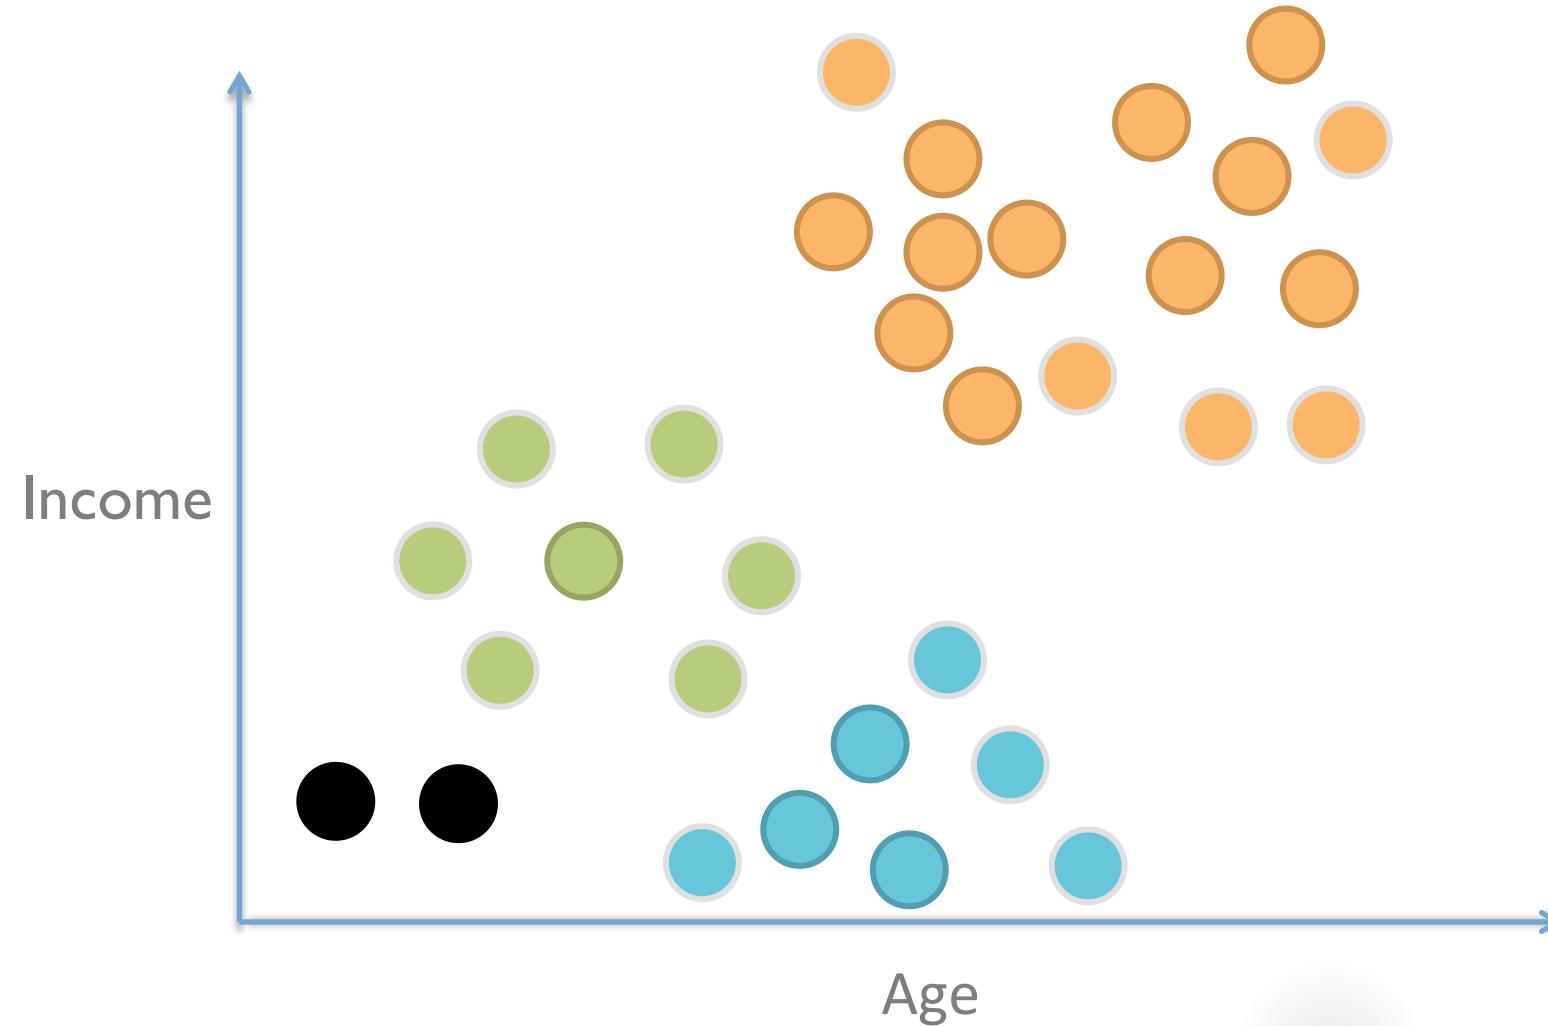
The effect of the parameters



epsilon = 1.25  
n\_clu = 3

# DBSCAN

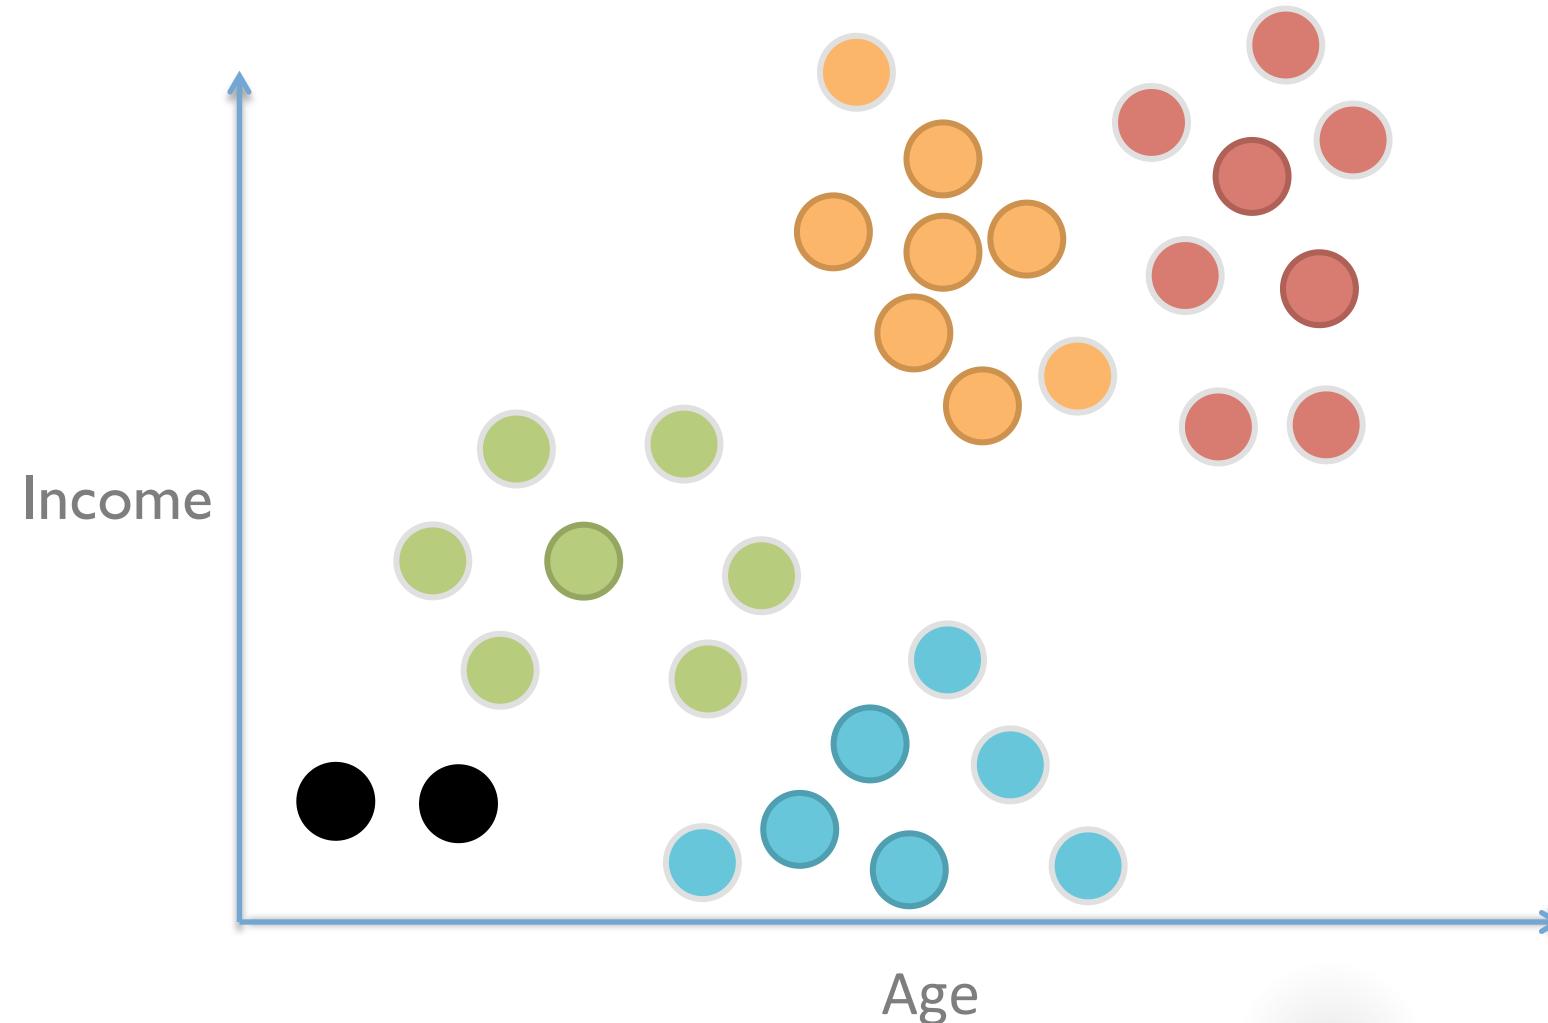
The effect of the parameters



epsilon = 1.00  
n\_clu = 3

# DBSCAN

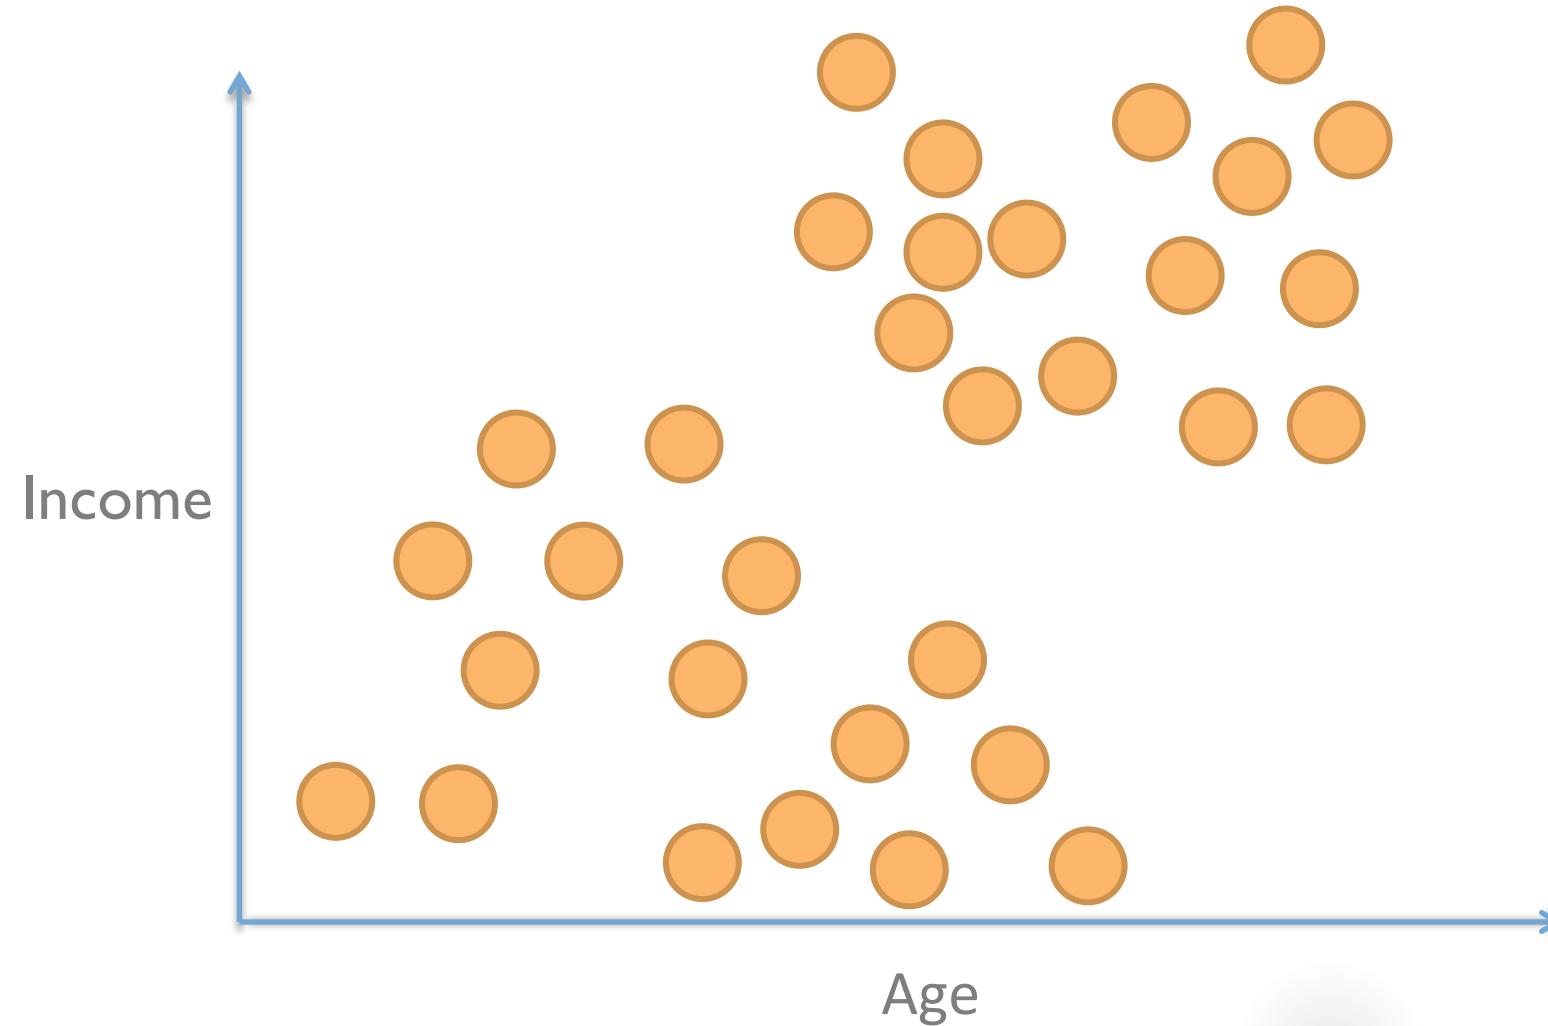
The effect of the parameters



epsilon = 3.50  
n\_clu = 3

# DBSCAN

The effect of the parameters



```
from sklearn.cluster import DBSCAN
```

```
DBSCAN( eps = 0.5, min_samples = 5, metric = "euclidean" )
```

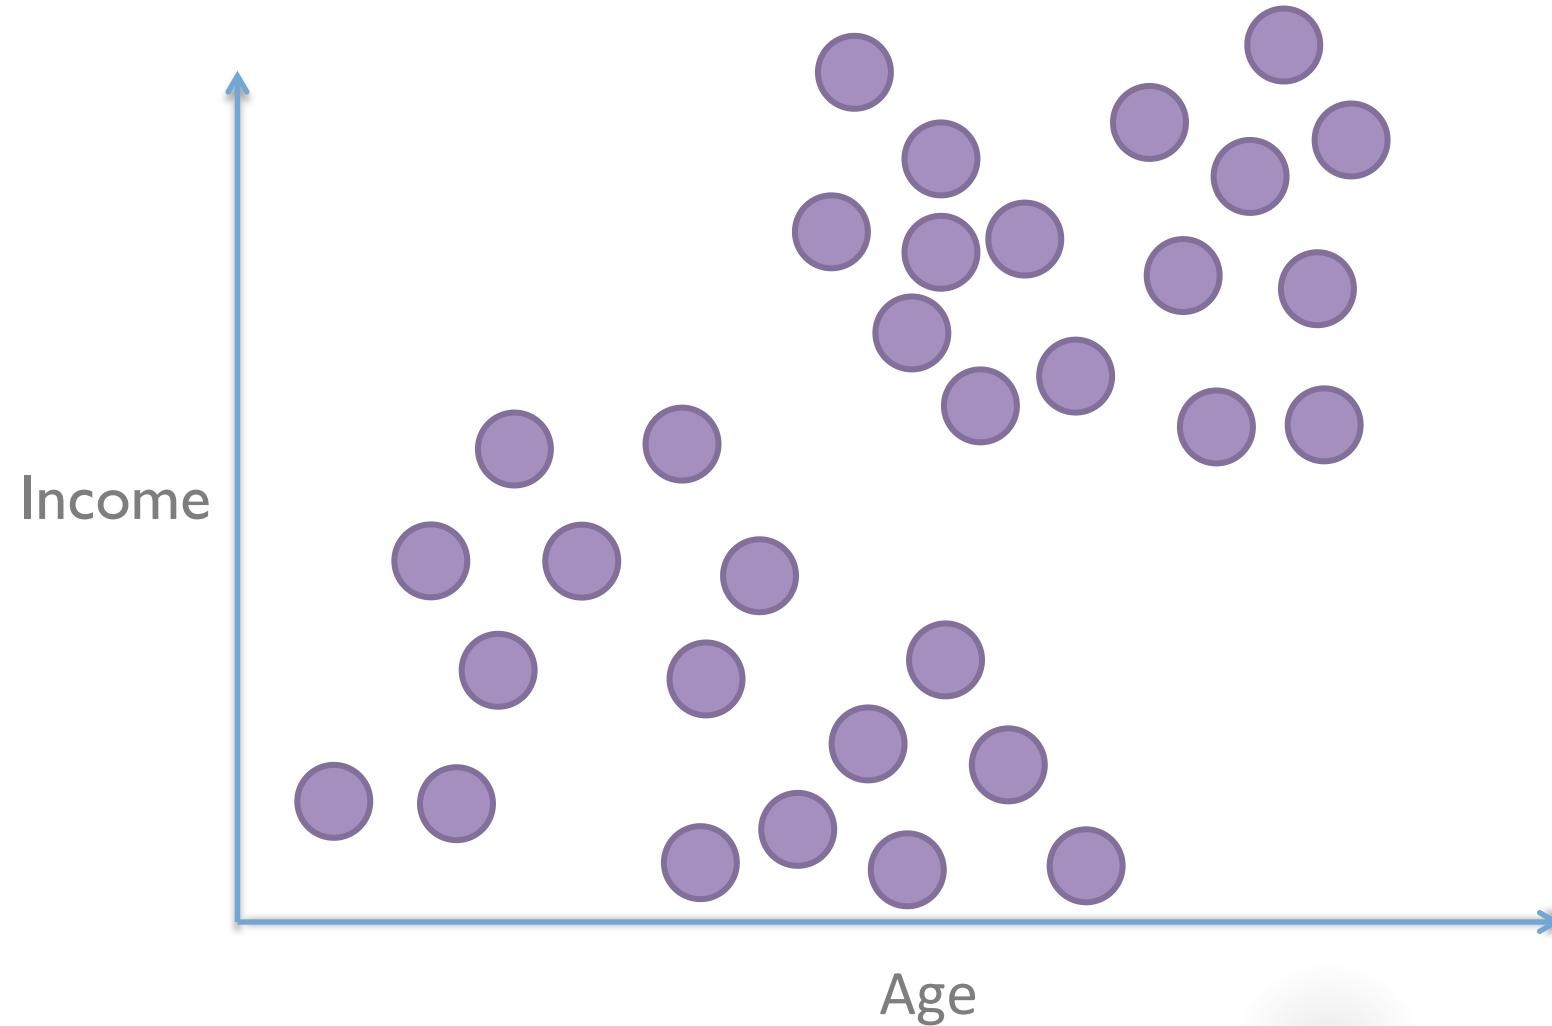




# Mean Shift

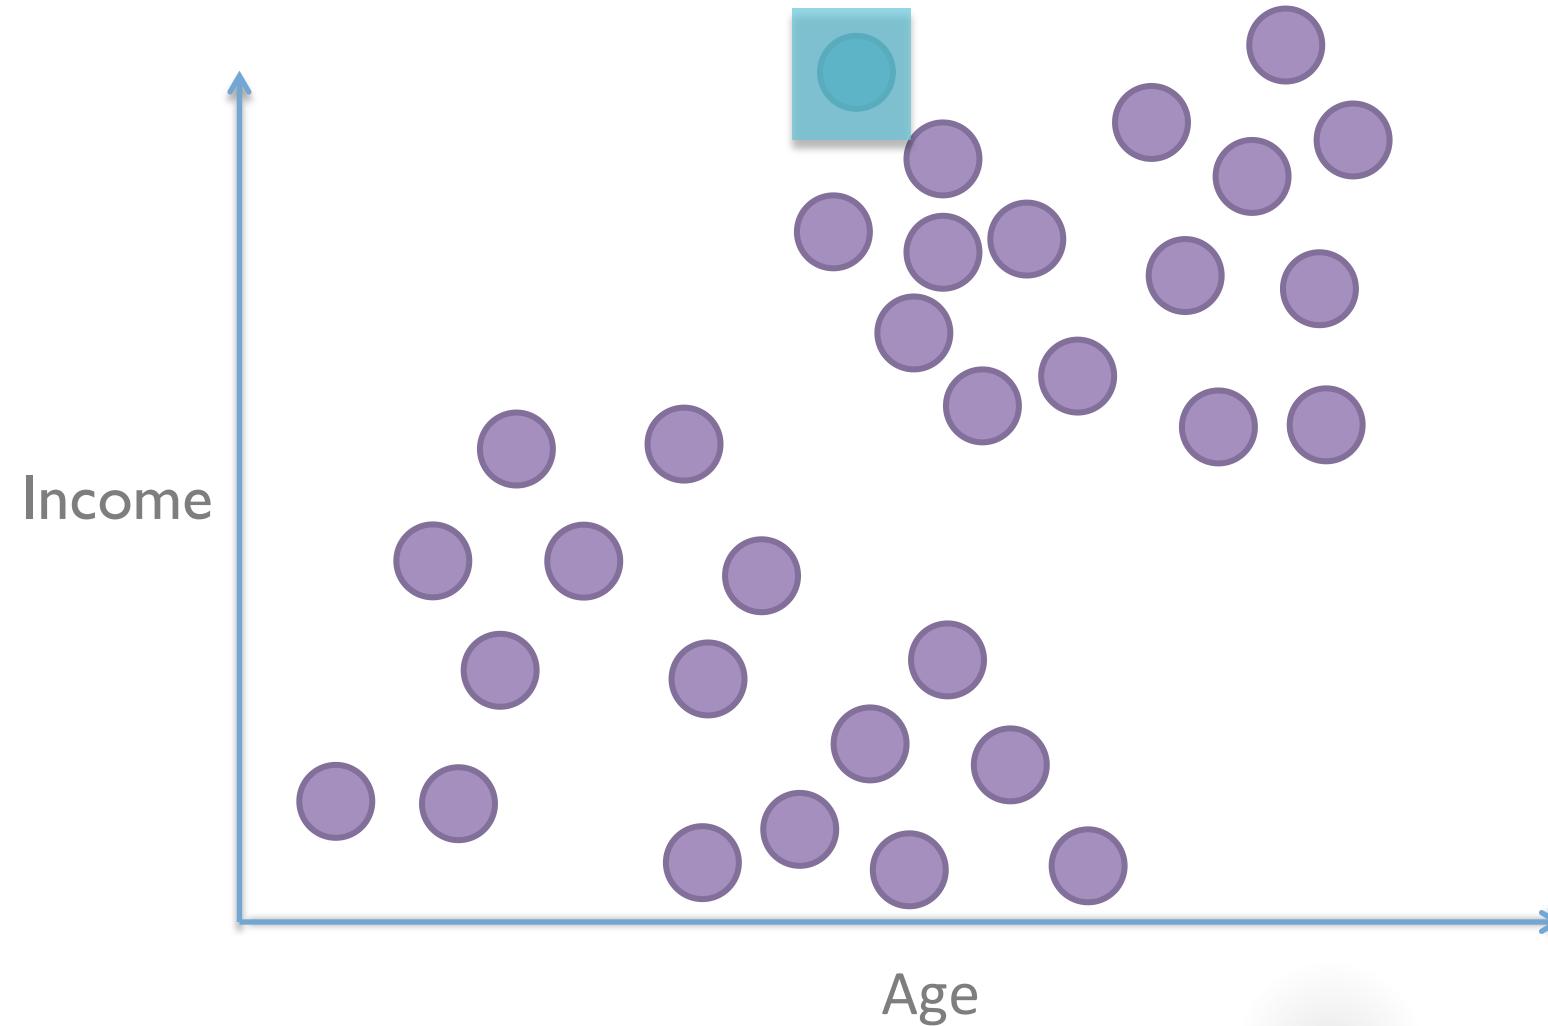
---

# Mean Shift



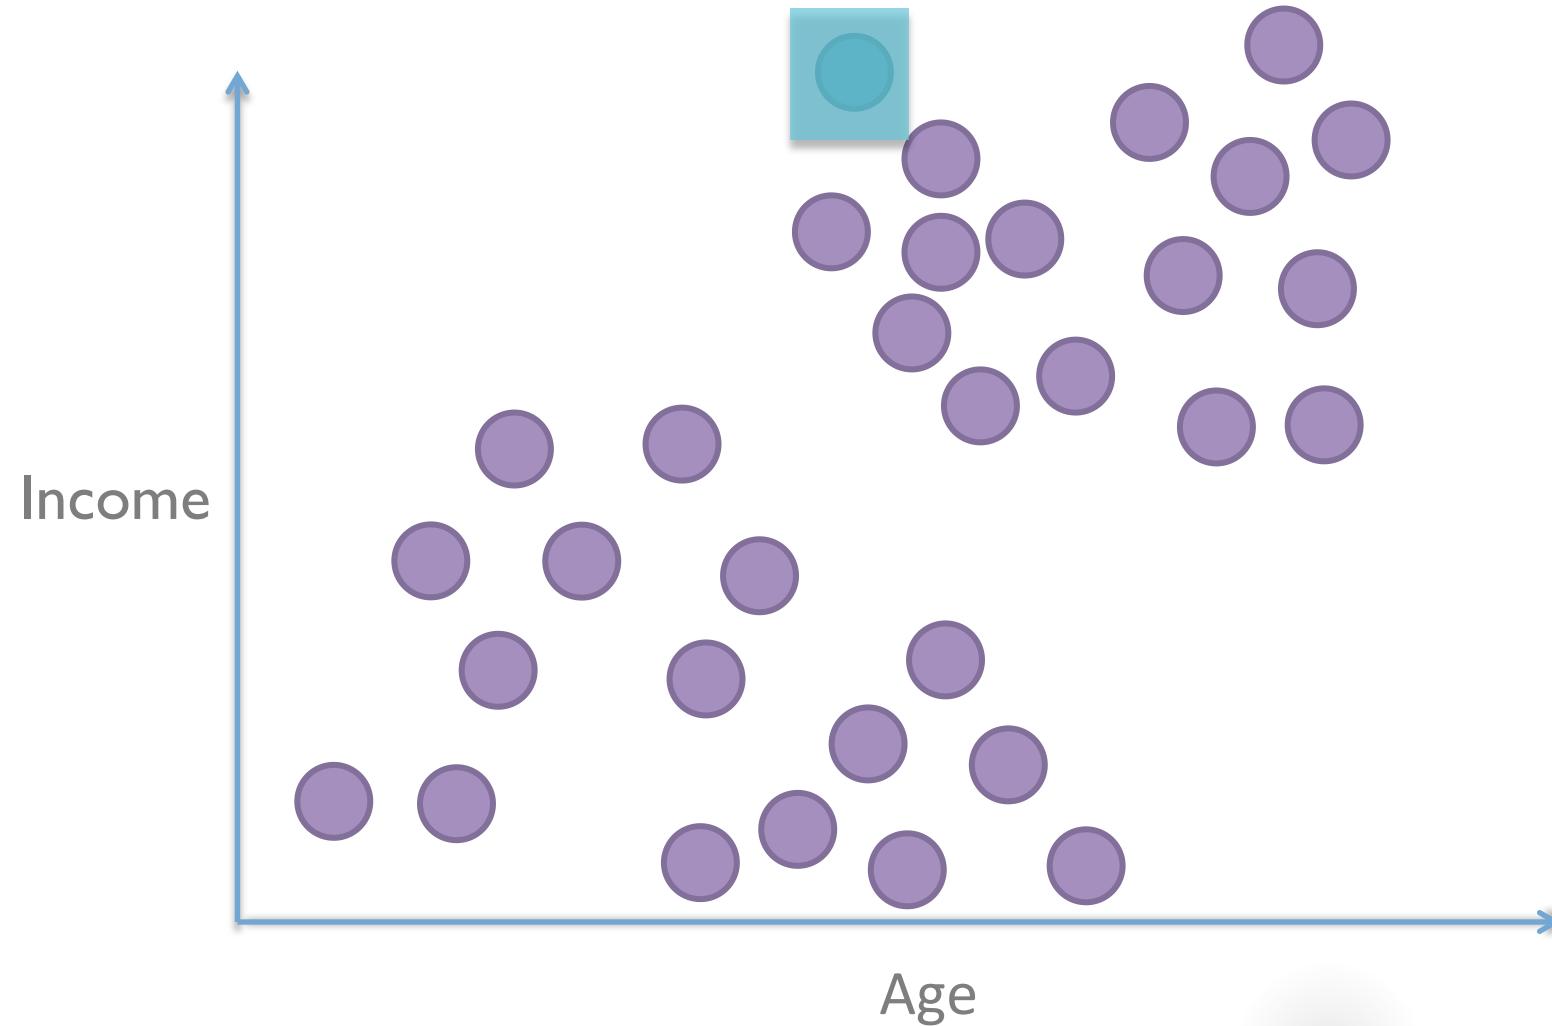
# Mean Shift

Start with a centroid at a point



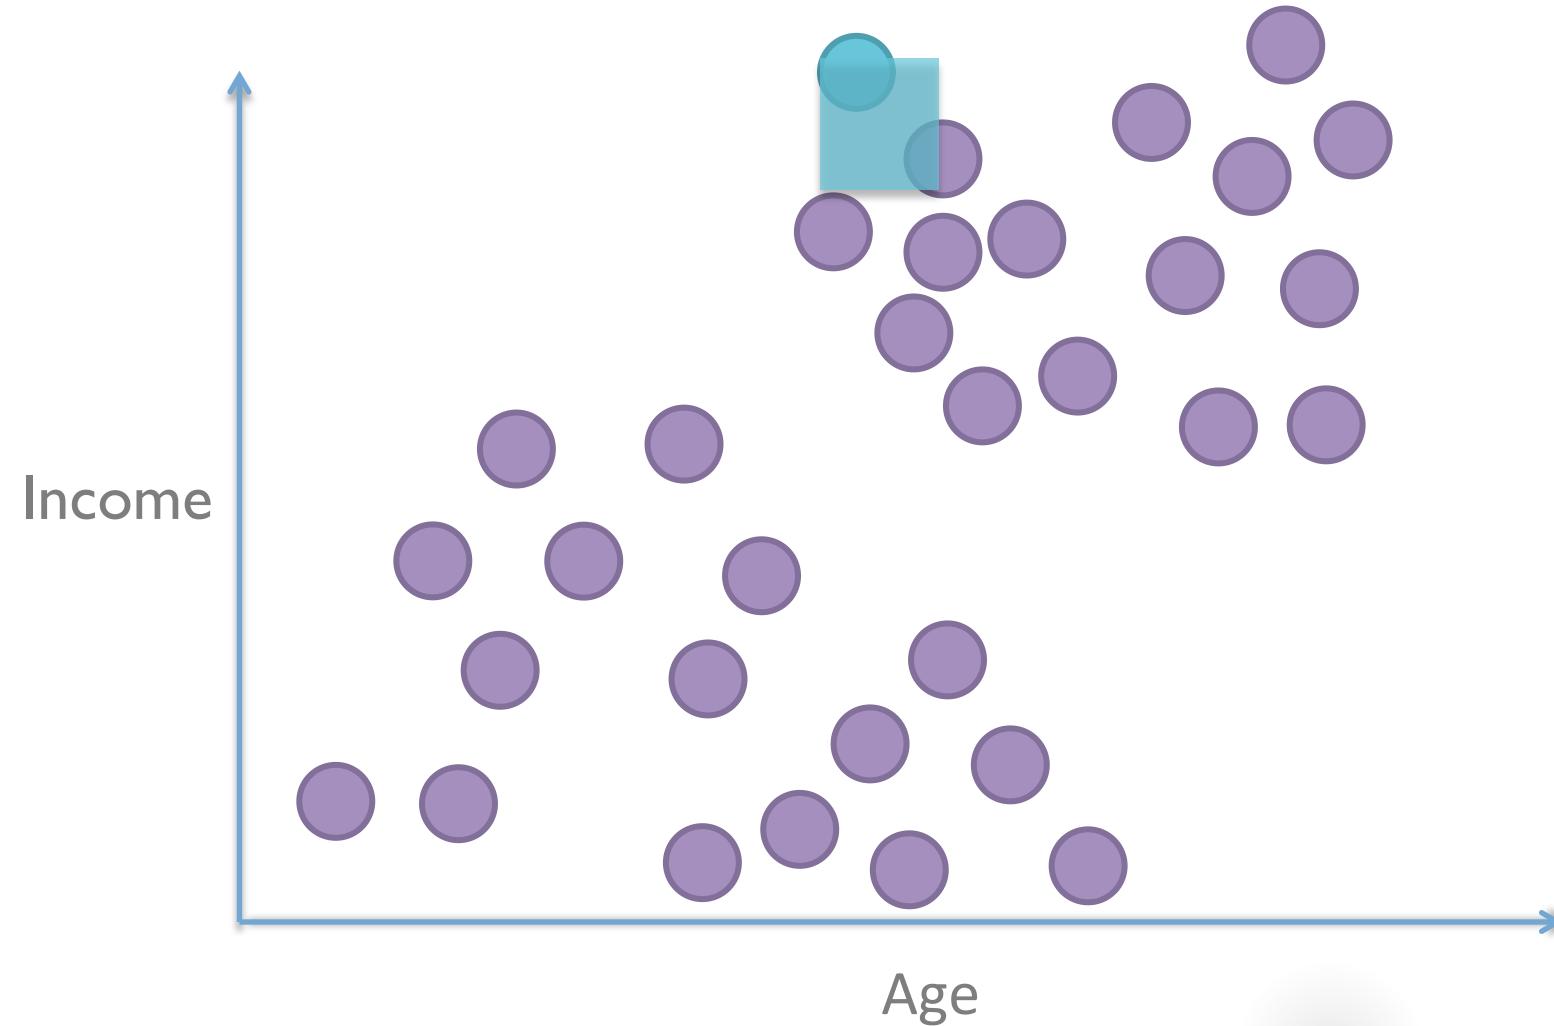
# Mean Shift

Sample local density, follow gradient towards denser direction



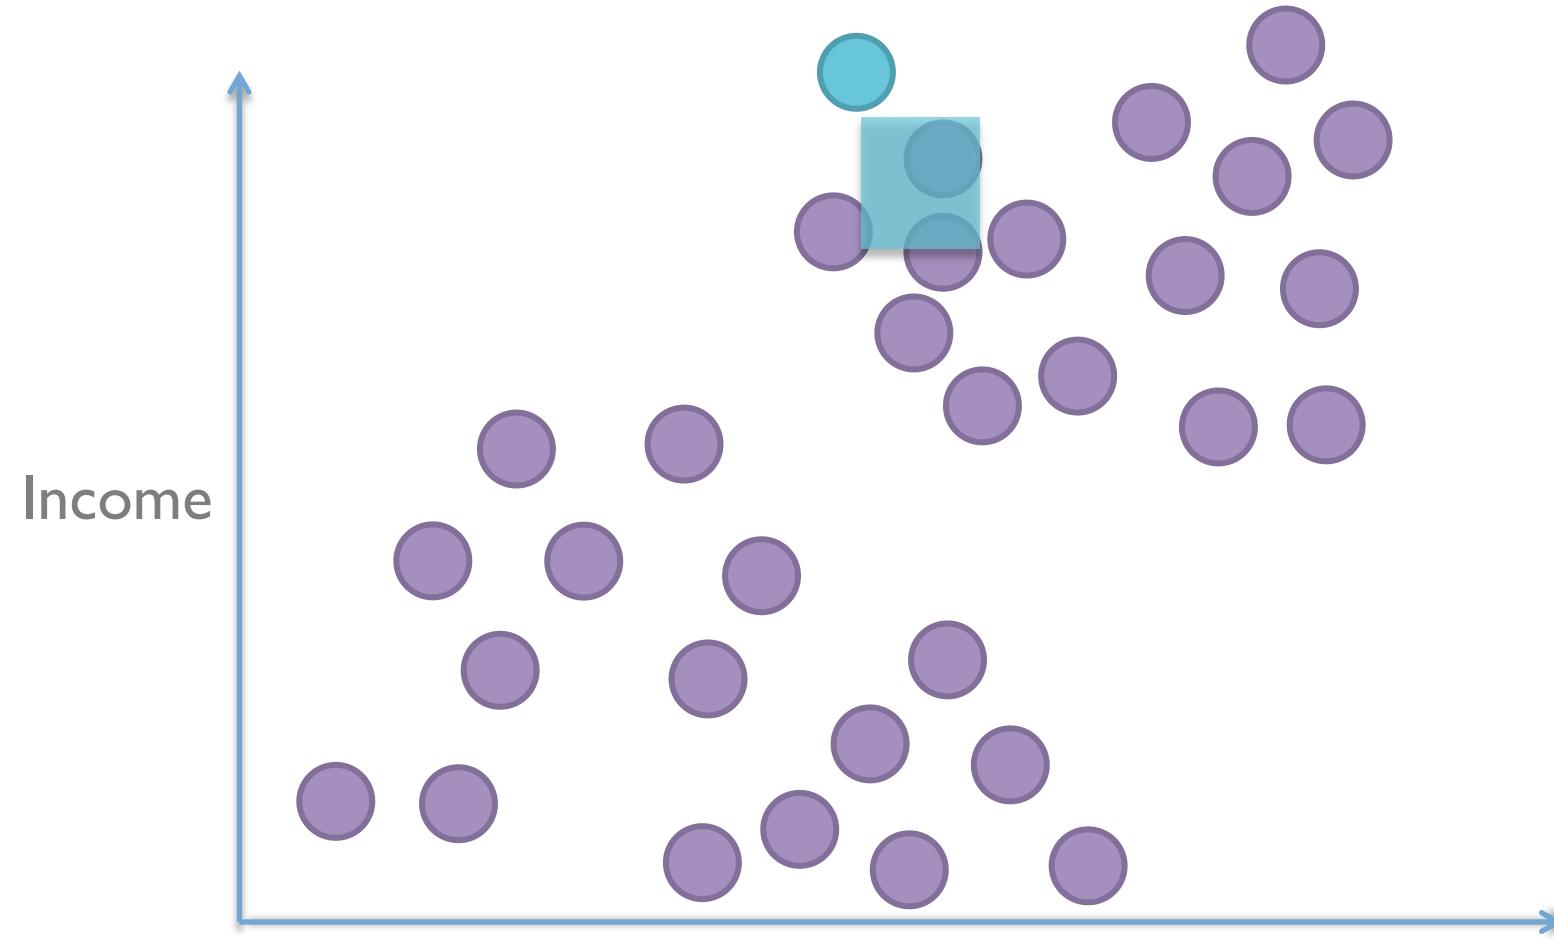
# Mean Shift

Sample local density, follow gradient towards denser direction



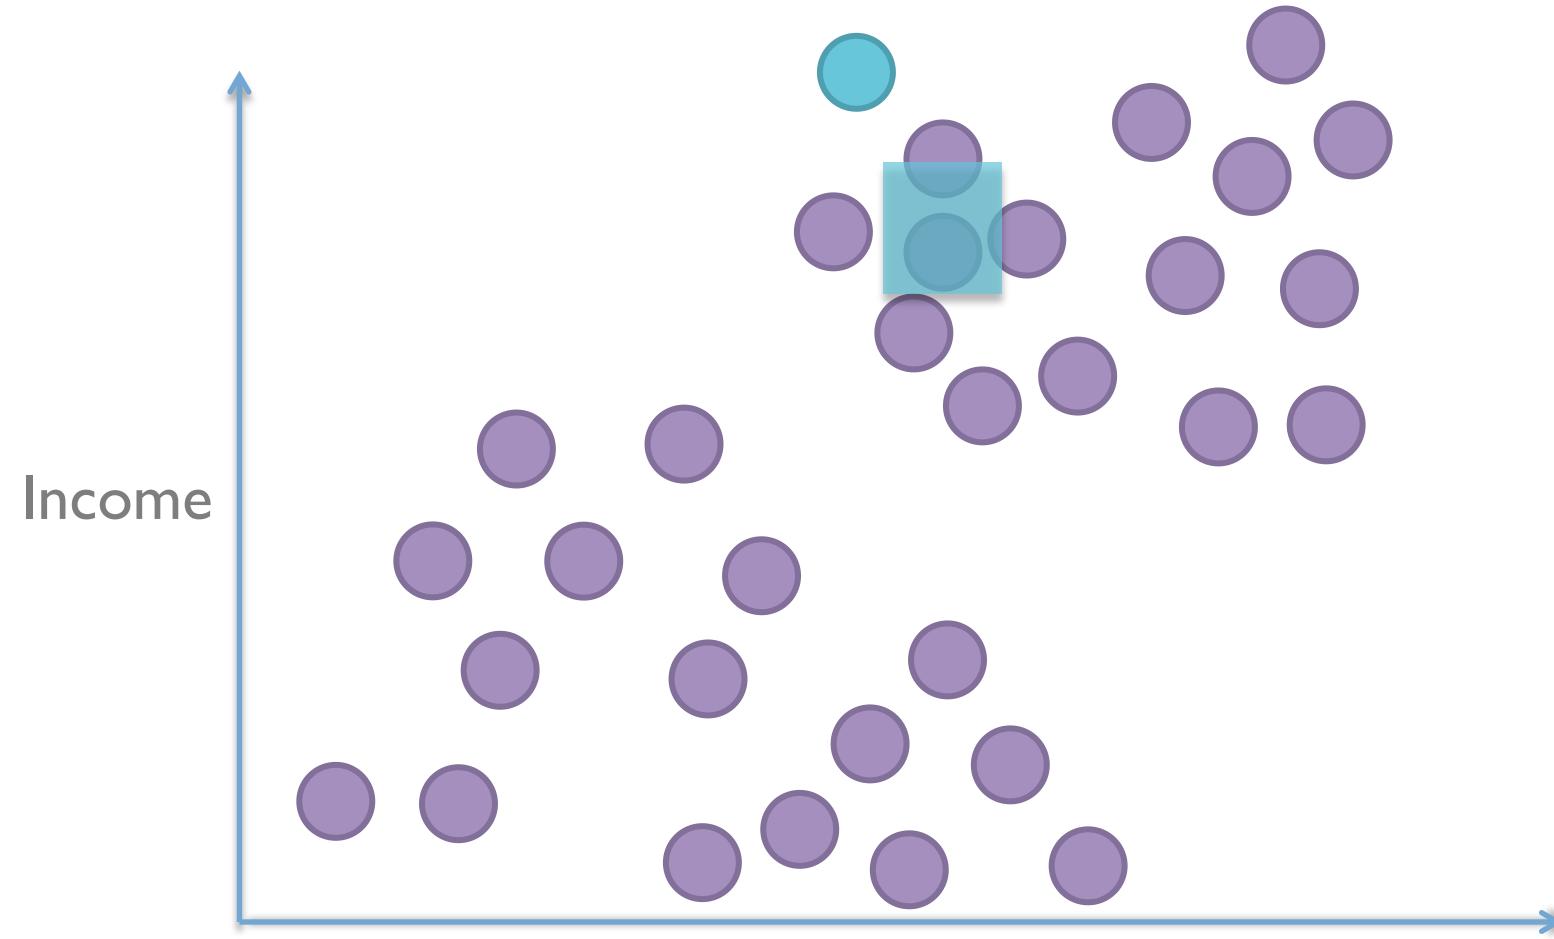
# Mean Shift

Sample local density, follow gradient towards denser direction



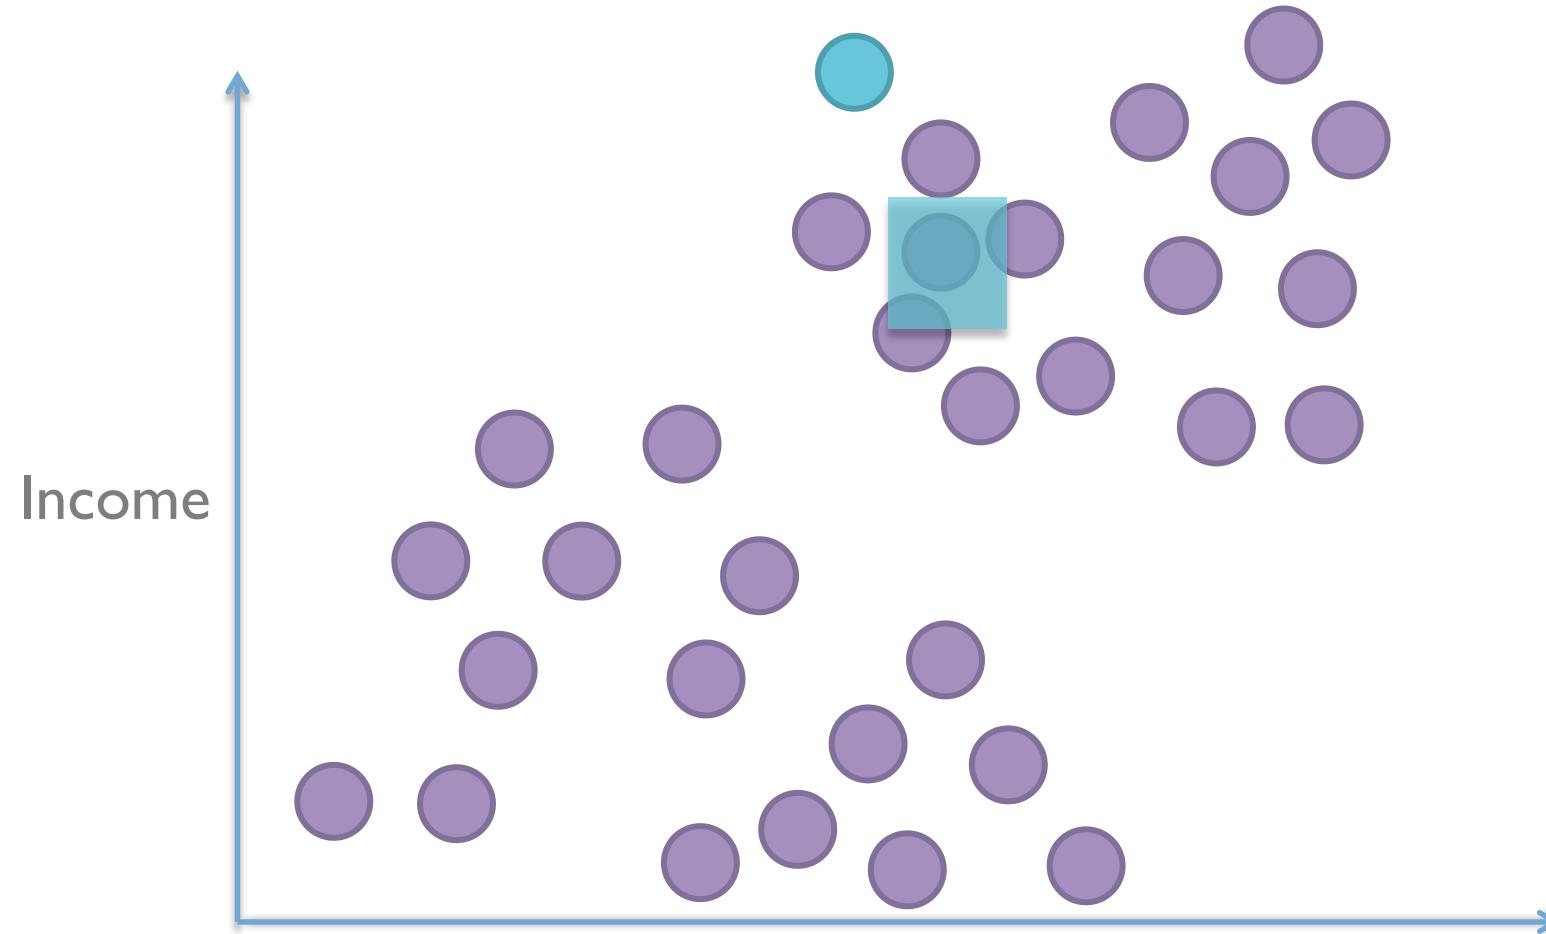
# Mean Shift

Sample local density, follow gradient towards denser direction



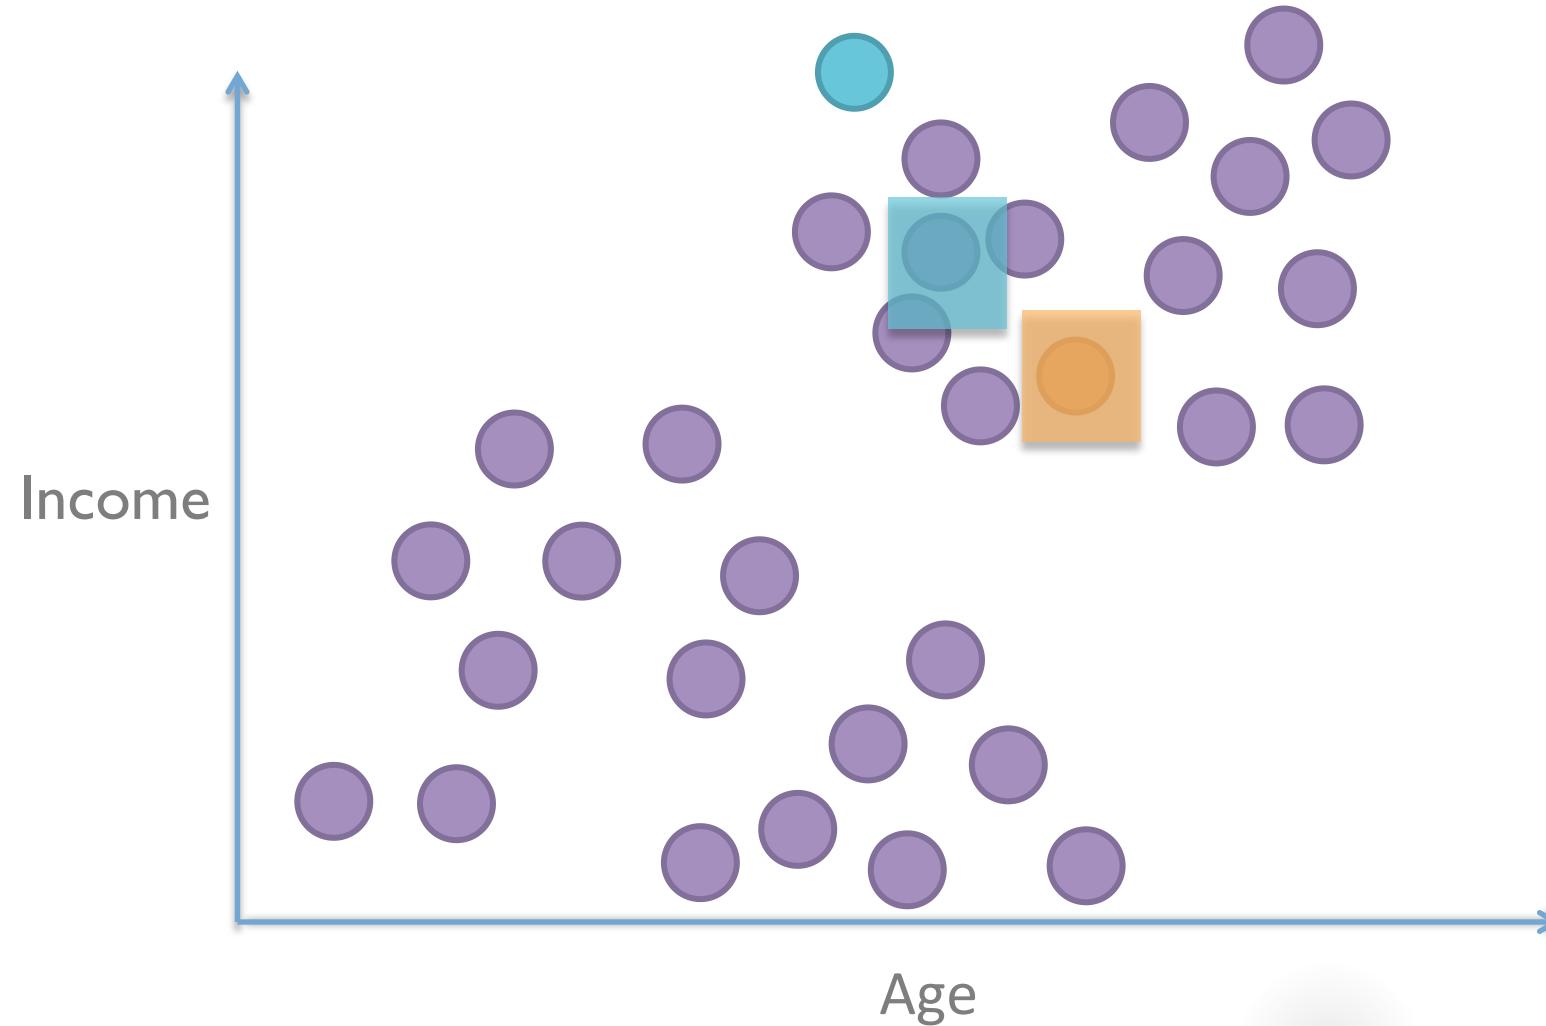
# Mean Shift

Found local density maximum! Stop.



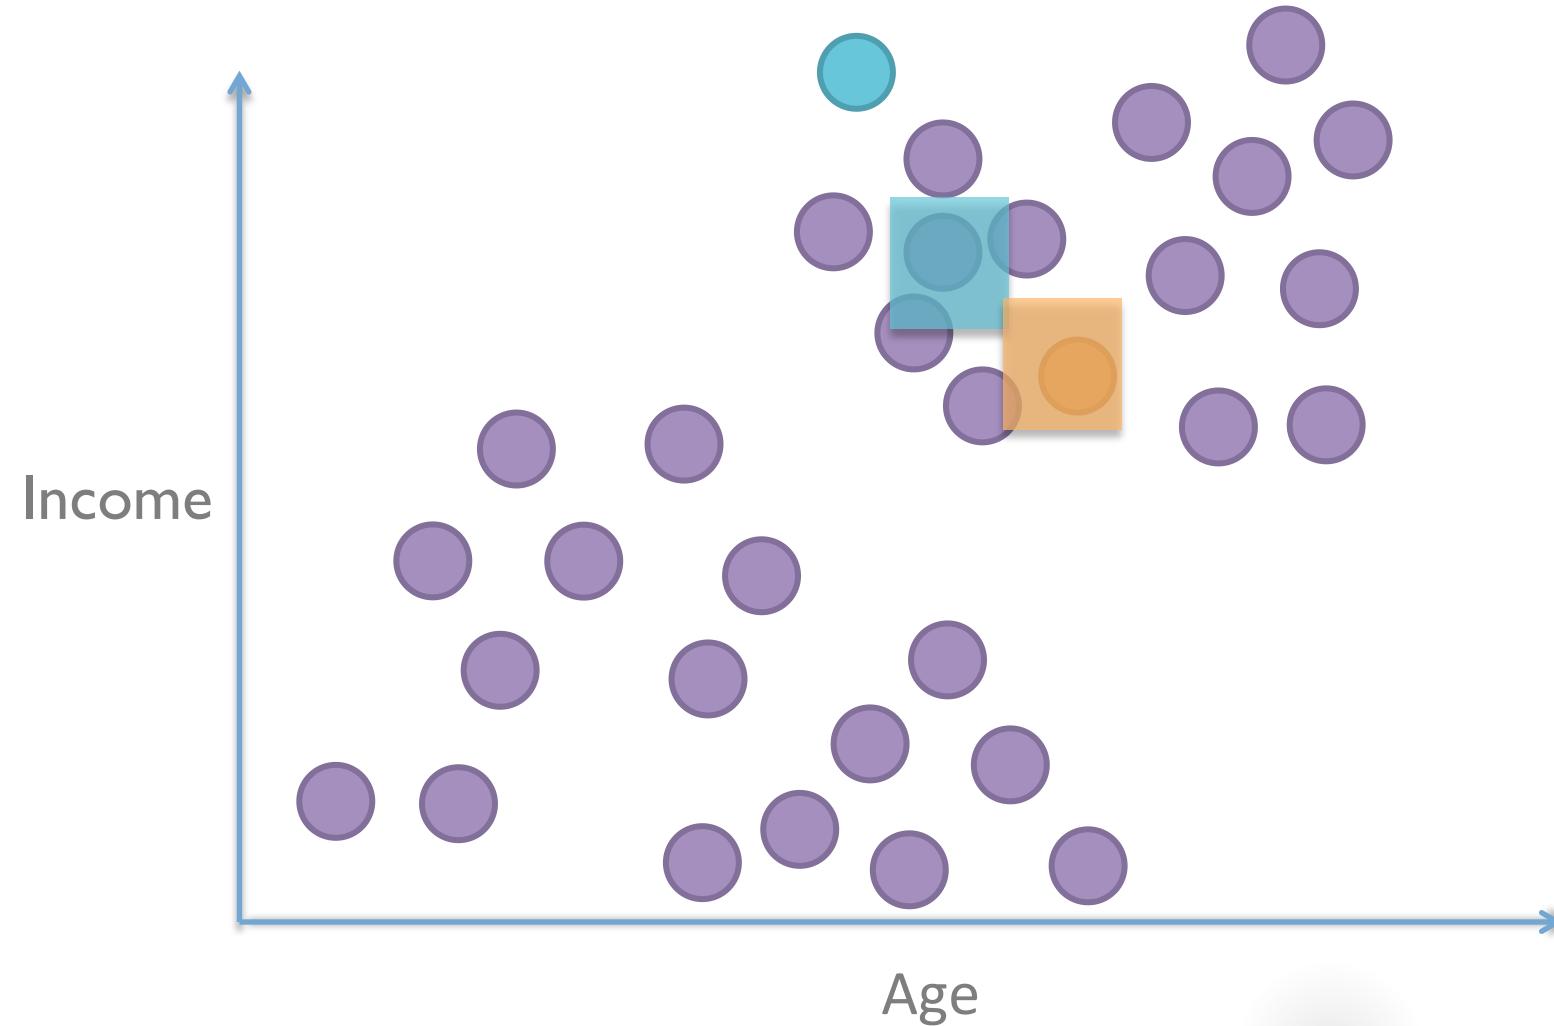
# Mean Shift

Start at another point.



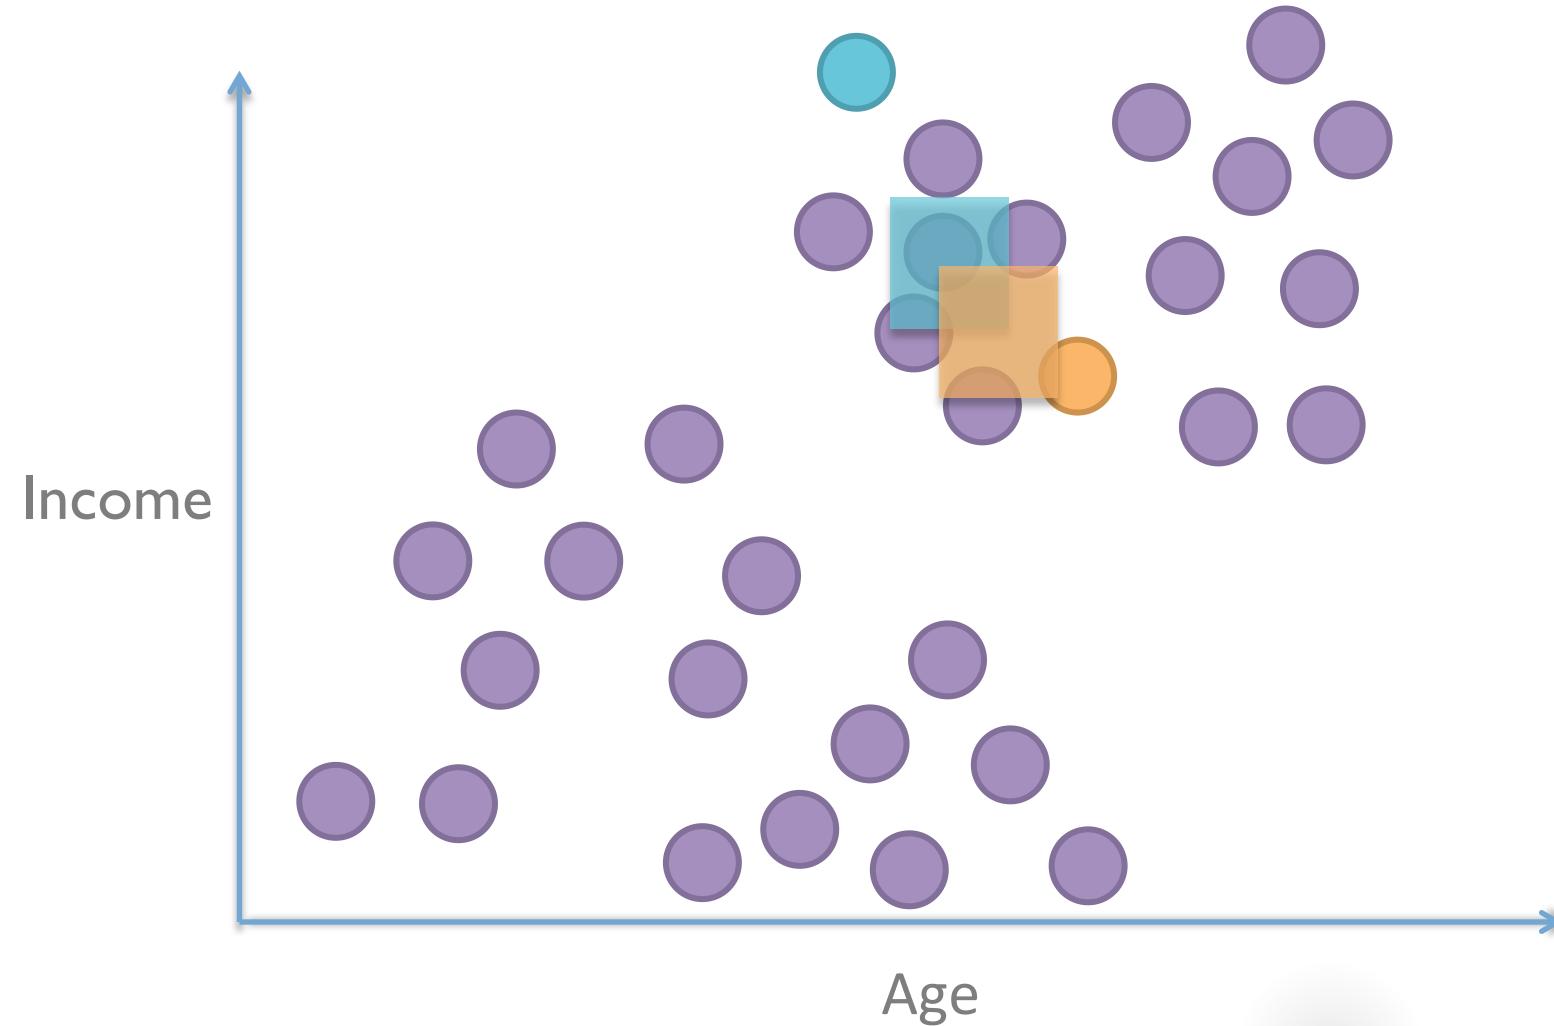
# Mean Shift

Sample local density, follow gradient towards denser direction



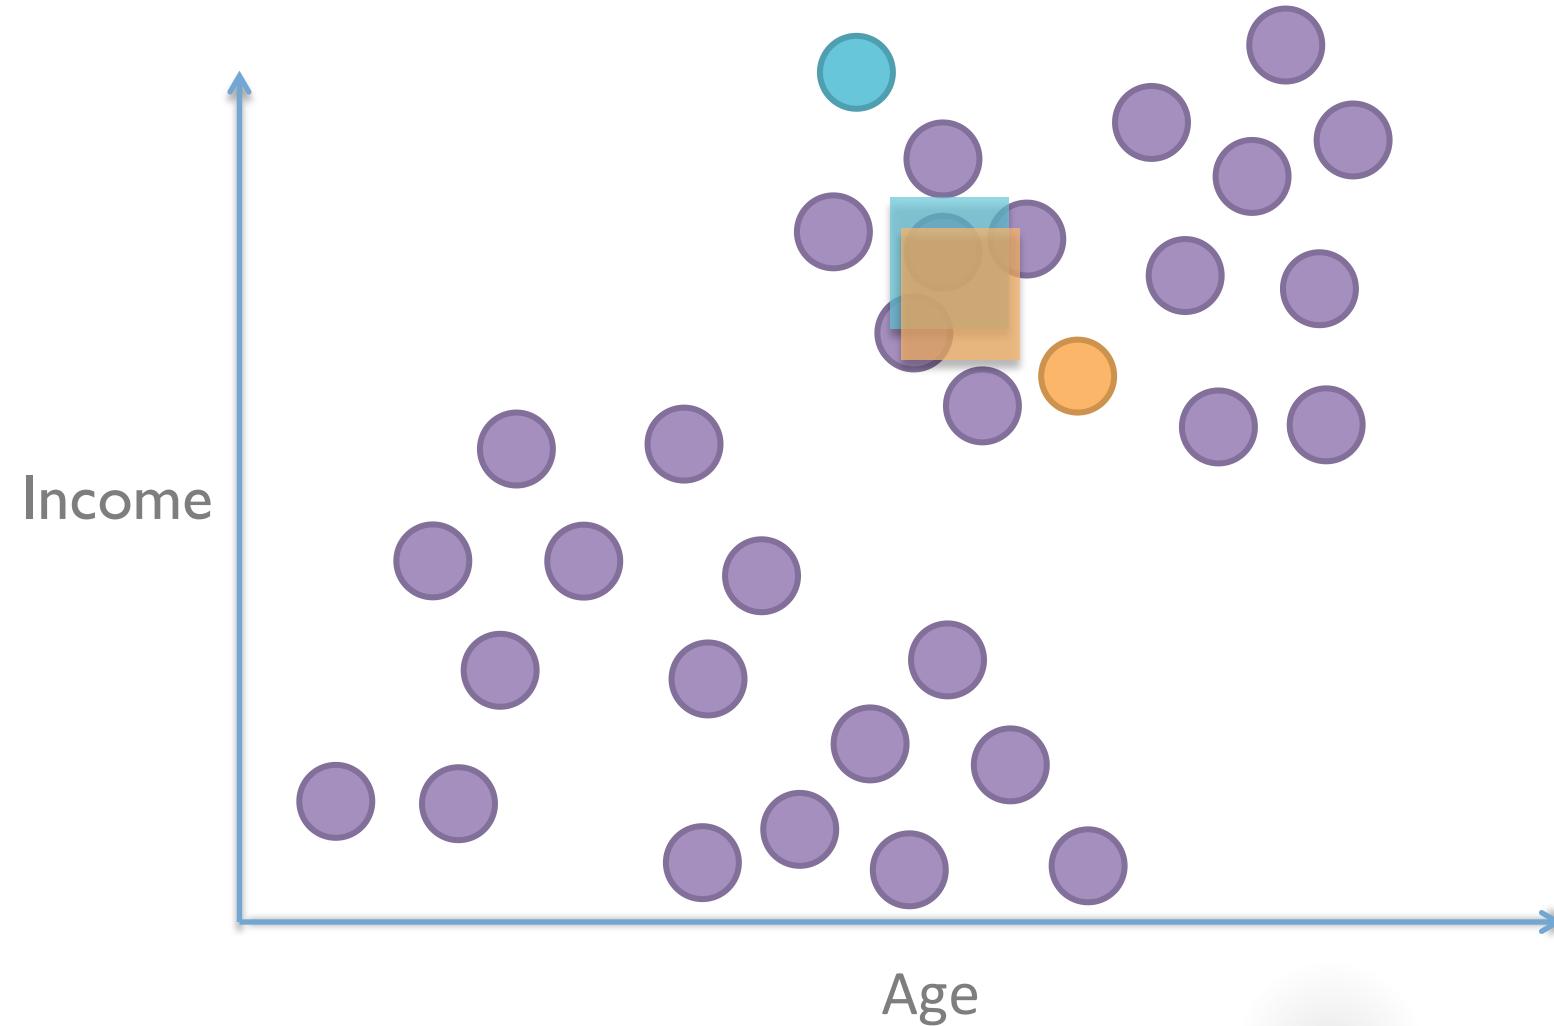
# Mean Shift

Sample local density, follow gradient towards denser direction



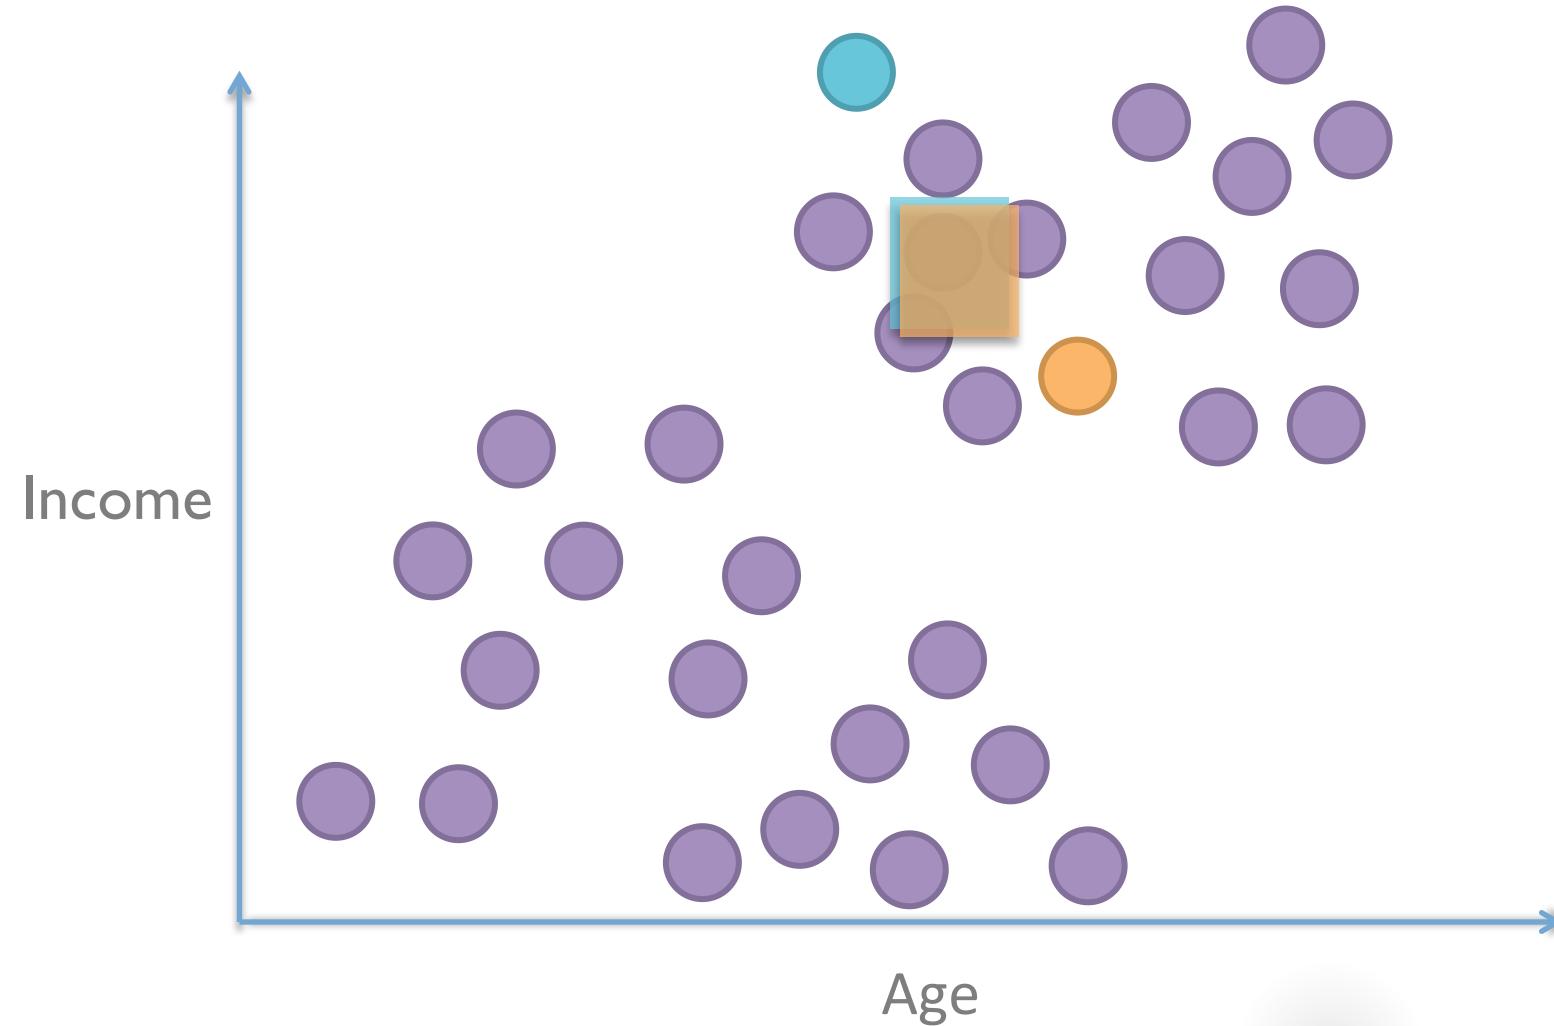
# Mean Shift

Sample local density, follow gradient towards denser direction



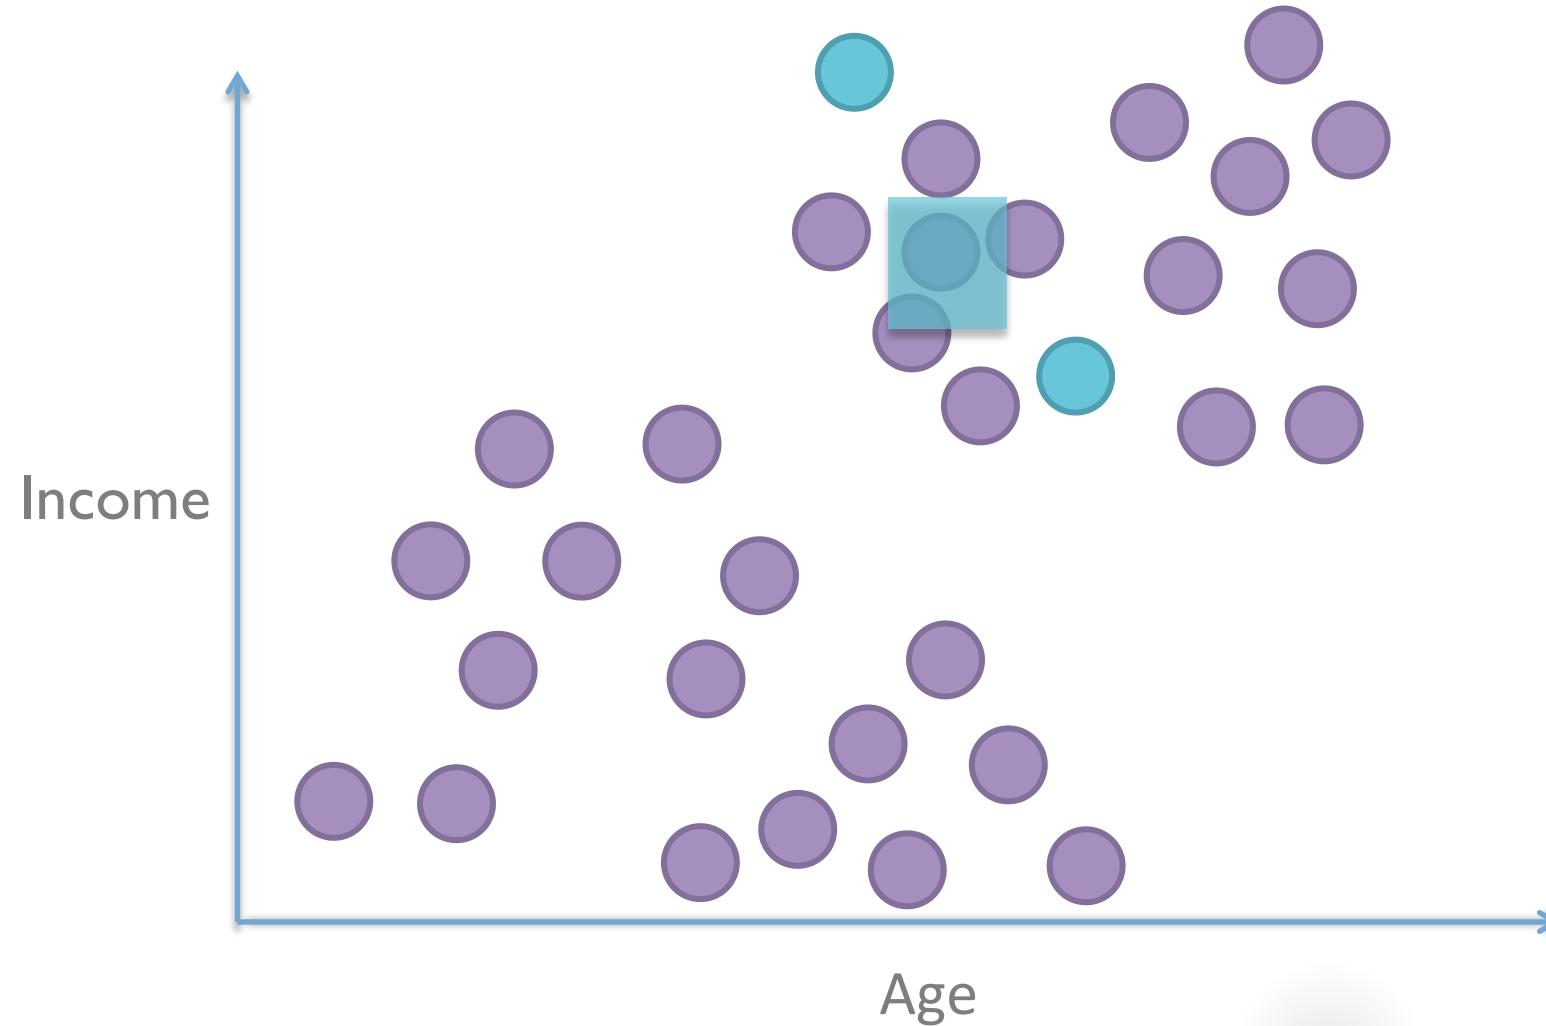
# Mean Shift

Sample local density, follow gradient towards denser direction



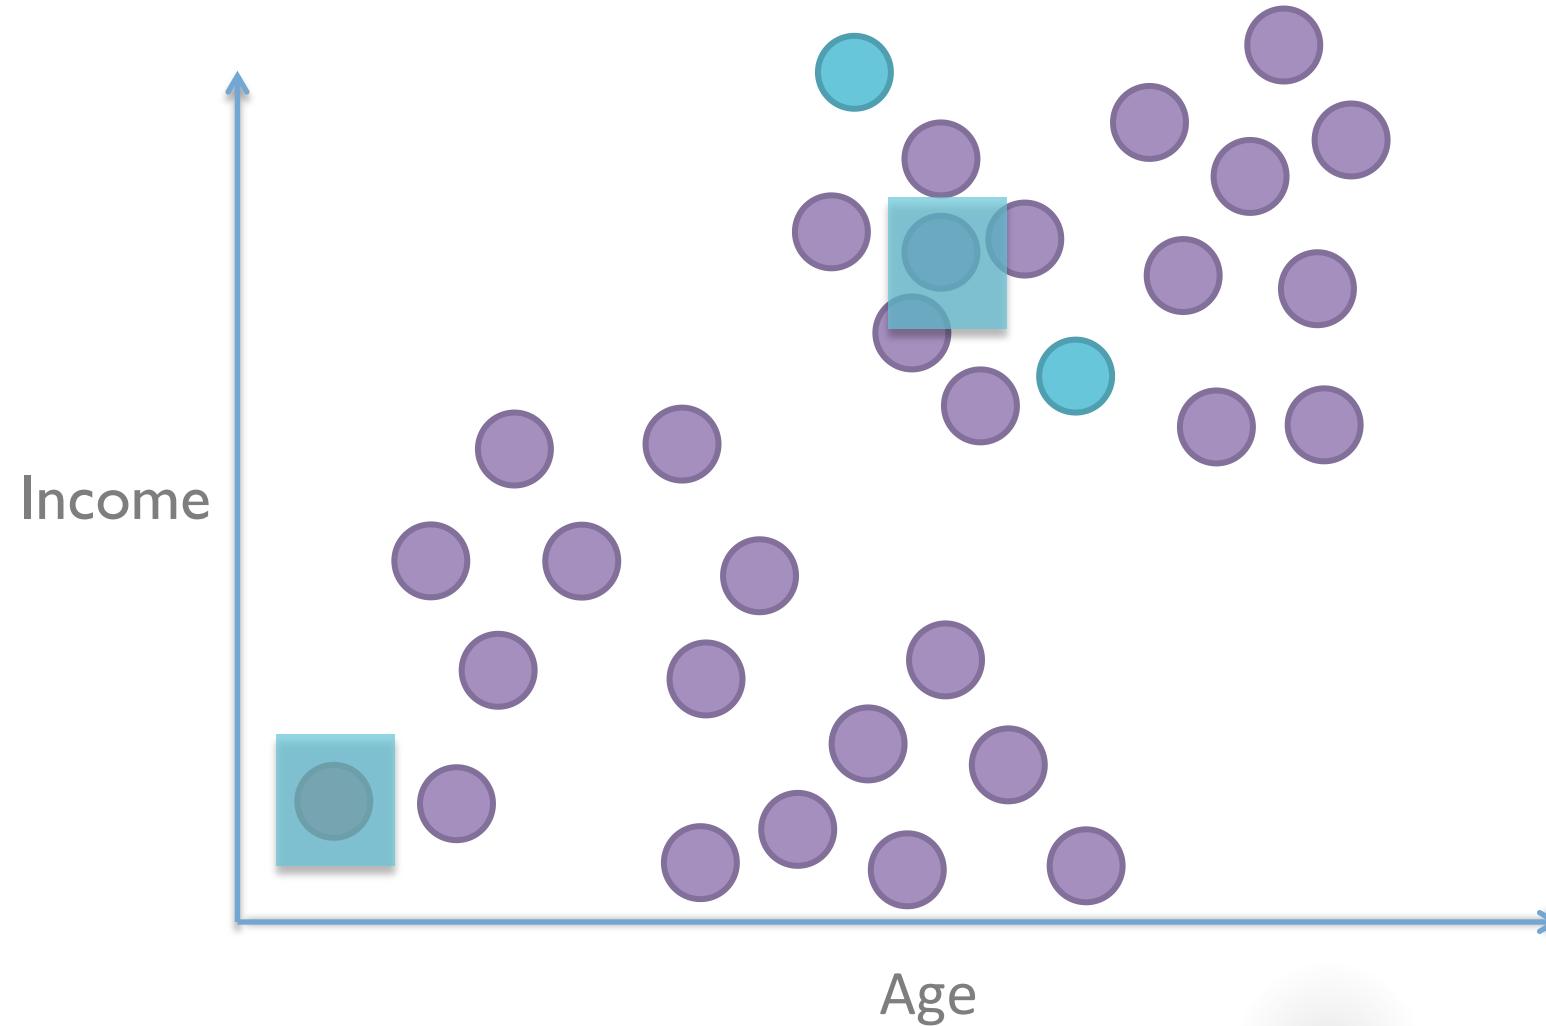
# Mean Shift

Found the same (close) local maximum, same cluster



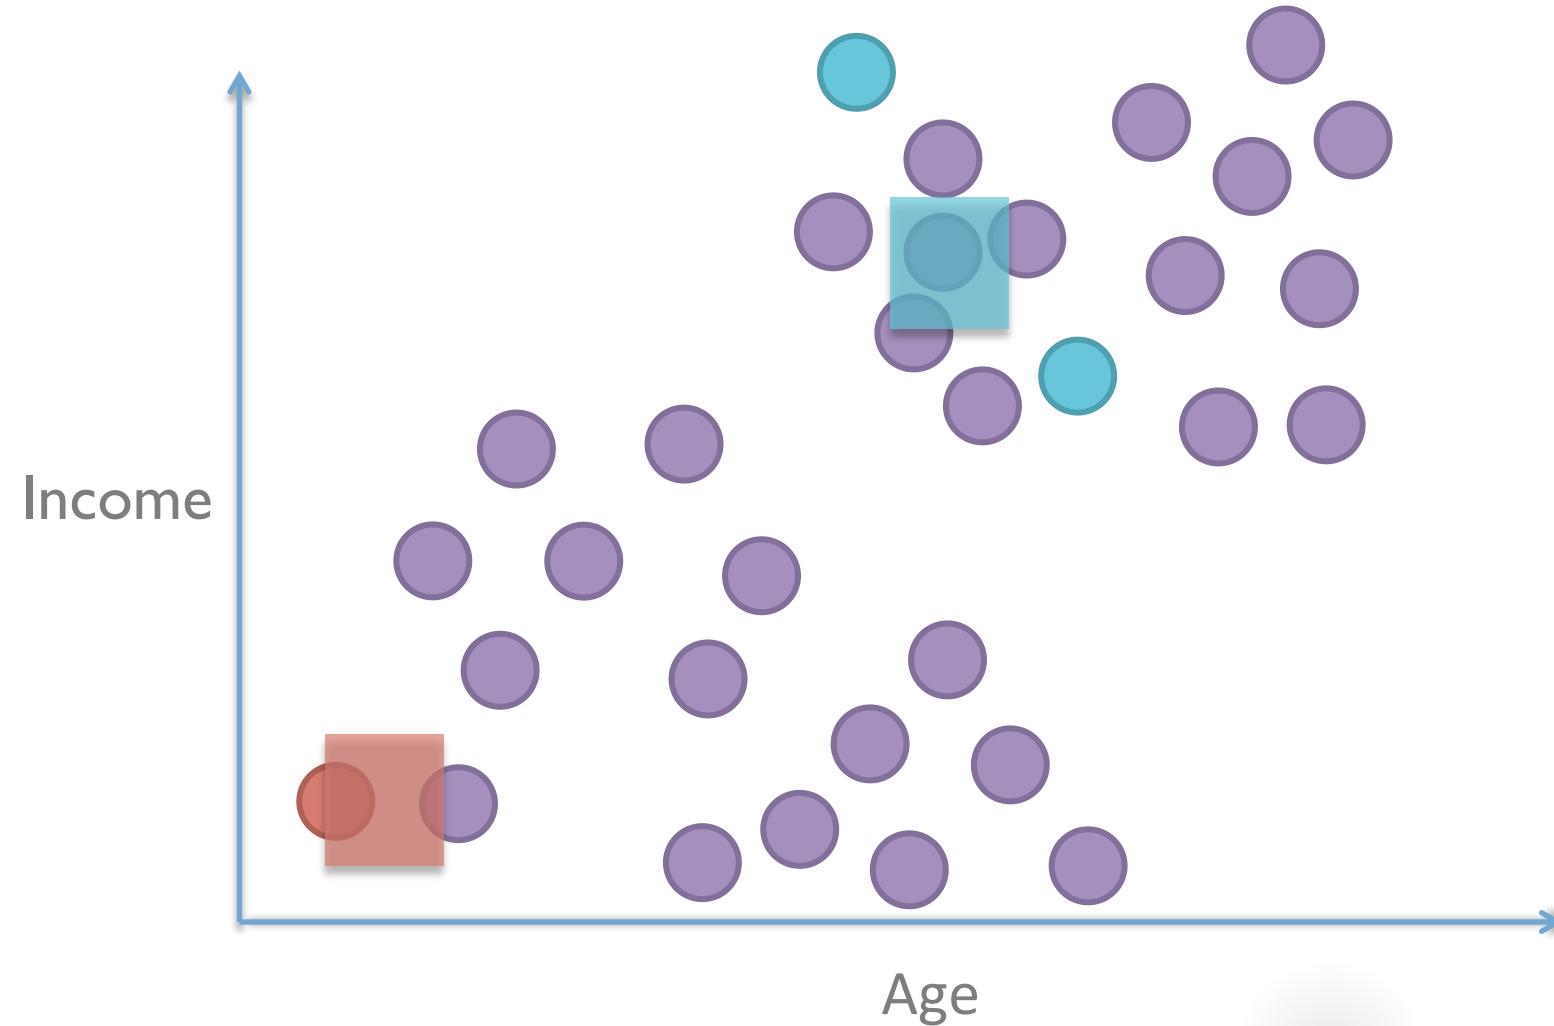
# Mean Shift

Start at another point



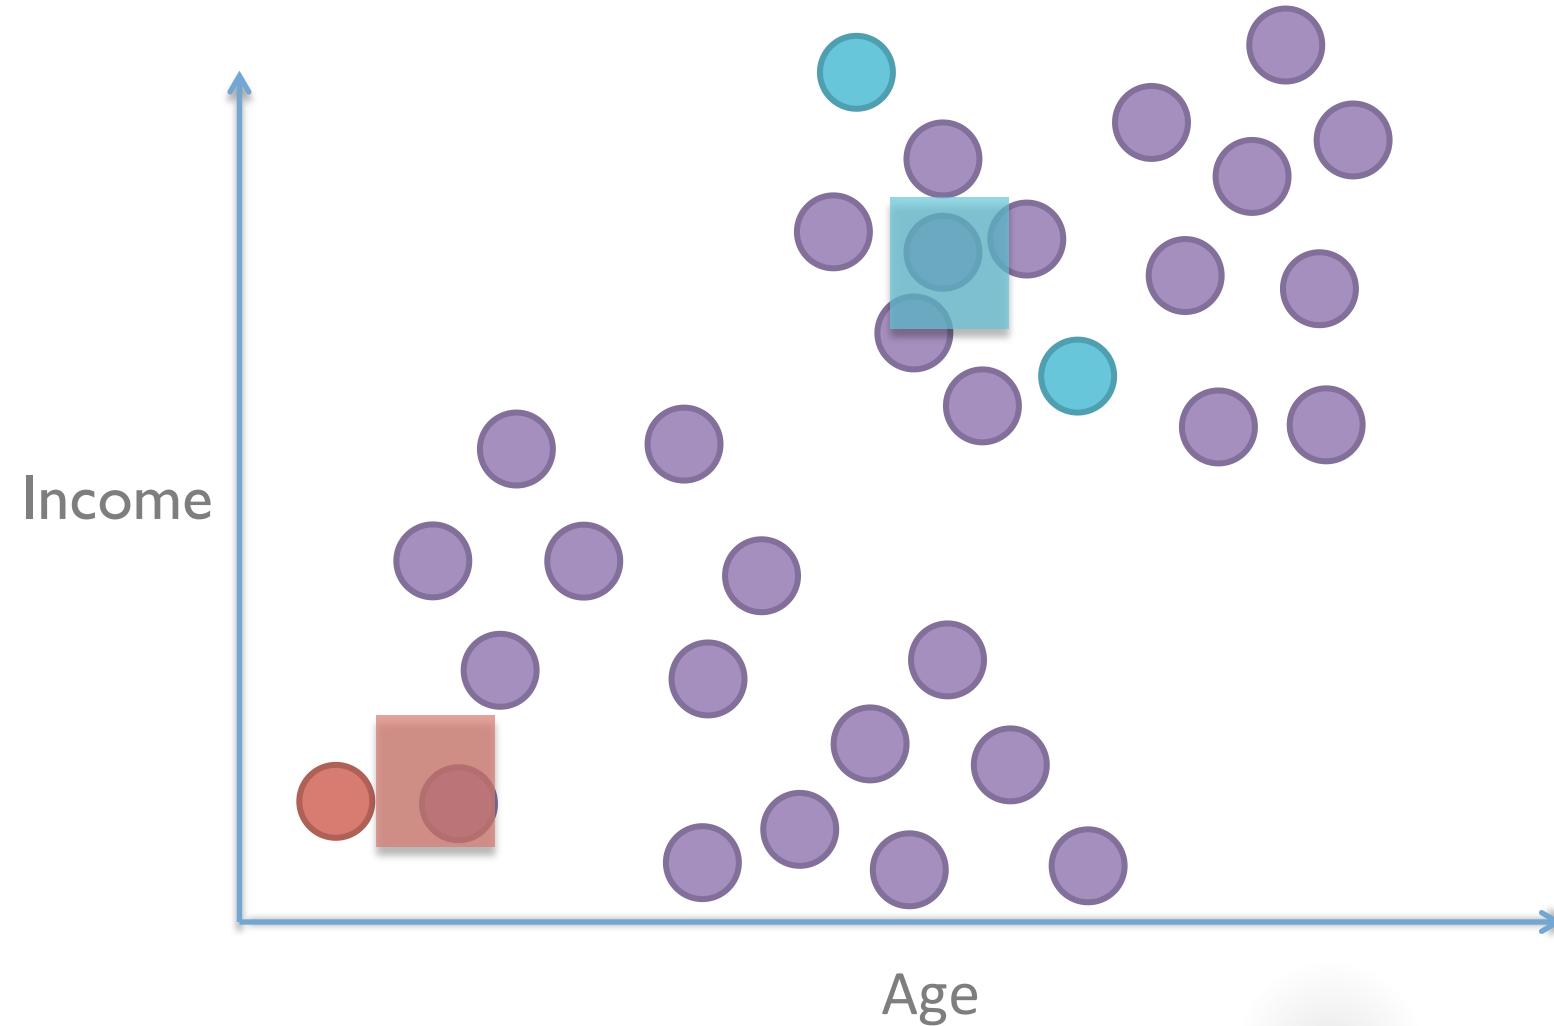
# Mean Shift

Sample local density, follow gradient towards denser direction



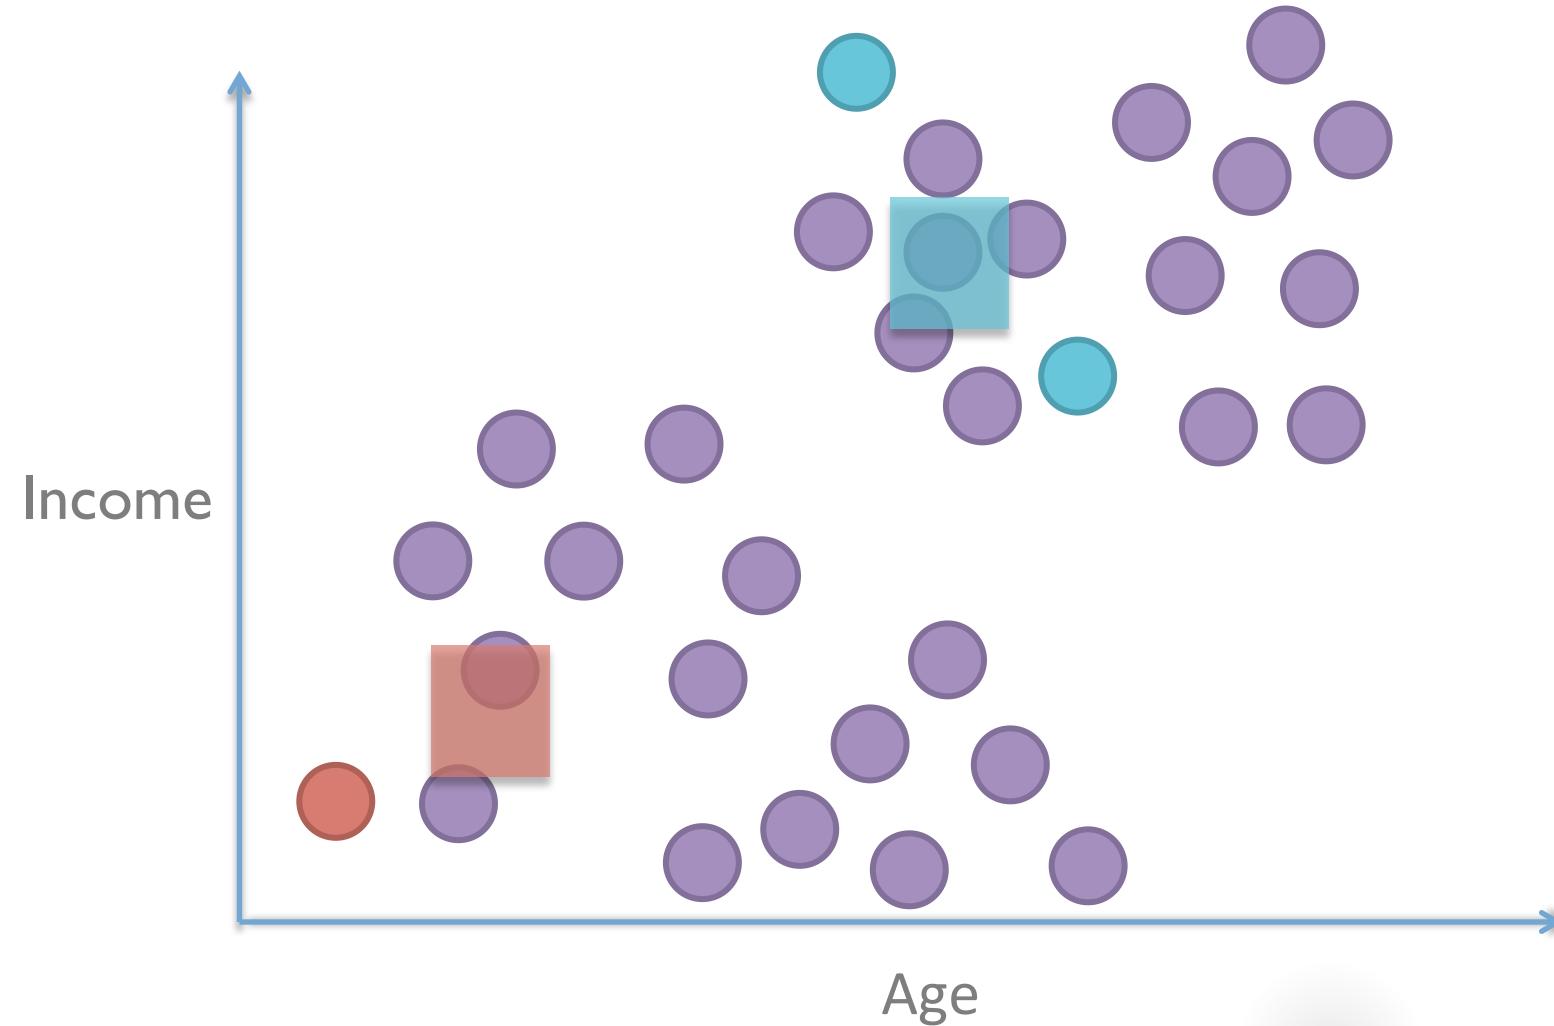
# Mean Shift

Sample local density, follow gradient towards denser direction



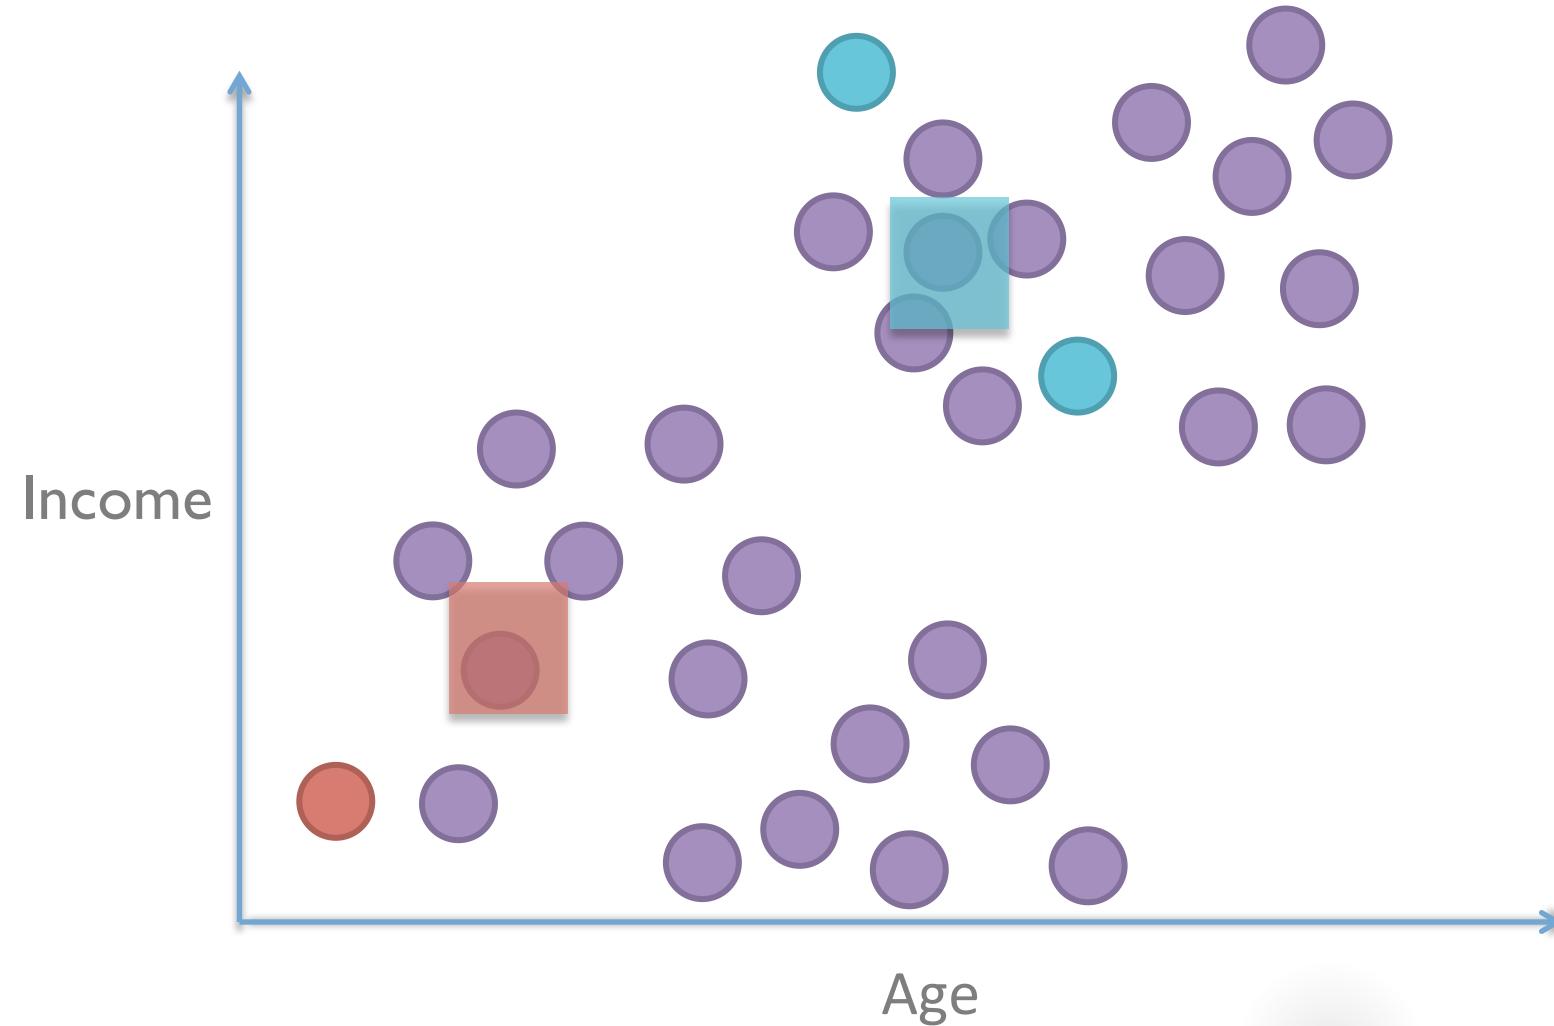
# Mean Shift

Sample local density, follow gradient towards denser direction



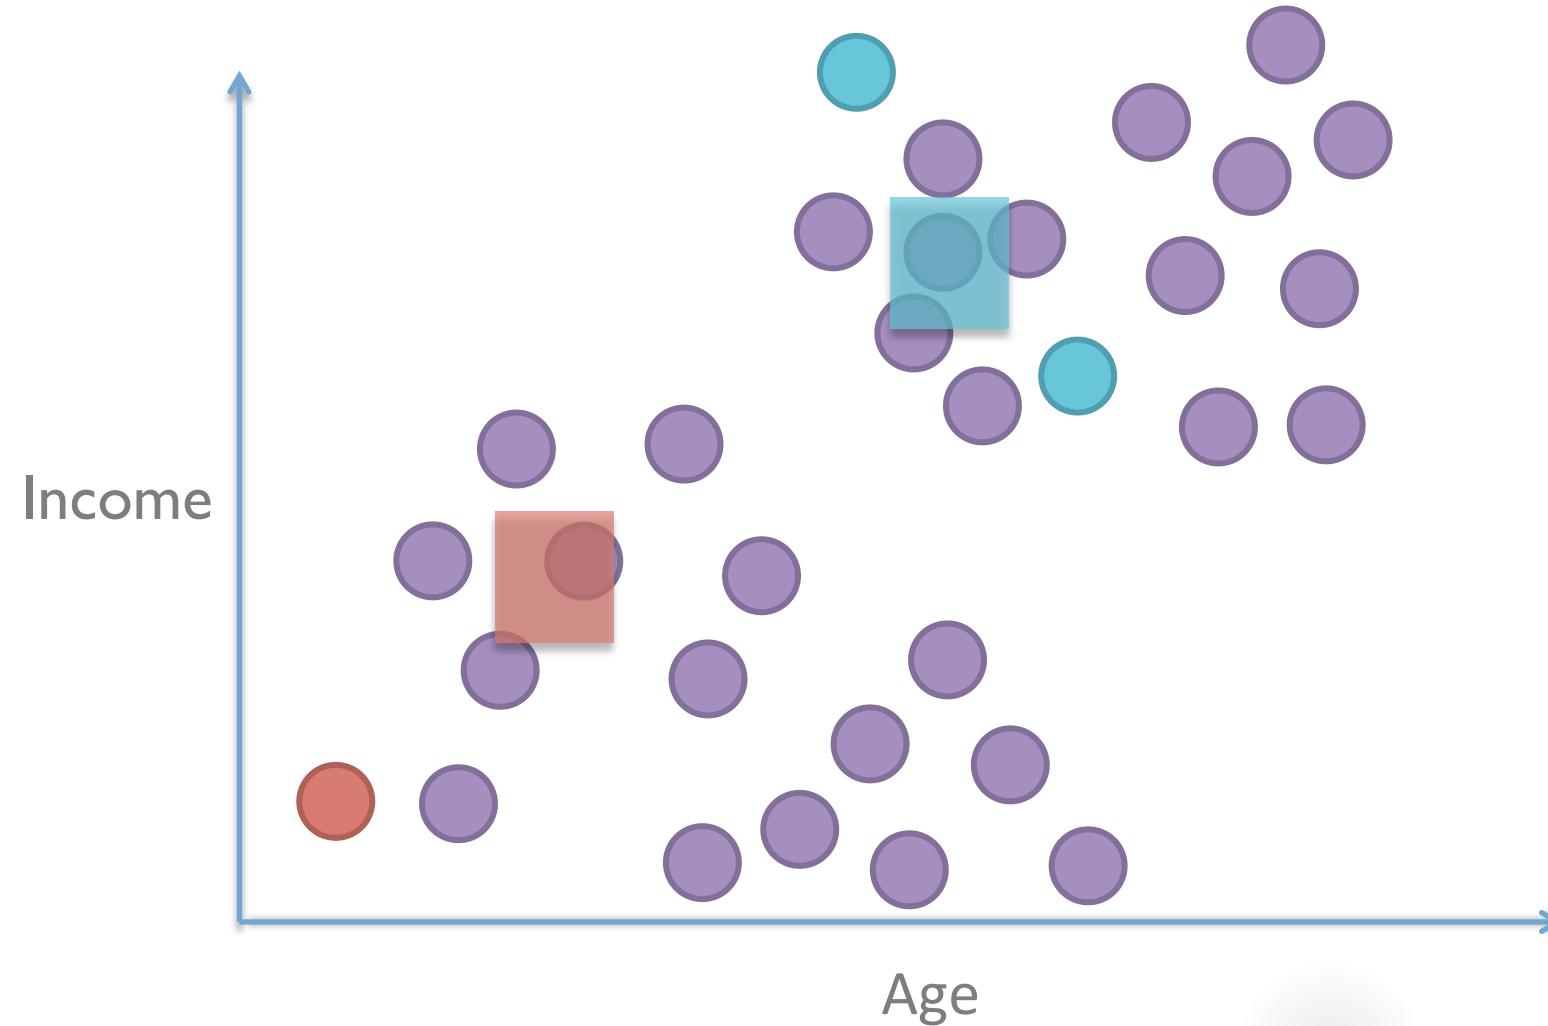
# Mean Shift

Sample local density, follow gradient towards denser direction



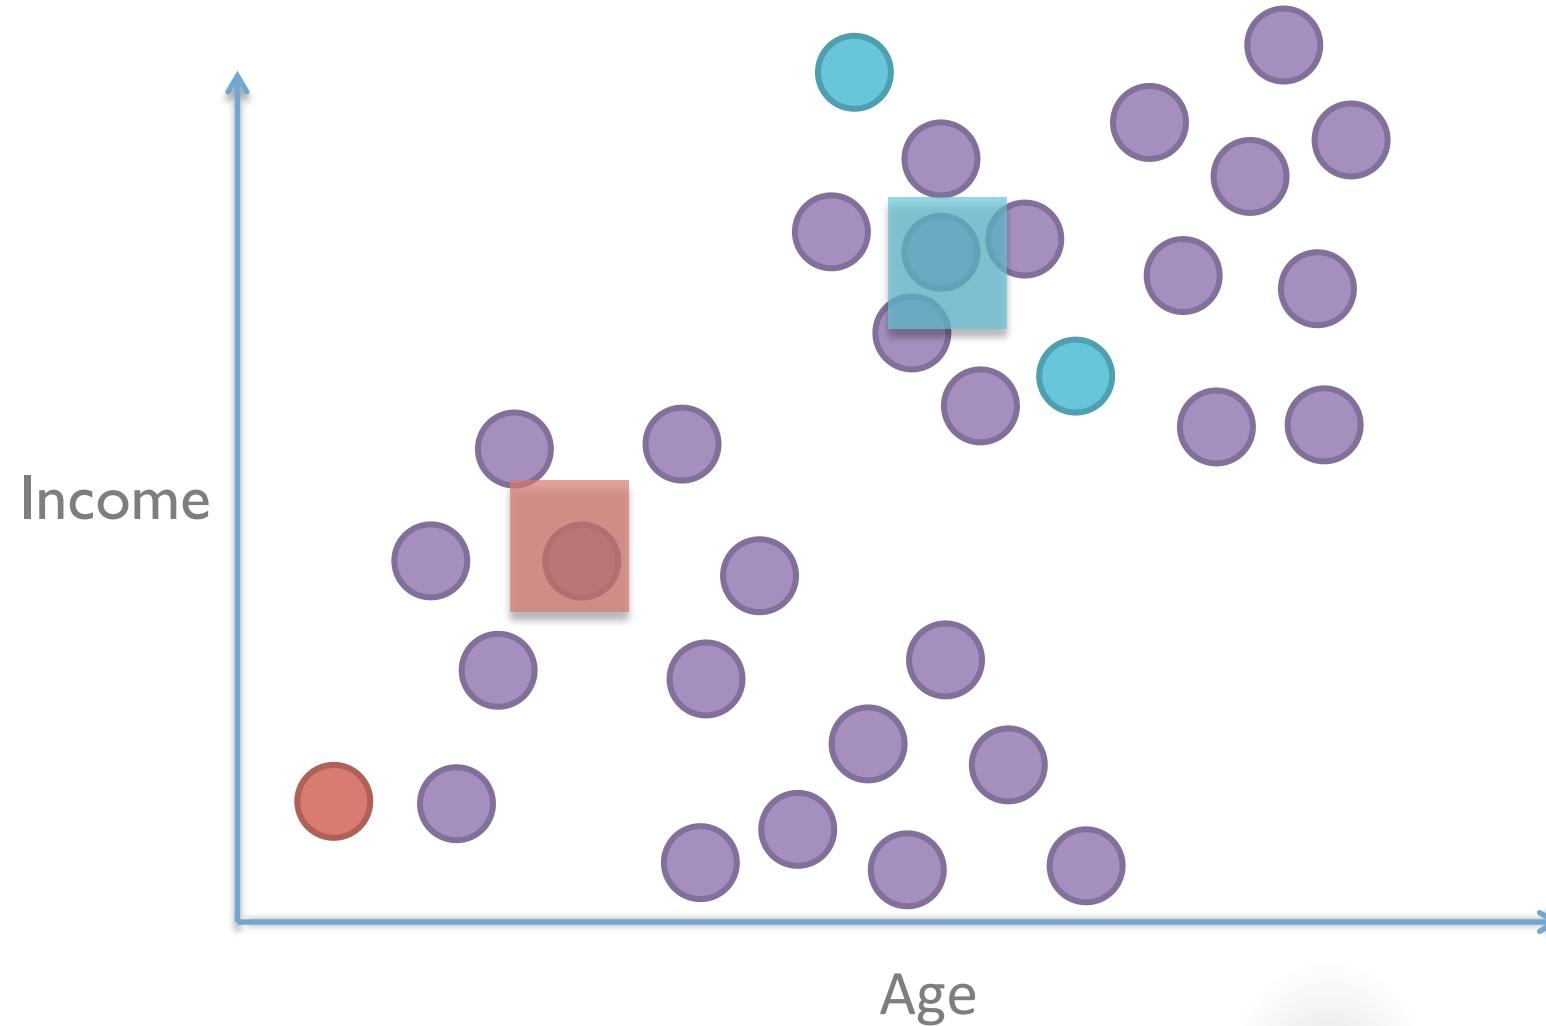
# Mean Shift

Sample local density, follow gradient towards denser direction



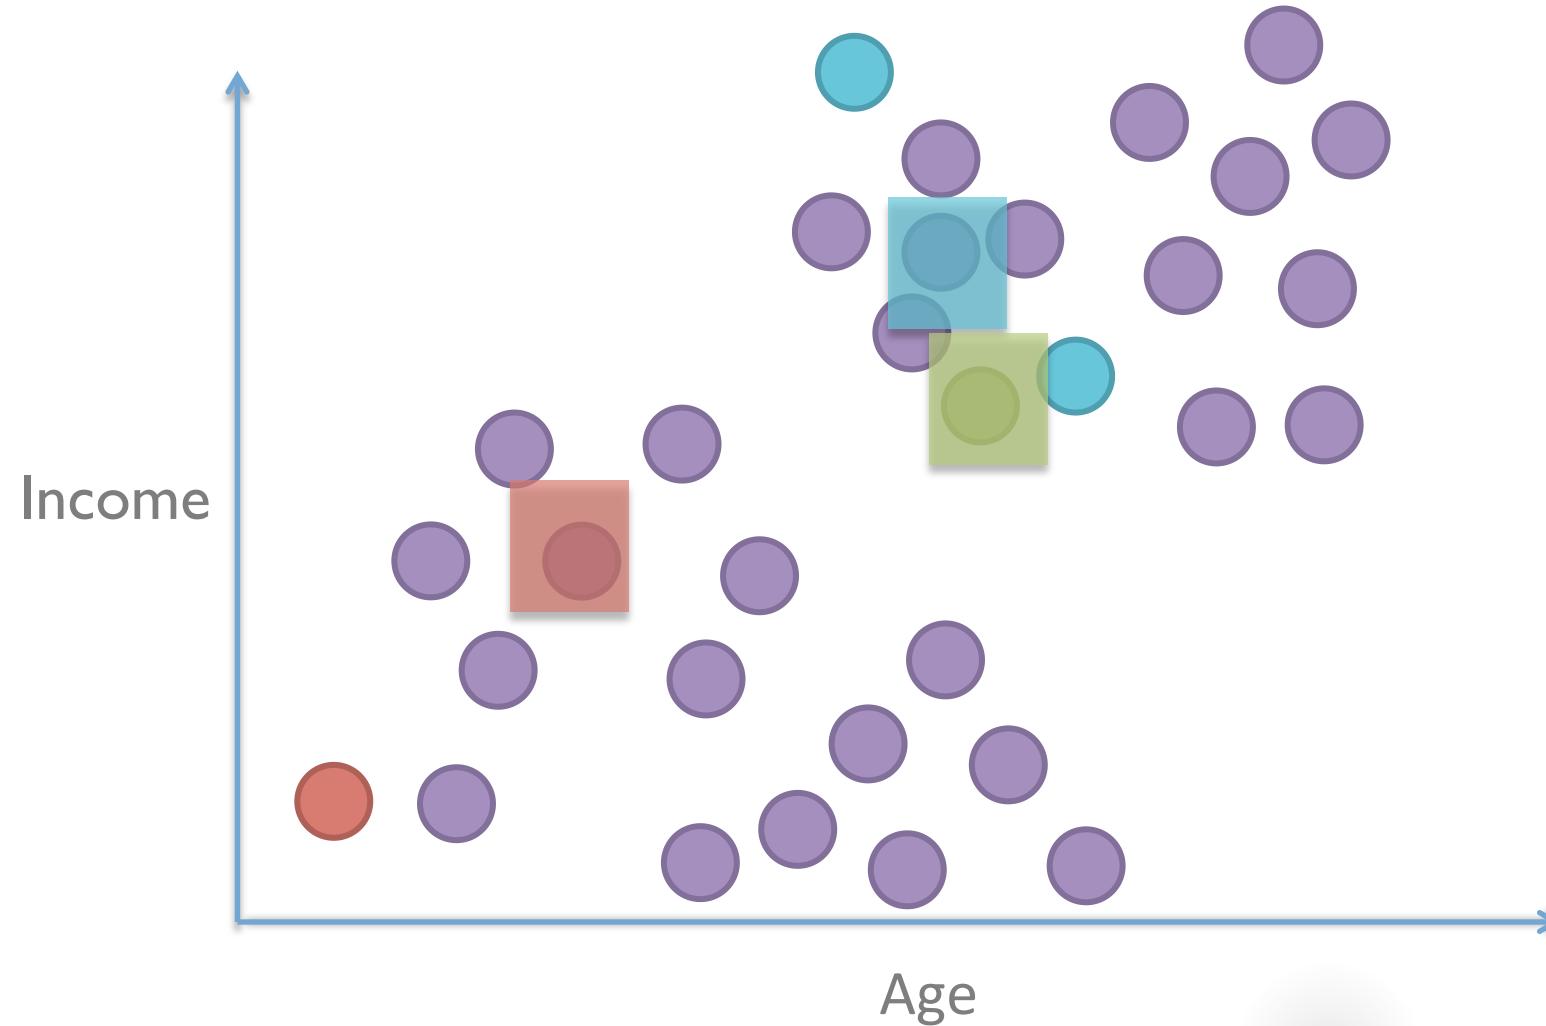
# Mean Shift

Found local maximum. Stop.



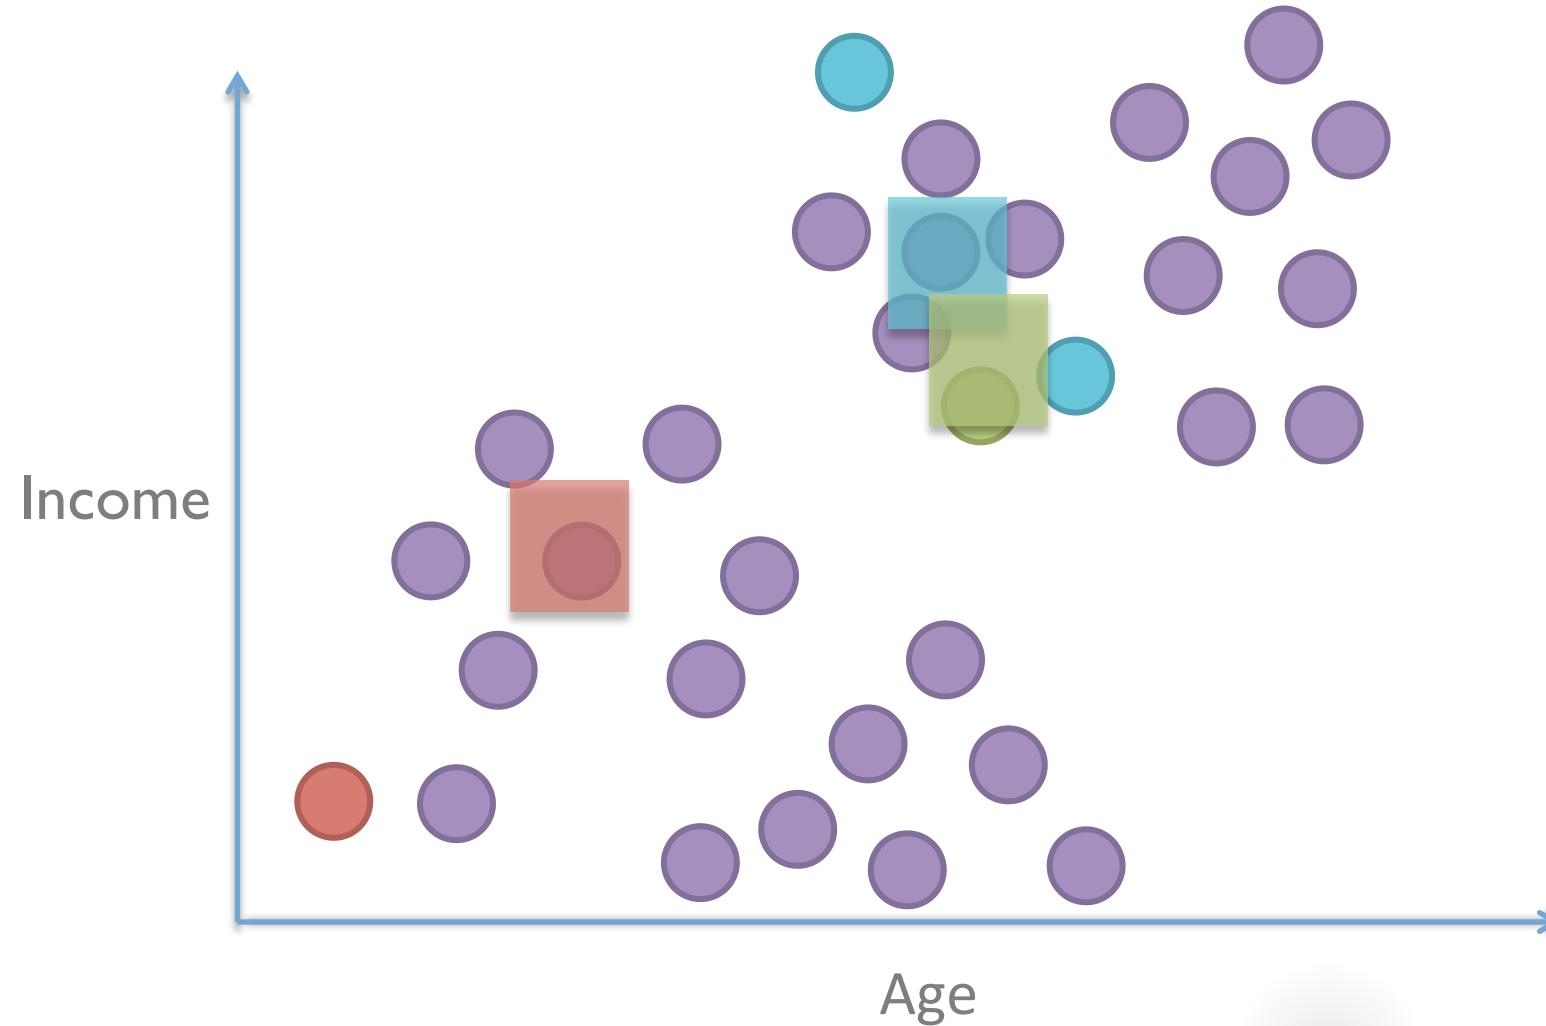
# Mean Shift

Start at another point



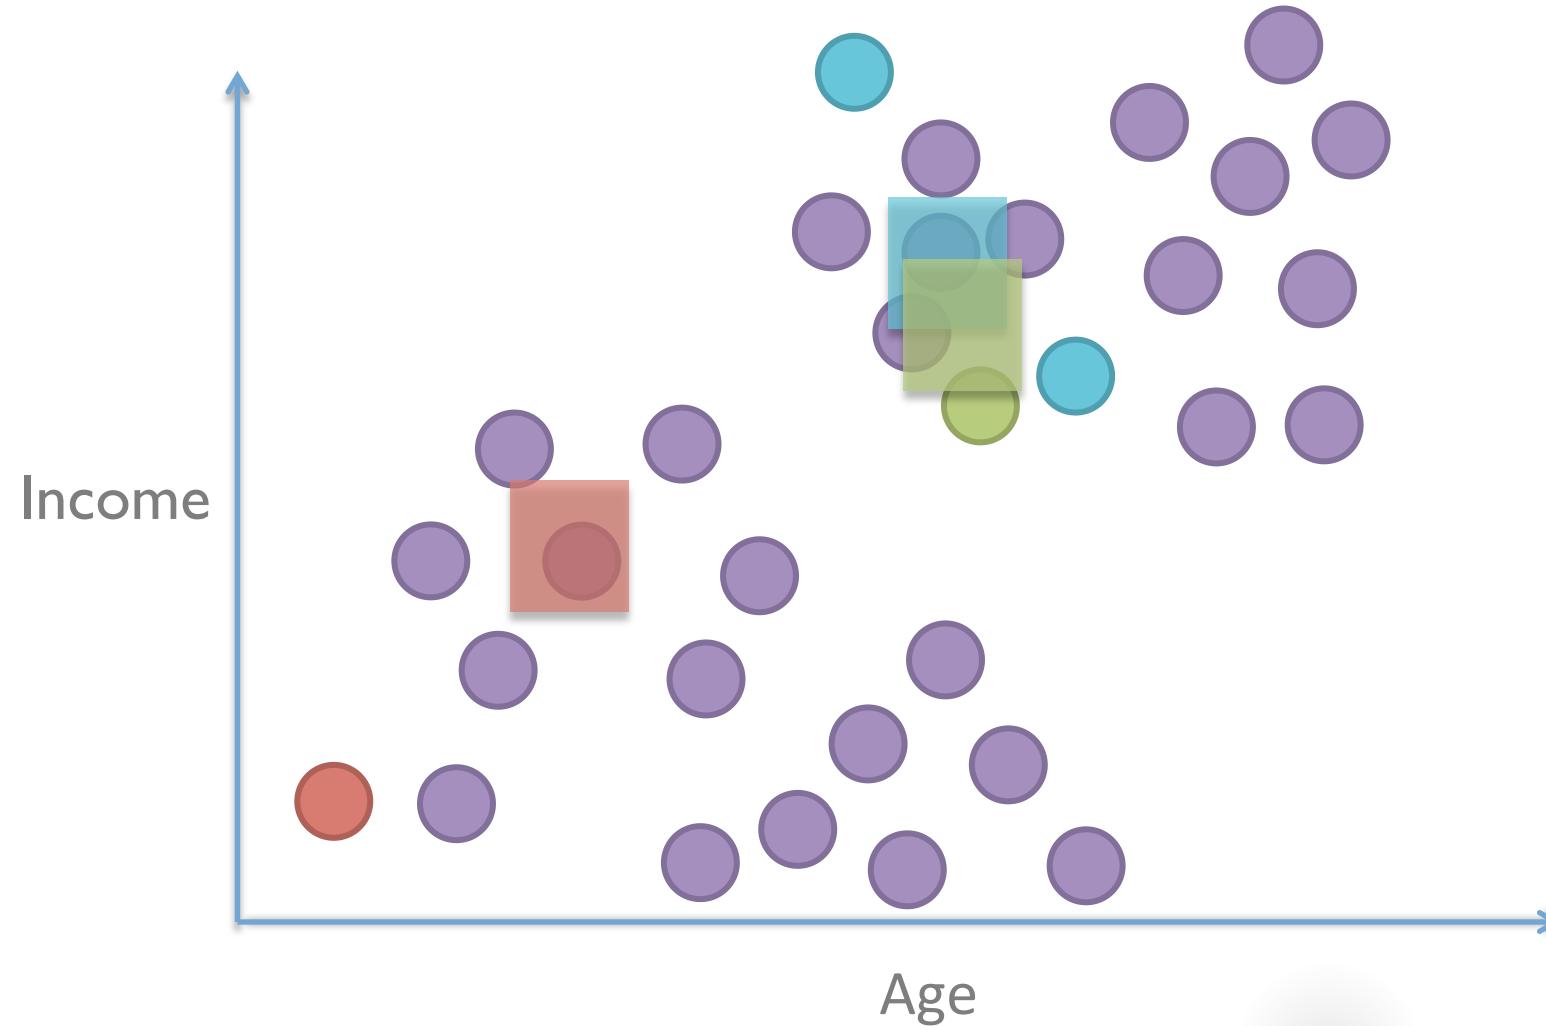
# Mean Shift

Search for local density maximum



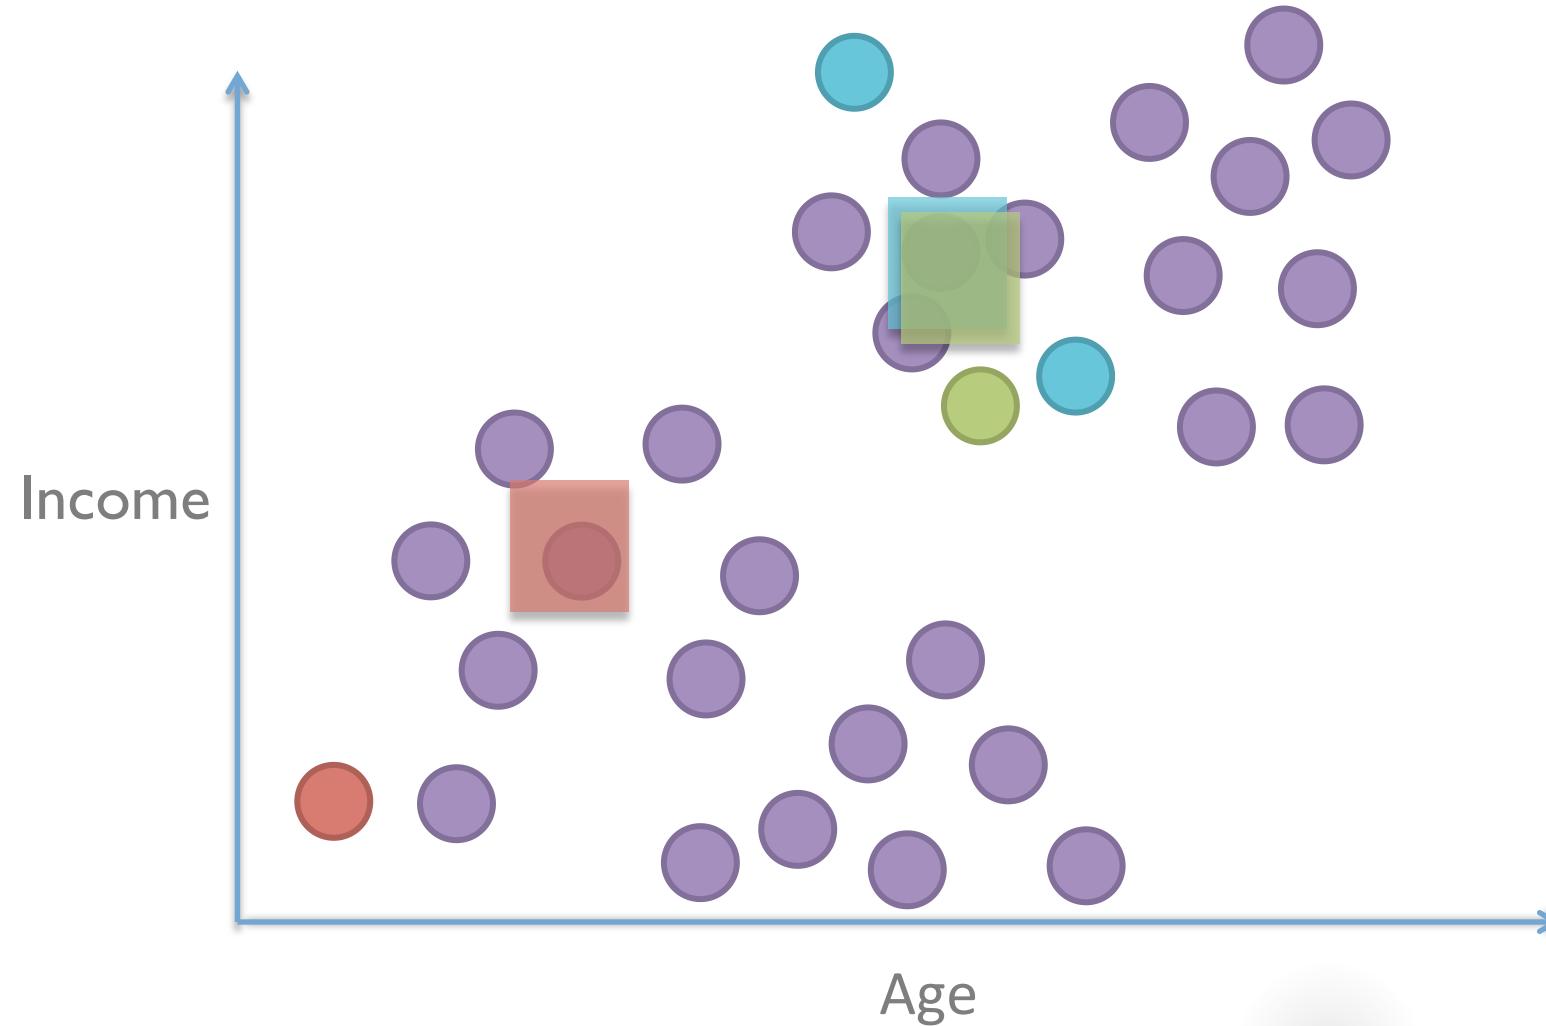
# Mean Shift

Search for local density maximum



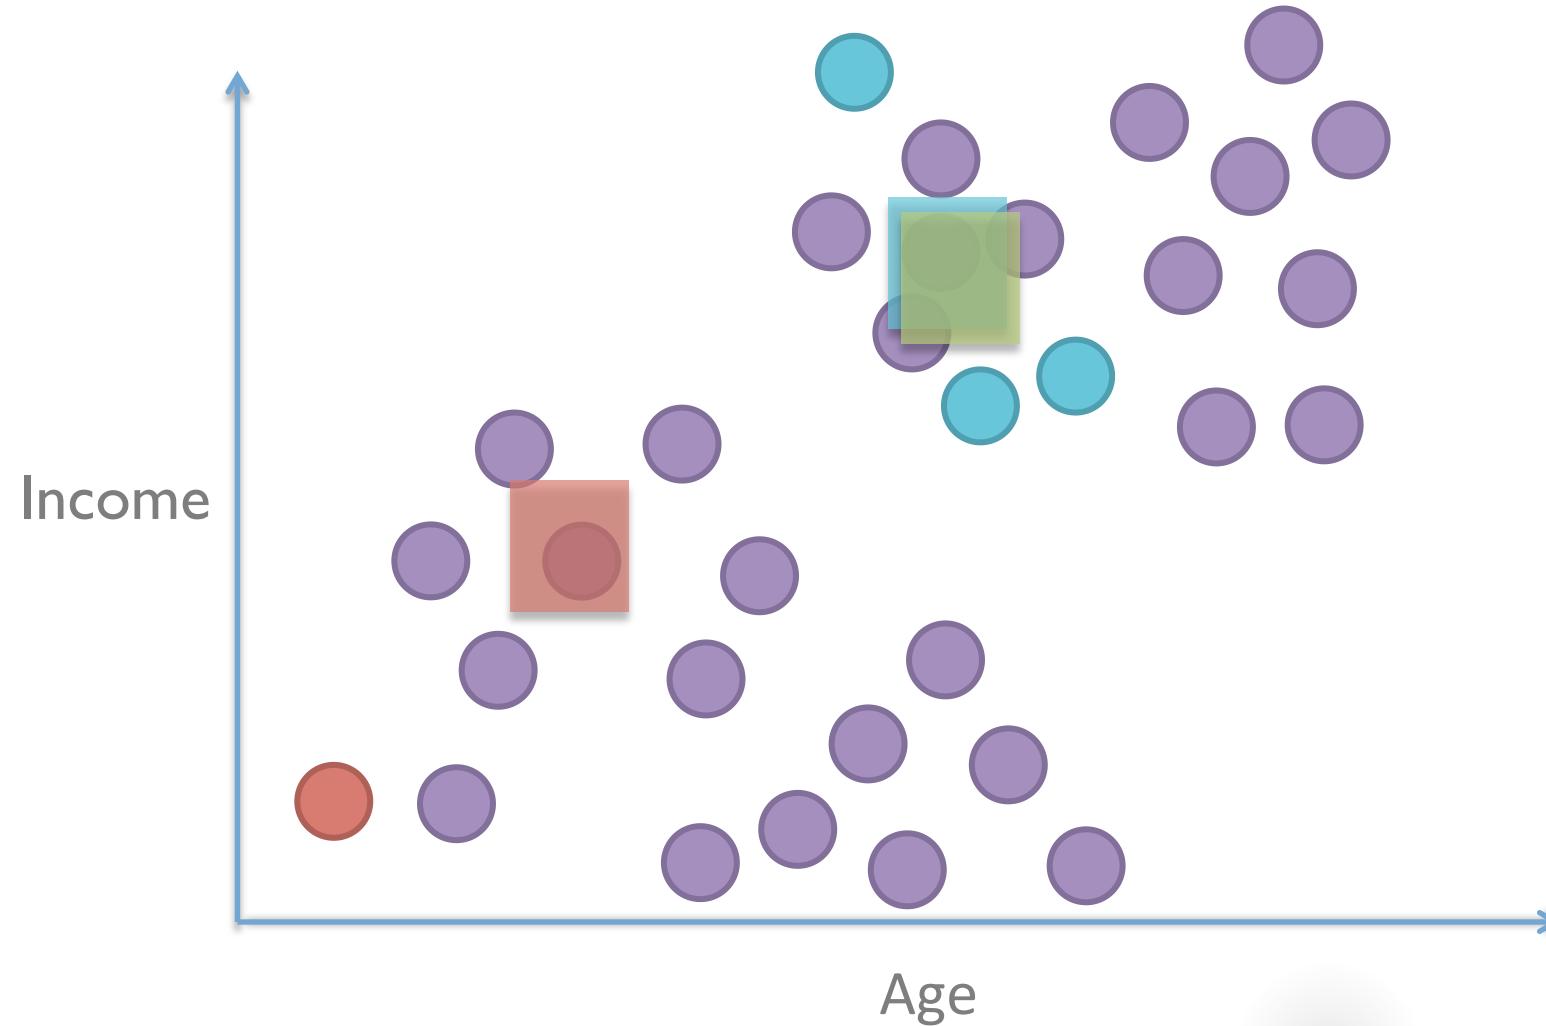
# Mean Shift

Found same maximum, same cluster



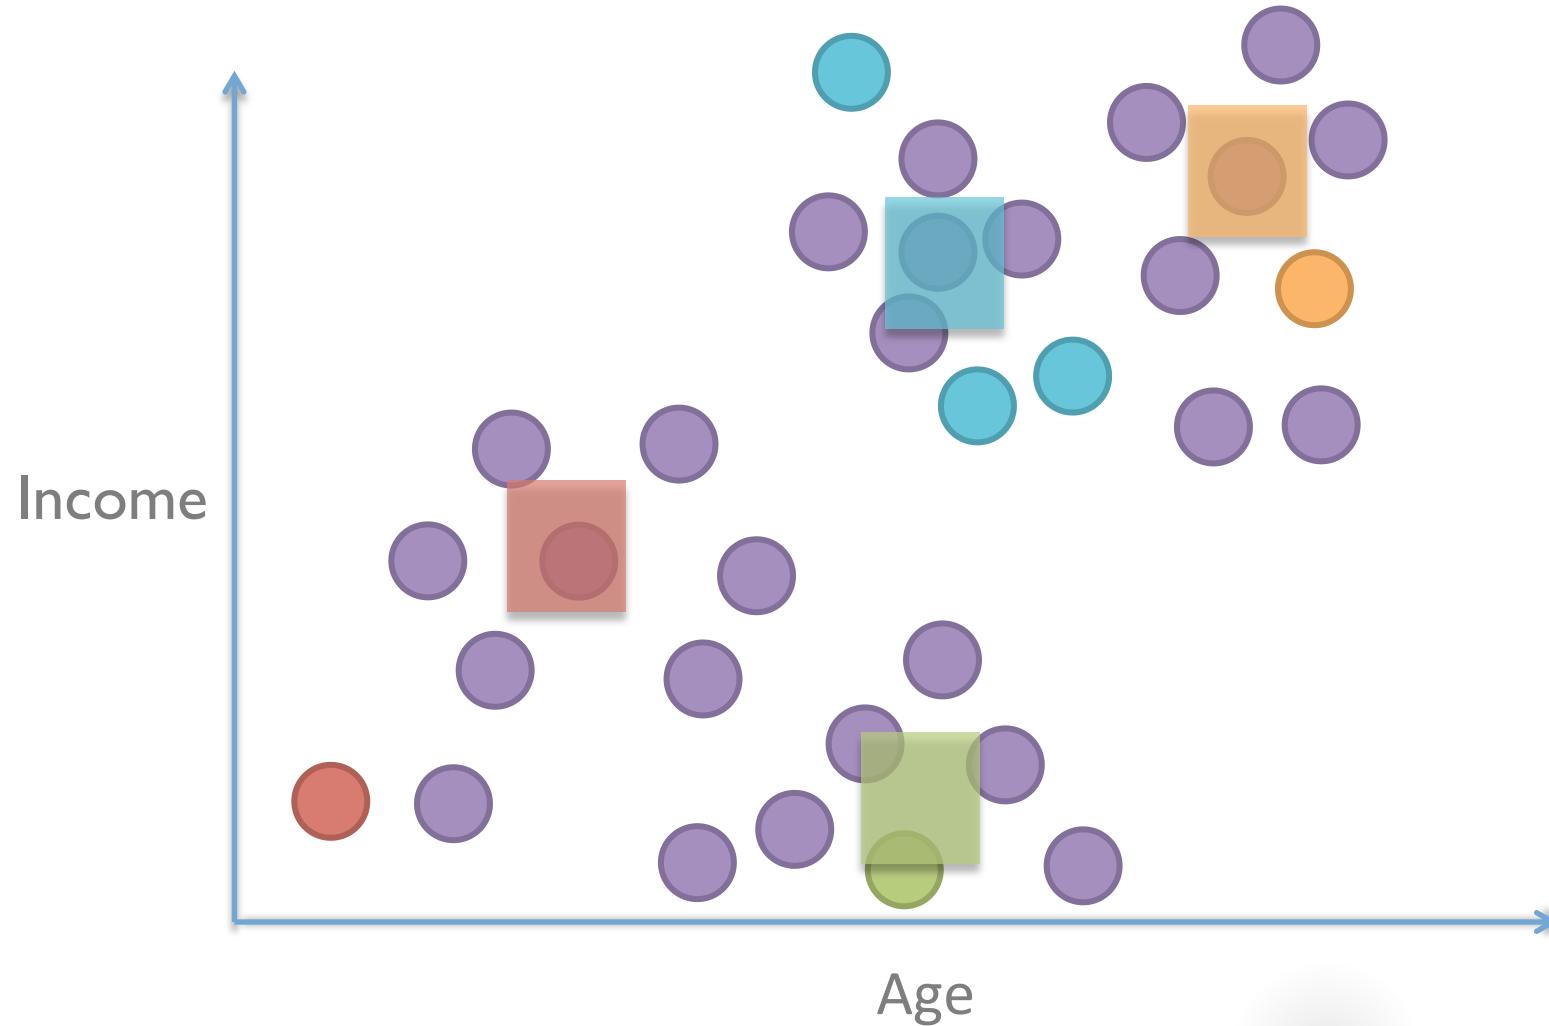
# Mean Shift

Found same maximum, same cluster



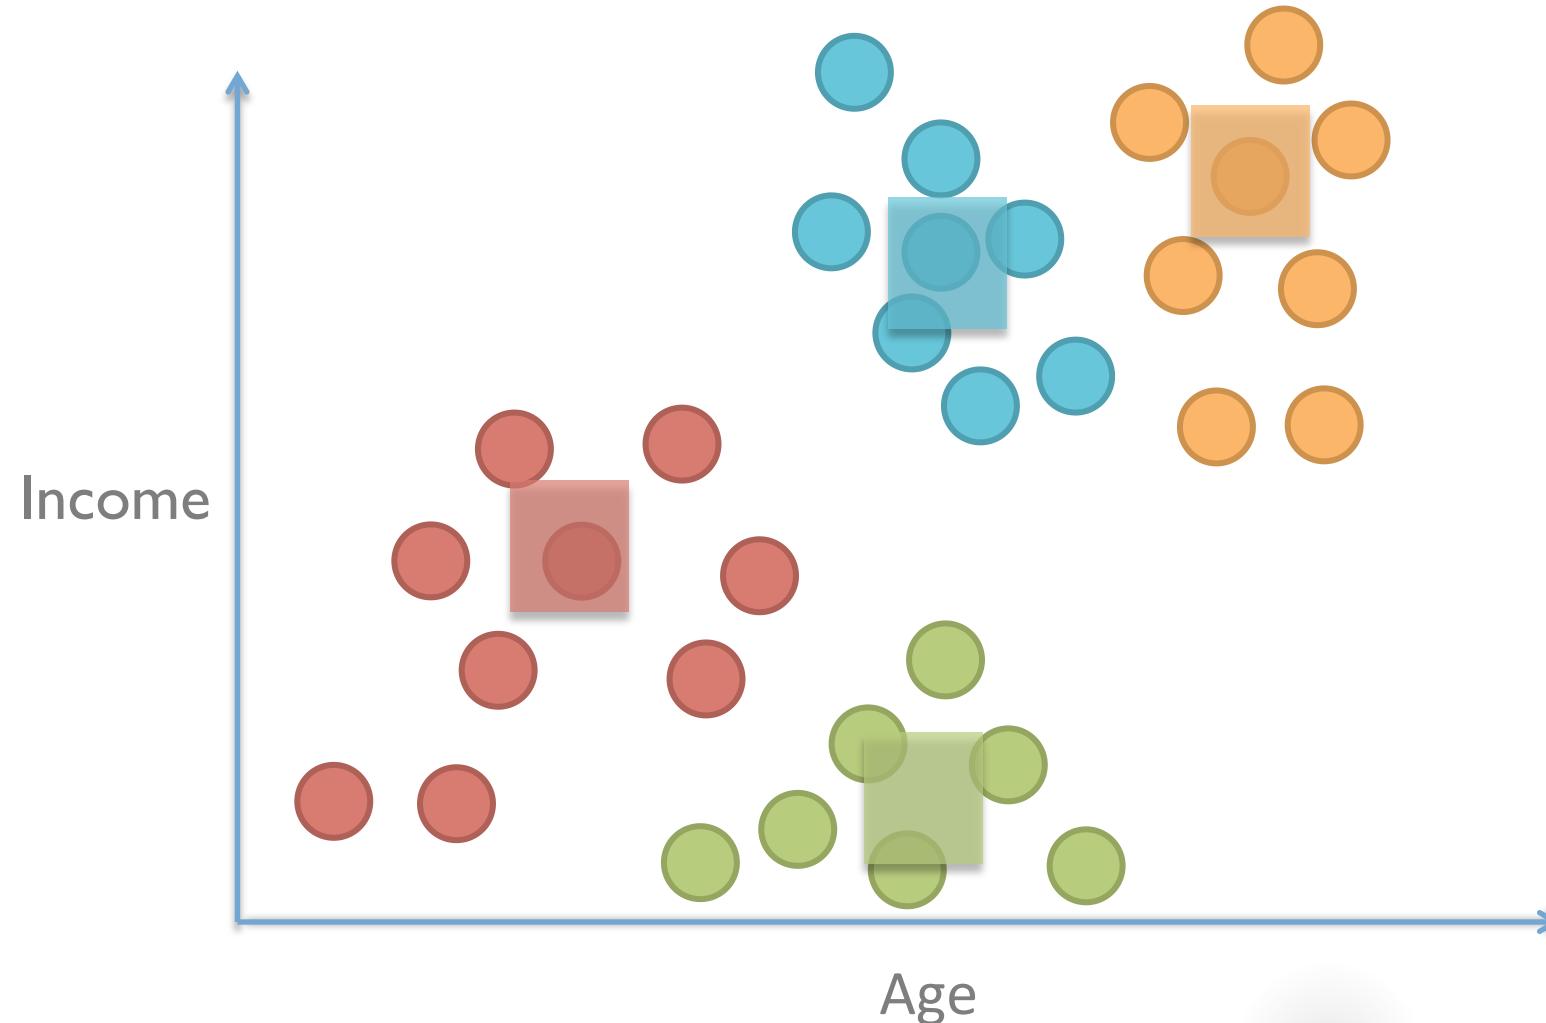
# Mean Shift

Keep going like this, eventually it finds four unique local maxima



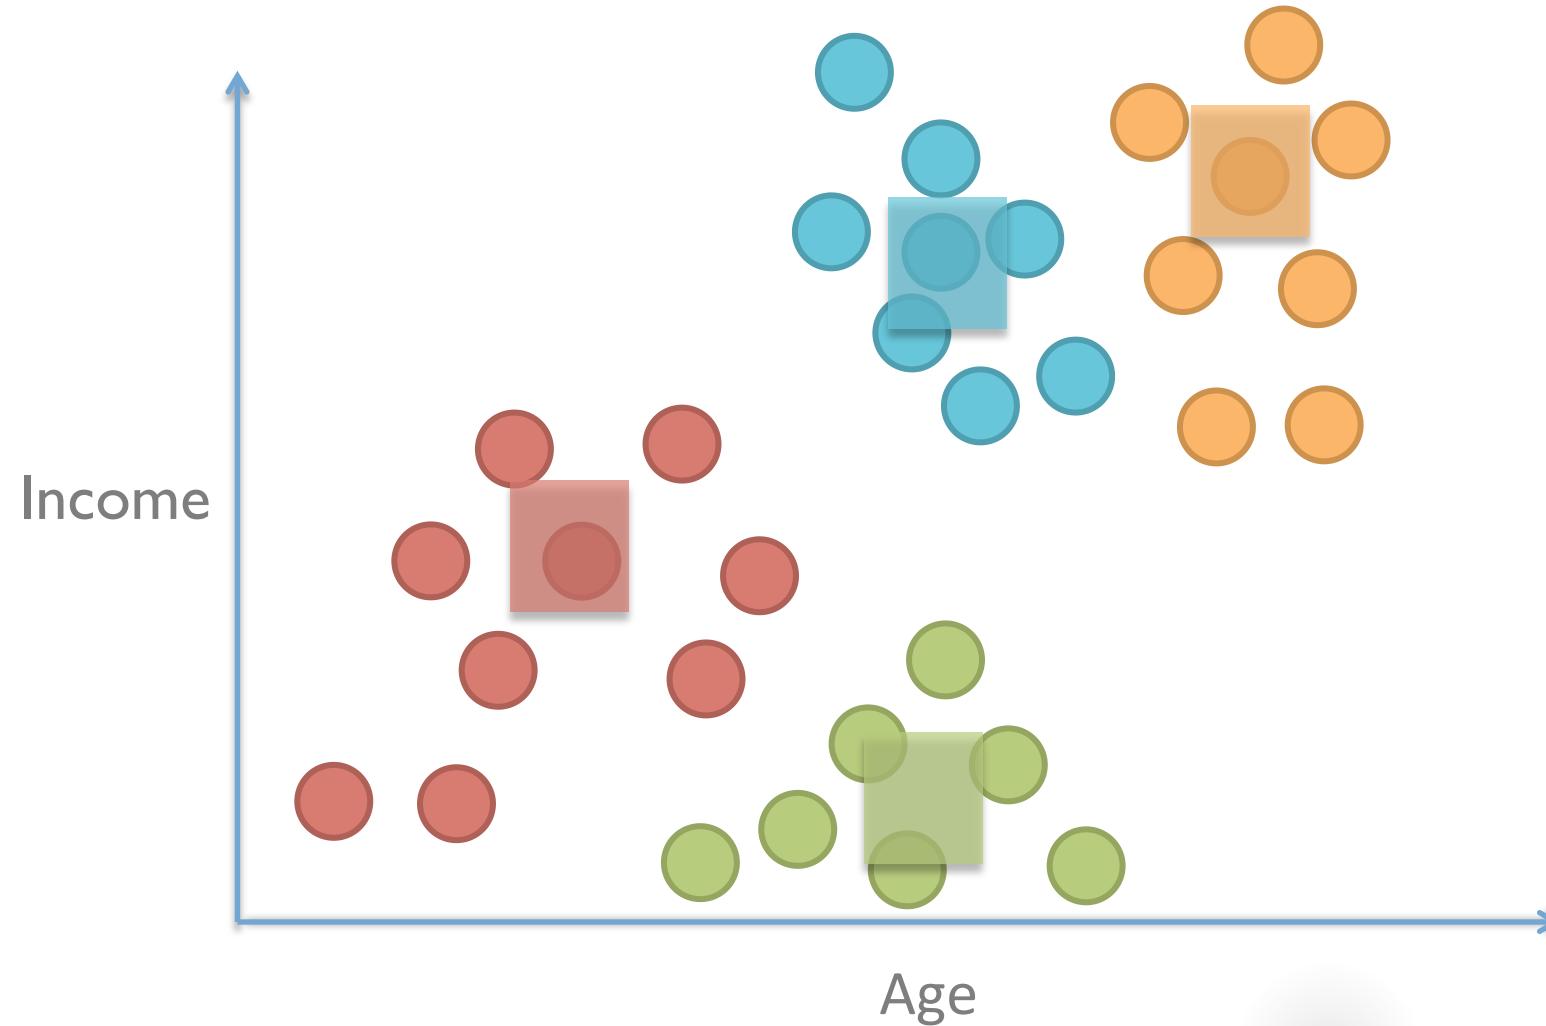
# Mean Shift

Assigns points to centroids they fall to



# Mean Shift

No cluster number or distance parameters



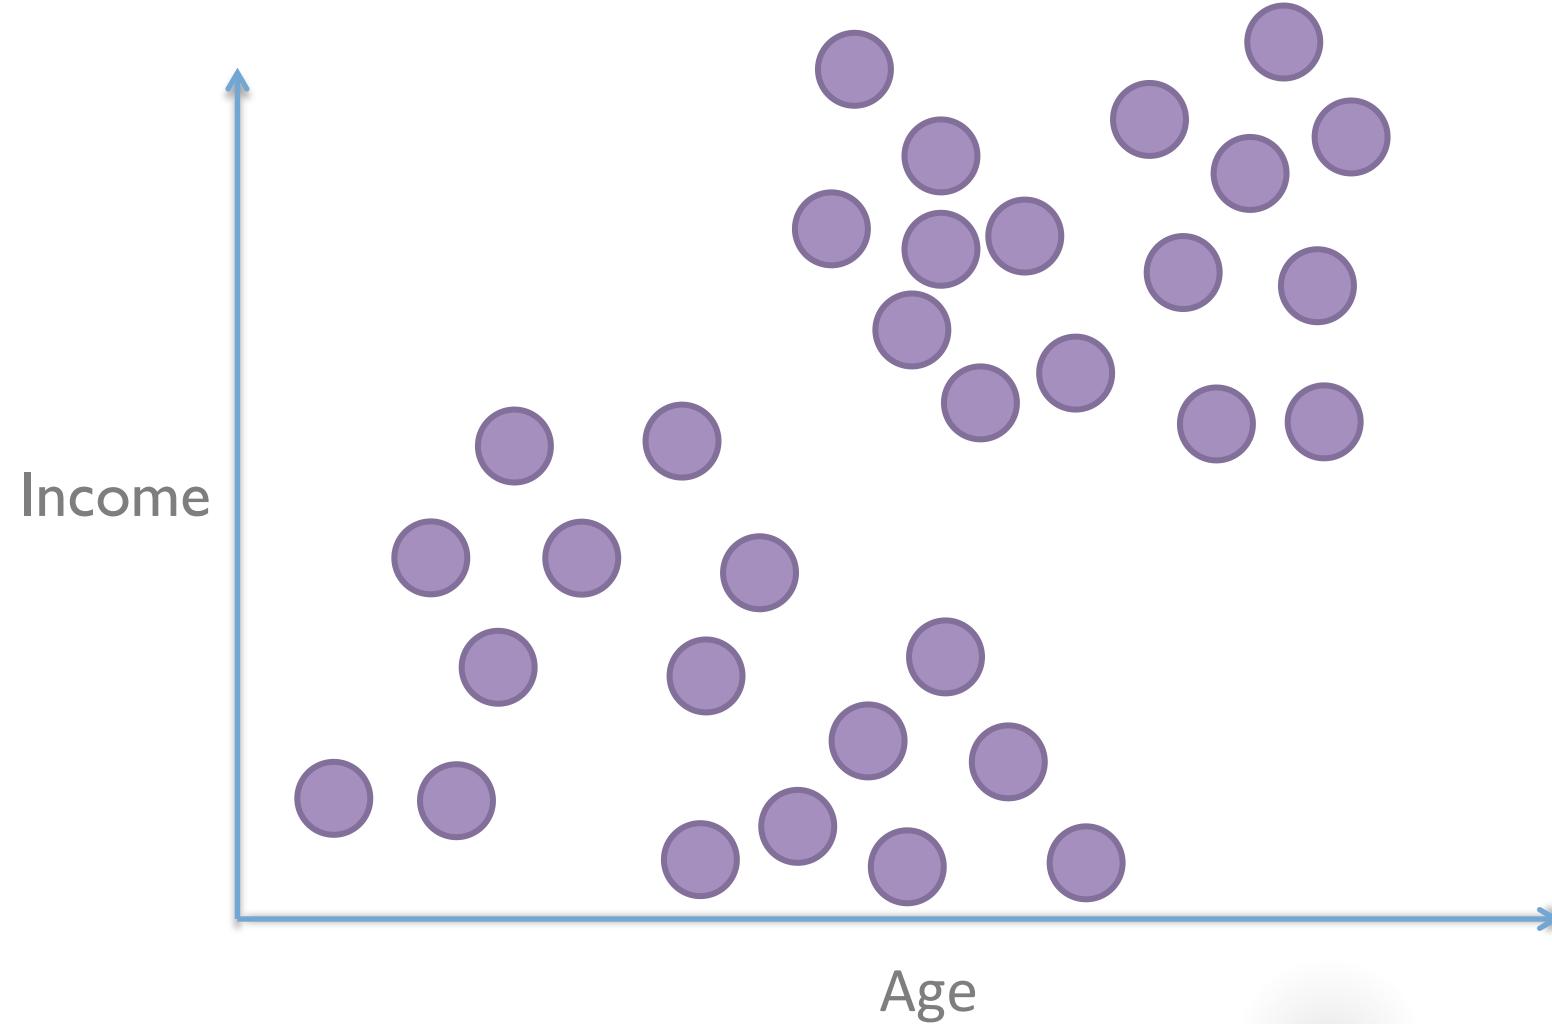
```
from sklearn.cluster import MeanShift  
MeanShift()
```



# Spectral Clustering “kernel k-means”

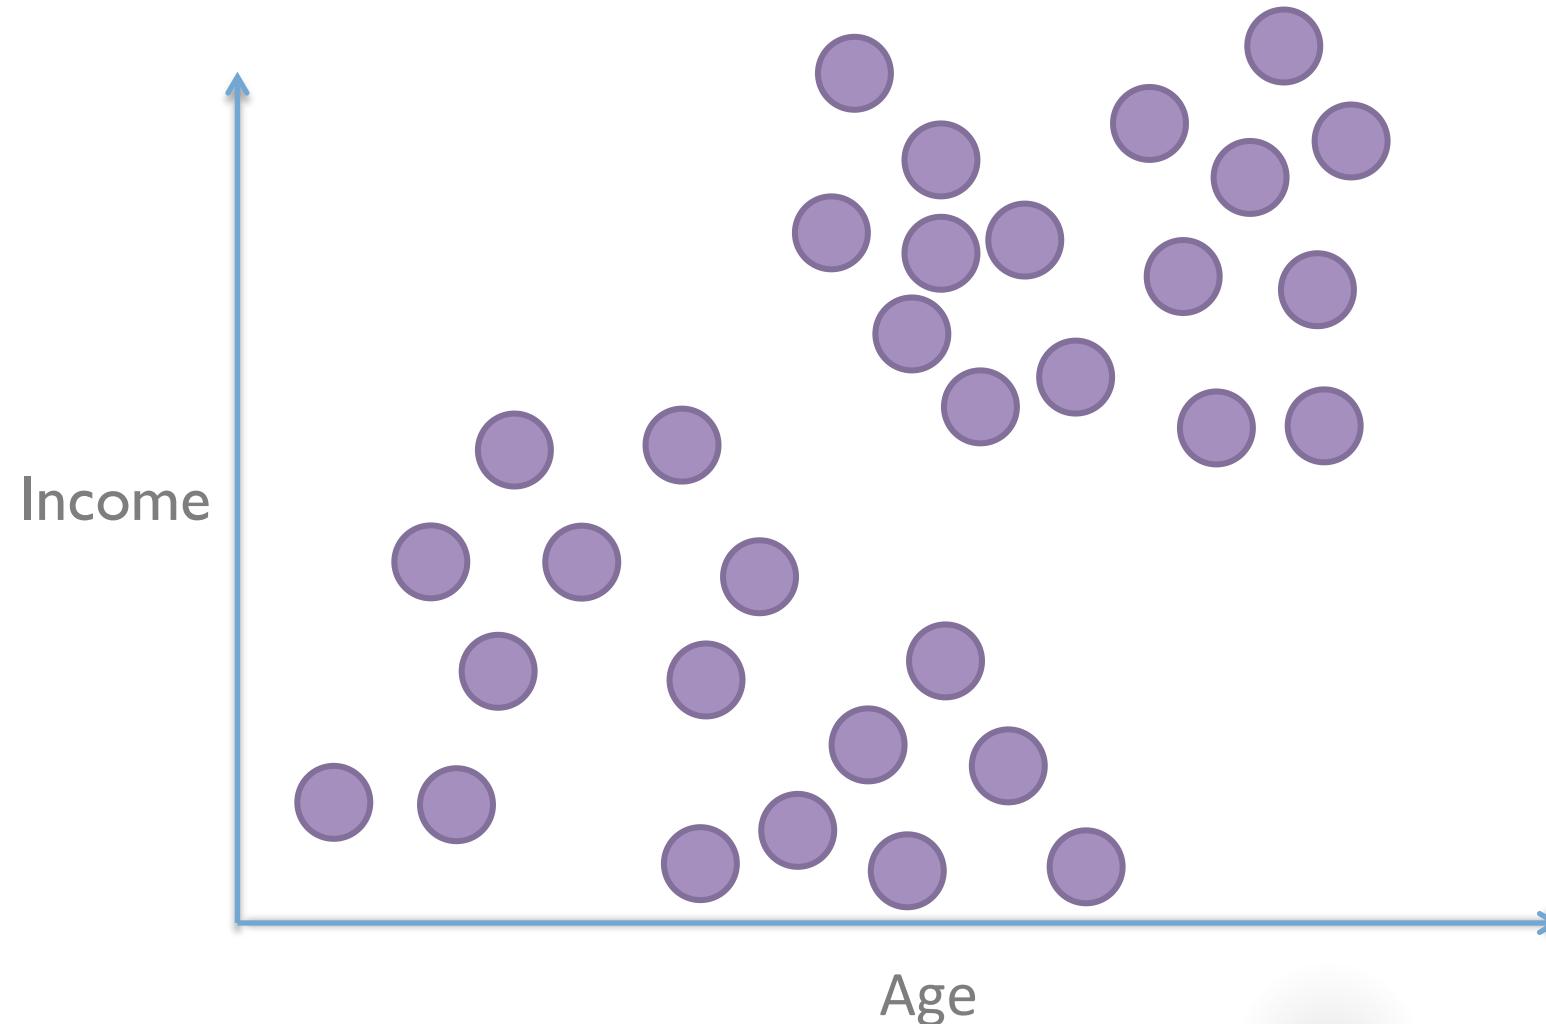


# Spectral Clustering



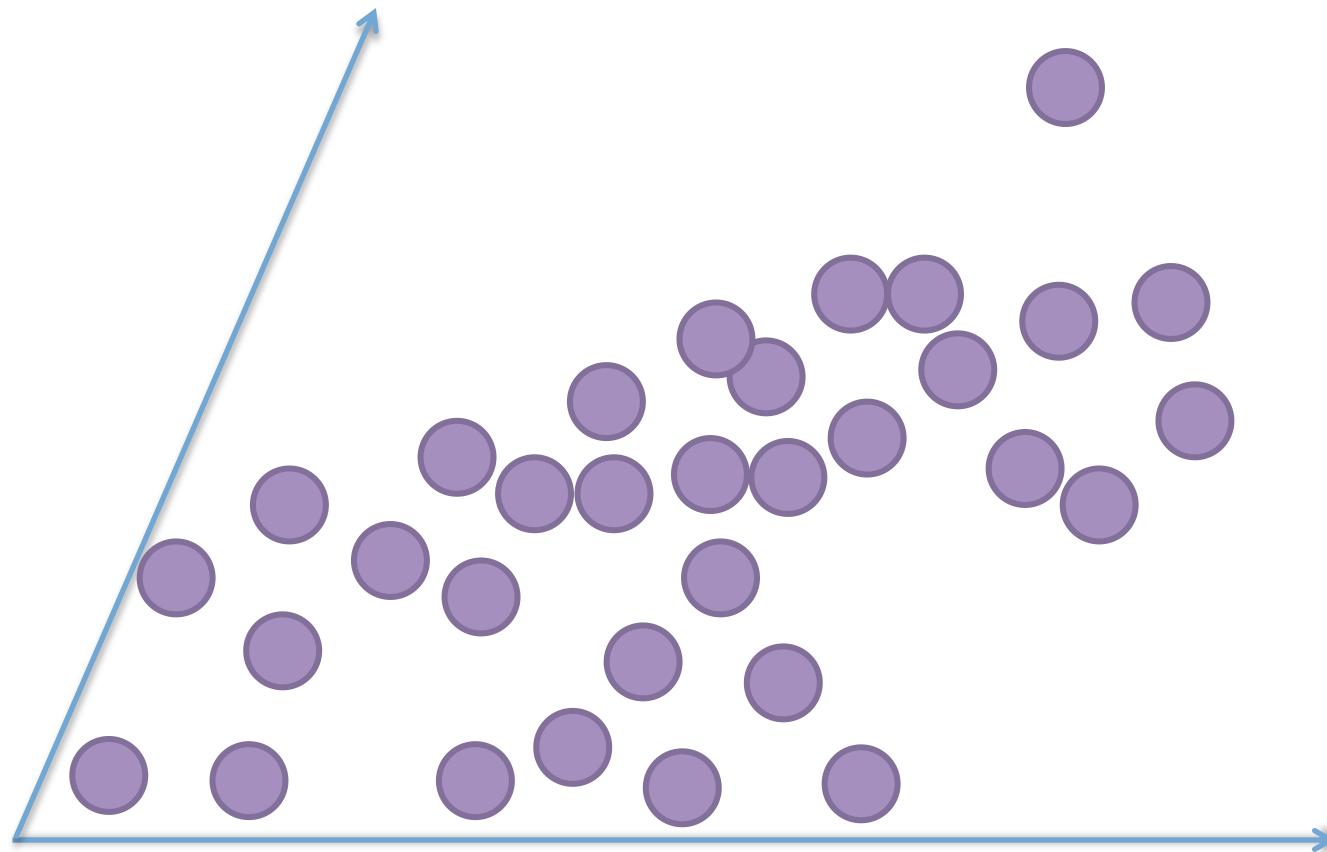
# Spectral Clustering

Transform the space into fewer dimensions by taking eigenvectors of the distances matrix



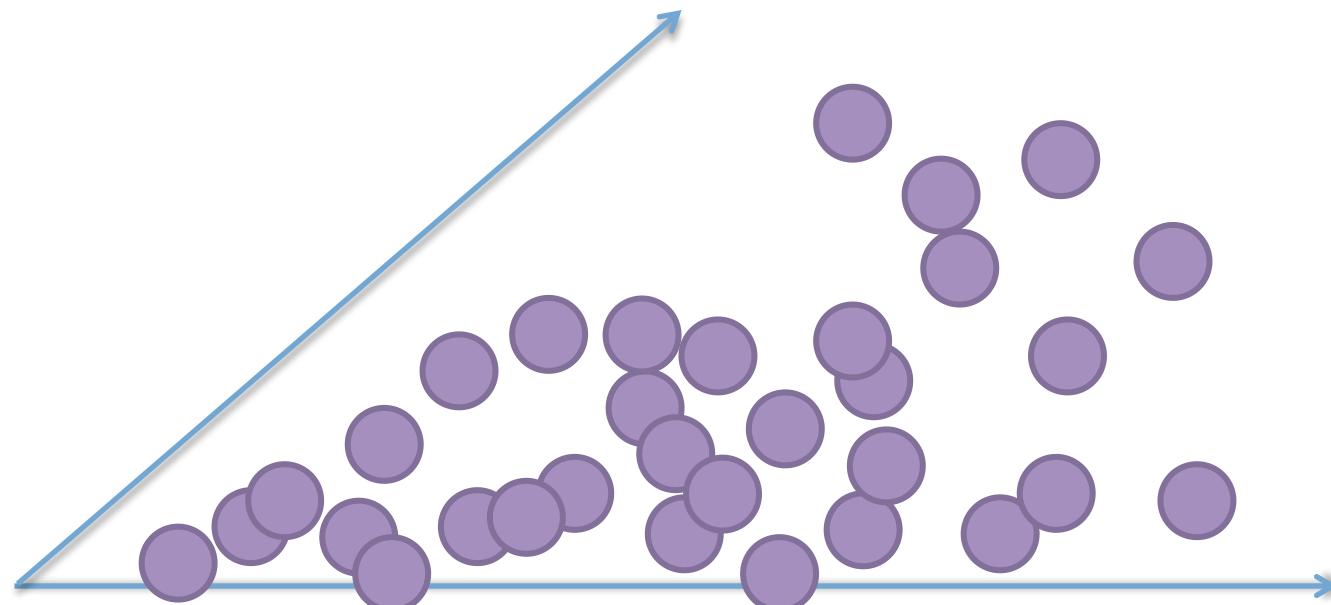
# Spectral Clustering

Transform the space into fewer dimensions by taking eigenvectors of the distances matrix



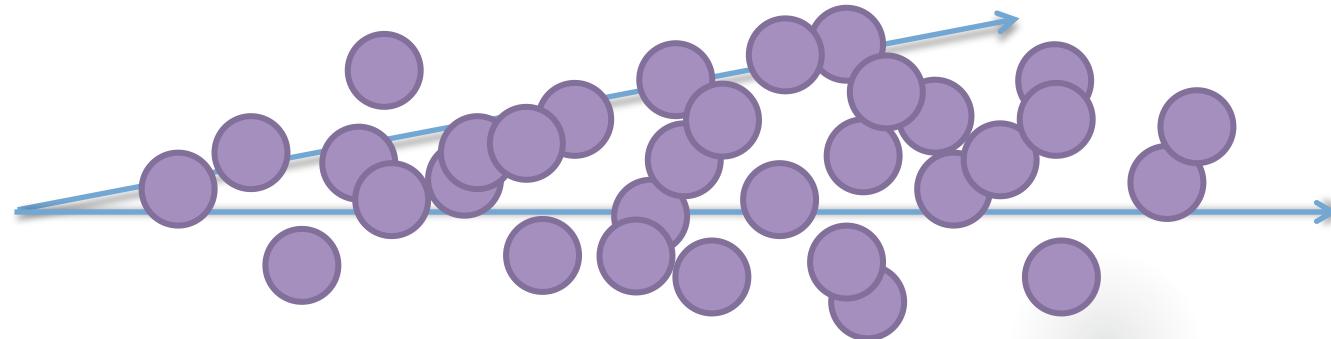
# Spectral Clustering

Transform the space into fewer dimensions by taking eigenvectors of the distances matrix



# Spectral Clustering

Transform the space into fewer dimensions by taking eigenvectors of the distances matrix



# Spectral Clustering

Transform the space into fewer dimensions by taking eigenvectors of the distances matrix



# Spectral Clustering

Now perform K-Means on the low dimension space ( $k=3$ )



# Spectral Clustering

Now perform K-Means on the low dimension space



# Spectral Clustering

Now perform K-Means on the low dimension space



# Spectral Clustering

Now perform K-Means on the low dimension space



# Spectral Clustering

Now perform K-Means on the low dimension space



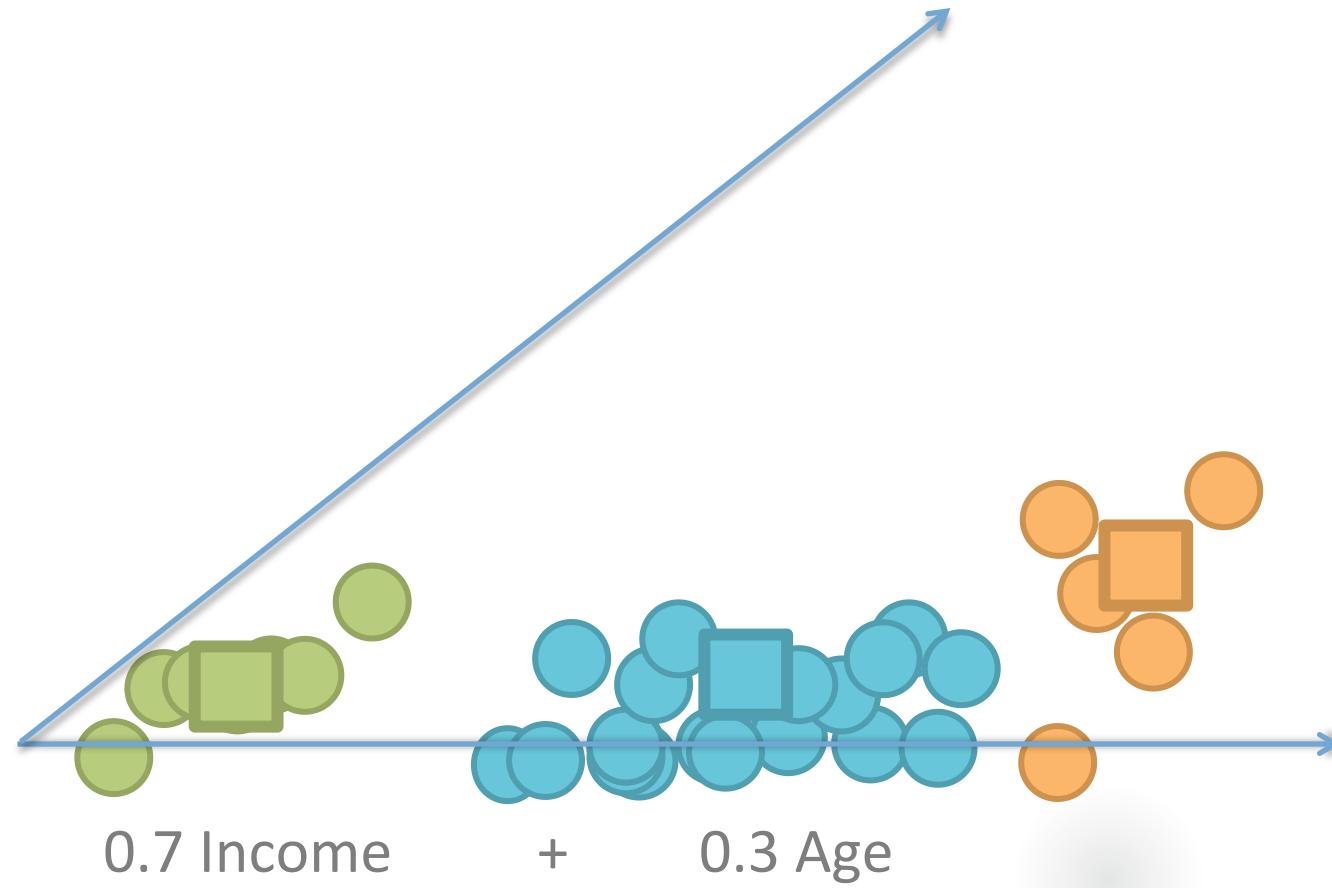
# Spectral Clustering

Look at them in the normal space



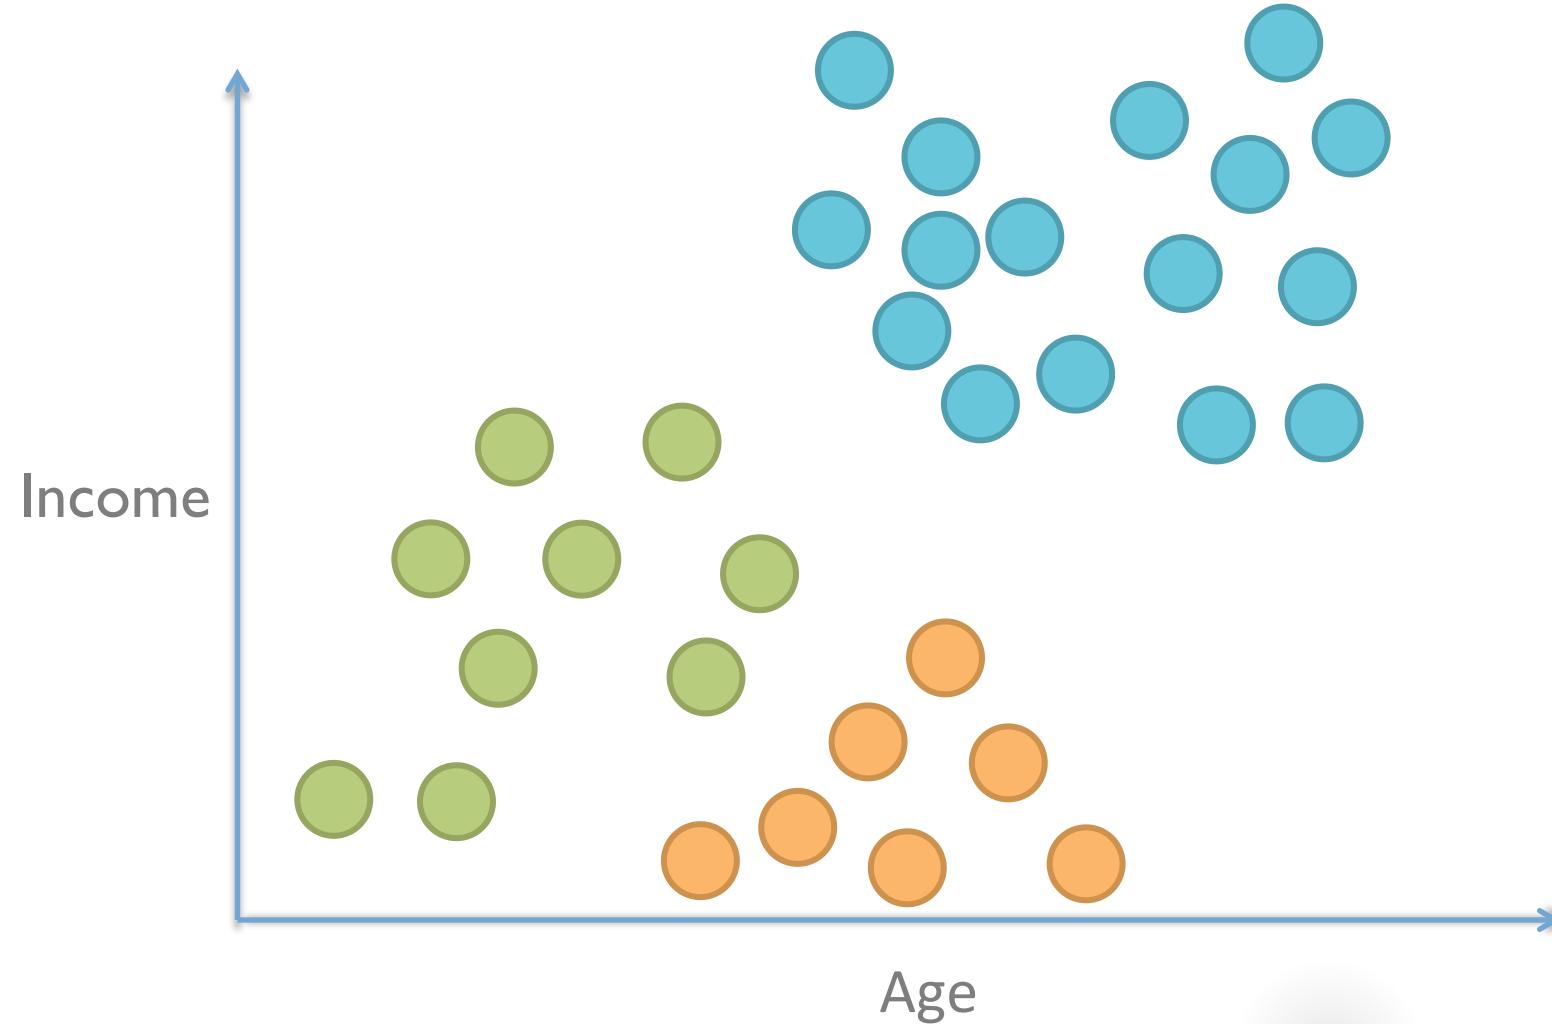
# Spectral Clustering

Look at them in the normal space



# Spectral Clustering

Look at them in the normal space



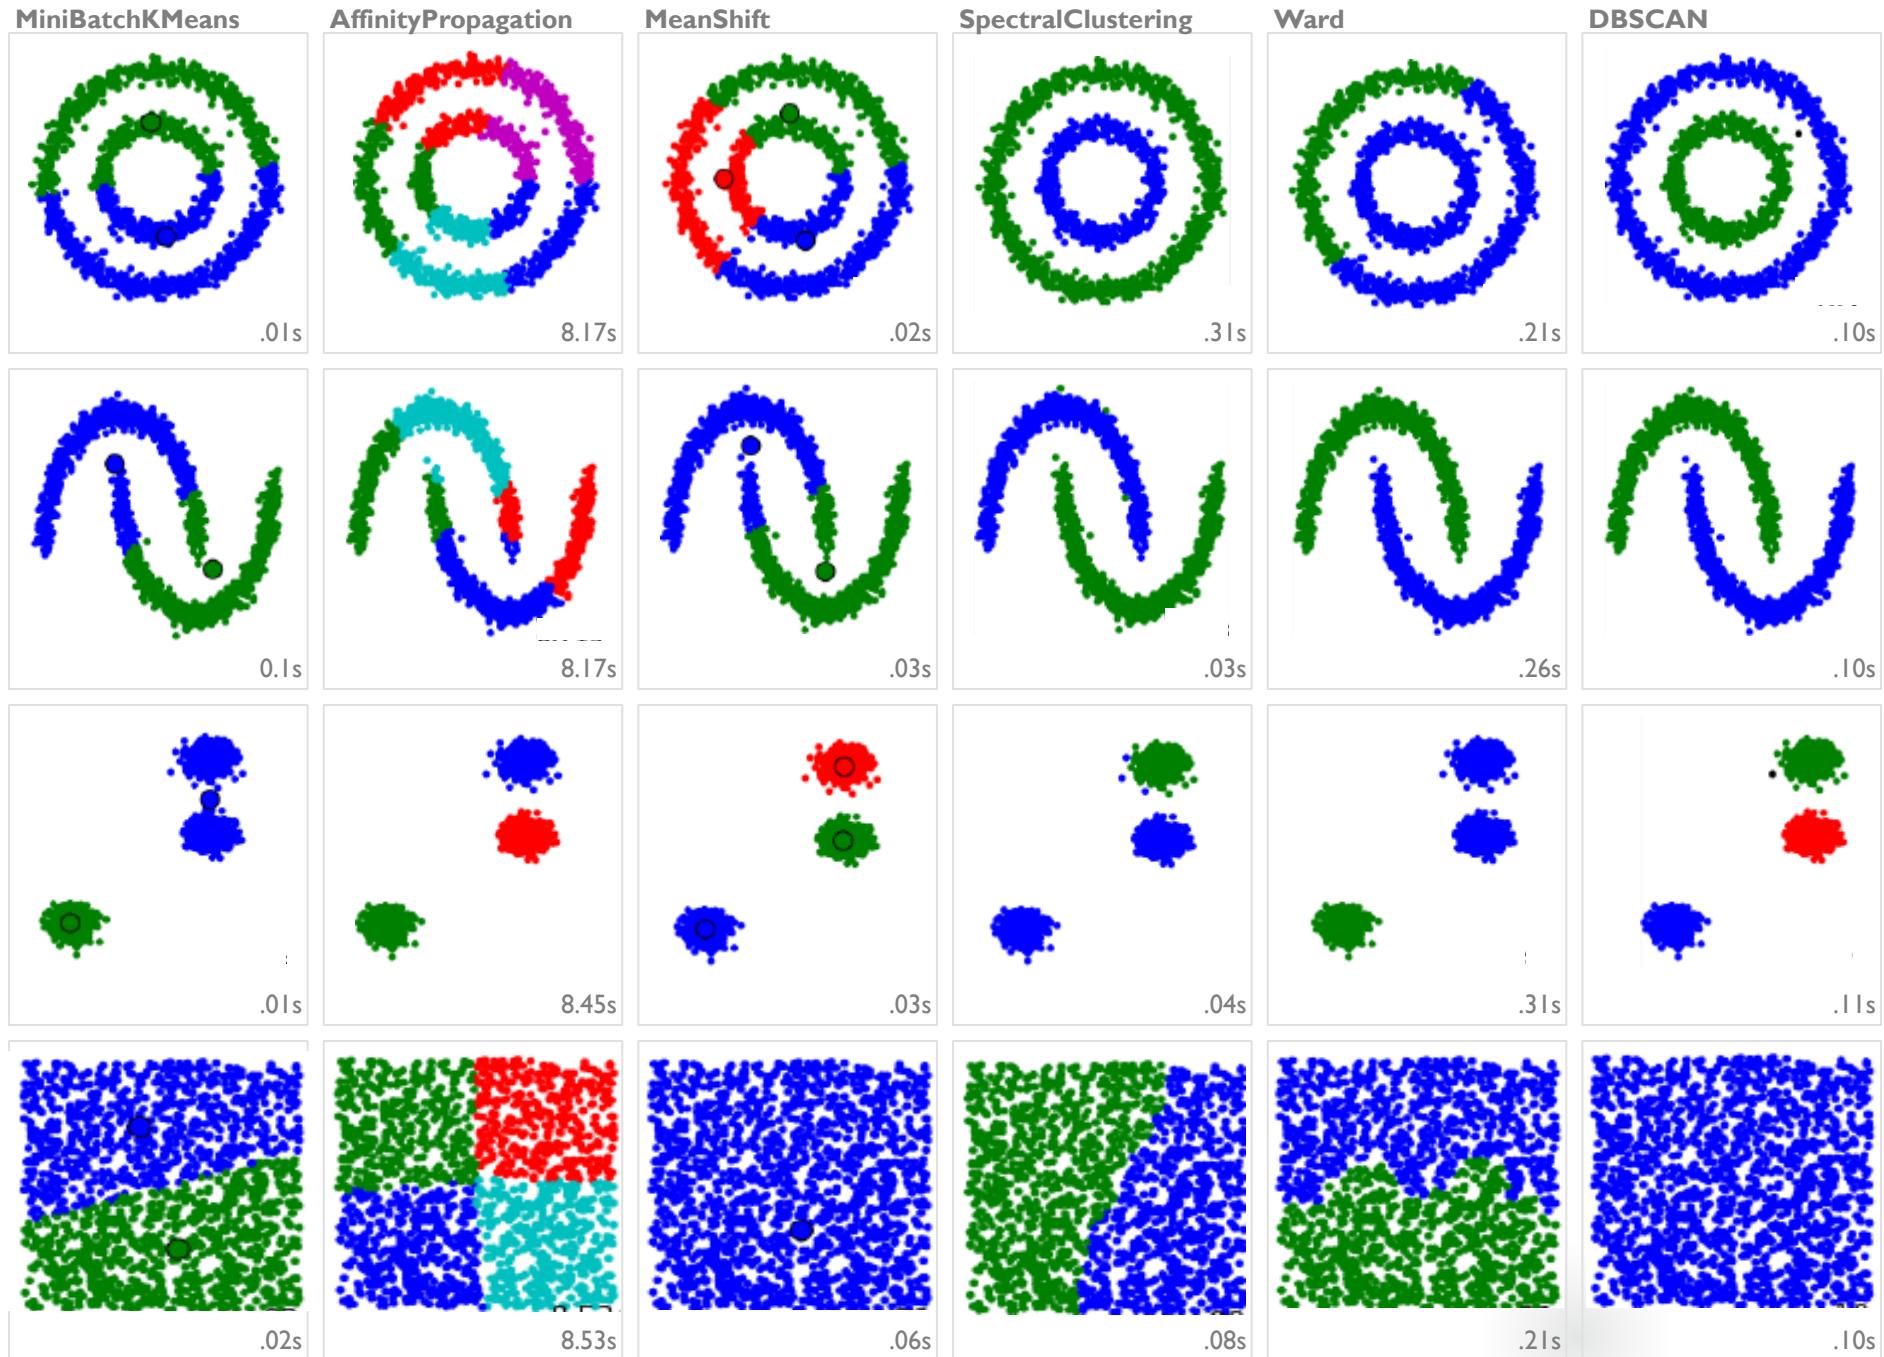
```
from sklearn.cluster import  
SpectralClustering  
  
SpectralClustering( n_clusters=8, n_init=10 )
```





# Properties, Advantages, Disadvantages

---



Method name	Parameters	Scalability	Usecase	Geometry (metric used)
<a href="#">K-means</a>	Number of clusters	Very large $n_{samples}$ medium $n_{clusters}$ with <a href="#">MiniBatch code</a>	General purpose, even cluster size, flat geometry, not too many clusters	Distances between points
<a href="#">Affinity propagation</a>	Damping, sample preferences	Not scalable with $n_{samples}$	Many clusters, uneven cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
<a href="#">Mean-shift</a>	Bandwidth	Not scalable with $n_{samples}$	Many clusters, uneven cluster size, non-flat geometry	Distances between points
<a href="#">Spectral clustering</a>	Number of clusters	Medium $n_{samples}$ , small $n_{clusters}$	Few clusters, even cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
<a href="#">Hierarchical clustering</a>	Number of clusters	Large $n_{samples}$ and $n_{clusters}$	Many clusters, possibly connectivity constraints	Distances between points
<a href="#">DBSCAN</a>	Neighborhood size	Very large $n_{samples}$ , medium $n_{clusters}$	Non-flat geometry, uneven cluster sizes	Distances between nearest points

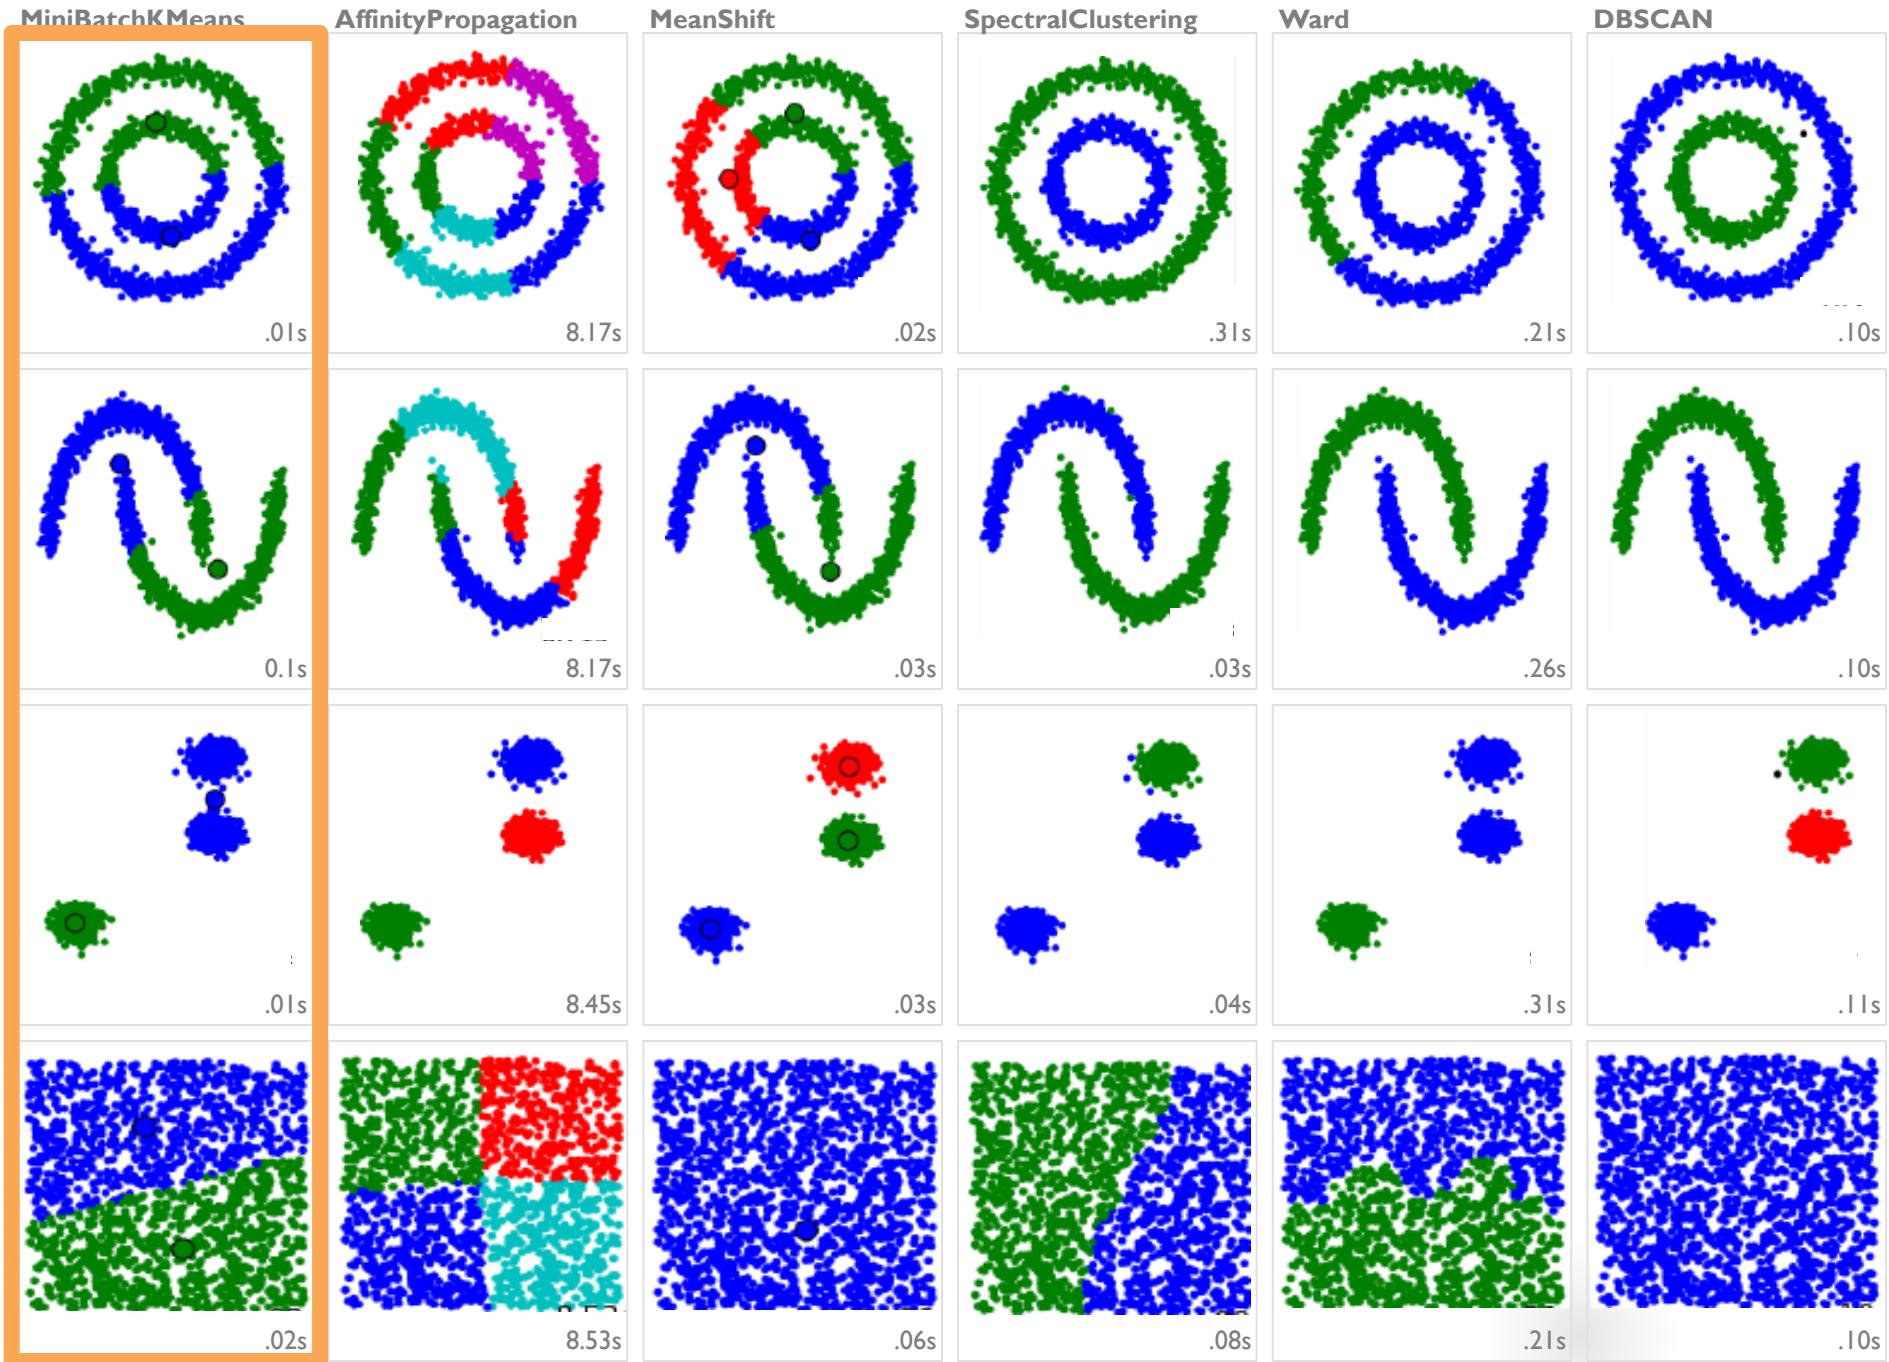


Method name	Parameters	Scalability	Usecase	Geometry (metric used)
K-means	Number of clusters	Very large $n_{samples}$ medium $n_{clusters}$ with <a href="#">MiniBatch code</a>	General purpose, even cluster size, flat geometry, not too many clusters	Distances between points
Affinity propagation	Damping, sample preferences	Not scalable with $n_{samples}$	Many clusters, uneven cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
Mean-shift	Bandwidth	Not scalable with $n_{samples}$	Many clusters, uneven cluster size, non-flat geometry	Distances between points
Spectral clustering	Number of clusters	Medium $n_{samples}$ , small $n_{clusters}$	Few clusters, even cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
Hierarchical clustering	Number of clusters	Large $n_{samples}$ and $n_{clusters}$	Many clusters, possibly connectivity constraints	Distances between points
DBSCAN	Neighborhood size	Very large $n_{samples}$ , medium $n_{clusters}$	Non-flat geometry, uneven cluster sizes	Distances between nearest points

## K Means:

- MiniBatch version is super fast (big data).
- Have to try k values (k not too big).
- Tends to find even sized.
- Bad with non-spherical cluster shapes.



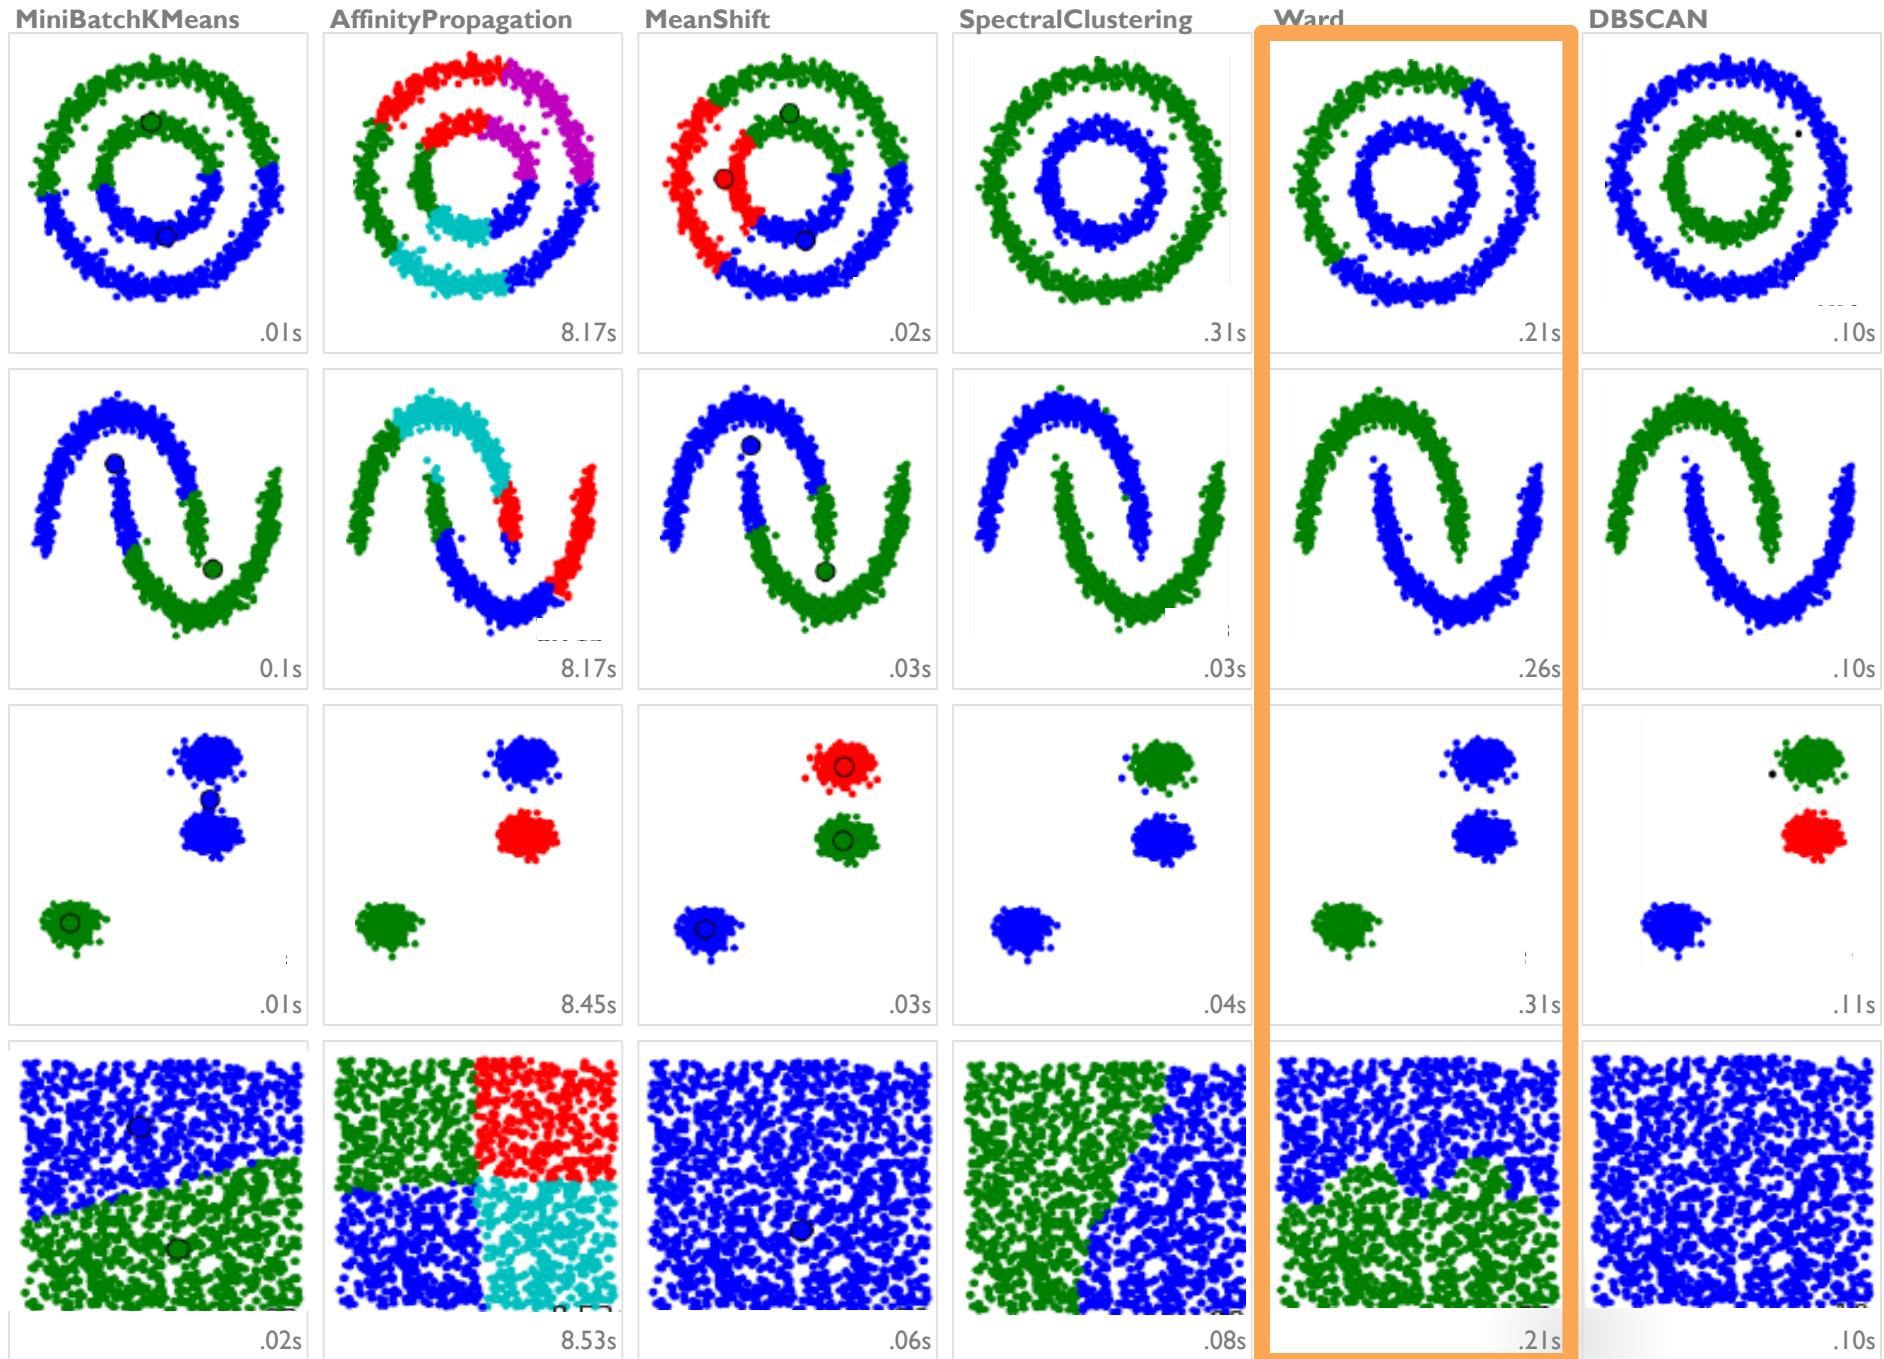


Method name	Parameters	Scalability	Usecase	Geometry (metric used)
K-means	Number of clusters	Very large $n_{samples}$ medium $n_{clusters}$ with <a href="#">MiniBatch code</a>	General purpose, even cluster size, flat geometry, not too many clusters	Distances between points
Affinity propagation	Damping, sample preferences	Not scalable with $n_{samples}$	Many clusters, uneven cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
Mean-shift	Bandwidth	Not scalable with $n_{samples}$	Many clusters, uneven cluster size, non-flat geometry	Distances between points
Spectral clustering	Number of clusters	Medium $n_{samples}$ , small $n_{clusters}$	Few clusters, even cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
Hierarchical clustering	Number of clusters	Large $n_{samples}$ and $n_{clusters}$	Many clusters, possibly connectivity constraints	Distances between points
DBSCAN	Neighborhood size	Very large $n_{samples}$ , medium $n_{clusters}$	Non-flat geometry, uneven cluster sizes	Distances between nearest points

## Hierarchical Clustering:

- You get a full **hierarchy tree**. Useful for some problems.
- Have to try k values.
- Finds **uneven cluster sizes** (one is big, some are tiny, etc.). Due to the “the chains get chainier effect.”
- A lot of **distance metric** and linkage **options**.
- Slow:  $O(n^2)$ .





Method name	Parameters	Scalability	Usecase	Geometry (metric used)
K-means	Number of clusters	Very large $n_{samples}$ medium $n_{clusters}$ with <a href="#">MiniBatch code</a>	General purpose, even cluster size, flat geometry, not too many clusters	Distances between points
Affinity propagation	Damping, sample preferences	Not scalable with $n_{samples}$	Many clusters, uneven cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
Mean-shift	Bandwidth	Not scalable with $n_{samples}$	Many clusters, uneven cluster size, non-flat geometry	Distances between points
Spectral clustering	Number of clusters	Medium $n_{samples}$ , small $n_{clusters}$	Few clusters, even cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
Hierarchical clustering	Number of clusters	Large $n_{samples}$ and $n_{clusters}$	Many clusters, possibly connectivity constraints	Distances between points
DBSCAN	Neighborhood size	Very large $n_{samples}$ , medium $n_{clusters}$	Non-flat geometry, uneven cluster sizes	Distances between nearest points

## DBSCAN

Density based, on the money with the right parameters.  
Have to try  $\text{epsilon}$  (and  $\text{num\_clu}$ ) values.

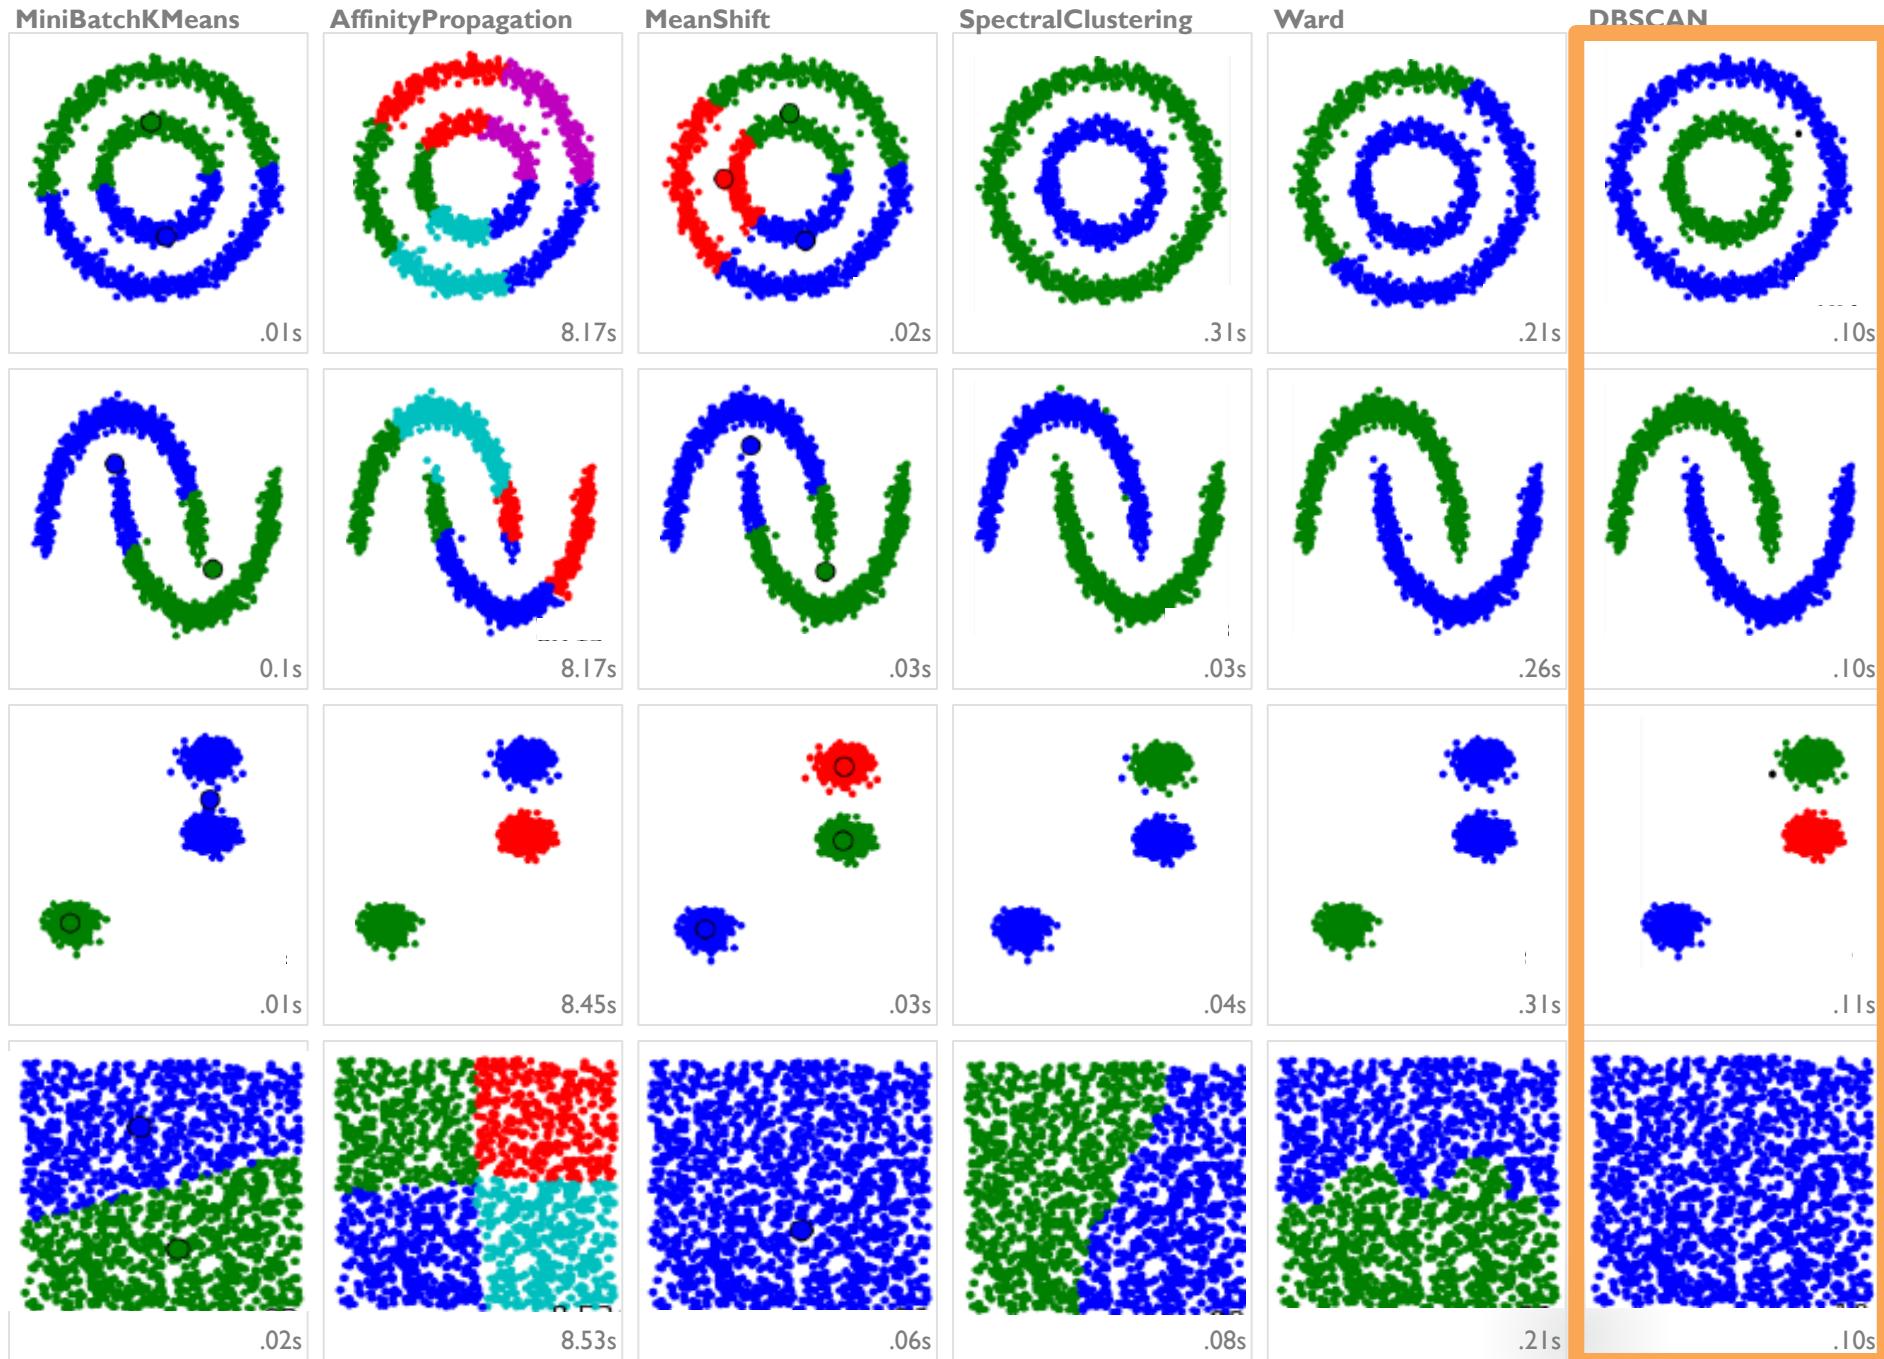
Can find **uneven cluster sizes**.

Full **distance metric options**.

Can handle tons of data and weird shapes.

Too small  $\text{epsilon}$  (too many clusters) is a bit shaky.



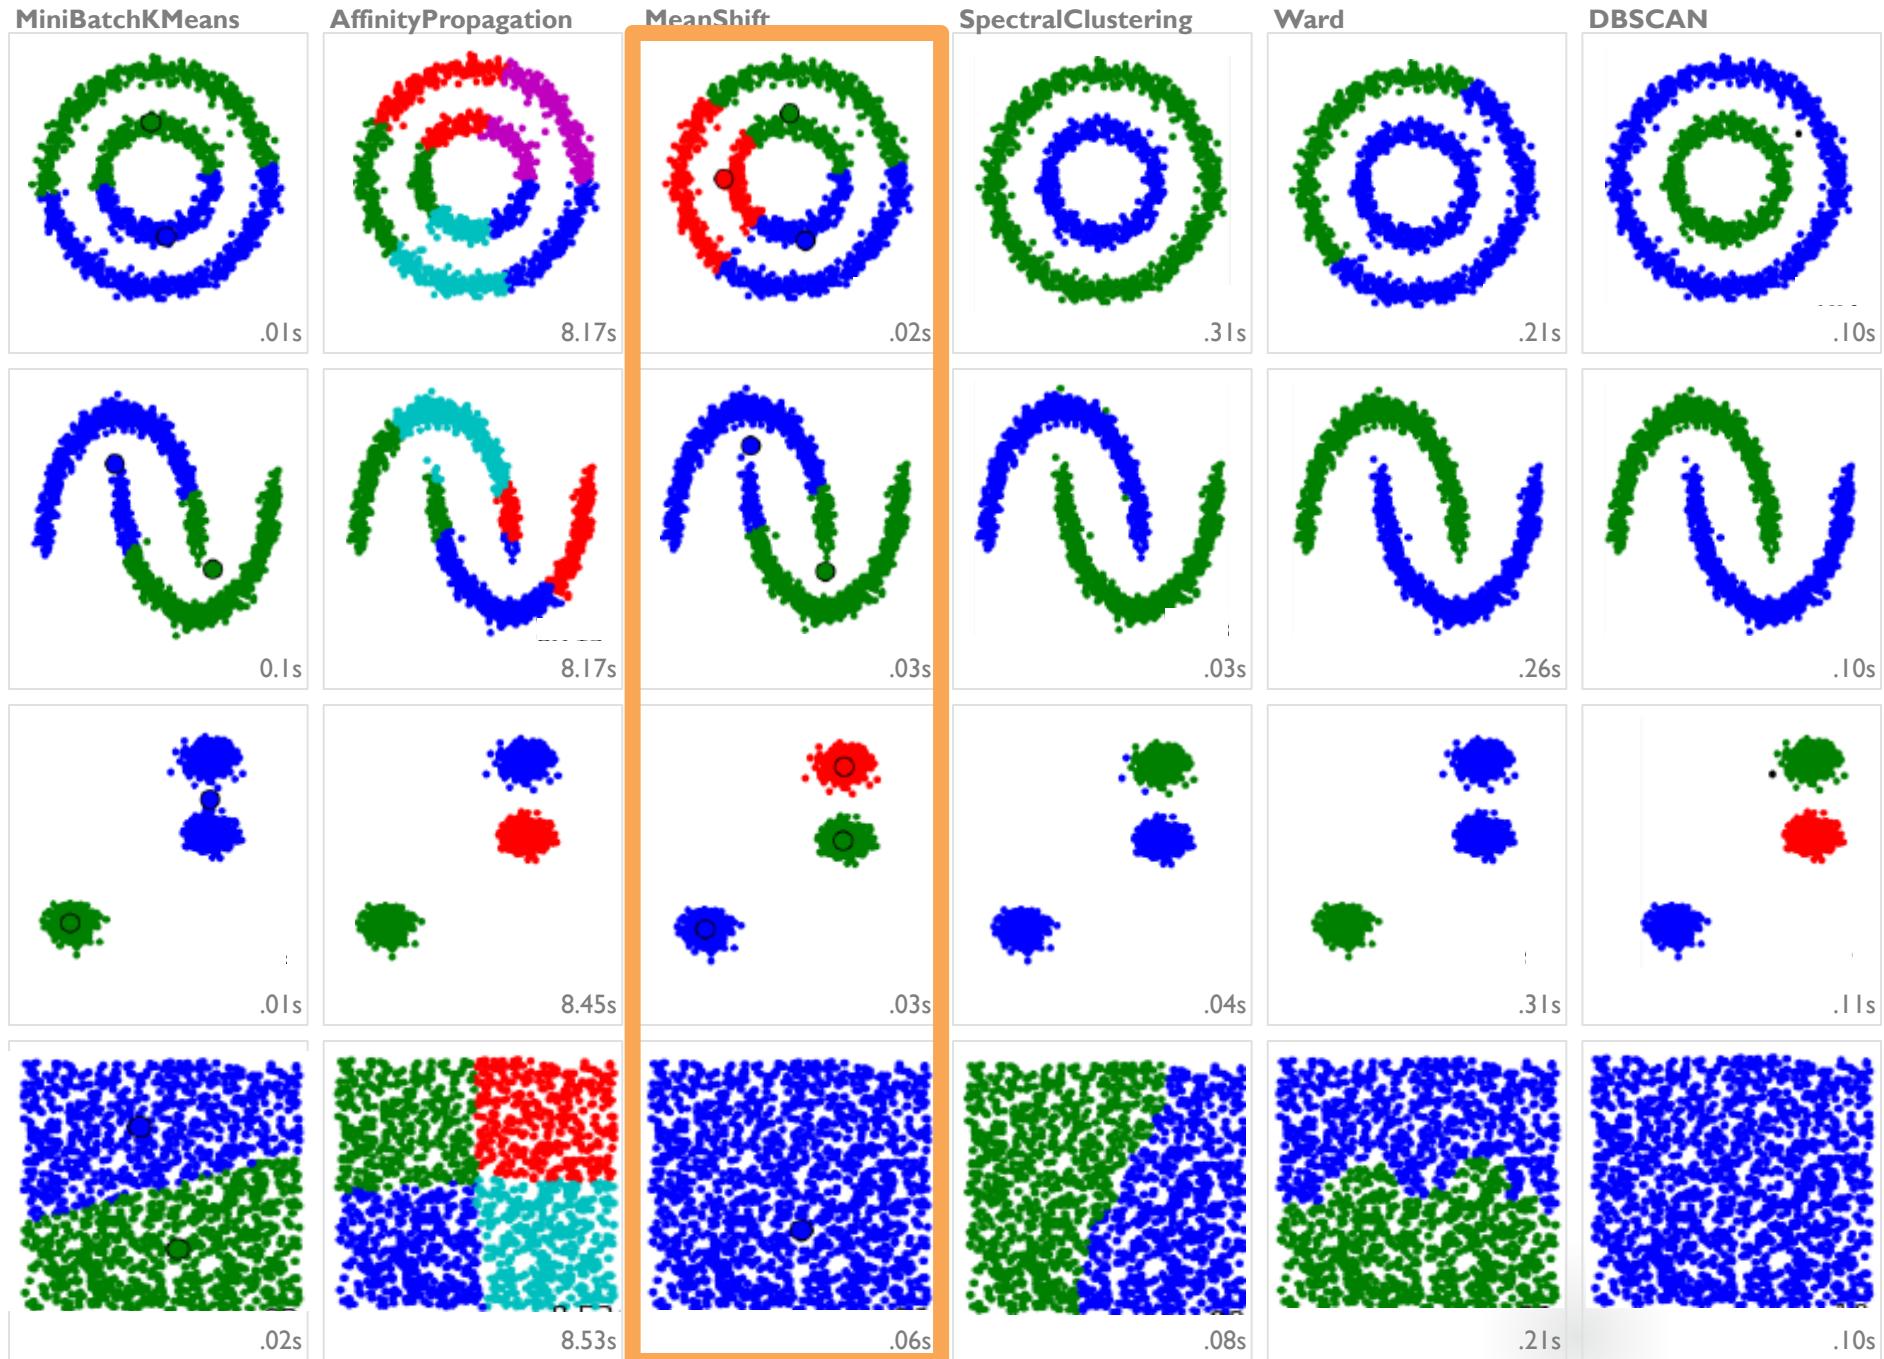


Method name	Parameters	Scalability	Usecase	Geometry (metric used)
K-means	Number of clusters	Very large $n_{samples}$ medium $n_{clusters}$ with <a href="#">MiniBatch code</a>	General purpose, even cluster size, flat geometry, not too many clusters	Distances between points
Affinity propagation	Damping, sample preferences	Not scalable with $n_{samples}$	Many clusters, uneven cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
Mean-shift	Bandwidth	Not scalable with $n_{samples}$	Many clusters, uneven cluster size, non-flat geometry	Distances between points
Spectral clustering	Number of clusters	Medium $n_{samples}$ , small $n_{clusters}$	Few clusters, even cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
Hierarchical clustering	Number of clusters	Large $n_{samples}$ and $n_{clusters}$	Many clusters, possibly connectivity constraints	Distances between points
DBSCAN	Neighborhood size	Very large $n_{samples}$ , medium $n_{clusters}$	Non-flat geometry, uneven cluster sizes	Distances between nearest points

## Mean-shift

- You DON'T have to guess k or other parameters. Fire & forget.
  - Can find uneven cluster sizes.
  - Slow with a lot of data, but lots of clusters no problem
  - Doesn't handle weird shapes well.
- Euclidean distance only.



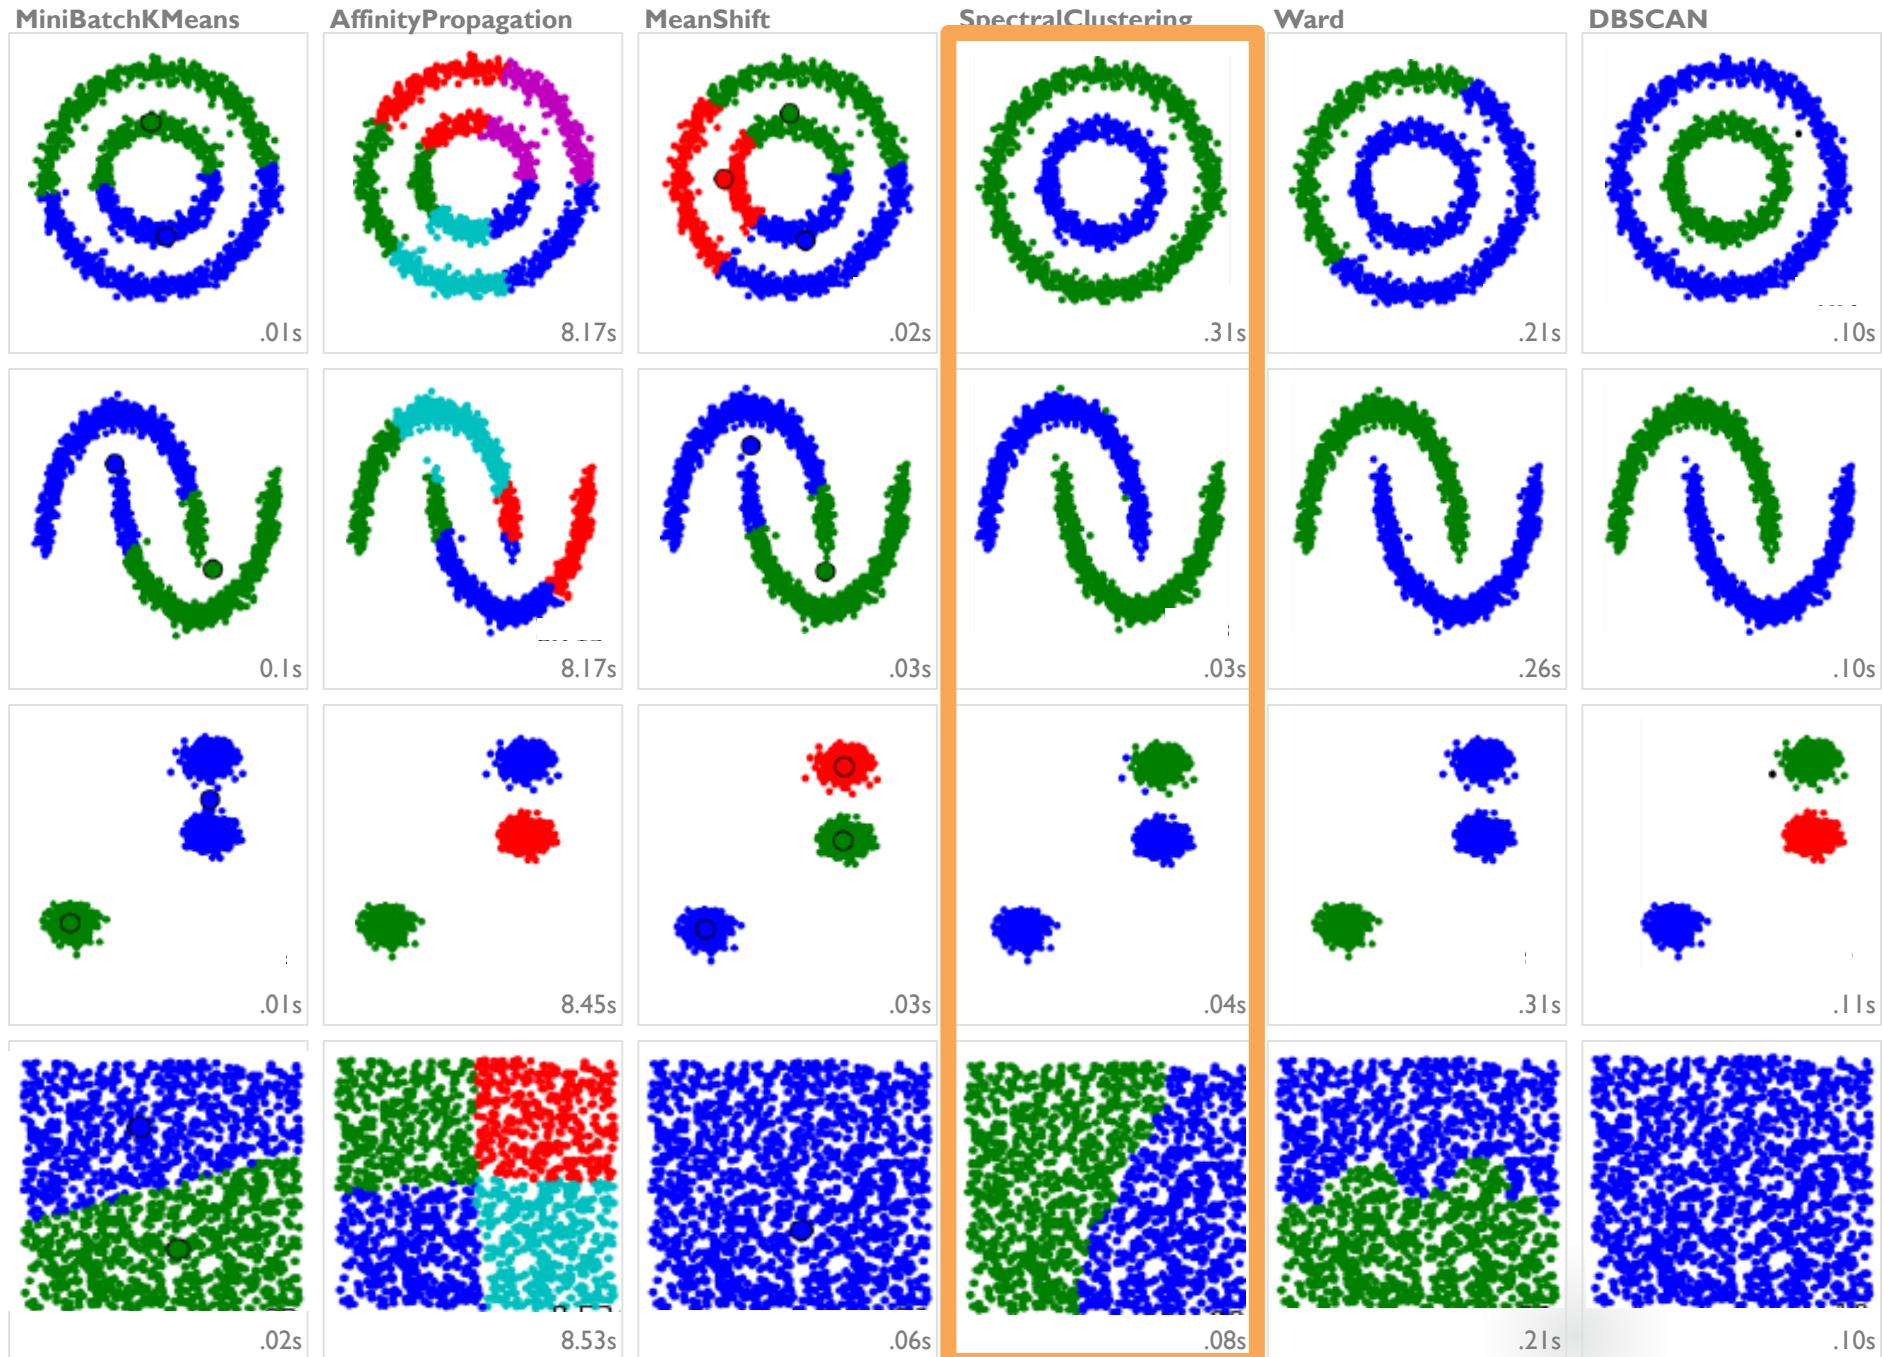


Method name	Parameters	Scalability	Usecase	Geometry (metric used)
K-means	Number of clusters	Very large $n_{samples}$ medium $n_{clusters}$ with <a href="#">MiniBatch code</a>	General purpose, even cluster size, flat geometry, not too many clusters	Distances between points
Affinity propagation	Damping, sample preferences	Not scalable with $n_{samples}$	Many clusters, uneven cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
Mean-shift	Bandwidth	Not scalable with $n_{samples}$	Many clusters, uneven cluster size, non-flat geometry	Distances between points
Spectral clustering	Number of clusters	Medium $n_{samples}$ , small $n_{clusters}$	Few clusters, even cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
Hierarchical clustering	Number of clusters	Large $n_{samples}$ and $n_{clusters}$	Many clusters, possibly connectivity constraints	Distances between points
DBSCAN	Neighborhood size	Very large $n_{samples}$ , medium $n_{clusters}$	Non-flat geometry, uneven cluster sizes	Distances between nearest points

## Spectral clustering:

- Best with **sparse distances**.
- Tends to find **even sized clusters**.
- Good with **only a few clusters**, rubbish with many.
- Can handle **weirdest connectivity shapes** (like concentric circles)
- **Euclidean distance only**.





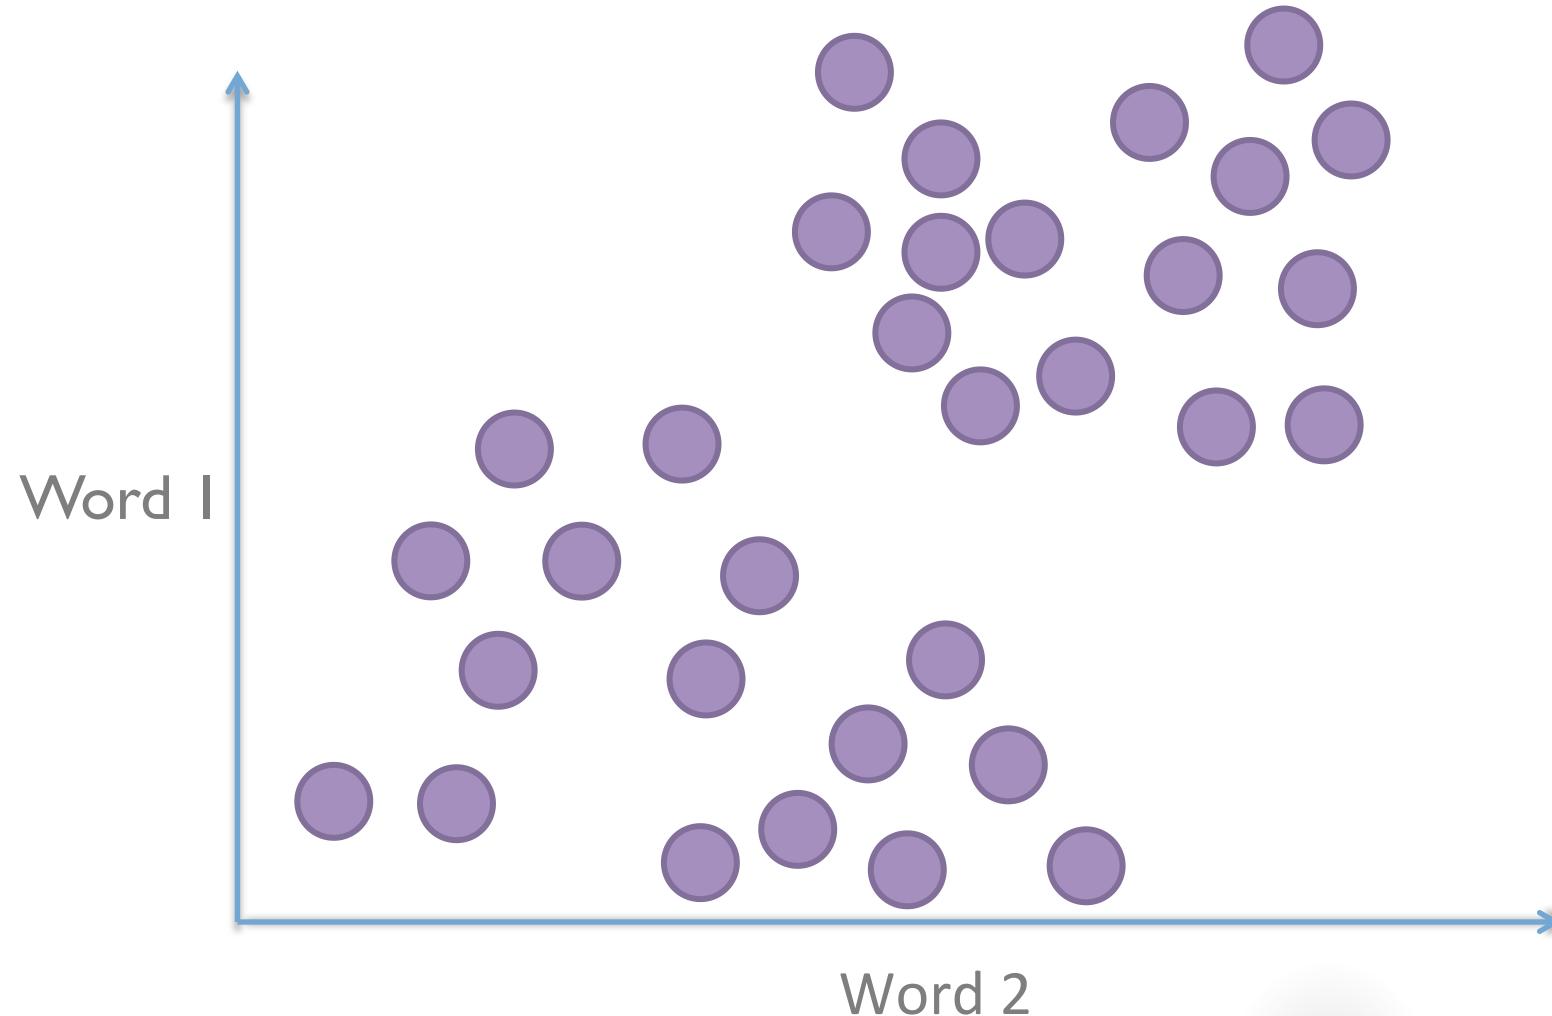


# What does clustering with cosine distance look like?

---

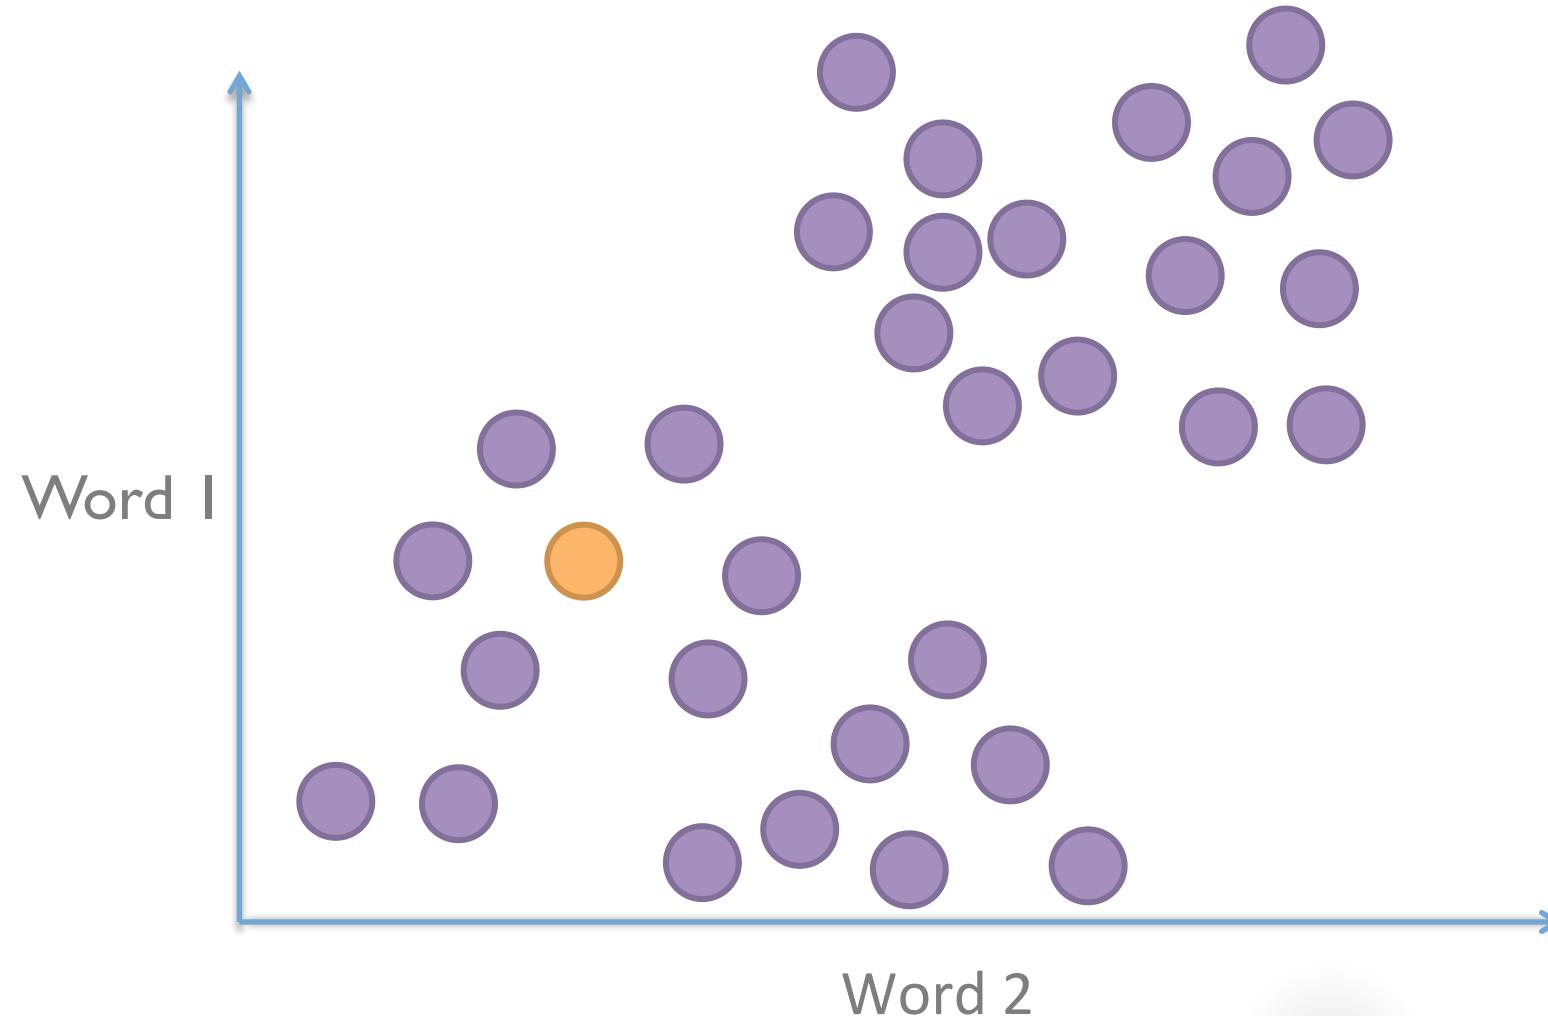
# DBSCAN

With cosine distance



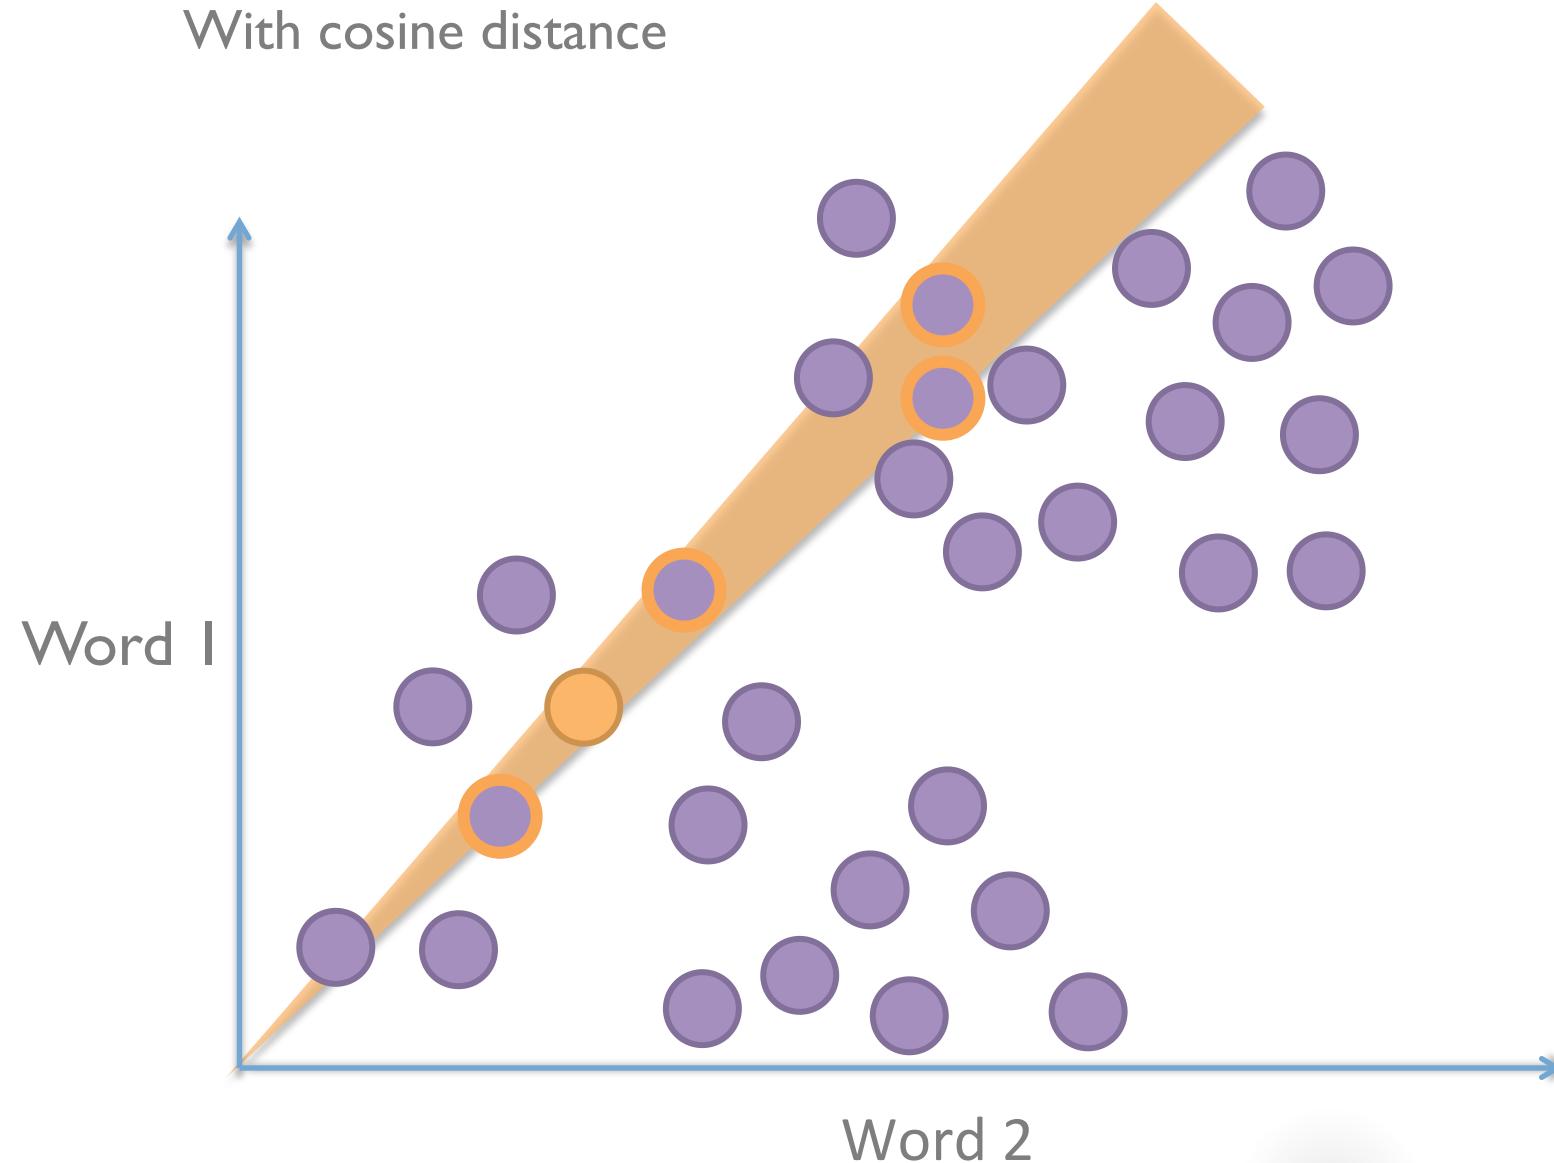
# DBSCAN

With cosine distance

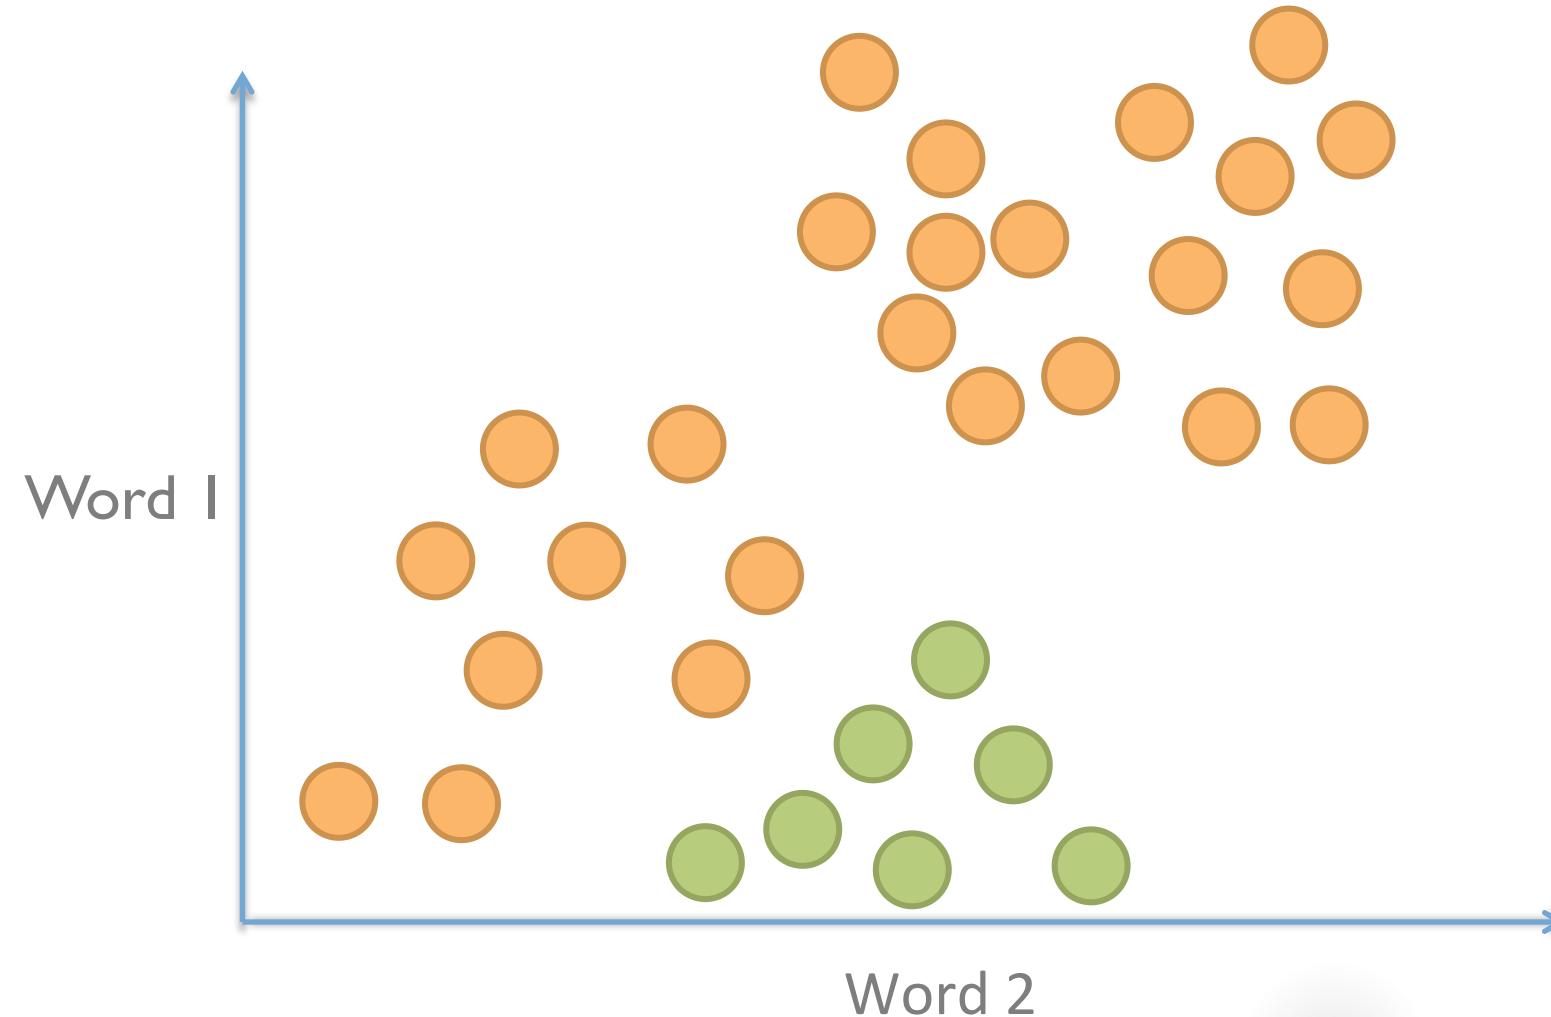


# DBSCAN

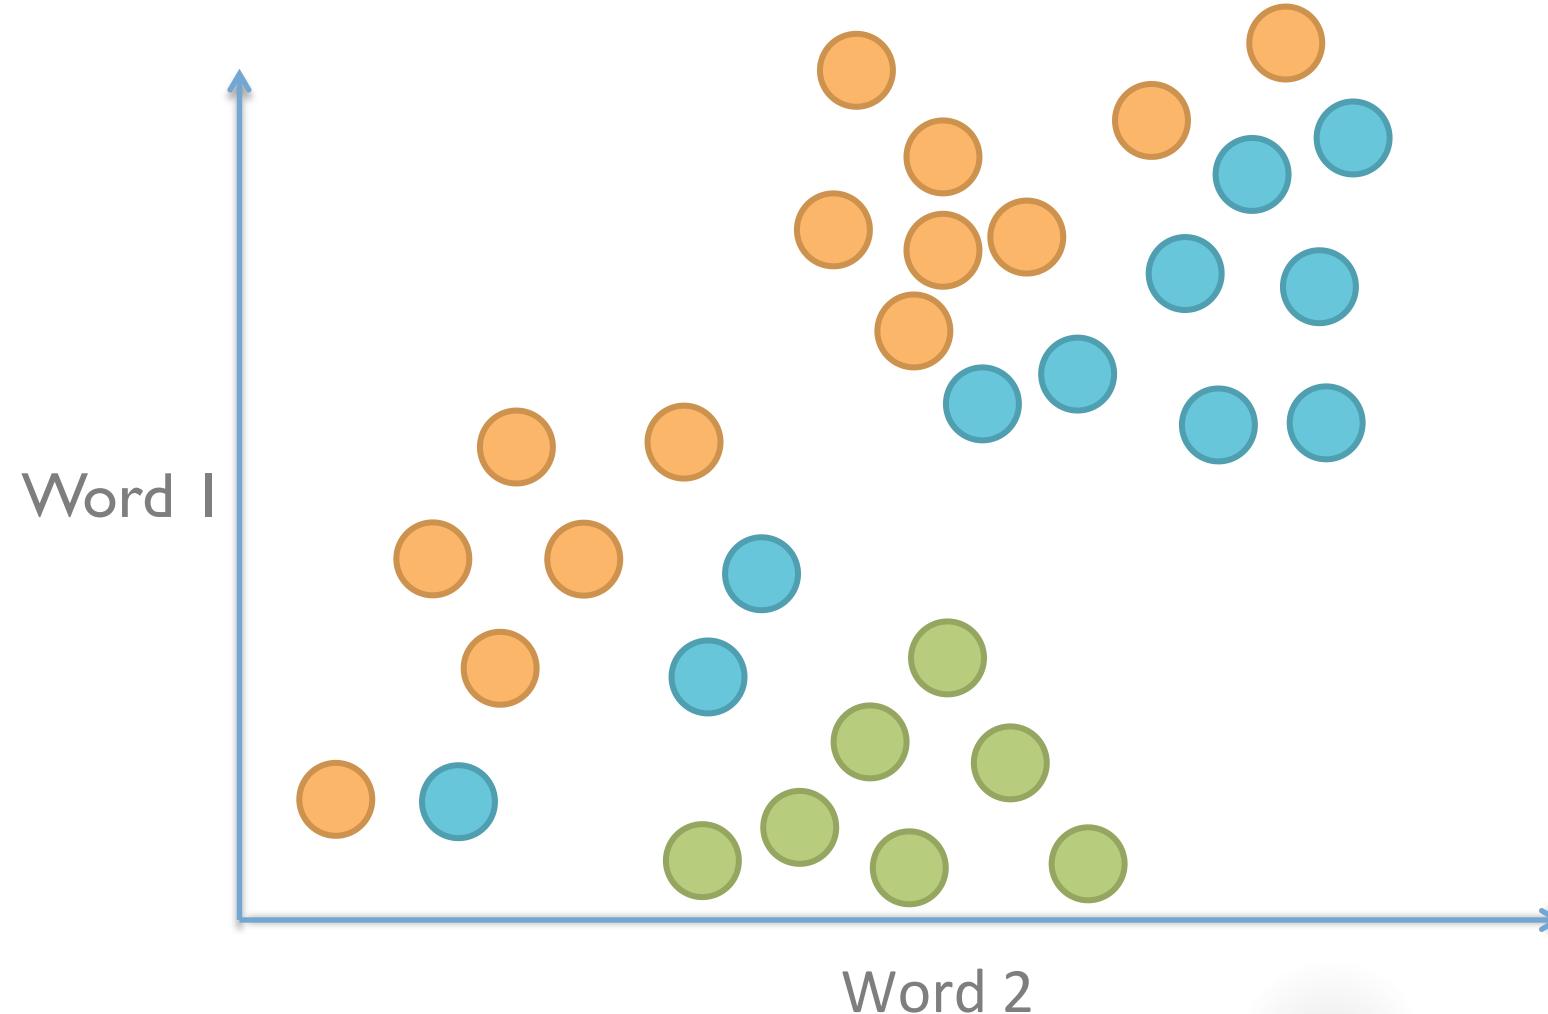
With cosine distance



# Potential clusters with cosine distance



# Potential clusters with cosine distance



# Potential clusters with cosine distance

